

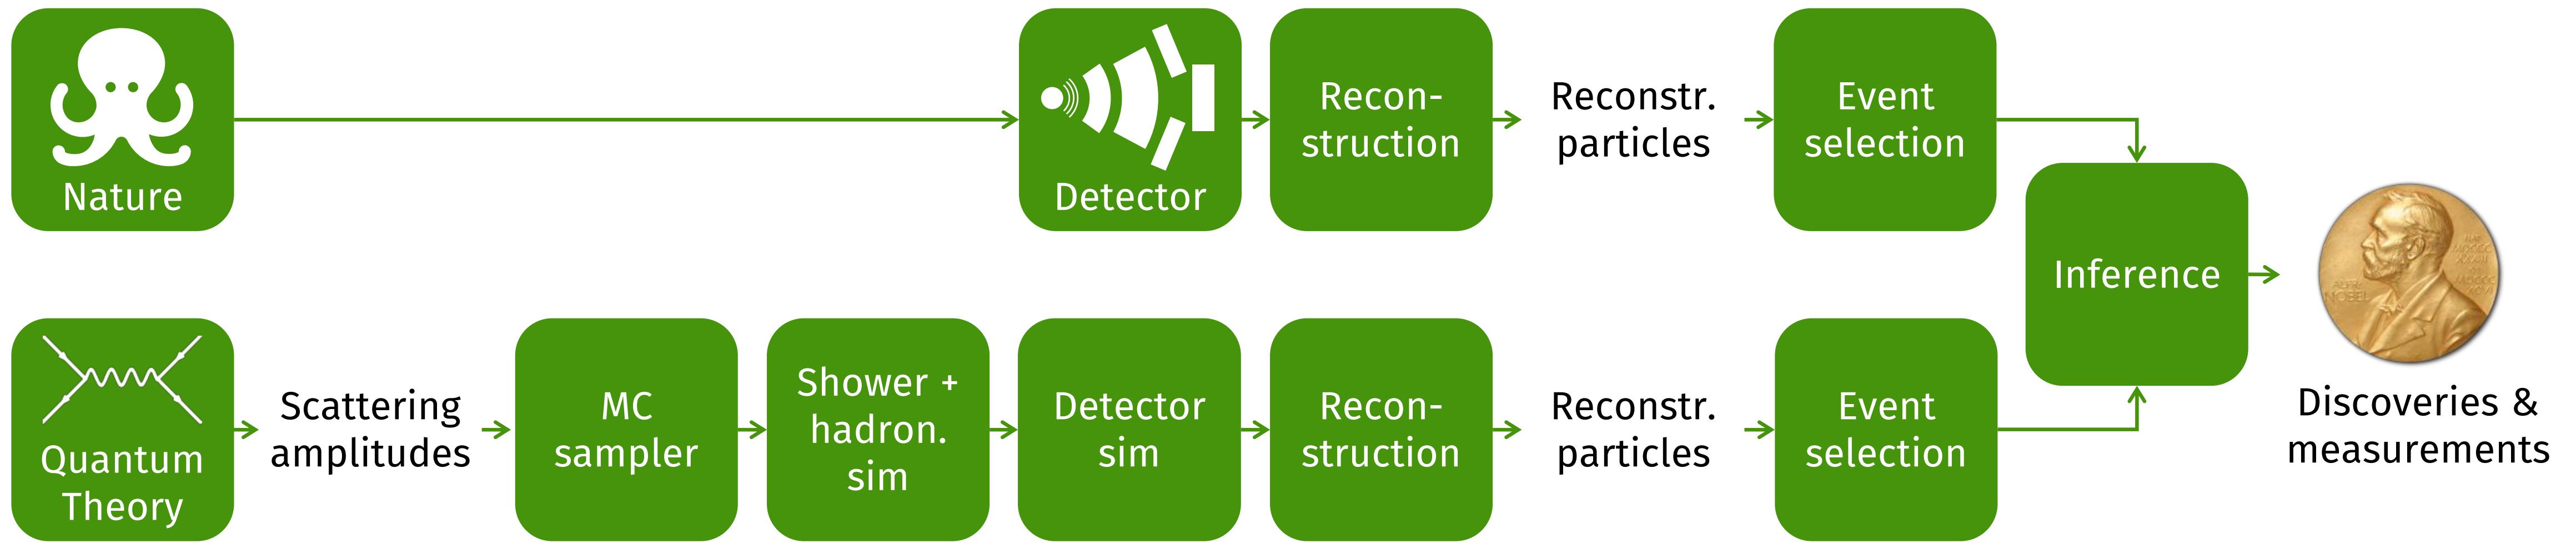
# L-GATr

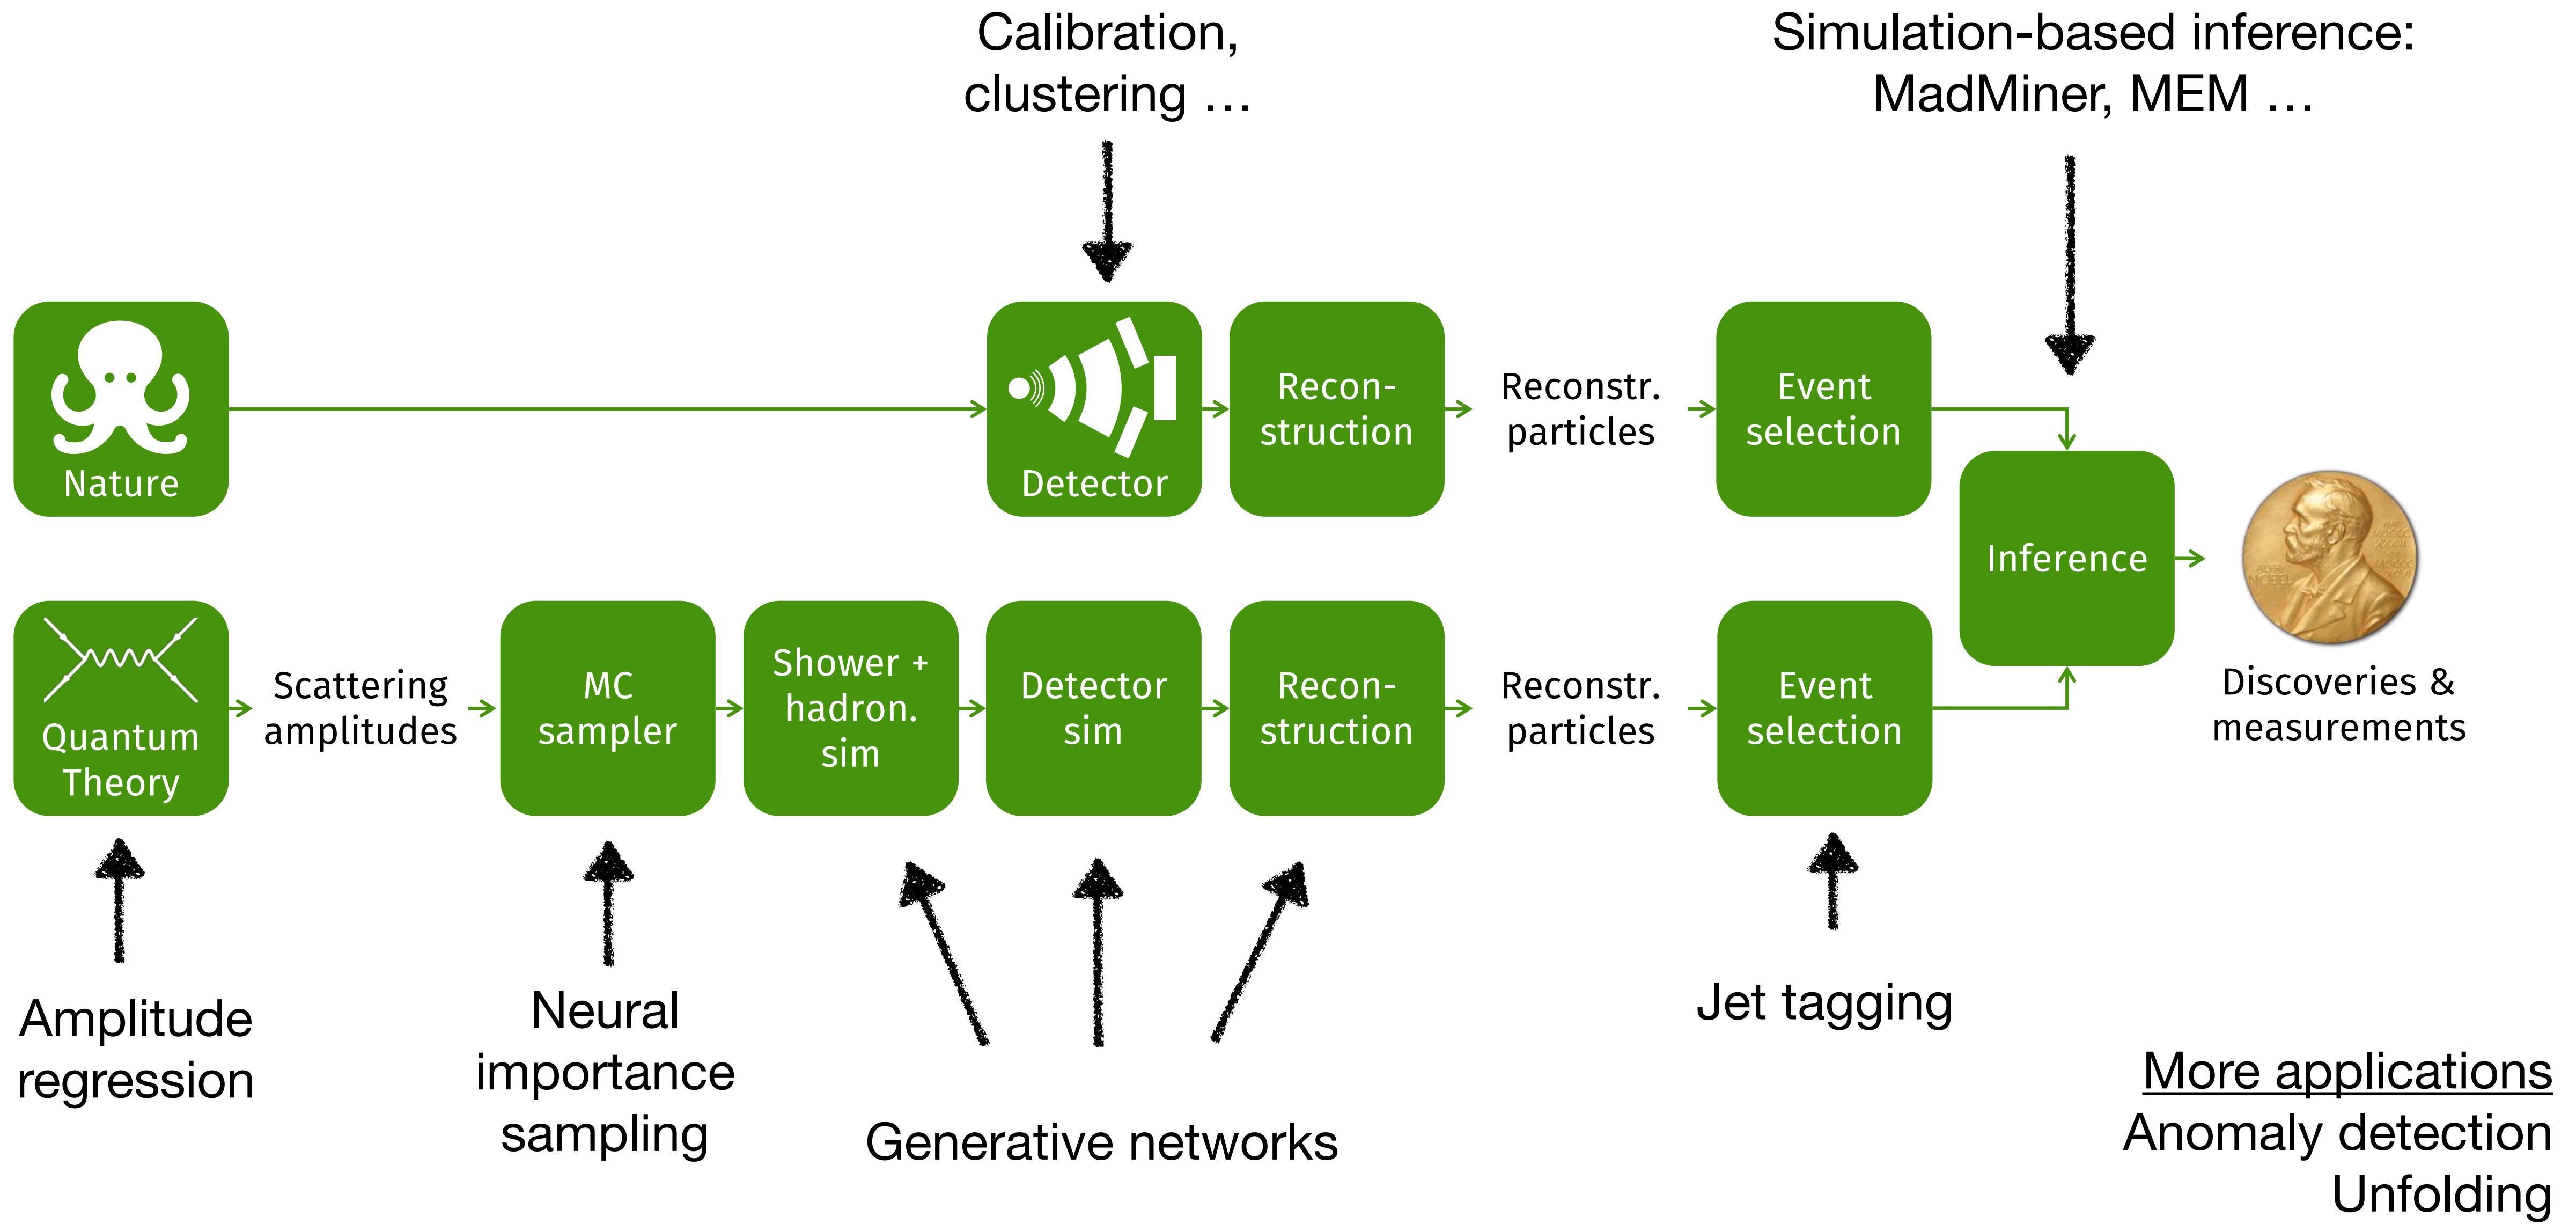
Jet Tagging with Lorentz-Equivariant  
Geometric Algebra Transformers

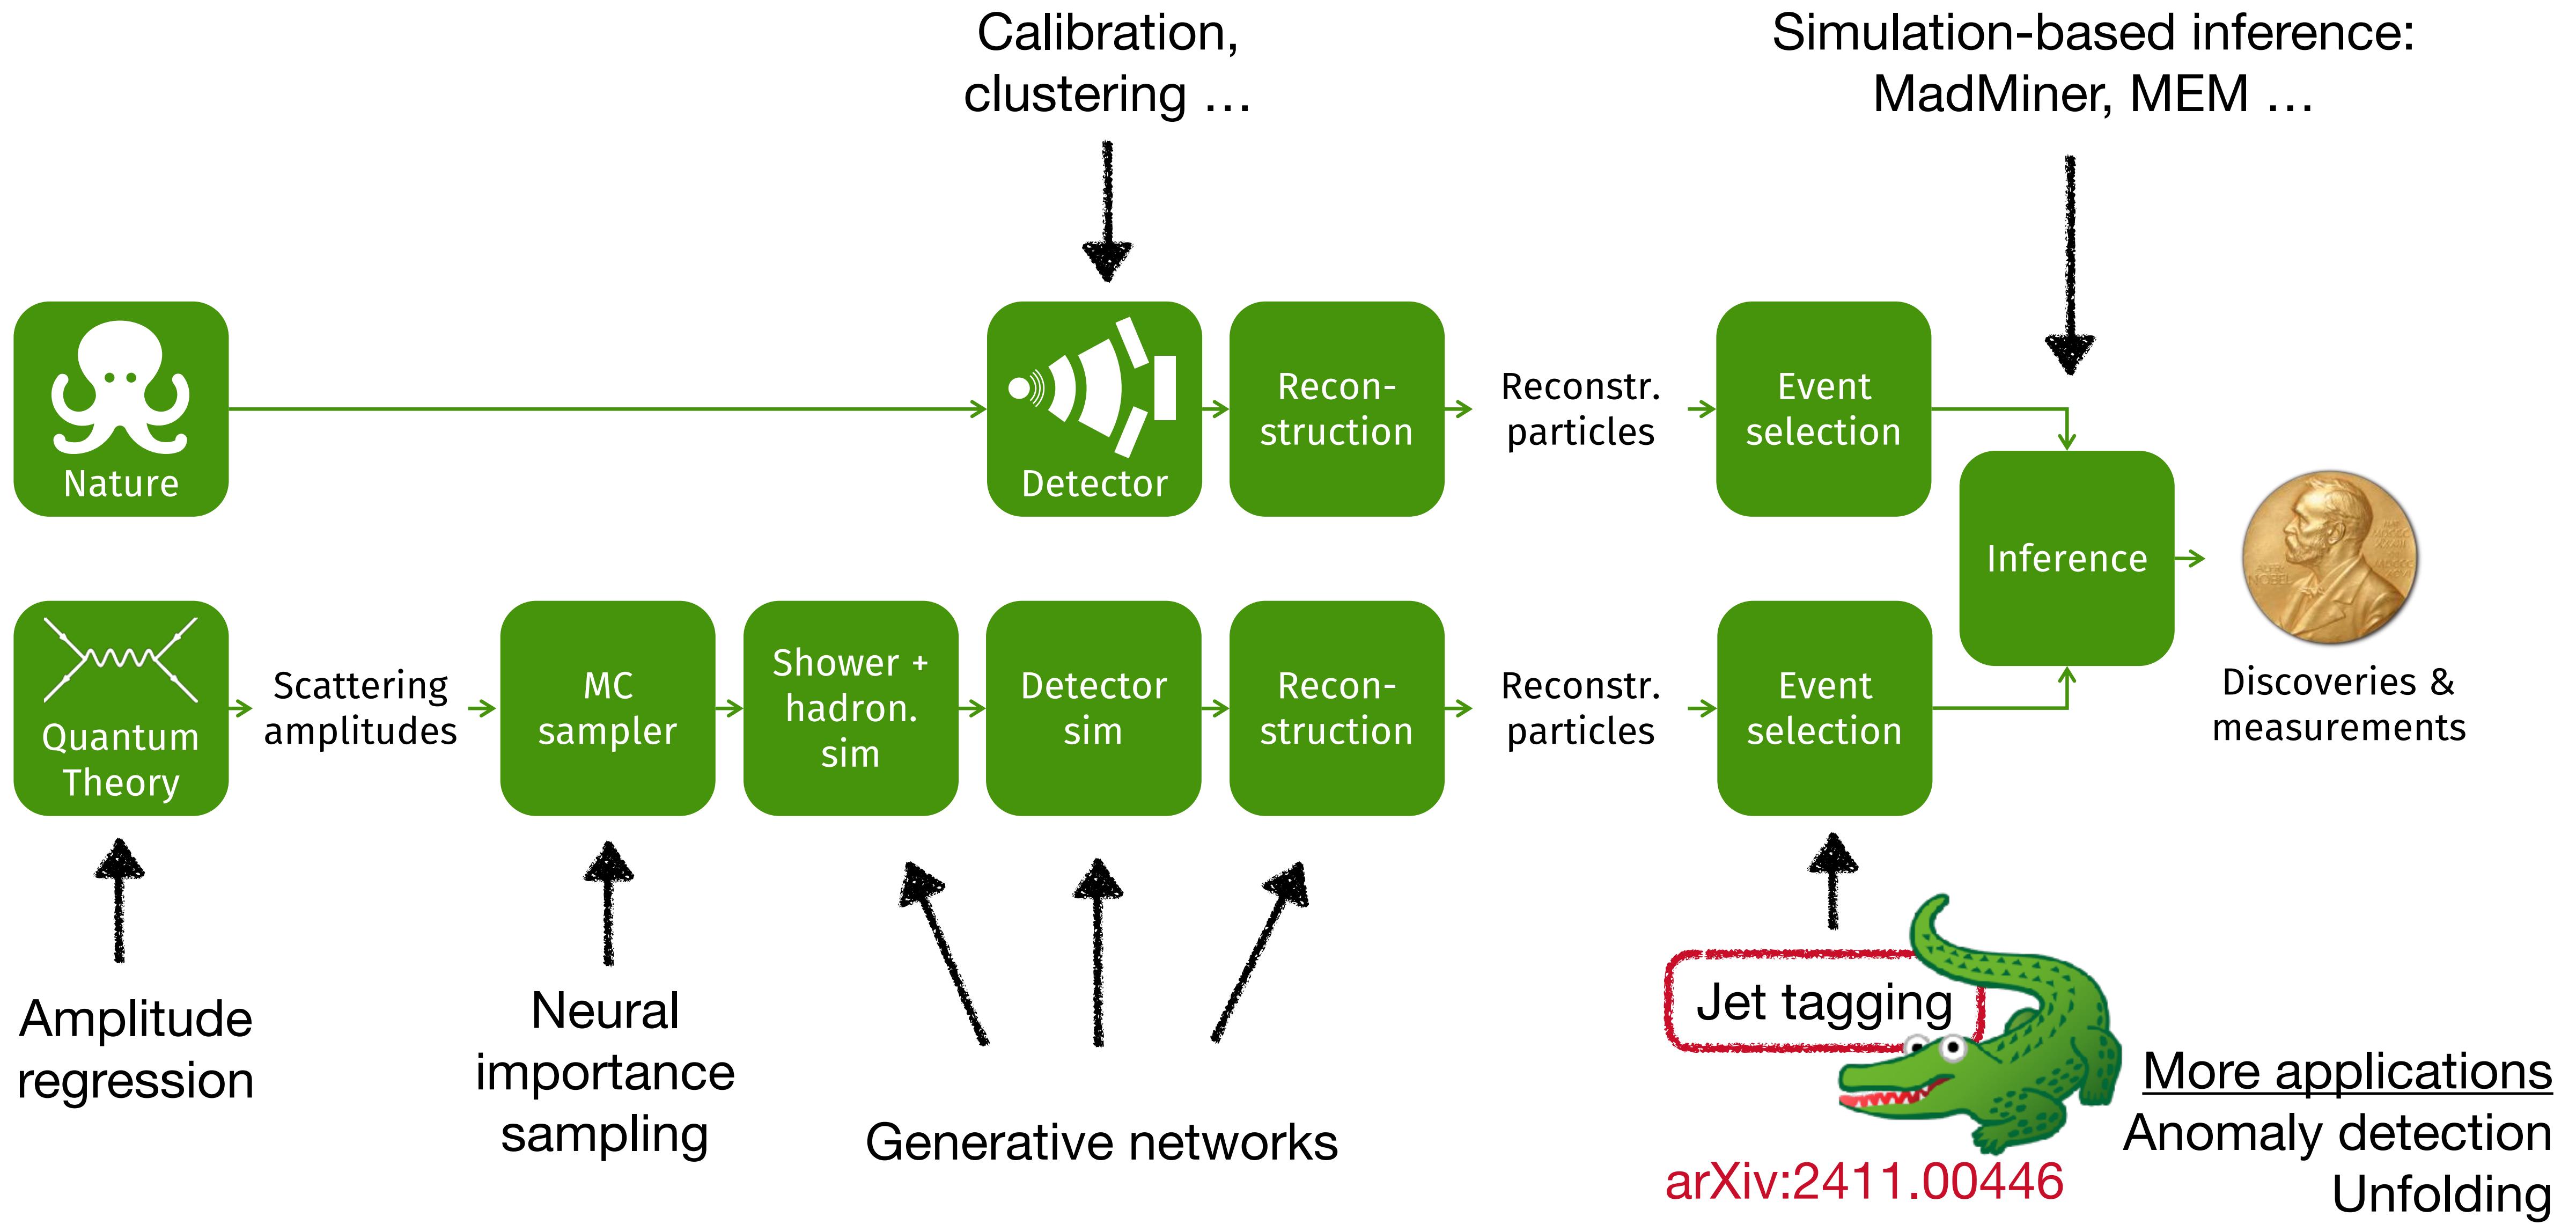
Johann Brehmer, Víctor Bresó,  
Pim de Haan, Tilman Plehn, Huilin Qu,  
Jonas Spinner, Jesse Thaler  
arXiv:2405.14806, arXiv:2411.00446

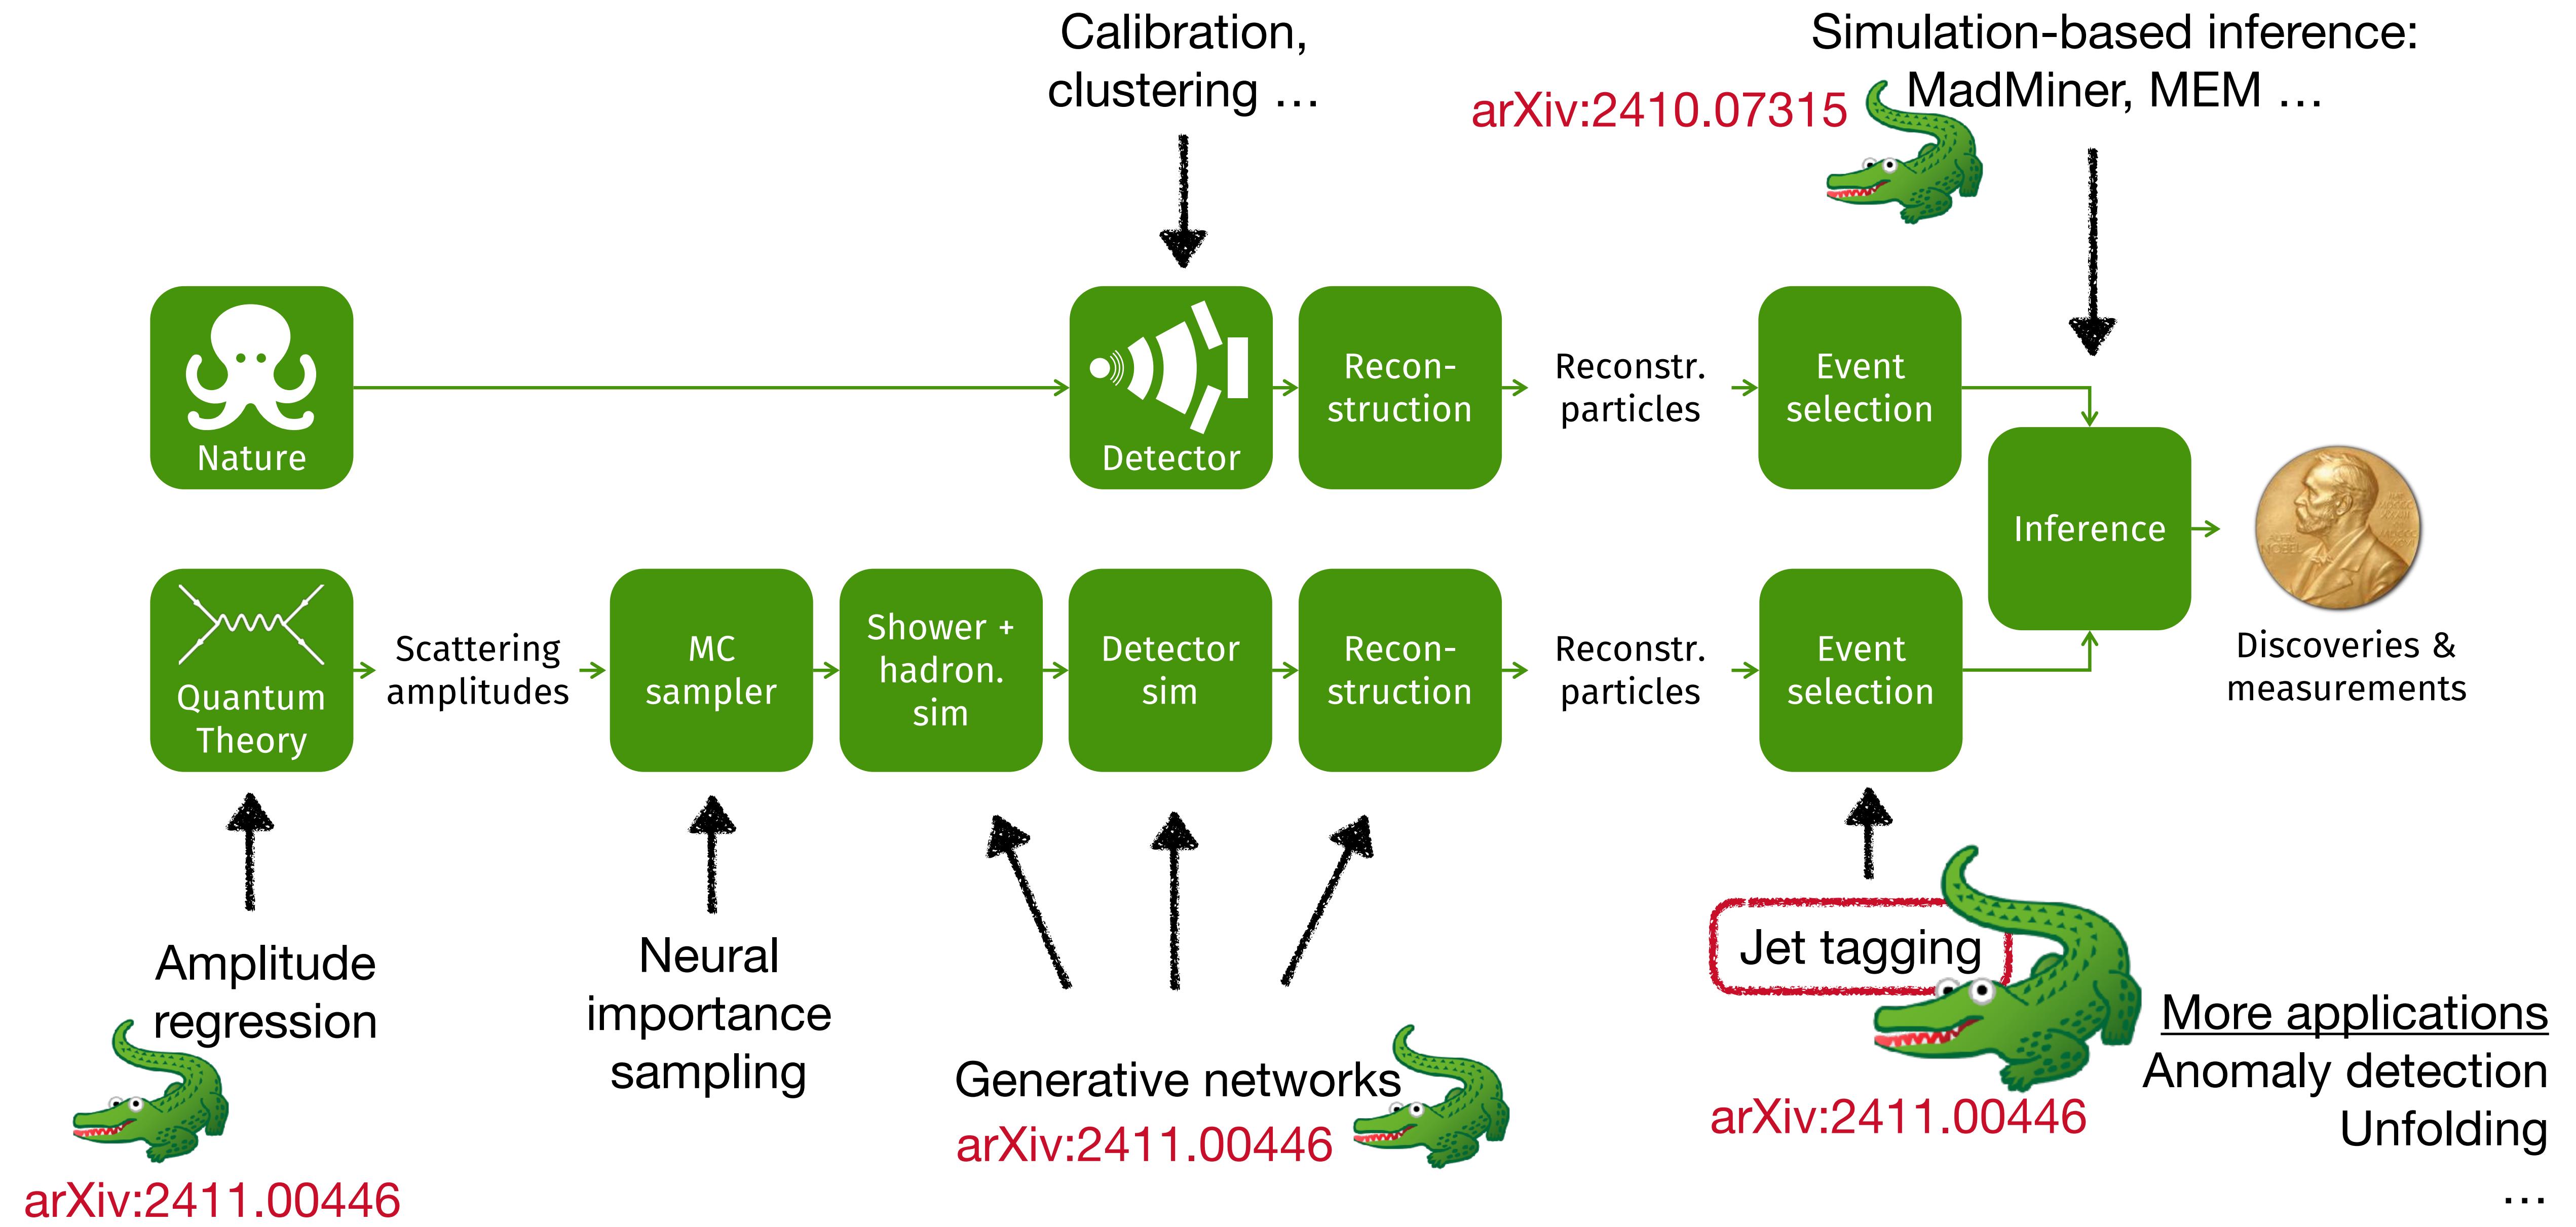


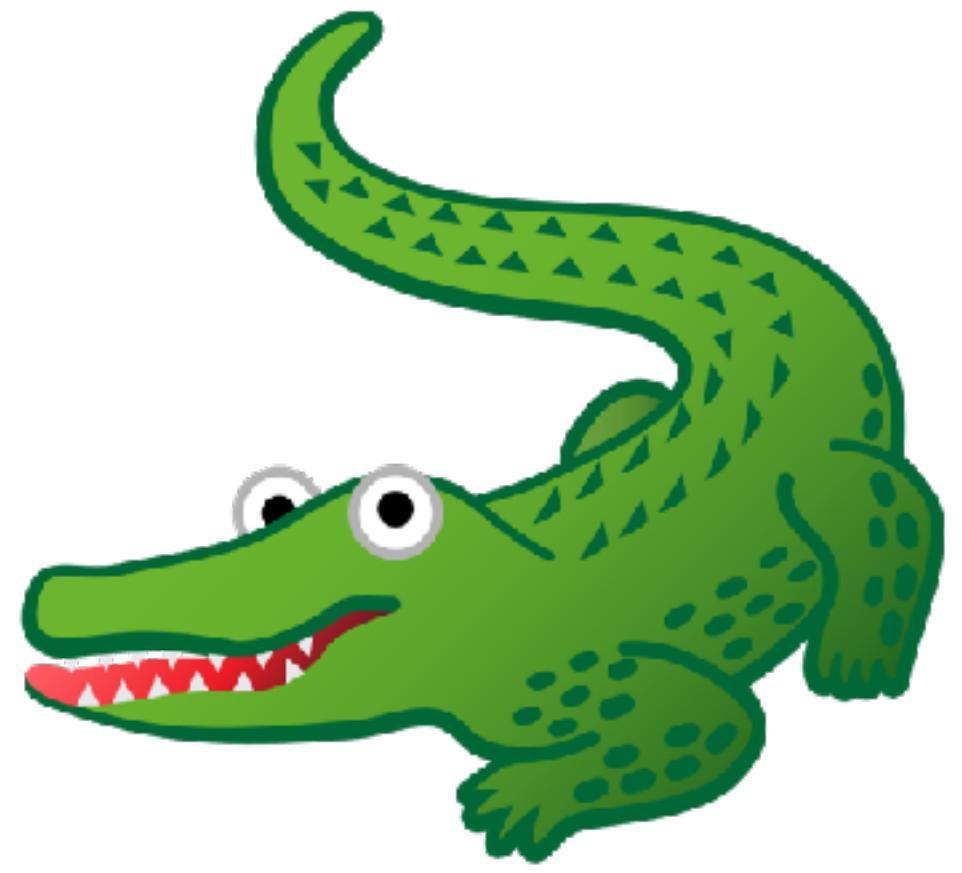
UNIVERSITÄT  
HEIDELBERG  
ZUKUNFT  
SEIT 1386



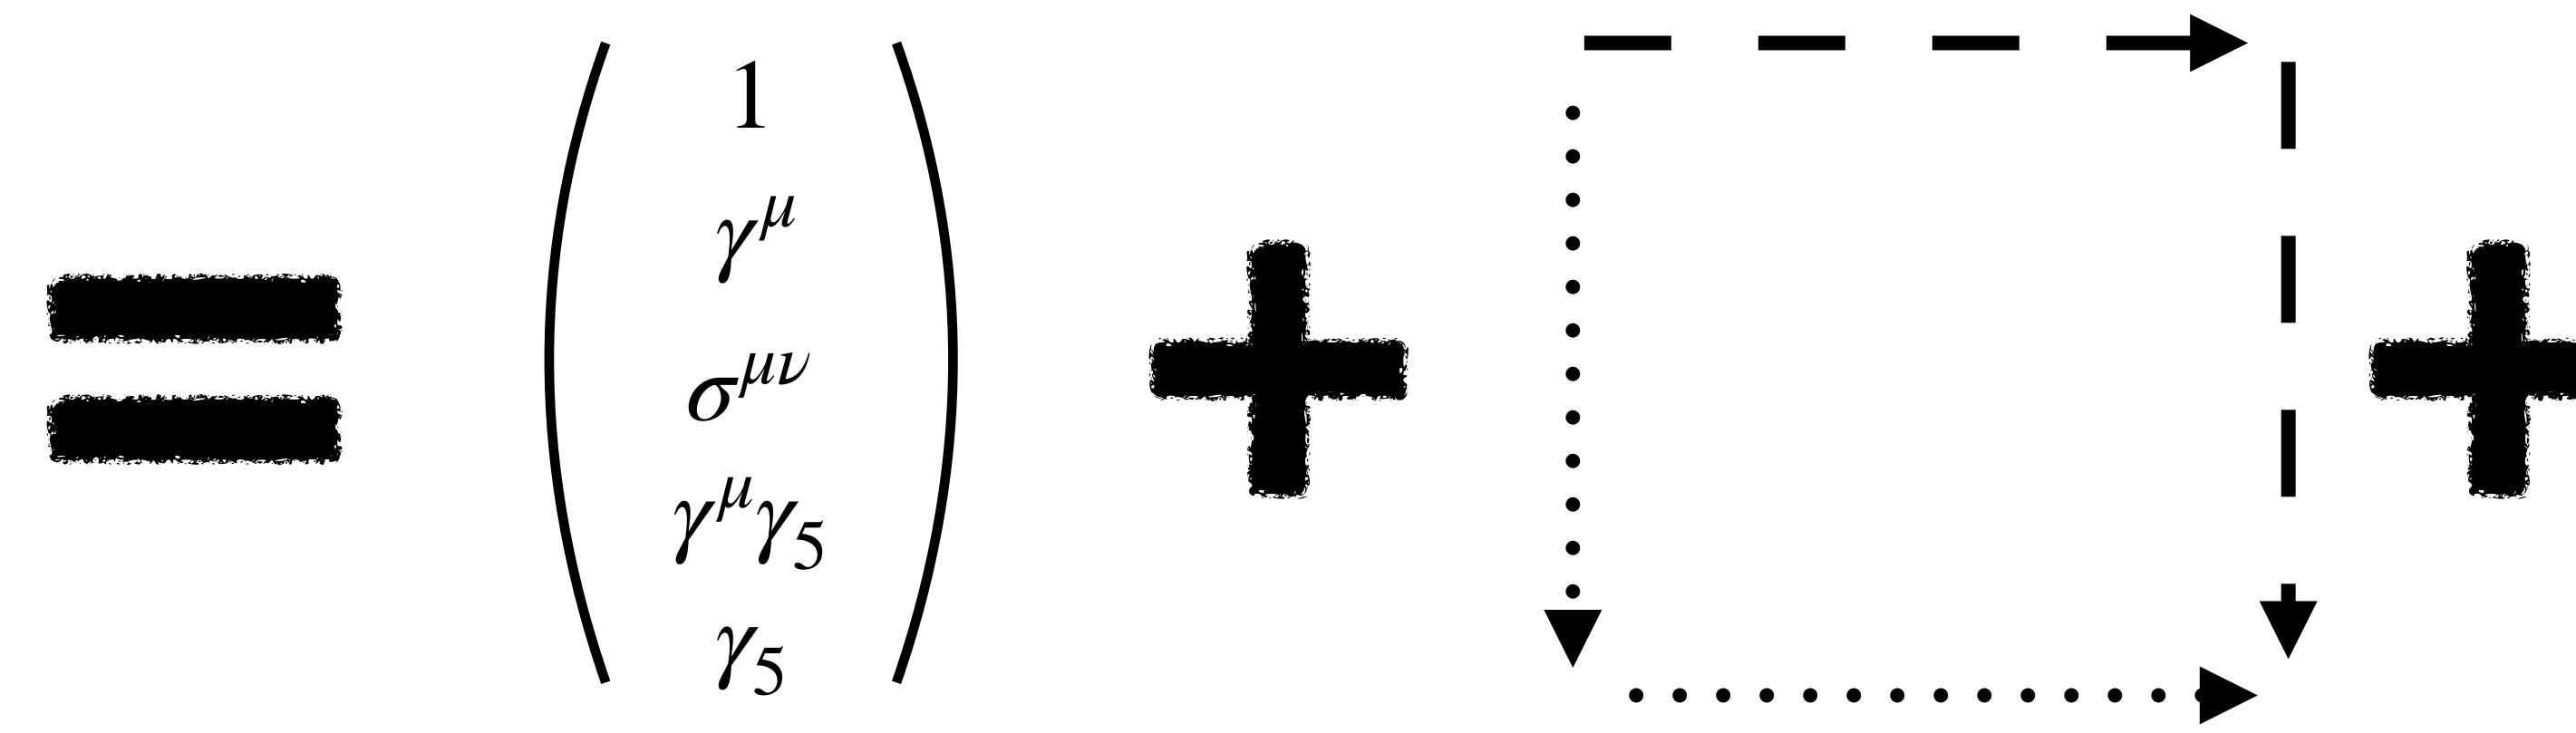








Lorentz-Equivariant  
Geometric **A**lgebra  
Transformer



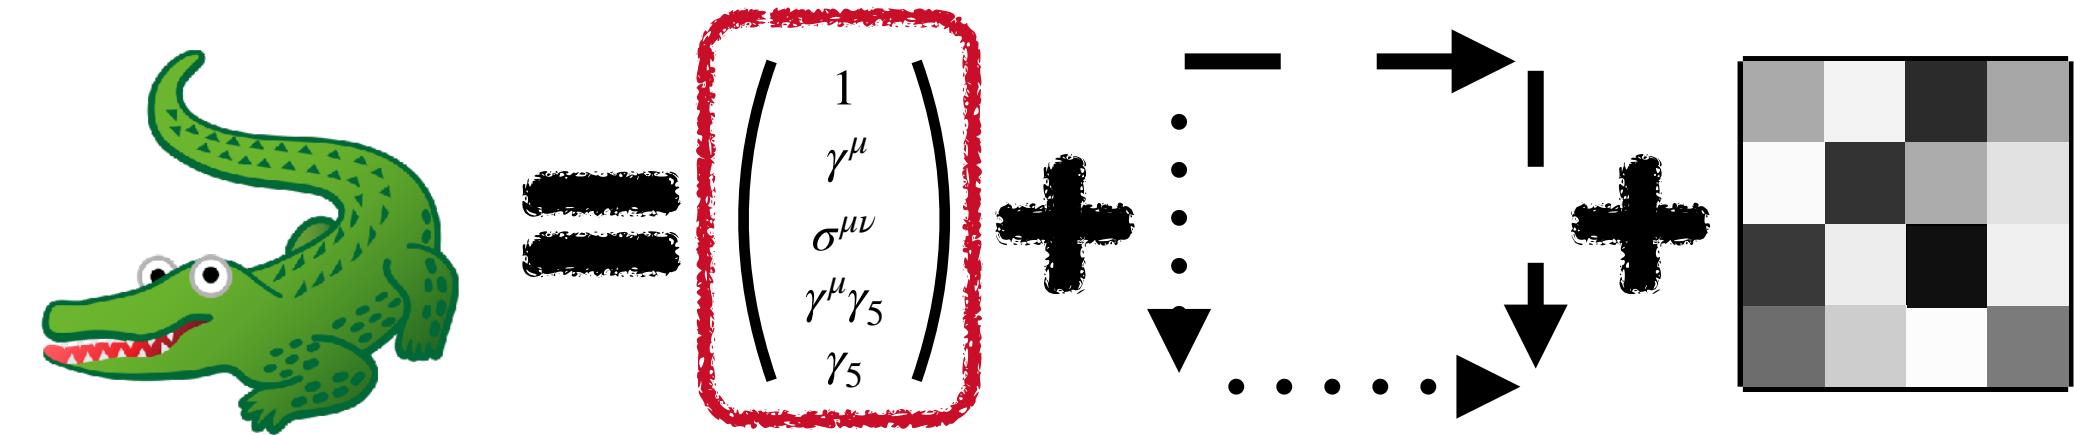
**Geometric algebra**  
representations

**Lorentz-Equivariant**  
layers

**Transformer**  
architecture

GATr was originally  
developed for E(3)  
arXiv:2305.18415

# L-GATr

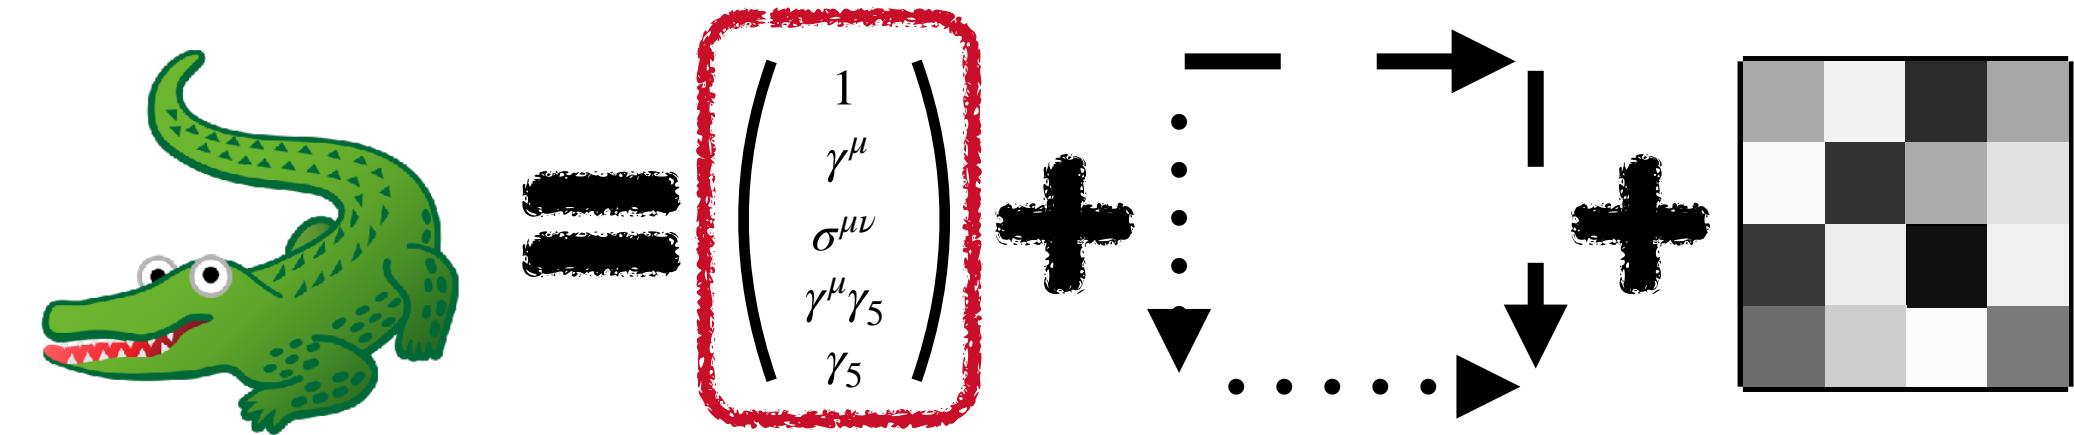


**Geometric algebra = Clifford algebra**

Geometric algebra = Vector space + geometric product  $xy = \frac{\{x, y\}}{2} + \frac{[x, y]}{2}$

- Symmetric part  $\{x, y\}$ : scalar/inner product
- Antisymmetric part  $[x, y]$ : outer product (yields higher-order objects)

# L-GATr



## Geometric algebra = Clifford algebra

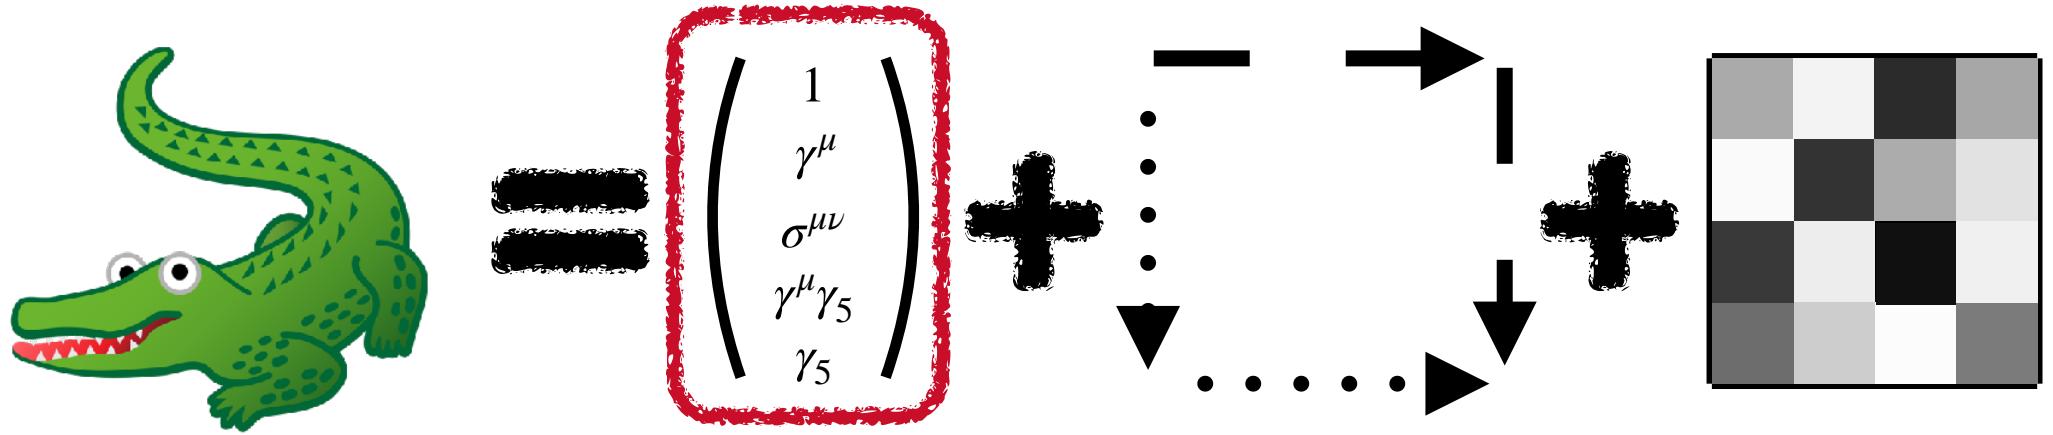
Geometric algebra = Vector space + geometric product  $xy = \frac{\{x, y\}}{2} + \frac{[x, y]}{2}$

- Symmetric part  $\{x, y\}$ : scalar/inner product
- Antisymmetric part  $[x, y]$ : outer product (yields higher-order objects)

Spacetime geometric algebra: Geometric algebra over vector space  $\mathbb{R}^4$  with Minkowski metric  $g = \text{diag}(1, -1, -1, -1)$

- Basis elements  $\gamma^\mu$  are orthonormal:  $\{\gamma^\mu, \gamma^\nu\} = 2g^{\mu\nu}$
- Dirac algebra is the same up to  $\mathbb{R} \rightarrow \mathbb{C}$

# L-GATr



## Building multivectors

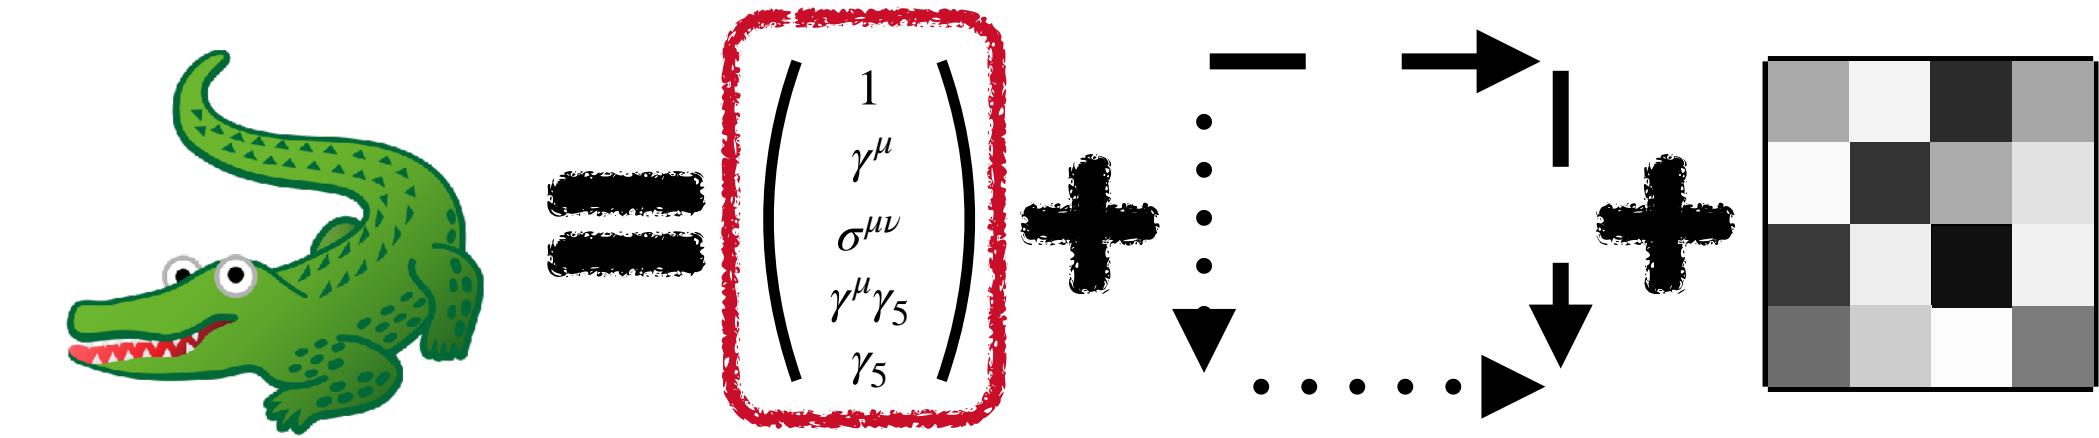
- Scalar and vectors  $1, \gamma^\mu$  (1+4 objects)
- Product of two vectors:  $\gamma^\mu \gamma^\nu = \frac{\{\gamma^\mu, \gamma^\nu\}}{2} + \frac{[\gamma^\mu, \gamma^\nu]}{2} = g^{\mu\nu} + \sigma^{\mu\nu}$  (6 new objects)
- Axial vector:  $\epsilon_{\mu\nu\rho\sigma} \gamma^\nu \gamma^\rho \gamma^\sigma \propto \gamma_\mu \gamma^5$  (4 new objects)
- Pseudoscalar:  $\gamma^5 = \gamma^0 \gamma^1 \gamma^2 \gamma^3 = \frac{1}{4!} \epsilon_{\mu\nu\rho\sigma} \gamma^\mu \gamma^\nu \gamma^\rho \gamma^\sigma$  (1 new object)

Multivector:  $x = x^S 1 + x_\mu^V \gamma^\mu + x_{\mu\nu}^B \sigma^{\mu\nu} + x_\mu^A \gamma^\mu \gamma^5 + x^P \gamma^5$  with  $(x^S, x_\mu^V, x_{\mu\nu}^B, x_\mu^A, x^P) \in \mathbb{R}^{16}$

Particle:  $x^V = (E, p_x, p_y, p_z), \quad x^S = \text{PID}$

# L-GATr

## Geometric algebra representations



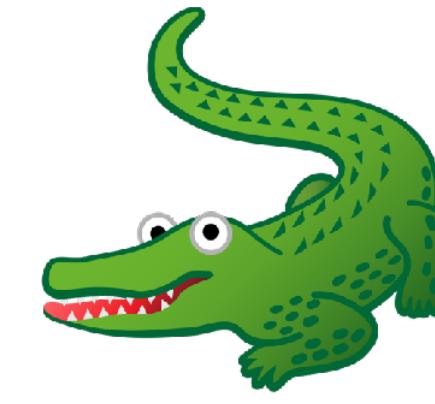
$$x^S \in \mathbb{R} \rightarrow \begin{pmatrix} x^S \\ x_\mu^V \\ x_{\mu\nu}^B \\ x_\mu^A \\ x_\mu^P \end{pmatrix} \in \mathbb{R}^{16}$$

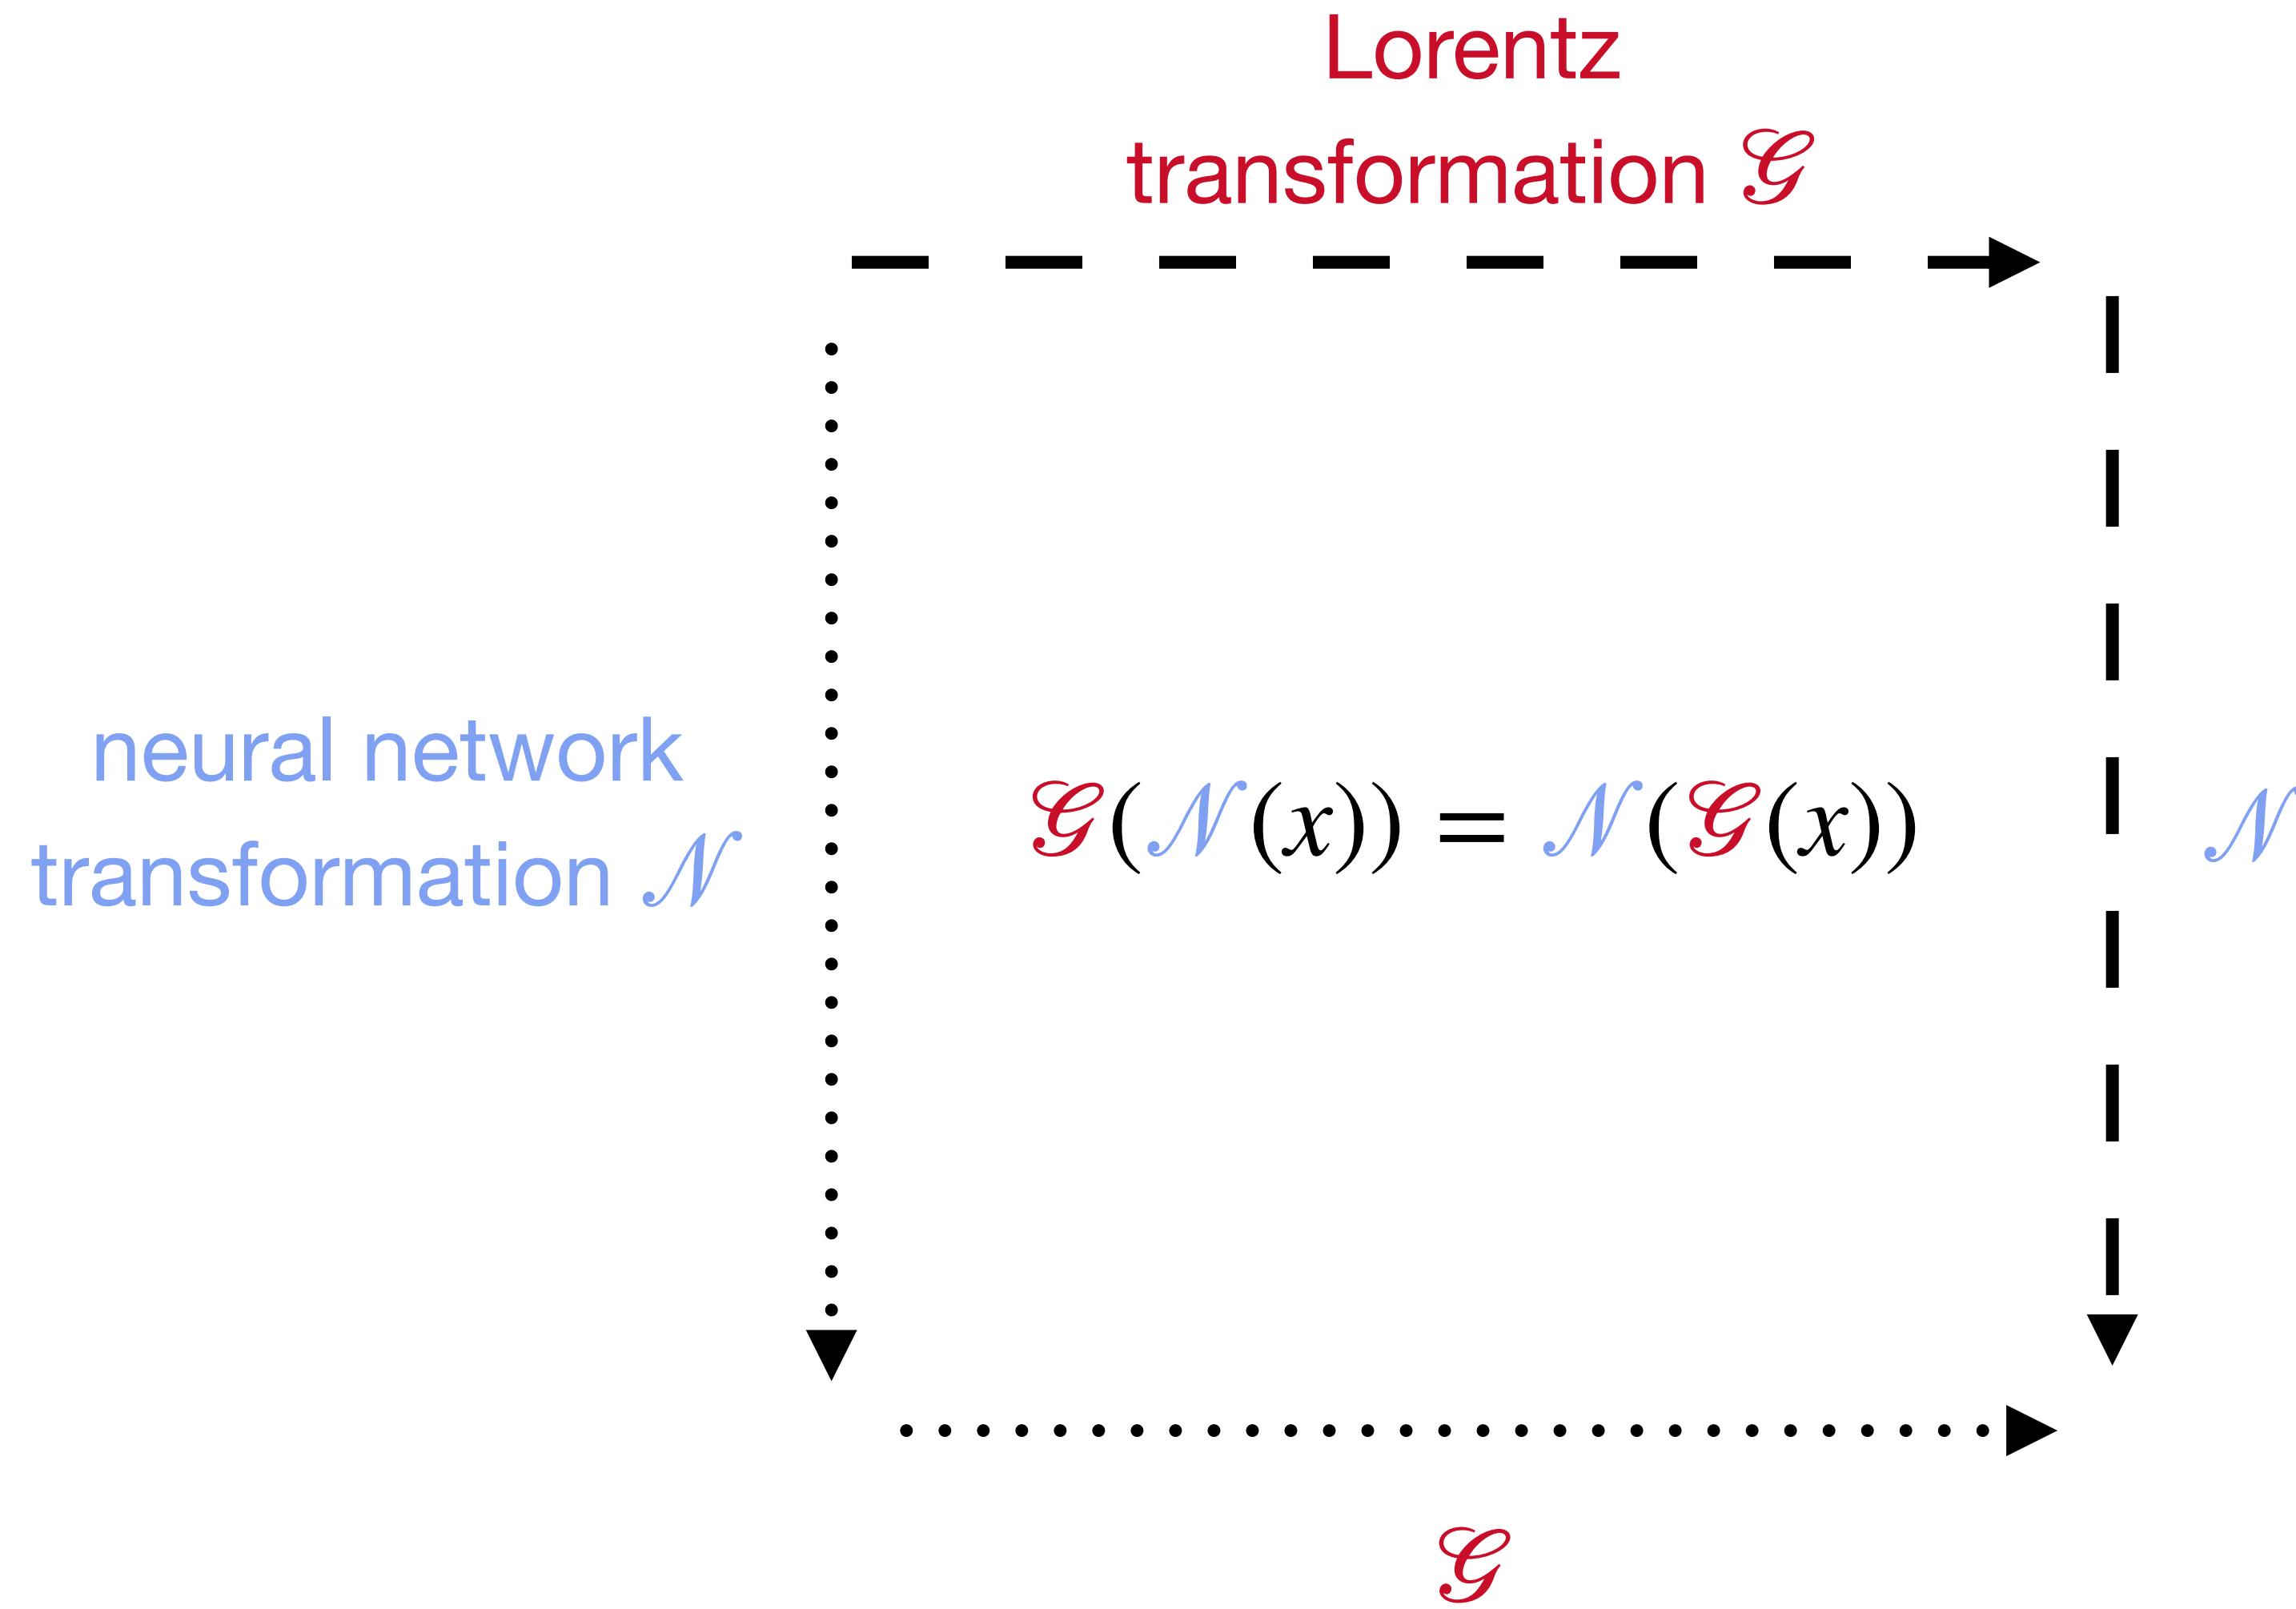
scalar channels

multivector channels

# L-GATr

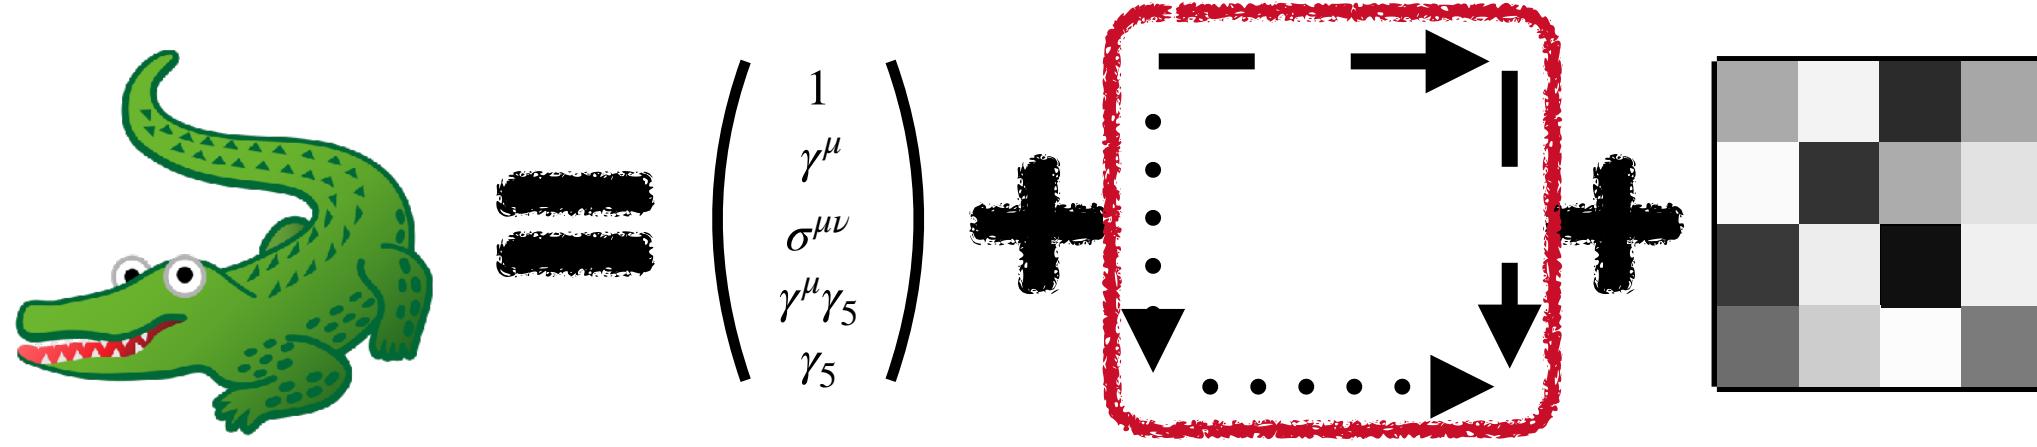
## Lorentz Equivariance


$$= \begin{pmatrix} 1 \\ \gamma^\mu \\ \sigma^{\mu\nu} \\ \gamma^\mu \gamma_5 \\ \gamma_5 \end{pmatrix} + \boxed{\dots} + \boxed{+} \quad \text{Equivariance = Covariance}$$



# L-GATr

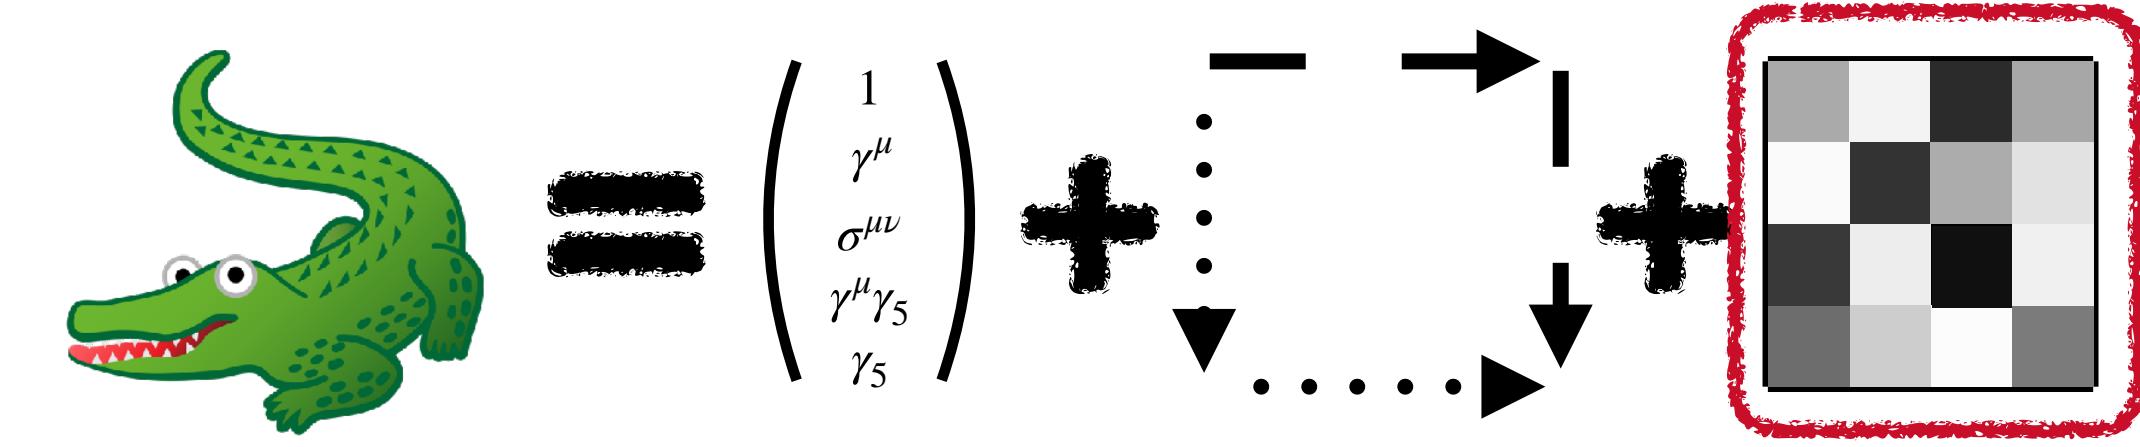
## GATr-ing all transformer layers



Layer type	Transformer	L-GATr
Linear( $x$ )	$vx + w$	$\sum_{k=0}^4 v_k \langle x \rangle_k \left( + \sum_{k=0}^4 w_k \gamma^5 \langle x \rangle_k \right)$
Attention( $q, k, v$ ) <sub>ic</sub>	$\sum_{j=1}^{n_t} \text{Softmax}_j \left( \sum_{c'=1}^{n_c} \frac{q_{ic'} k_{jc'}}{\sqrt{n_c}} \right) v_{jc}$	$\sum_{j=1}^{n_t} \text{Softmax}_j \left( \sum_{c'=1}^{n_c} \frac{\langle q_{ic'}, k_{jc'} \rangle}{\sqrt{16n_c}} \right) v_{jc}$
LayerNorm( $x$ )	$x \left[ \frac{1}{n_c} \sum_{c=1}^{n_c} x_c^2 + \epsilon \right]^{-1/2}$	$x \left[ \frac{1}{n_c} \sum_{c=1}^{n_c} \sum_{k=0}^4 \left  \langle \langle x_c \rangle_k, \langle x_c \rangle_k \rangle \right  + \epsilon \right]^{-1/2}$
Activation( $x$ )	GELU( $x$ )	GELU( $\langle x \rangle_0$ ) $x$
GP( $x, y$ )	—	$xy$

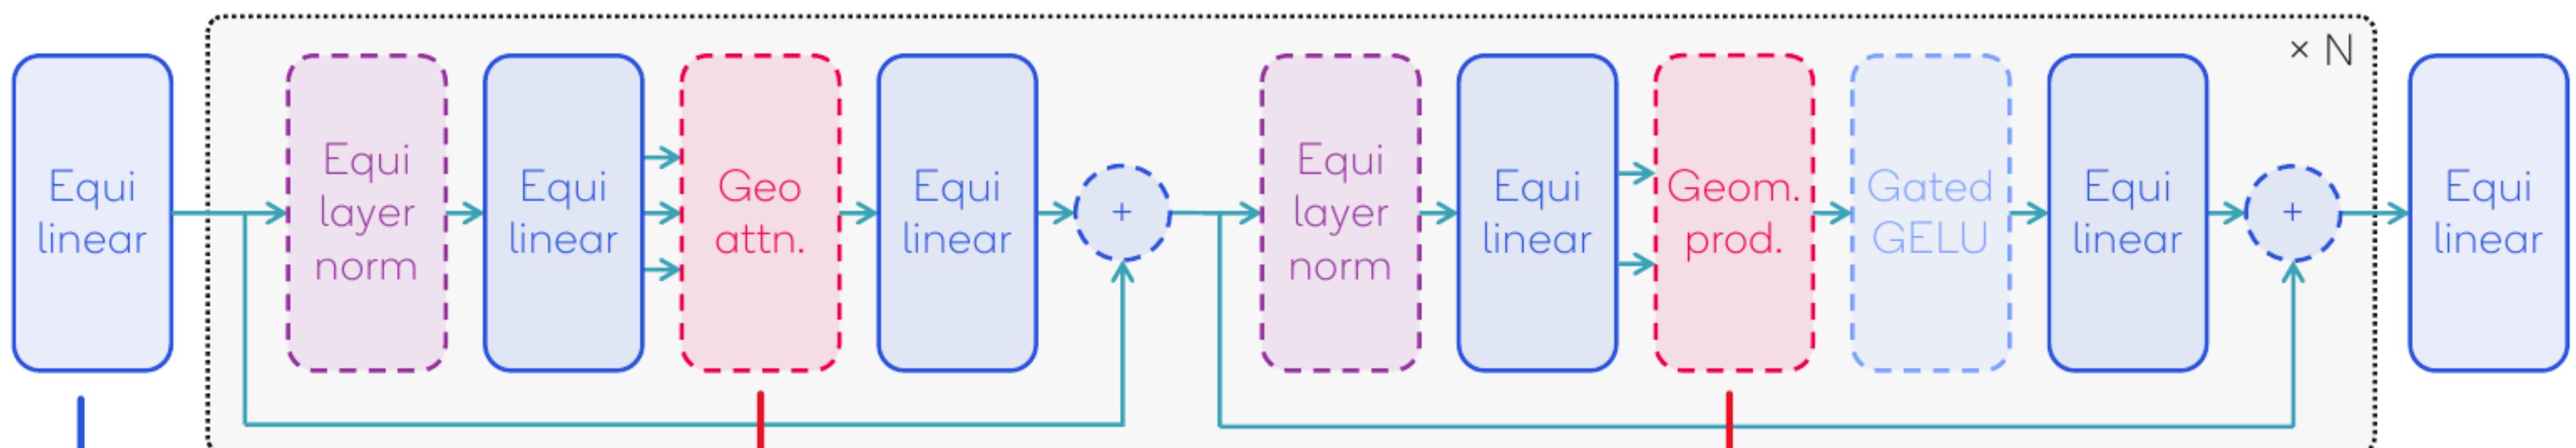
# L-GATr

## Full architecture



### Input and output data

can have one or multiple token dimensions



### Linear layers

between GA representations with equivariance constraint

**Geometric attention**  
generalizes scaled dot-product attention

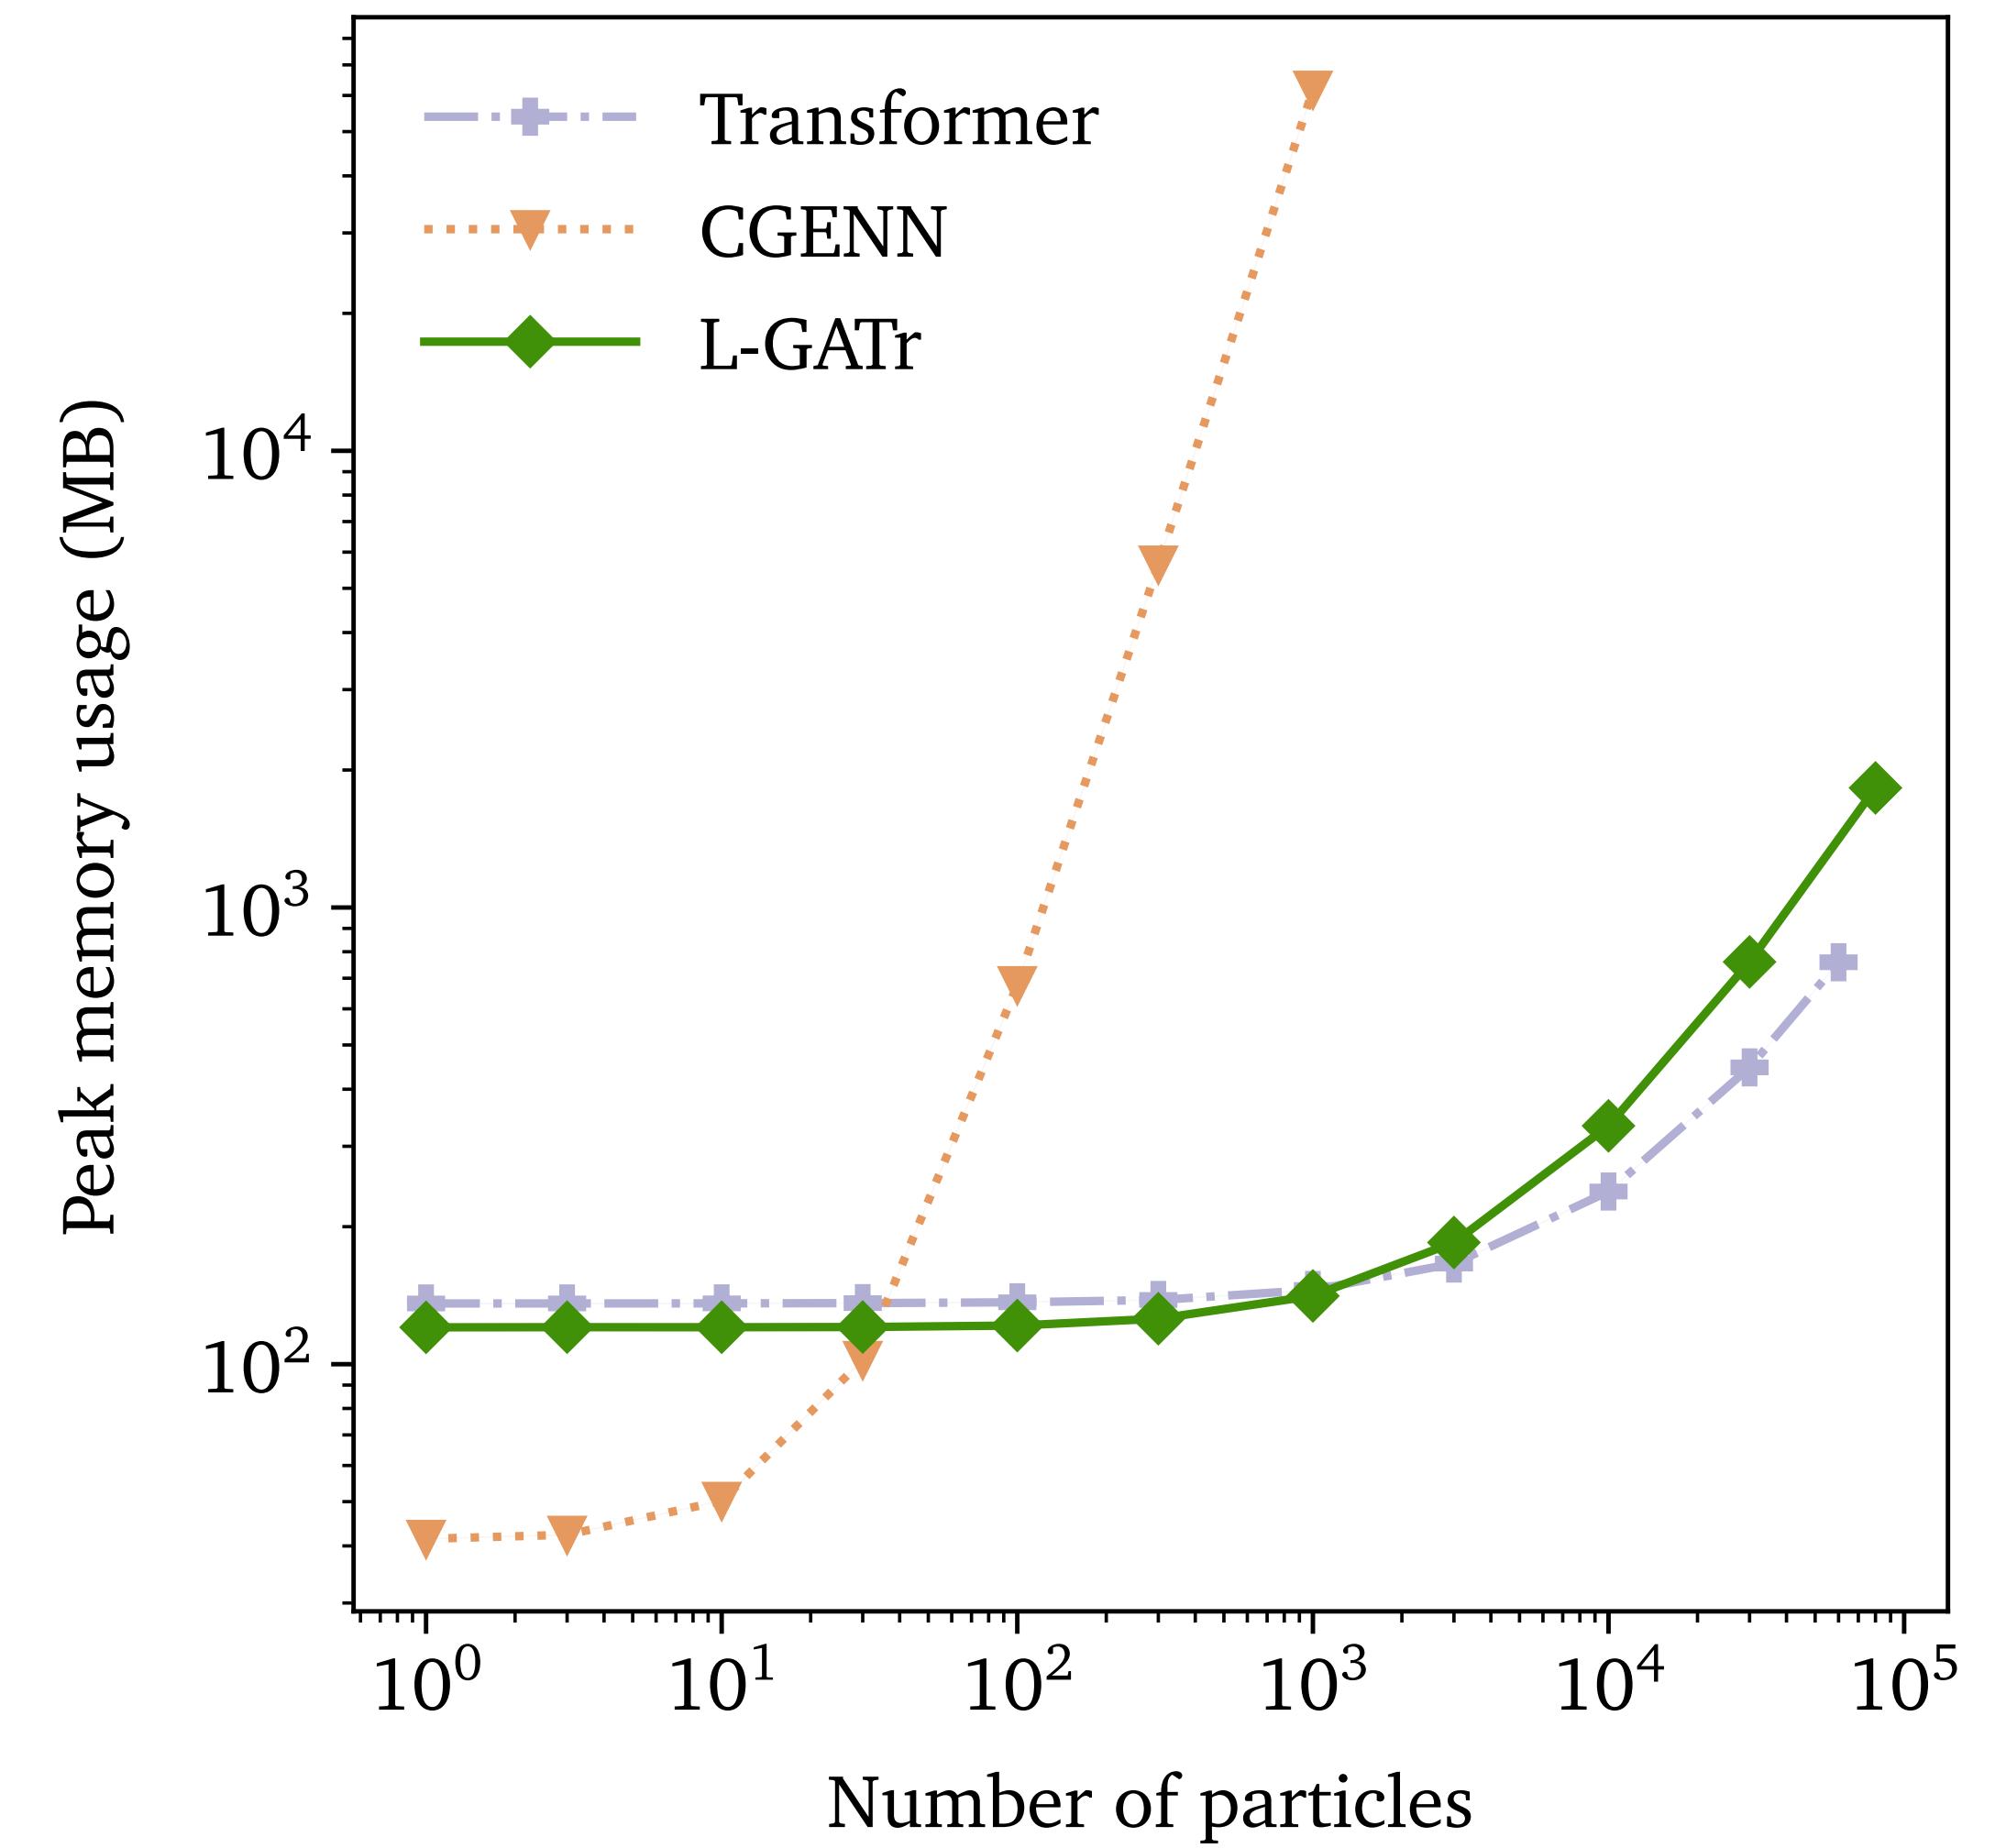
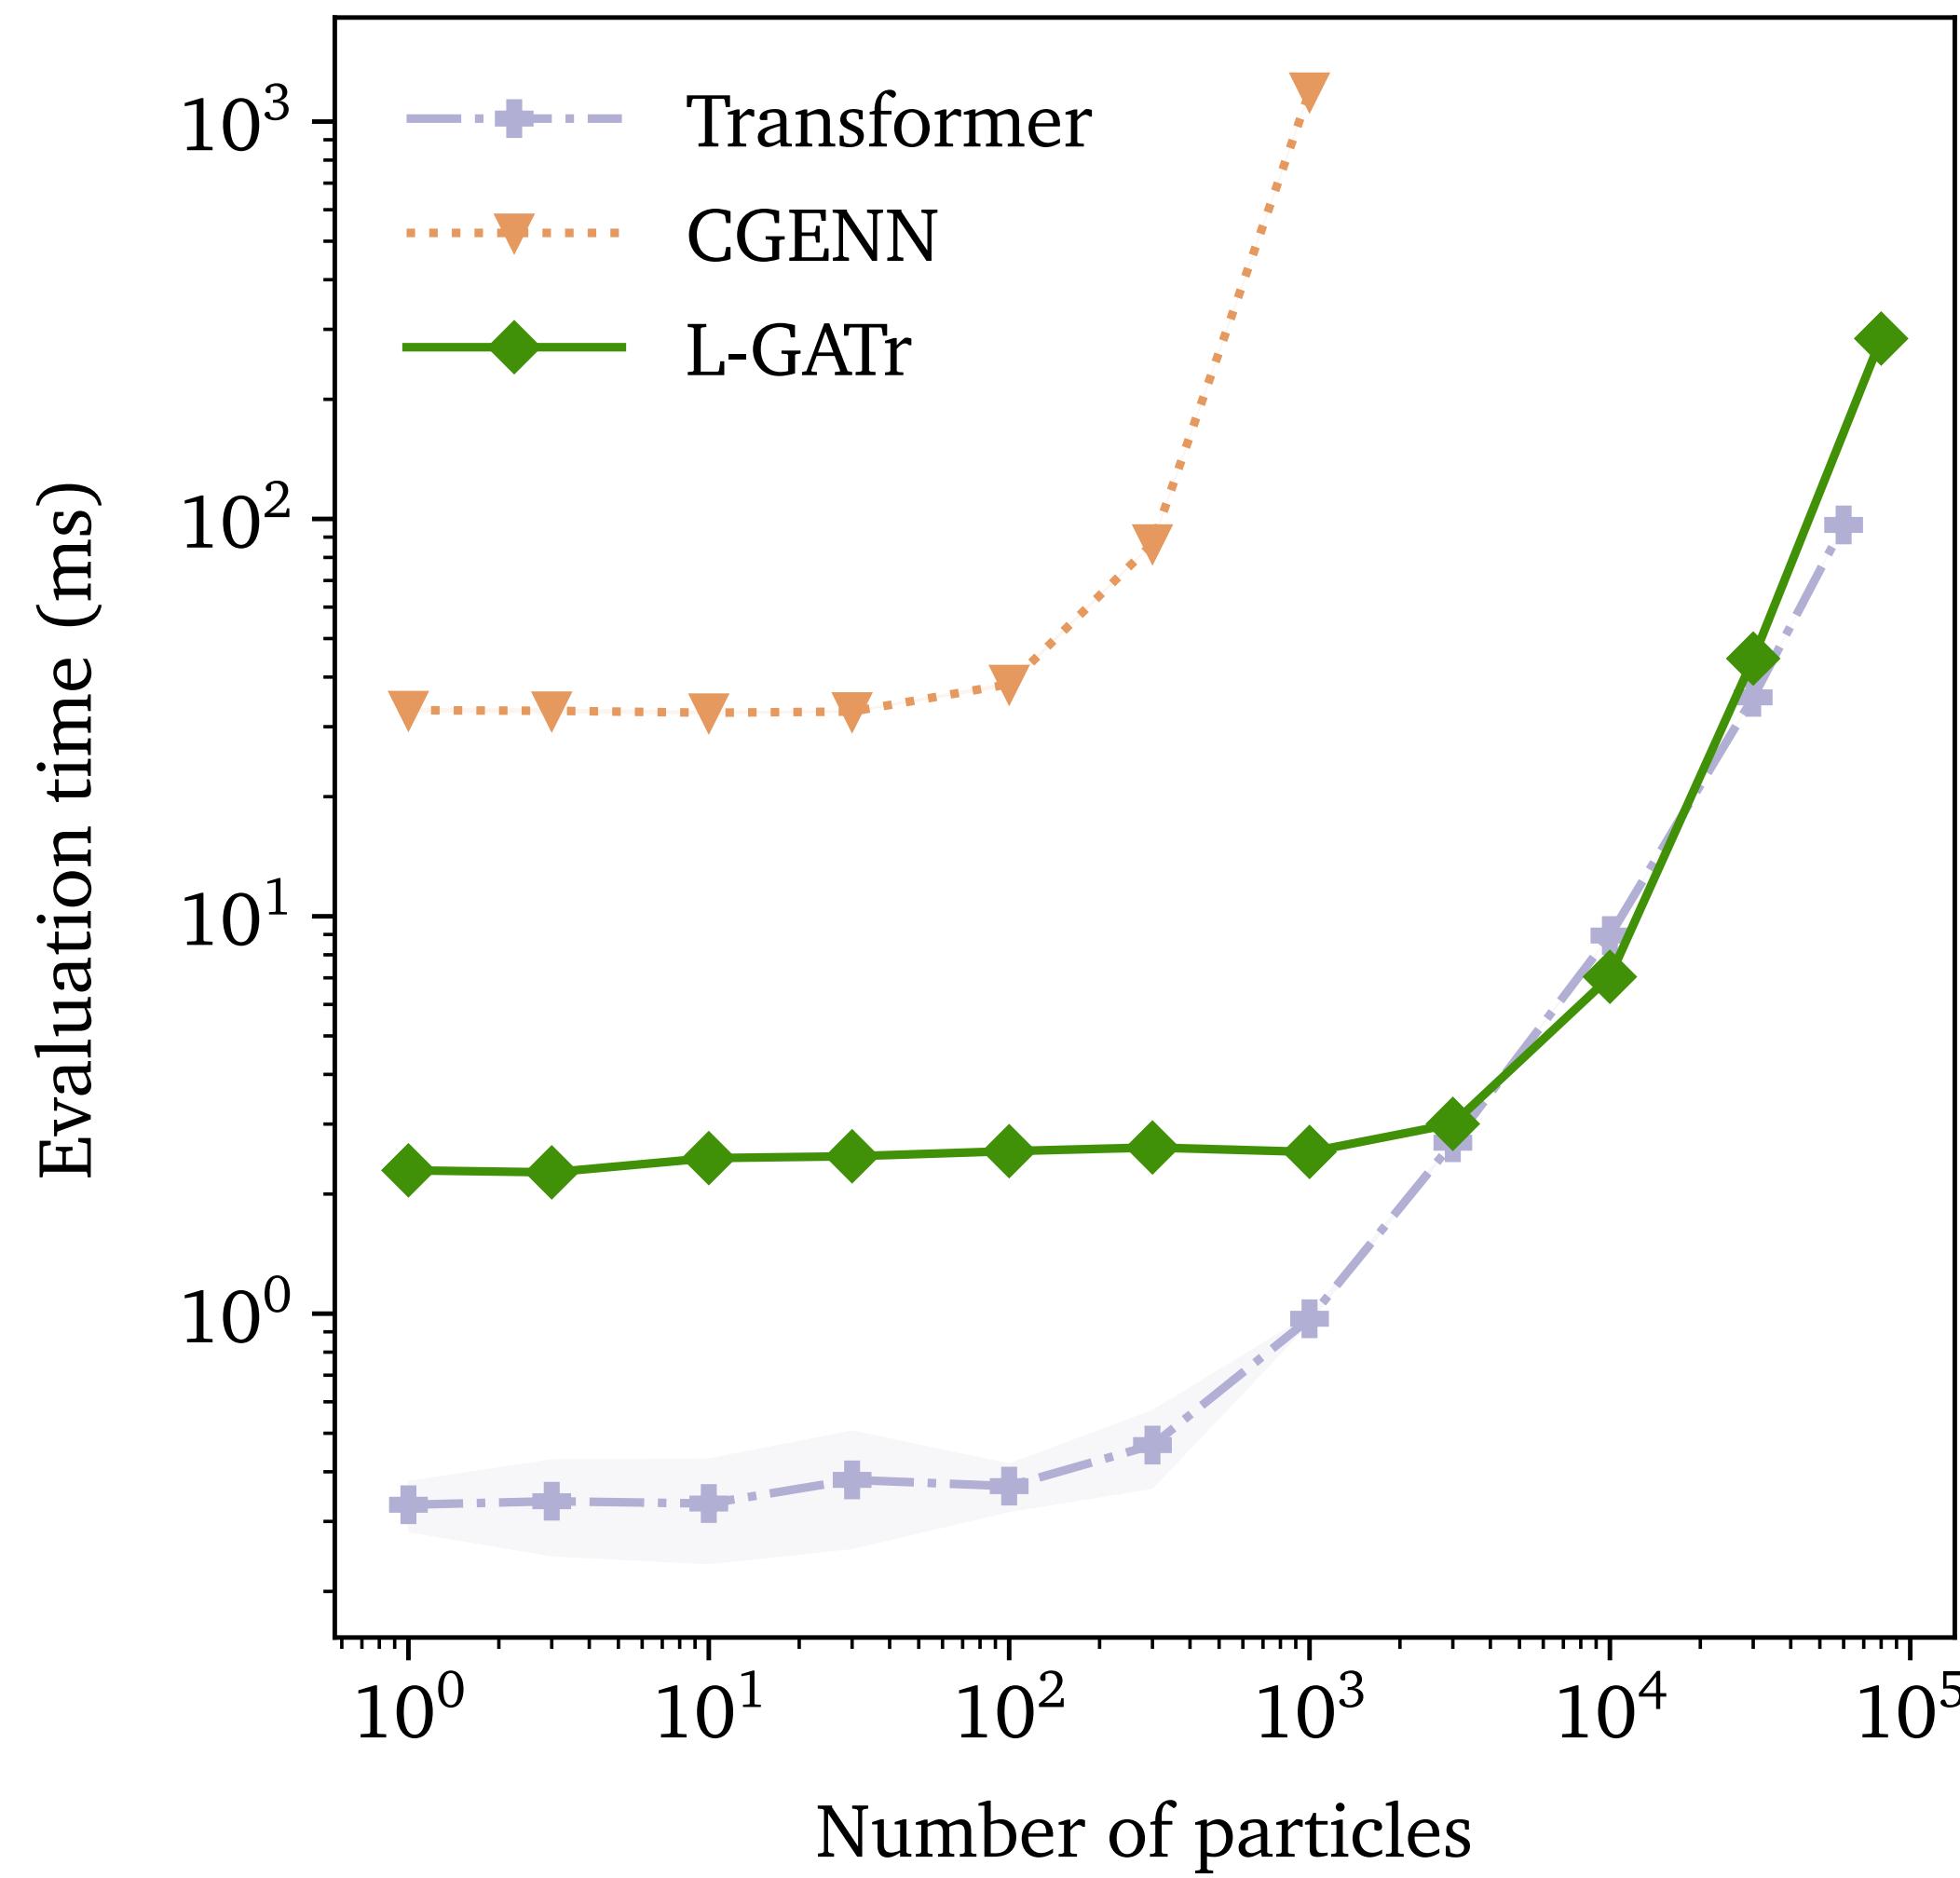
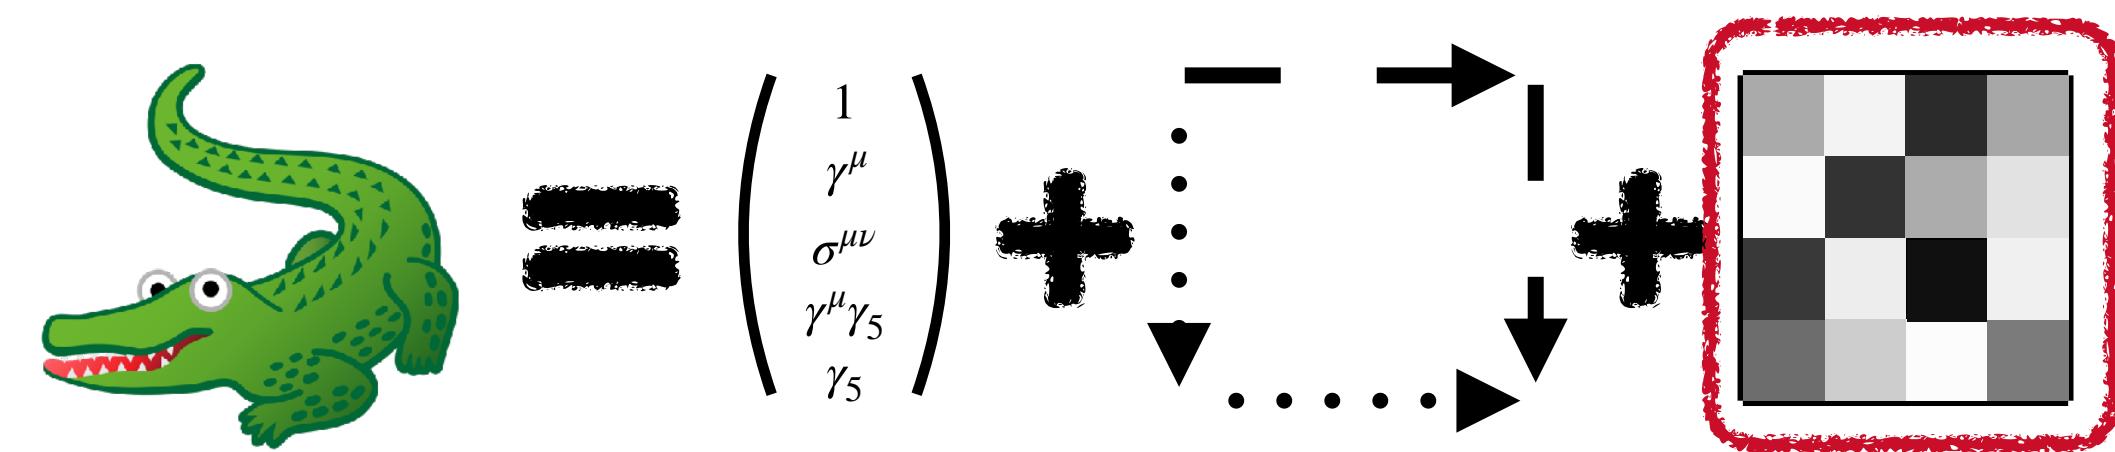
### Attention blocks

can be stacked to large depth, gradients are propagated efficiently

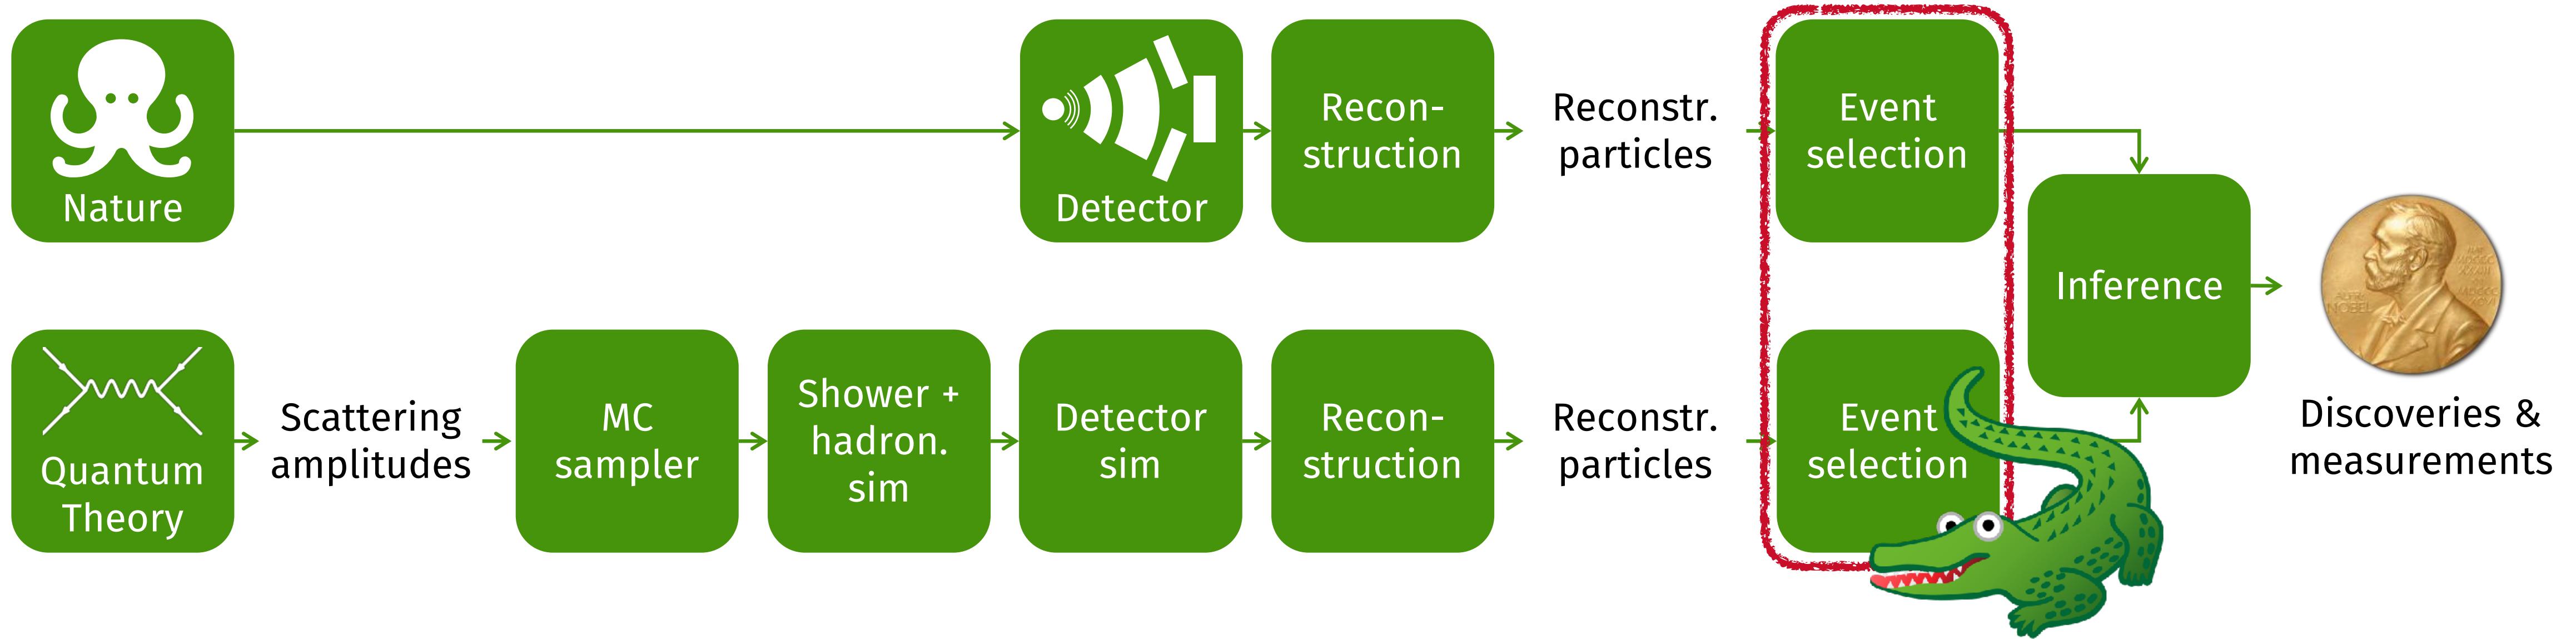
**Geometric product**  
allow for construction of new geometric types

# L-GATr

## Processing thousands of particles



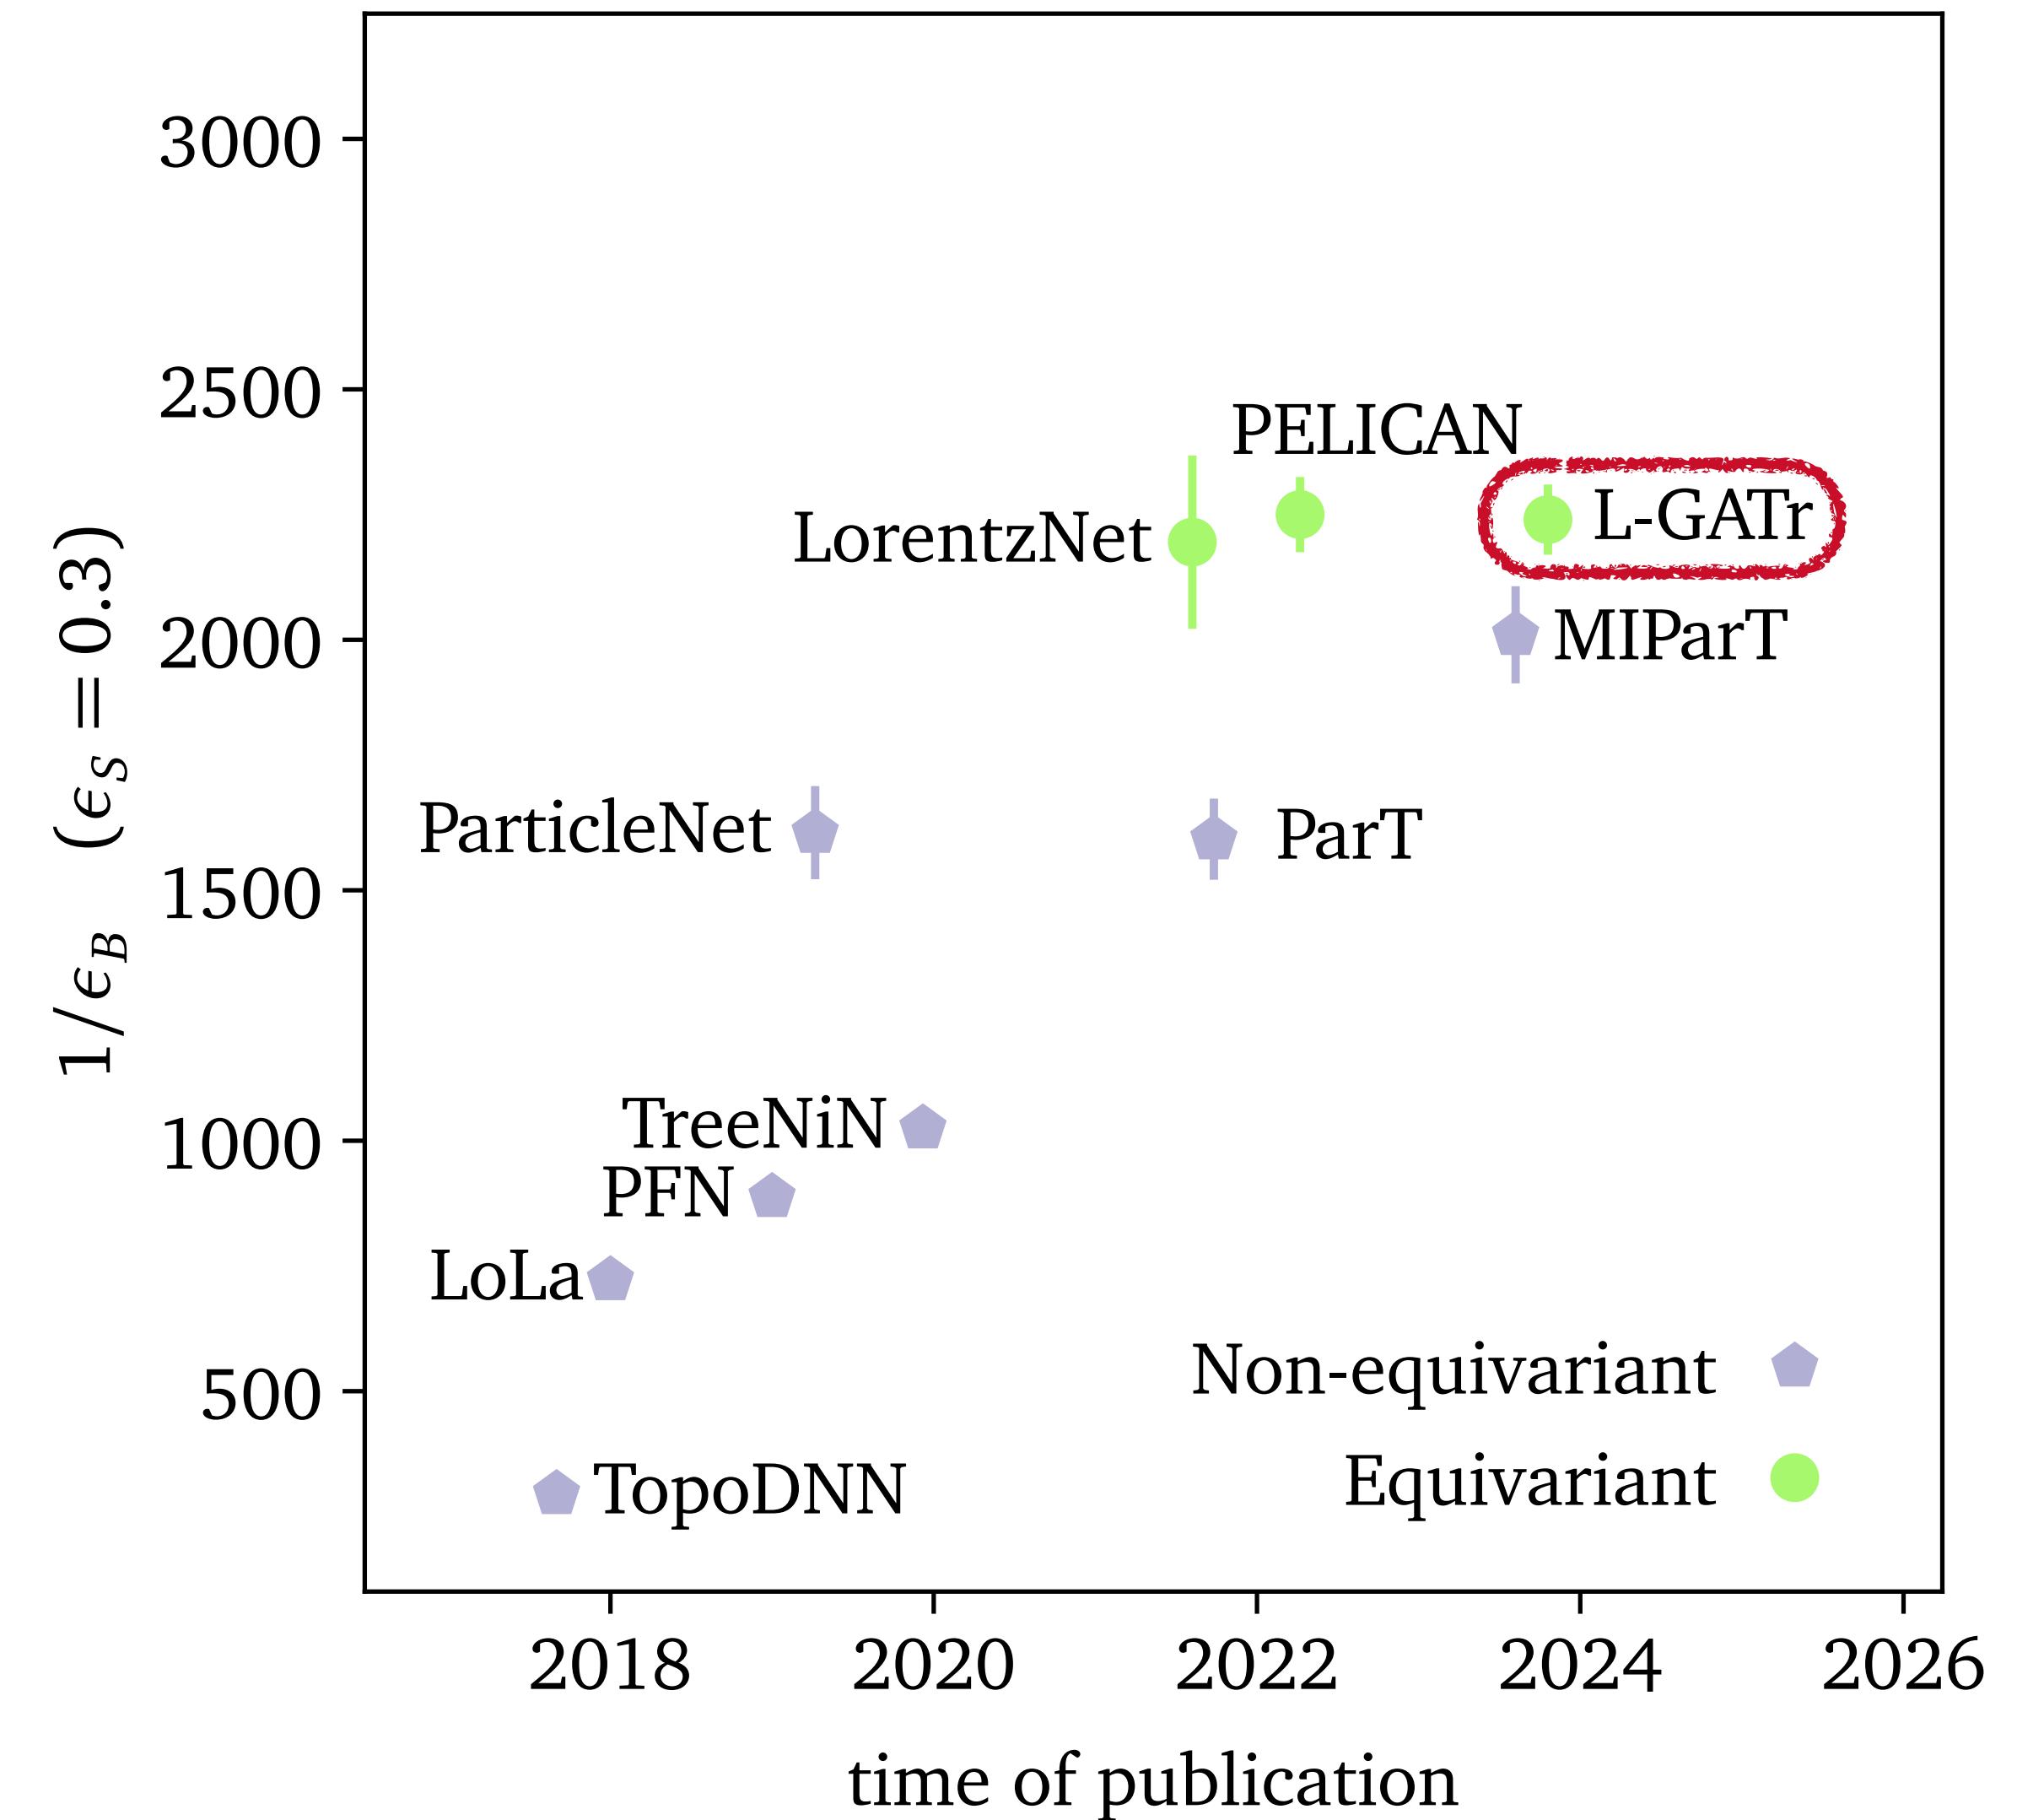
# Jet Tagging with L-GATr



# Jet Tagging with L-GATr

## The history of top tagging

All top-performing taggers  
use Lorentz symmetry

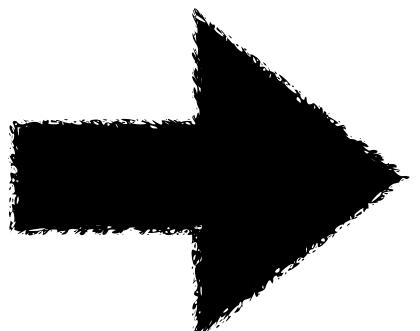


# Jet Tagging with L-GATr

## Symmetry breaking with multivectors

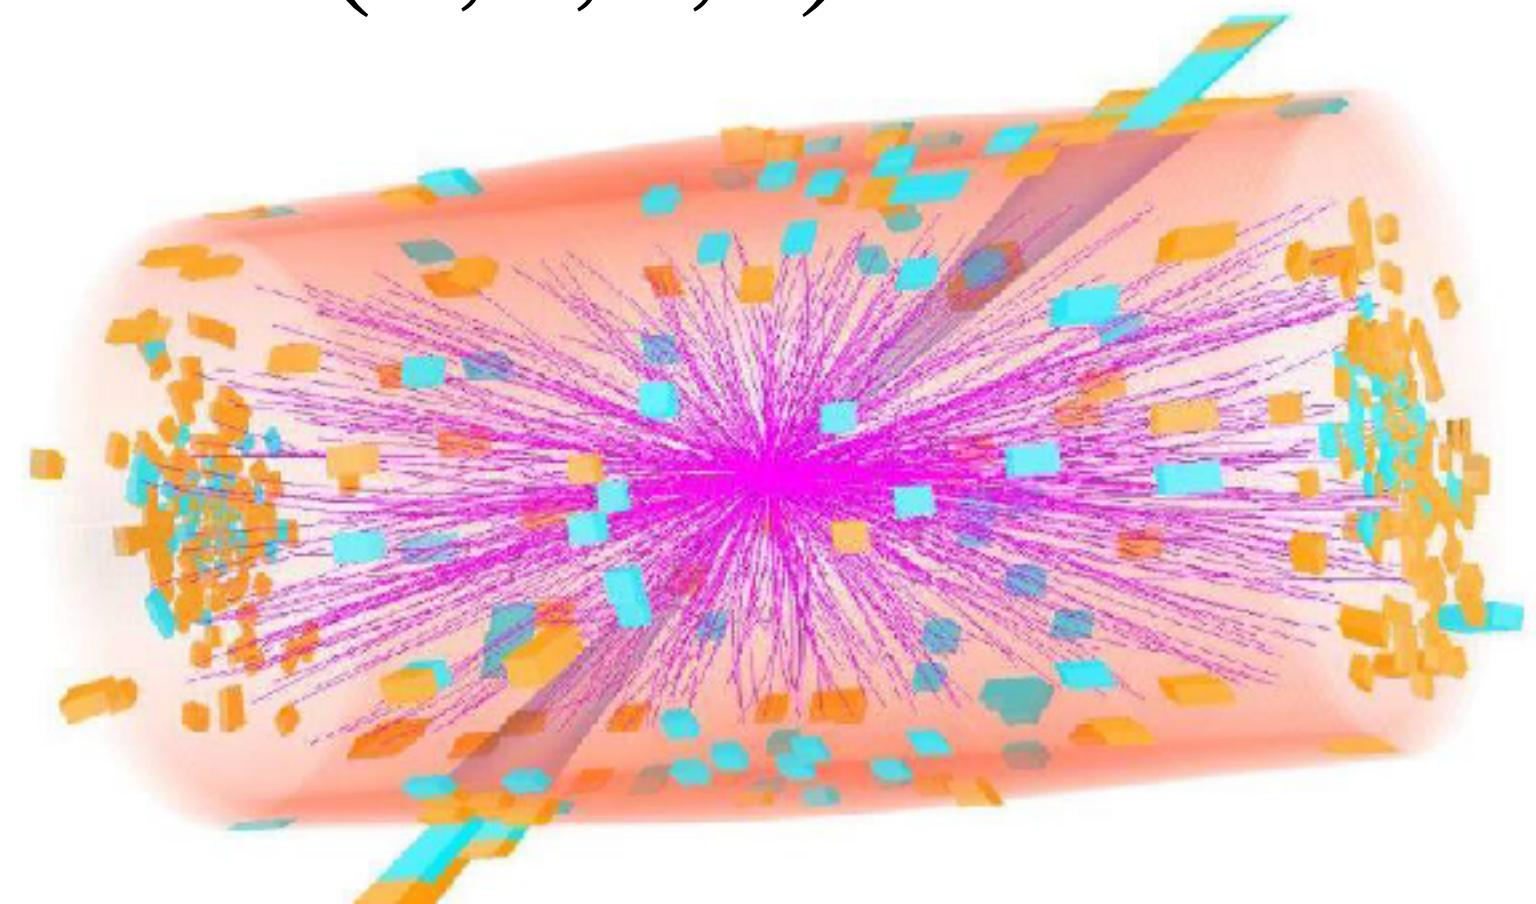
The jet tagging score is not Lorentz-invariant

- Beam direction breaks invariance under rotations around the x- and y-axis
- Detector breaks boost invariance



Add reference multivectors as extra particles to break Lorentz symmetry

- Beam reference multivector:  
 $x^V = (0,0,0, \pm 1)$  or  $x_{12}^B = 1$
- Time reference multivector:  
 $x^V = (1,0,0,0)$



All Lorentz-invariant taggers use symmetry-breaking inputs

# Jet Tagging with L-GATr

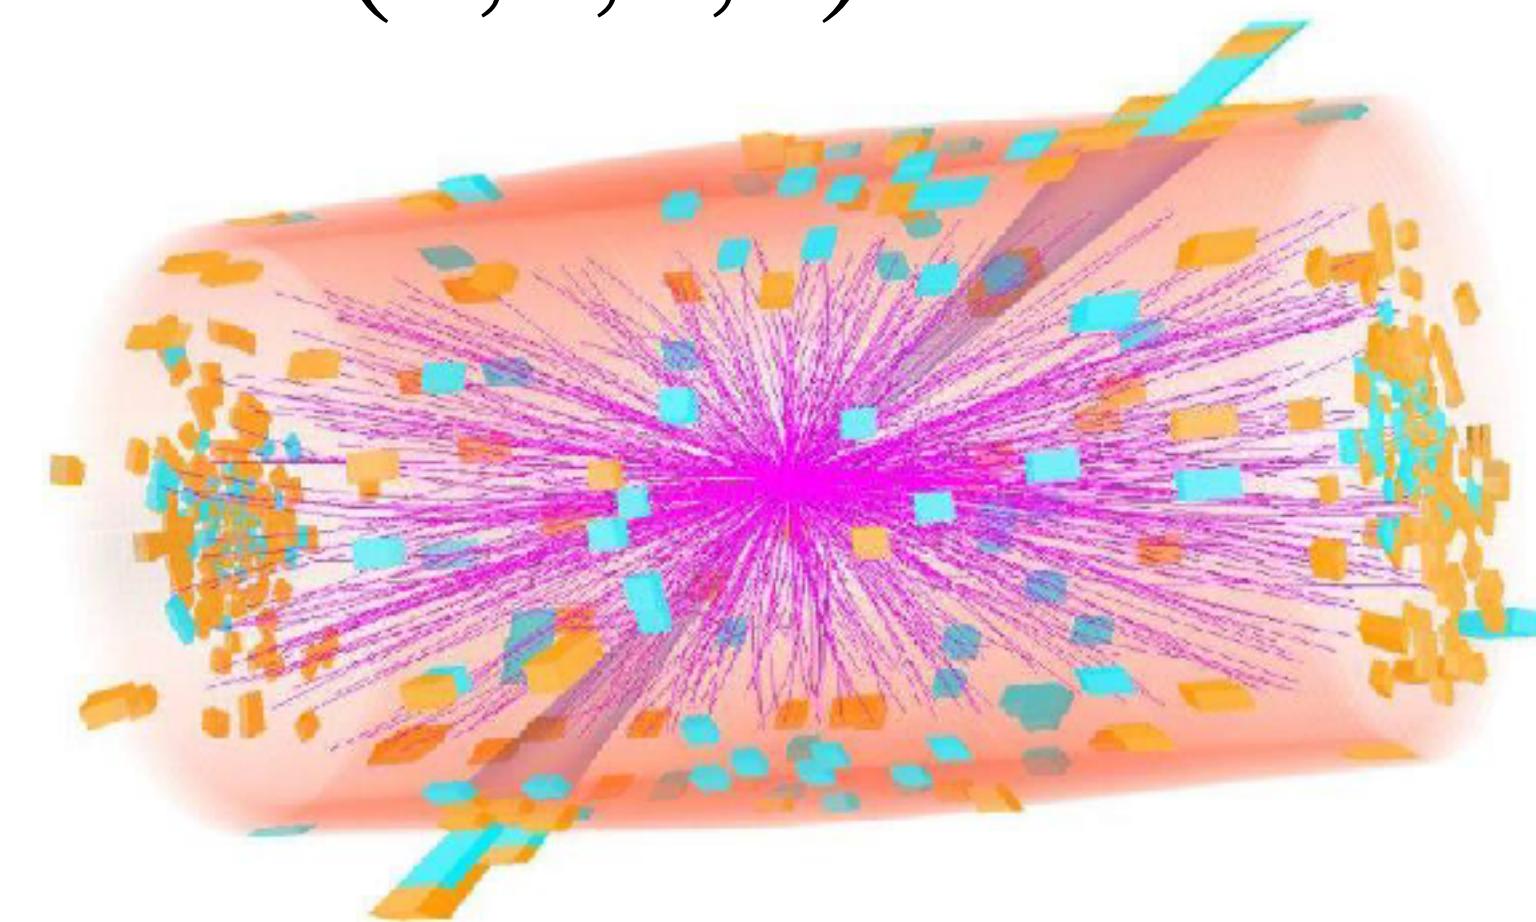
## Symmetry breaking with multivectors

Beam	Time	Embedding	AUC	$1/\epsilon_B$ ( $\epsilon_s = 0.3$ )
-	-	Token	0.9844	1422
$x_3^V = \pm 1$	-	Token	0.9850	1905
-	$x_0^V = 1$	Token	0.9865	1923
$x_{12}^B = x_{13}^B = x_{23}^B = 1$	$x_0^V = 1$	Token	0.9863	2009
$x_{12}^B = 1$	$x_0^V = 1$	Channel	0.9865	2060
$x_0^V = 1, x_3^V = \pm 1$	$x_0^V = 1$	Token	0.9869	2114
$x_3^V = \pm 1$	$x_0^V = 1$	Token	0.9869	2152
$x_{12}^B = 1$	$x_0^V = 1$	Token	0.9870	2240

Several choices give  
essentially the same performance

Add reference multivectors as extra particles to break Lorentz symmetry

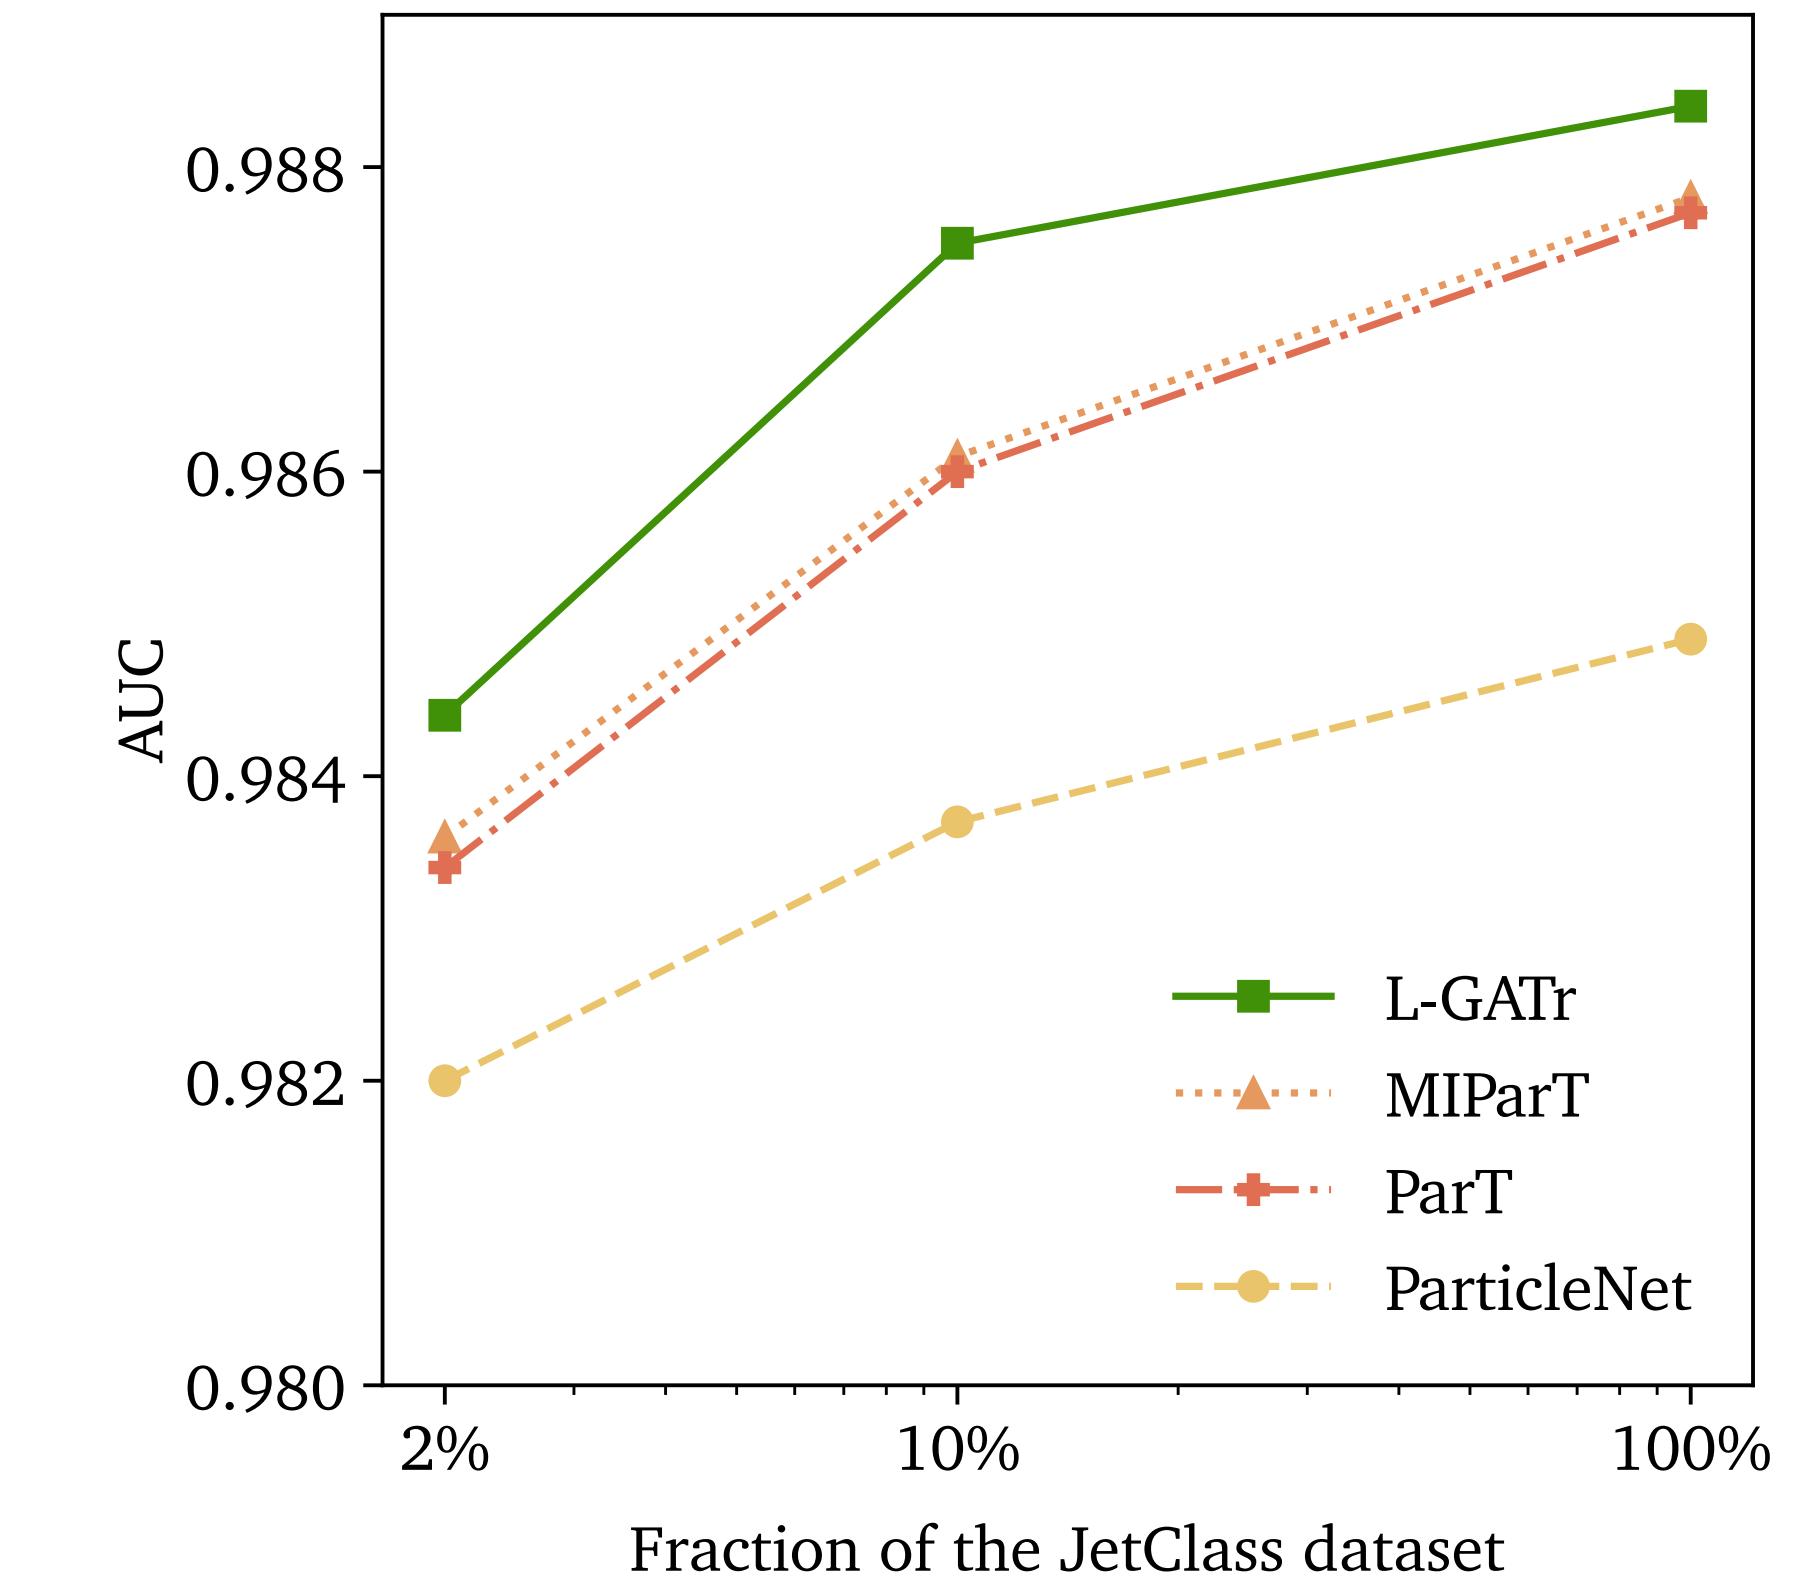
- Beam reference multivector:  
 $x^V = (0,0,0, \pm 1)$  or  $x_{12}^B = 1$
- Time reference multivector:  
 $x^V = (1,0,0,0)$



# Jet Tagging with L-GATr

Training on 100M jets

Transformers significantly outperform graph networks

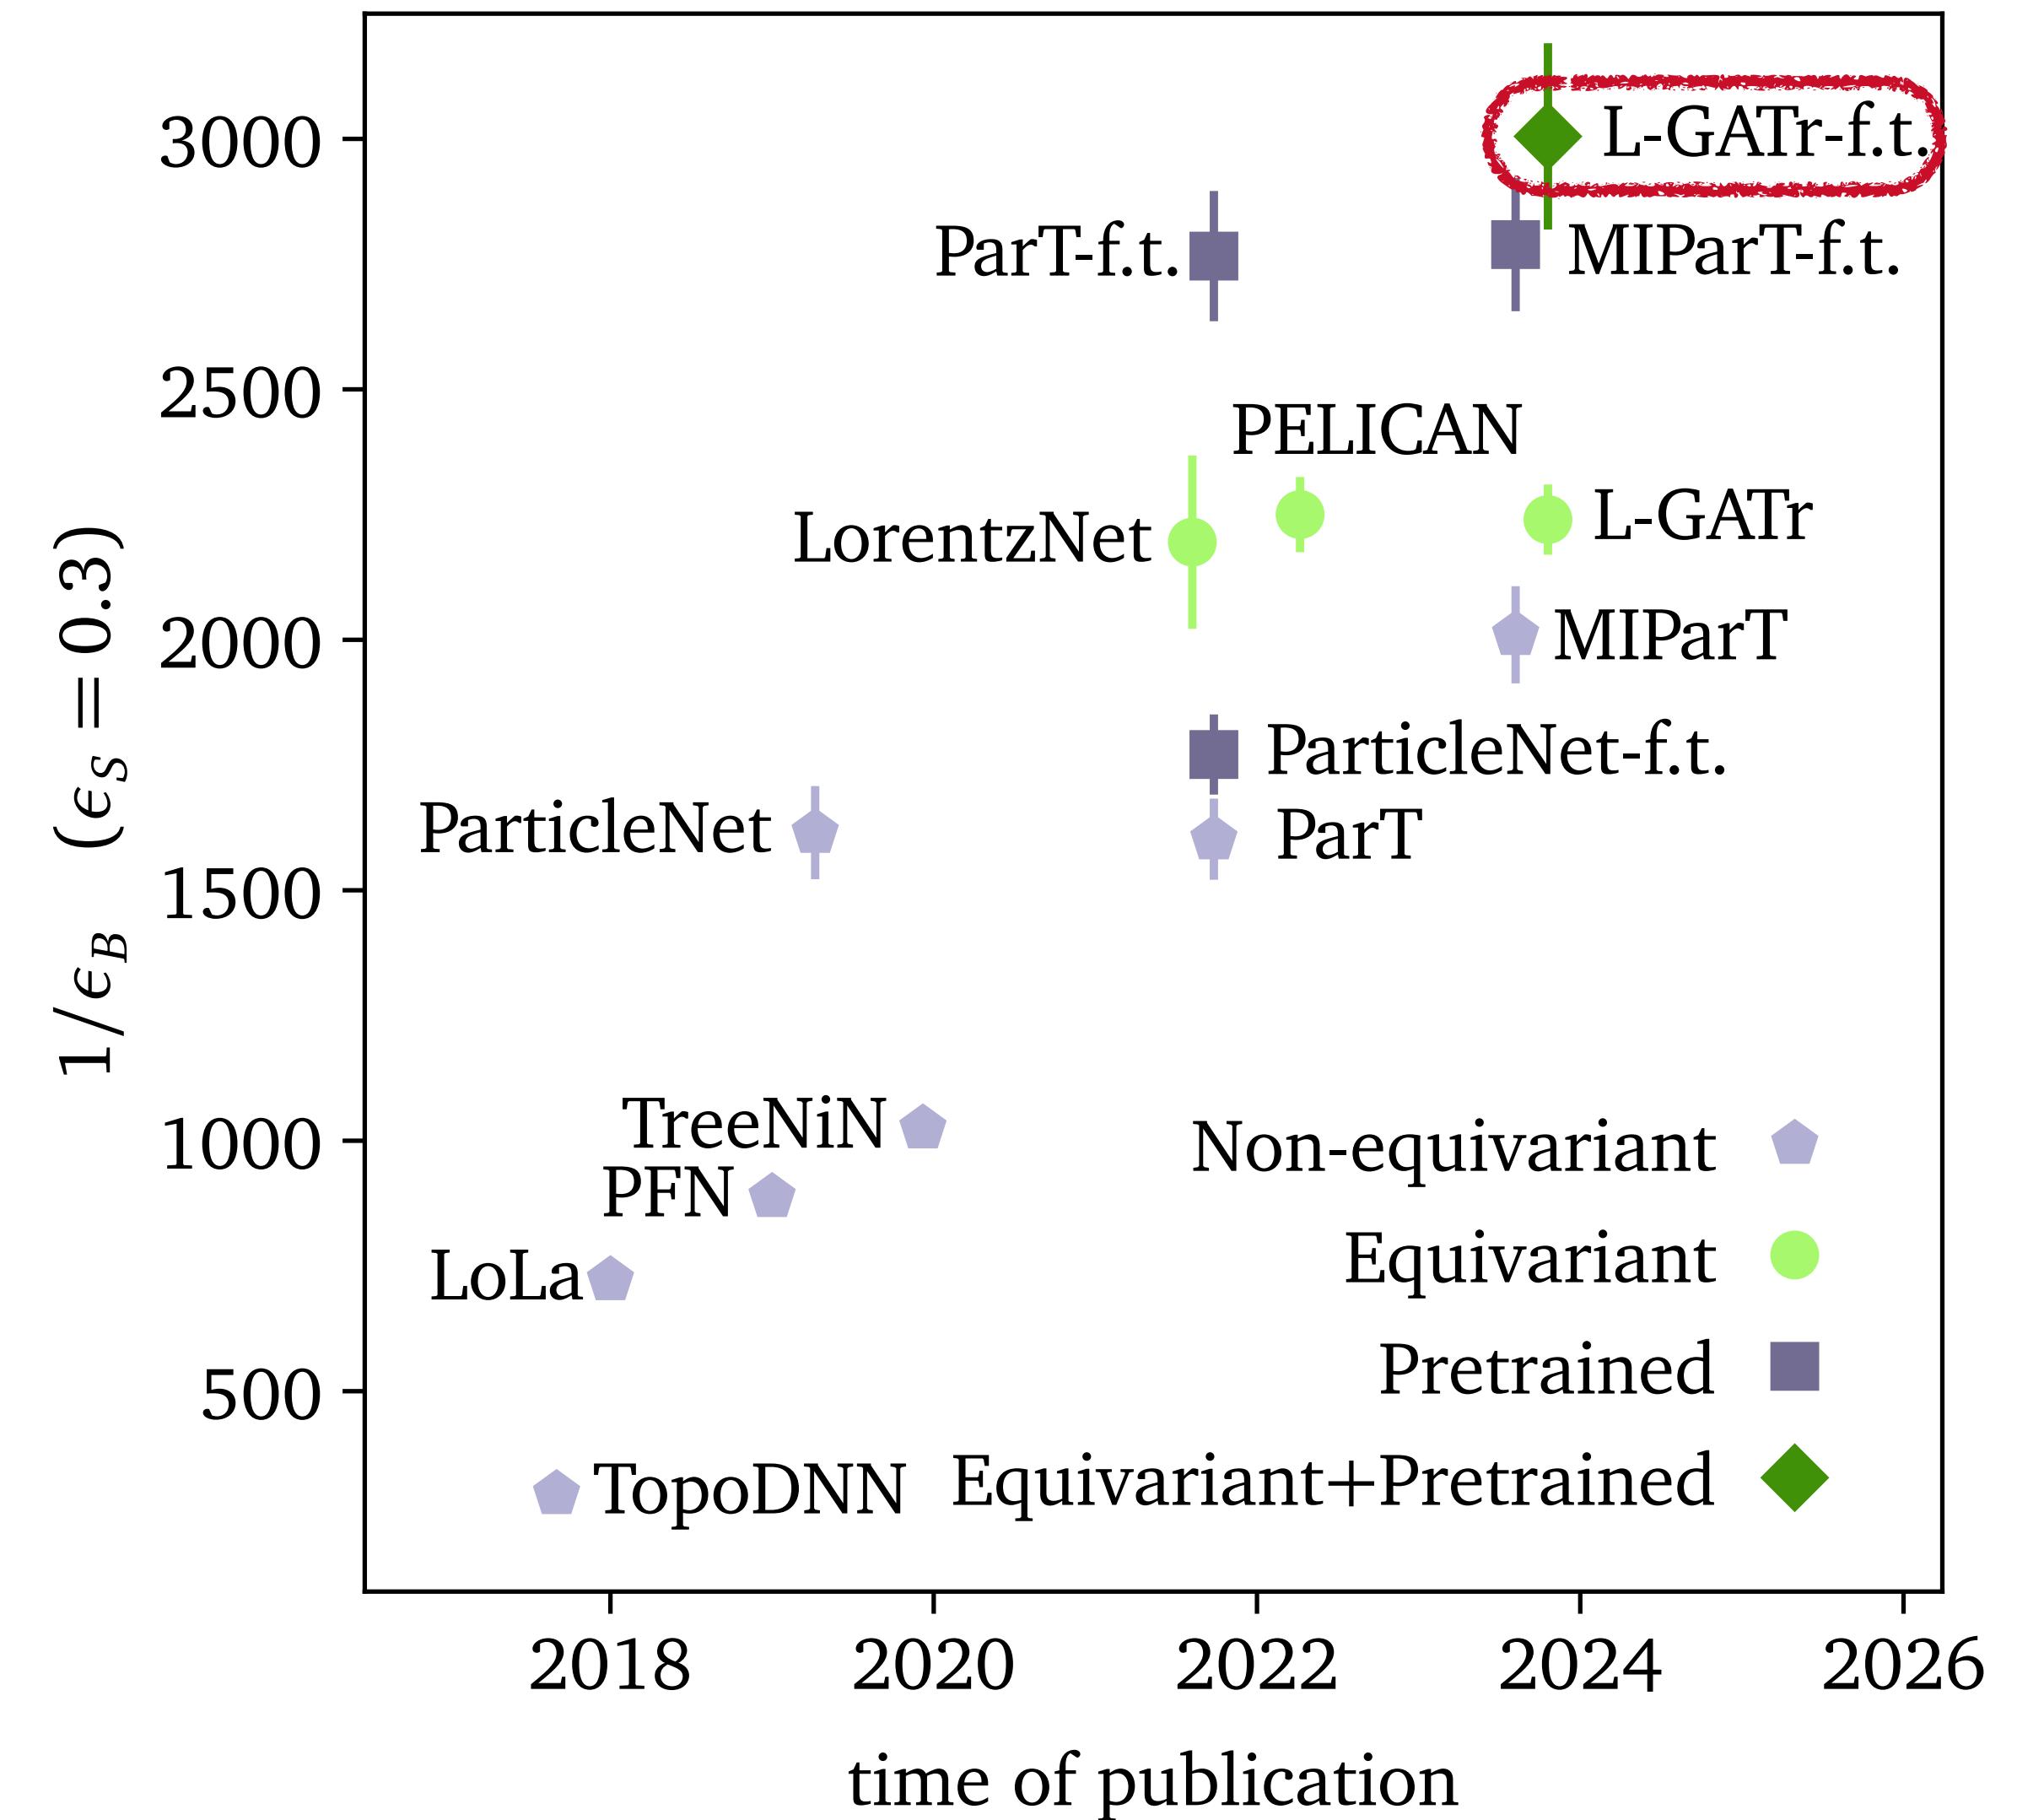


	All classes Accuracy	AUC	$H \rightarrow b\bar{b}$ Rej <sub>50%</sub>	$H \rightarrow c\bar{c}$ Rej <sub>50%</sub>	$H \rightarrow gg$ Rej <sub>50%</sub>	$H \rightarrow 4q$ Rej <sub>50%</sub>	$H \rightarrow l\nu q\bar{q}'$ Rej <sub>99%</sub>	$t \rightarrow bq\bar{q}'$ Rej <sub>50%</sub>	$t \rightarrow bl\nu$ Rej <sub>99.5%</sub>	$W \rightarrow q\bar{q}'$ Rej <sub>50%</sub>	$Z \rightarrow q\bar{q}$ Rej <sub>50%</sub>
ParticleNet [58]	0.844	0.9849	7634	2475	104	954	3339	10526	11173	347	283
ParT [59]	0.861	0.9877	10638	4149	123	1864	5479	32787	15873	543	402
MIParT [60]	0.861	0.9878	10753	4202	123	1927	5450	31250	16807	542	402
L-GATr	0.865	0.9884	12195	4819	128	2304	5764	37736	19231	580	427

# Jet Tagging with L-GATr

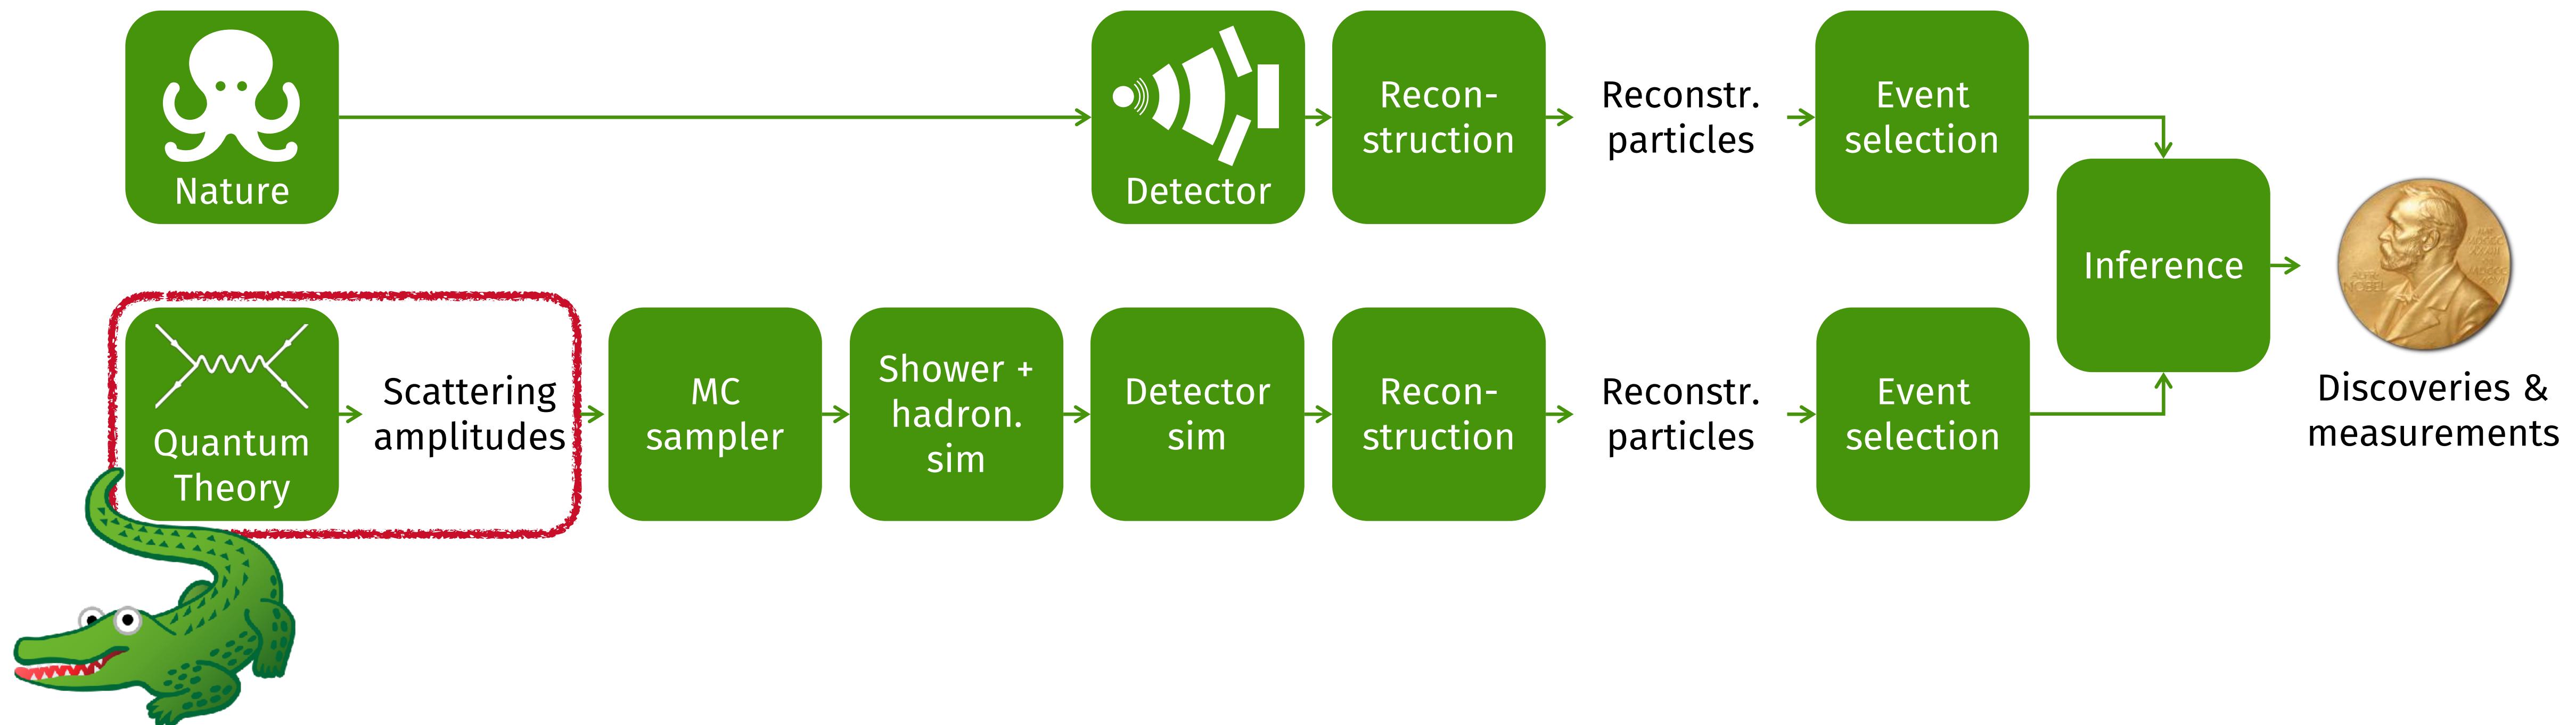
## Transfer learning for top tagging

L-GATr combines benefits from pretraining and Lorentz equivariance



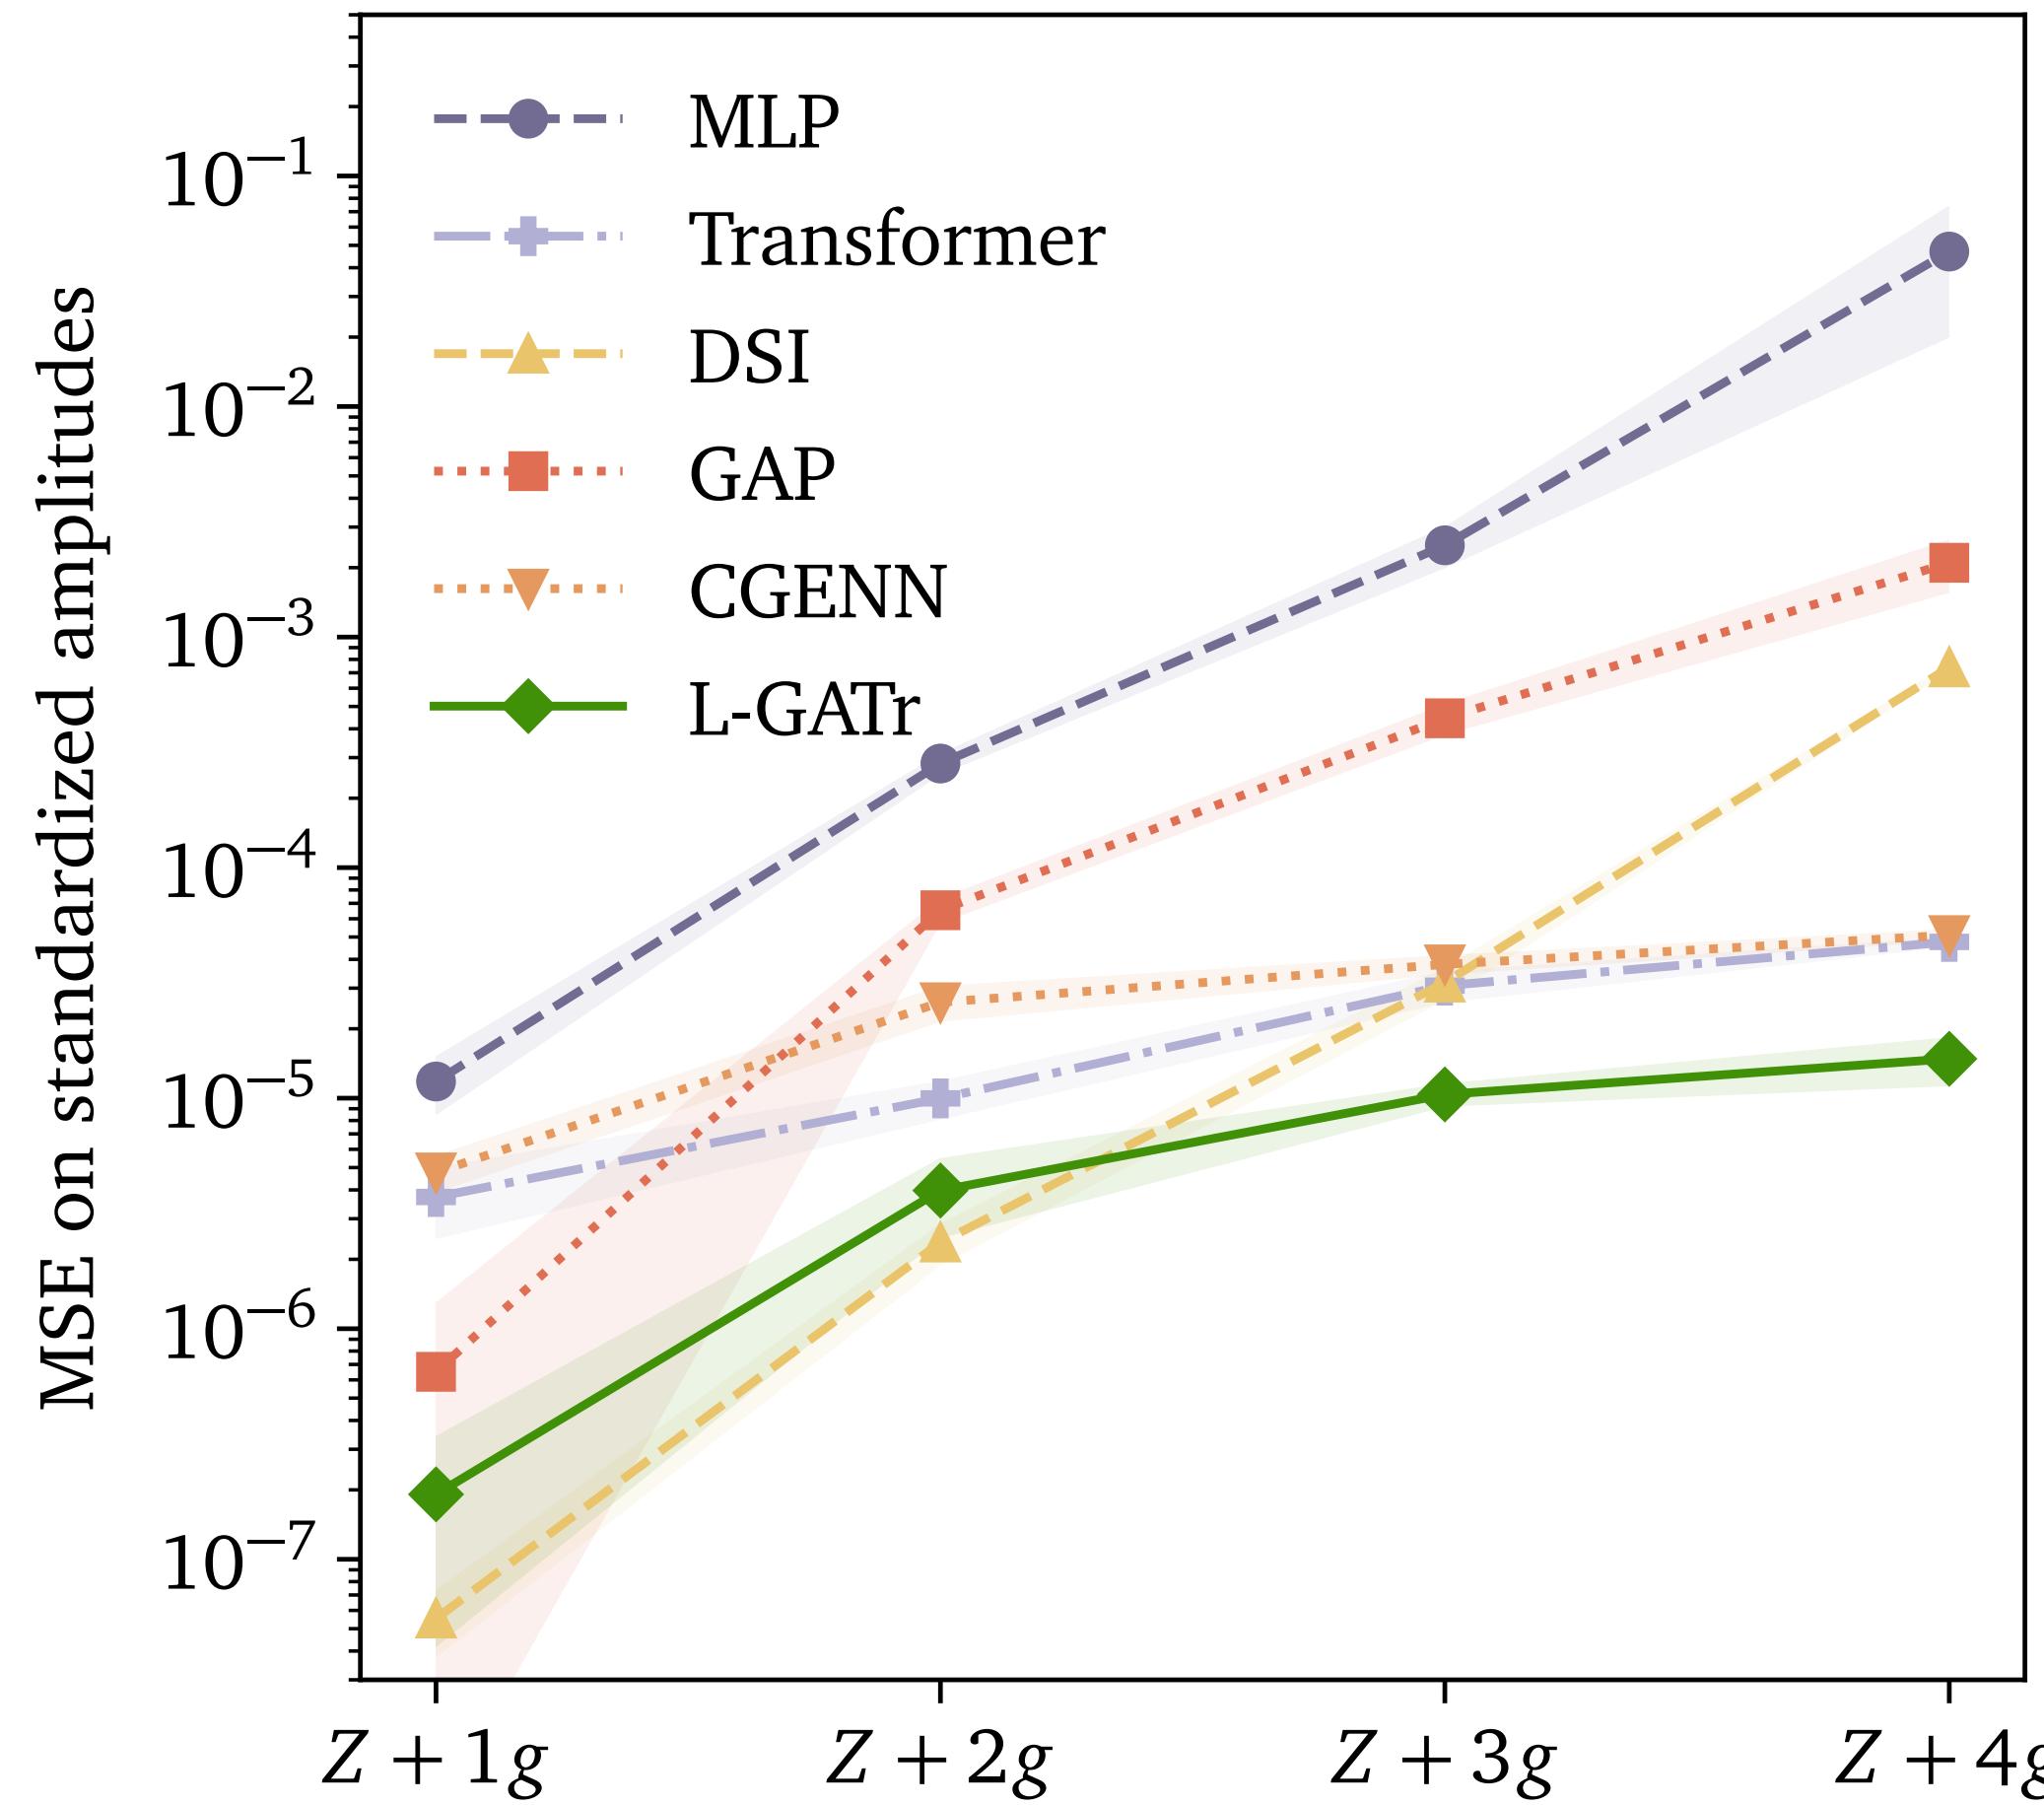
# L-GATr beyond jet tagging

## QFT amplitude regression



# L-GATr beyond jet tagging

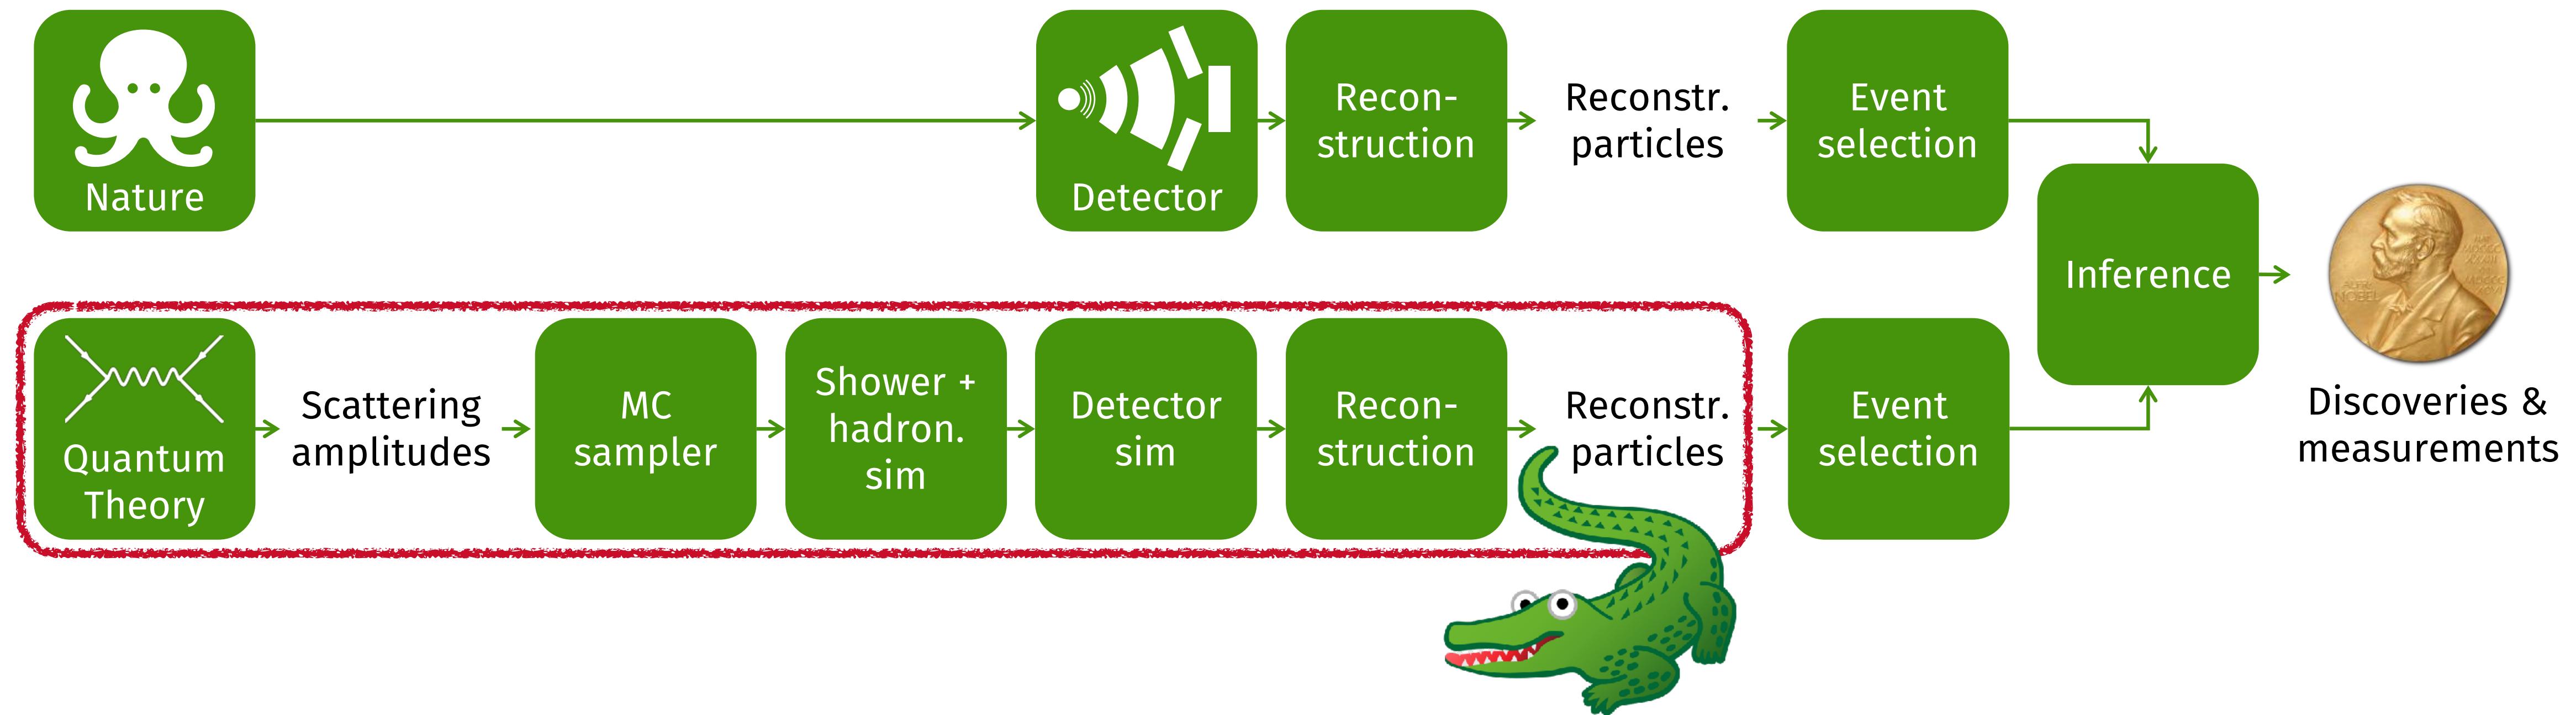
## QFT amplitude regression



The combination Lorentz-equivariance  
+ permutation-equivariant Transformer  
scales best to high multiplicity

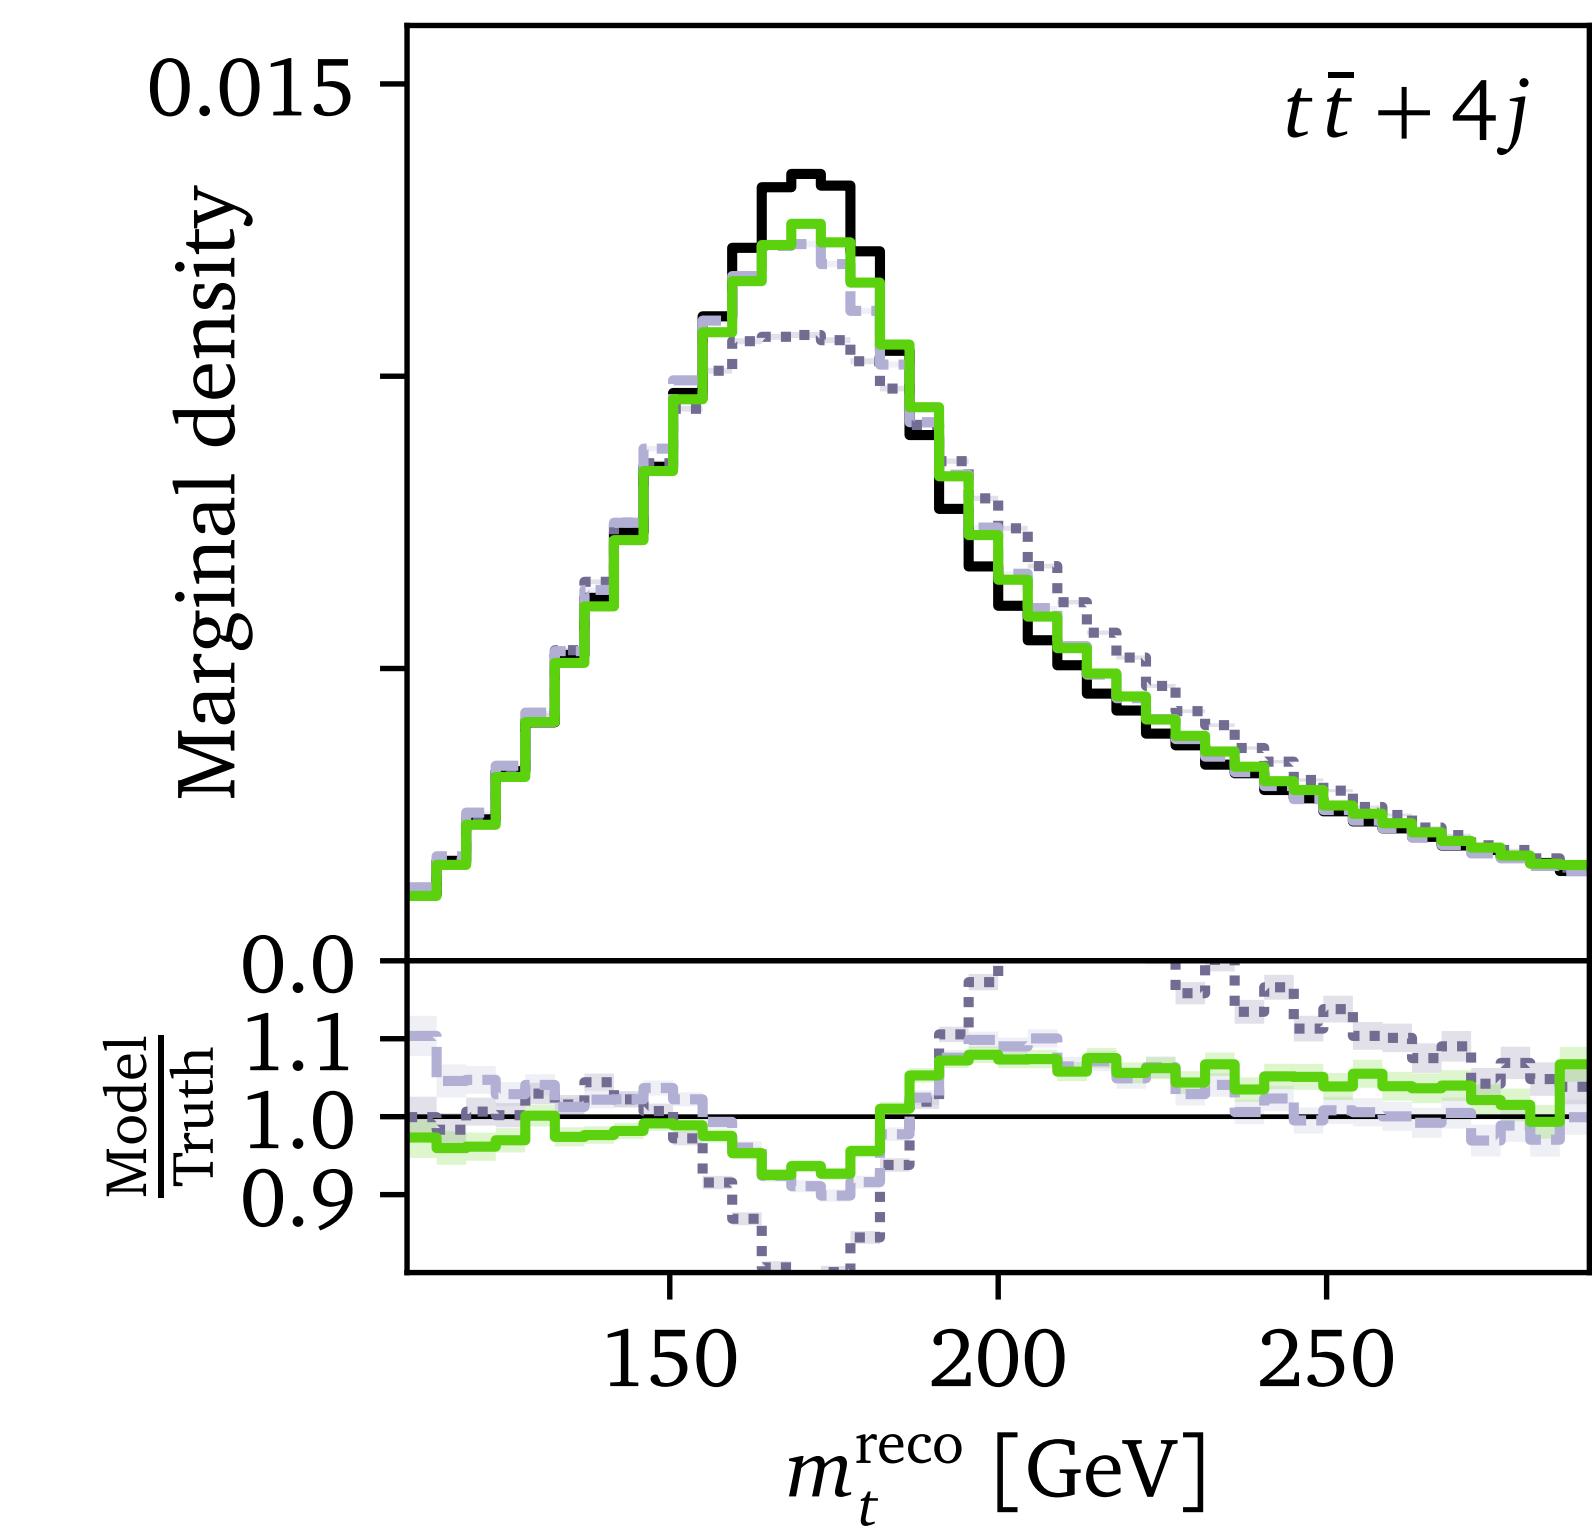
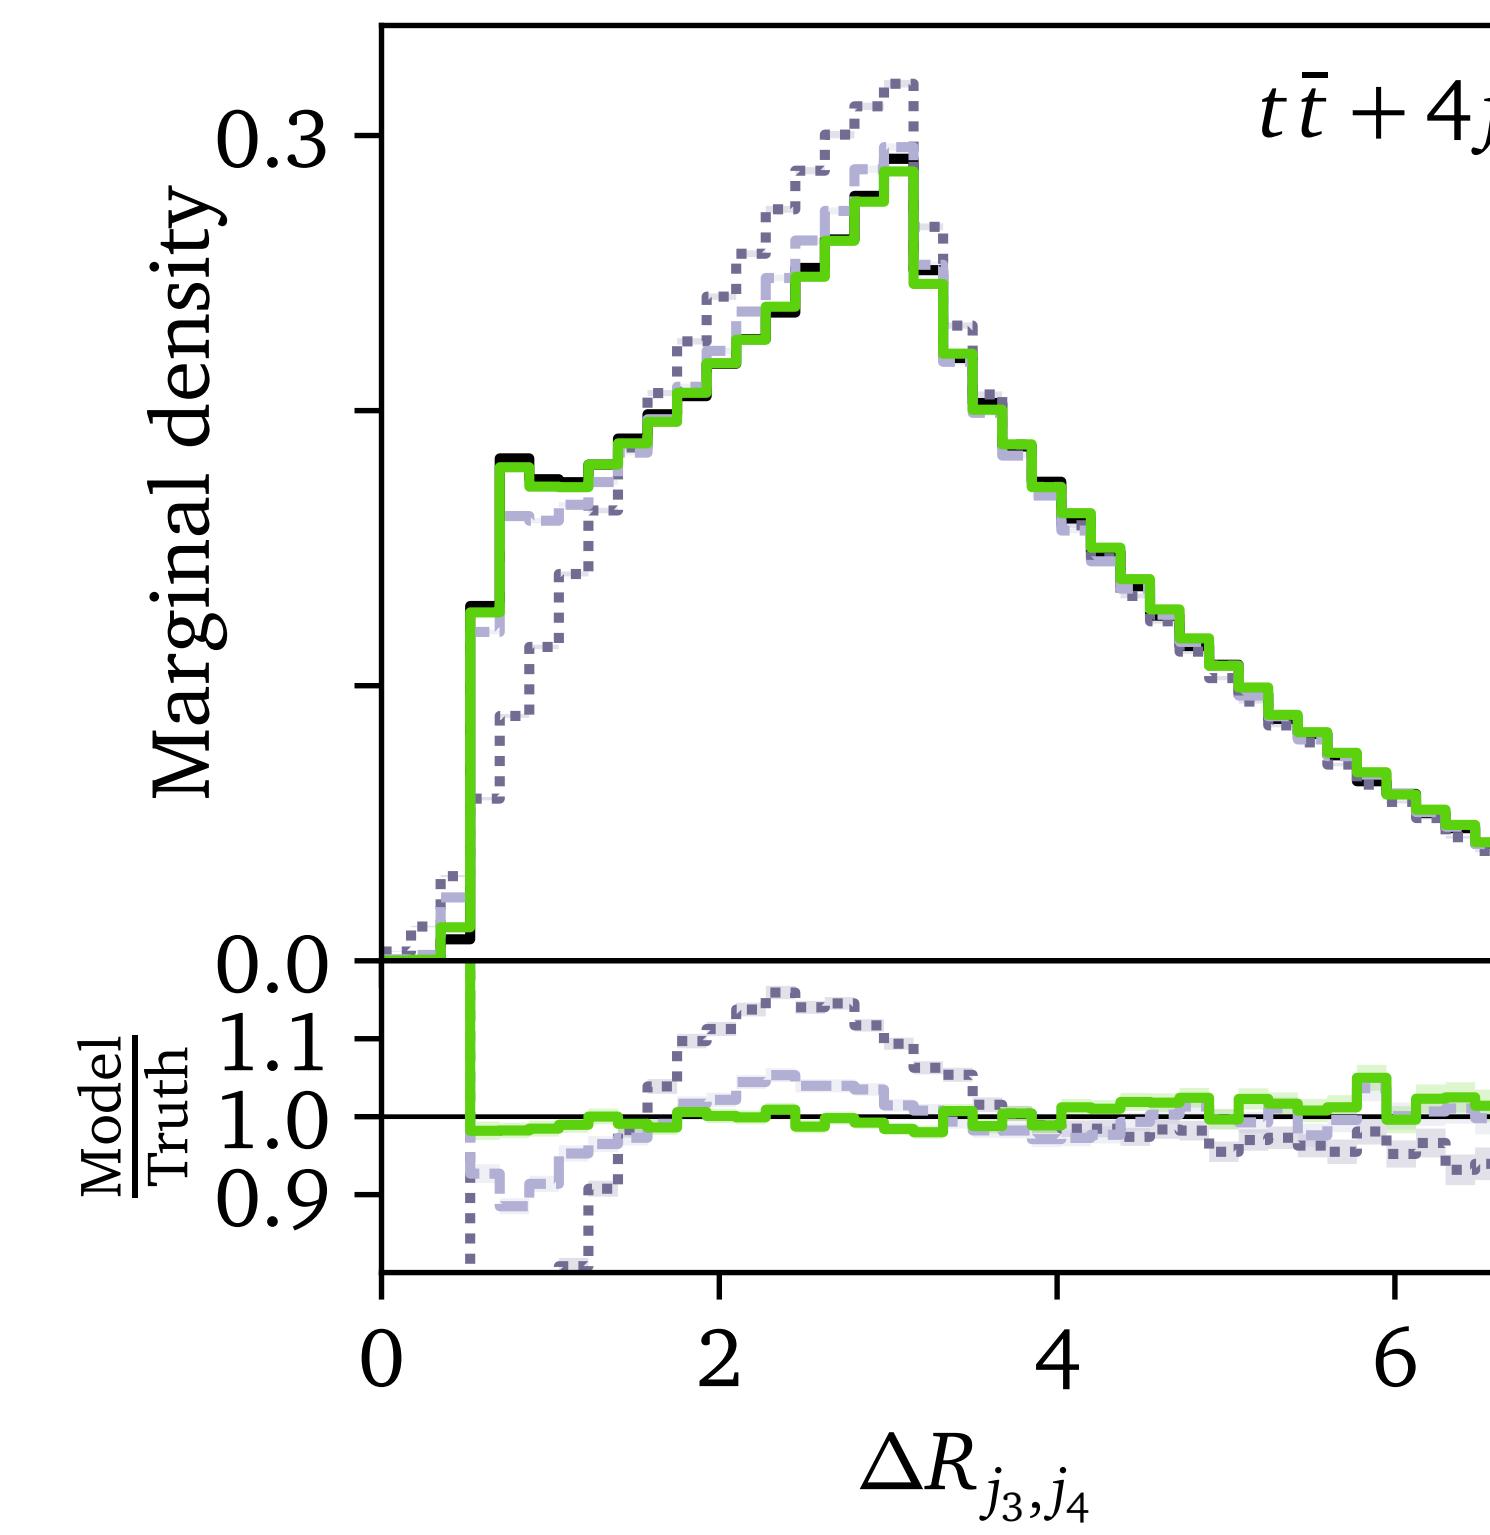
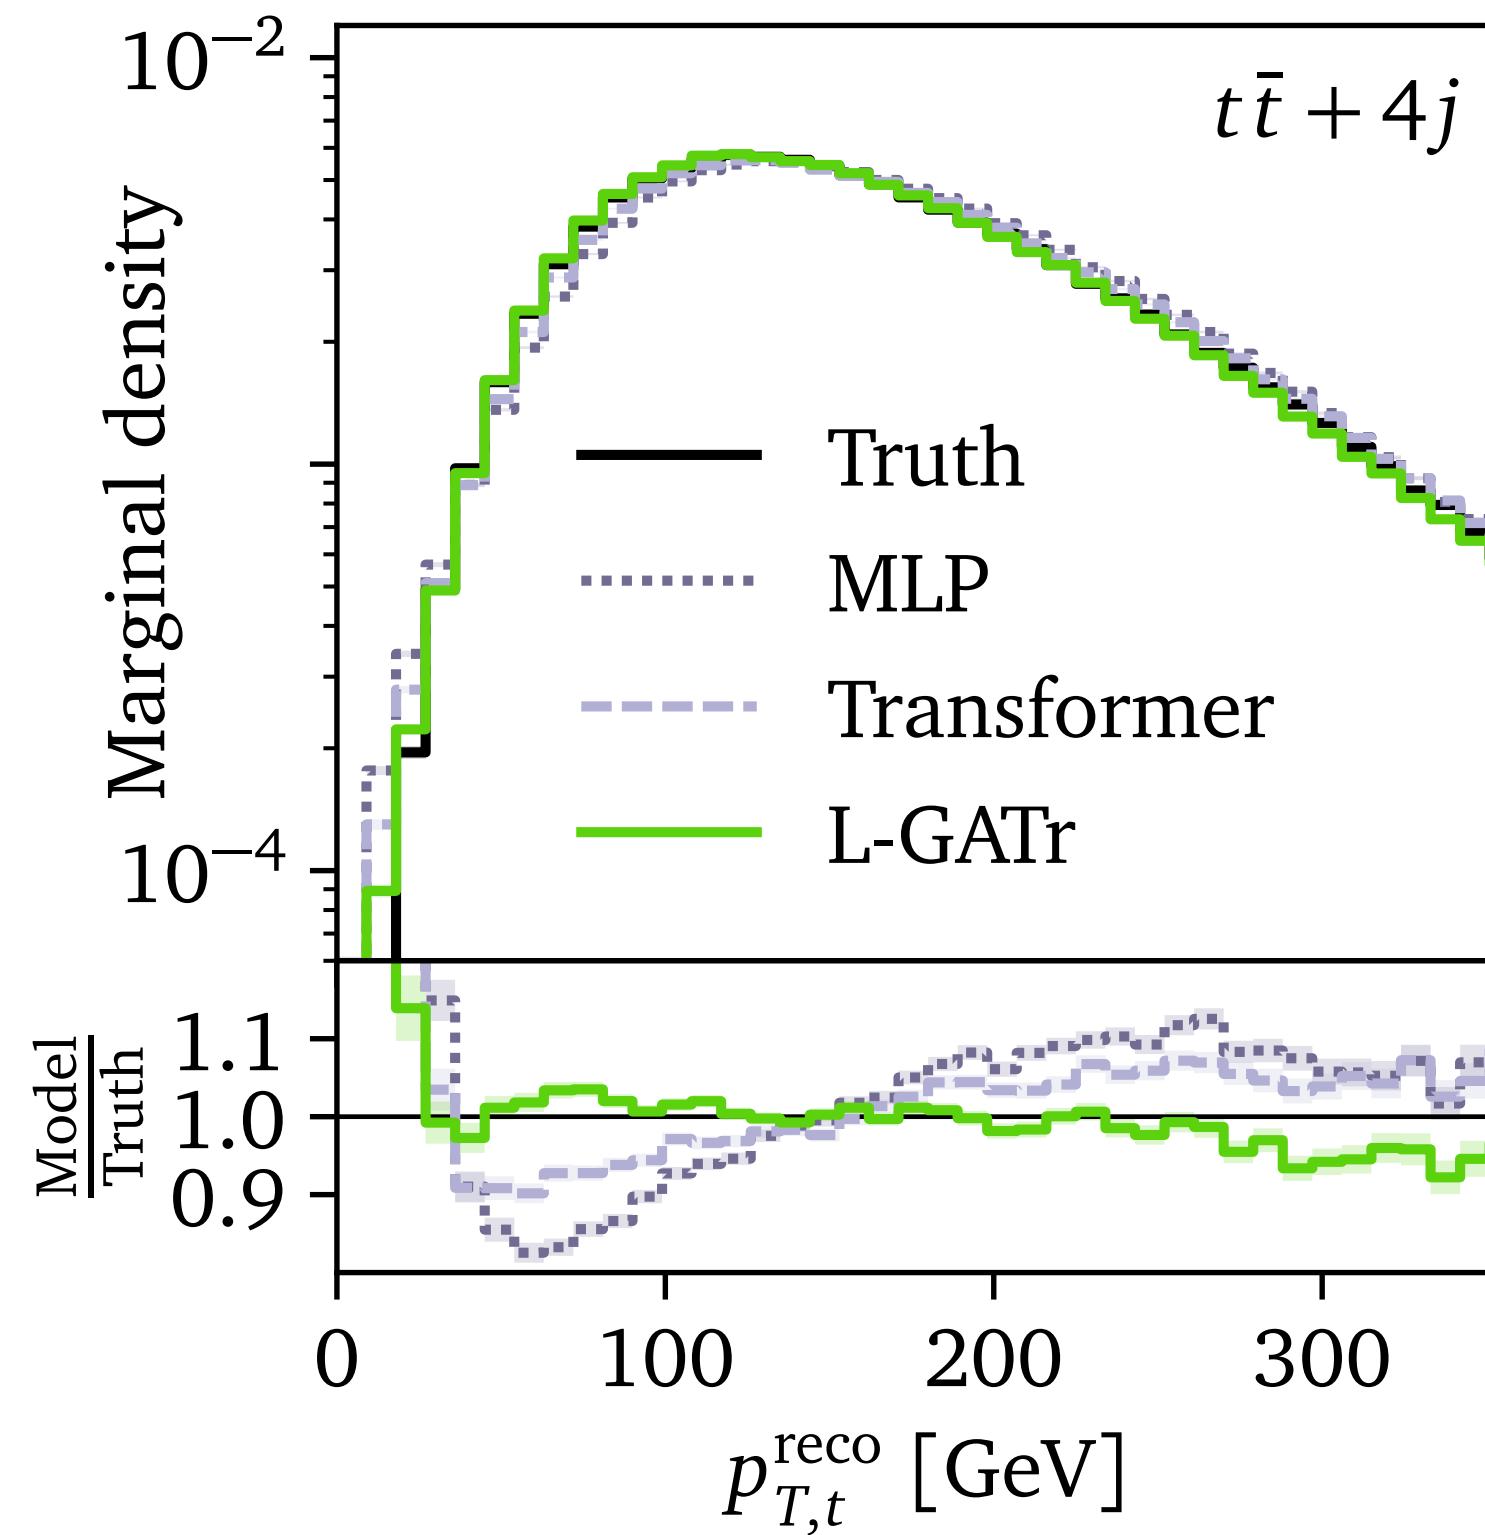
# L-GATr beyond jet tagging

## LHC event generation



# L-GATr beyond jet tagging

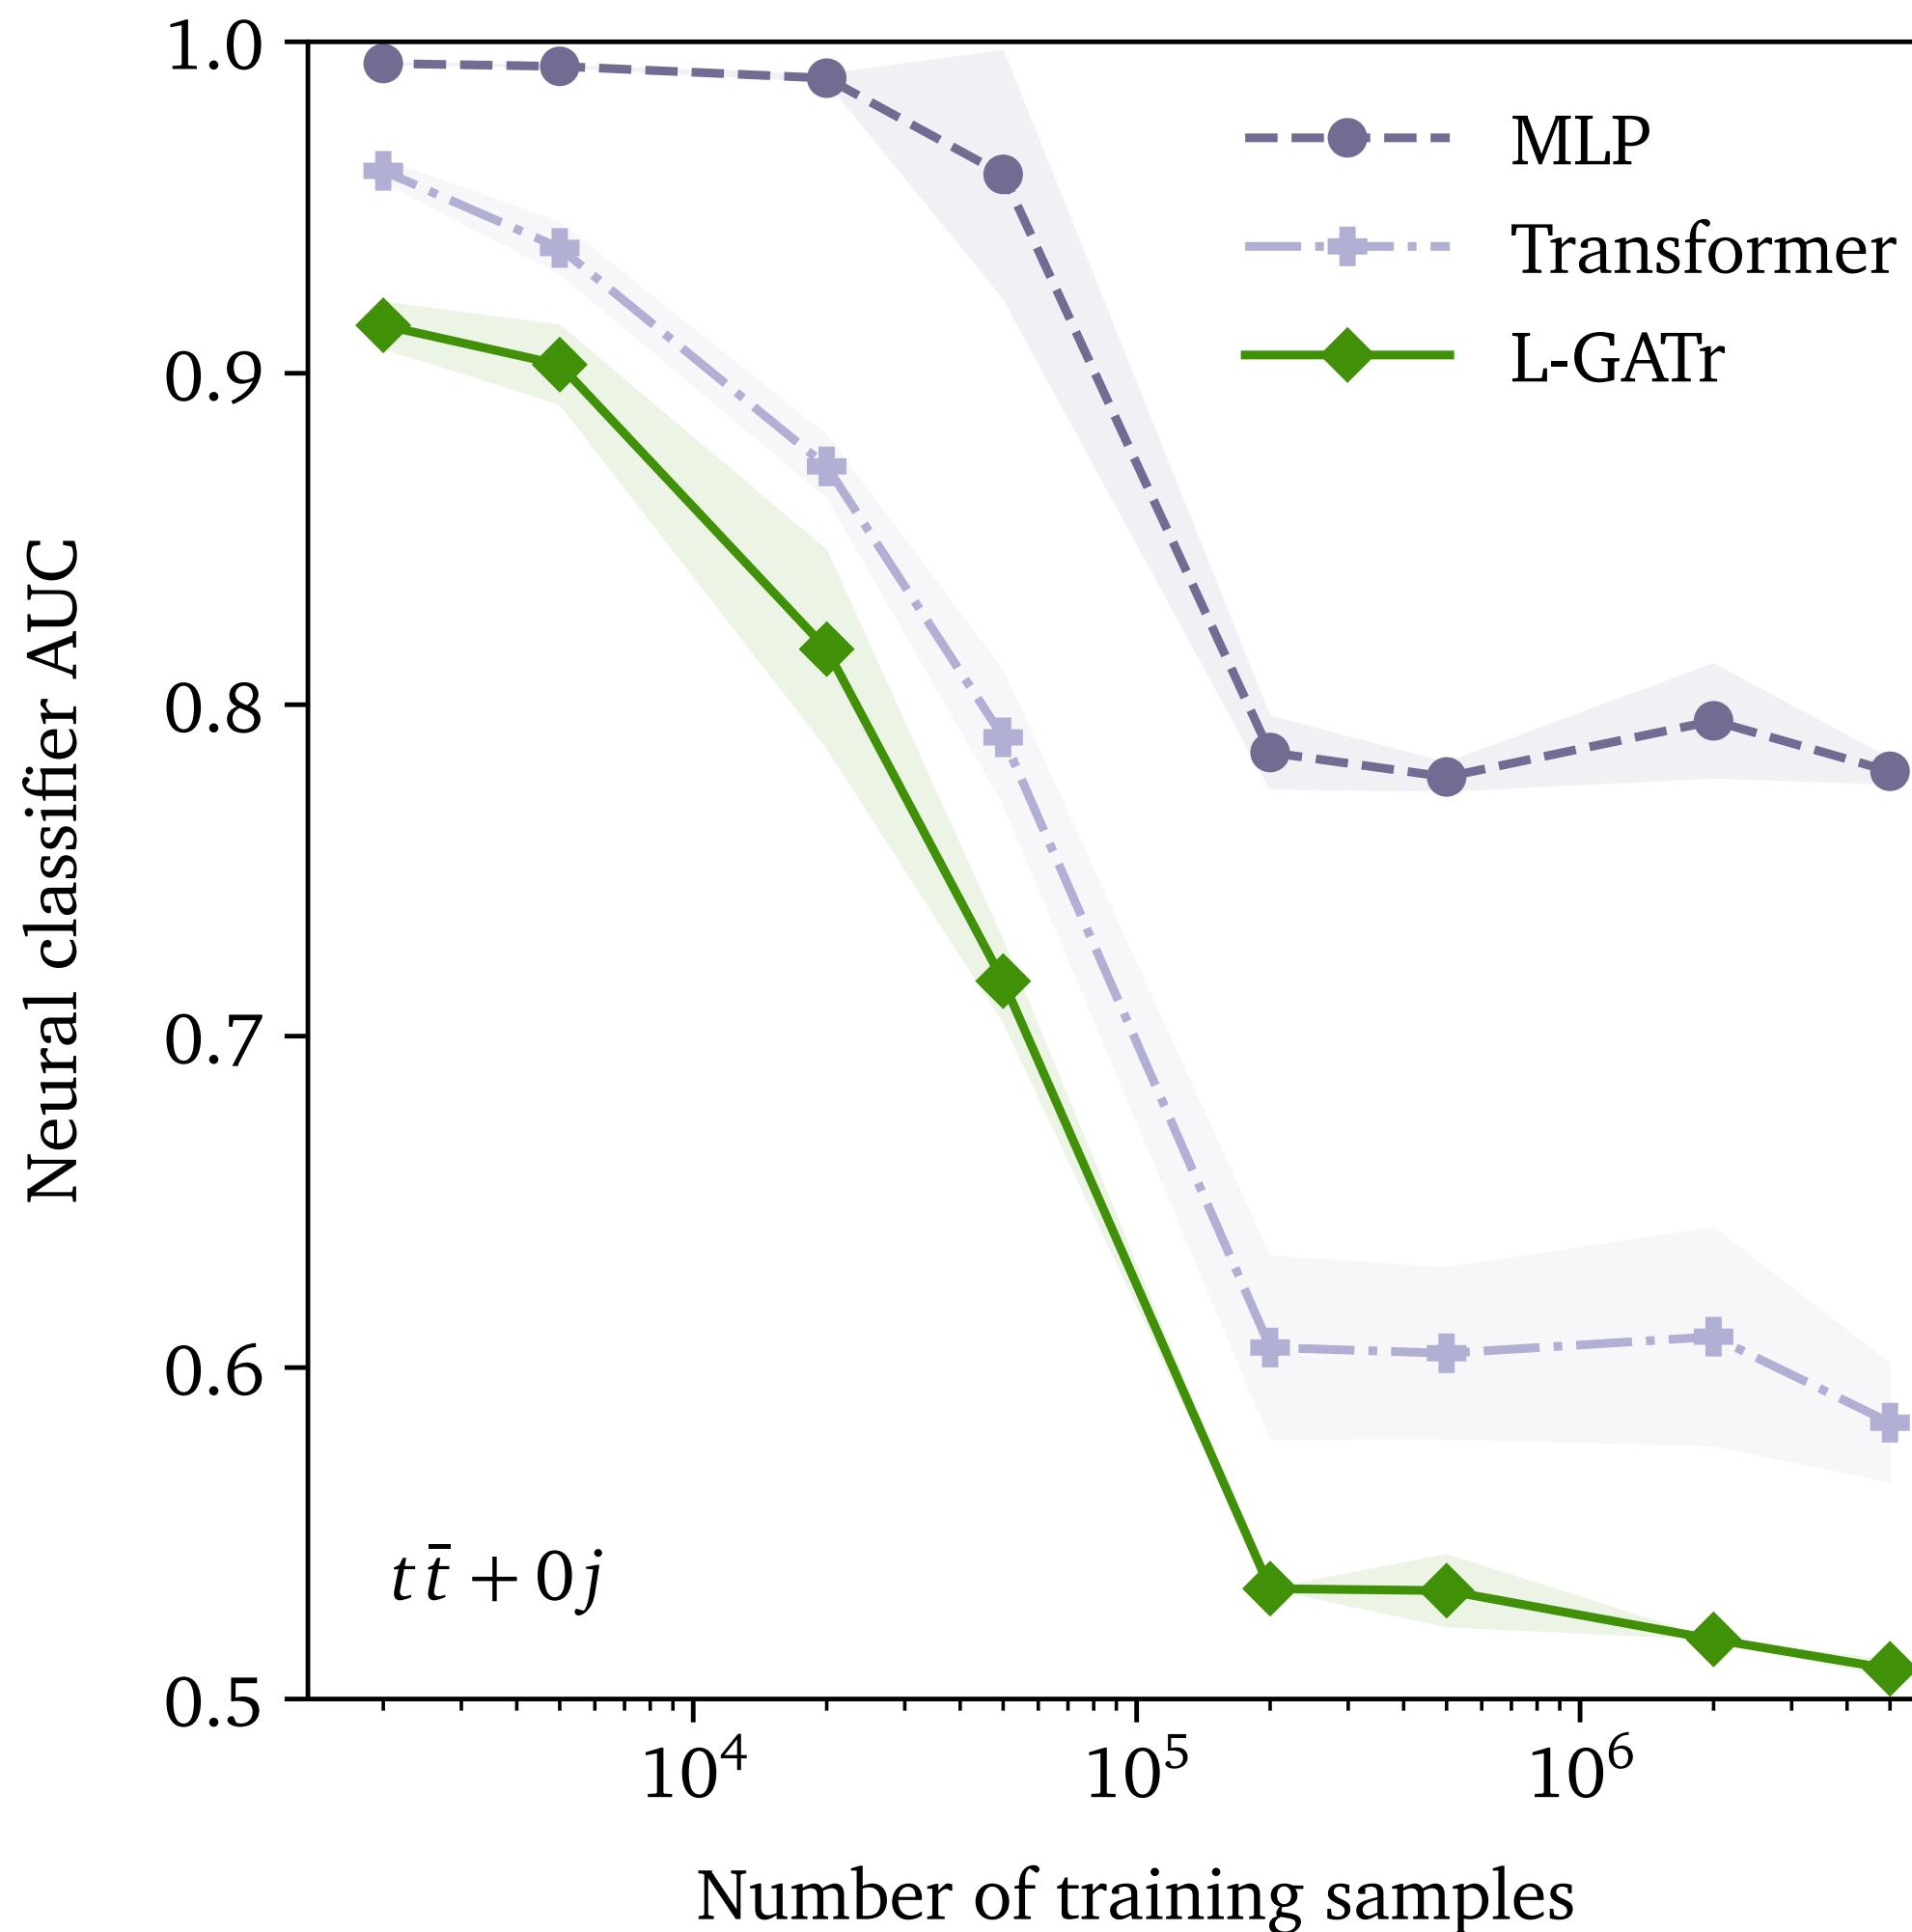
## LHC event generation



Equivariance helps,  
especially for angular correlations

# L-GATr beyond jet tagging

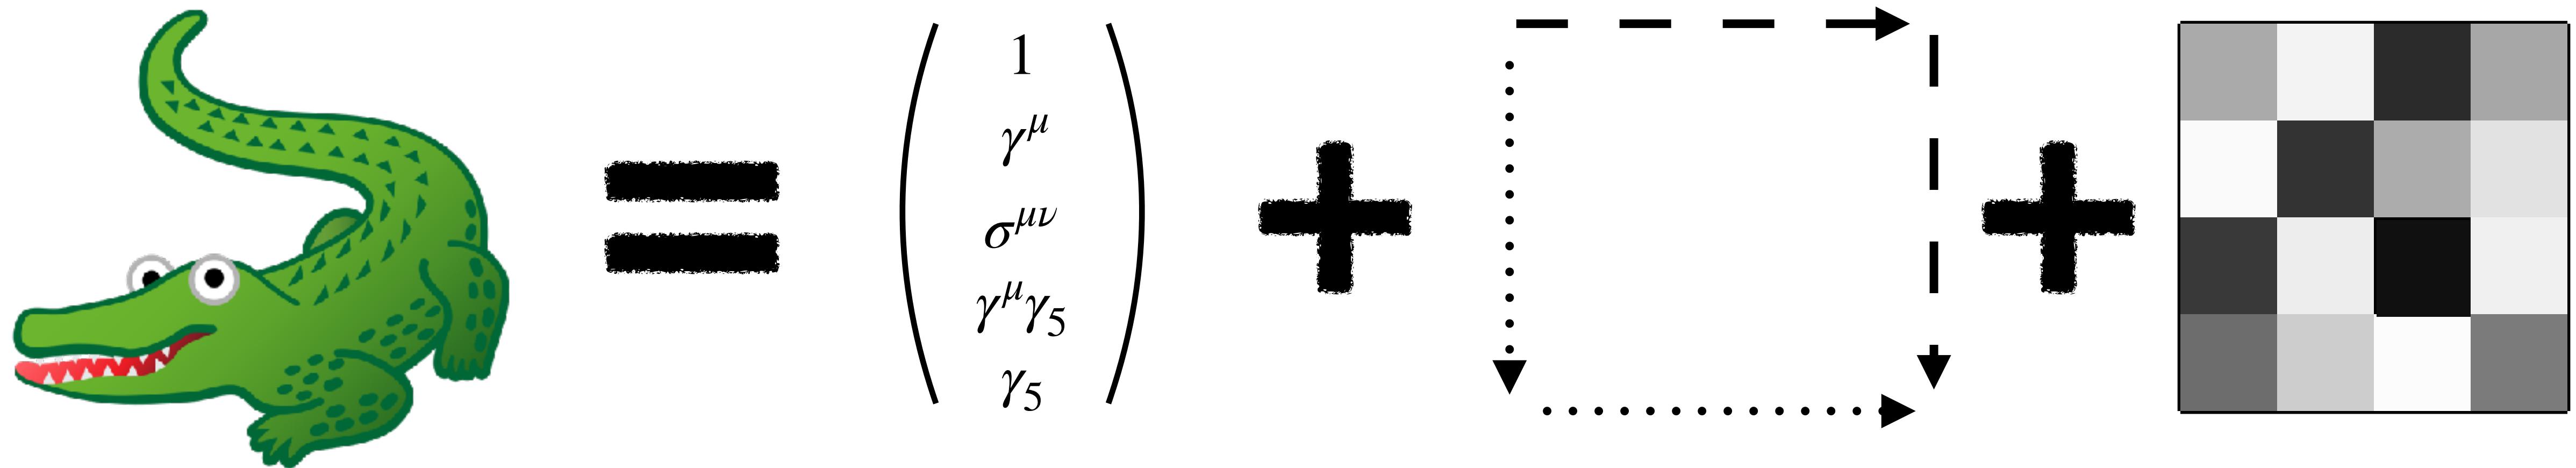
## LHC event generation



L-GATr generates samples that a classifier can barely distinguish from the ground truth

# Messages

- L-GATr is a multi-purpose architecture for LHC physics
- In jet tagging, L-GATr combines the benefits from equivariance + pretraining
- Including symmetry-breaking inputs is essential



Lorentz-Equivariant  
Geometric Algebra  
Transformer

**Geometric algebra**  
representations

**Lorentz-Equivariant**  
layers

**Transformer**  
architecture



# L-GATr is easy to use

```
from gatr import GATr, SelfAttentionConfig, MLPConfig
from gatr.interface import embed_vector, extract_scalar, embed_spurions
import torch

class ExampleWrapper(torch.nn.Module):
    """Example wrapper around a L-GATr model.

Parameters
-----
num_blocks : int
    Number of transformer blocks
hidden_mv_channels : int
    Number of hidden multivector channels
hidden_s_channels : int
    Number of hidden scalar channels
.....
def __init__(self, blocks=6, hidden_mv_channels=16, hidden_s_channels=32):
    super().__init__()
    self.gatr = GATr(
        in_mv_channels=3,
        out_mv_channels=1,
        hidden_mv_channels=hidden_mv_channels,
        in_s_channels=None,
        out_s_channels=None,
        hidden_s_channels=hidden_s_channels,
        num_blocks=num_blocks,
        attention=SelfAttentionConfig(), # Use default parameters for attention
        mlp=MLPConfig(), # Use default parameters for MLP
    )
```

```
def forward(self, fourmomenta):
    """Forward pass.

Parameters
-----
fourmomenta : torch.Tensor with shape (batchsize, num_points, 4)
    fourmomentum point cloud input data

Returns
-----
outputs : torch.Tensor with shape (batchsize, 1)
    Model prediction: a single scalar for the whole point cloud.
.....
batchsize, num_points, _ = fourmomenta.shape

# Embed fourmomentum point cloud inputs in GA
multivectors = embed_vector(fourmomenta).unsqueeze(-2) # (batchsize, num_points, 1, 1)

# Append spurions for symmetry breaking (optional)
spurions = embed_spurions(bean_reference="xyplane", add_time_reference=True, device=fo
spurions = spurions[None, None, ...].repeat(batchsize, num_points, 1, 1) # (batchsize
multivectors = torch.cat((multivectors, spurions), dim=-2) # (batchsize, num_points,

# Pass data through GATr
multivector_outputs, _ = self.gatr(multivectors, scalars=None) # (batchsize, num_poin

# Extract scalar outputs
outputs = extract_scalar(multivector) # (batchsize, num_points, 1)

# Mean aggregation to extract a single scalar for the whole point cloud
score = outputs.mean(dim=1)

return score
```



Victor Bresó



Pim de Haan



Tilman Plehn



Huilin Qu



Jesse Thaler



Johann Brehmer

## Lorentz-Equivariant Geometric Algebra Transformer for High-Energy Physics

Jonas Spinner\*, Víctor Bresó\*, Pim de Haan,  
Tilman Plehn, Jesse Thaler, Johann Brehmer  
NeurIPS 2024, arXiv:2405.14806



CS paper



HEP paper

## A Lorentz-Equivariant Transformer for all of the LHC

Johann Brehmer, Víctor Bresó,  
Pim de Haan, Tilman Plehn,  
Huilin Qu, Jonas Spinner, Jesse Thaler  
arXiv:2411.00446



L-GATr code

For what will **you** use L-GATr?

# Bonus material

Network	Accuracy	AUC	$1/\epsilon_B$ ( $\epsilon_S = 0.5$ )	$1/\epsilon_B$ ( $\epsilon_S = 0.3$ )
TopoDNN [54]	0.916	0.972	–	$295 \pm 5$
LoLa [9]	0.929	0.980	–	$722 \pm 17$
<i>N</i> -subjettiness [55]	0.929	0.981	–	$867 \pm 15$
PFN [56]	0.932	0.9819	$247 \pm 3$	$888 \pm 17$
TreeNiN [57]	0.933	0.982	–	$1025 \pm 11$
ParticleNet [58]	0.940	0.9858	$397 \pm 7$	$1615 \pm 93$
ParT [59]	0.940	0.9858	$413 \pm 16$	$1602 \pm 81$
MIParT [60]	0.942	0.9868	$505 \pm 8$	$2010 \pm 97$
LorentzNet* [10]	0.942	0.9868	$498 \pm 18$	$2195 \pm 173$
CGENN* [14]	0.942	0.9869	500	2172
PELICAN* [42]	$0.9426 \pm 0.0002$	$0.9870 \pm 0.0001$	–	$2250 \pm 75$
L-GATr* [35]	$0.9423 \pm 0.0002$	$0.9870 \pm 0.0001$	$540 \pm 20$	$2240 \pm 70$
ParticleNet-f.t. [60]	0.942	0.9866	$487 \pm 9$	$1771 \pm 80$
ParT-f.t. [60]	0.944	0.9877	$691 \pm 15$	$2766 \pm 130$
MIParT-f.t. [60]	0.944	0.9878	$640 \pm 10$	$2789 \pm 133$
L-GATr-f.t.* (new)	$0.9442 \pm 0.0002$	$0.98792 \pm 0.00004$	$661 \pm 24$	$3005 \pm 186$

# Event generation

## Conditional Flow Matching

Continuous normalising flow (CNF)

connect a simple base density  
to a complex target density  
through a neural differential equation

$$\frac{d}{dt}x = v_t(x)$$

Continuous normalising flows  
arXiv:1806.07366

Conditional flow matching (CFM)

is a simple way to train CNFs  
by comparing the learned velocity  $v_t(x)$   
to a conditional target velocity  $u_t(x | x_1)$

$$\mathcal{L} = \left\langle (v_t(x) - u_t(x | x_1))^2 \right\rangle$$

Conditional flow matching  
arXiv:2210.02747

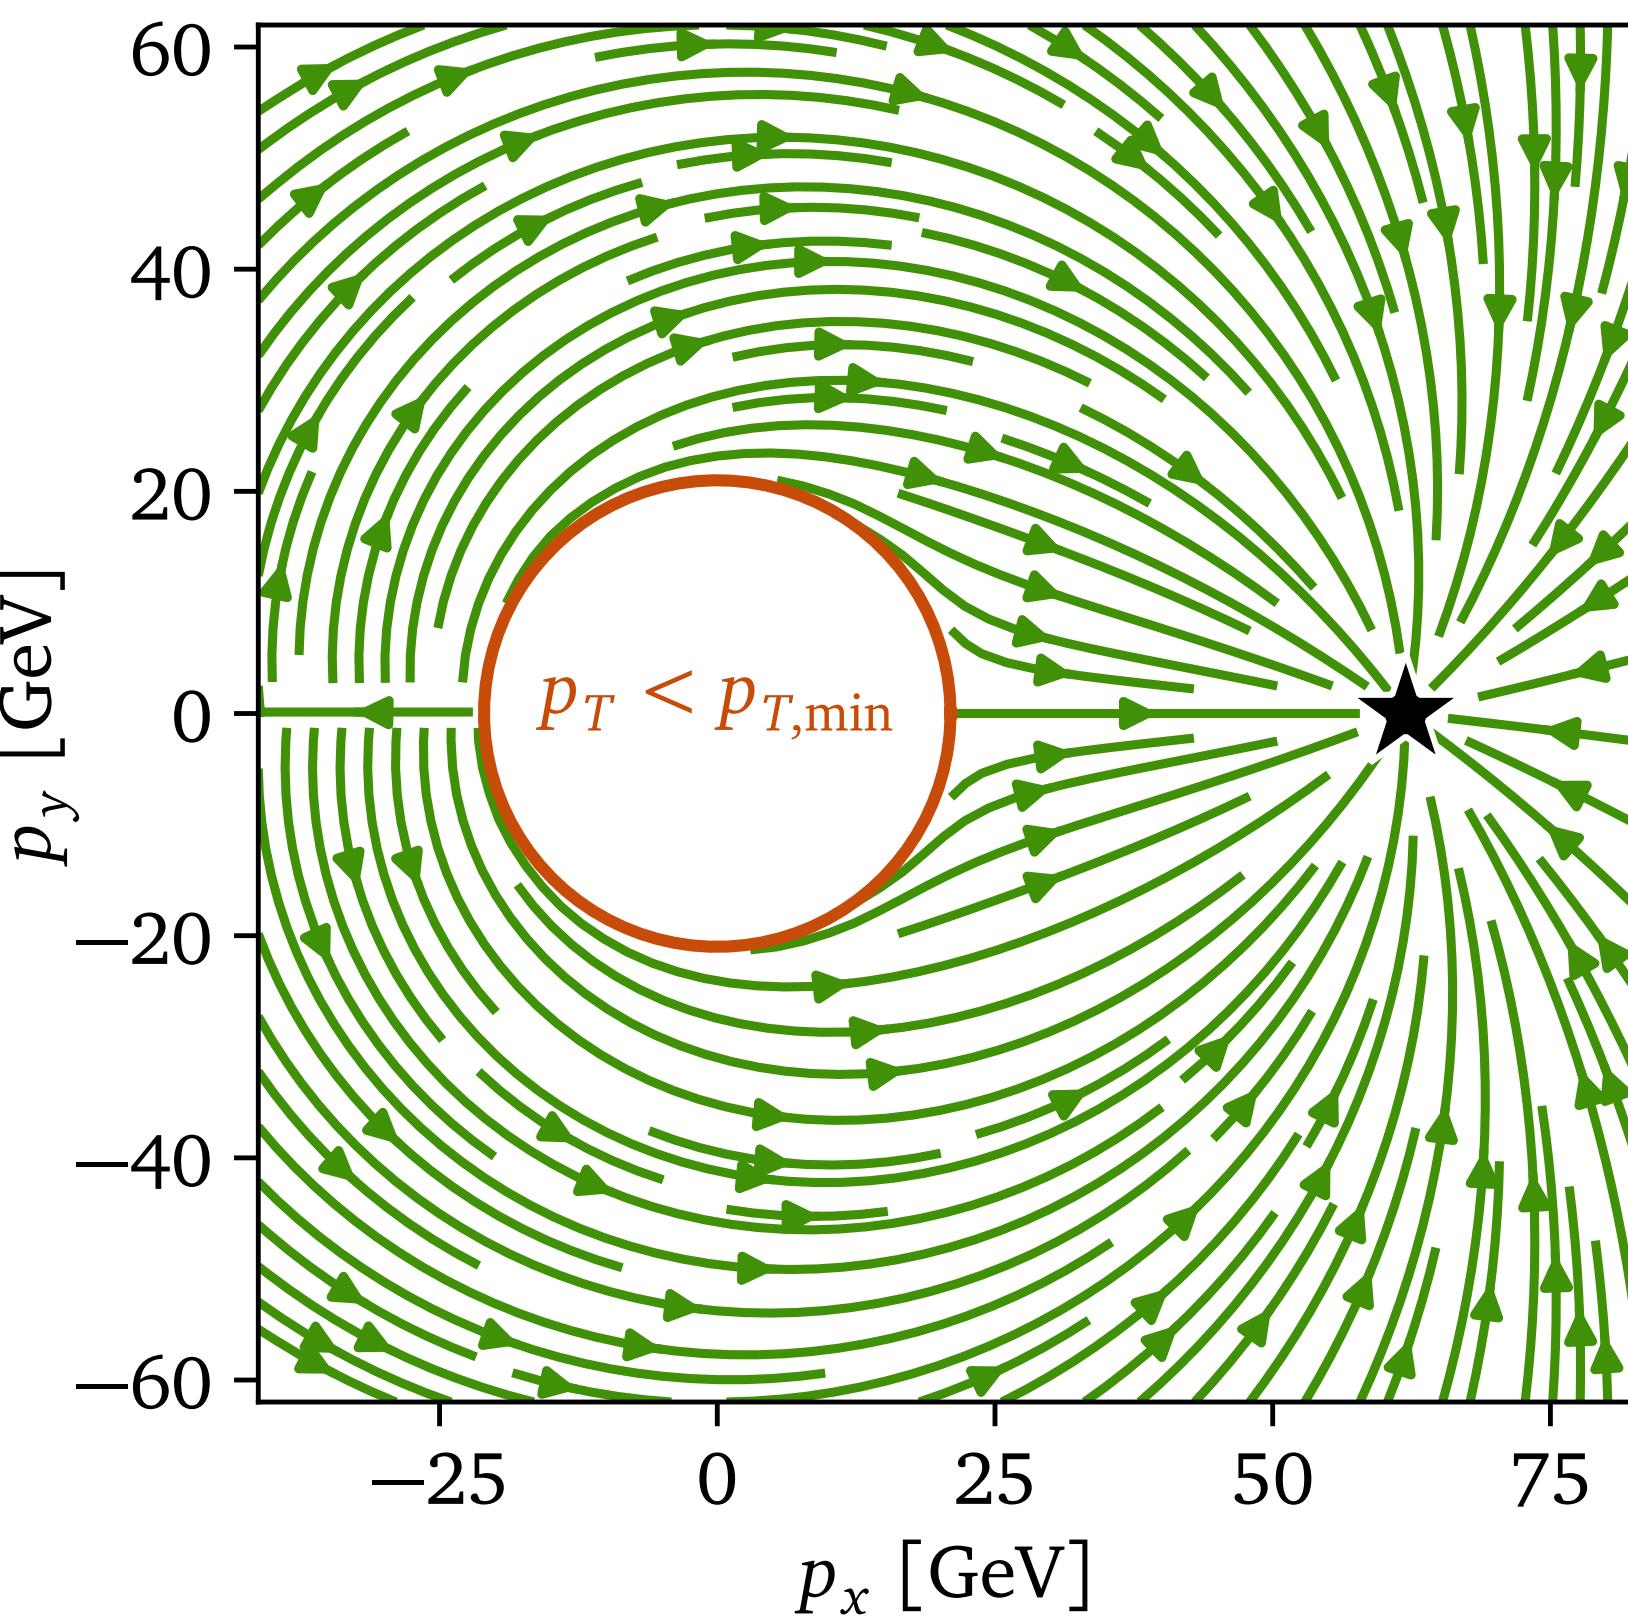
How to pick the target velocity  $u_t(x | x_1)$ ?

# Event generation

## Physics-inspired target trajectories

Straight trajectories in ‘modified jet momenta’  $x$ :

$$p = \begin{pmatrix} E \\ p_x \\ p_y \\ p_z \end{pmatrix} \rightarrow f^{-1}(p) = x = \begin{pmatrix} x_p \\ x_m \\ x_\eta \\ x_\phi \end{pmatrix} = \begin{pmatrix} \log(p_T - p_{T,\min}) \\ \log m^2 \\ \eta \\ \phi \end{pmatrix}$$

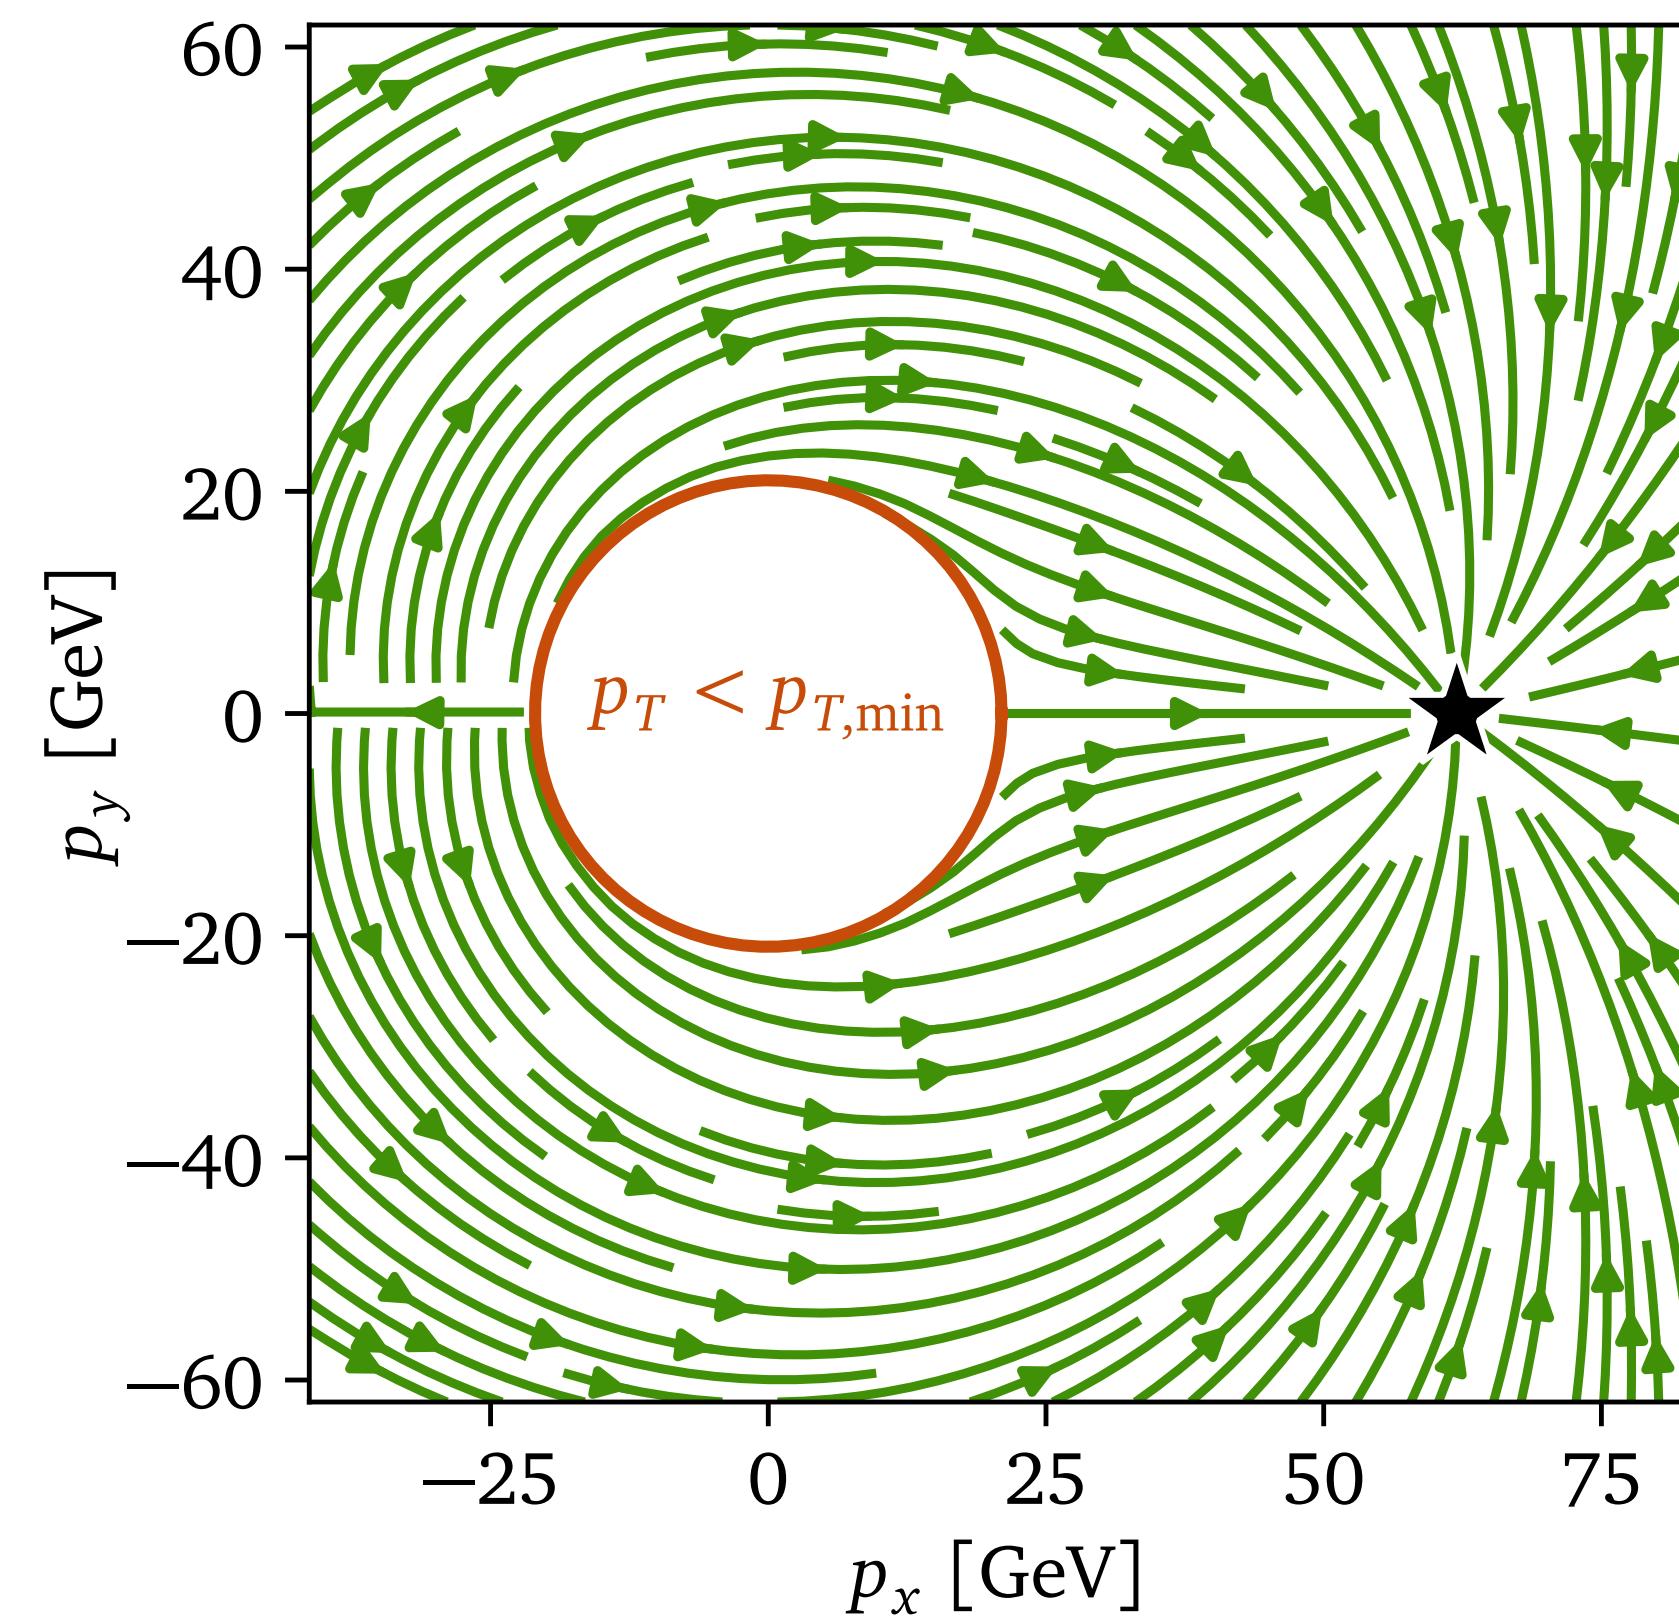


# Event generation

## Physics-inspired target trajectories

Straight trajectories in ‘modified jet momenta’  $x$ :

$$p = \begin{pmatrix} E \\ p_x \\ p_y \\ p_z \end{pmatrix} \rightarrow f^{-1}(p) = x = \begin{pmatrix} x_p \\ x_m \\ x_\eta \\ x_\phi \end{pmatrix} = \begin{pmatrix} \log(p_T - p_{T,\min}) \\ \log m^2 \\ \eta \\ \phi \end{pmatrix}$$



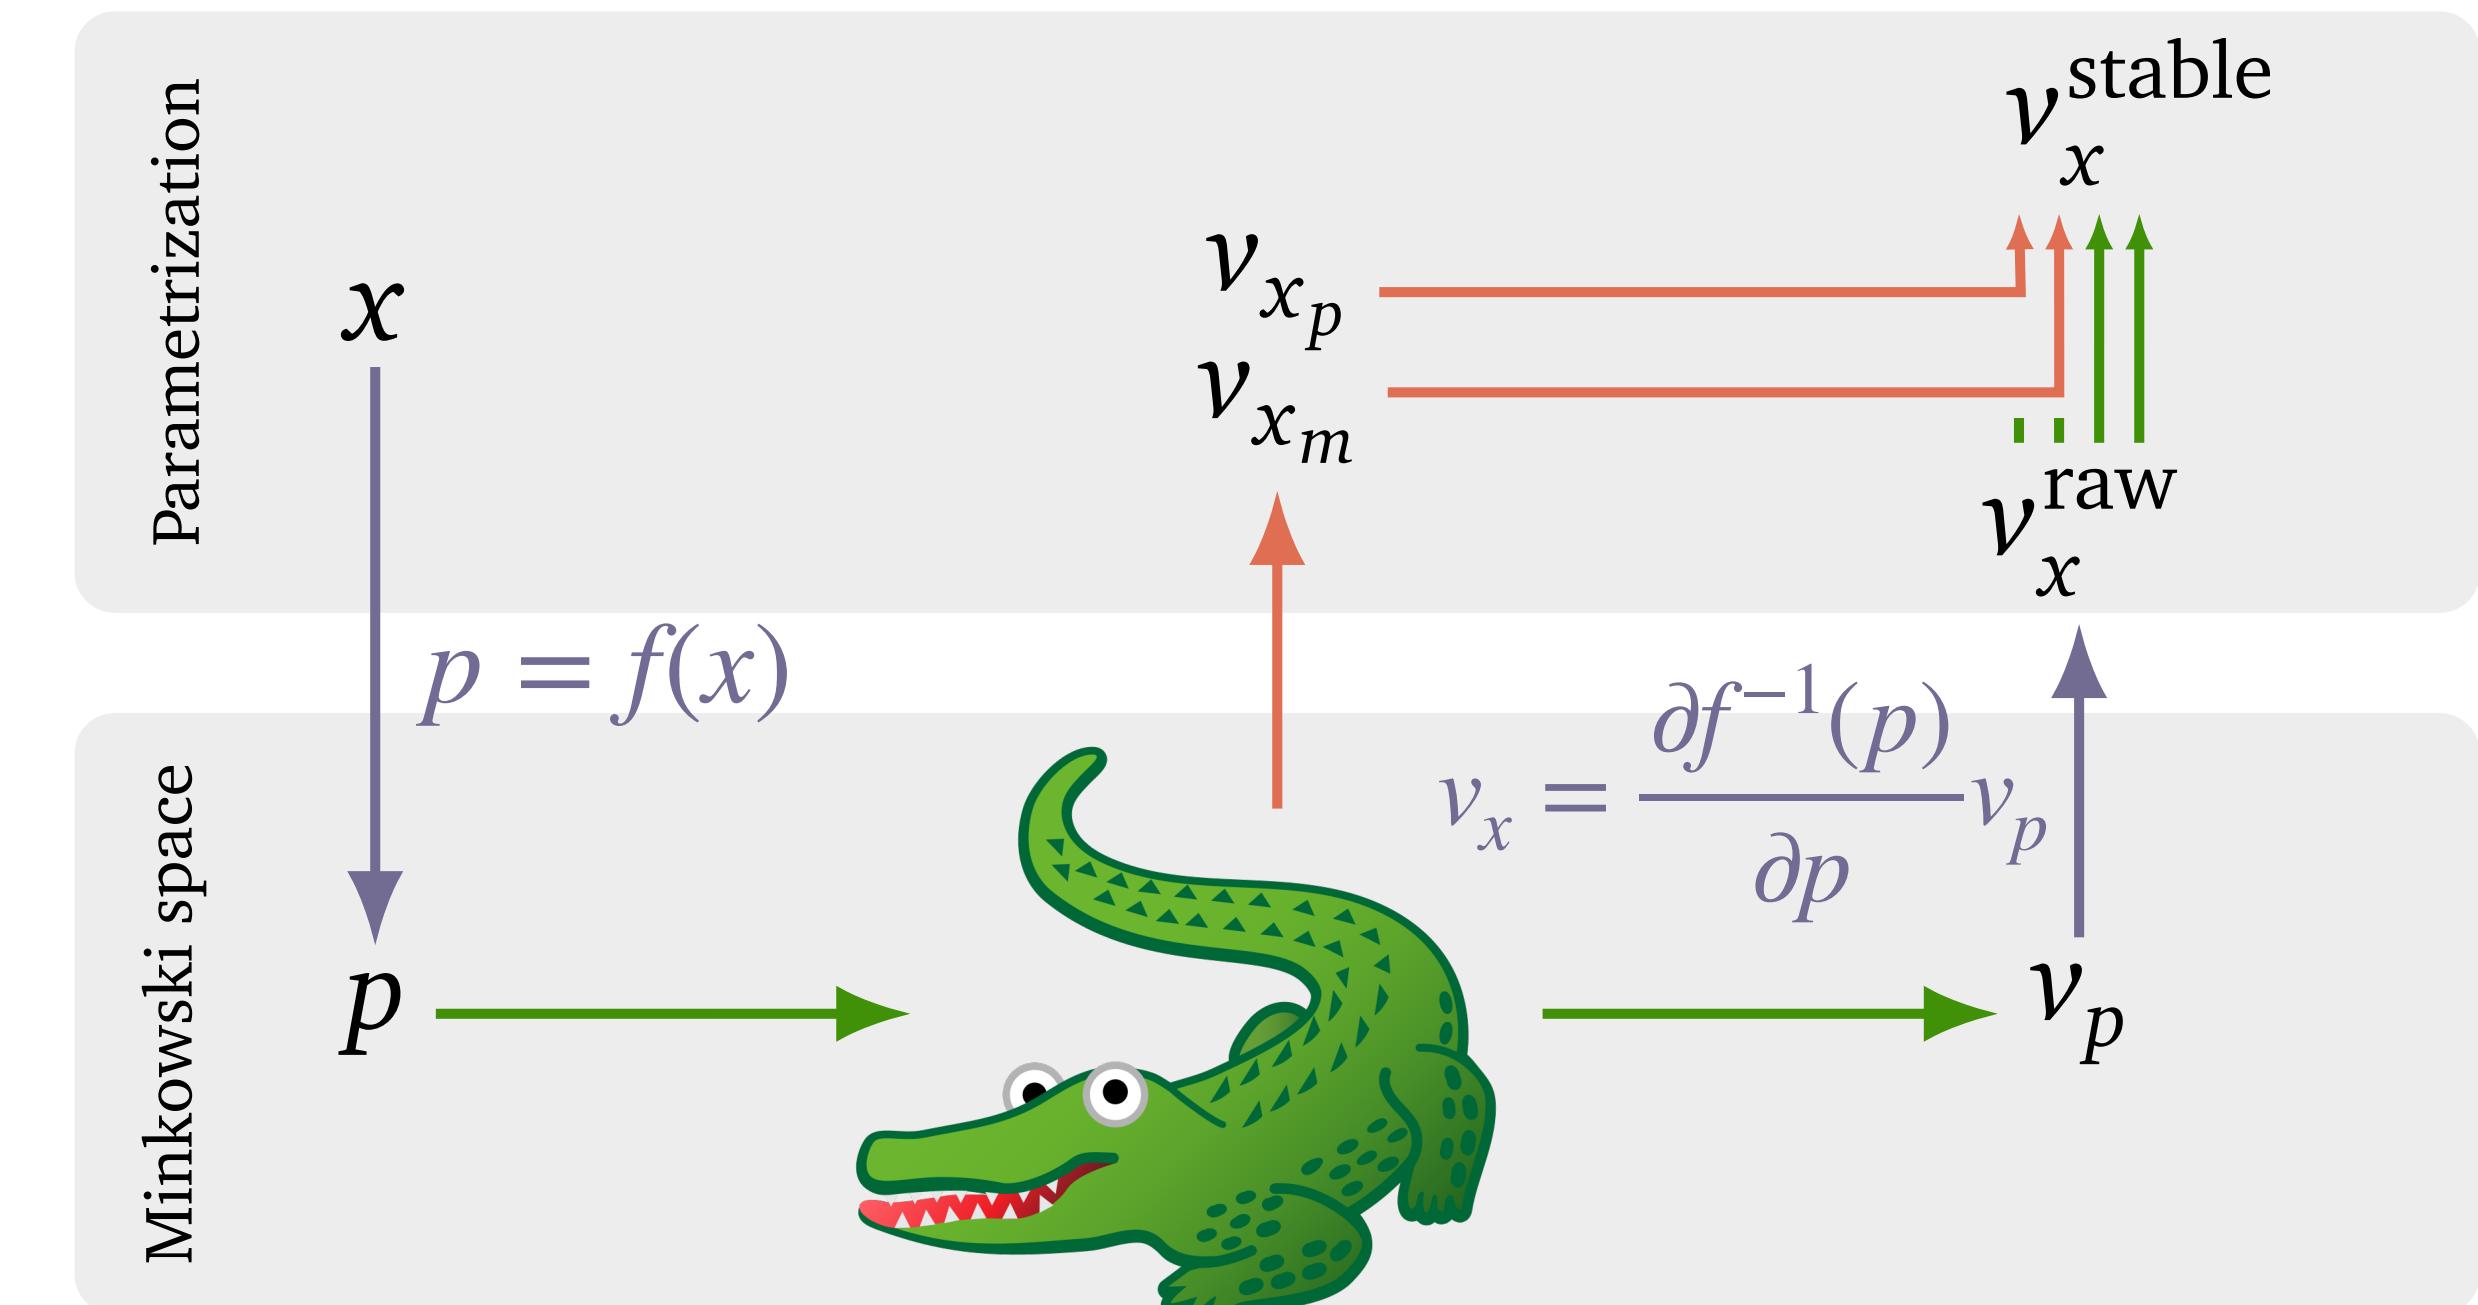
Data	Architecture	Base distribution	Periodic	Neg. log-likelihood	AUC	
$p$	L-GATr	rejection sampling	✓	$-30.80 \pm 0.17$	$0.945 \pm 0.004$	
$x$	MLP	rejection sampling	✓	$-32.13 \pm 0.05$	$0.780 \pm 0.003$	
$x$	L-GATr	rejection sampling	✗	$-32.57 \pm 0.05$	$0.530 \pm 0.017$	
$x$	L-GATr	no rejection sampling	✓	$-32.58 \pm 0.04$	$0.523 \pm 0.014$	
(default)	$x$	L-GATr	rejection sampling	✓	$-32.65 \pm 0.04$	$0.515 \pm 0.009$

# Event generation

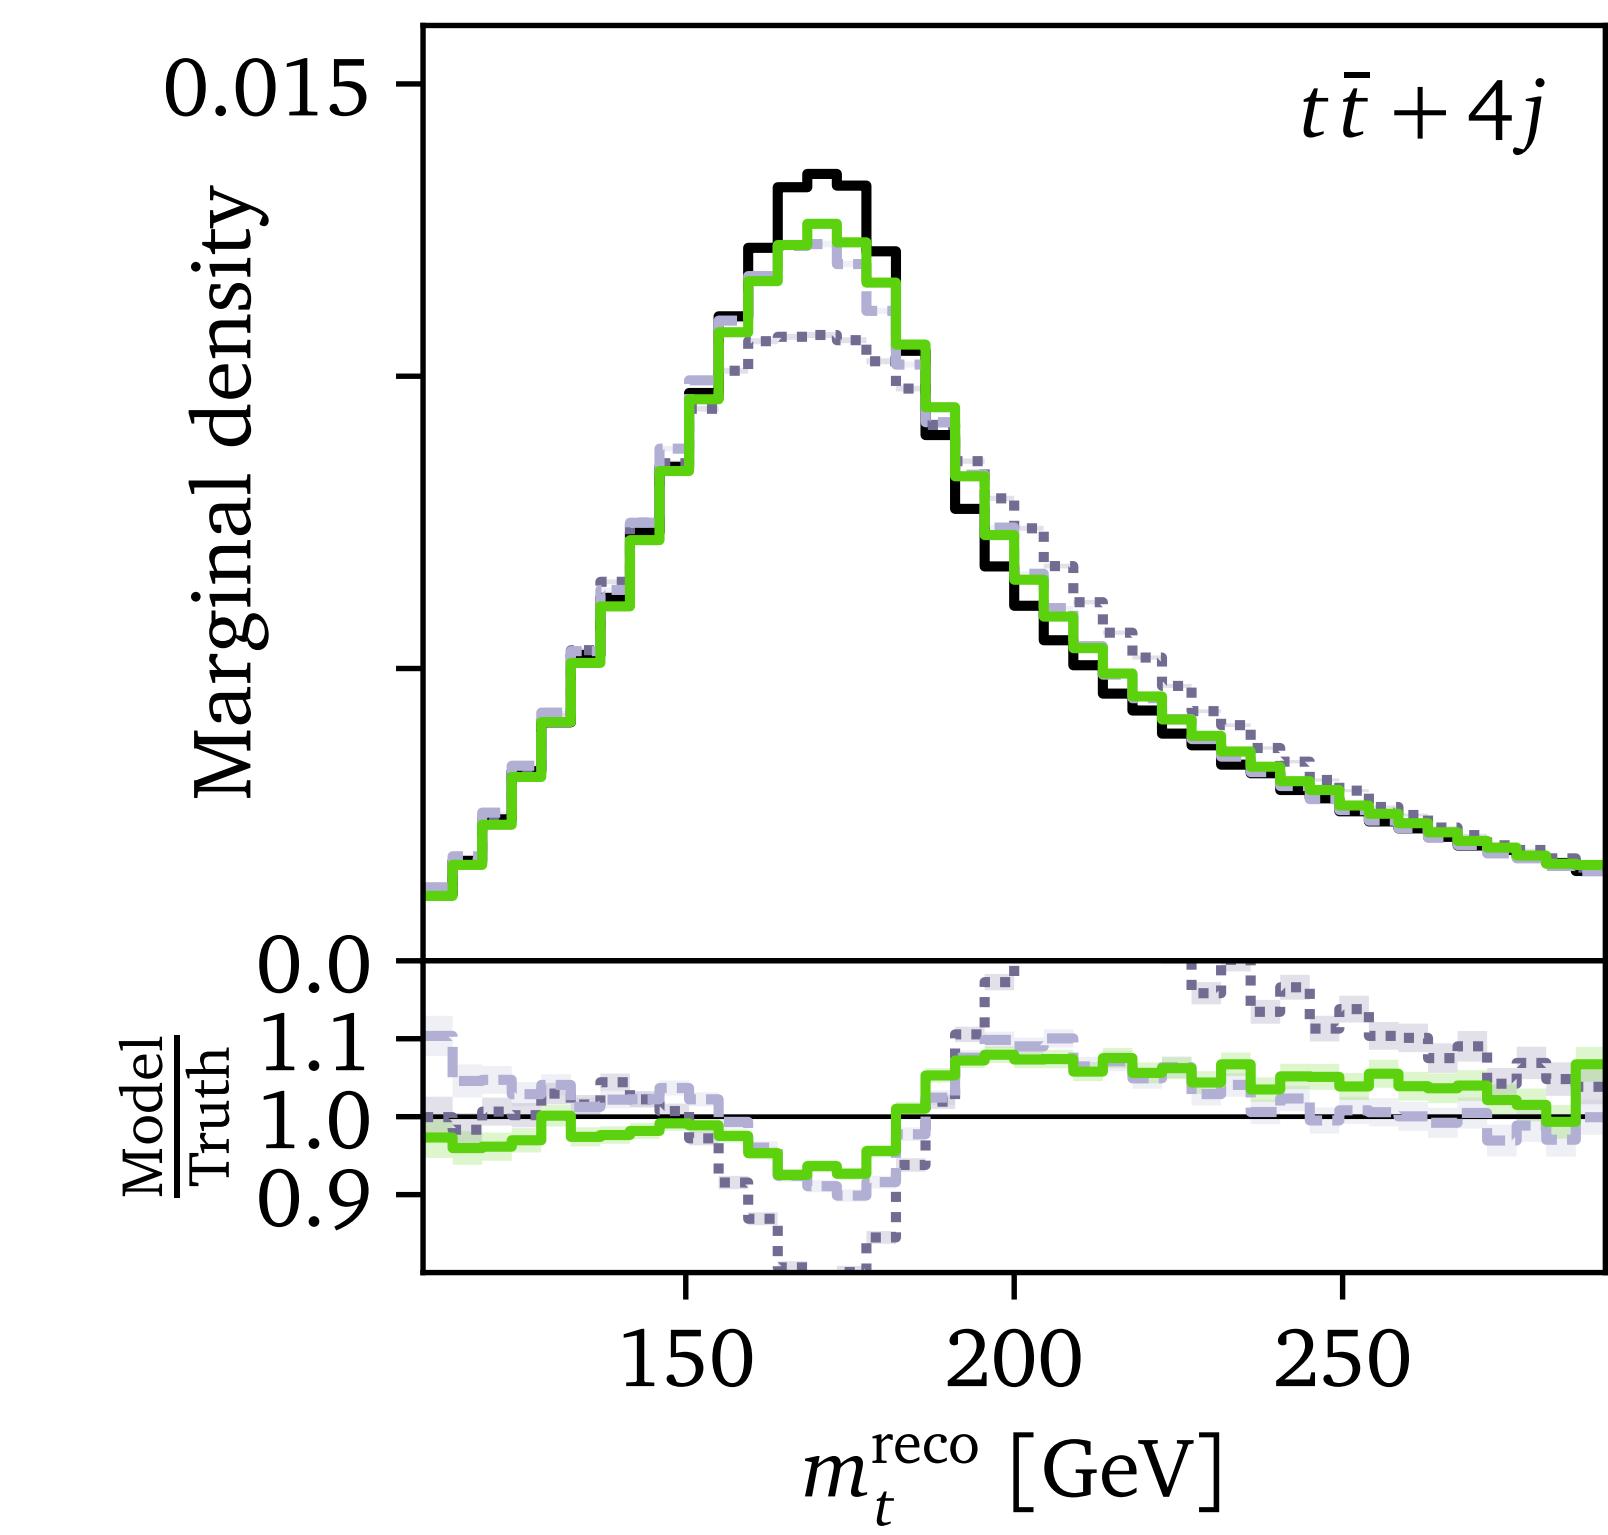
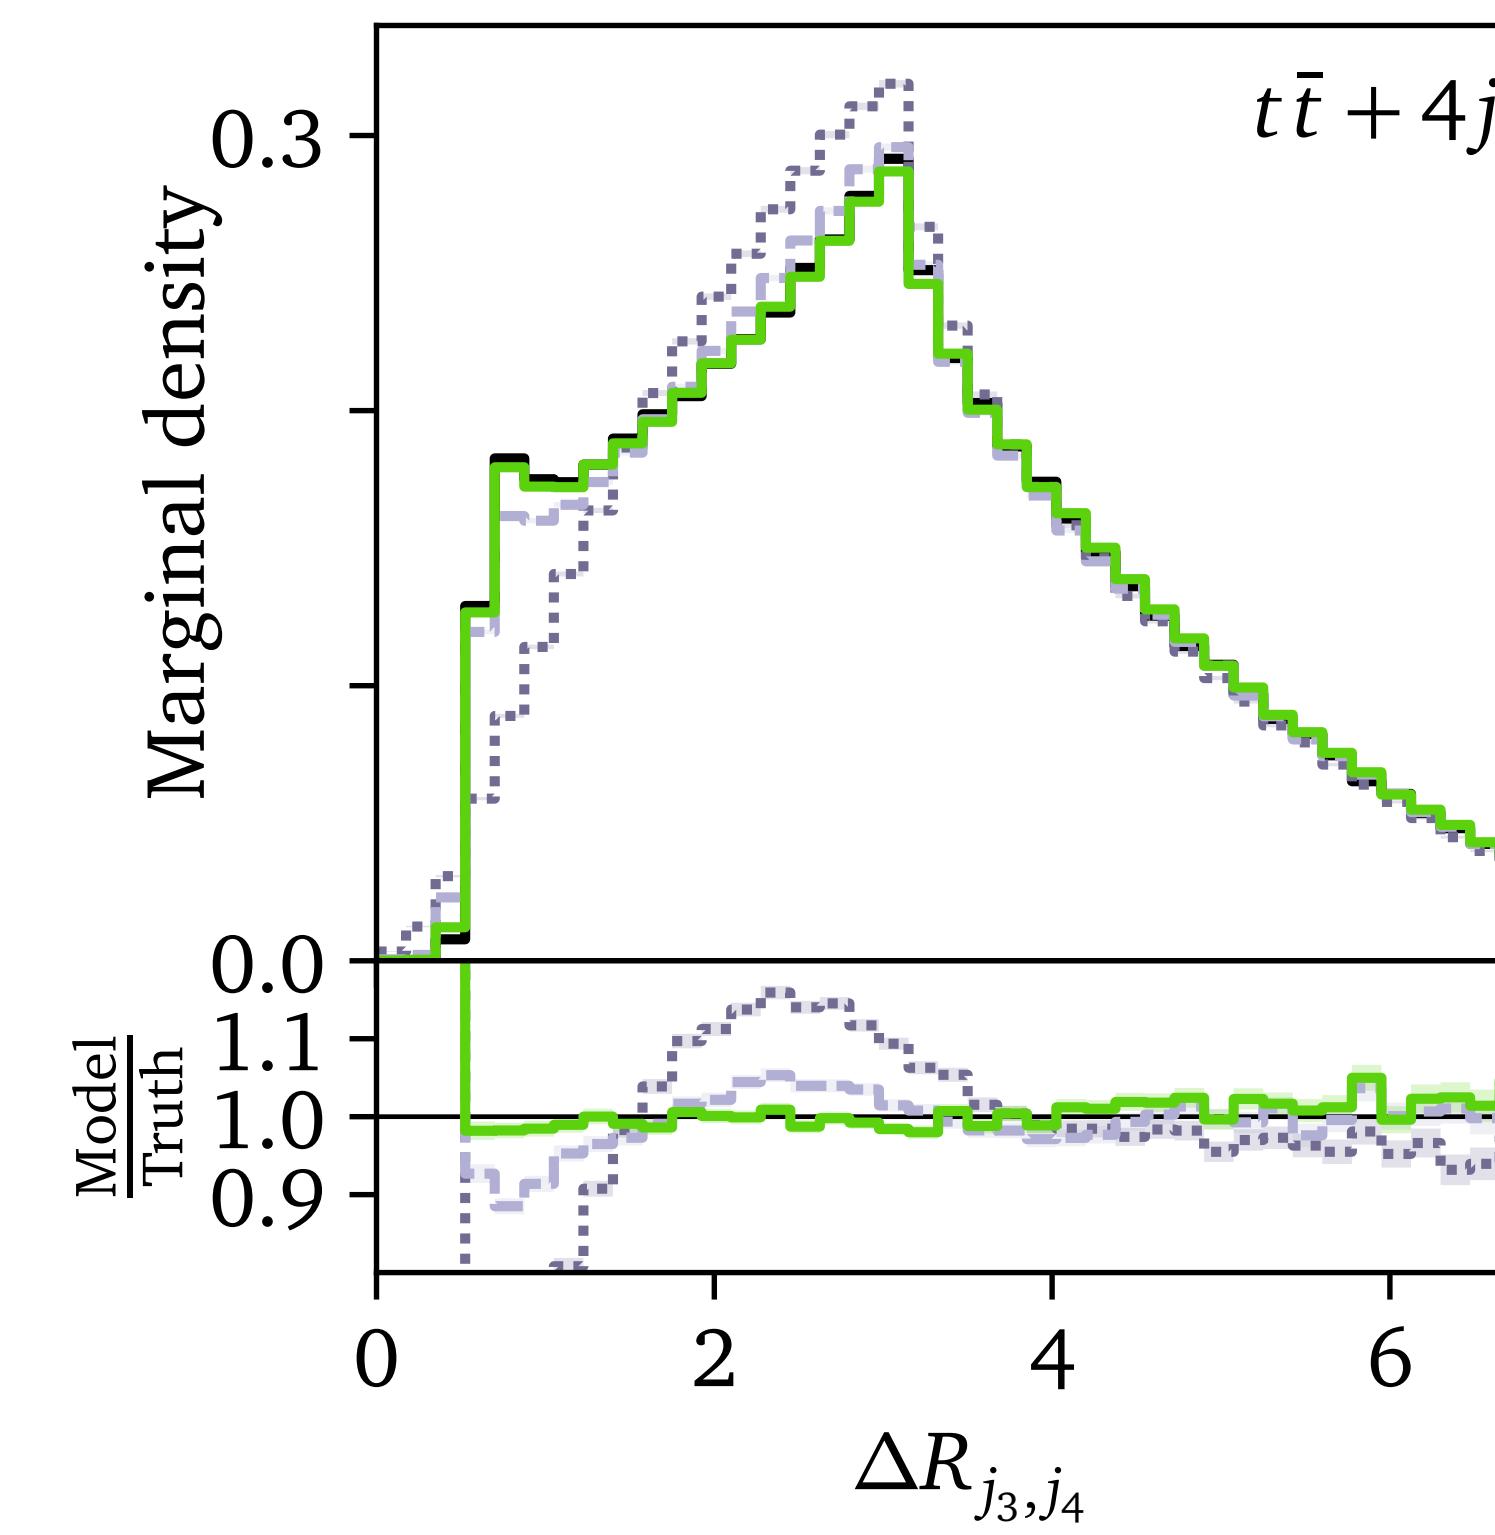
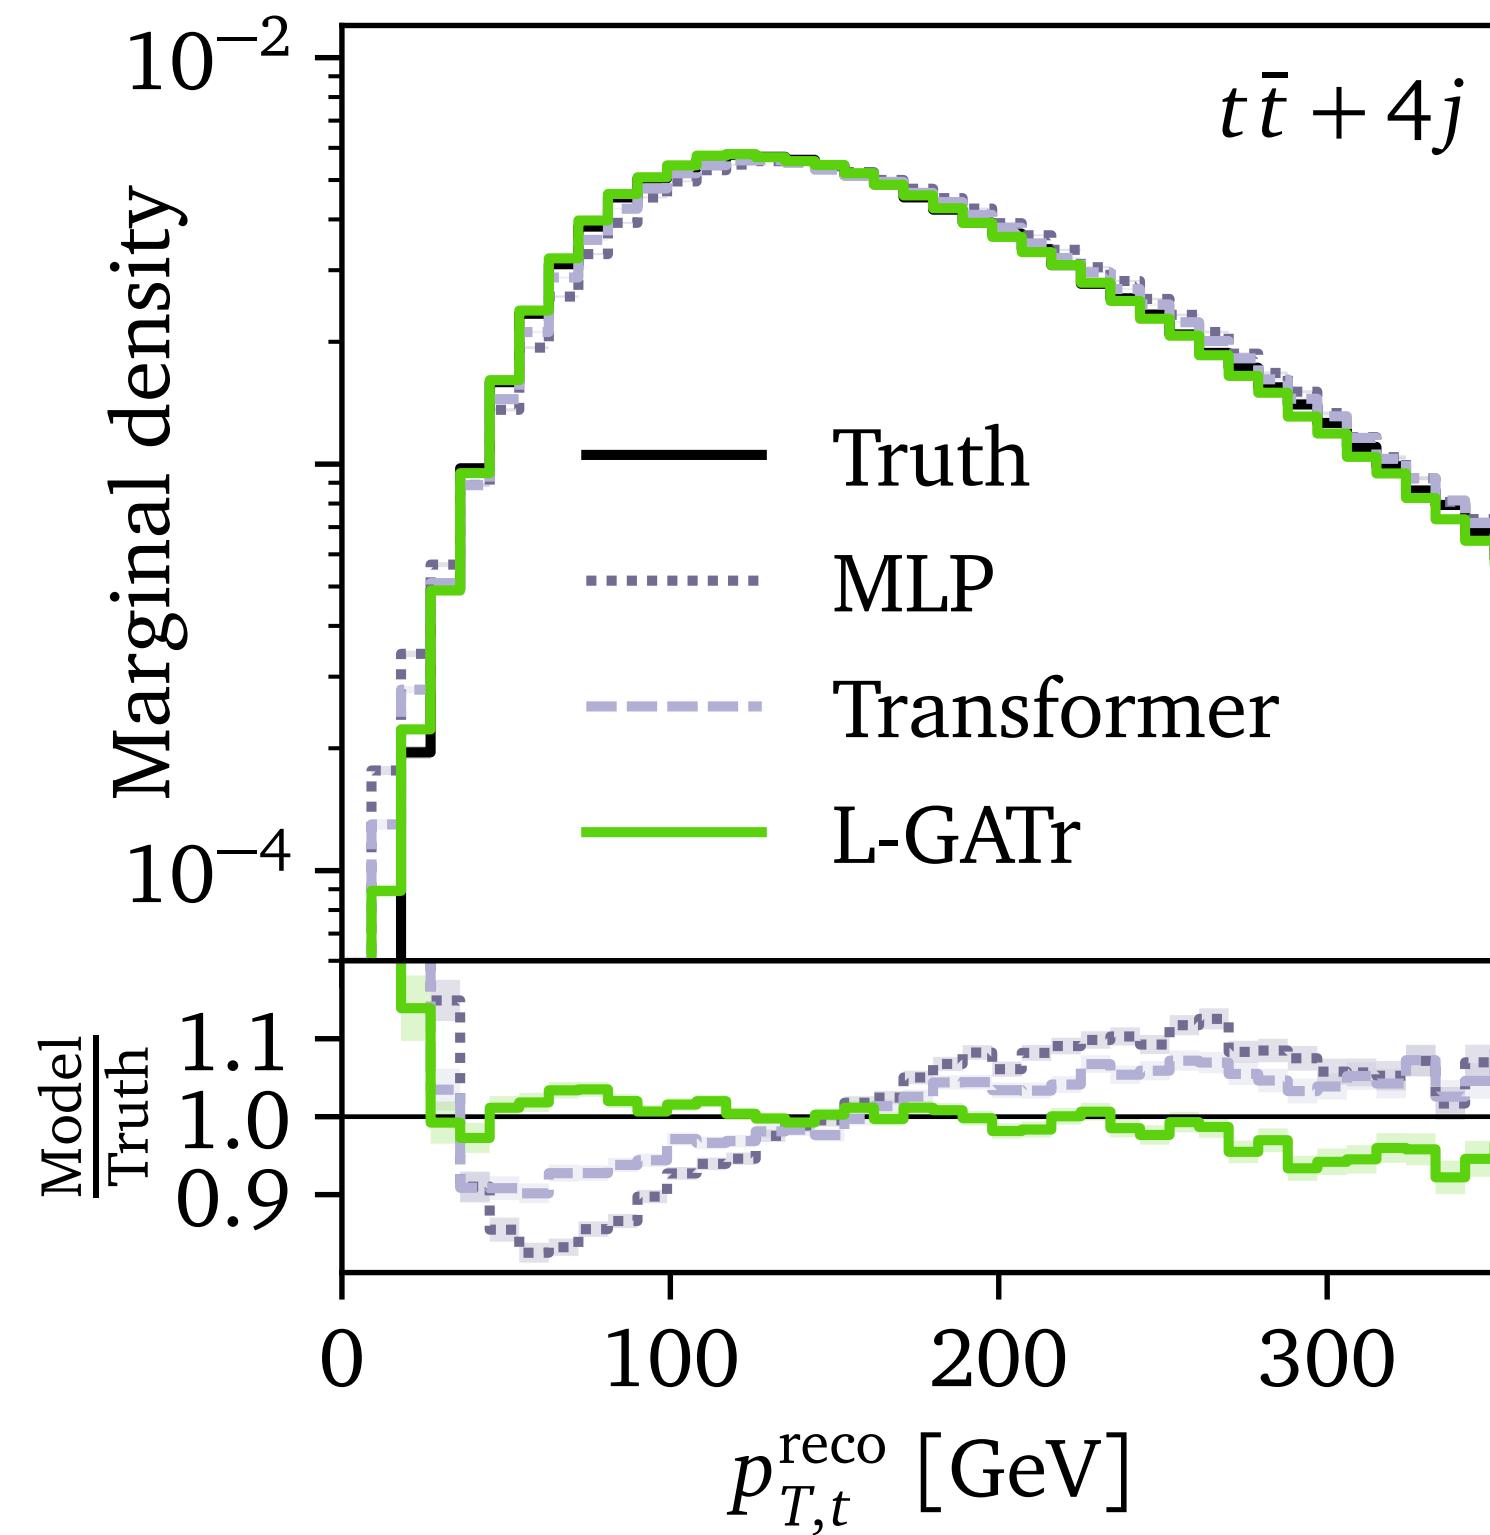
How to extract the CFM velocity field  $v_x(x)$ ?

Extend standard CFM workflow  
with L-GATr:

- Transformations  $f(x)$  between Minkowski space  $p$  and the parametrization  $x$
- Equivariant operations using multivectors
- Symmetry-breaking operations using scalars (required for numerical stability)

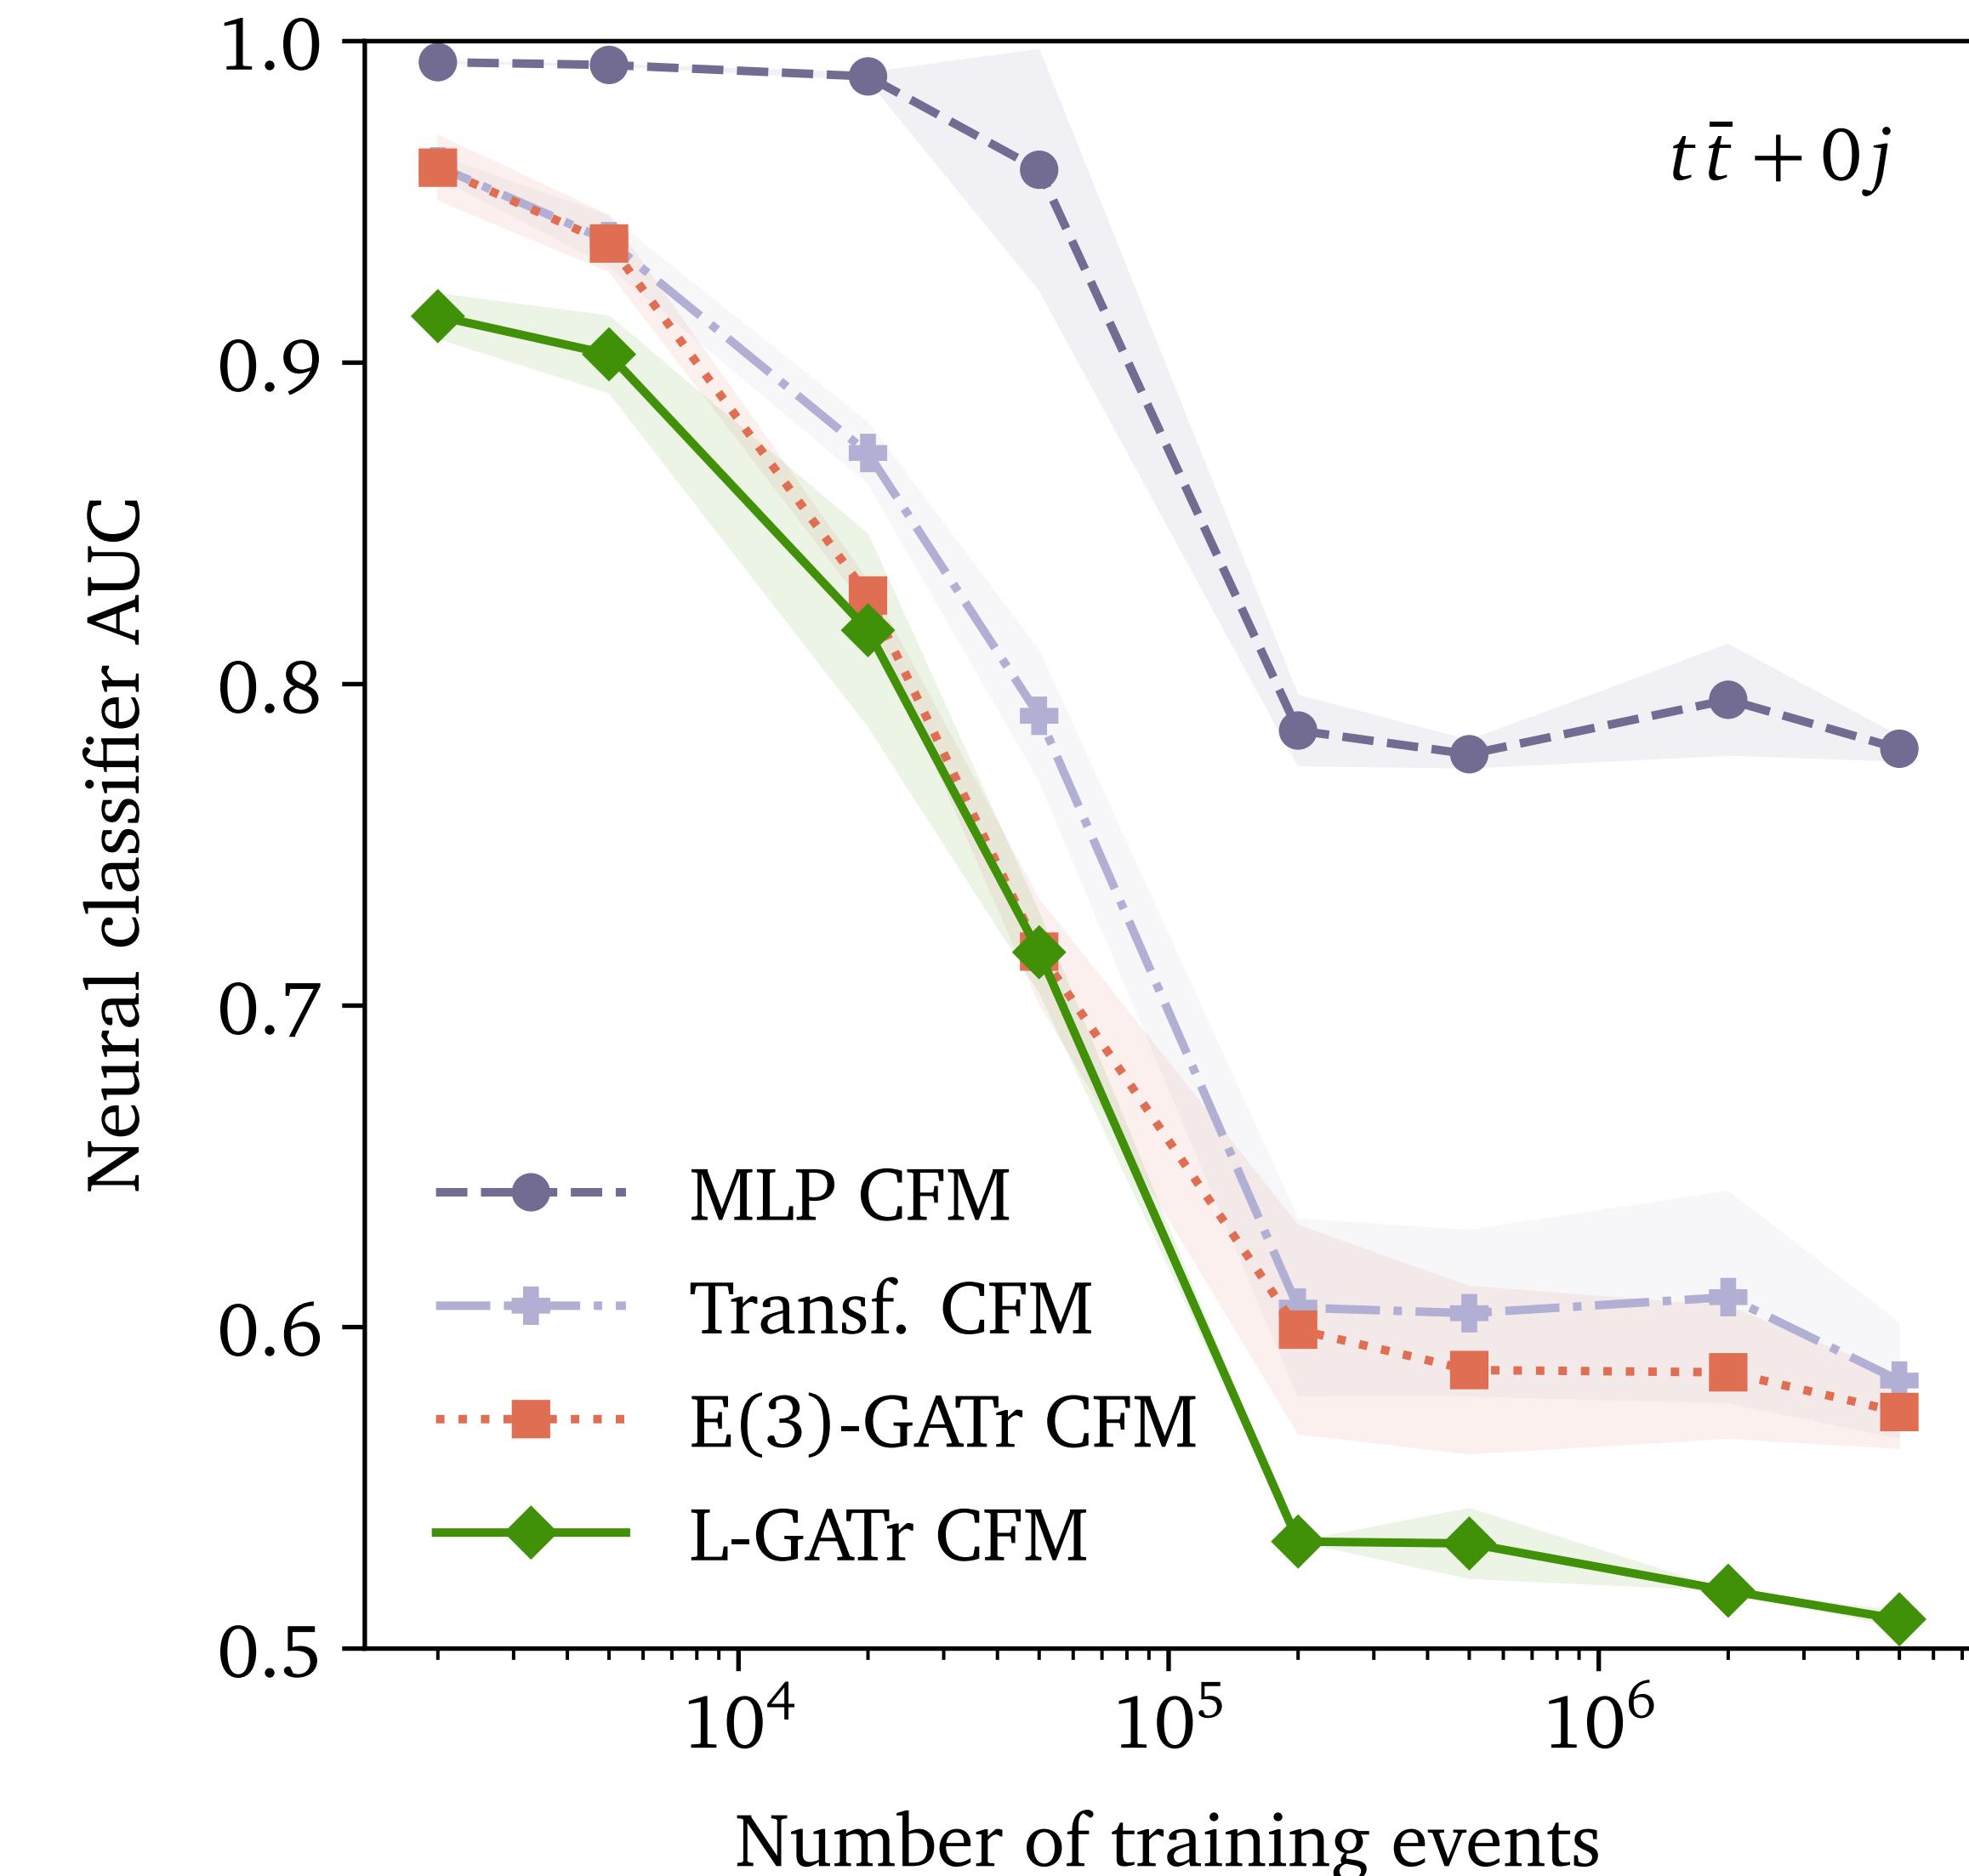
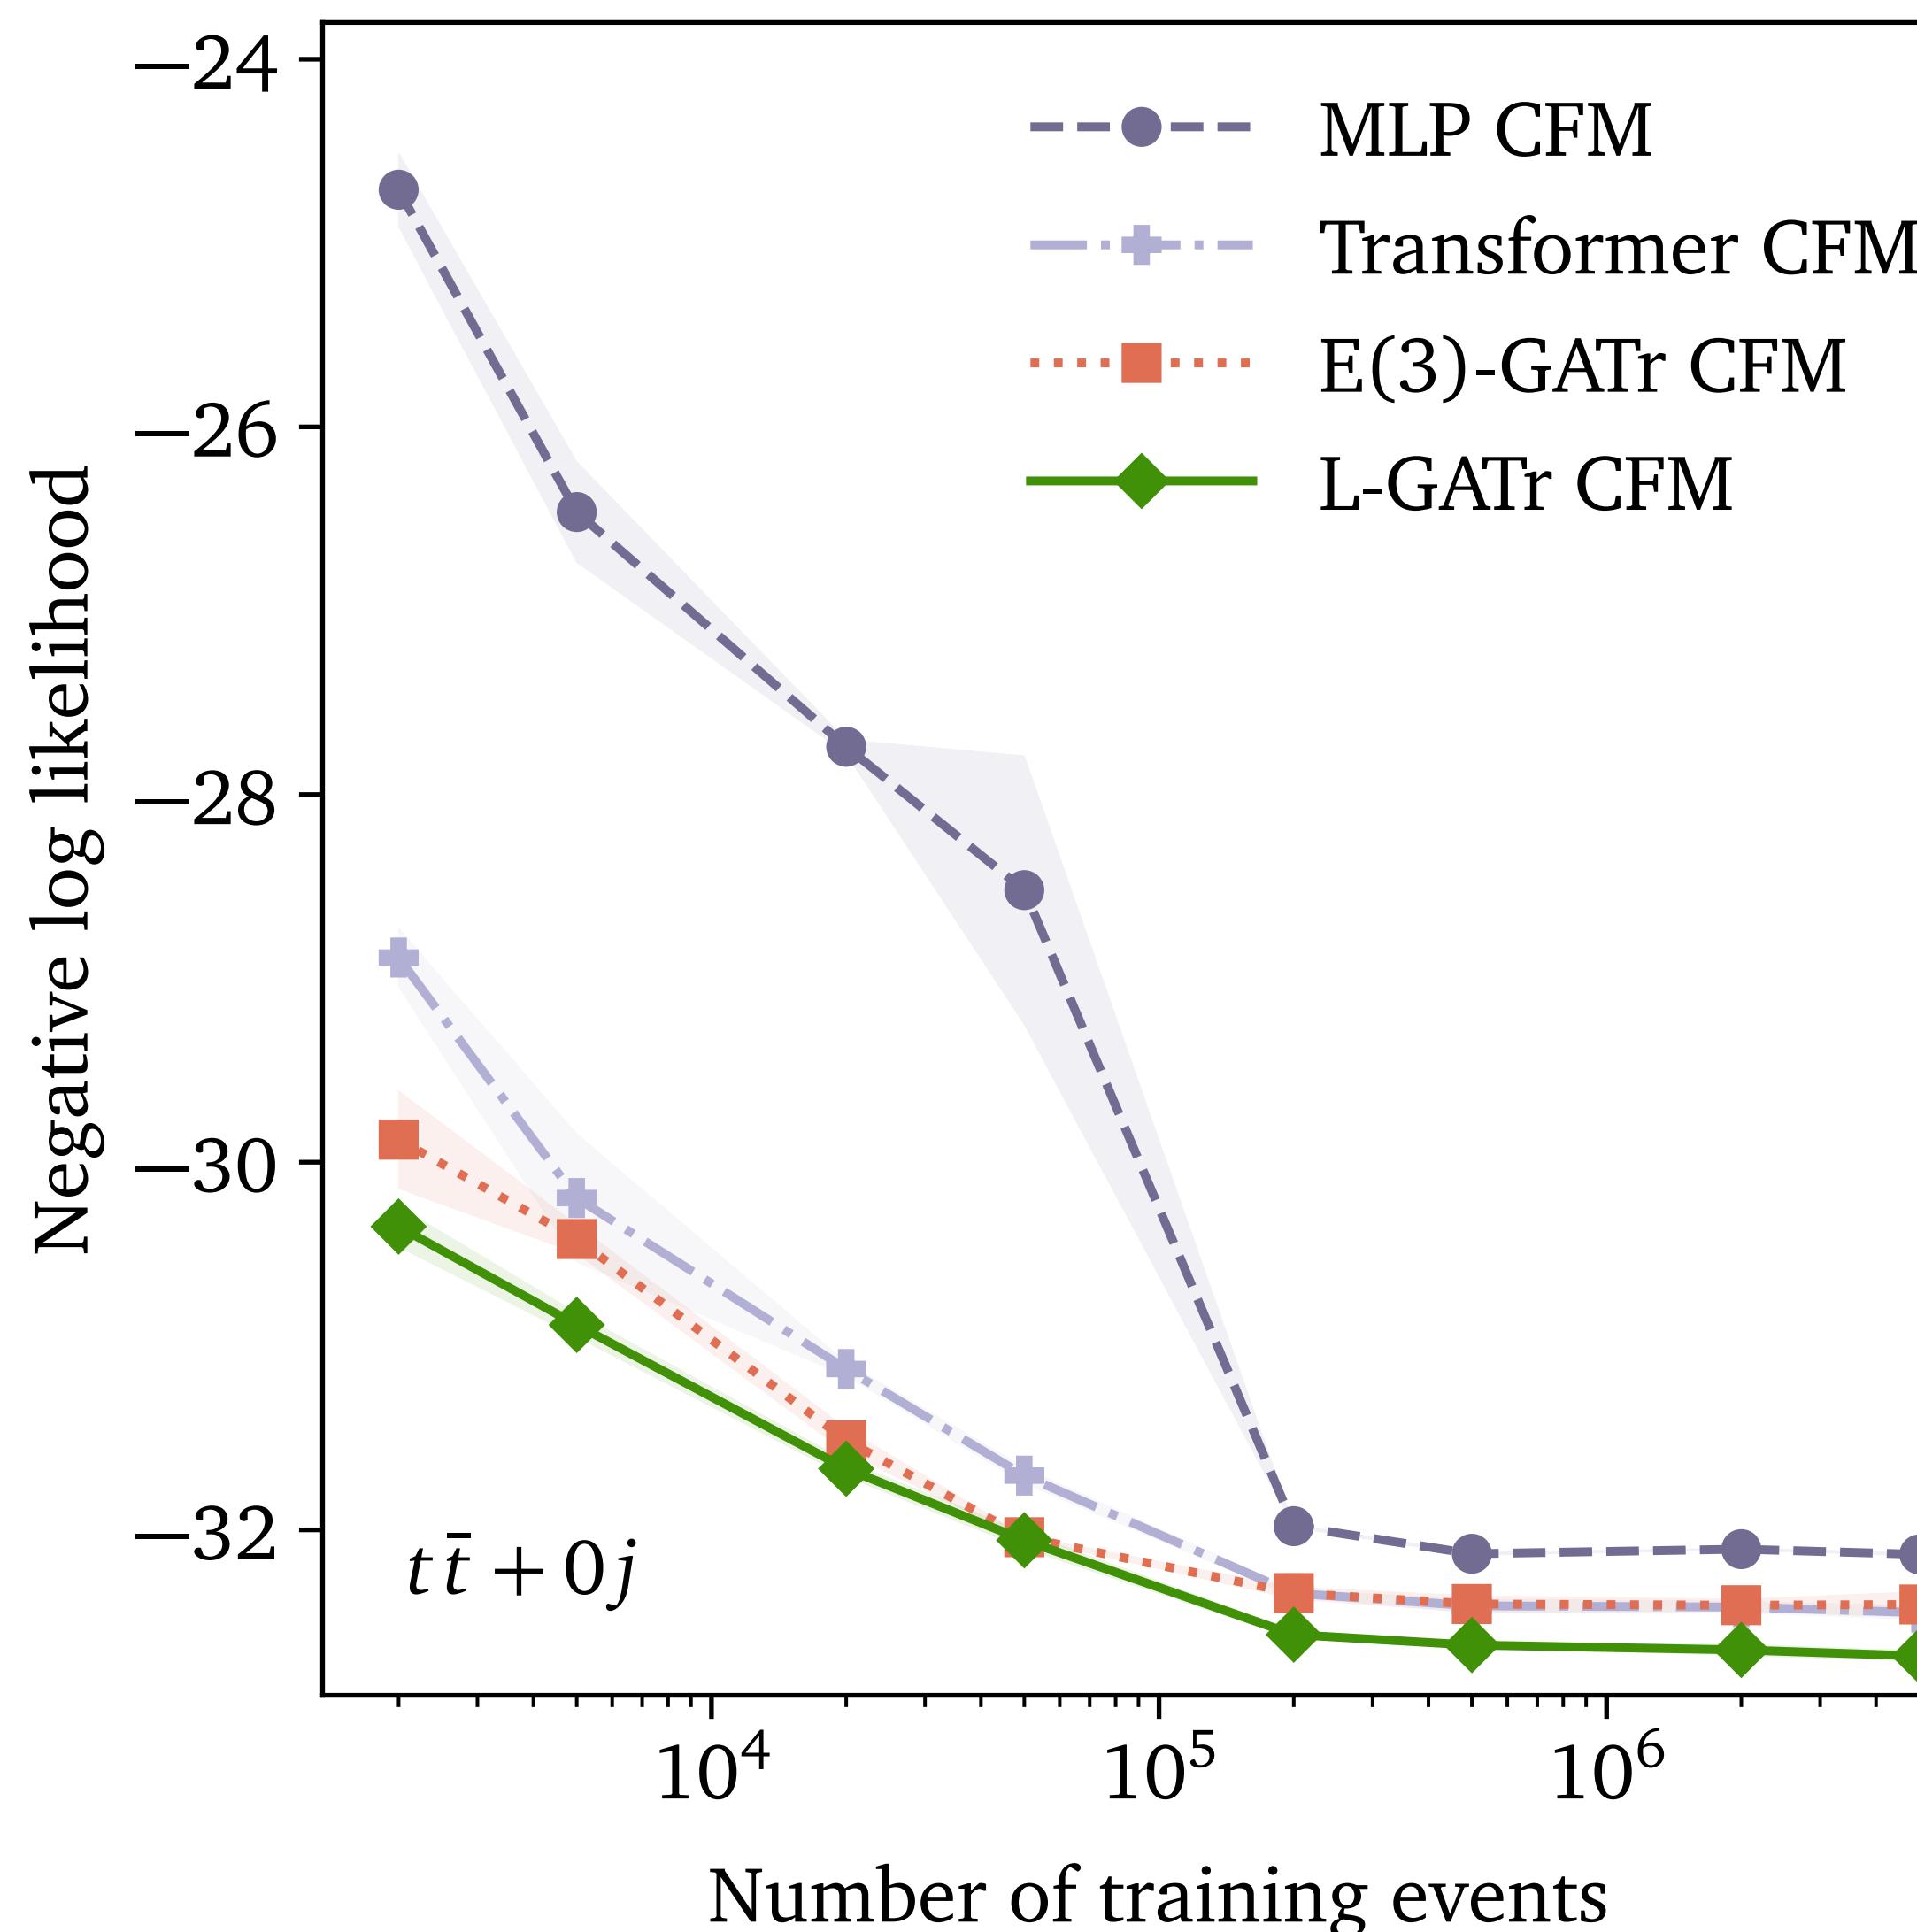


# Event generation

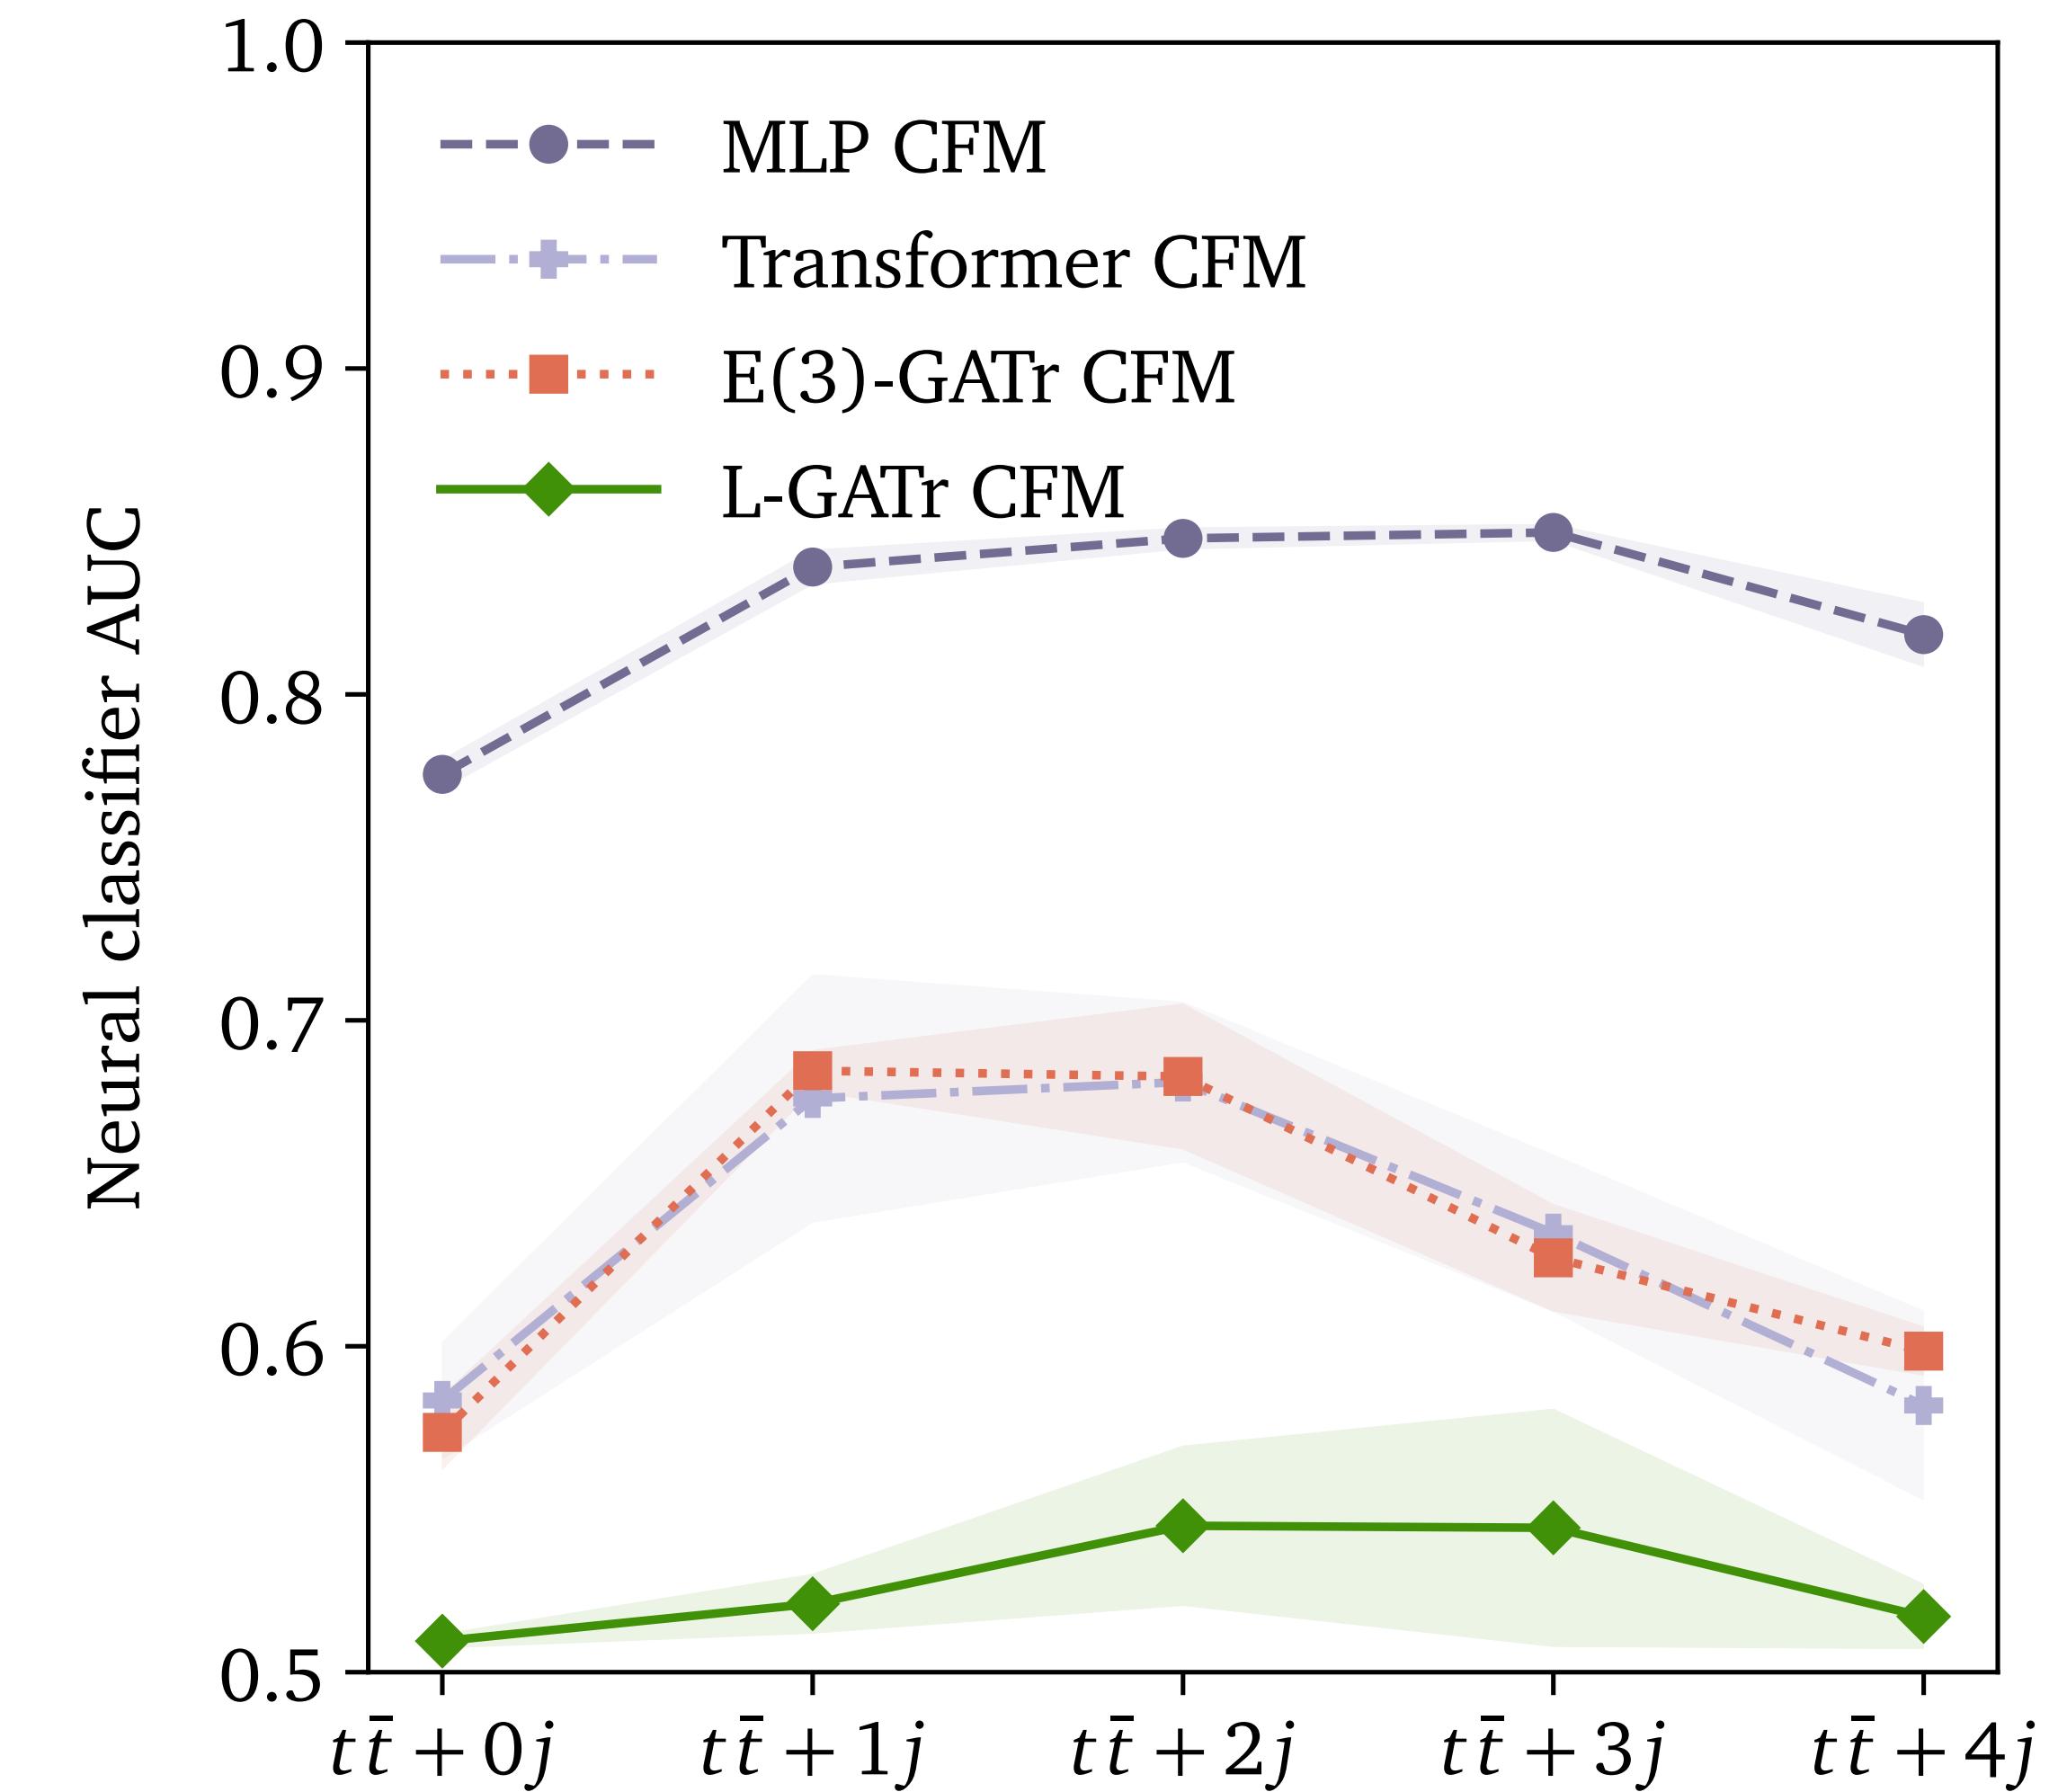
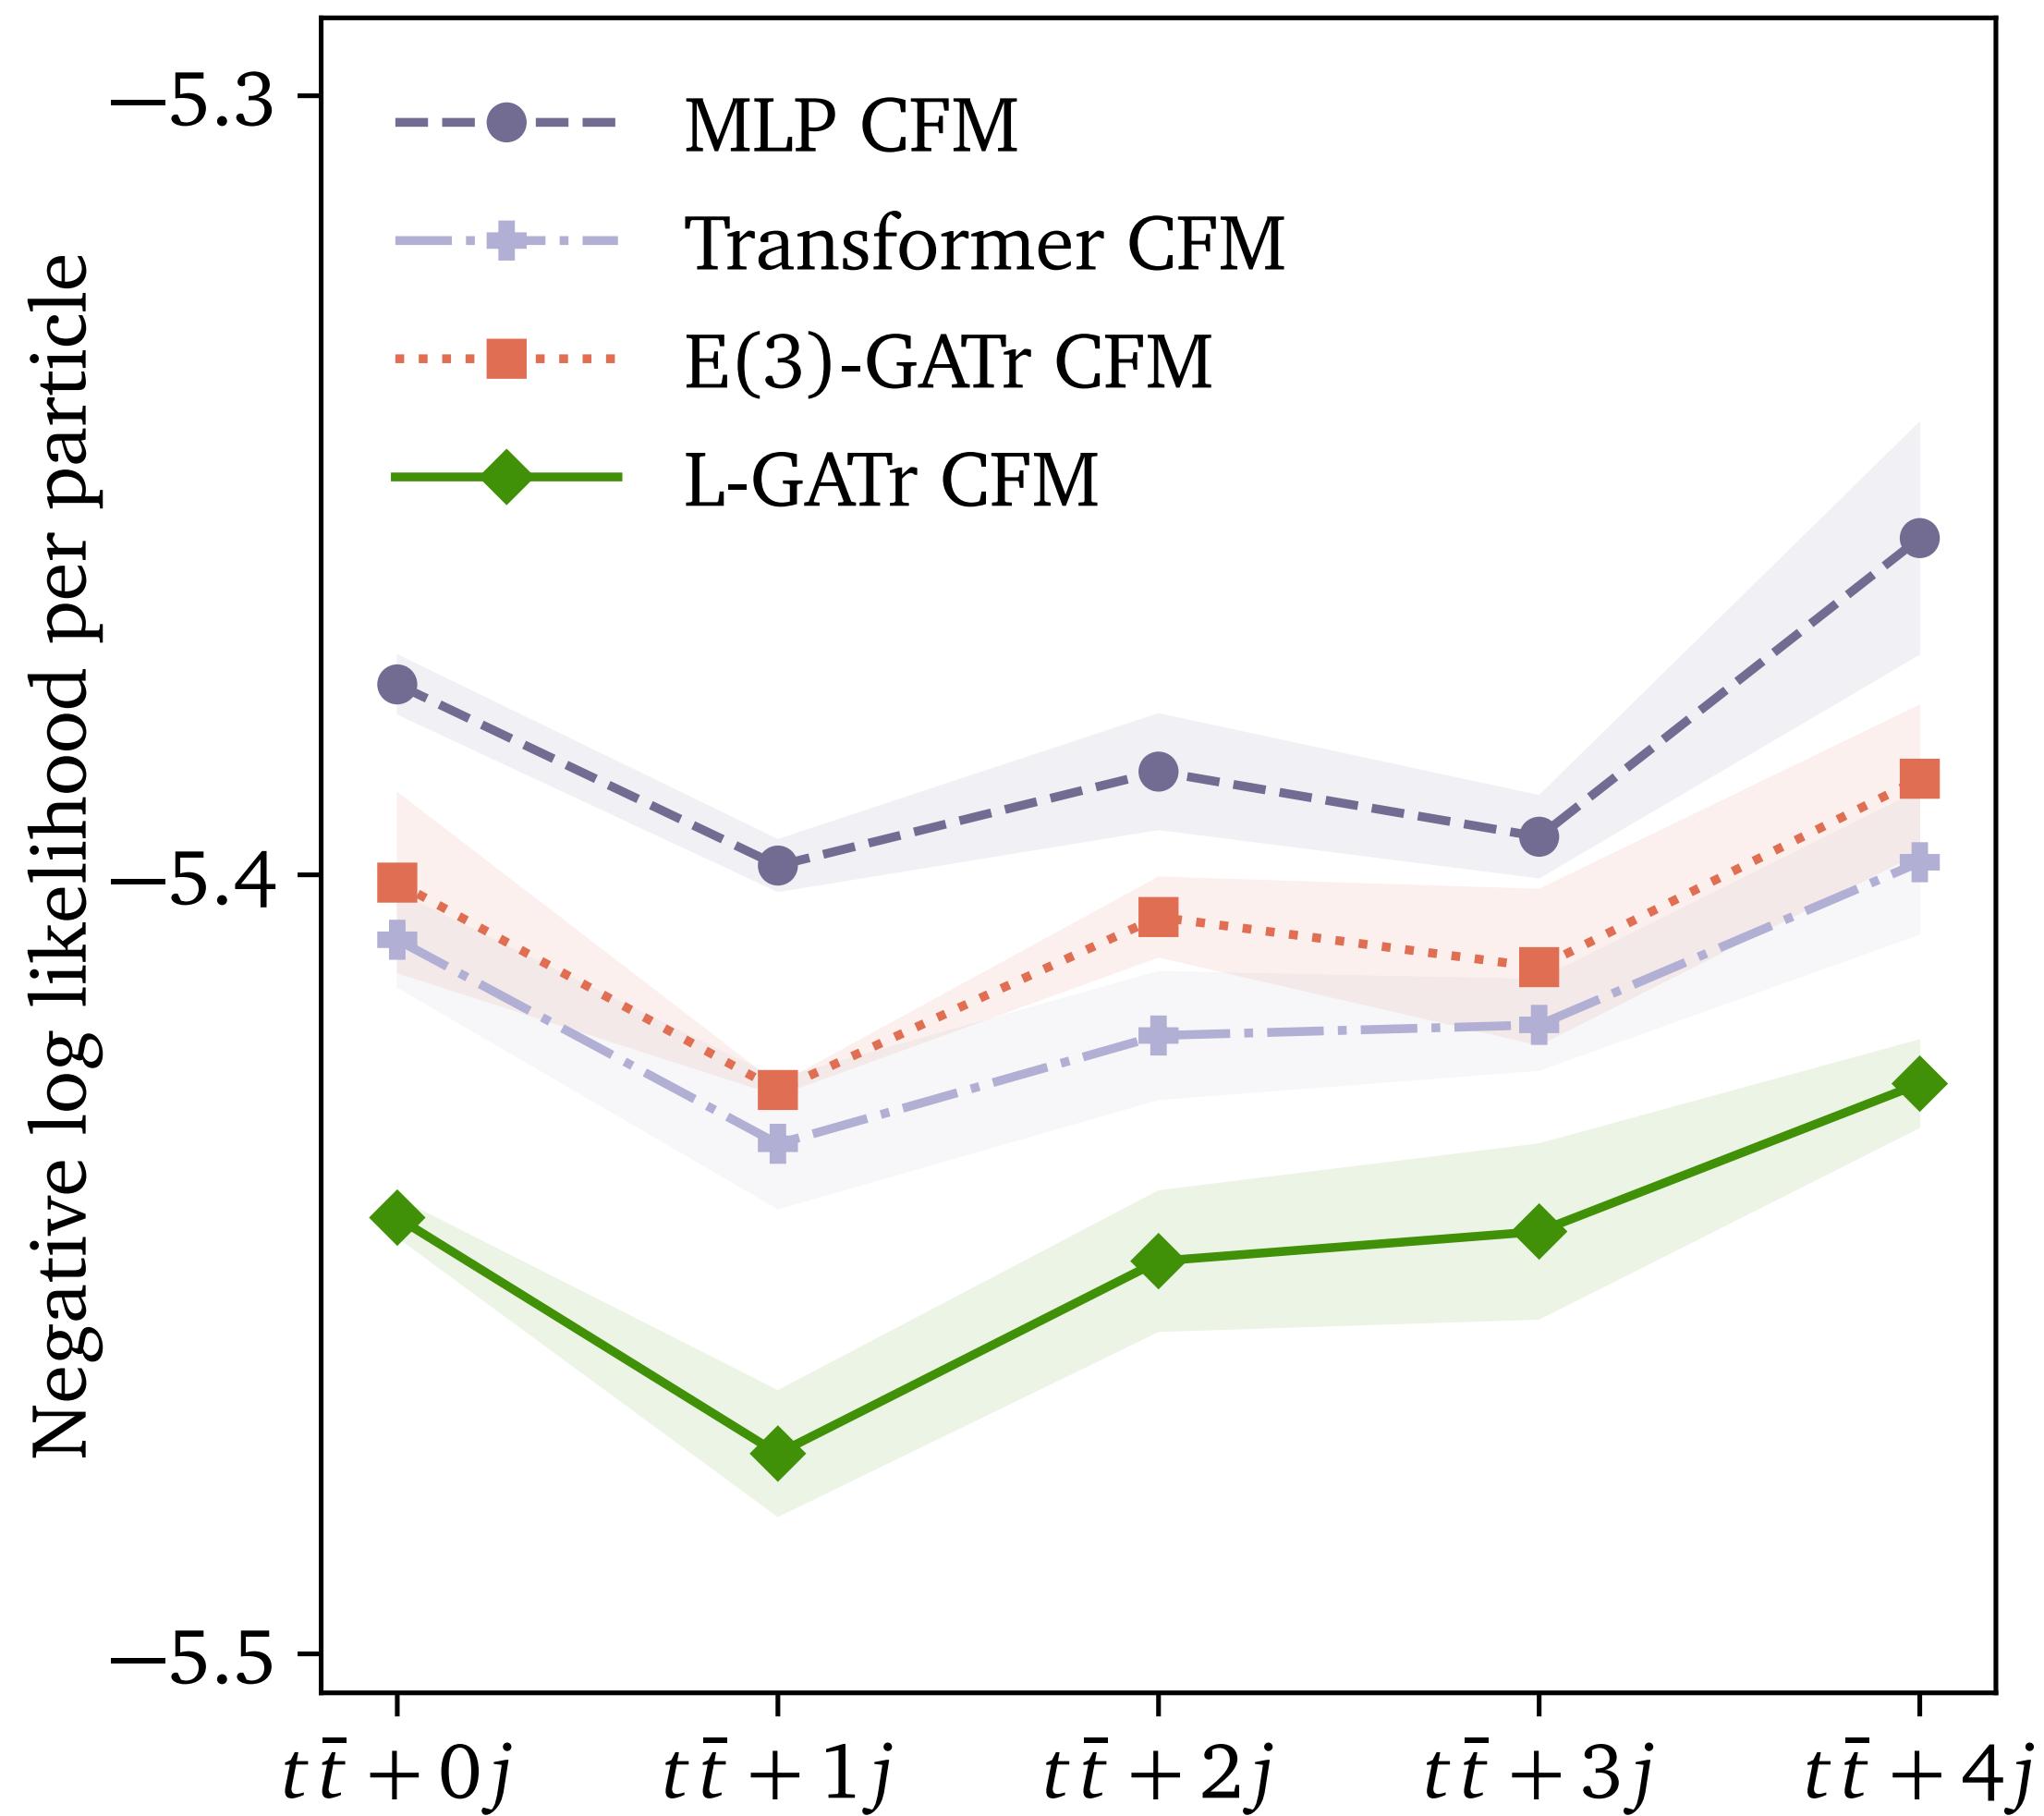


Equivariance helps,  
especially for angular correlations

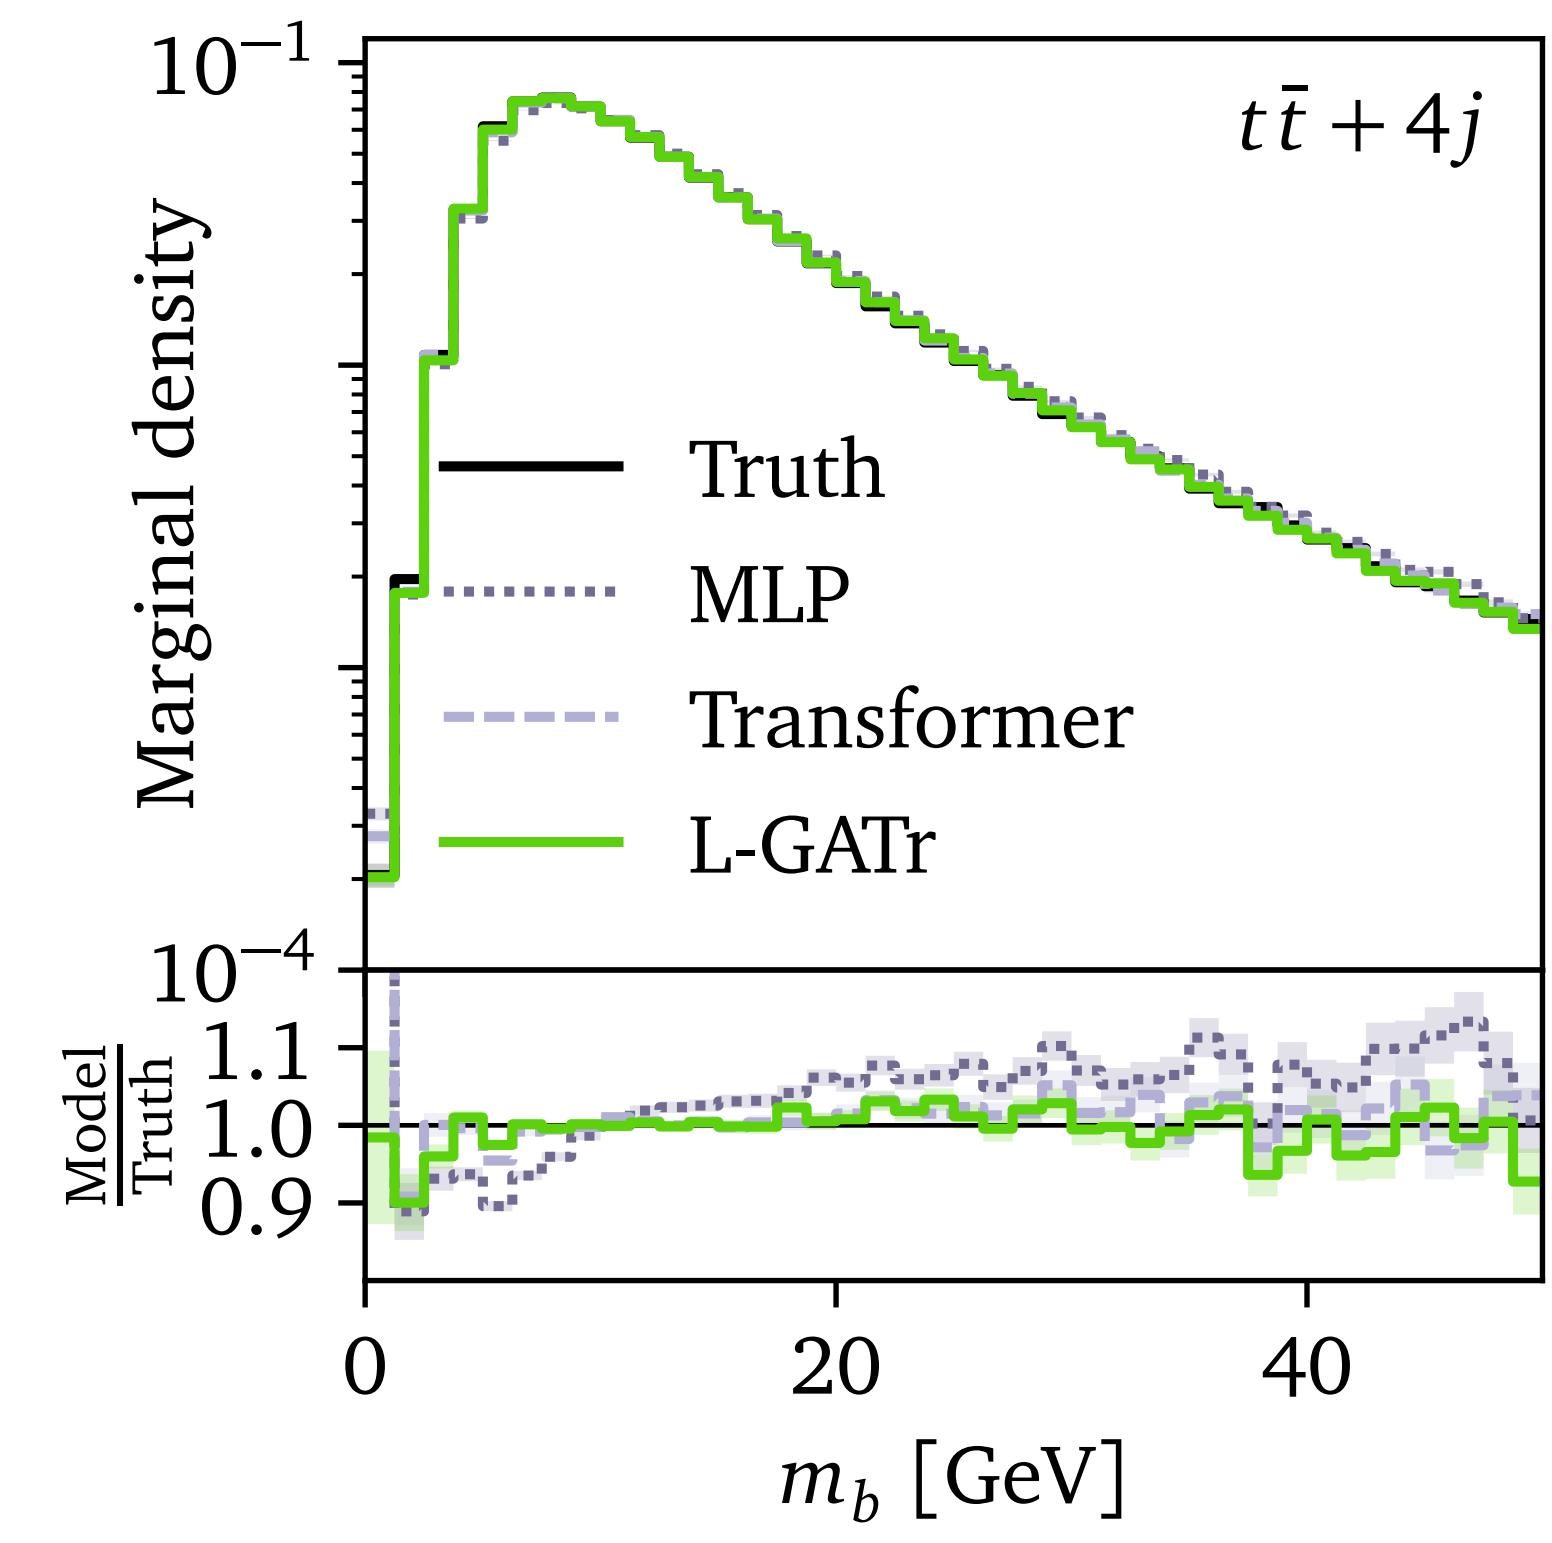
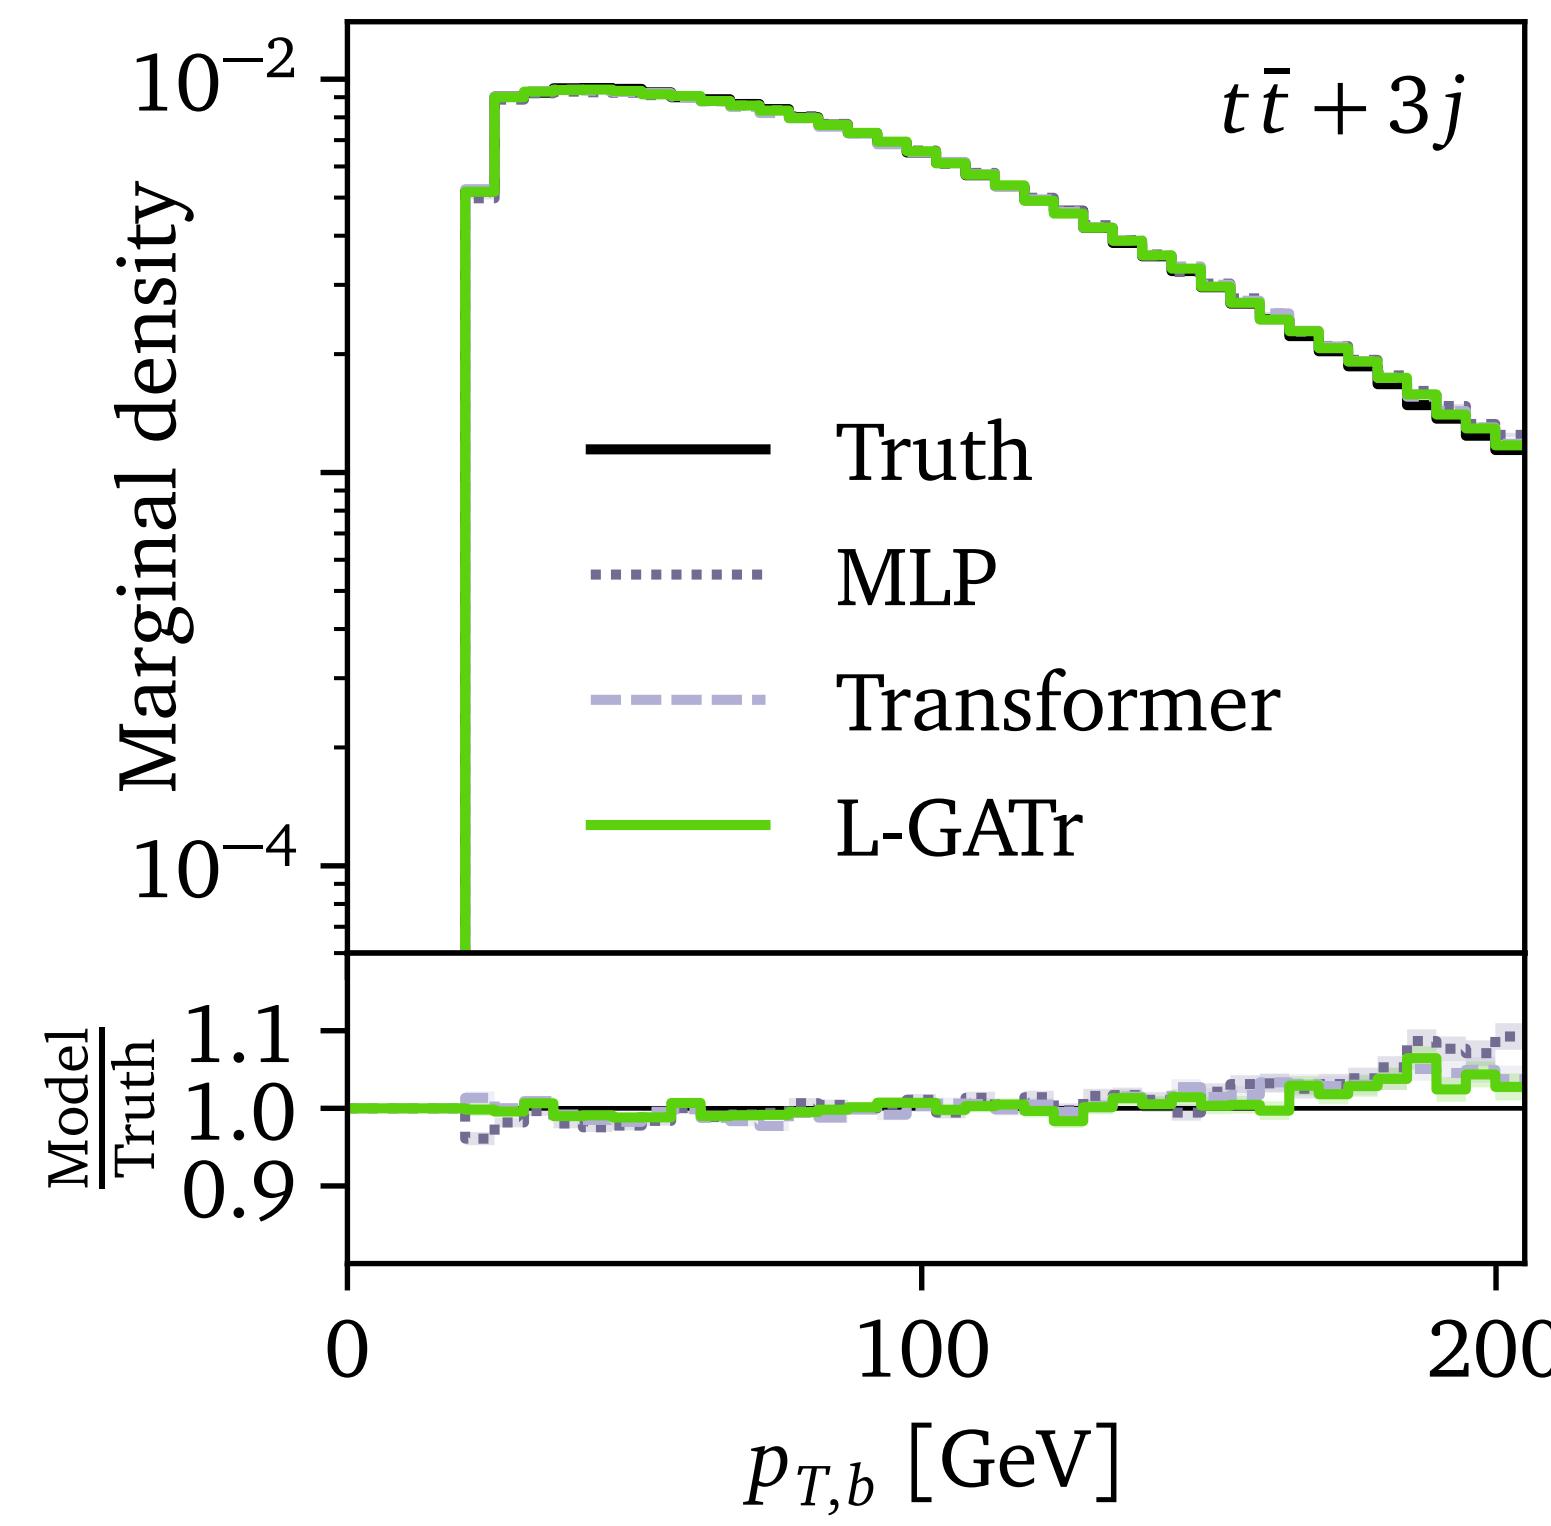
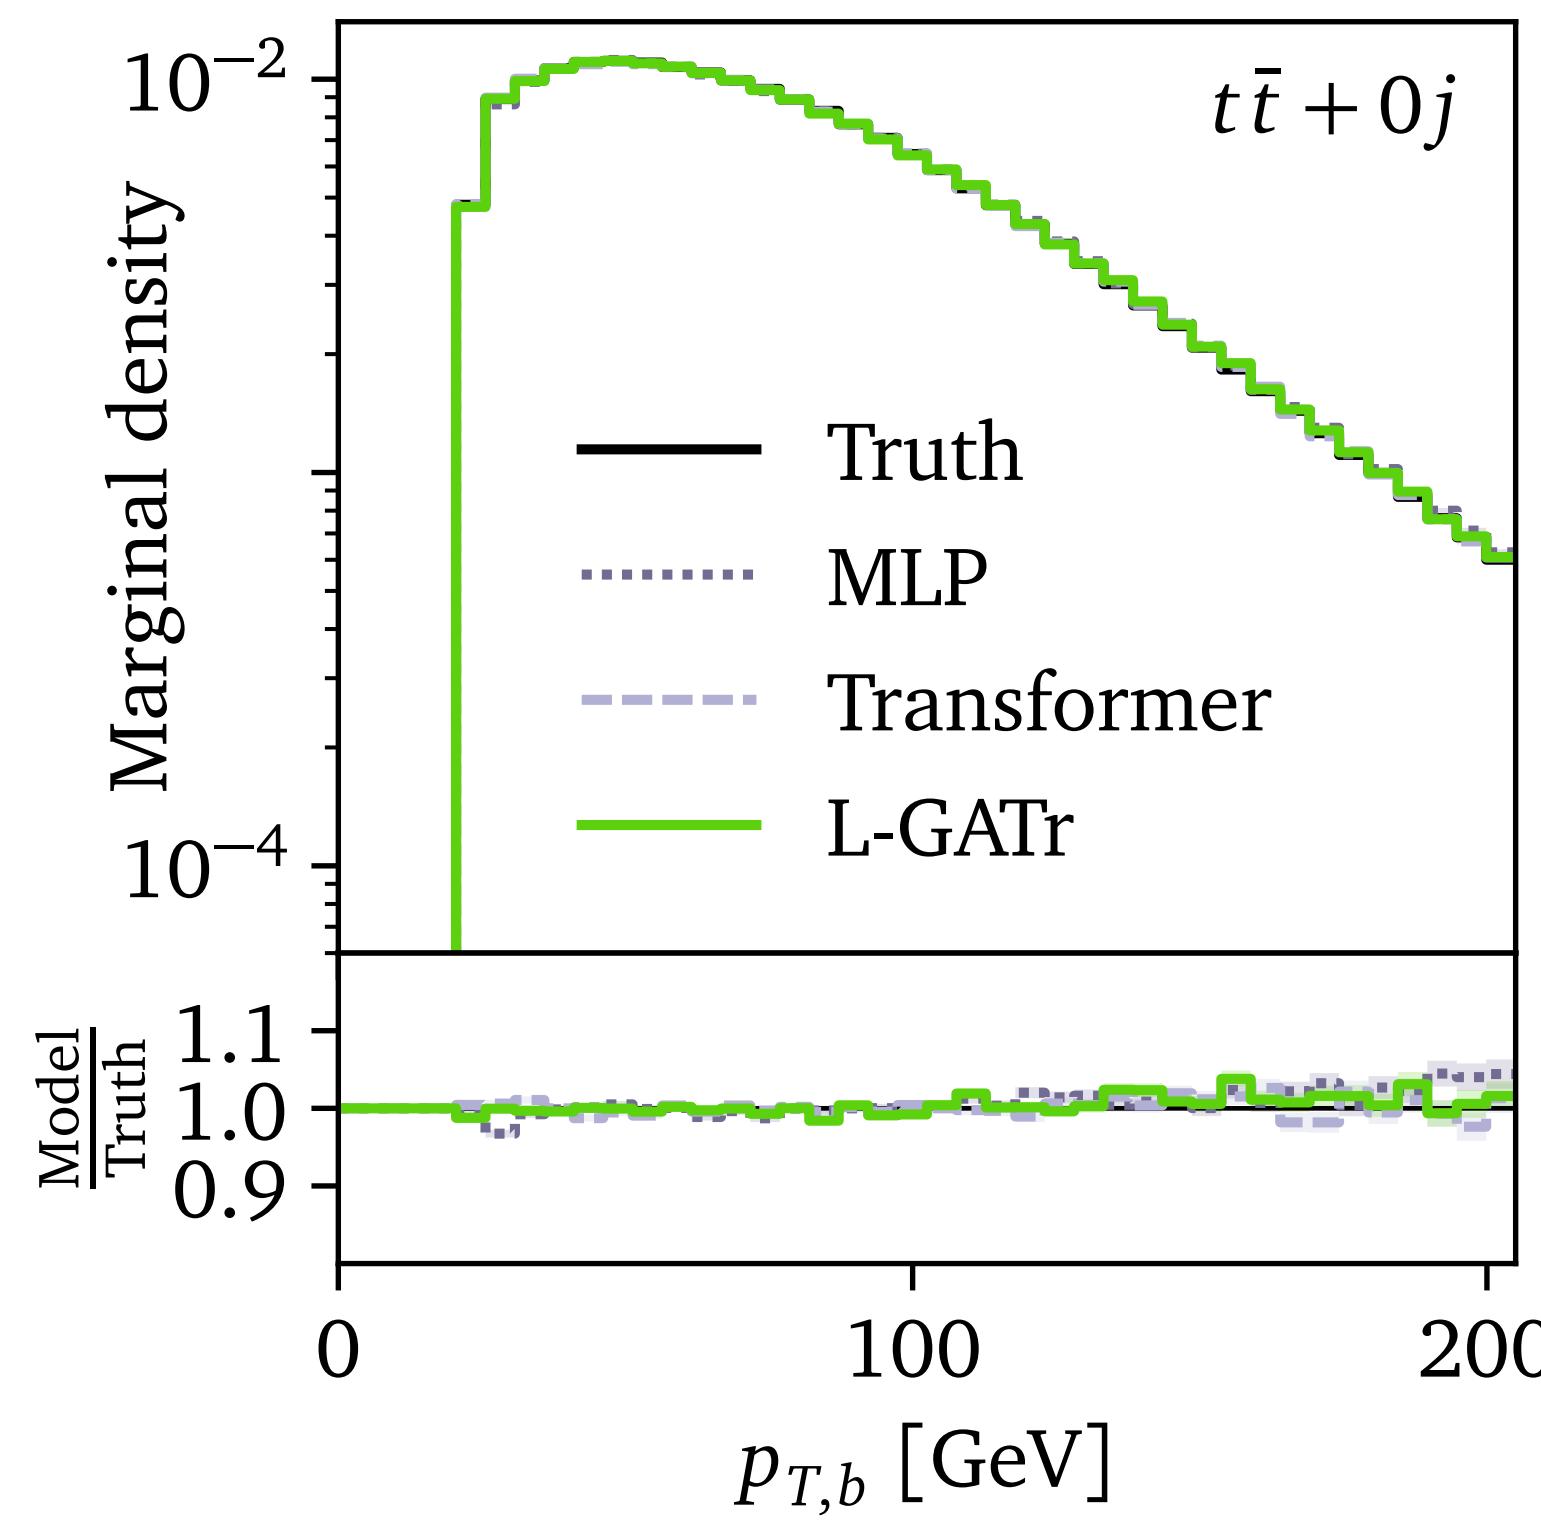
# Event generation



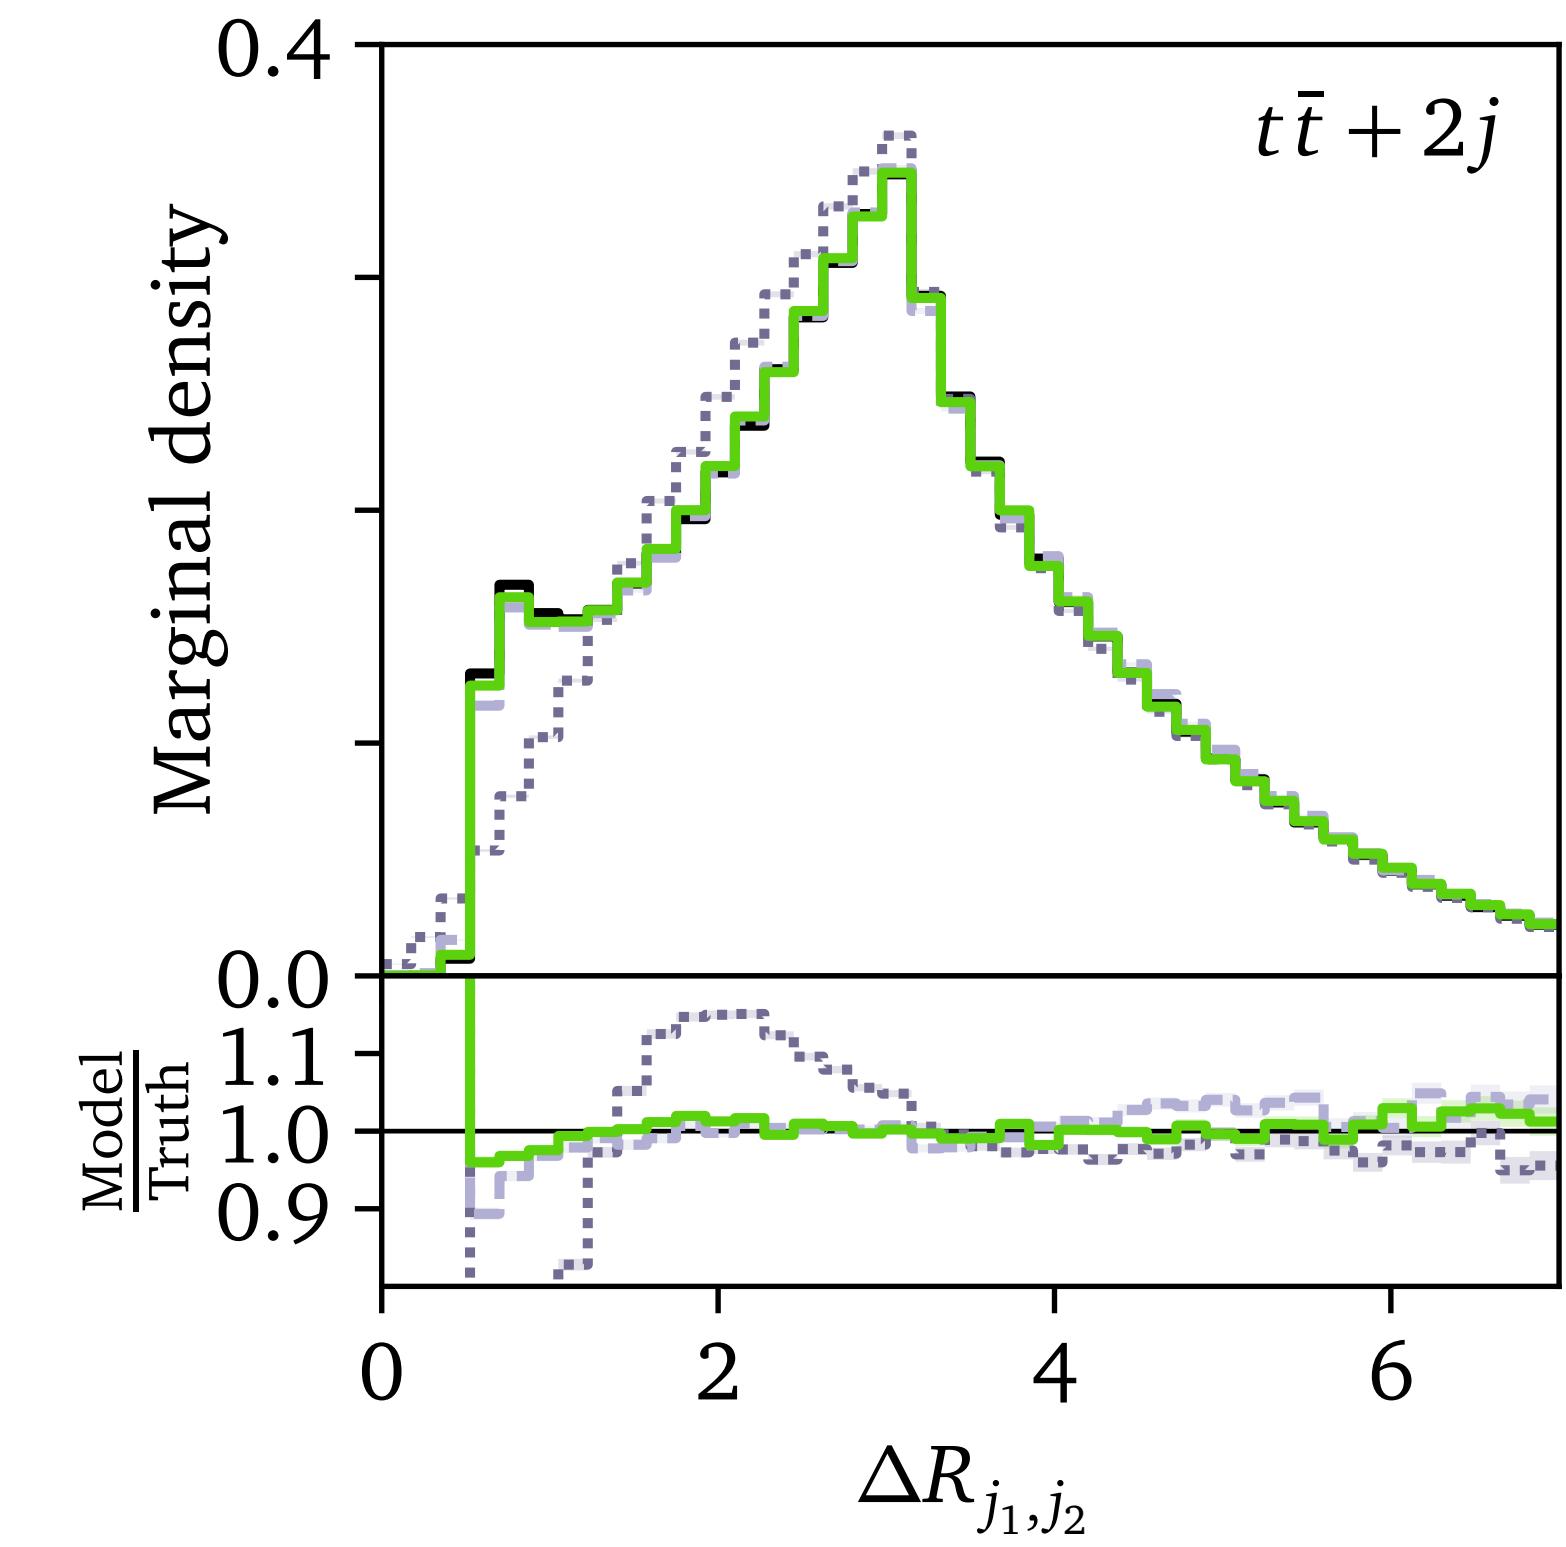
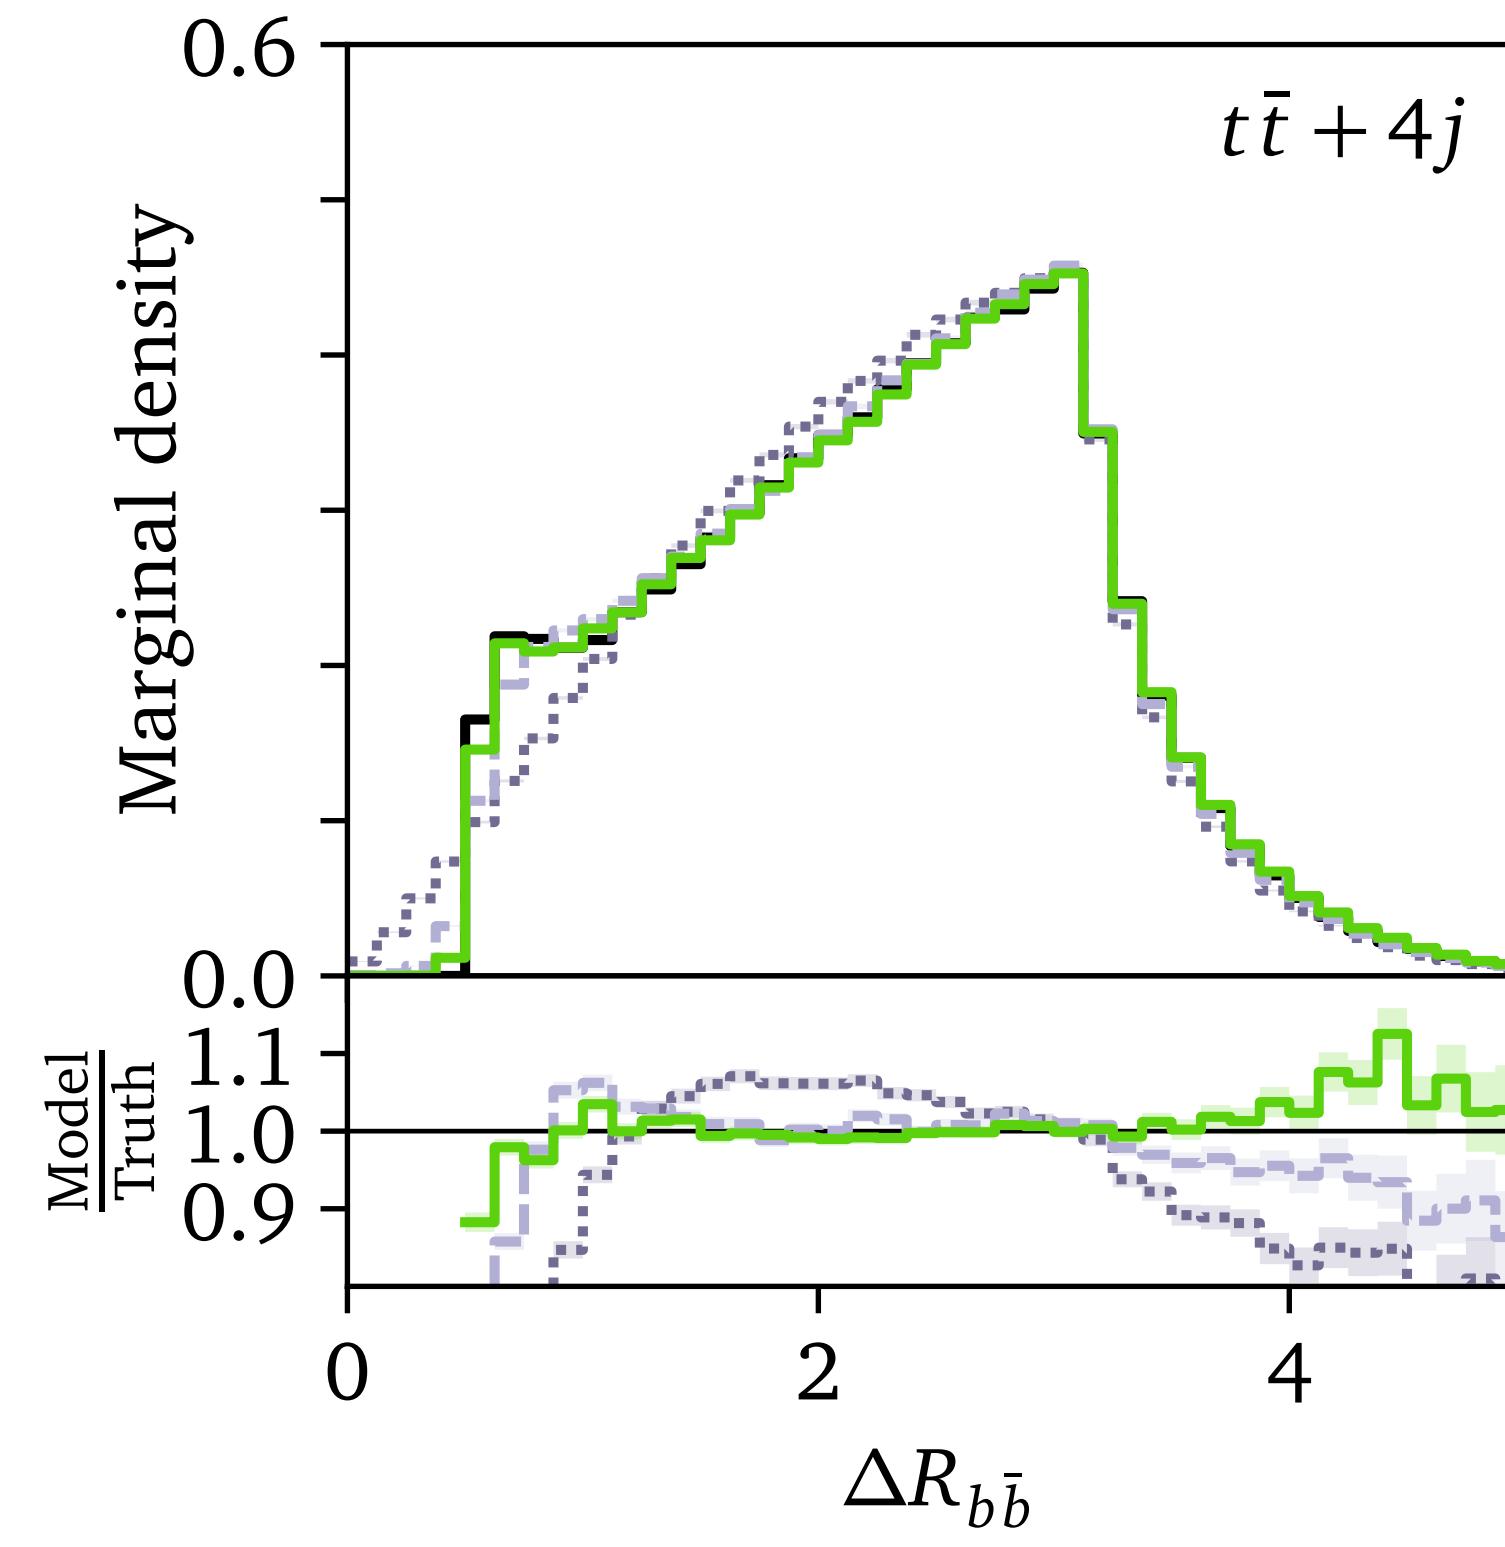
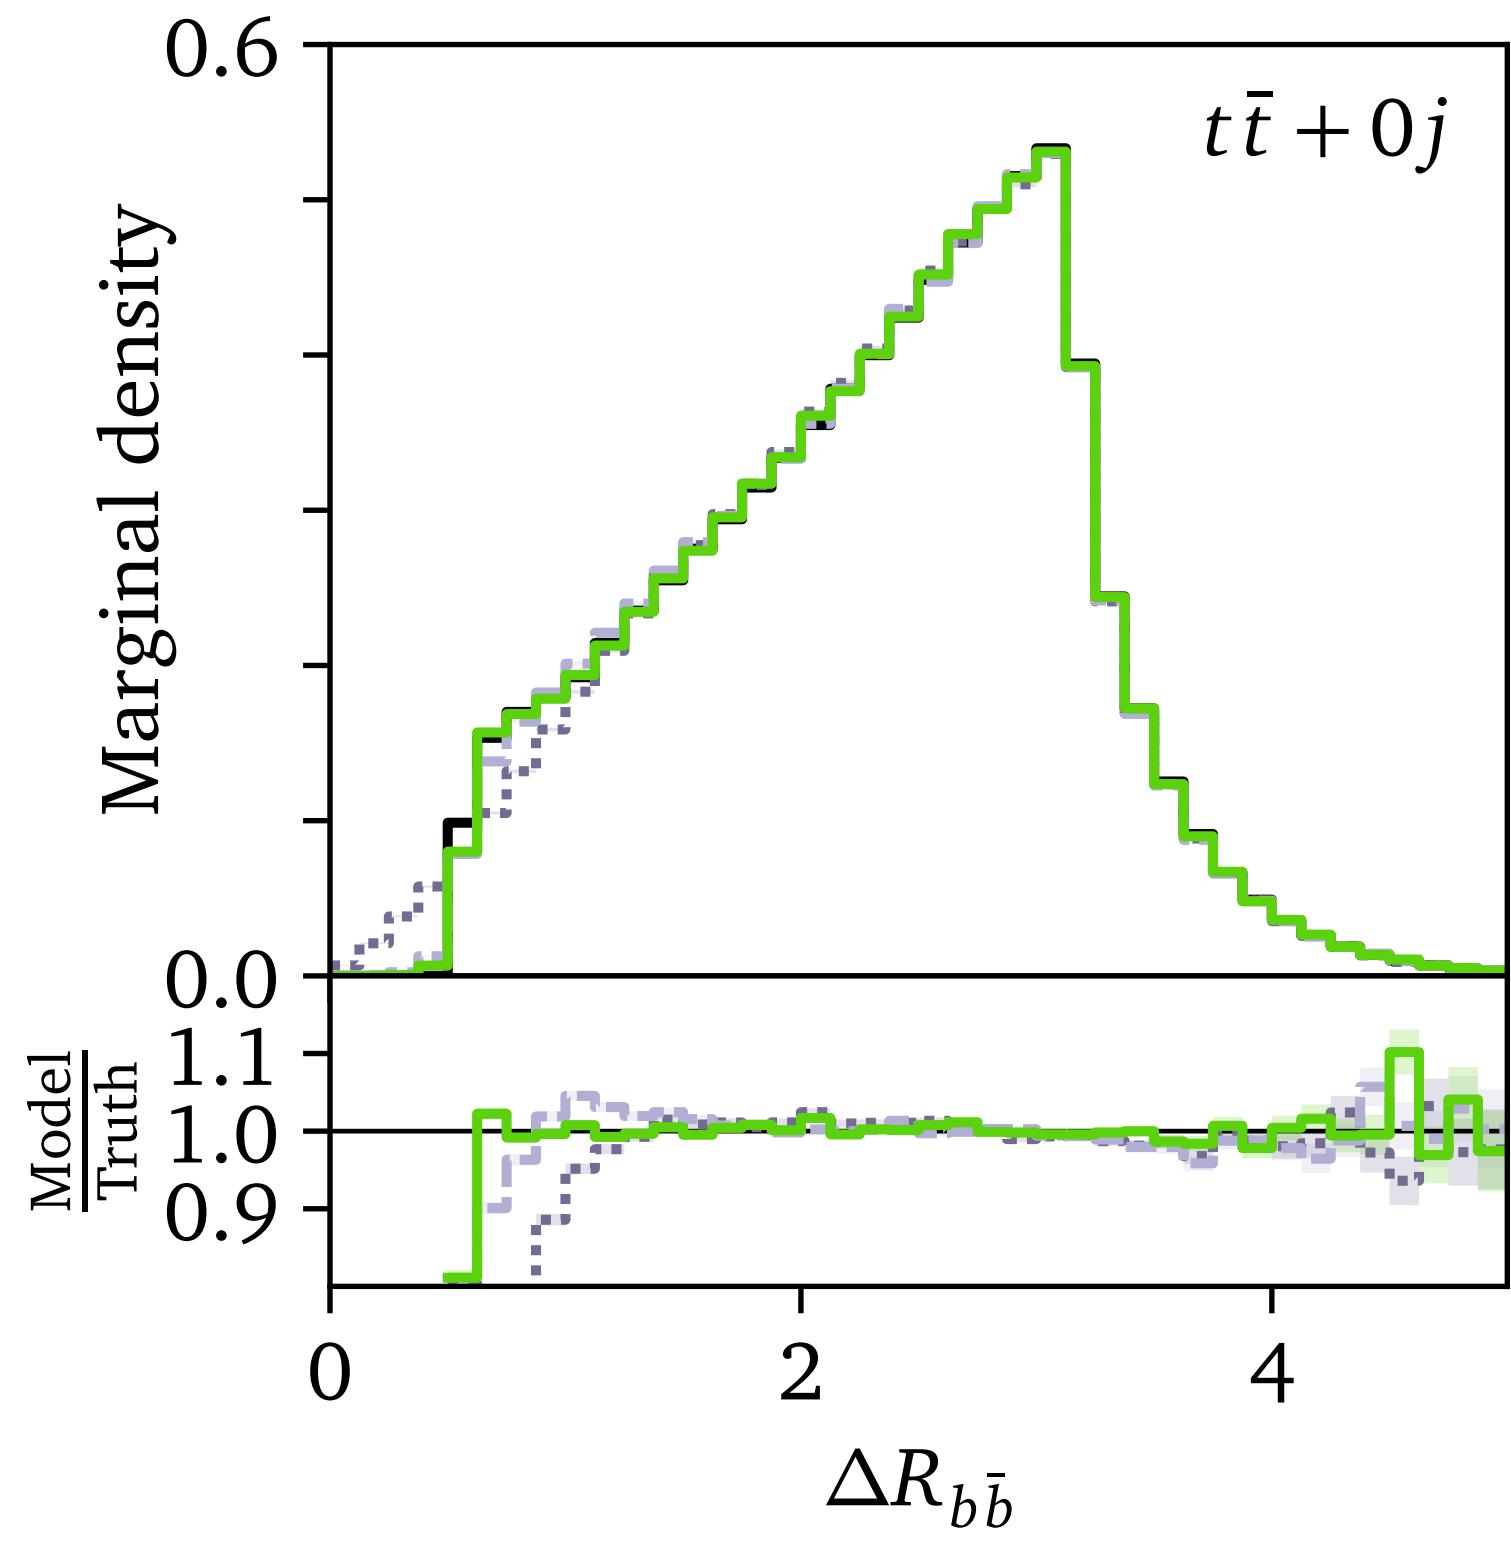
# Event generation



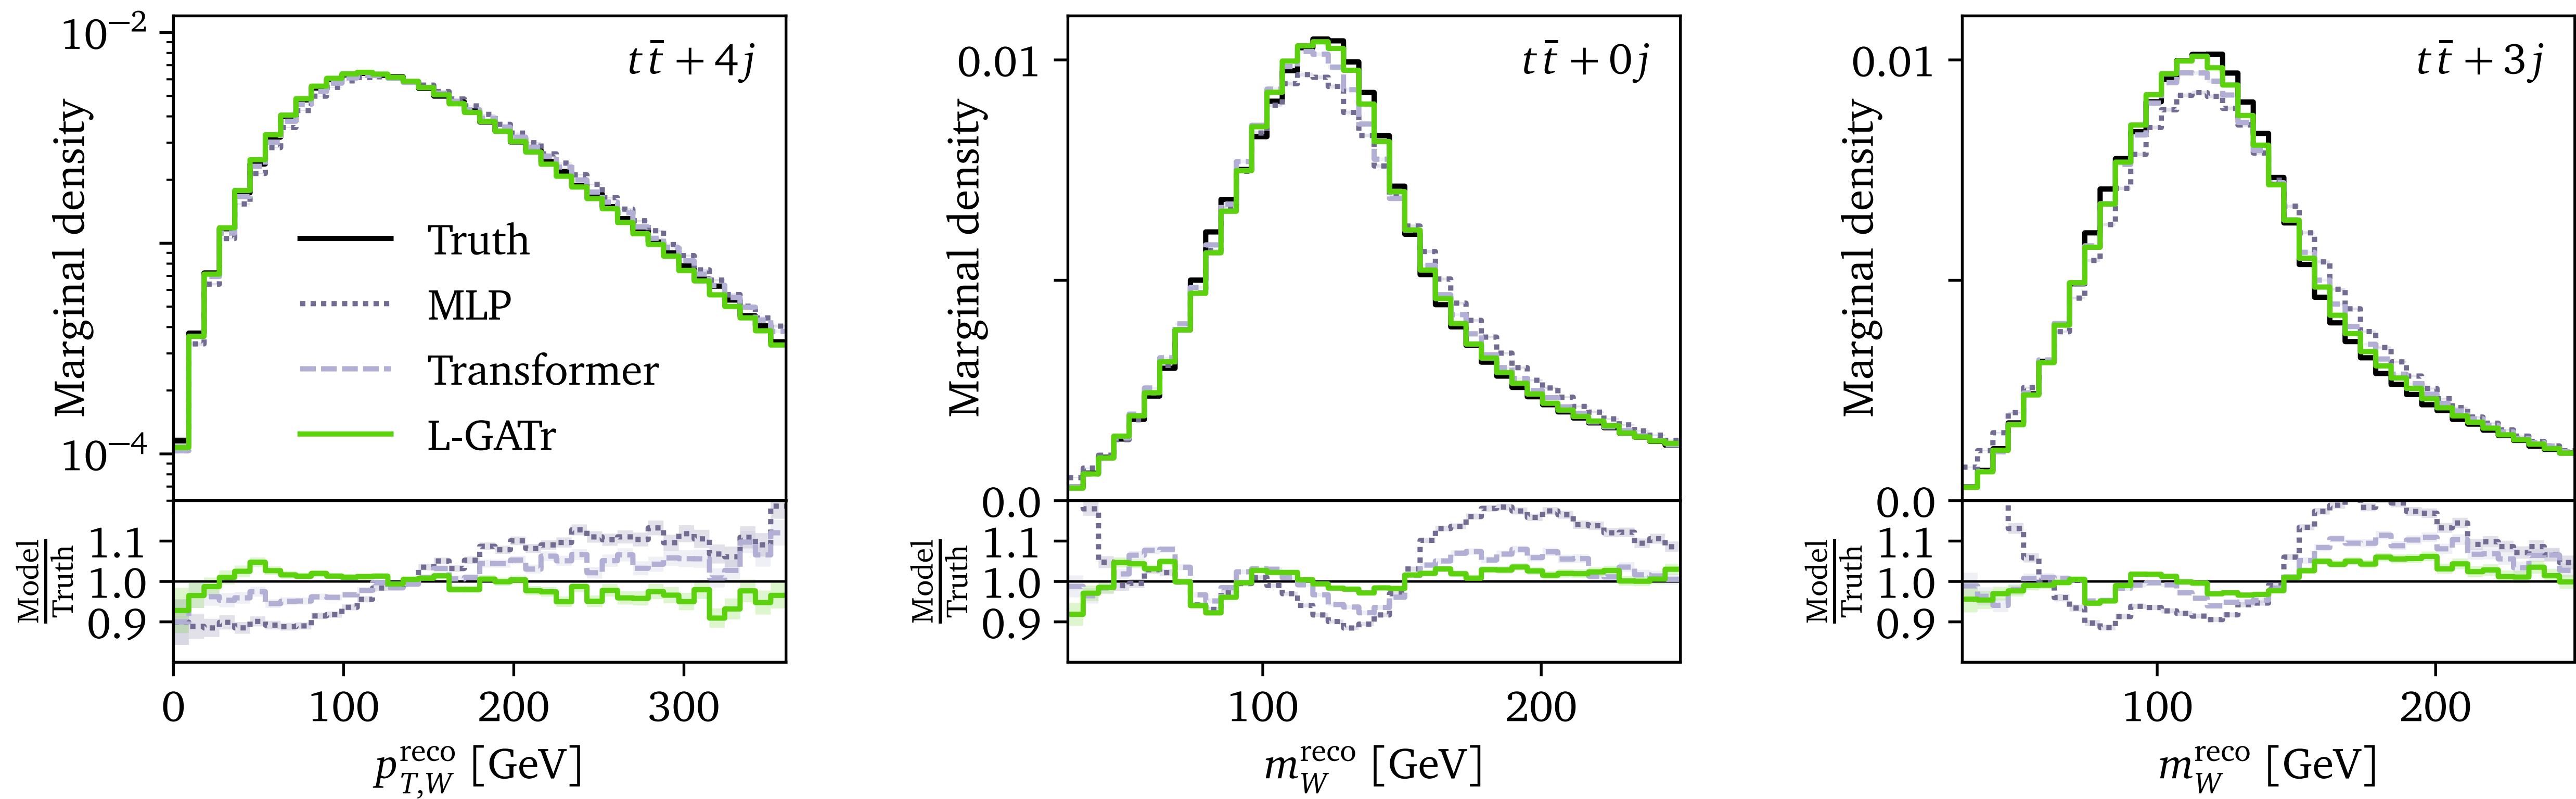
# Event generation



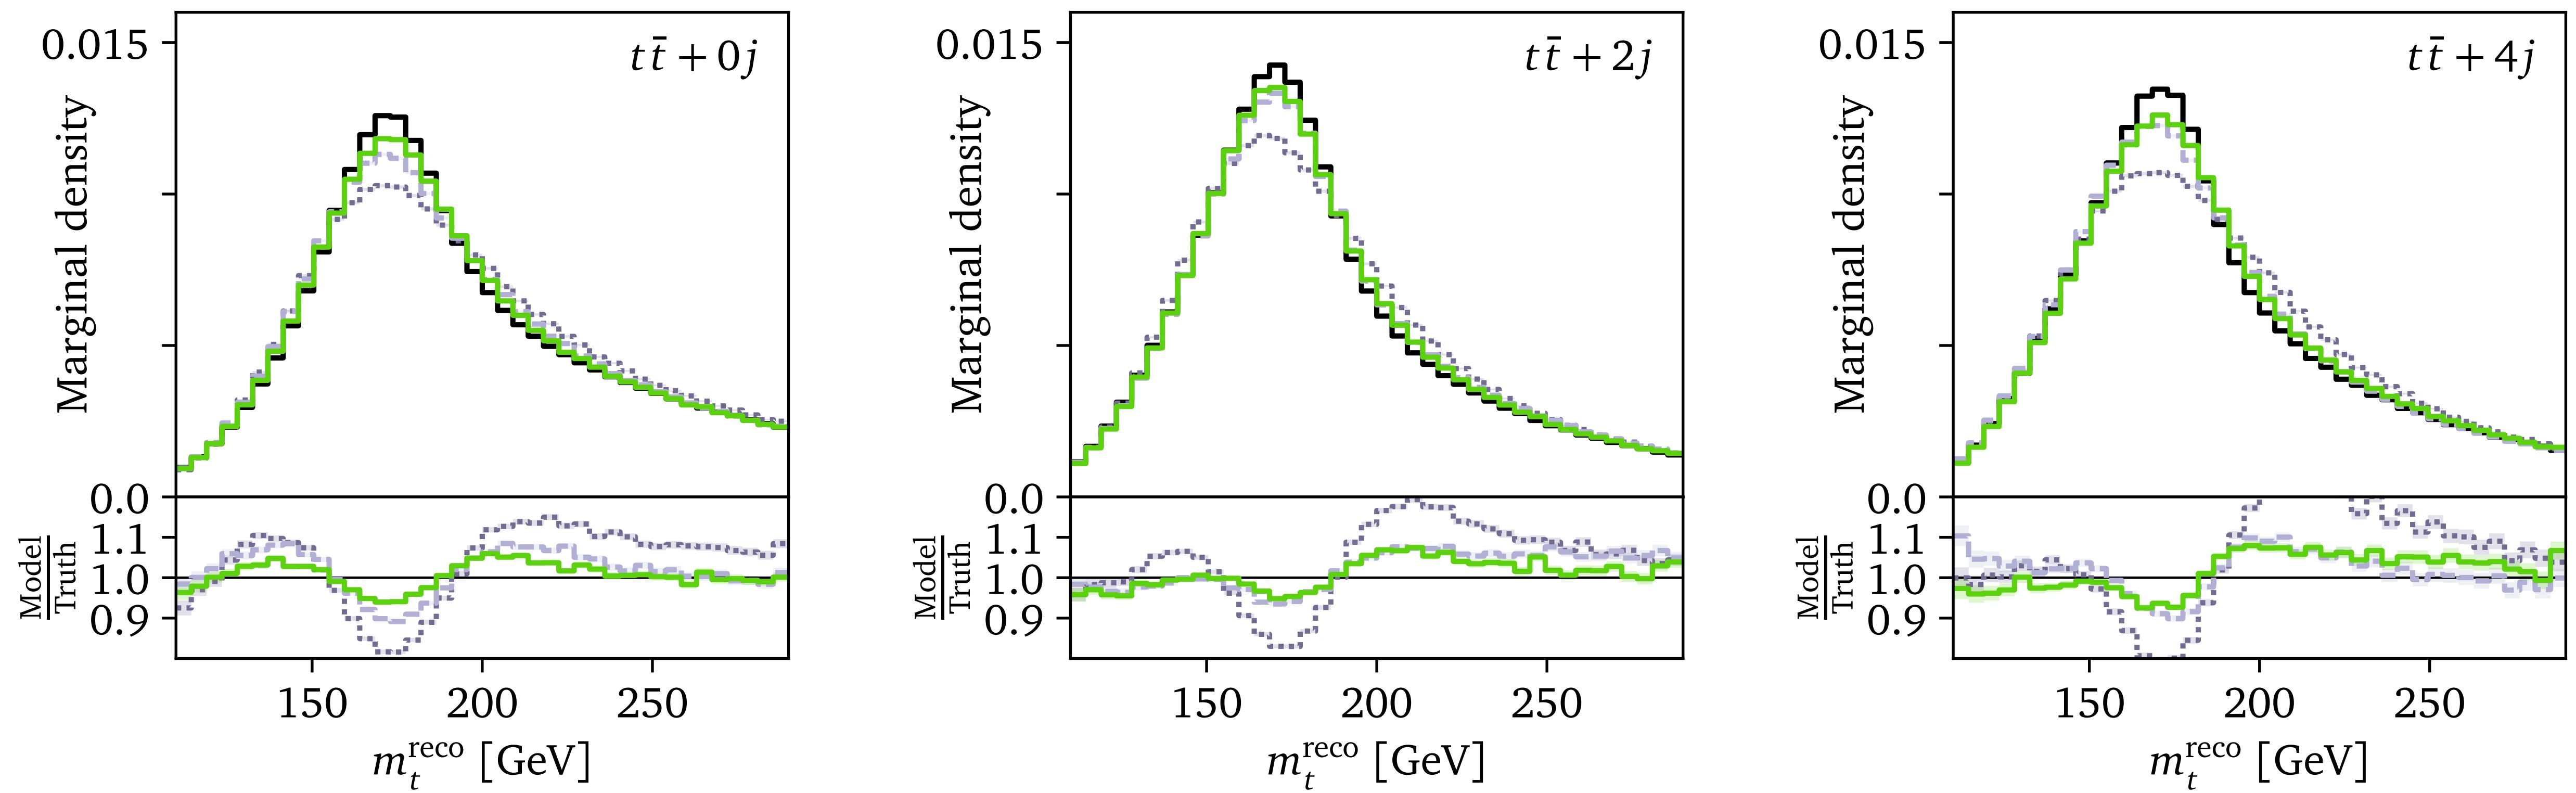
# Event generation



# Event generation

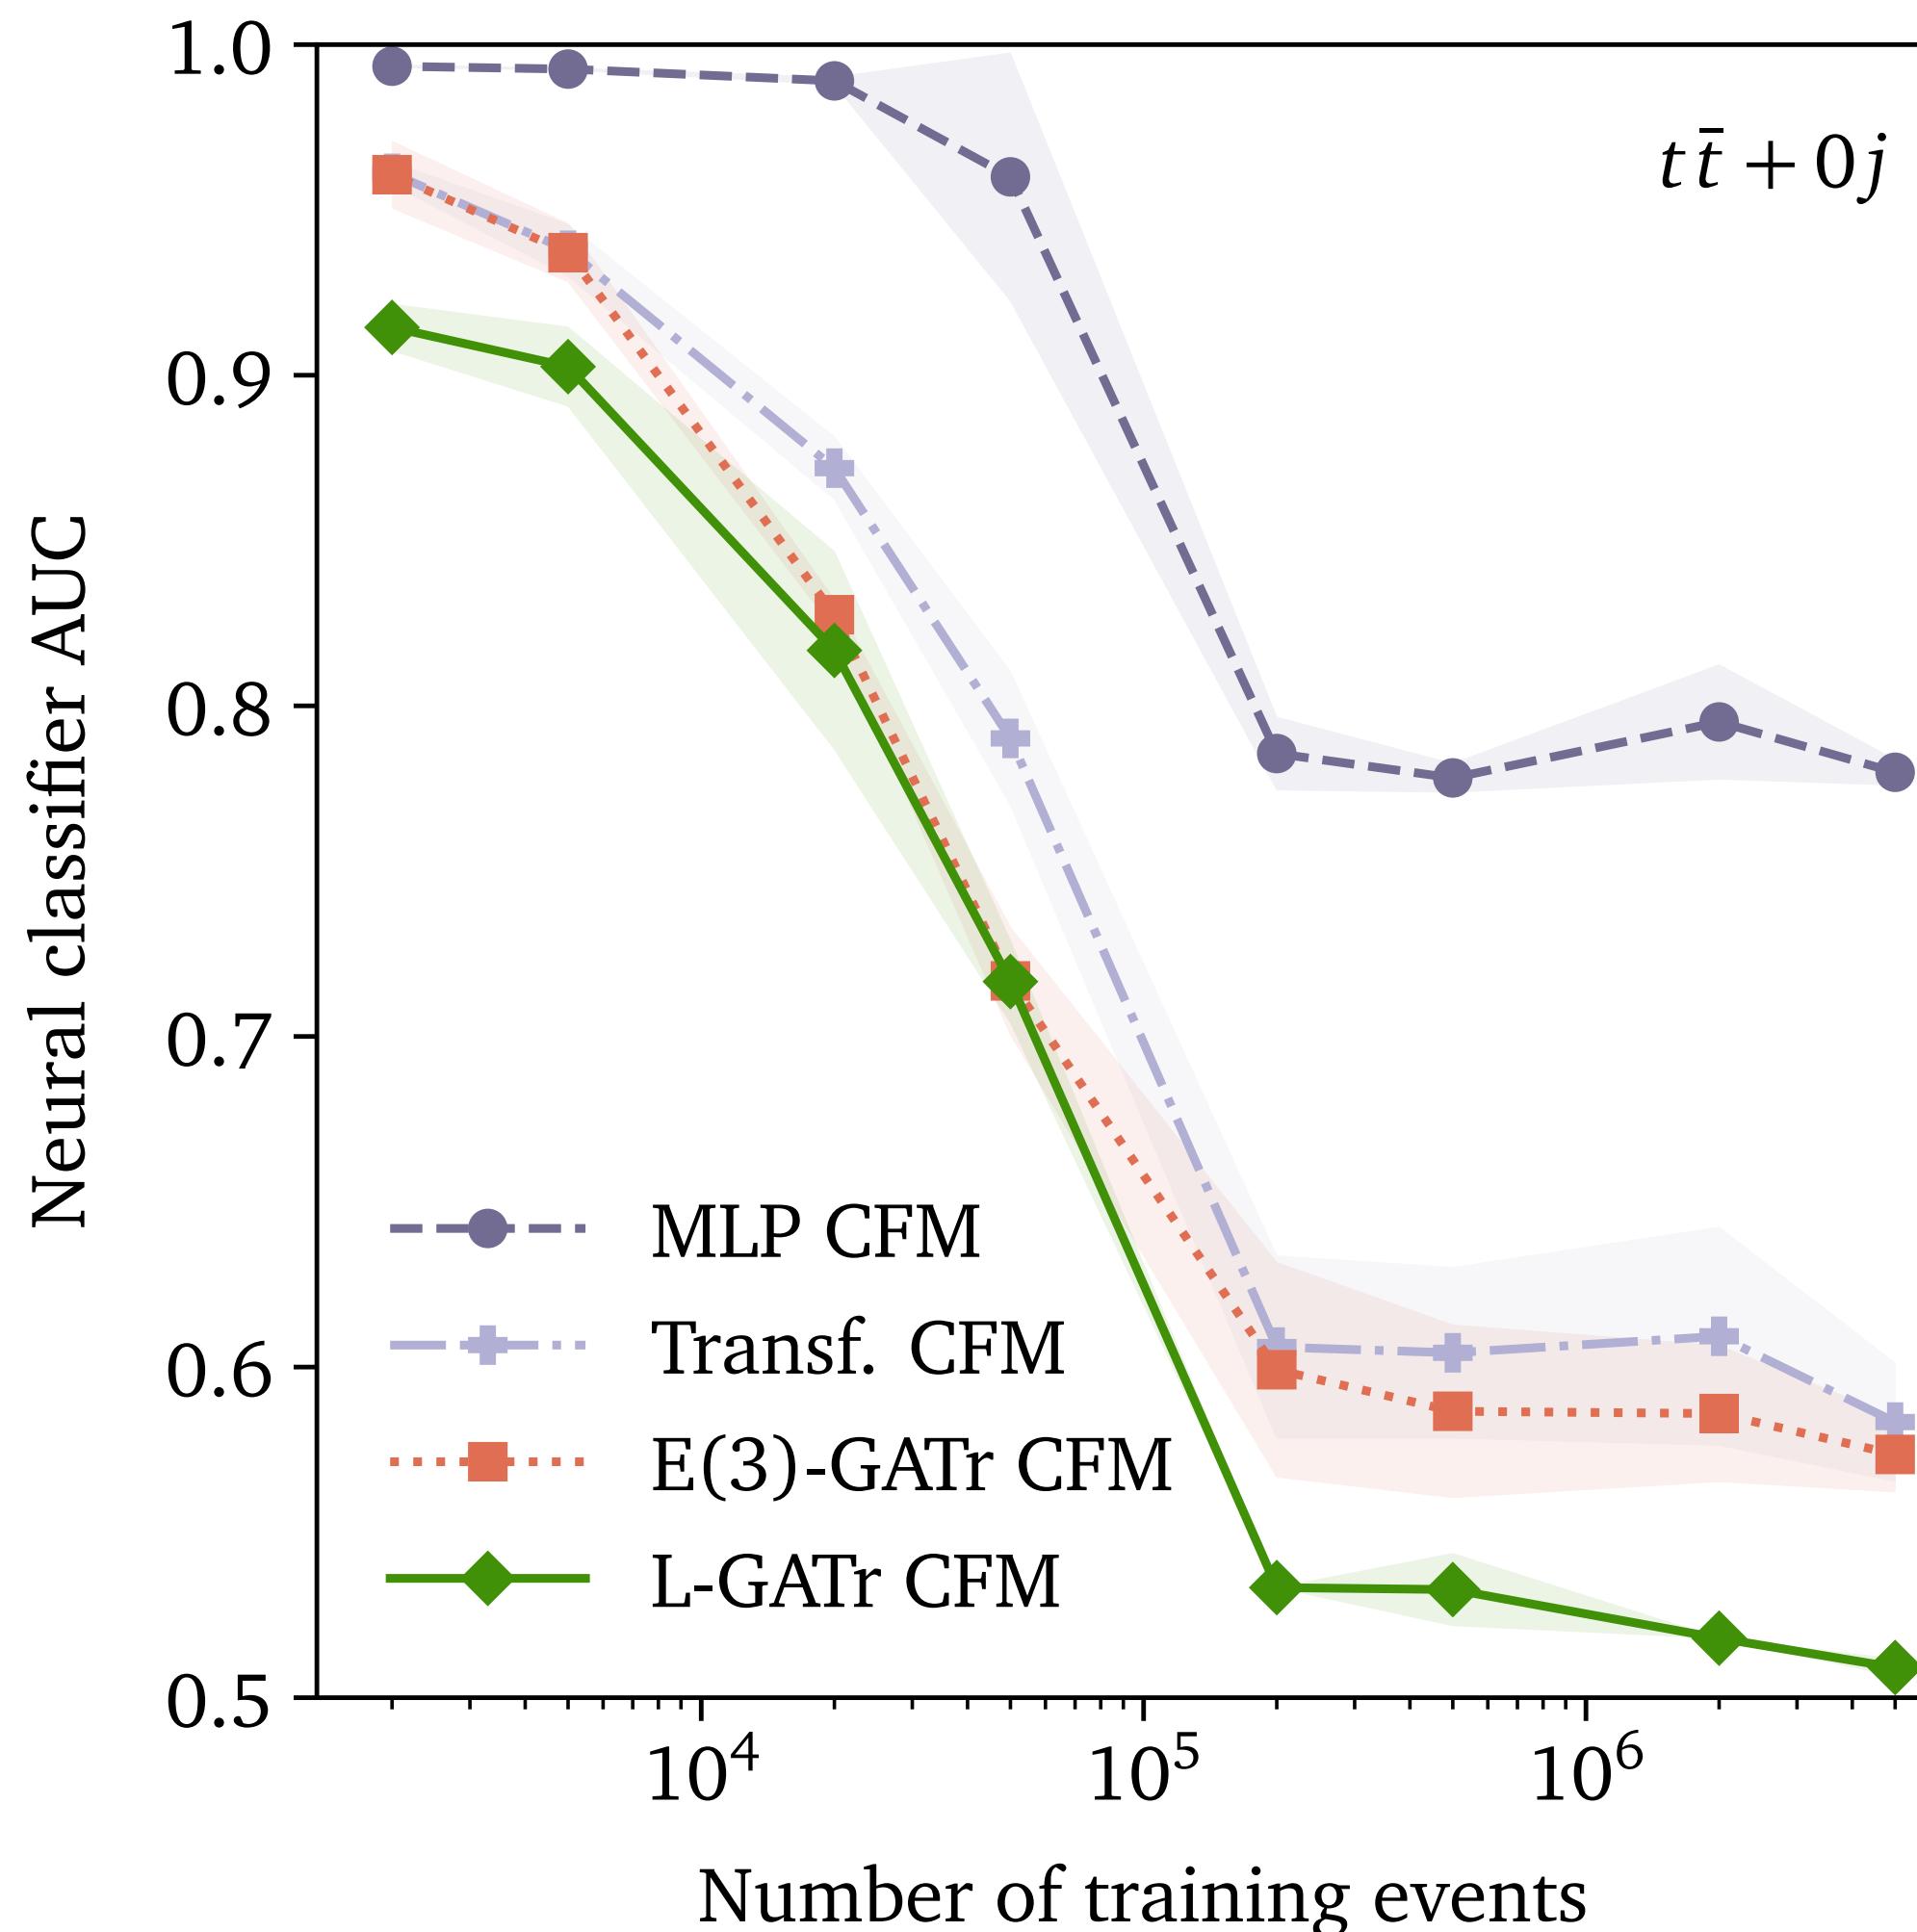


# Event generation



# Event generation

## Boost equivariance lets L-GATr shine



L-GATr and E(3)-GATr are both not boost equivariant:

- E(3)-GATr: No boost equivariance at all
- L-GATr: Boost equivariance broken by reference multivectors

Equiariant networks  
with full symmetry breaking  
outperform non-equivariant networks