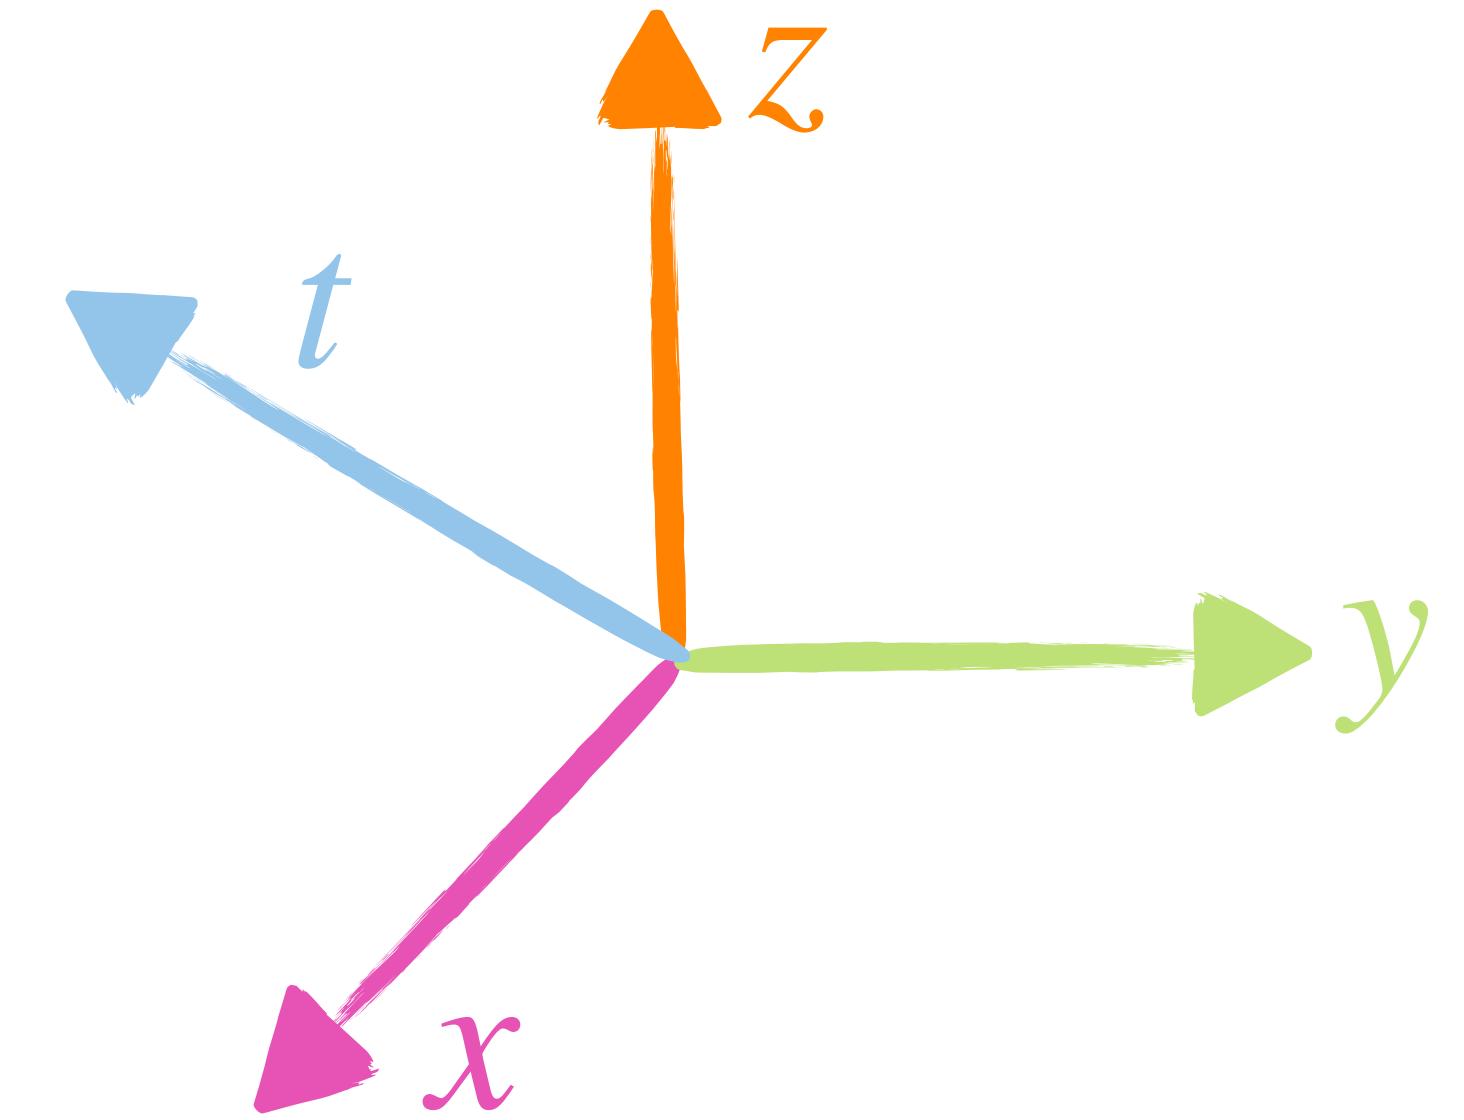


# Lorentz Local Canonicalization

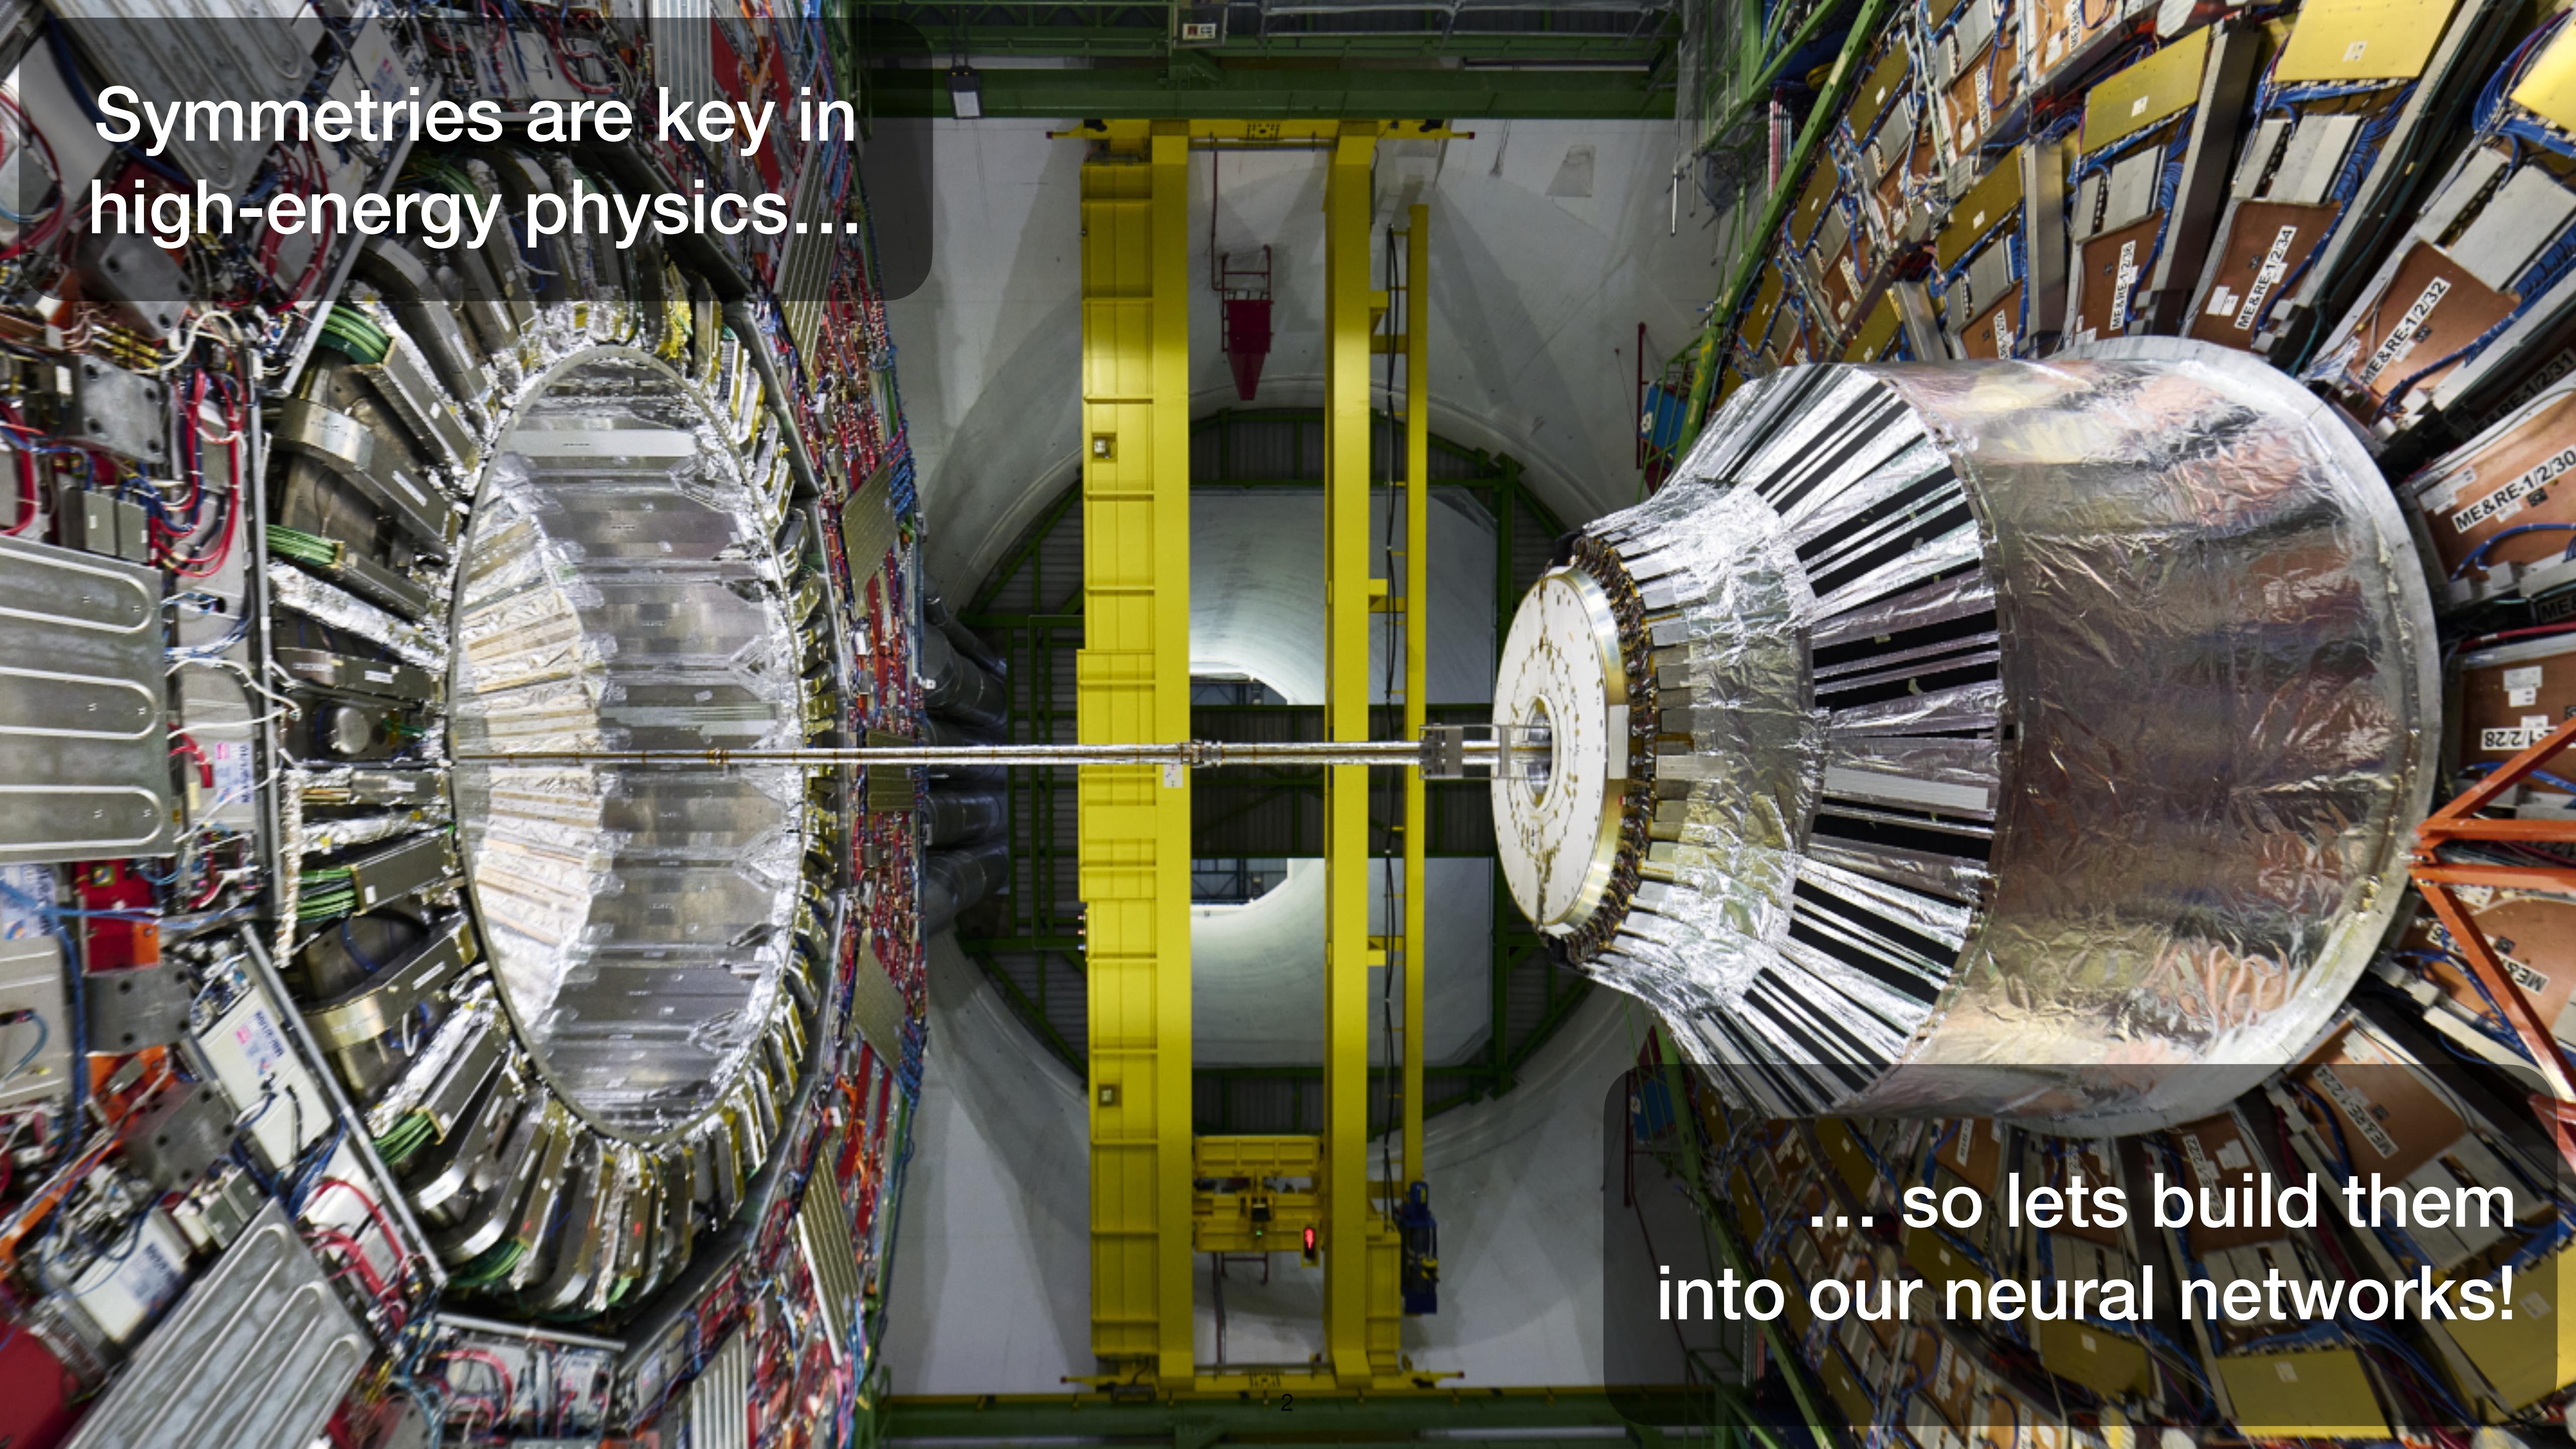
How to make any Network  
Lorentz-equivariant

Jonas Spinner\*, Luigi Favaro\*, Peter Lippmann\*,  
Sebastian Pitz, Gerrit Gerhartz,  
Huilin Qu, Tilman Plehn, Fred A. Hamprecht  
arXiv:2505.20280 [stat.ML]  
arXiv:2508.????? [hep-ph]

IAIFI Summer Workshop  
14.08.2025



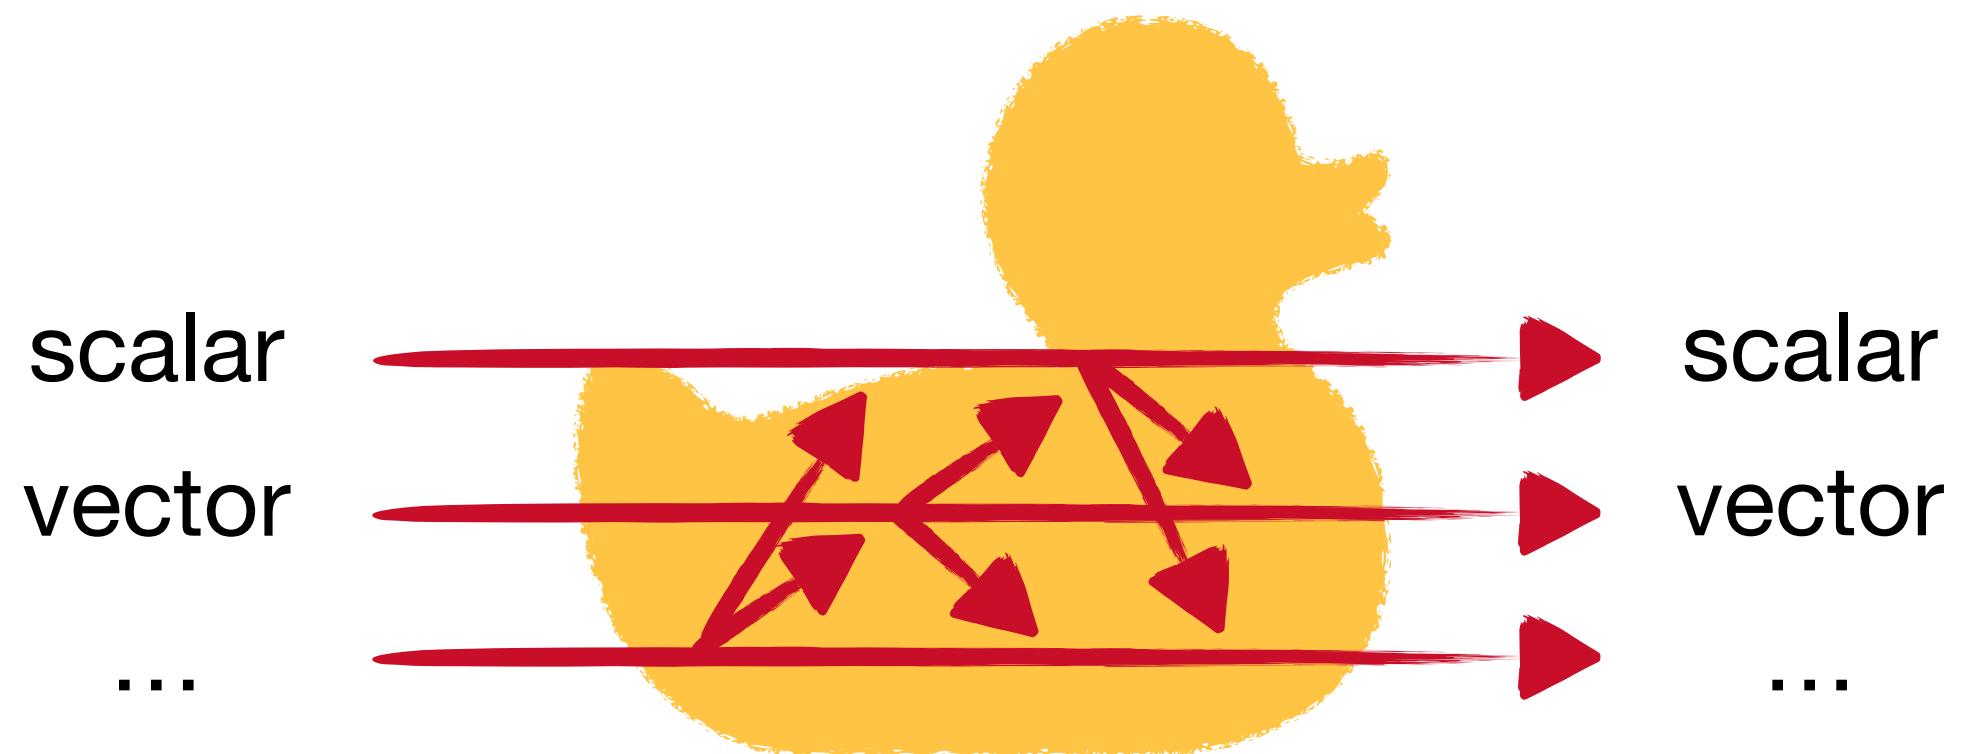
UNIVERSITÄT  
HEIDELBERG  
ZUKUNFT  
SEIT 1386



Symmetries are key in  
high-energy physics...

... so lets build them  
into our neural networks!

# Lorentz-equivariance with specialised layers



**PELICAN**

arXiv:2211.00454  
arXiv:2307.16506

**LorentzNet**

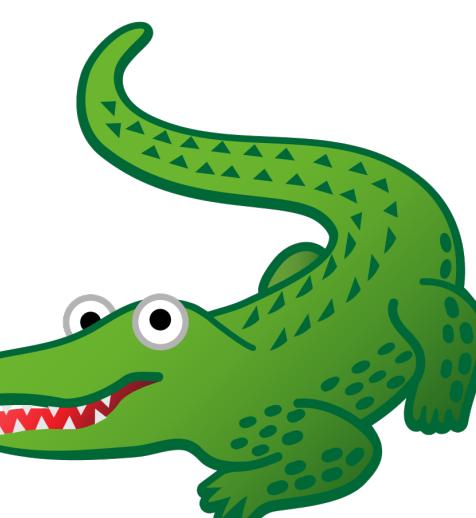
arXiv:2201.08187

**L-GATr**

arXiv:2405.14806

arXiv:2411.00446

pip install lgatr

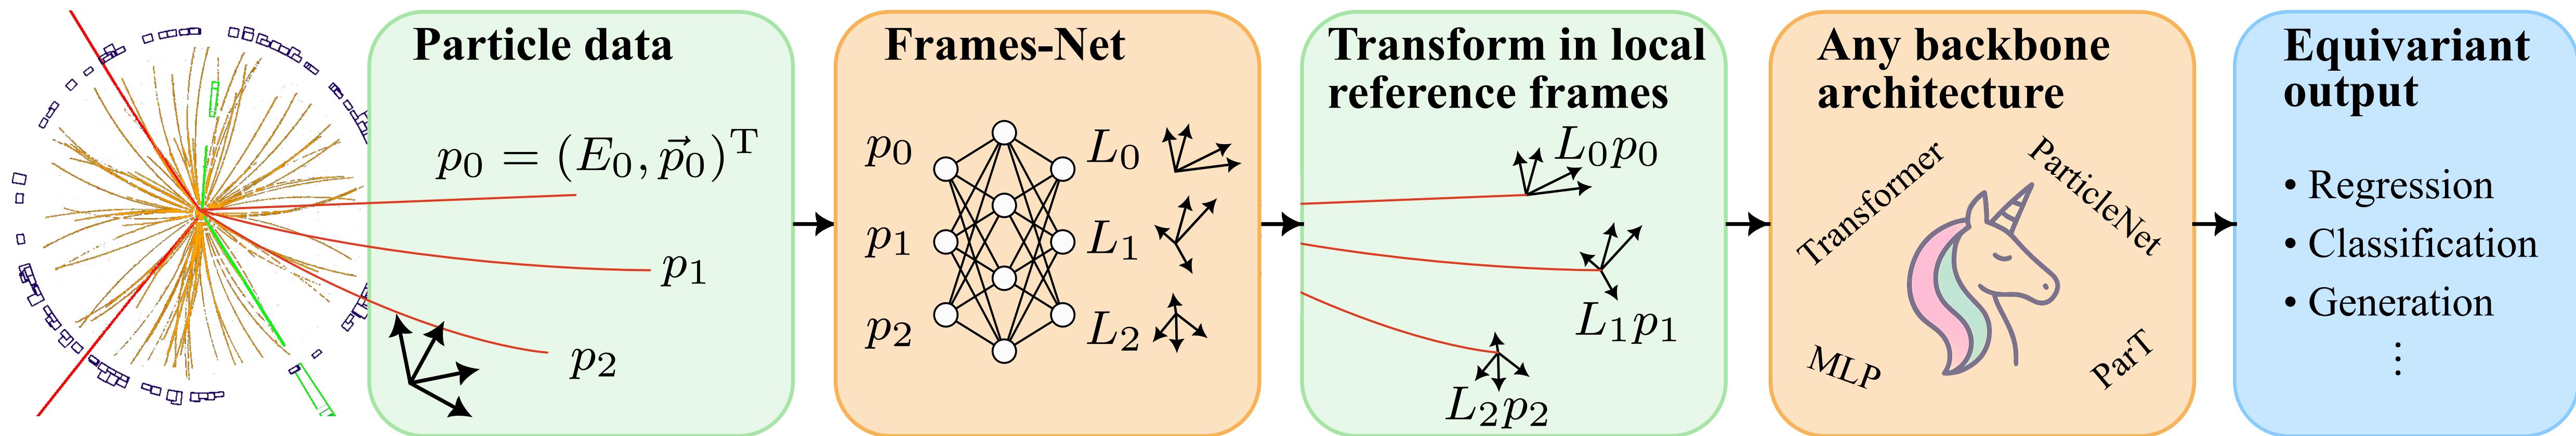


What about **higher-order**  
internal representations?

**Efficient** Lorentz-equivariance?

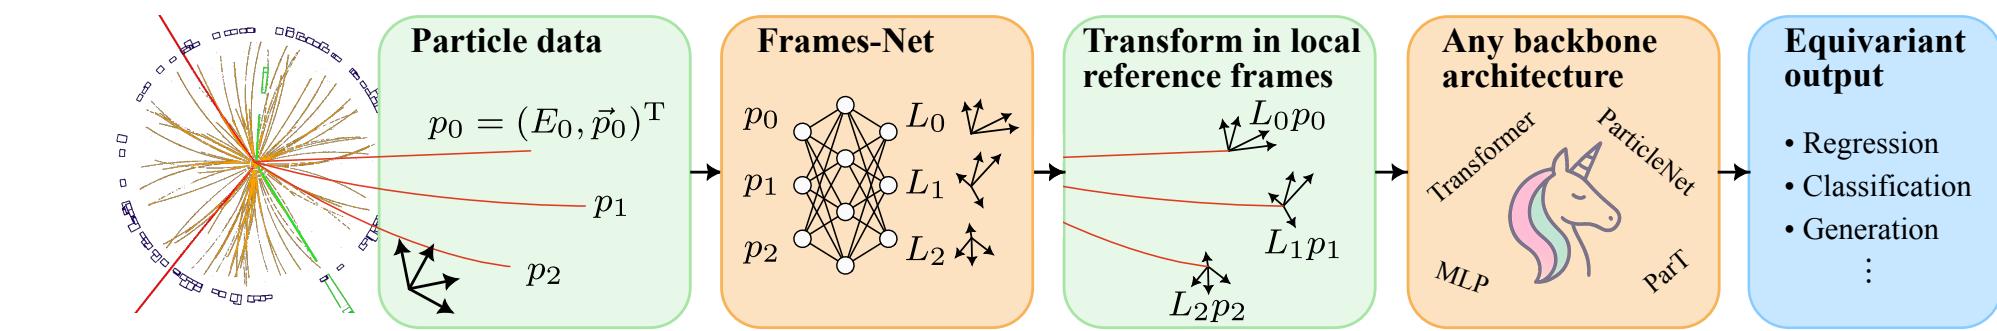
How to make my **domain-specific**  
architecture Lorentz-equivariant?

# Lorentz Local Canonicalization



# LLoCa

## Local frames



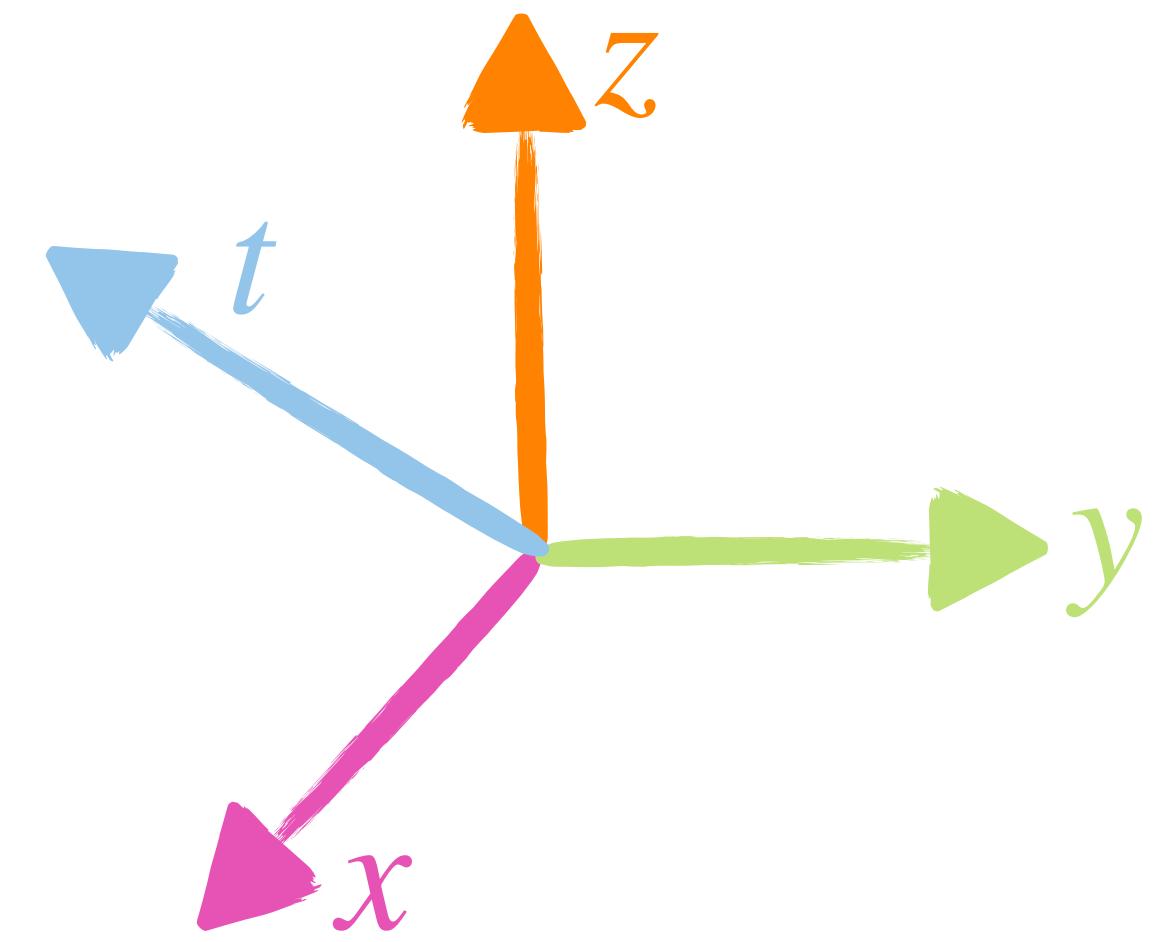
Features  $x_L = Lx$  in local frame are **invariant** by definition:

$$x_L = Lx \xrightarrow{\Lambda} L'x' = L'\Lambda x \stackrel{!}{=} x_L \quad \Rightarrow \quad L' = L\Lambda^{-1}$$

Achieve this by stacking **Lorentz vectors**  $u_a^\mu$ ,  $a = 0,1,2,3$ :

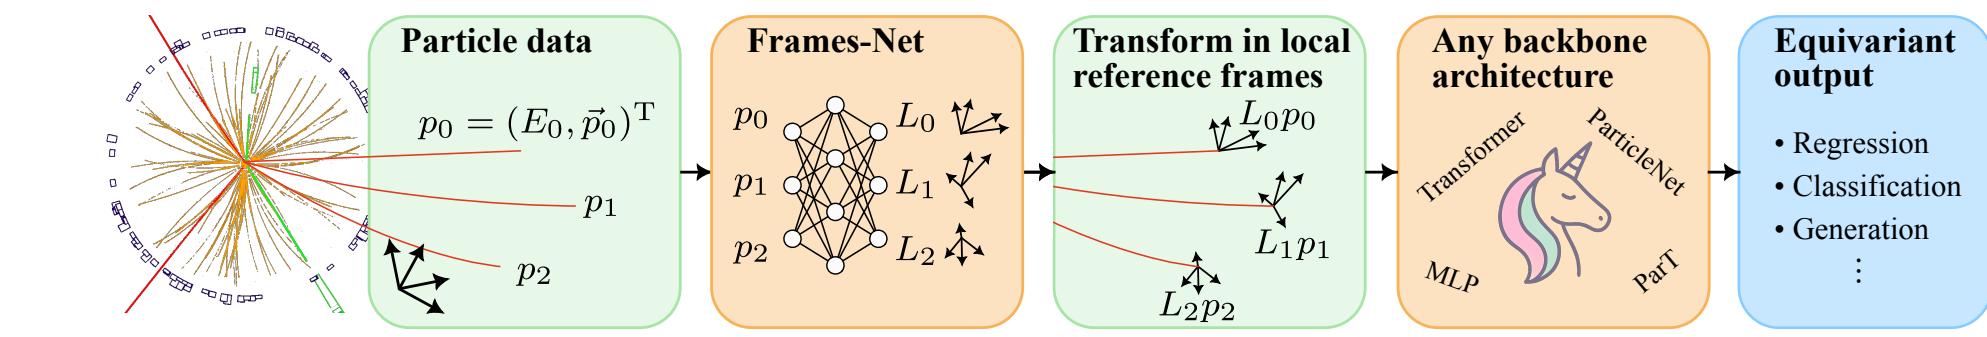
$$L = \begin{pmatrix} u_0^T g \\ u_1^T g \\ u_2^T g \\ u_3^T g \end{pmatrix} \xrightarrow{\Lambda} L' = \begin{pmatrix} u_0^T \Lambda^T g \\ u_1^T \Lambda^T g \\ u_2^T \Lambda^T g \\ u_3^T \Lambda^T g \end{pmatrix} = L\Lambda^{-1}$$

$$\sum_{\mu, \nu} u_a^\mu u_b^\nu g_{\mu\nu} = g_{ab}$$



# LLoCa

## Equivariance from canonicalisation



Latent features in the local frame  $f_L = \rho(L)f$  are **invariant**

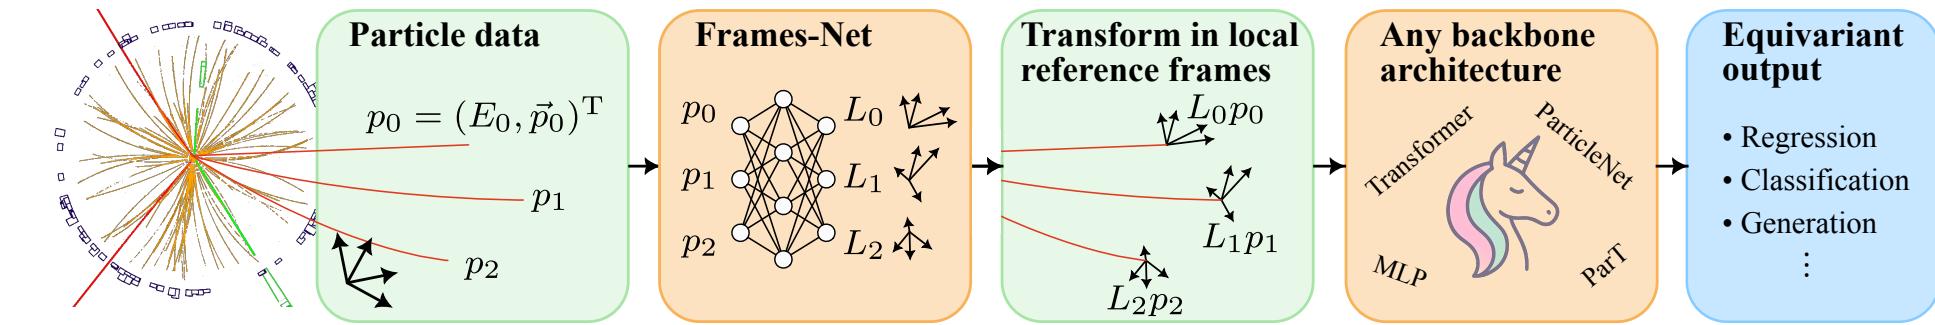
$$f_L \xrightarrow{\Lambda} f'_L = \rho(L')f' = \rho(L\Lambda^{-1})\rho(\Lambda)f = \rho(L\Lambda^{-1}\Lambda)f = \rho(L)f = f_L$$

Output features in the global frame  $y = \rho(L^{-1})y_L$  are **equivariant**

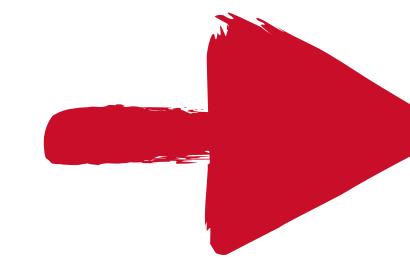
$$y \xrightarrow{\Lambda} y' = \rho(L'^{-1})y'_L = \rho(\Lambda L^{-1})y_L = \rho(\Lambda)\rho(L^{-1})y_L = \rho(\Lambda)y$$

# LLoCa

## Message-passing between local frames



**Standard message-passing**

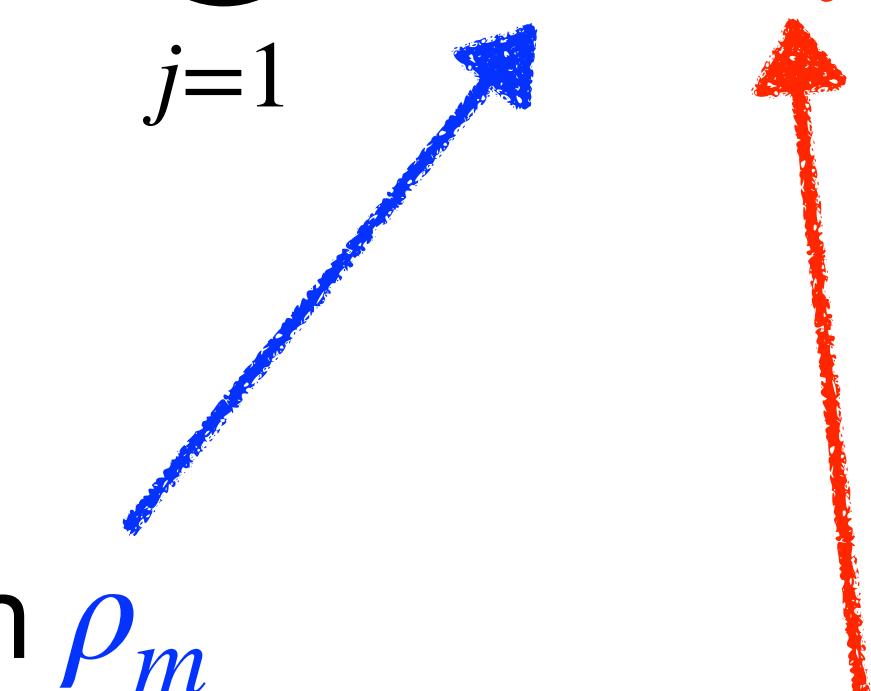


**Tensorial message-passing**  
between distinct local frames

$$f_i^{\text{updated}} = \psi\left(f_i, \bigoplus_{j=1}^N \phi(m_j)\right)$$

$$f_{i,\mathbf{L}_i}^{\text{updated}} = \psi\left(f_{i,\mathbf{L}_i}, \bigoplus_{j=1}^N \phi\left(\rho_m(\mathbf{L}_i \mathbf{L}_j^{-1}) m_{j,\mathbf{L}_j}\right)\right)$$

Message representation  $\rho_m$   
is a hyperparameter



Invariant frame  
transformation matrix  $\mathbf{L}_i \mathbf{L}_j^{-1}$

# LLoCa

## How to construct local frames

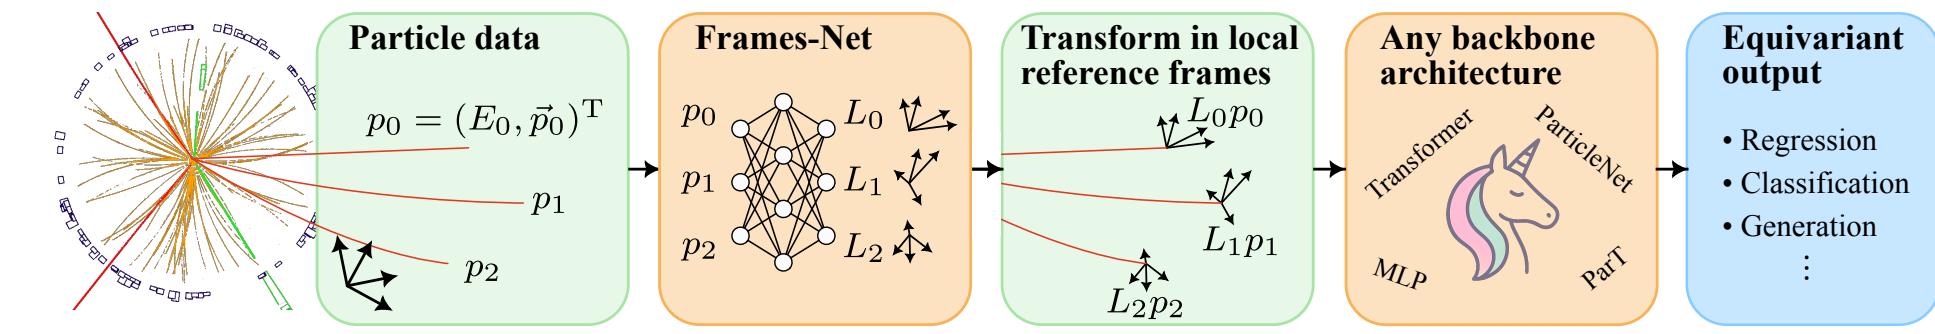
1) Predict vectors  $v_{i,k}$  with simple **Lorentz-equivariant network**

$$v_{i,k} = \sum_{j=1}^N \text{softmax}\left(\varphi_k(s_i, s_j, \langle p_i, p_j \rangle)\right)(p_i + p_j) \quad \text{for } k = 0, 1, 2.$$

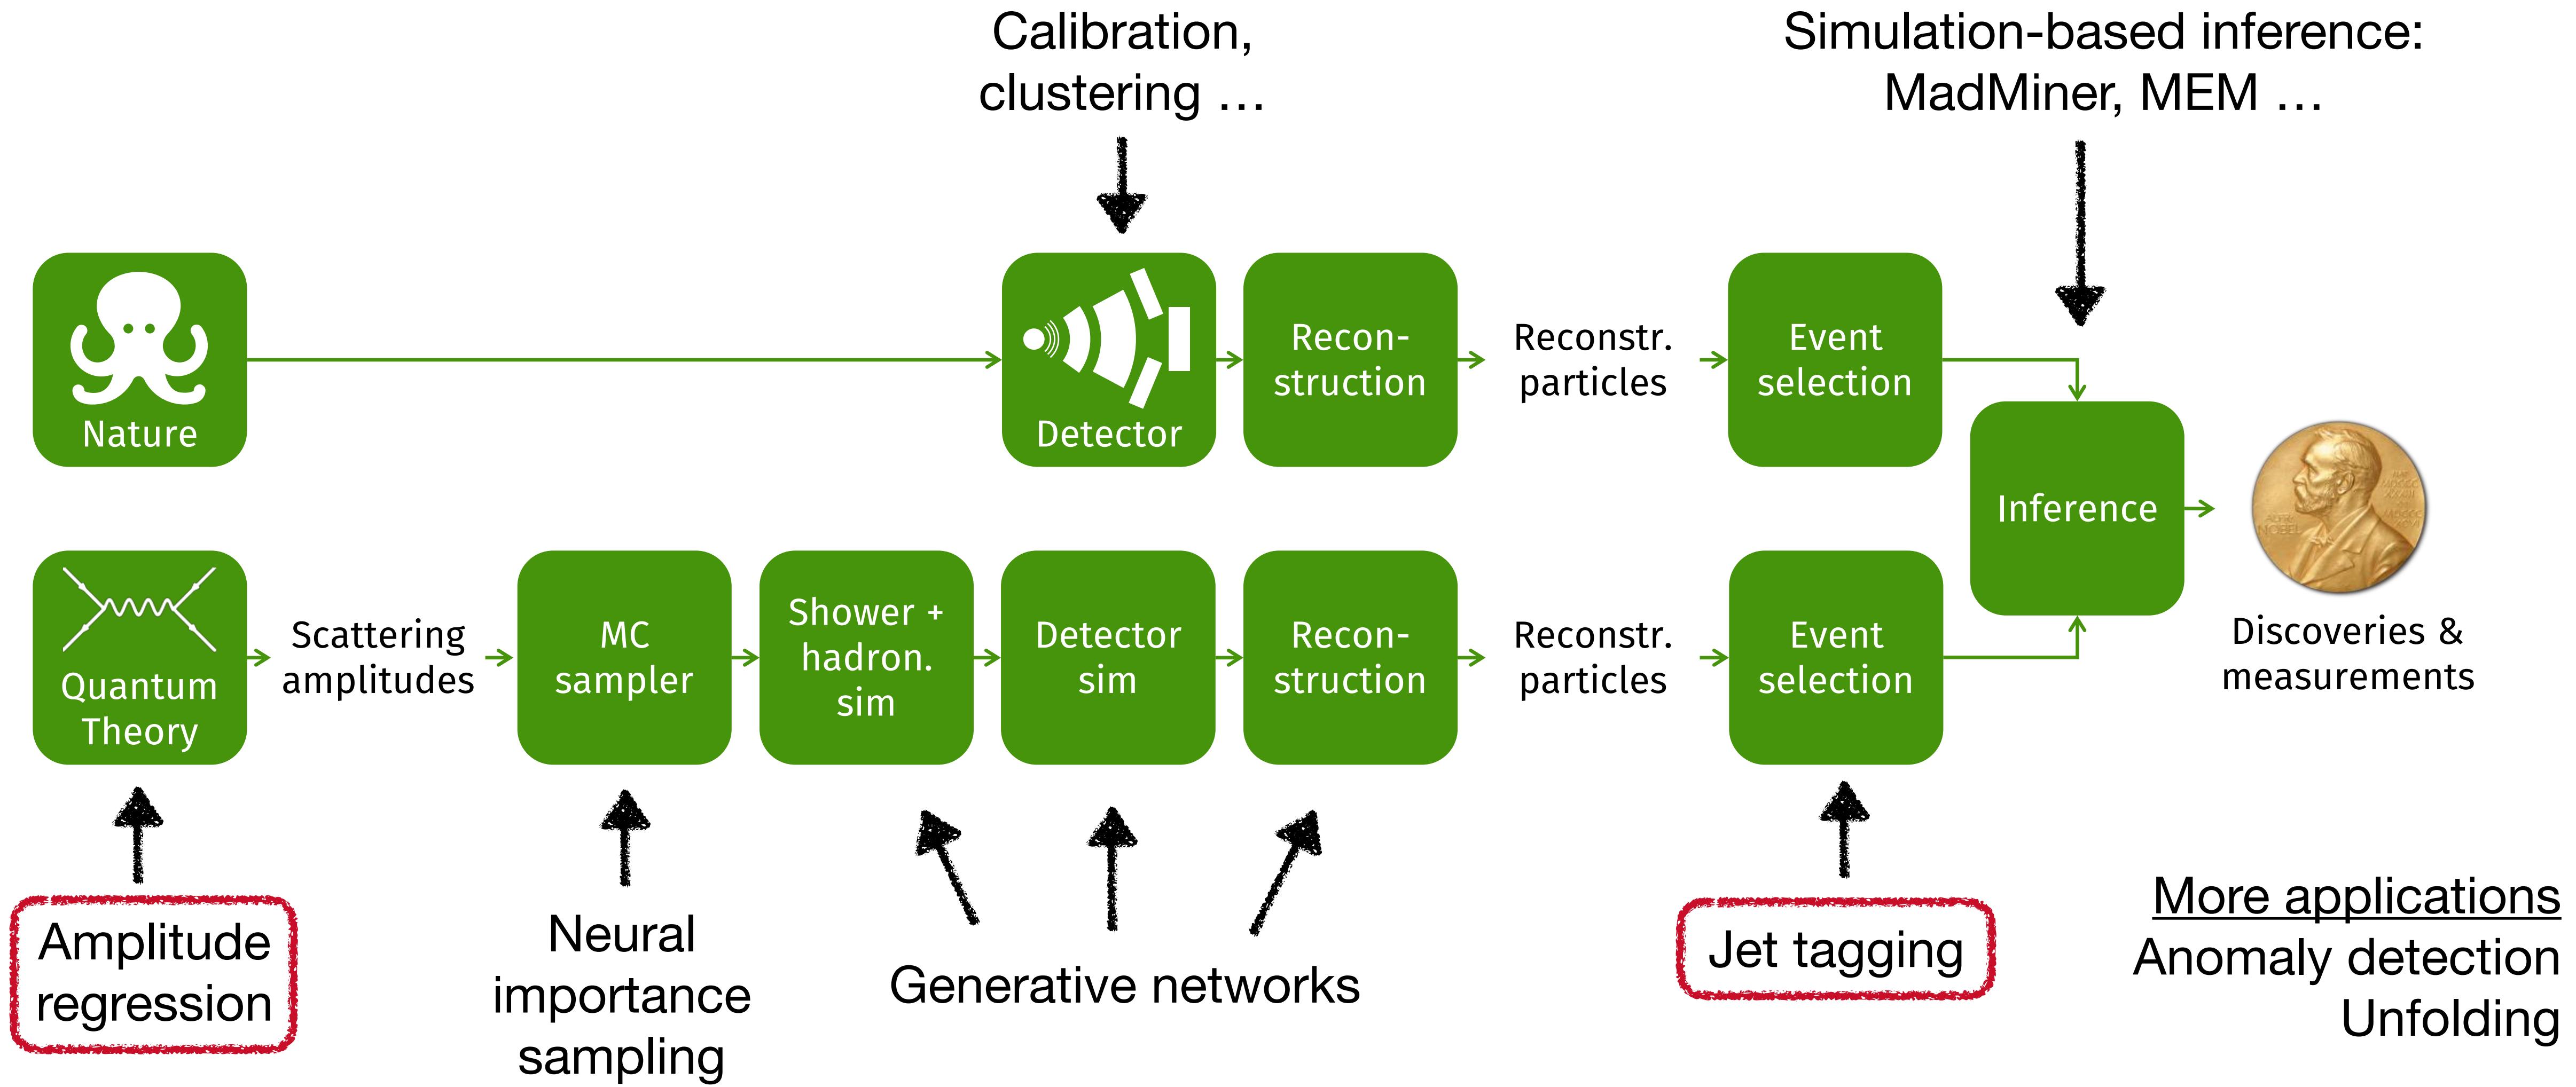


Frames-Net (MLP)

2) Orthonormalize vectors with **Gram-Schmidt algorithm** to obtain  $u_{i,a}$

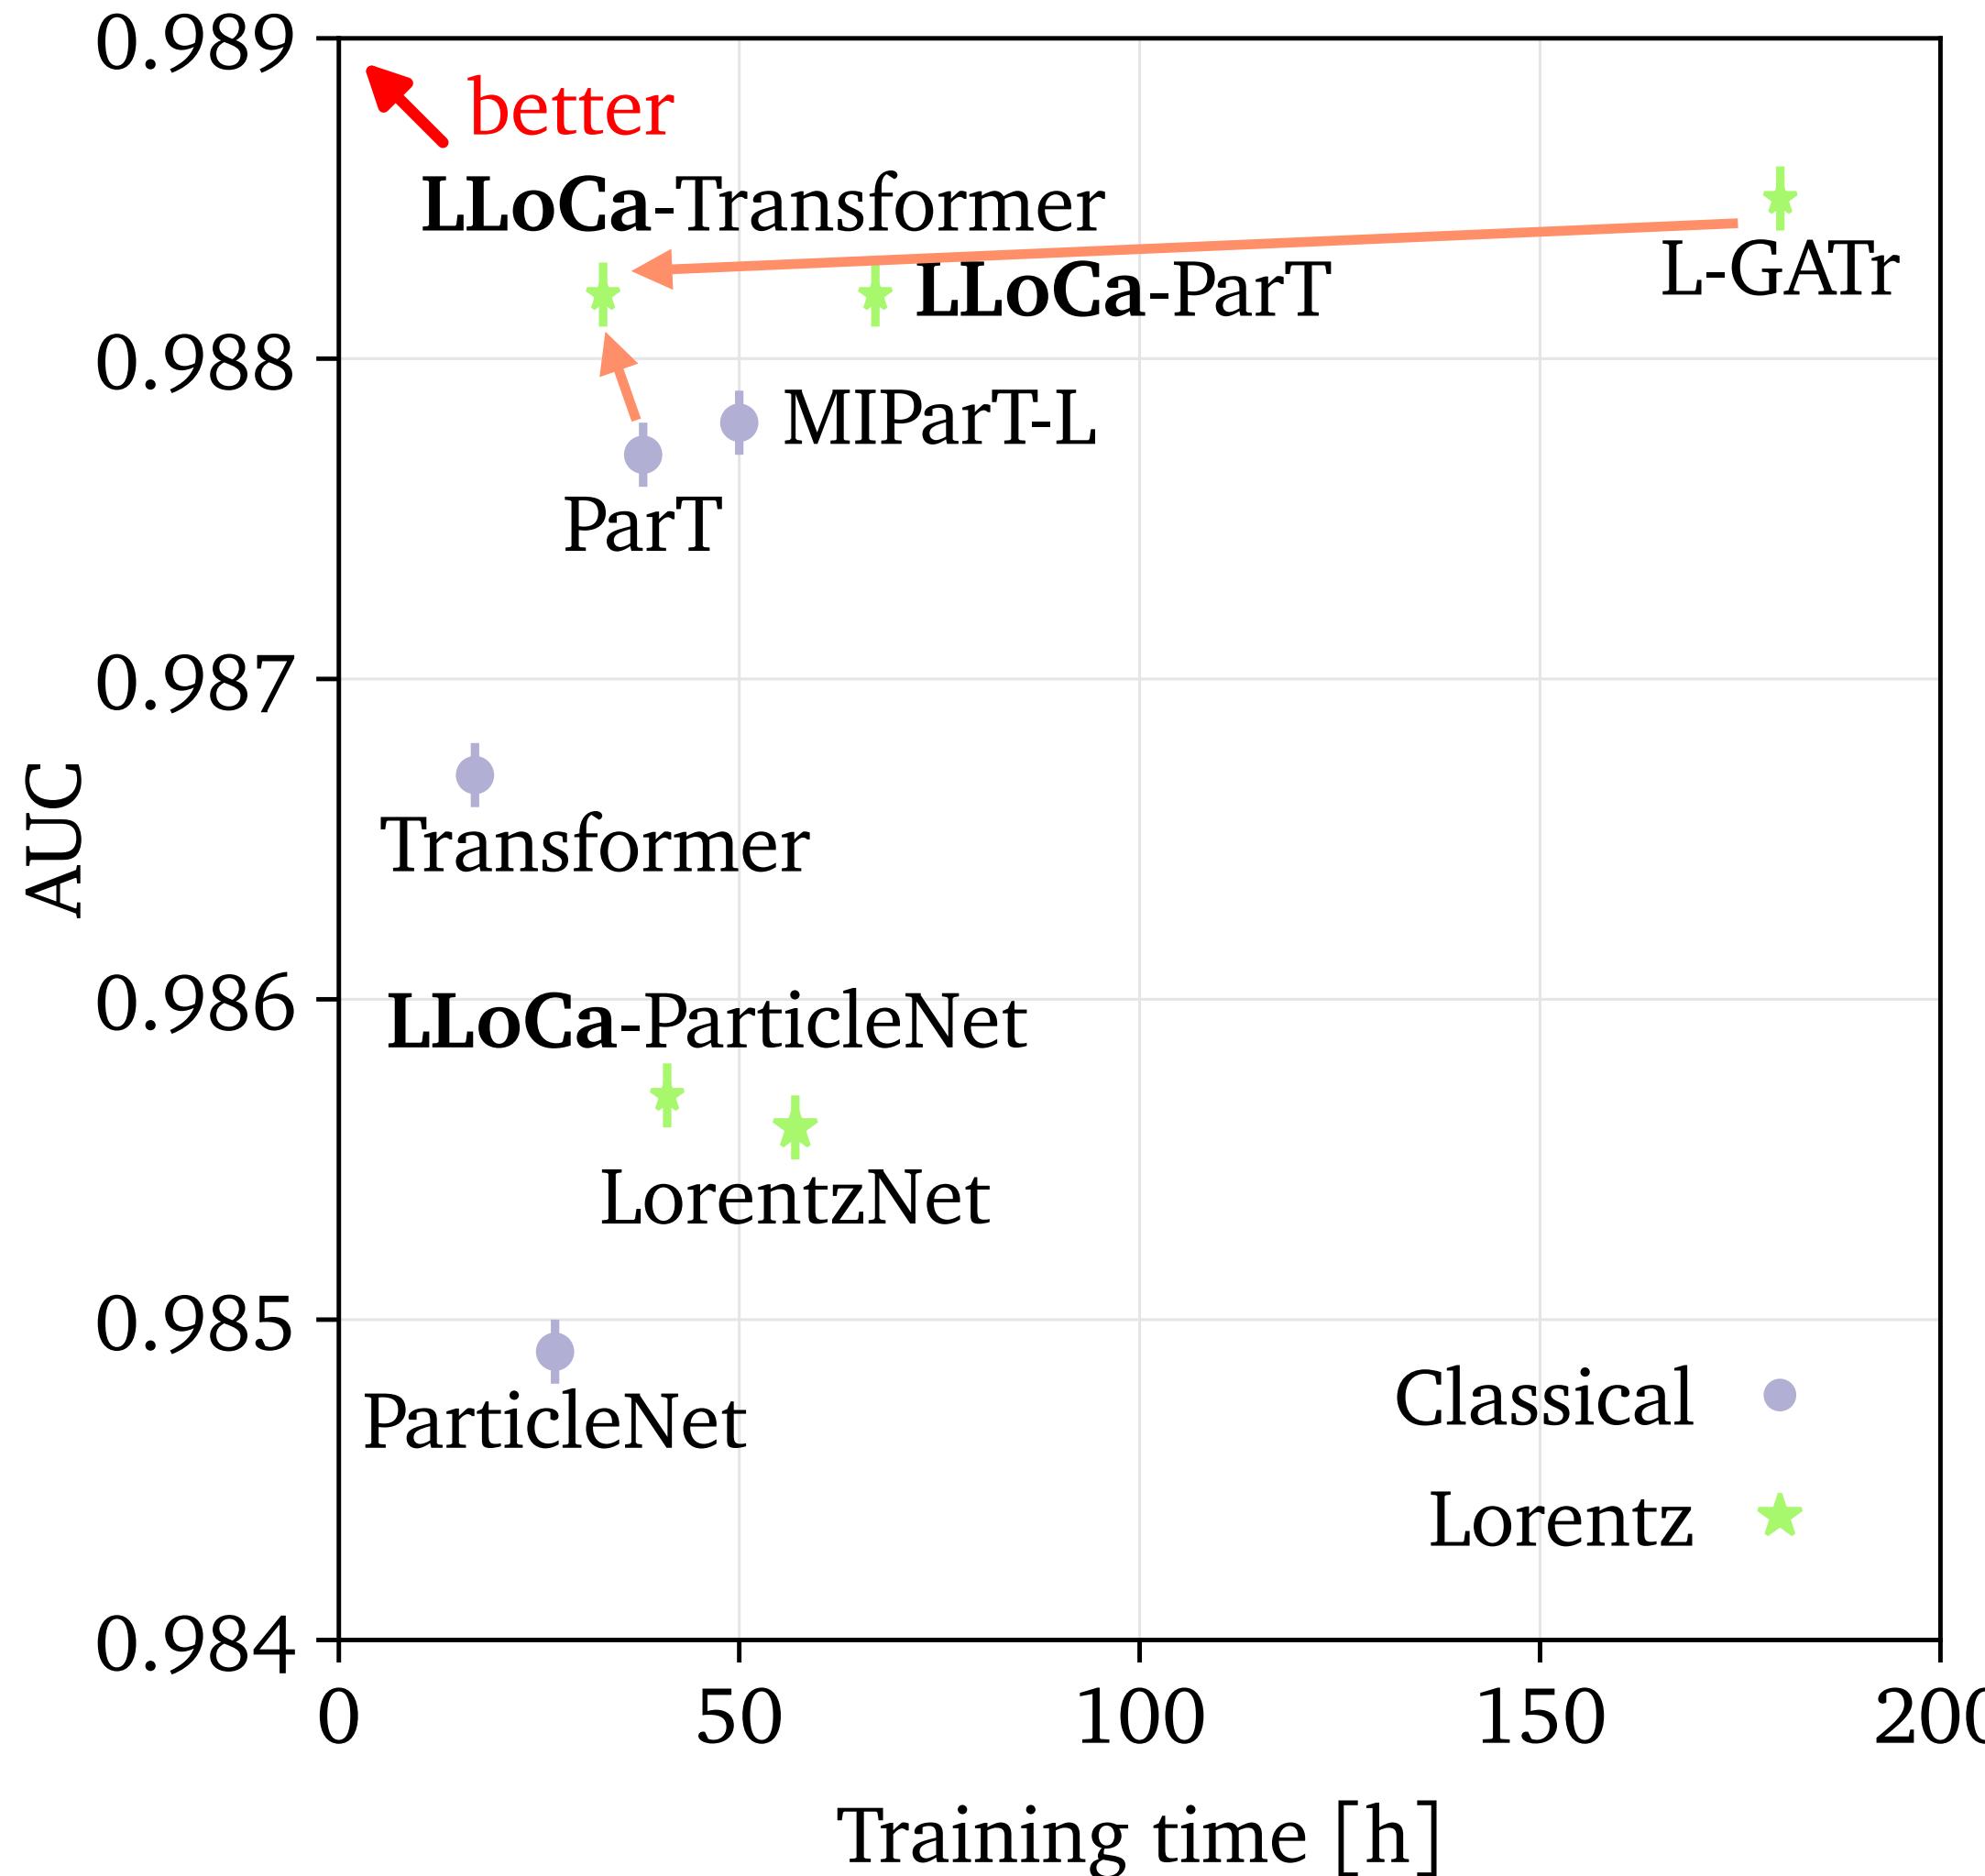


# The LHC ML workflow



# Jet Tagging

## Speed vs performance on JetClass

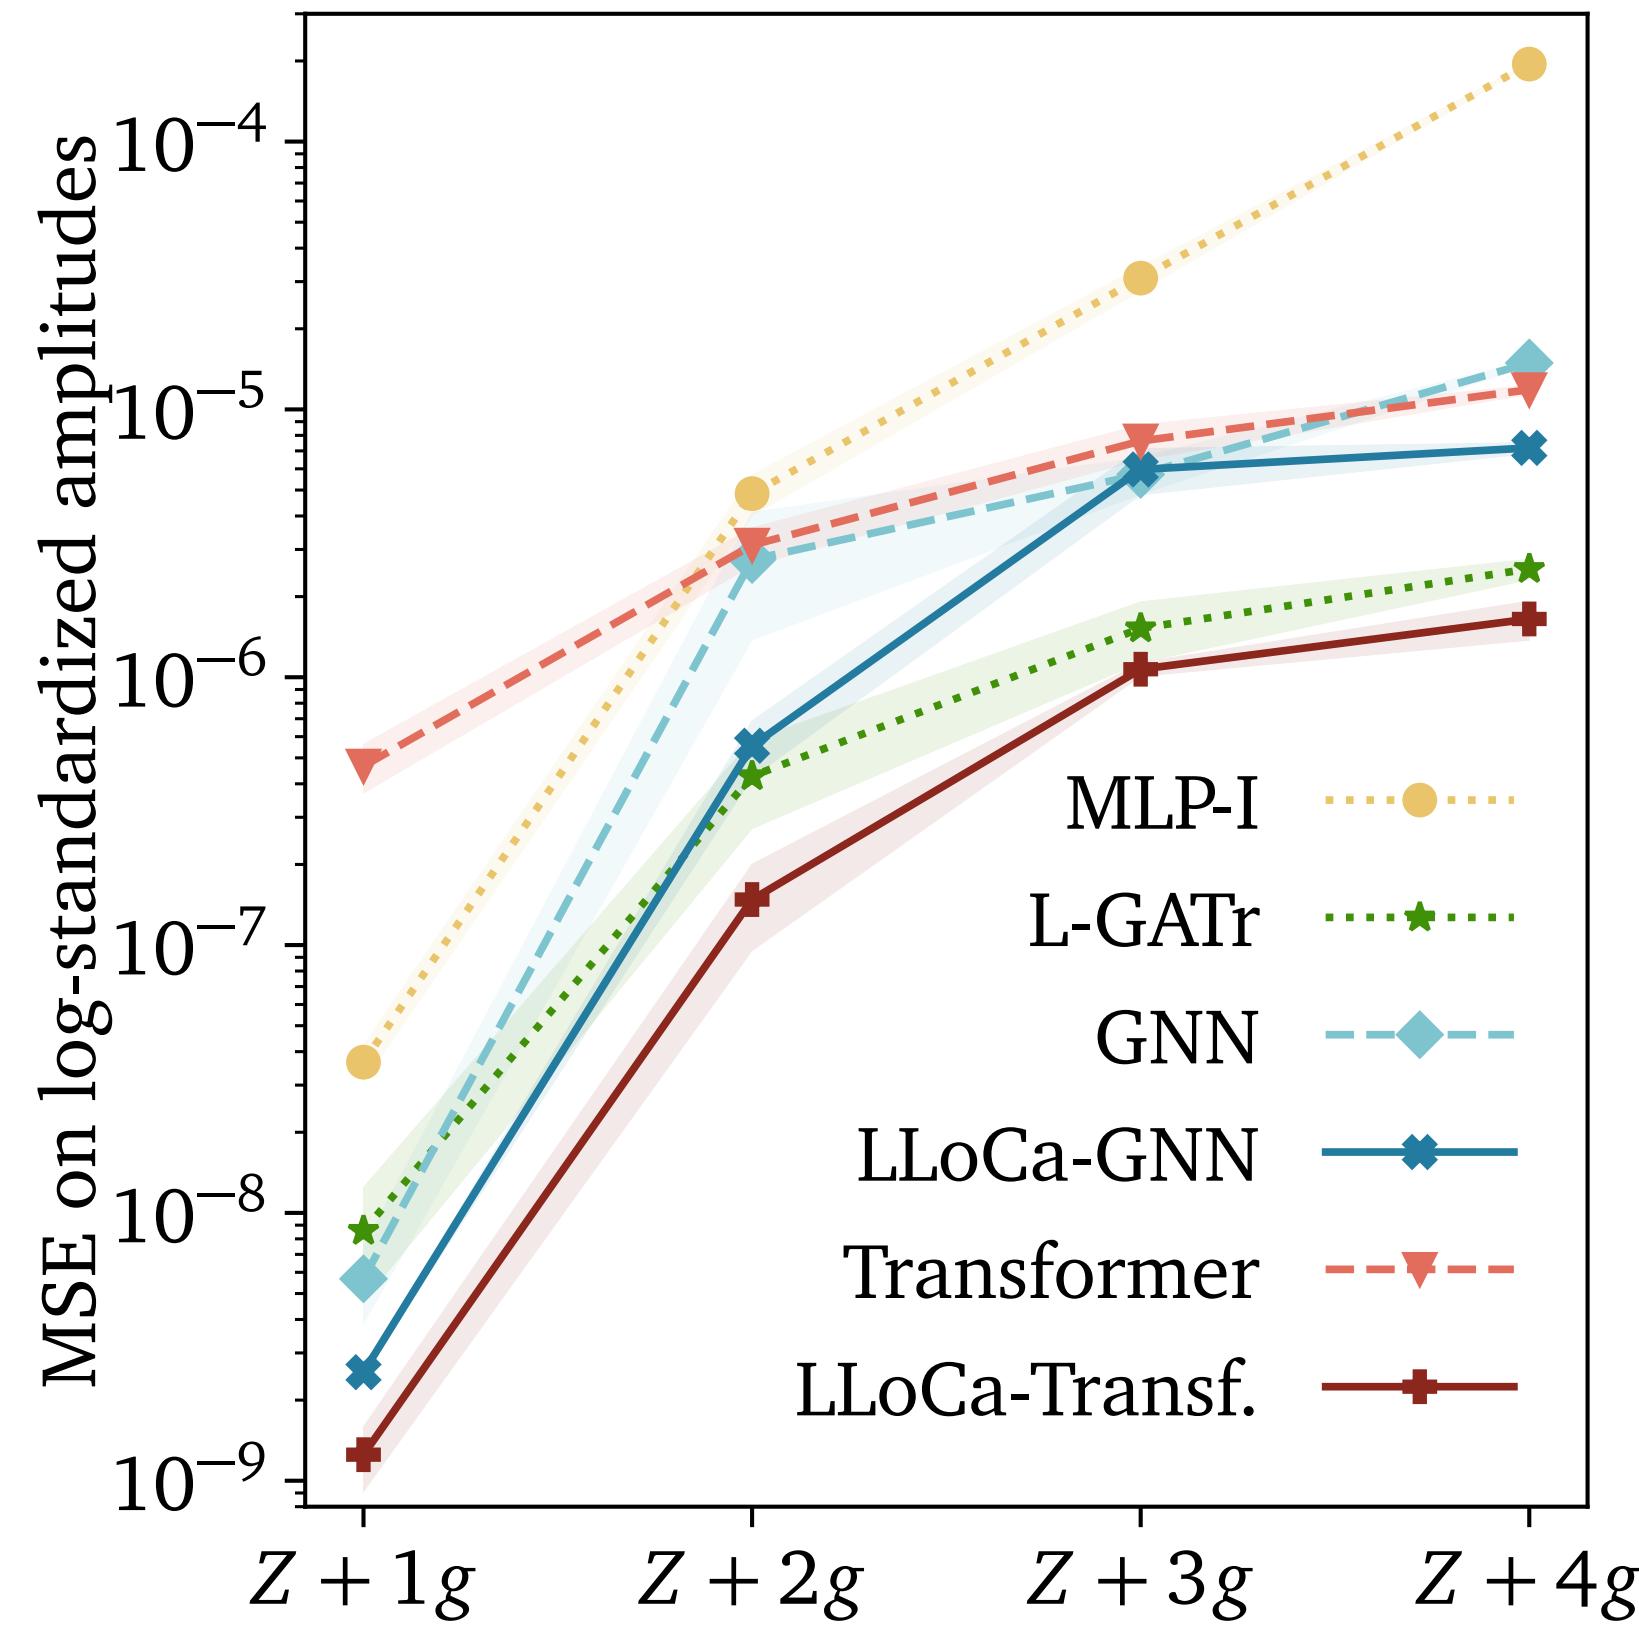


LLoCa-Transformer outperforms ParT while running 10% faster

LLoCa-Transformer matches L-GATr performance at 6x speedup

# Amplitude regression

## Speed vs performance



LLoCa-Transformer and L-GATr  
achieve similar performance

Model	MSE $\times 10^{-6}$ ( $\downarrow$ )	FLOPs	Time
MLP-I [38]	195.0 $\pm$ 2	0.1M	0.6h
L-GATr [38]	2.5 $\pm$ 0.2	1160.0M	19.5h
GNN	14.9 $\pm$ 0.3	9.9M	1.3h
LLoCa-GNN	7.2 $\pm$ 0.3	10.6M	2.3h
Transformer	11.9 $\pm$ 0.5	10.6M	3.5h
LLoCa-Transformer	1.5 $\pm$ 0.1	11.3M	4.5h

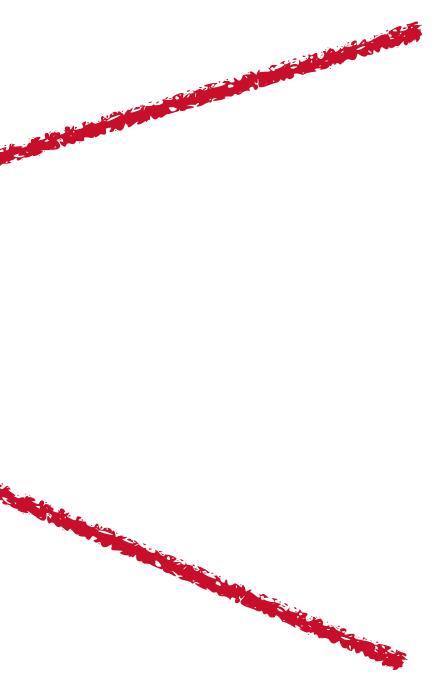
... but LLoCa-Transformer is 4x faster  
and requires 100x less FLOPs!

# Amplitude regression

Effect of internal representation  $\rho_m$

Method	MSE ( $\times 10^{-6}$ )
Non-equivariant	$11.9 \pm 0.8$
Global canonical.	$5.9 \pm 0.4$
LLoCa (16 scalars)	$54.0 \pm 3$
LLoCa (single 2-tensor)	$2.4 \pm 0.3$
LLoCa (4 vectors)	$2.0 \pm 0.2$
LLoCa (8 scalars, 2 vectors)	<b><math>1.5 \pm 0.1</math></b>

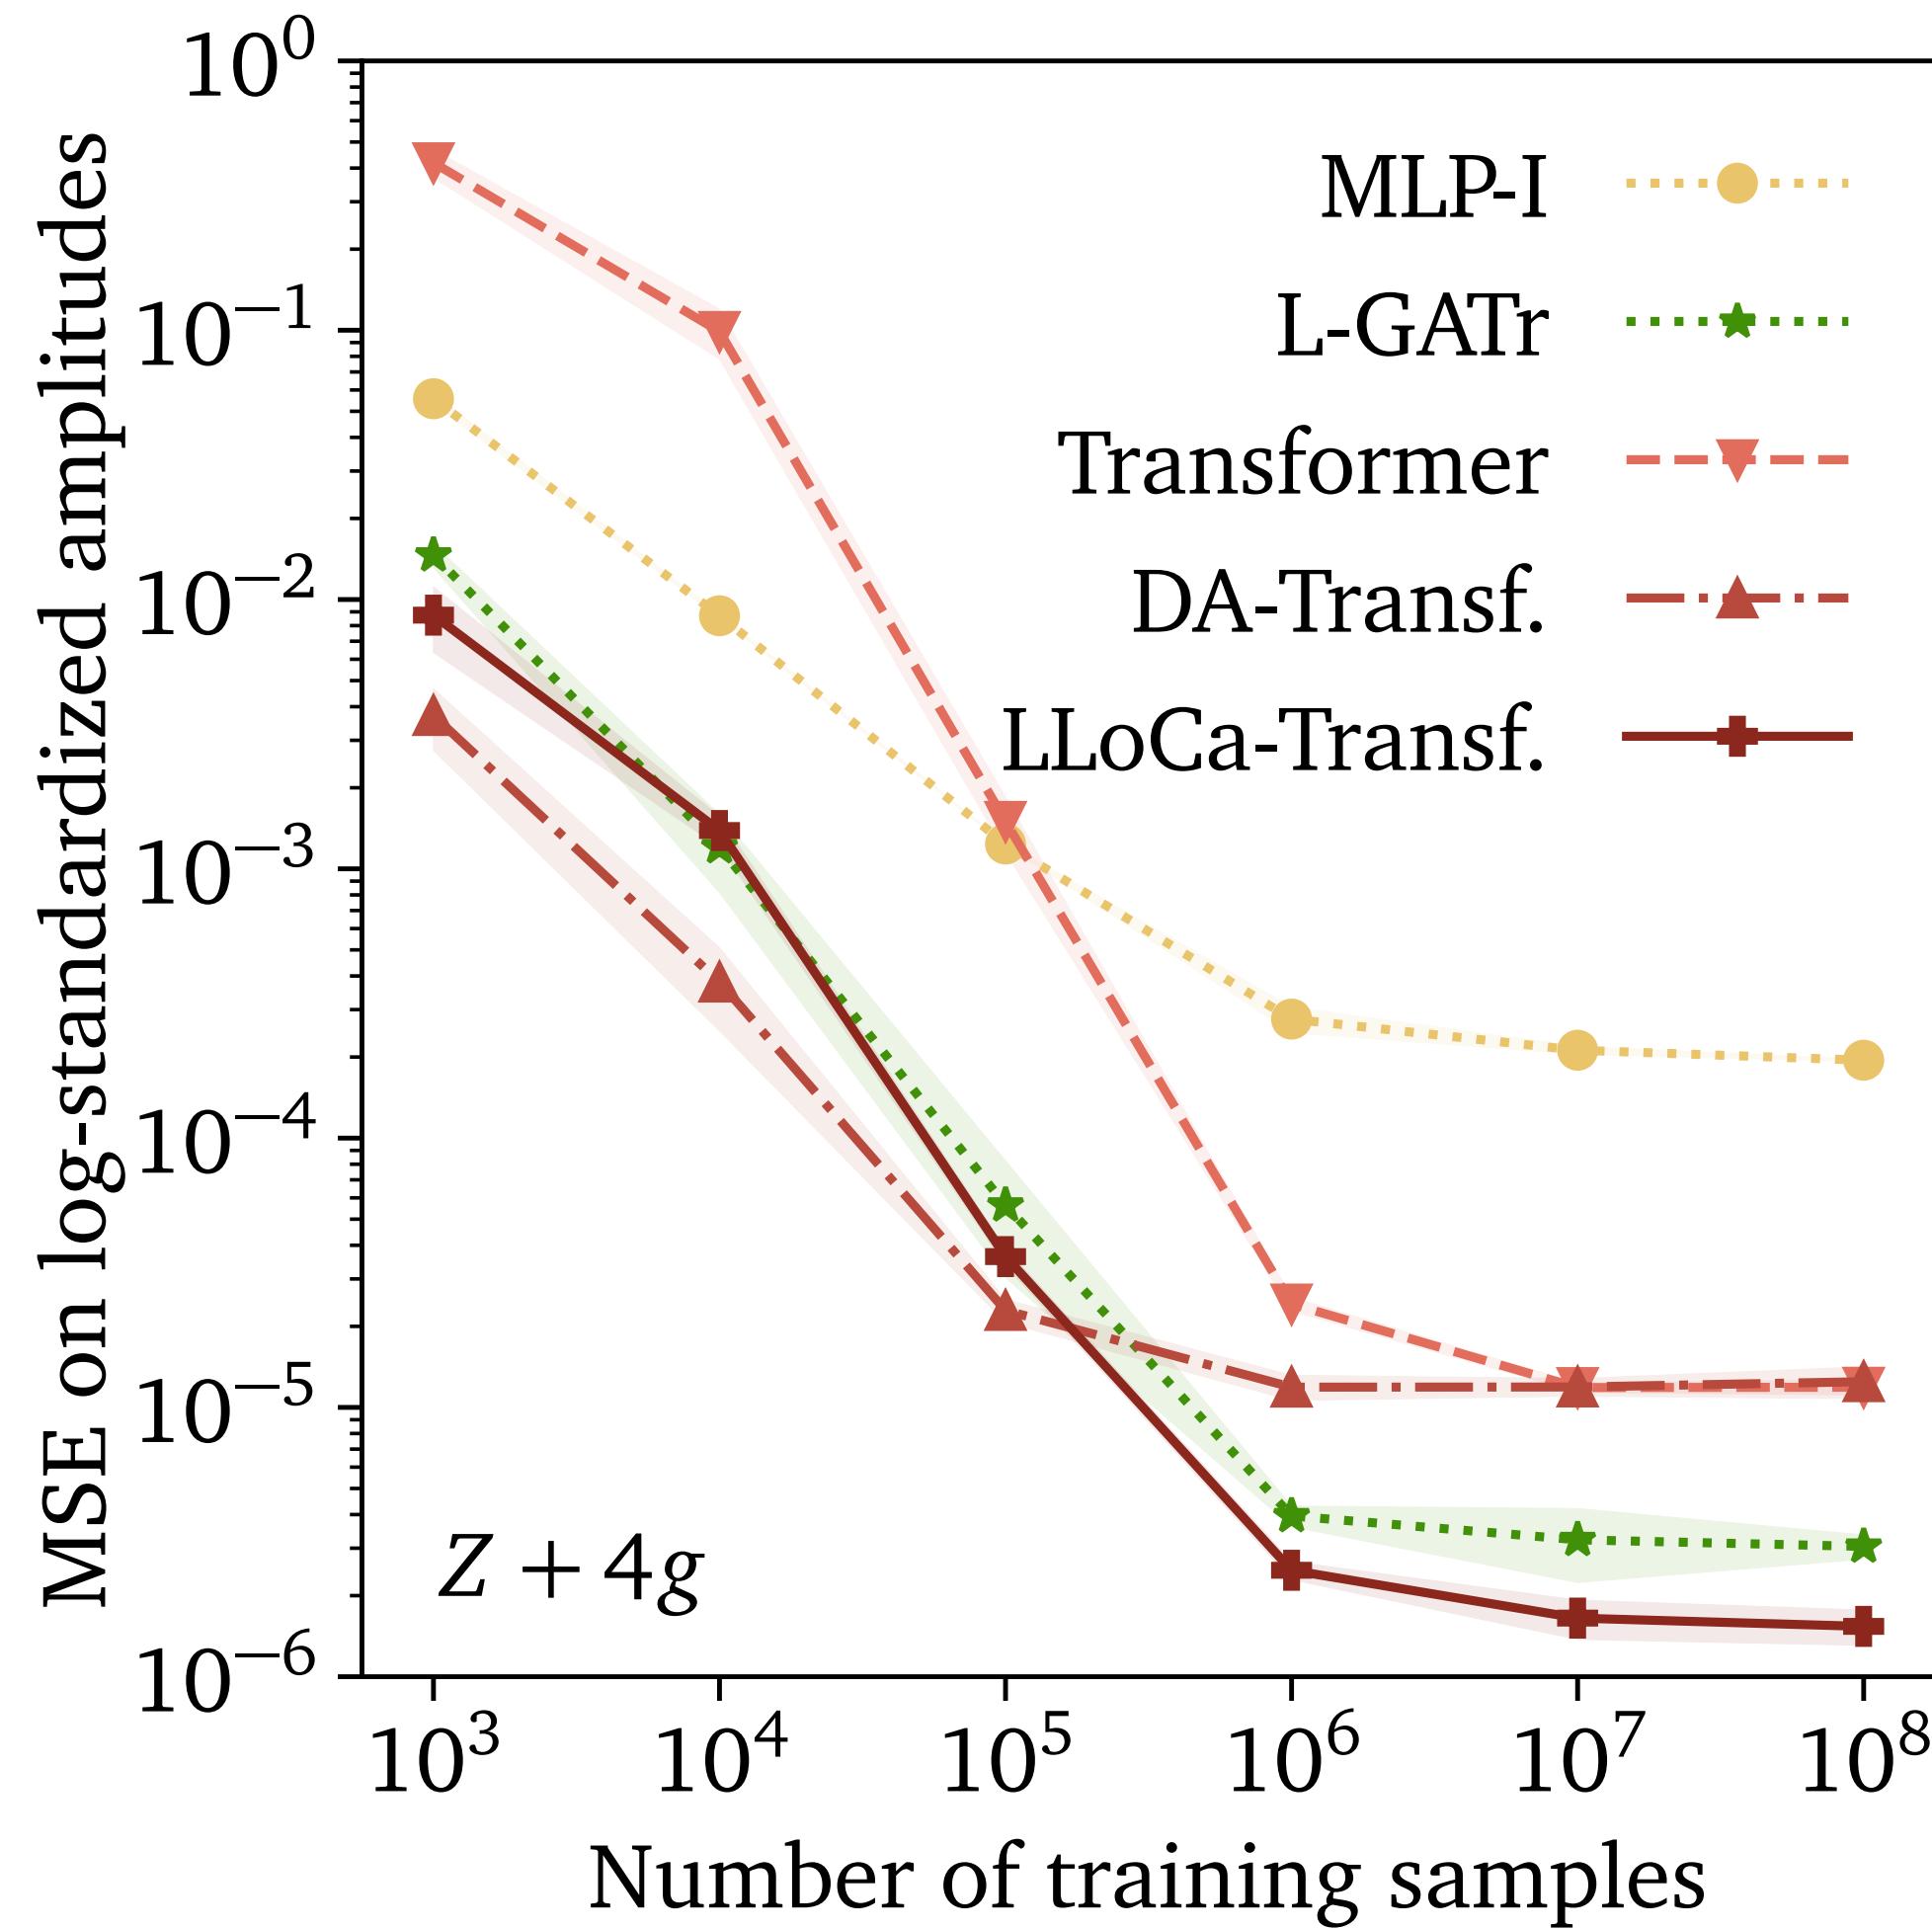
LLoCa handles arbitrary tensorial representations



Across our experiments,  
an equal mix of scalar and vector  
representations works best.

# Amplitude regression

## Equivariance vs data augmentation



Data augmentation (DA) emerges from LLoCa as a special choice of local frame

Lorentz-equivariance wins in the large-data regime

Data augmentation wins in the small-data regime

# Jet Tagging

## Lorentz symmetry breaking

Symmetry breaking Network	Architecture	Input	
	Accuracy AUC	Accuracy	AUC
Non-equivariant	0.855	0.9867	0.863    0.9880
SO(2)-equivariant	0.862	0.9878	0.864    0.9882
Lorentz-equivariant	0.856	0.9870	0.856    0.9870

Jet tagging score is  
only SO(2)-invariant

= LLoCa-Transformer

Lorentz-equivariance + symmetry breaking  
outperforms SO(2)-equivariance!

# Summary

- Canonicalization makes Lorentz-equivariance fast + flexible
  - More efficient than specialised layers
  - Arbitrary internal representations
  - Make any network Lorentz-equivariant
  - Can easily reduce the symmetry group
  - Data augmentation wins in the small-data regime



Luigi Favaro



Peter Lippmann



Sebastian Pitz



Gerrit Gerhartz



Huilin Qu



Tilman Plehn

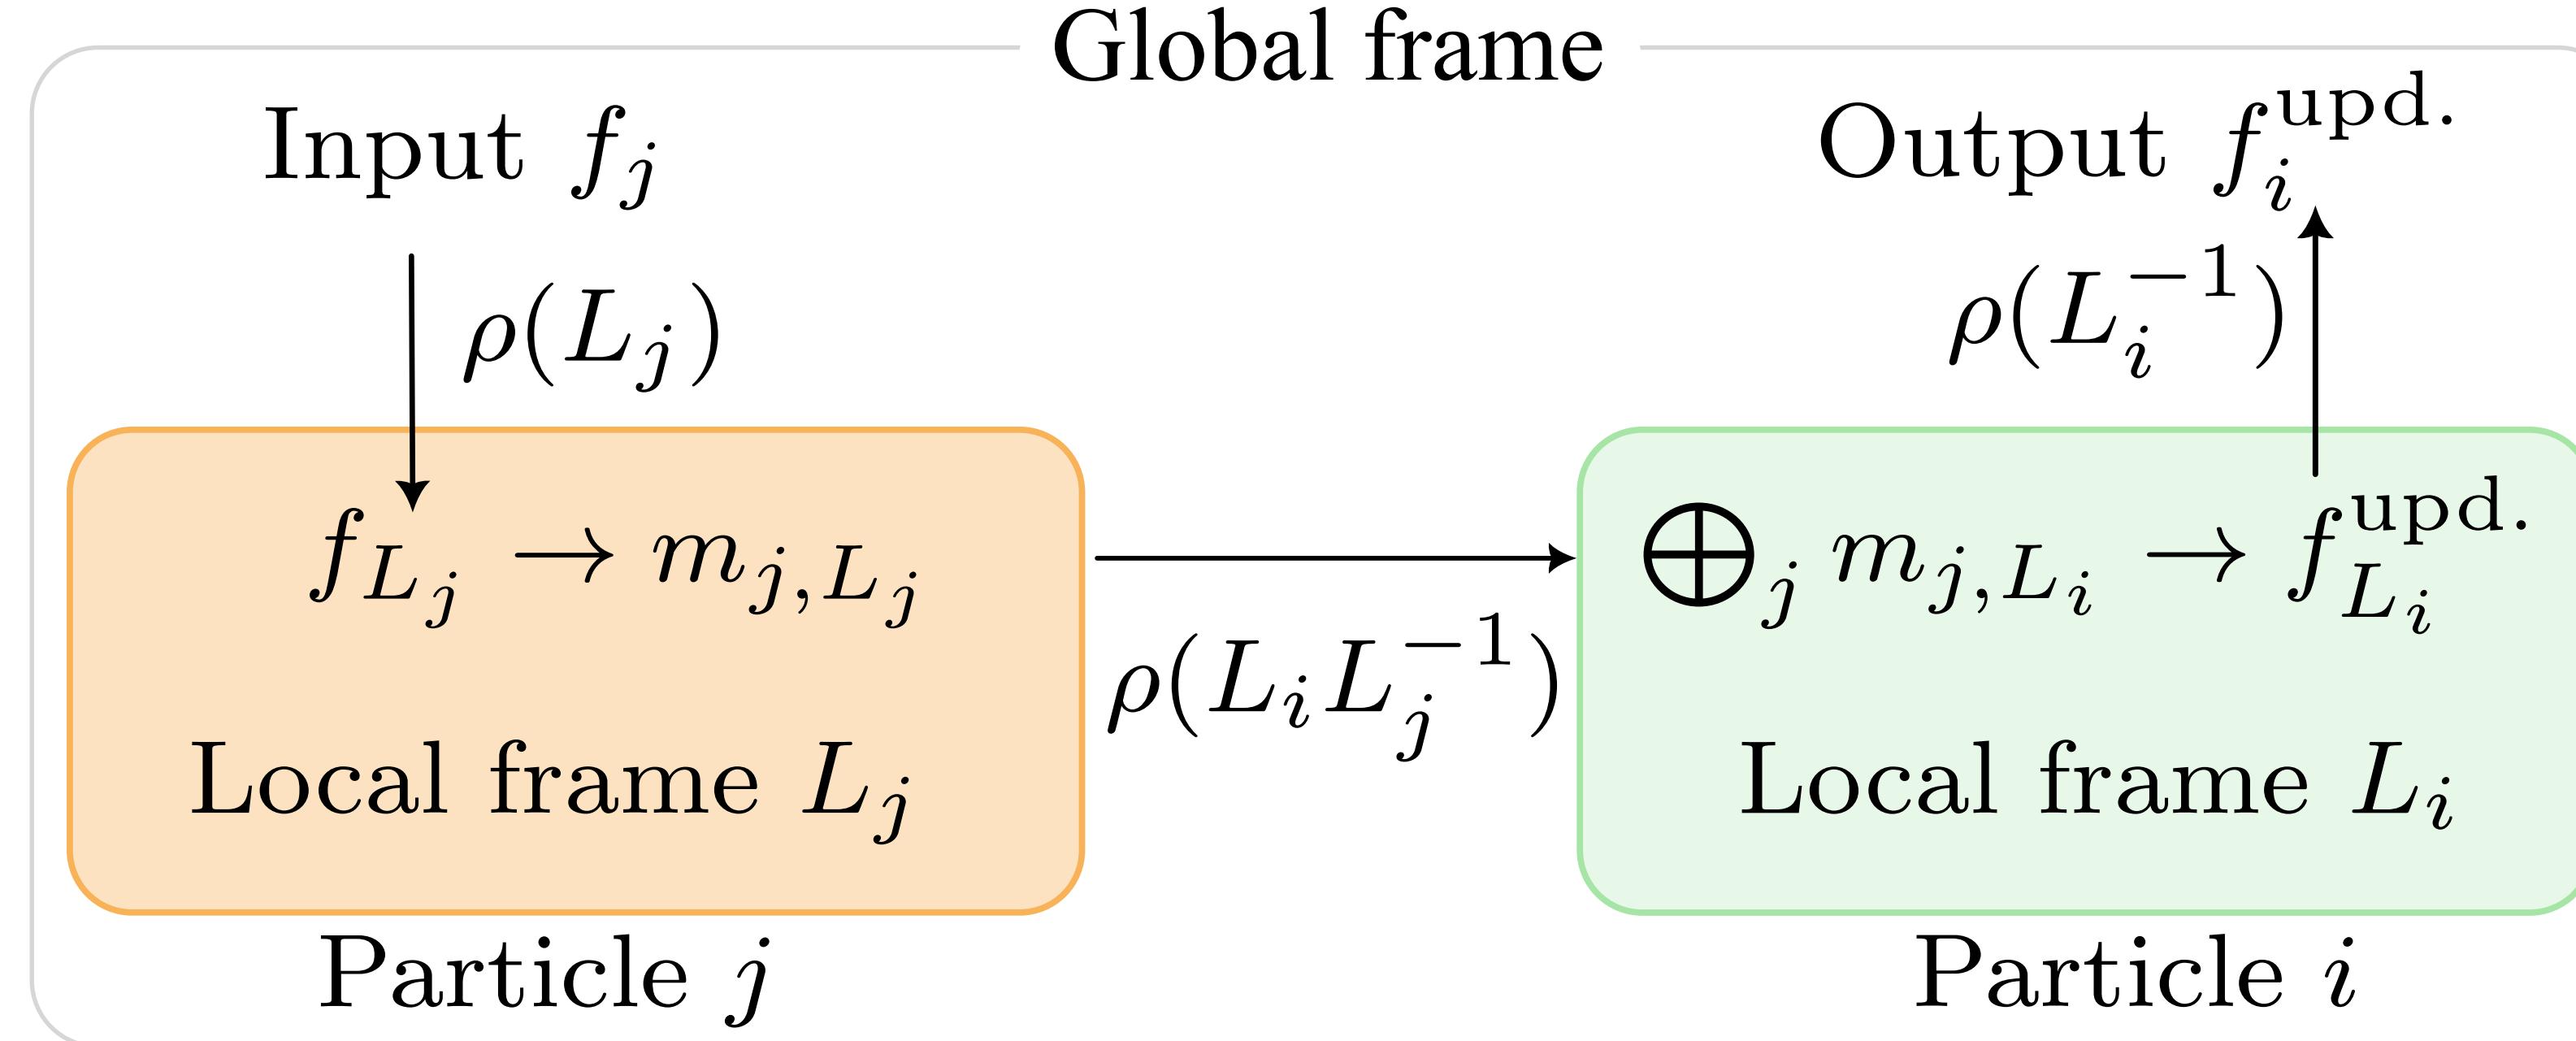


Fred Hamprecht

# Bonus material

# LLoCa

## Tensorial message-passing



# LLoCa

## Efficient Minkowski attention

$$f_{i,\mathbf{L}_i}^{\text{updated}} = \sum_{j=1}^N \text{softmax}_j \left( \frac{1}{\sqrt{d}} \left\langle q_{i,\mathbf{L}_i}, \rho_k(\mathbf{L}_i \mathbf{L}_j^{-1}) k_{j,\mathbf{l}_j} \right\rangle \right) \rho_v(\mathbf{L}_i \mathbf{L}_j^{-1}) v_{j,\mathbf{L}_j}$$

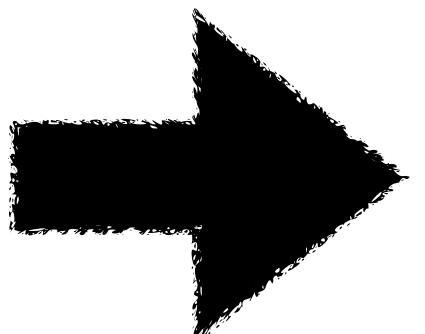
$$\left\langle q_{i,L_i}, \rho_k(L_i L_j^{-1}) k_{j,L_j} \right\rangle = \left\langle \rho_k(L_i^{-1}) q_{i,L_i}, \rho_k(L_j^{-1}) k_{j,L_j} \right\rangle$$

# Symmetry breaking

## Symmetry breaking with reference vectors

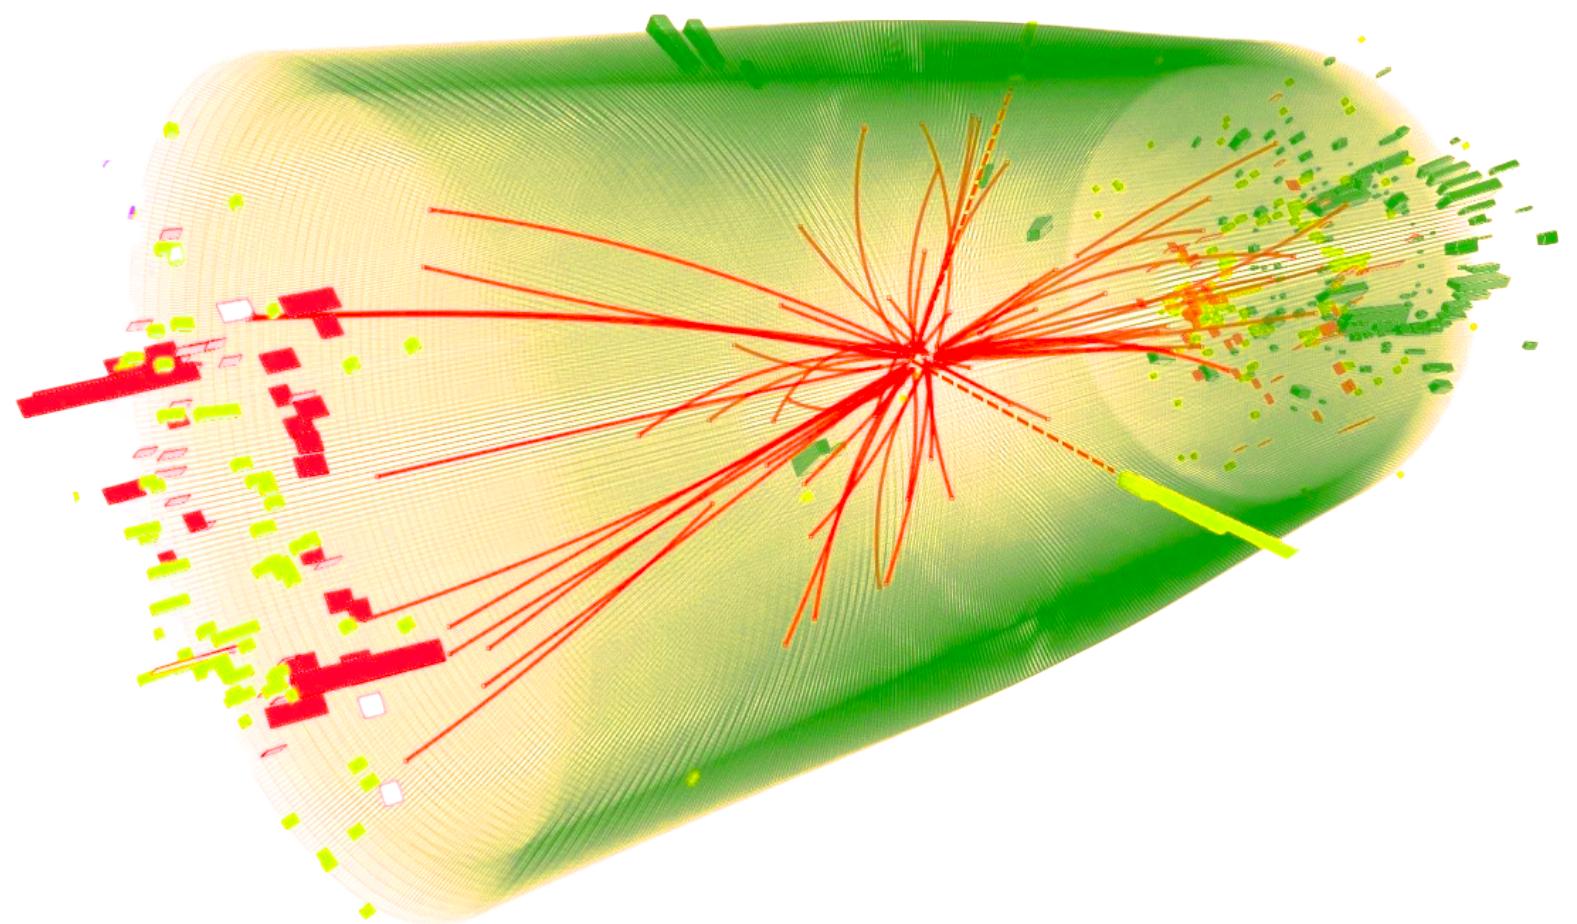
The jet tagging score is not Lorentz-equivariant

- Beam direction breaks invariance under rotations around the x- and y-axis
- Detector breaks boost invariance



Add reference vectors as extra particles to break Lorentz symmetry

- Beam reference vector:  
 $x^V = (1,0,0,1)$
- Time reference vector:  
 $x^V = (1,0,0,0)$



All Lorentz-equivariant taggers use symmetry-breaking inputs

# Jet tagging

## Top tagging dataset

Network	Accuracy	AUC	$1/\epsilon_B$ ( $\epsilon_S = 0.5$ )	$1/\epsilon_B$ ( $\epsilon_S = 0.3$ )
PFN [56]	0.932	0.9819	$247 \pm 3$	$888 \pm 17$
ParticleNet [55]	0.940	0.9858	$397 \pm 7$	$1615 \pm 93$
ParT [21]	0.940	0.9858	$413 \pm 16$	$1602 \pm 81$
MIParT [53]	0.942	0.9868	$505 \pm 8$	$2010 \pm 97$
LorentzNet* [11]	0.942	0.9868	$498 \pm 18$	$2195 \pm 173$
CGENN* [13]	0.942	0.9869	500	2172
PELICAN* [12]	$0.9426 \pm 0.0002$	$0.9870 \pm 0.0001$	–	$2250 \pm 75$
L-GATr* [14]	$0.9423 \pm 0.0002$	$0.9870 \pm 0.0001$	$540 \pm 20$	$2240 \pm 70$
LLoCa-Transformer*	$0.9416 \pm 0.0001$	$0.9866 \pm 0.0001$	$492 \pm 15$	$2150 \pm 130$

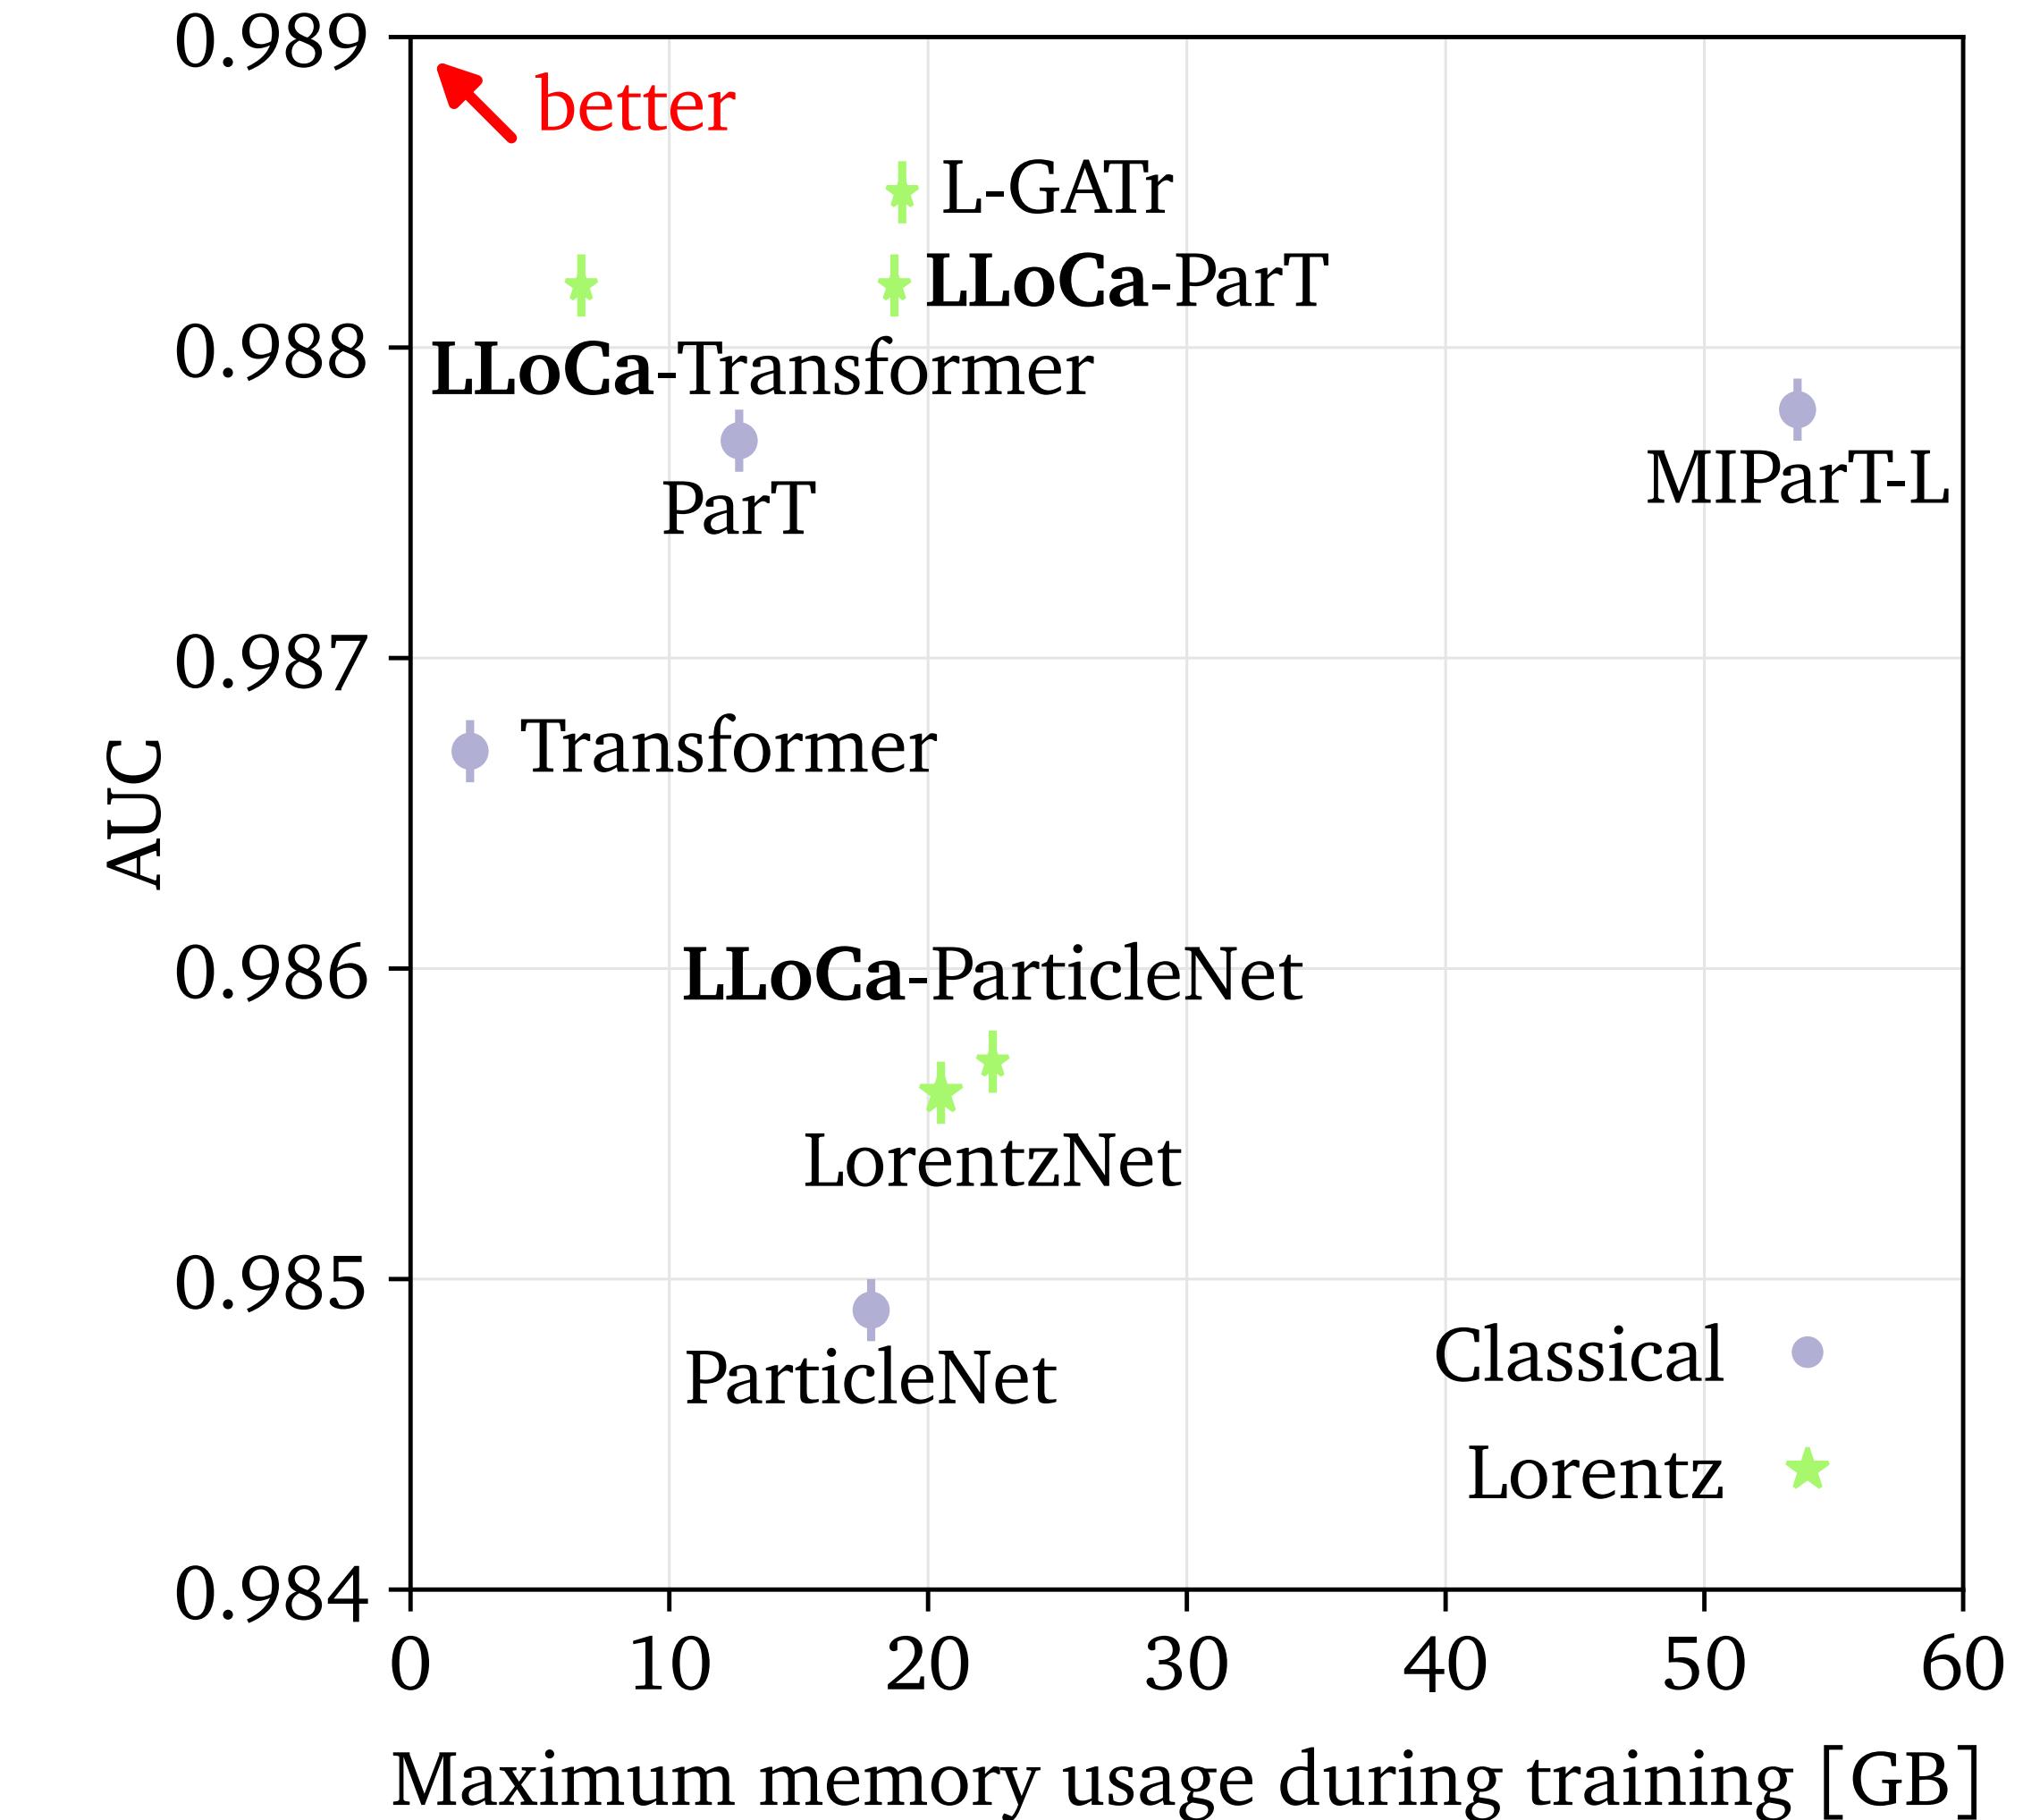
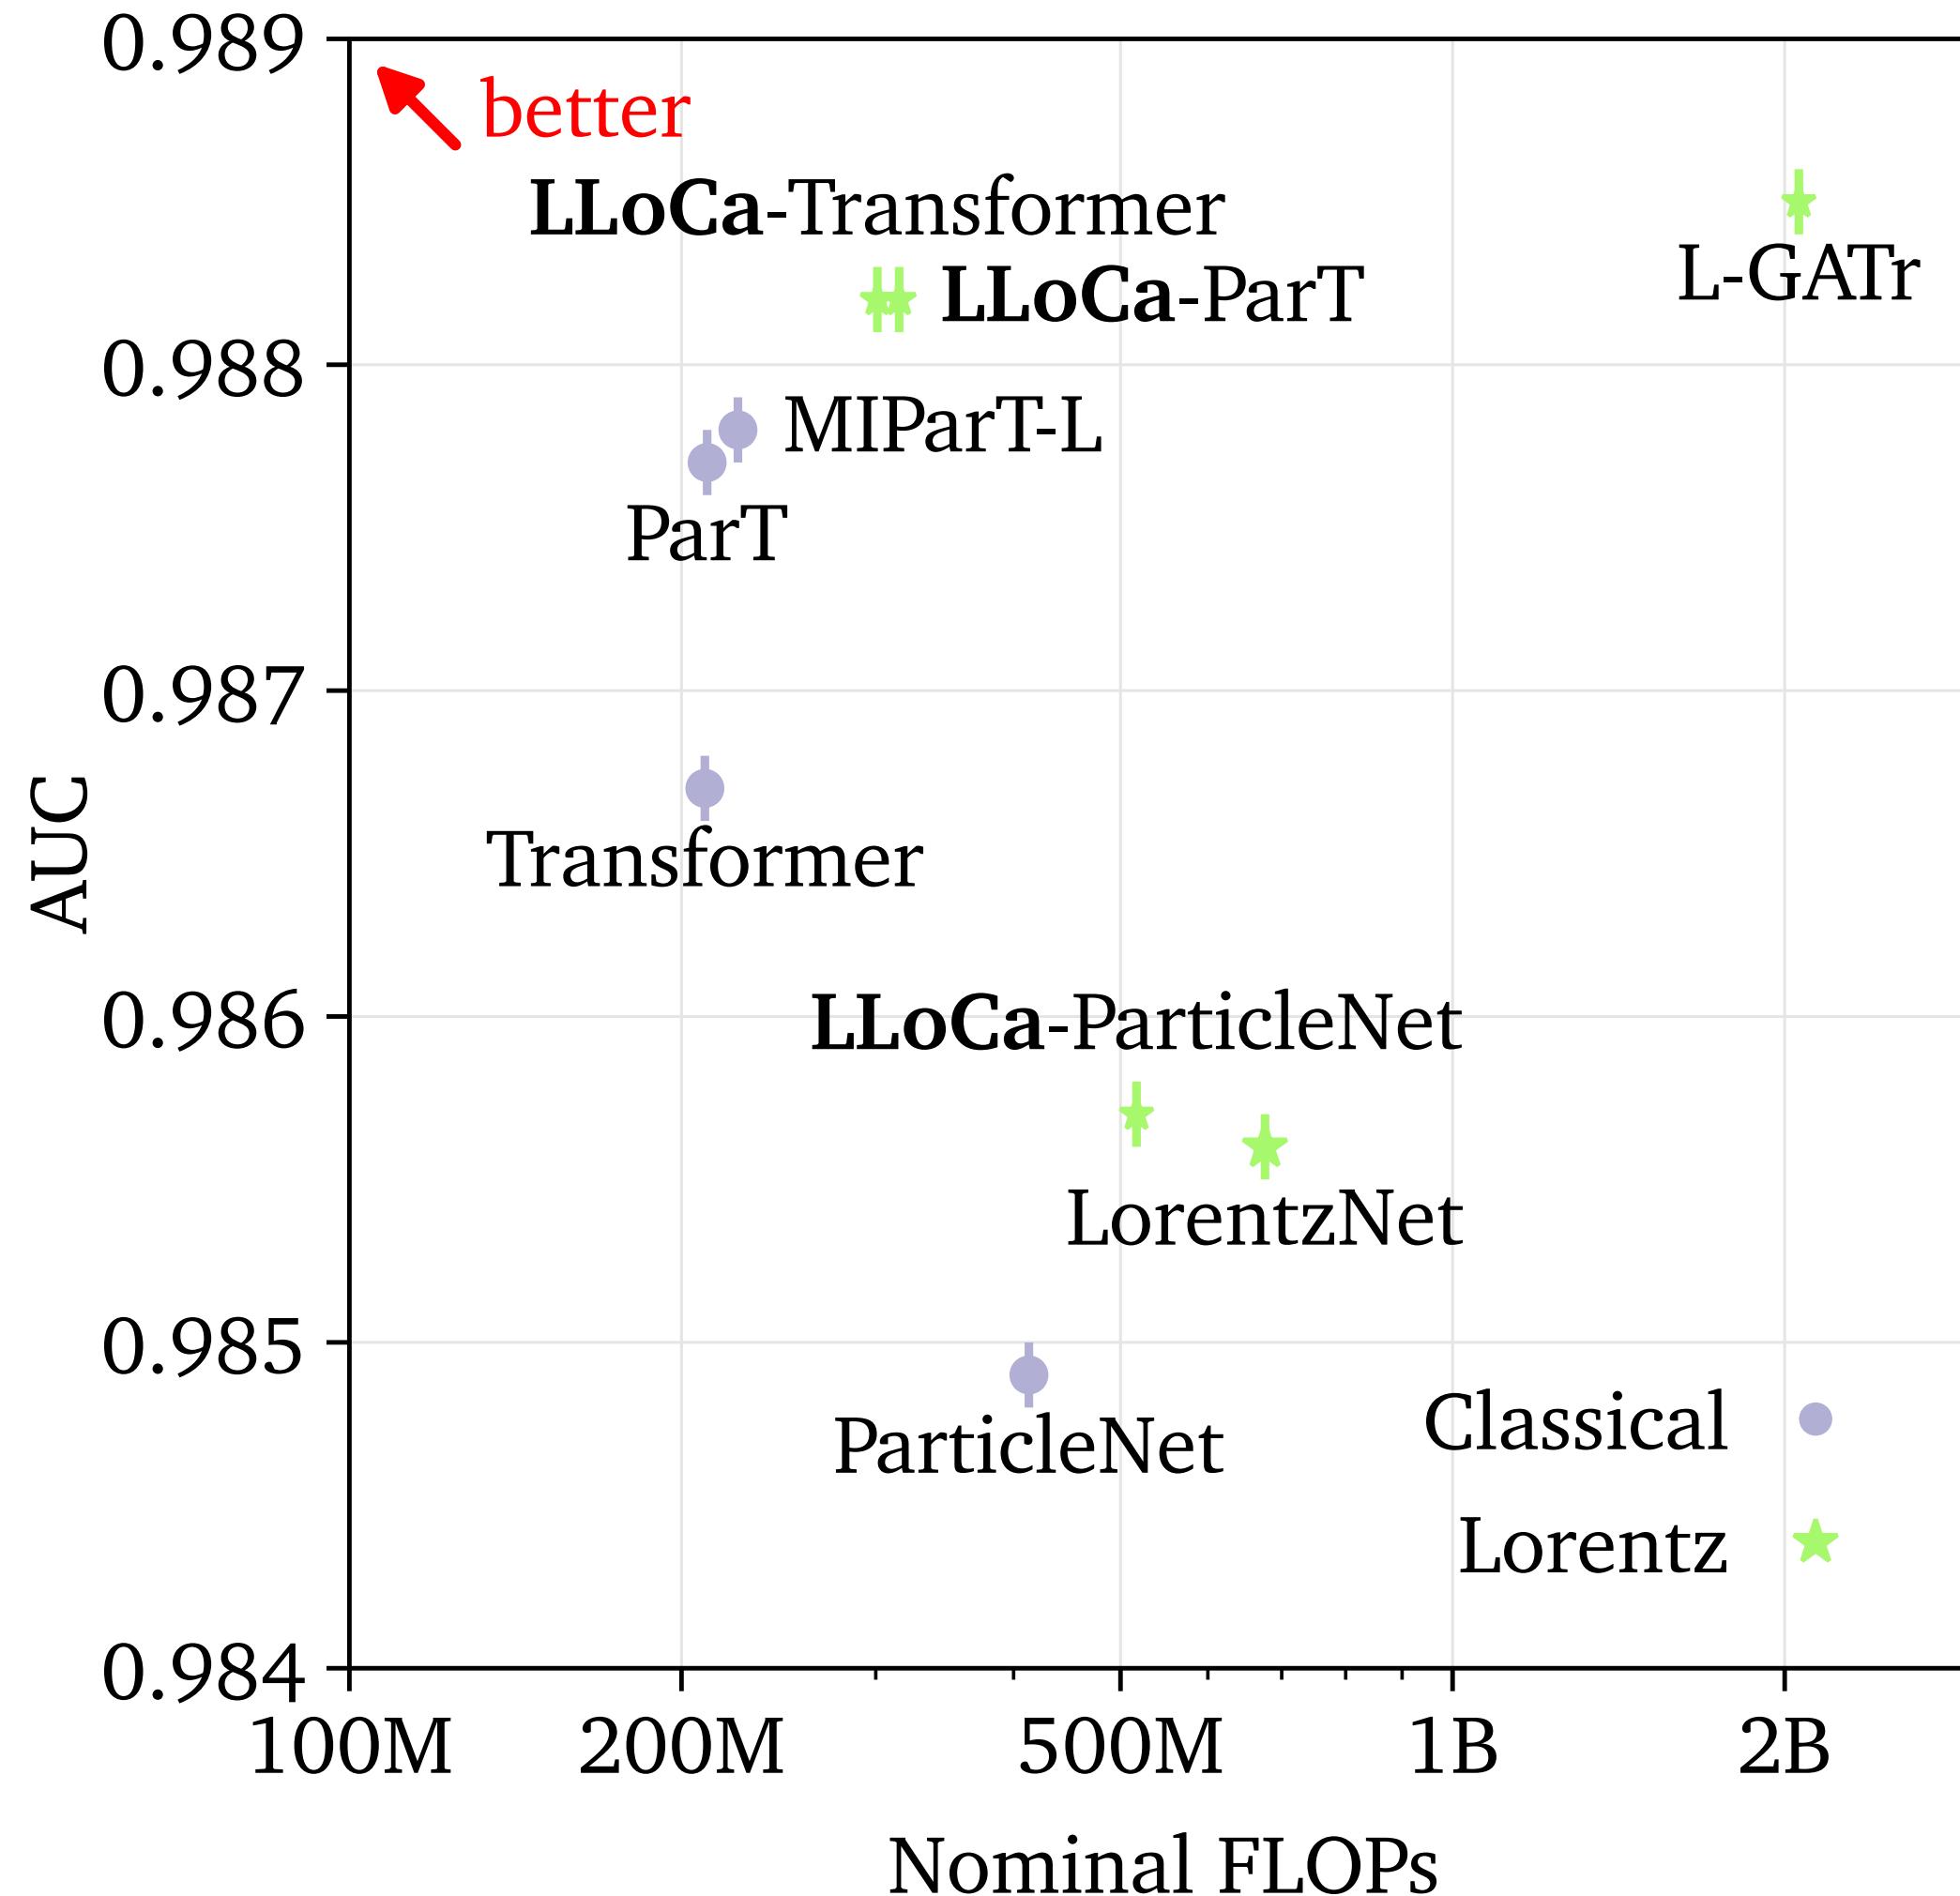
# Jet tagging

## JetClass dataset

Network	Accuracy	AUC	Time	FLOPs	Memory	Parameters
PFN [56]	0.772	0.9714	3h	3M	1.6G	86k
P-CNN [56]	0.809	0.9789	3h	12M	3.0G	354k
MIParT-L [53]	0.861	0.9878	47h	225M	53.6G	2380k
LorentzNet* [11]	0.847	0.9856	57h	676M	20.5G	223k
L-GATr* [15]	0.866	0.9885	180h	2060M	19.0G	1079k
ParticleNet [55]	0.844	0.9849	27h	413M	16.5G	366k
LLoCa-ParticleNet* [22]	0.848	0.9857	41h	517M	23.5G	385k
ParT [21]	0.861	0.9877	38h	211M	13.3G	2141k
LLoCa-ParT* [22]	0.864	0.9882	67h	315M	19.9G	2160k
Transformer [22]	0.855	0.9867	17h	210M	2.3G	1979k
LLoCa-Transformer* [22]	0.864	0.9882	33h	301M	6.9G	1998k

# Jet tagging

## FLOPs and memory



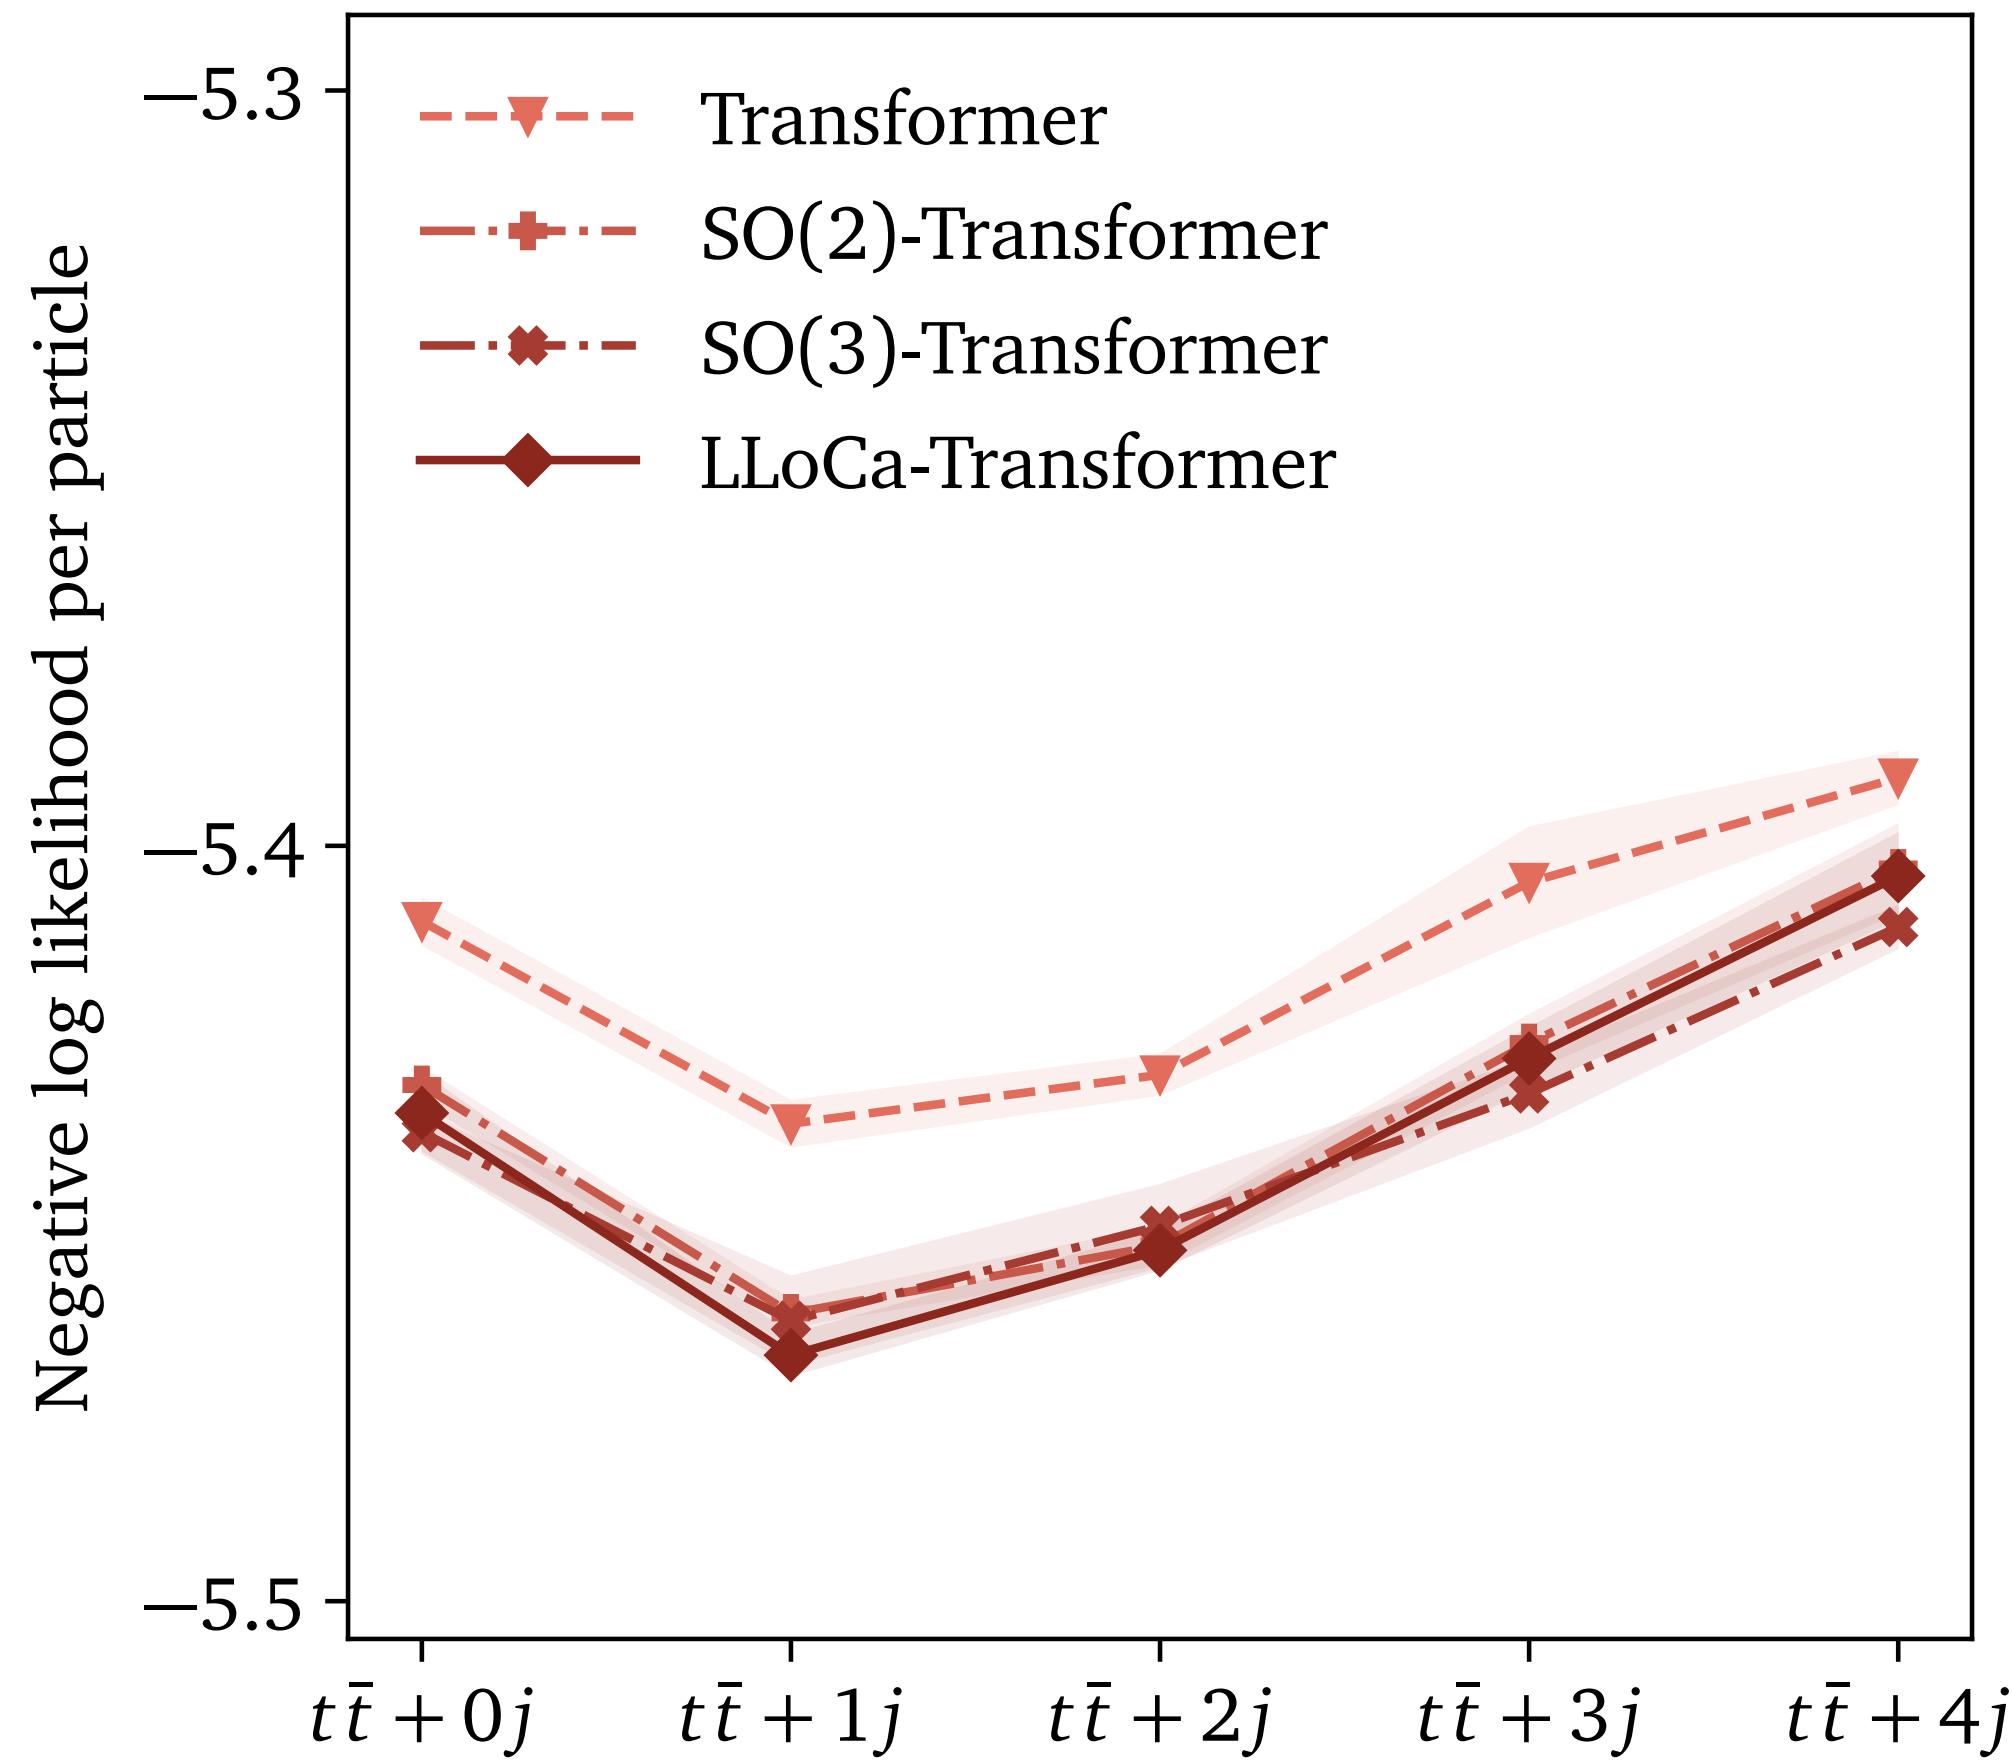
# Jet tagging

## Effect of internal representations

Network	Accuracy AUC	
Transformer	0.855	0.9867
ParT	0.861	0.9877
LLoCa-Transformer (global can.)	0.861	0.9878
LLoCa-Transformer (16 scalars)	0.851	0.9862
LLoCa-Transformer (4 vectors)	0.863	0.9880
LLoCa-Transformer (8 scalars, 2 vectors)	0.864	0.9882
LLoCa-Transformer (12 scalars, 1 vector)	0.864	0.9882

# Event generation

## Effect of internal representations



Lorentz-equivariance does  
not help for event generation