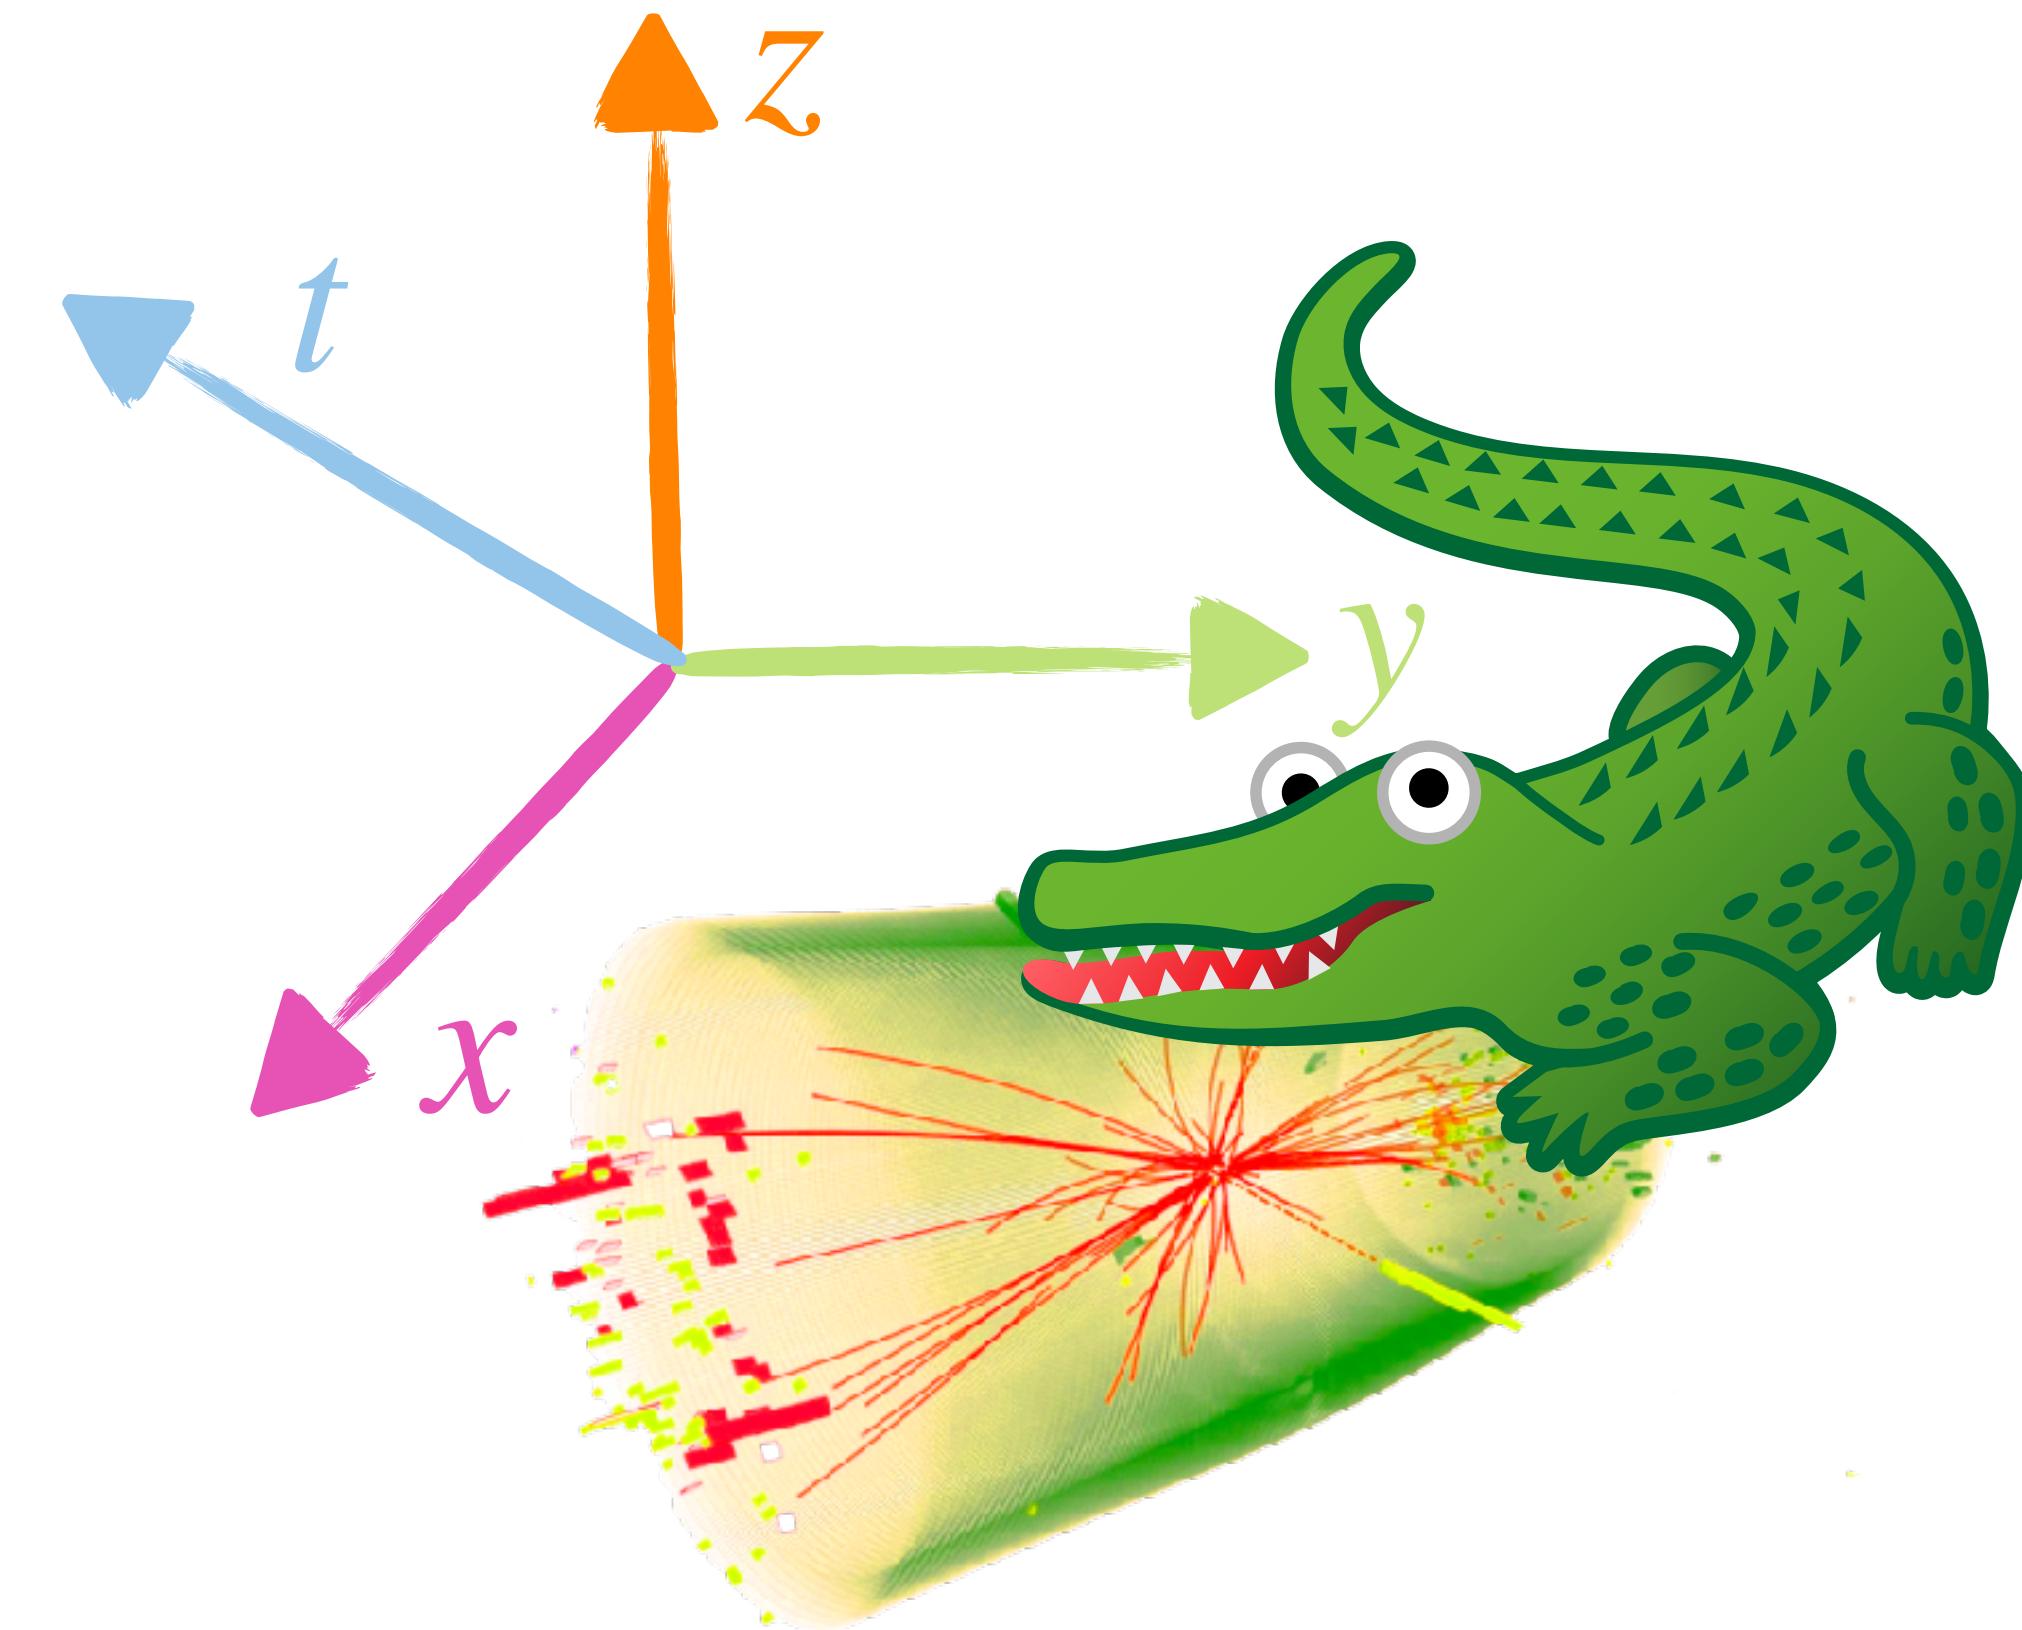


# Transformers and Symmetries

Jonas Spinner

in collaboration with Johann Brehmer, Víctor Bresó,  
Luigi Favaro, Gerrit Gerhartz, Pim de Haan,  
Fred A. Hamprecht, Peter Lippmann, Sebastian Pitz,  
Tilman Plehn, Huilin Qu and Jesse Thaler

Build Big or Build Smart Workshop  
05.09.2025



UNIVERSITÄT  
HEIDELBERG  
ZUKUNFT  
SEIT 1386

# Symmetries in LHC physics

## Permutation symmetry

## Lorentz symmetry

Gong et al, arXiv:2201.08187

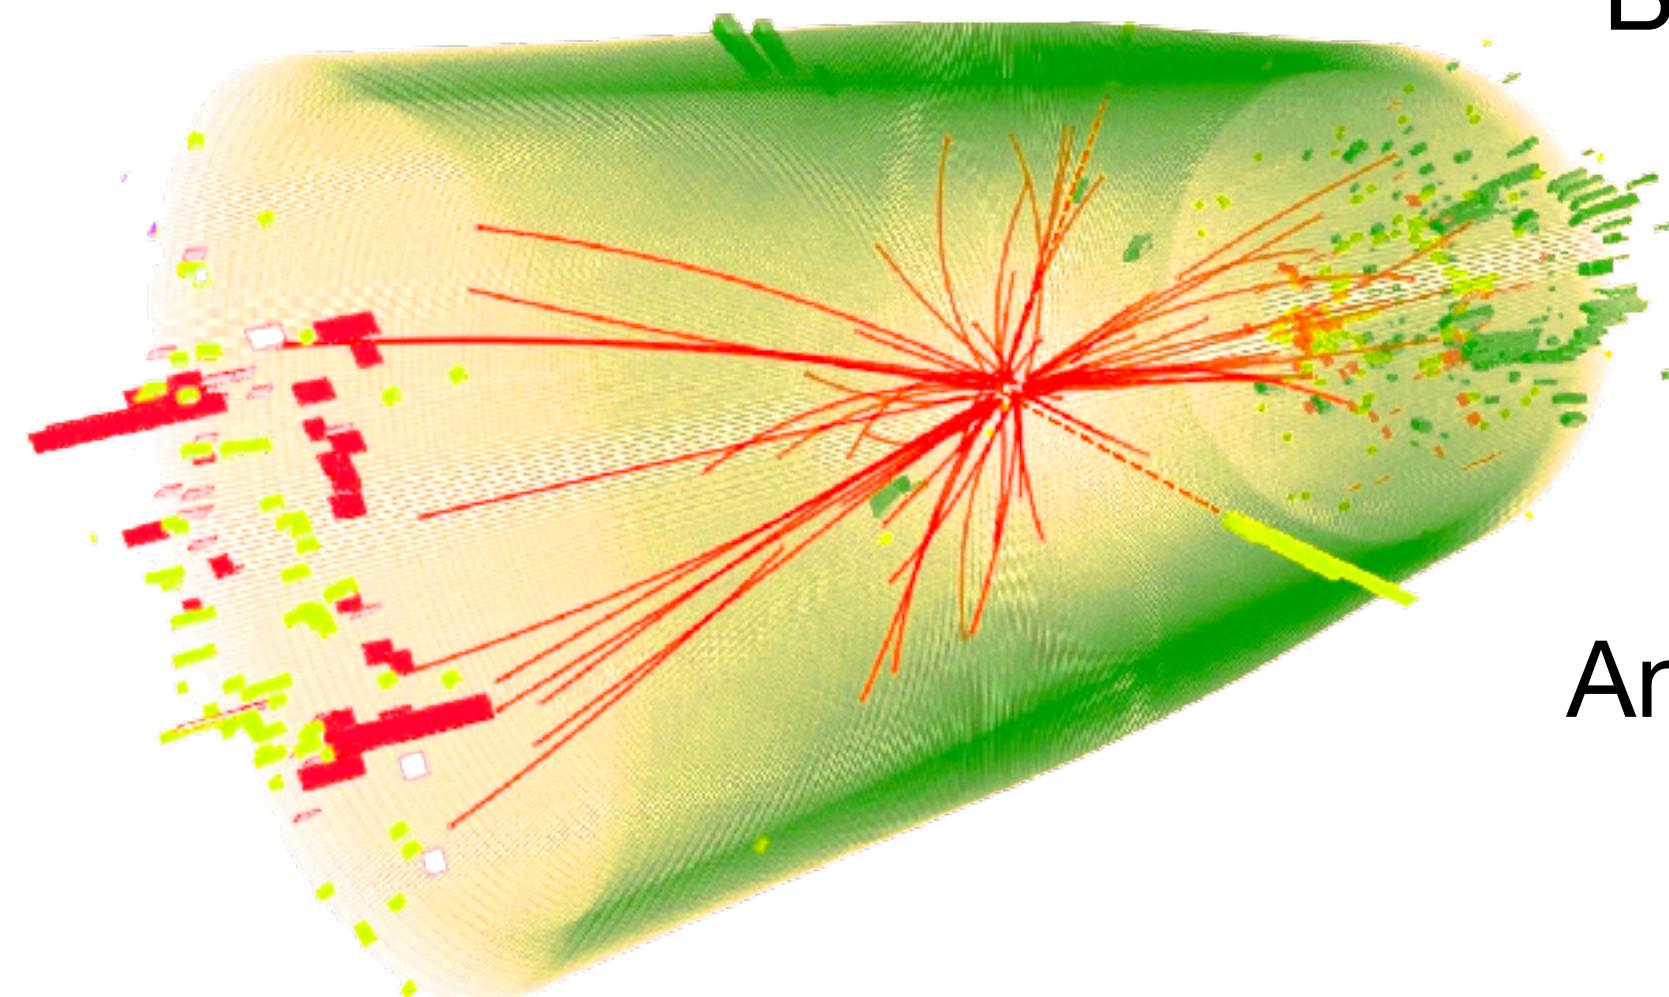
Bogatskiy et al, arXiv:2211.00454

Hao et al, arXiv:2212.07347

Ruhe et al, arXiv:2305.11141

Spinner et al, arXiv:2405.14806

Spinner et al, arXiv:2505.20280



## Rotations and translations in $(\eta, \phi)$

Dillon et al, arXiv:2108.04253

Favaro et al, arXiv:2312.03067

Bhardwaj et al, arXiv:2402.12449

## Universal splitting kernels

Andreassen et al, arXiv:1804.09720

Butter et al, arXiv:2412.12074

## Invariance under re-simulation

Harris et al, arXiv:1804.09720

More...

# Permutation symmetry

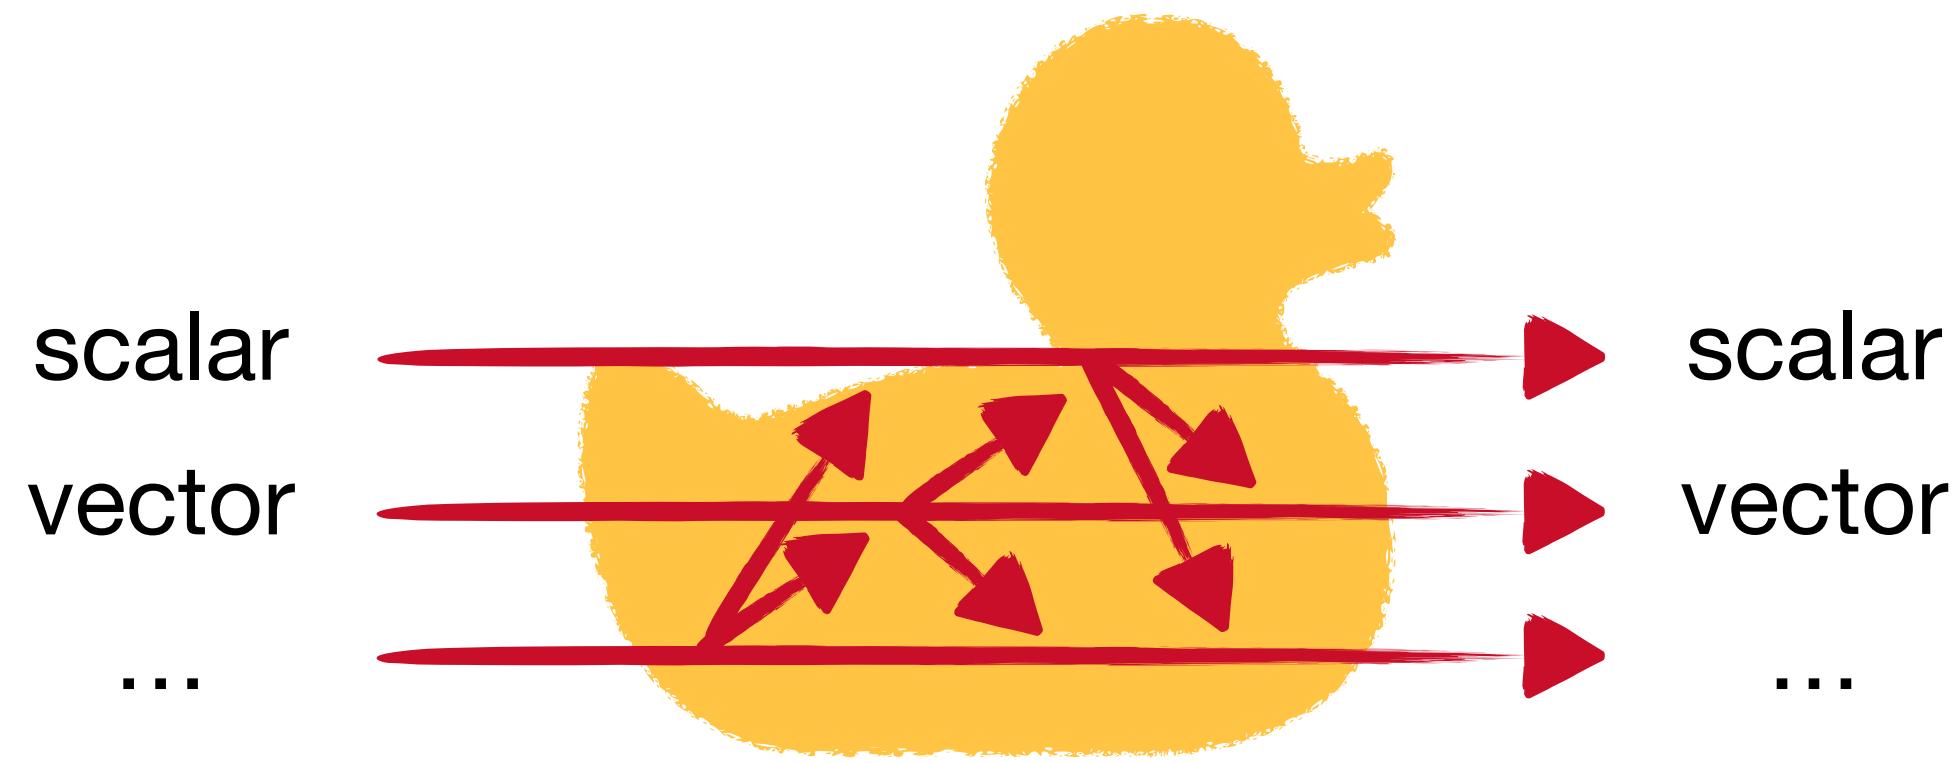
**Message-passing  
Graph Networks**  
 Formally more expressive

vs

**Transformers**  
 Better at scale

# Lorentz symmetry

Specialised layers



**PELICAN**

arXiv:2211.00454

arXiv:2307.16506

**LorentzNet**

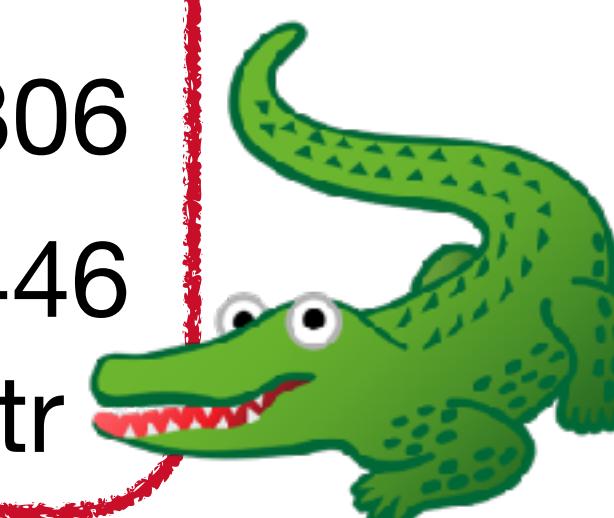
arXiv:2201.08187

**L-GATr**

arXiv:2405.14806

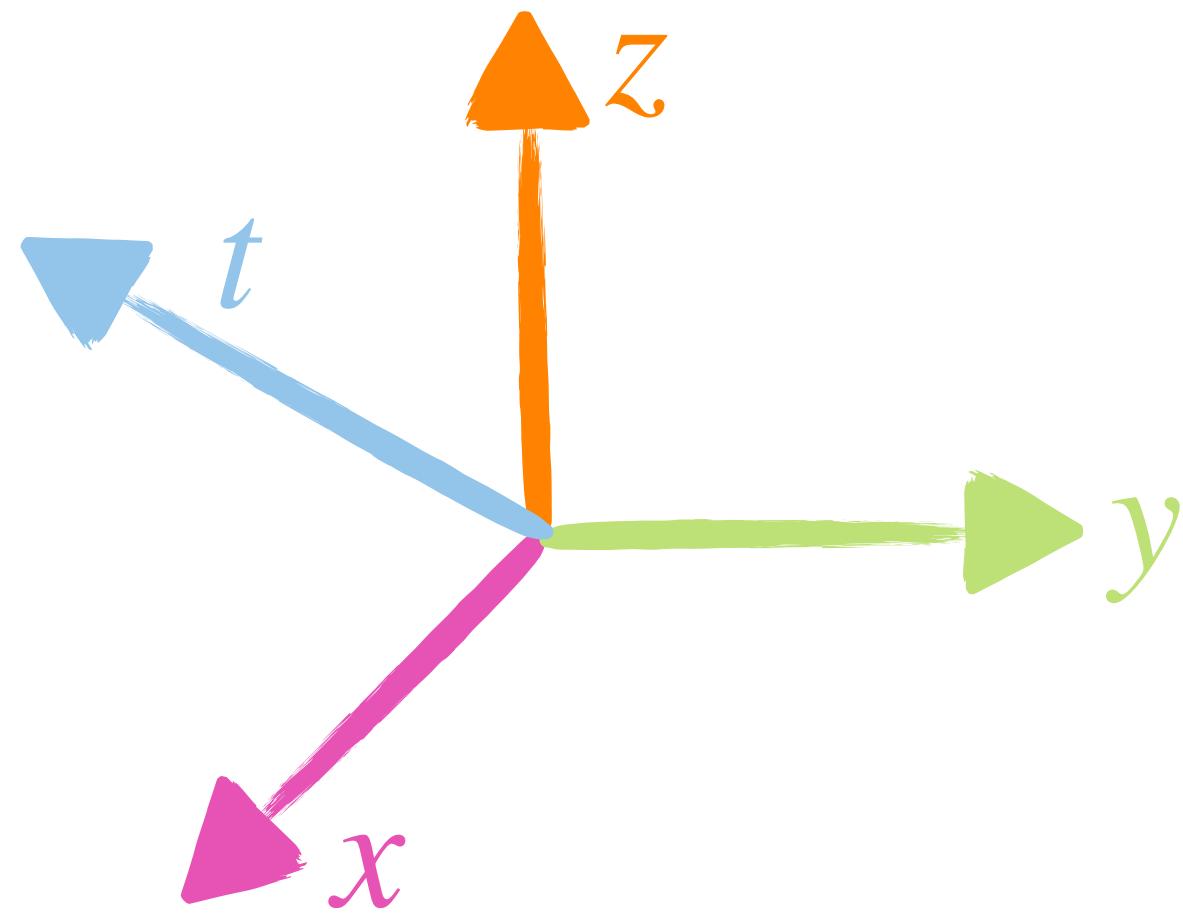
arXiv:2411.00446

pip install lgatr



vs

Canonicalization



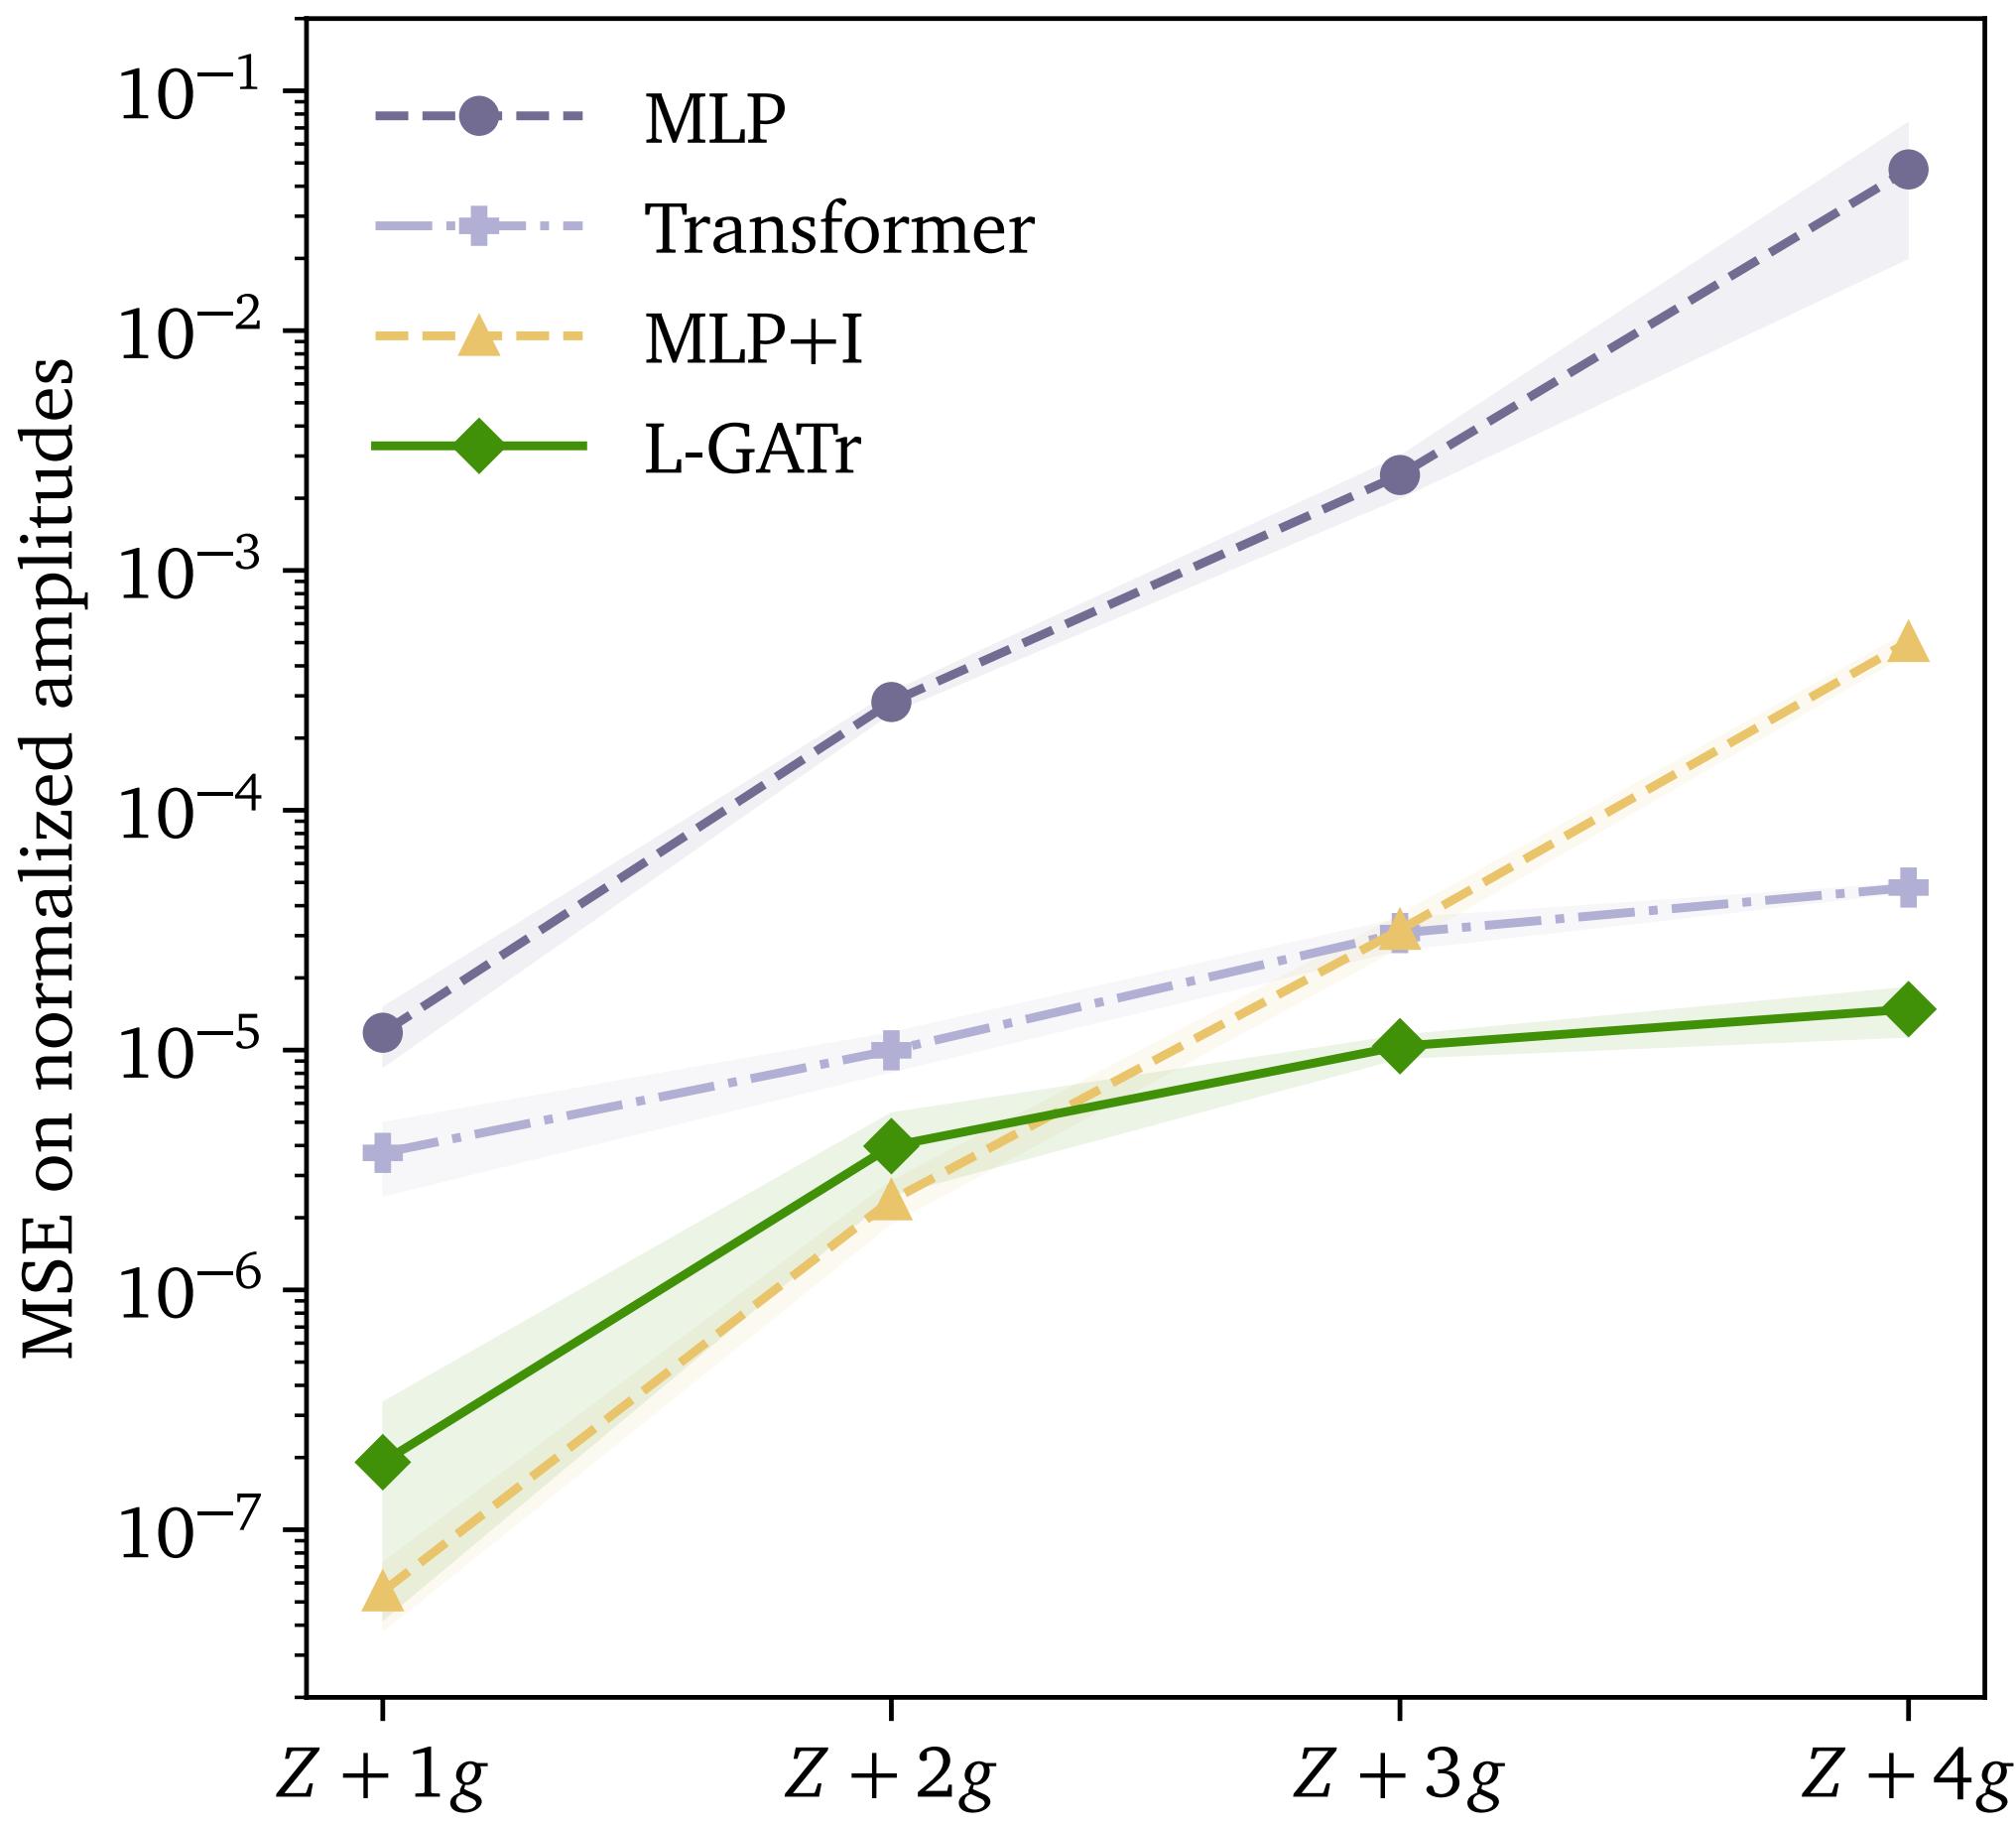
**LLoCa**

arXiv:2505.20280

arXiv:2508.14898

# Amplitude regression

## Gain from Lorentz and Permutation symmetry



Naive

$$\begin{pmatrix} E_1 \\ p_{x,1} \\ p_{y,1} \\ p_{z,1} \\ \text{PID}_1 \\ E_2 \\ p_{x,2} \\ p_{y,2} \\ p_{z,2} \\ \vdots \end{pmatrix}$$

MLP

Permutation symmetry

$$\begin{pmatrix} E_1 \\ p_{x,1} \\ p_{y,1} \\ p_{z,1} \\ \text{PID}_1 \\ E_2 \\ p_{x,2} \\ p_{y,2} \\ p_{z,2} \\ \vdots \\ \vdots \end{pmatrix}$$

Transformer

Lorentz invariants

$$\begin{pmatrix} p_1^\mu p_{1,\mu} \\ p_1^\mu p_{2,\mu} \\ p_1^\mu p_{3,\mu} \\ \vdots \\ p_2^\mu p_{2,\mu} \\ p_2^\mu p_{3,\mu} \\ \vdots \end{pmatrix}$$

MLP+I

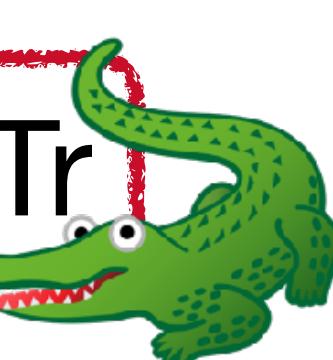
Permutation + Lorentz

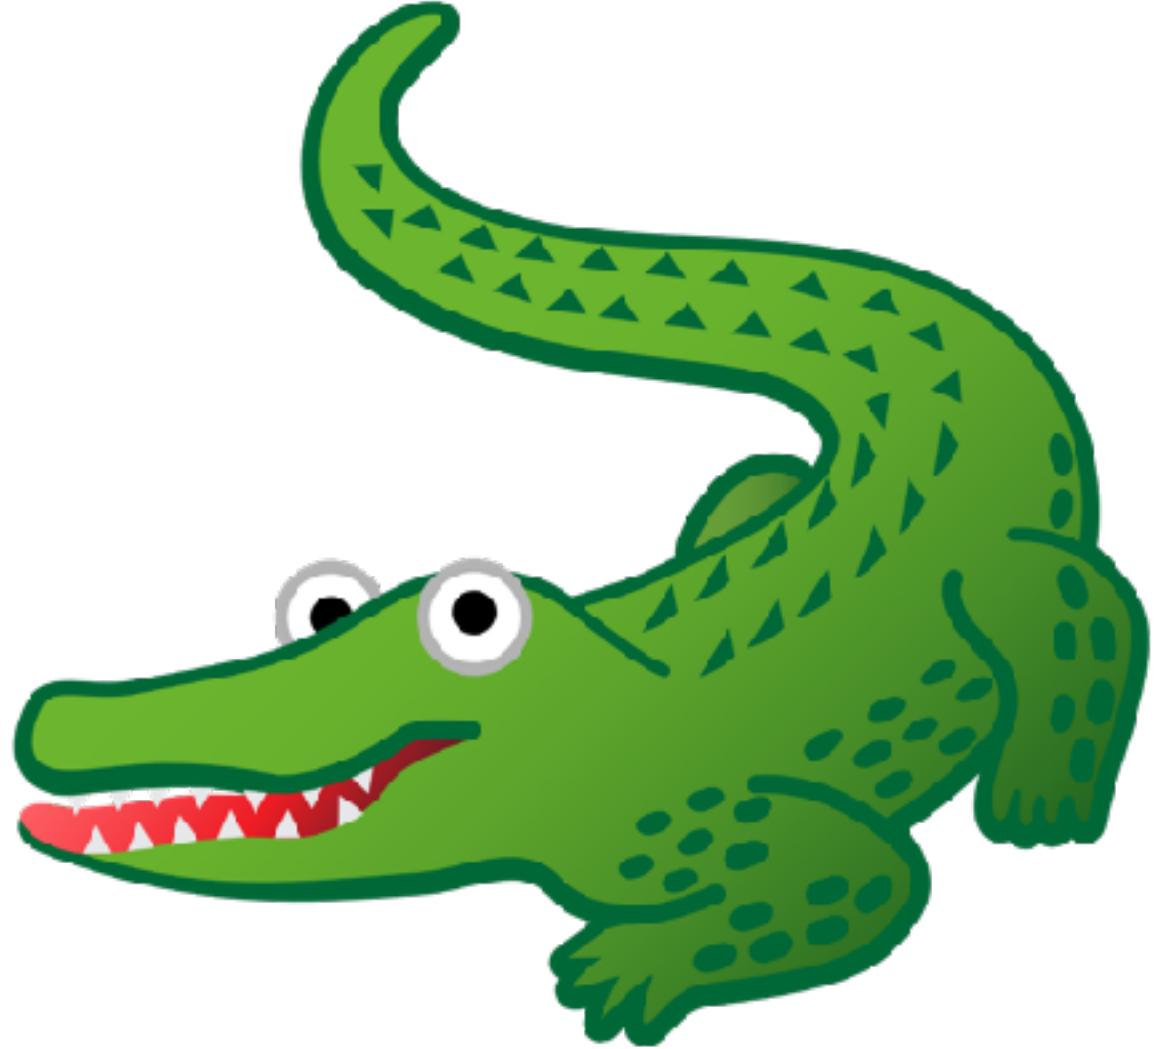
$$\begin{pmatrix} p_1^\mu \\ \text{PID}_1 \end{pmatrix}$$

$$\begin{pmatrix} p_2^\mu \\ \text{PID}_1 \end{pmatrix}$$

$\vdots$

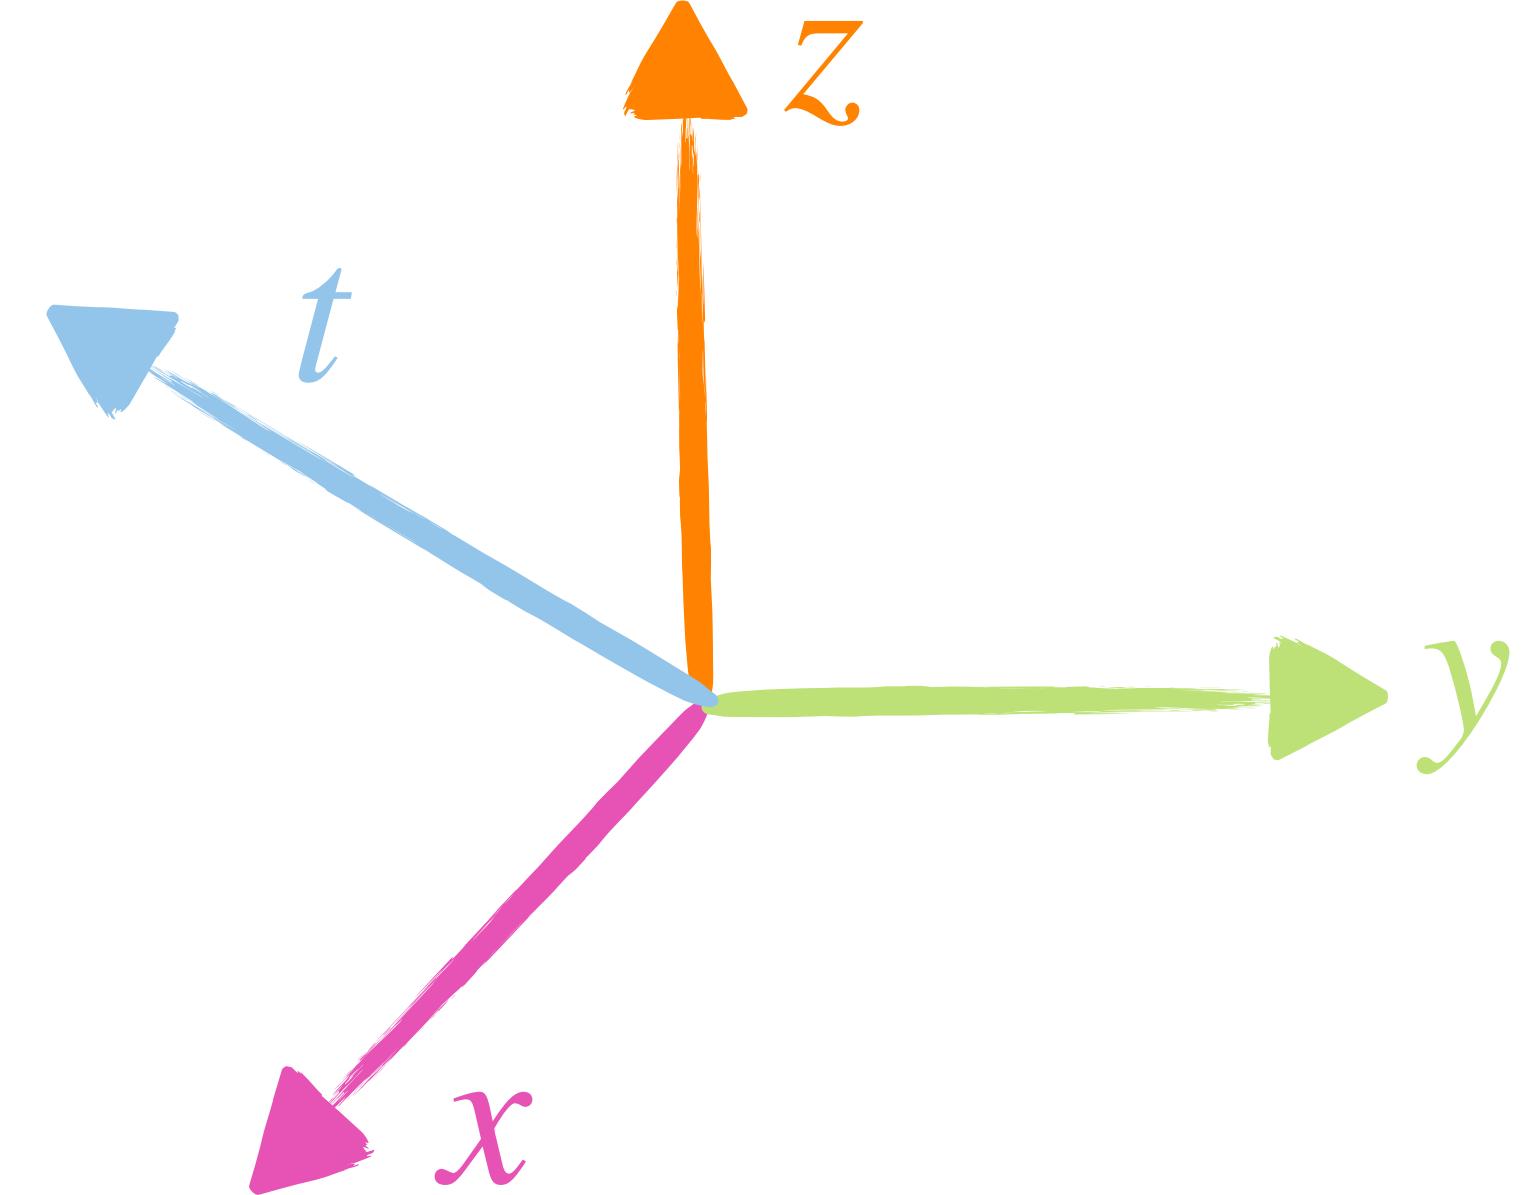
L-GATr





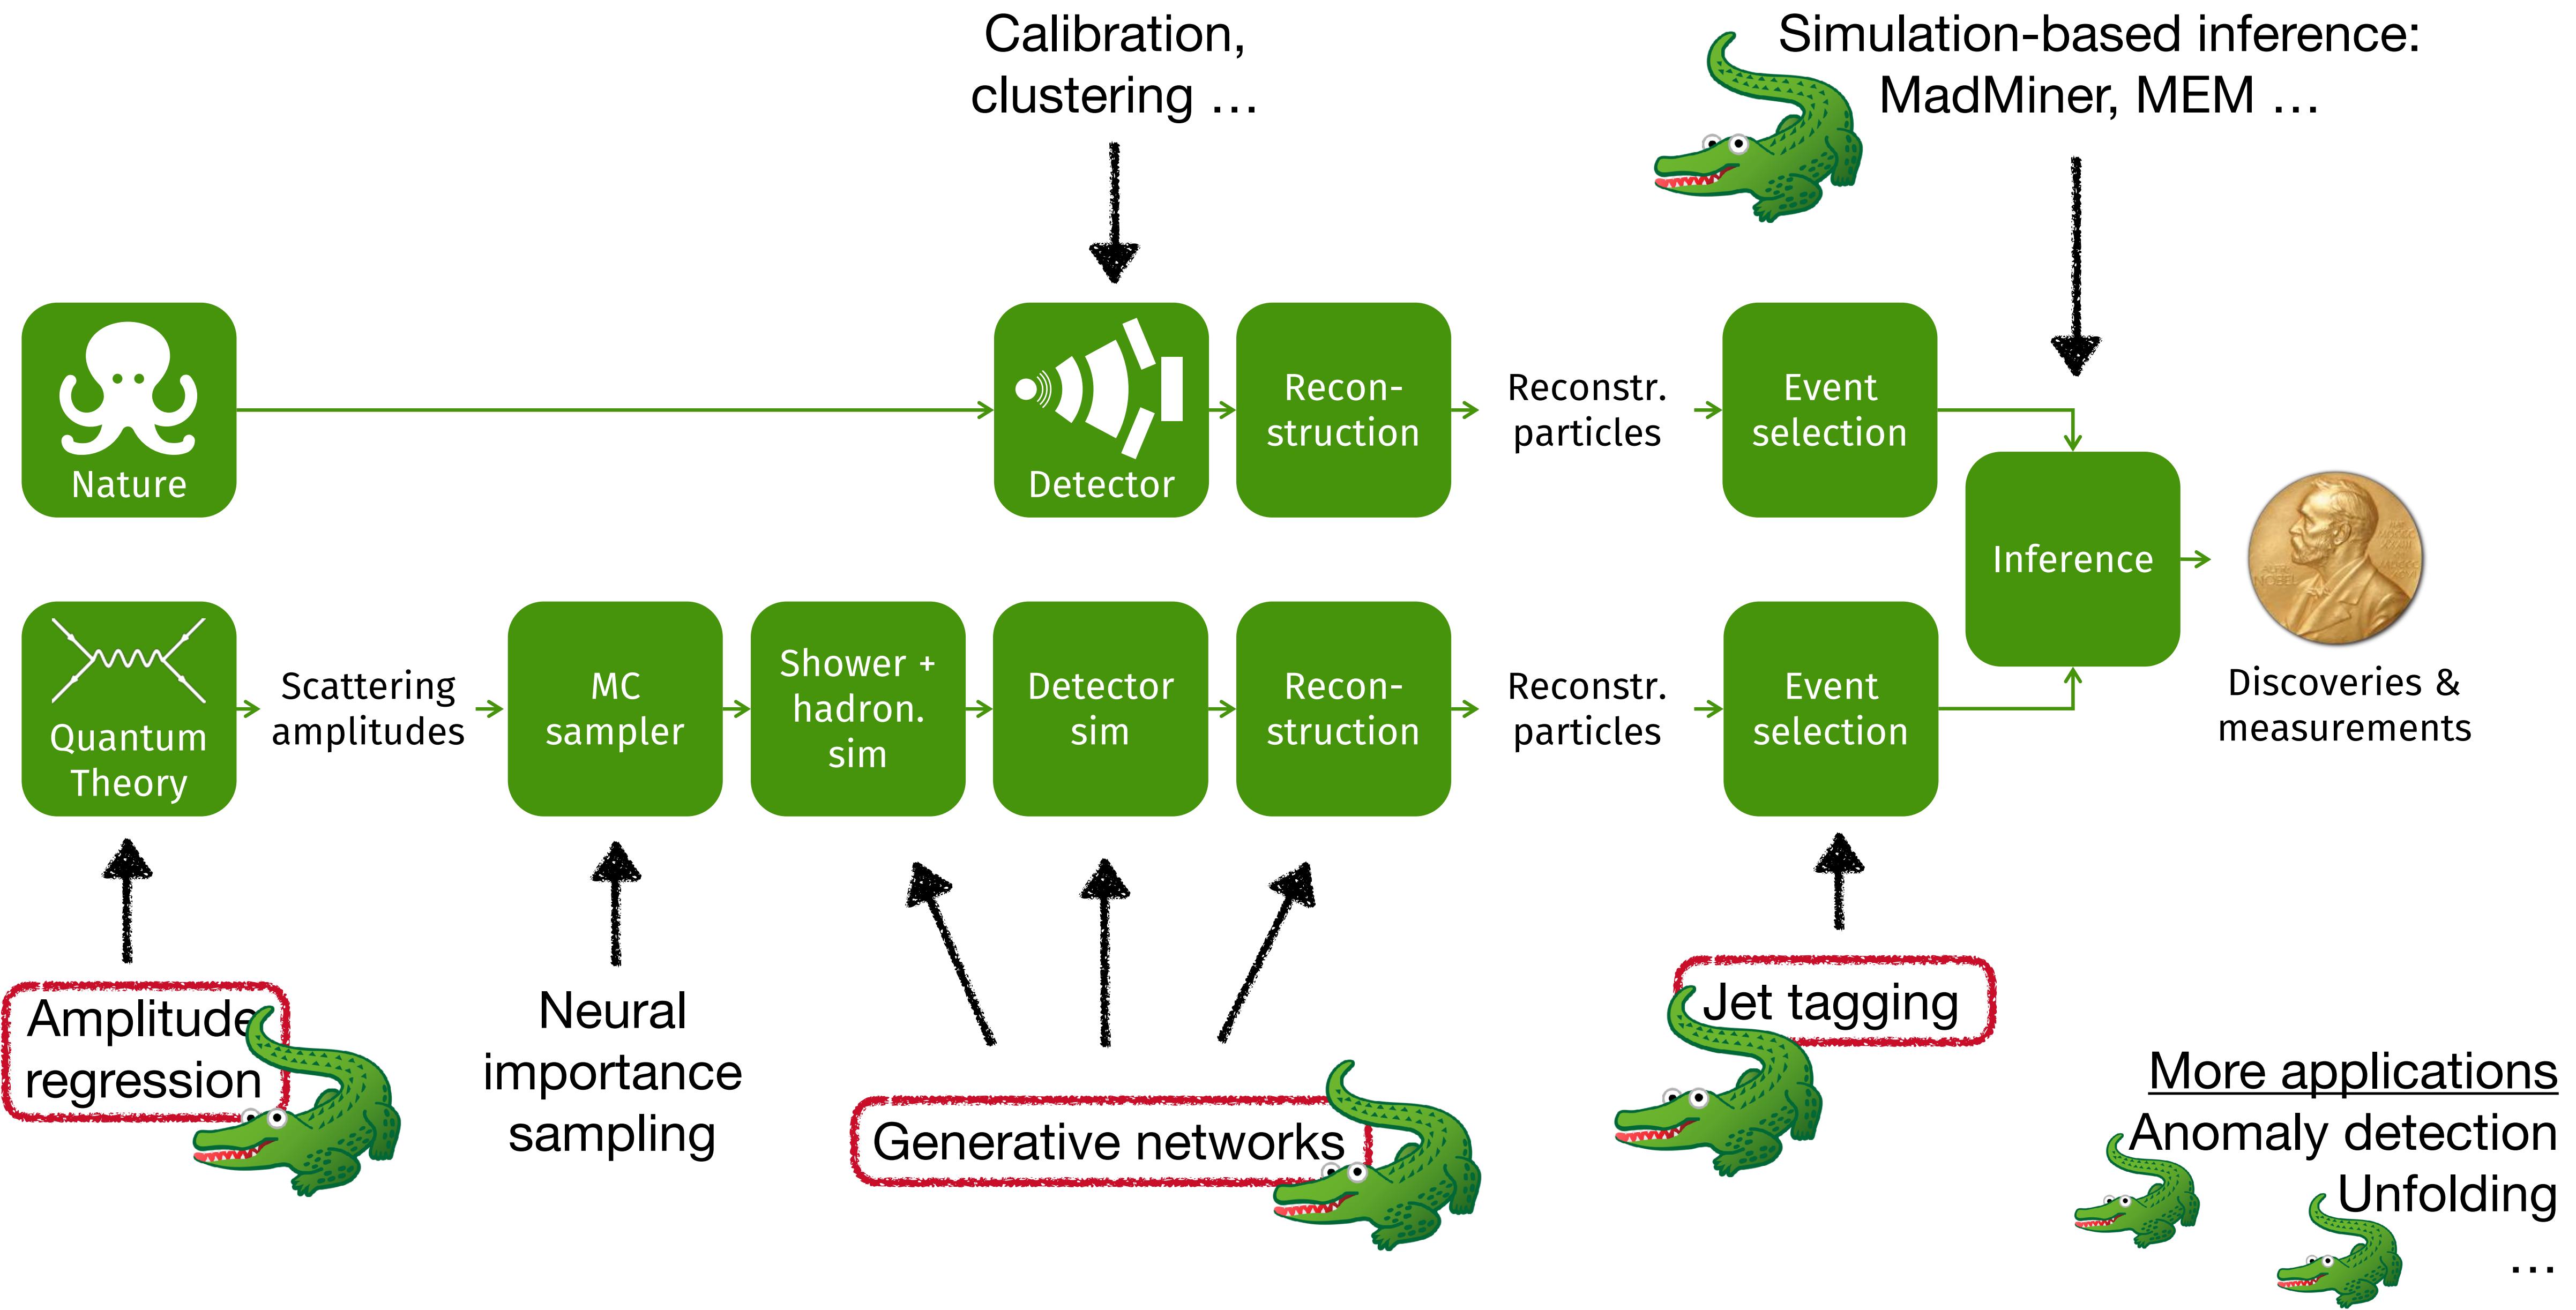
# Lorentz-Equivariant Geometric Algebra Transformer

- ✓ pip install lgatr
- ✓ Easy to use



Lorentz Local  
Canonicalization

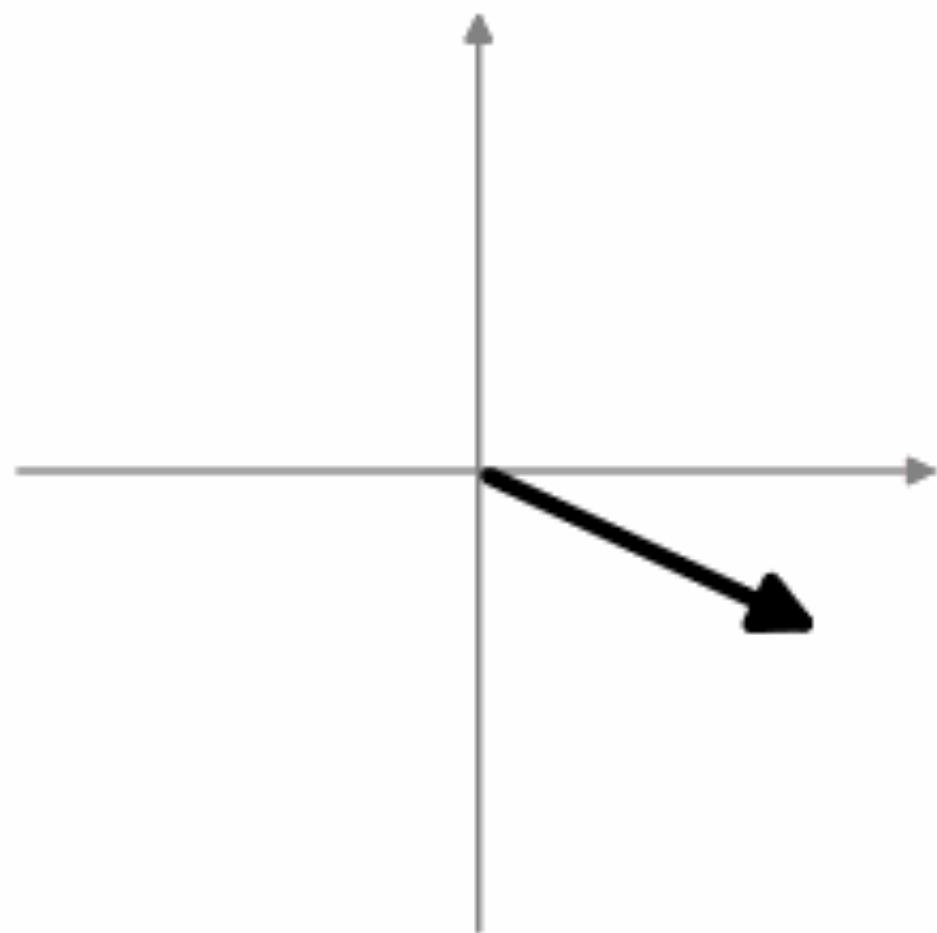
- ✓ Resource efficient
- ✓ Flexible



# Symmetries

## Invariance

Input



Output  
(not invariant)

0.8



Output  
(invariant)

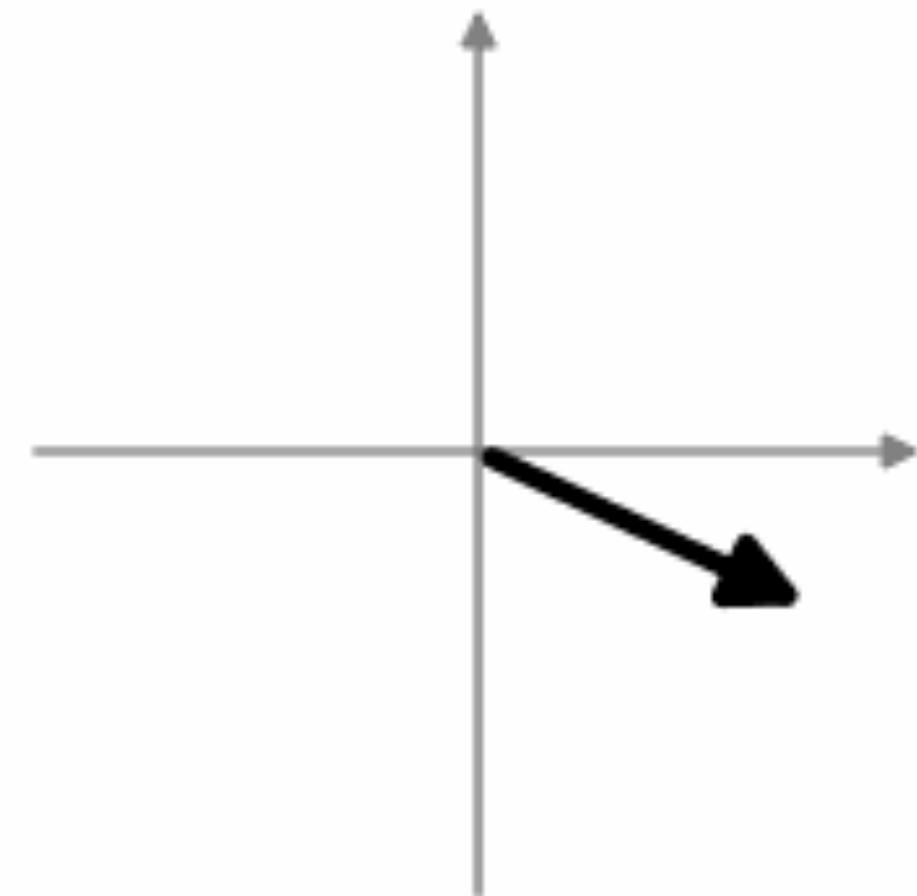
0.8



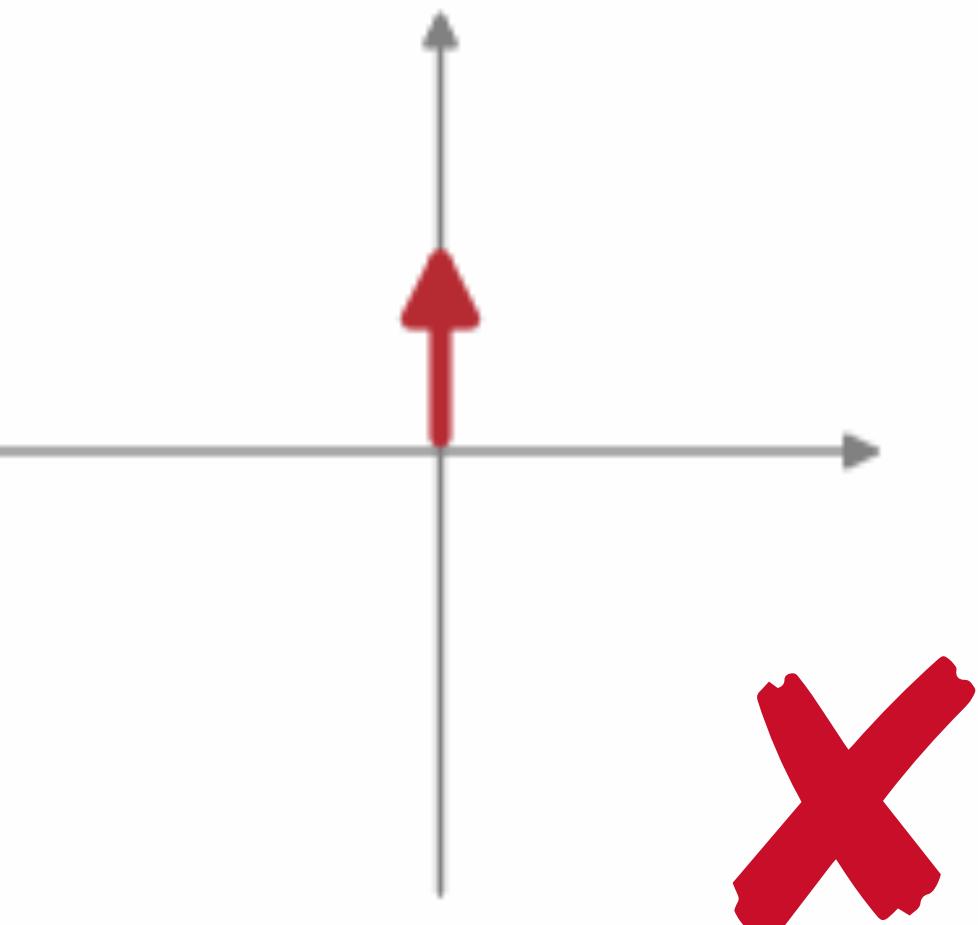
# Symmetries

**Equivariance = Covariance**

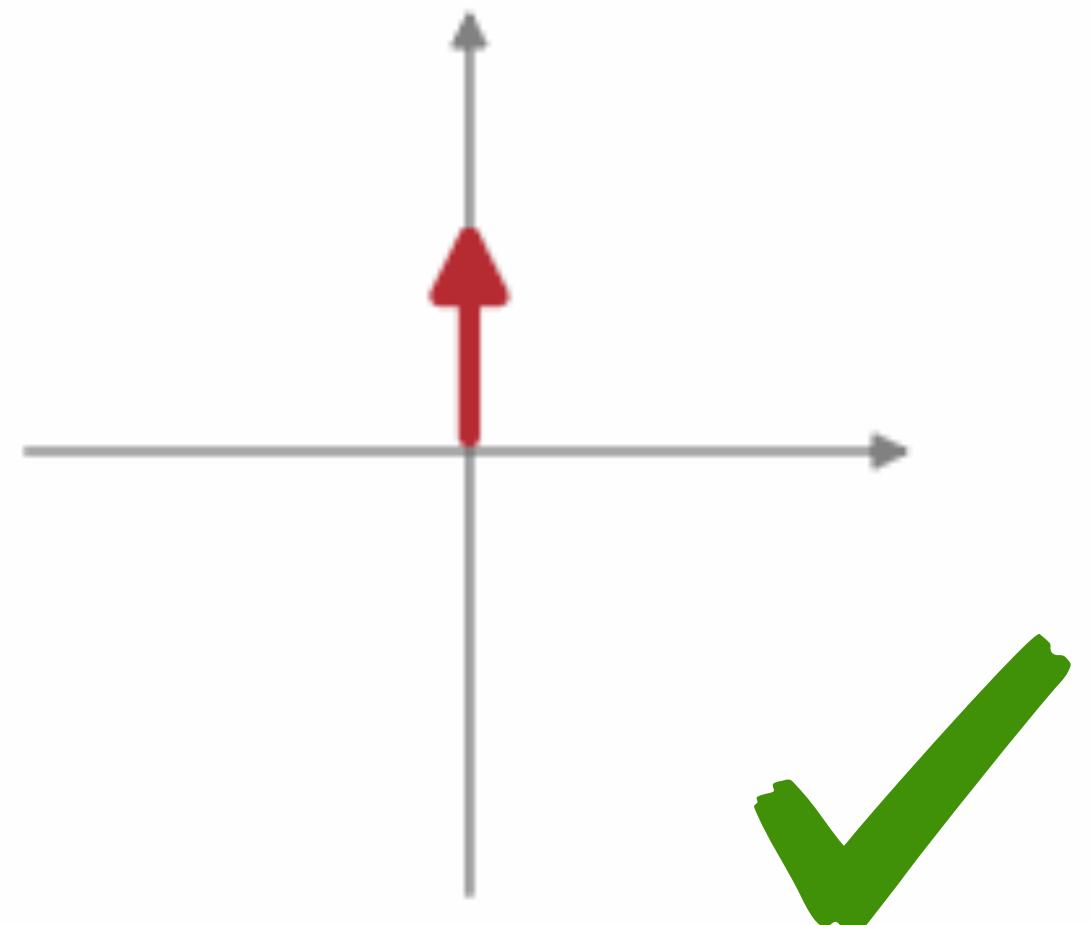
Input



Output  
(not equivariant)



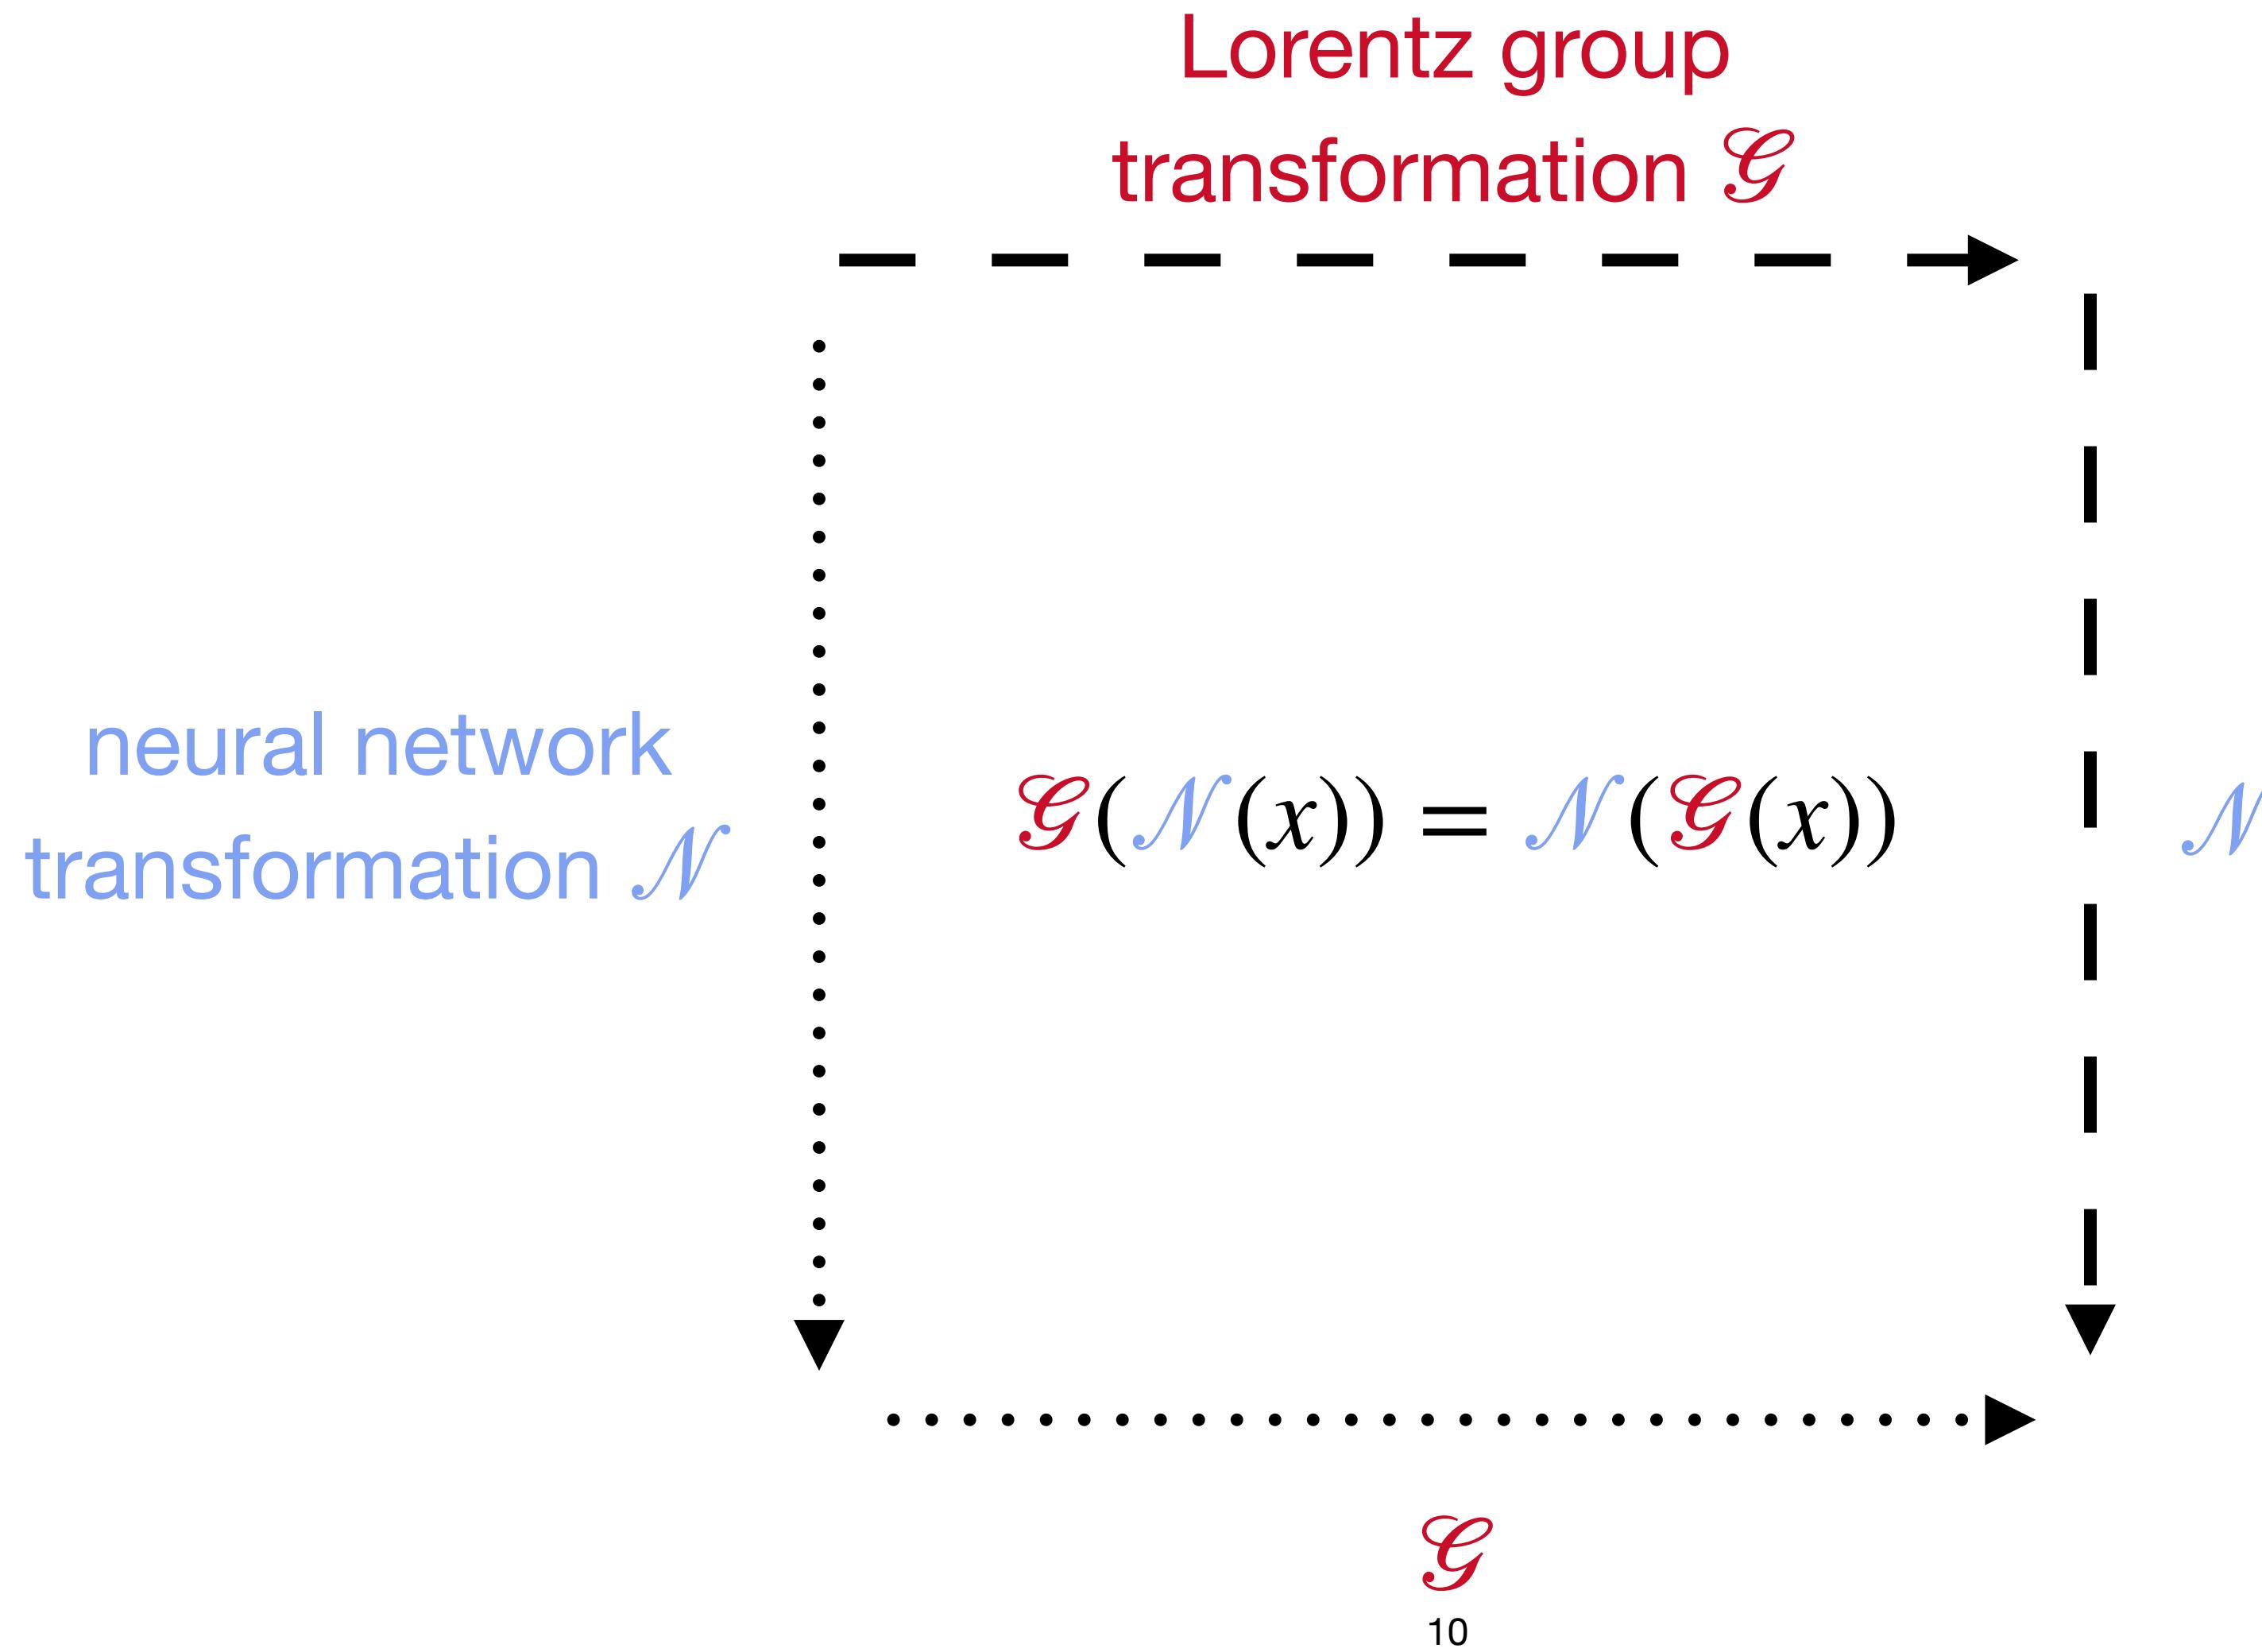
Output  
(equivariant)

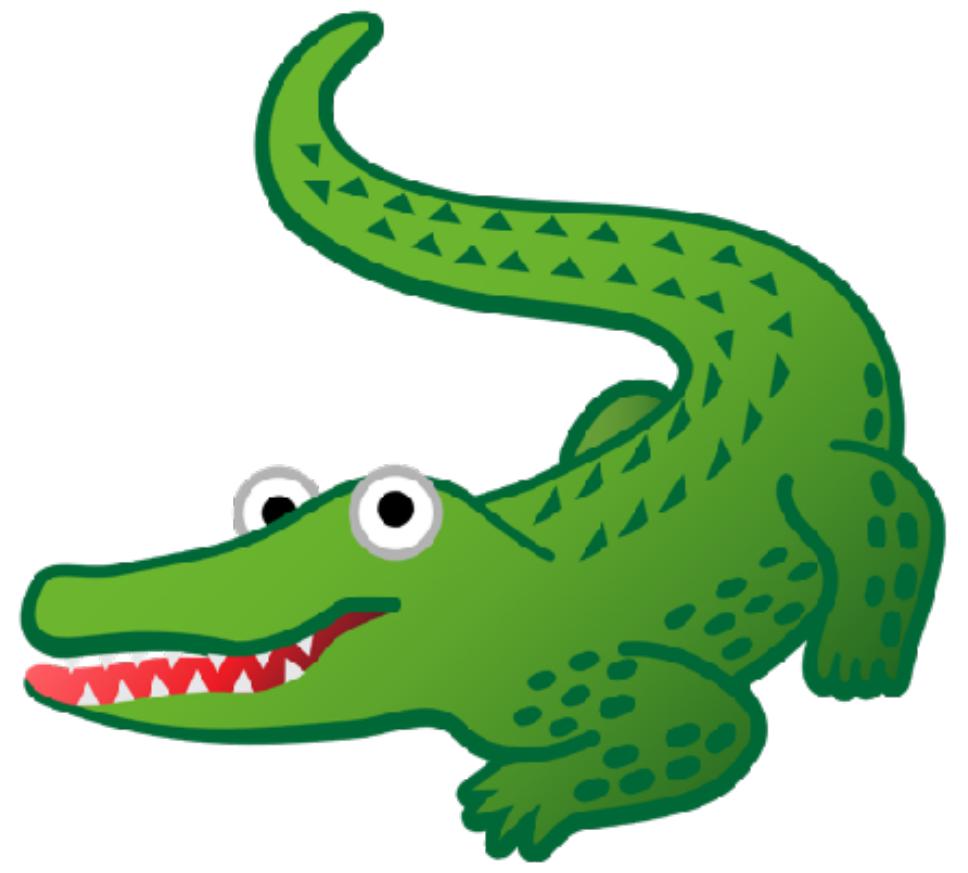


Invariance = Equivariance with scalar output

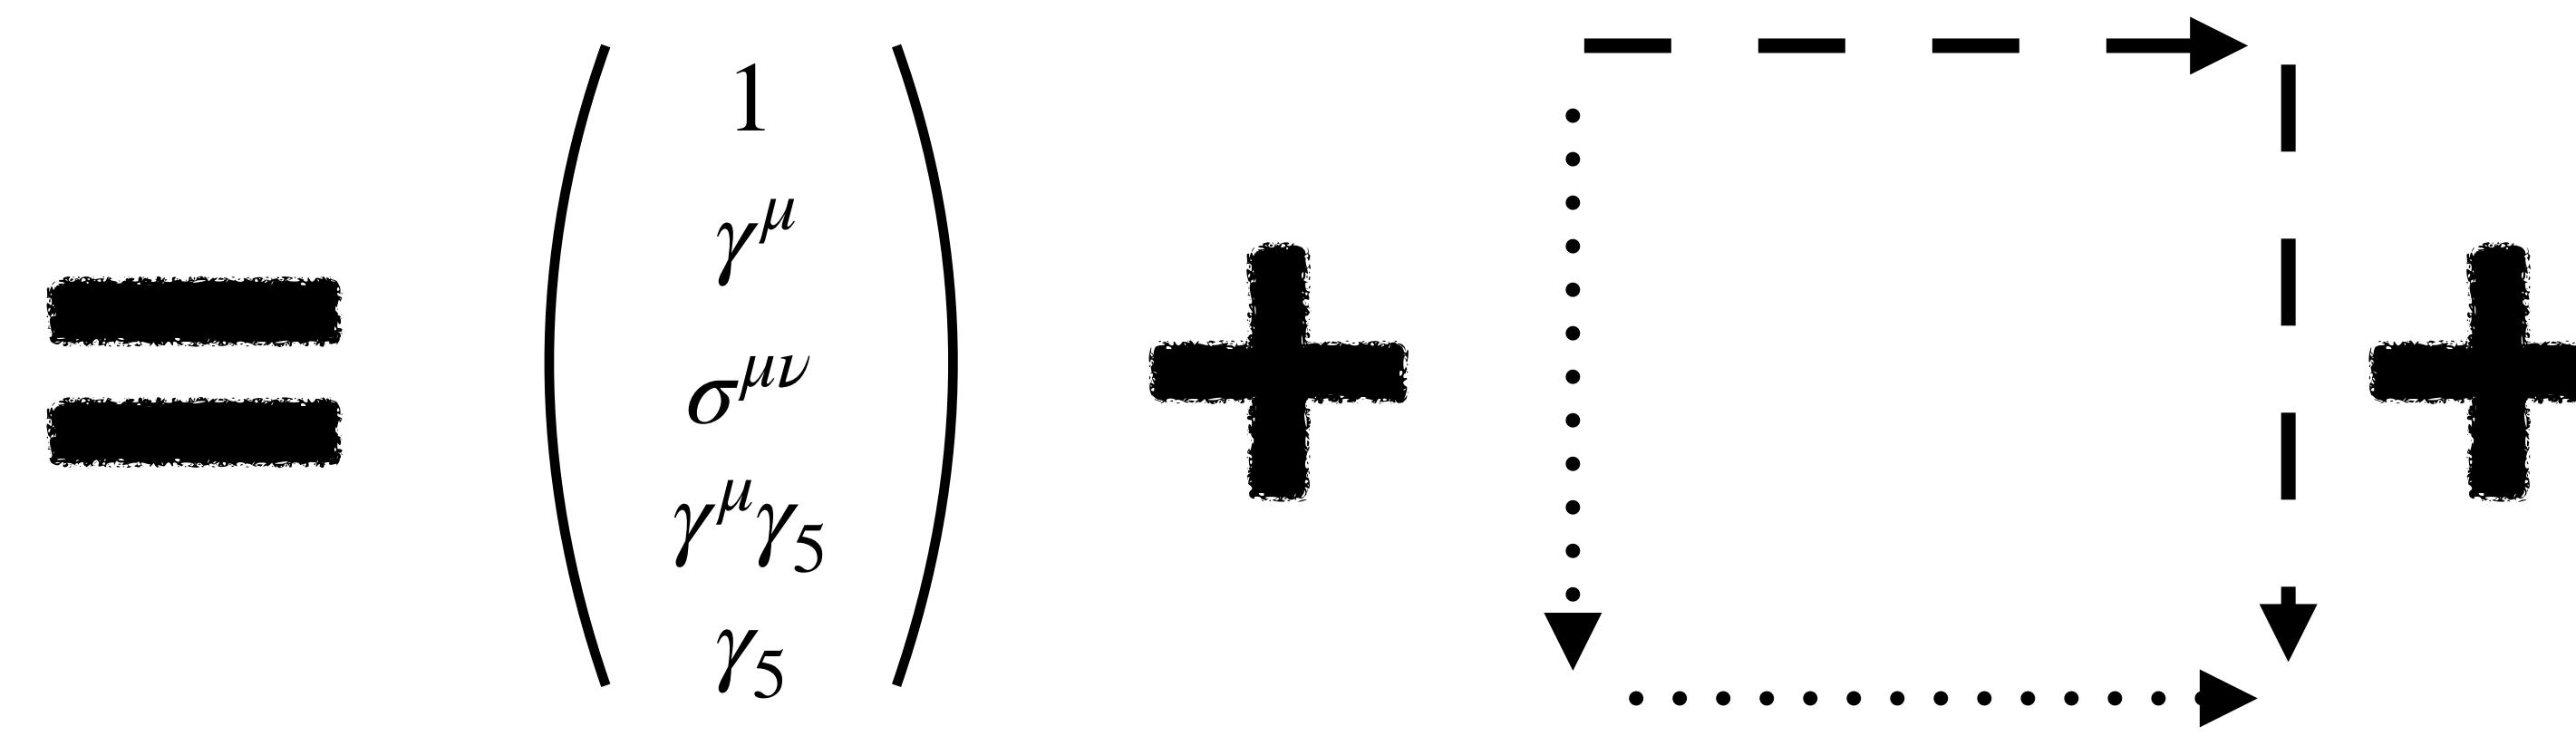
# Symmetries

## Equivariance - formalised



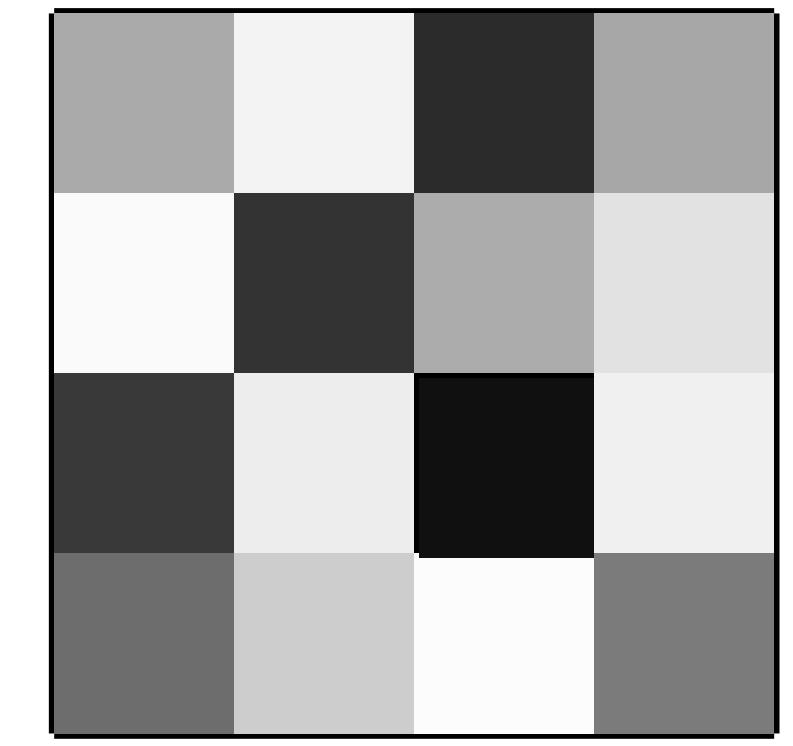


Lorentz-Equivariant  
Geometric **A**lgebra  
Transformer



**Geometric algebra**  
representations

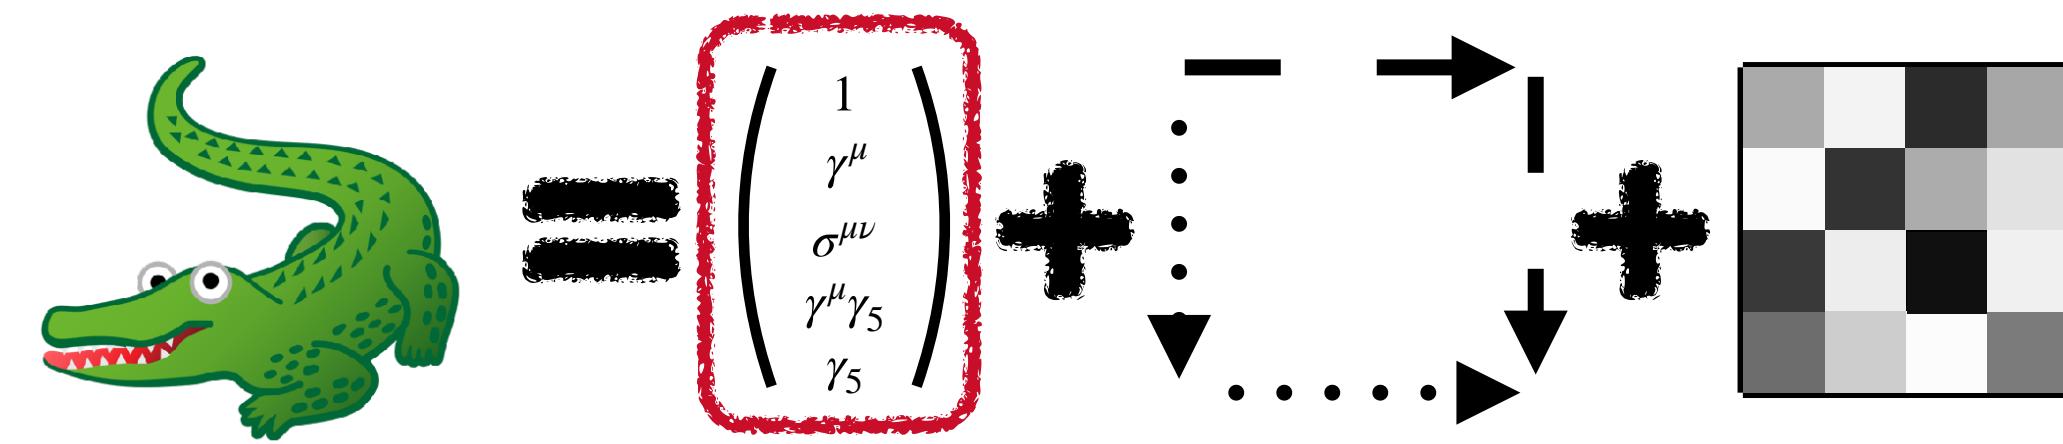
**Lorentz-Equivariant**  
layers



**Transformer**  
architecture

GATr was originally  
developed for E(3)  
arXiv:2305.18415

# L-GATr



**Geometric algebra = Clifford algebra**

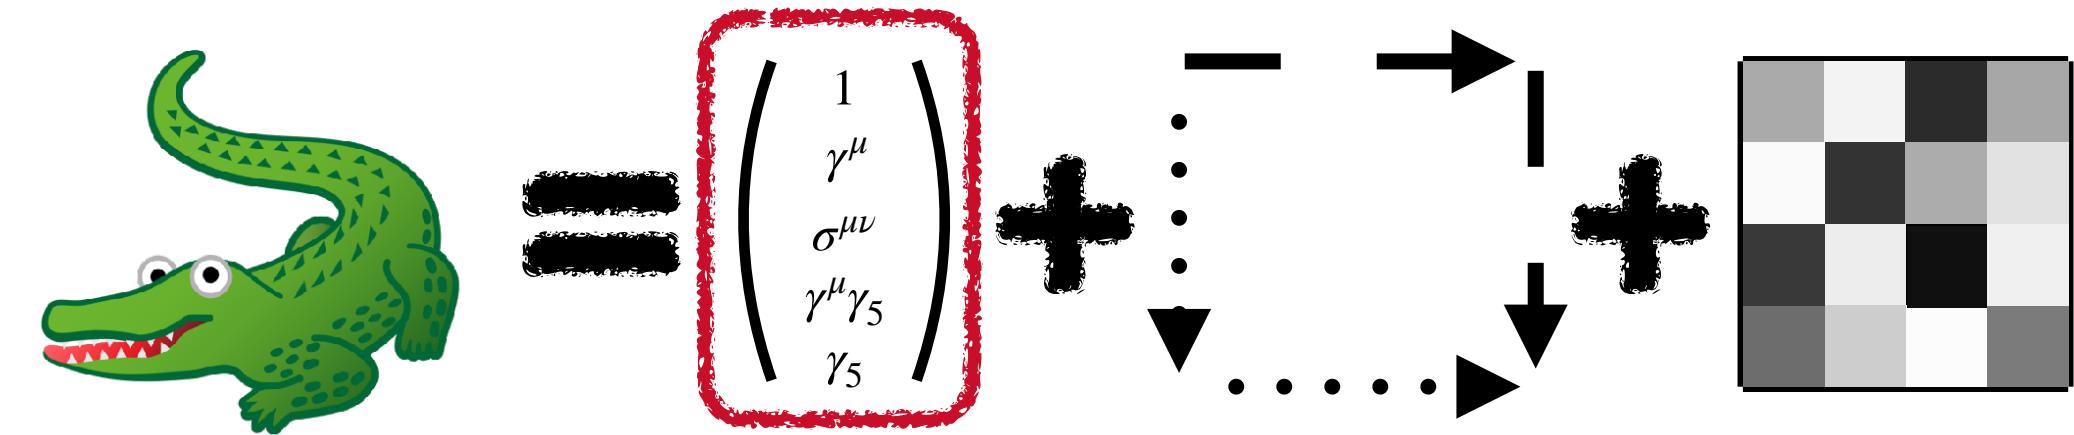
## Geometric algebra

= Vector space + geometric product  $xy = \frac{\{x, y\}}{2} + \frac{[x, y]}{2}$

## Spacetime geometric algebra

= Minkowski space +  $\{\gamma^\mu, \gamma^\nu\} = 2g^{\mu\nu}\mathbf{1}$  over  $\mathbb{R}$   
≠ Dirac algebra (uses  $\mathbb{C}$ )

# L-GATr



## Building the algebra

$$\{\gamma^\mu, \gamma^\nu\} = 2g^{\mu\nu}1$$

Linear combinations  
+ Geometric product

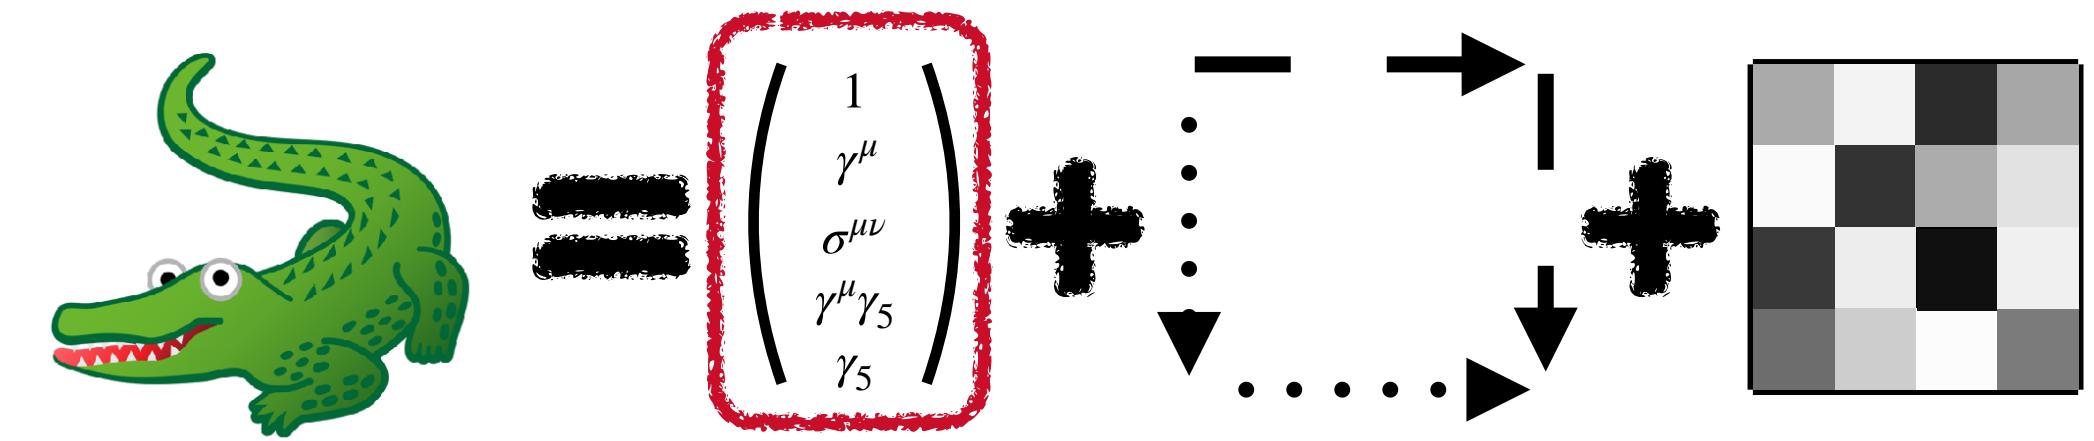
$$x = x^S 1 + x_\mu^V \gamma^\mu + x_{\mu\nu}^B \sigma^{\mu\nu} + x_\mu^A \gamma^\mu \gamma^5 + x^P \gamma^5$$

**Multivector** representation:  $(x^S, x_\mu^V, x_{\mu\nu}^B, x_\mu^A, x^P) \in \mathbb{R}^{16}$

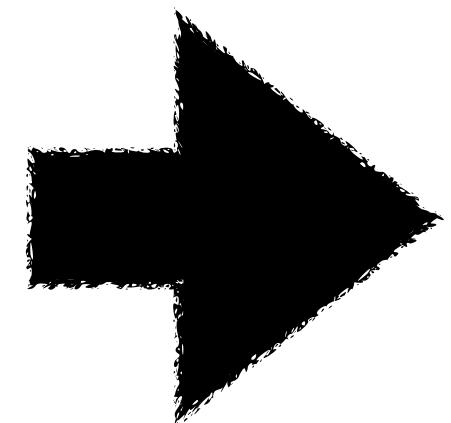
Particle:  $x^V = (E, p_x, p_y, p_z)$ ,  $x^S = \text{PID}$ ,  $x^{\text{others}} = 0$

# L-GATr

## Geometric algebra representations



$x^S \in \mathbb{R}$



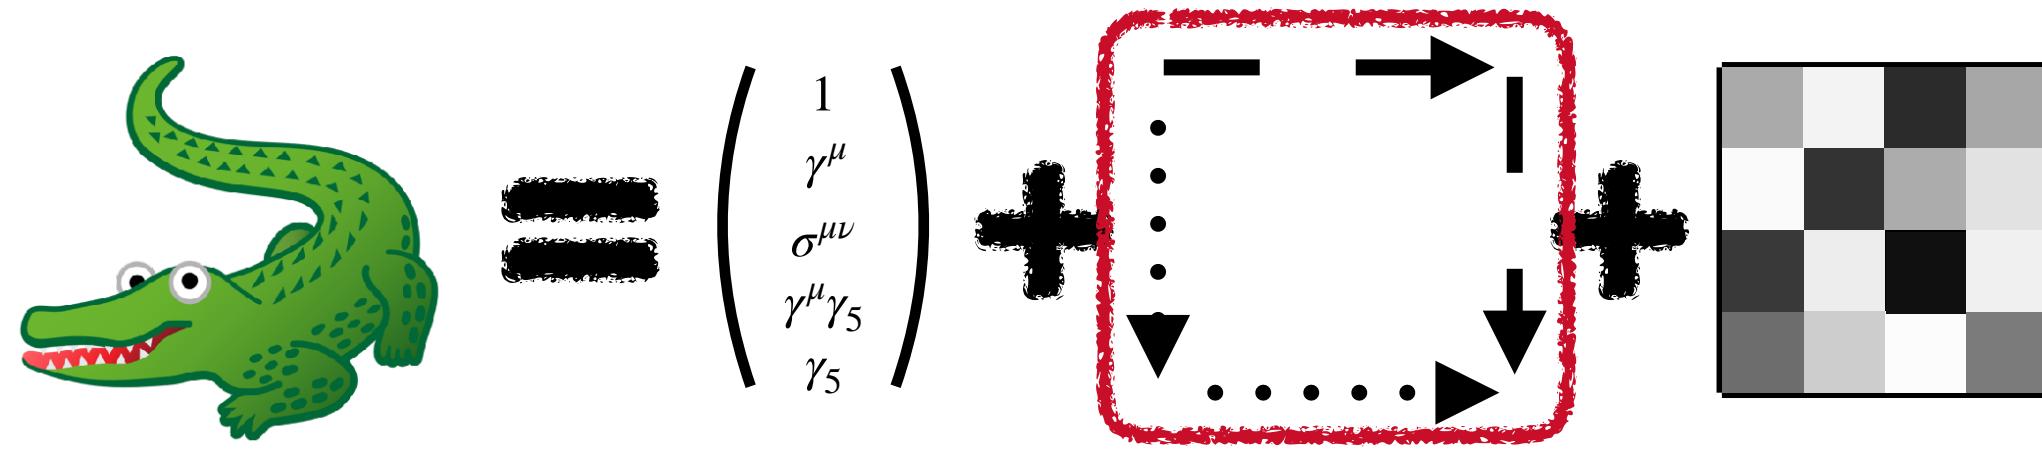
scalar channels

$$\begin{array}{c} x^S \\ x_\mu^V \\ x_{\mu\nu}^B \\ x_\mu^A \\ x_\nu^P \end{array}$$

$\in \mathbb{R}^{16}$

multivector channels

# L-GATr

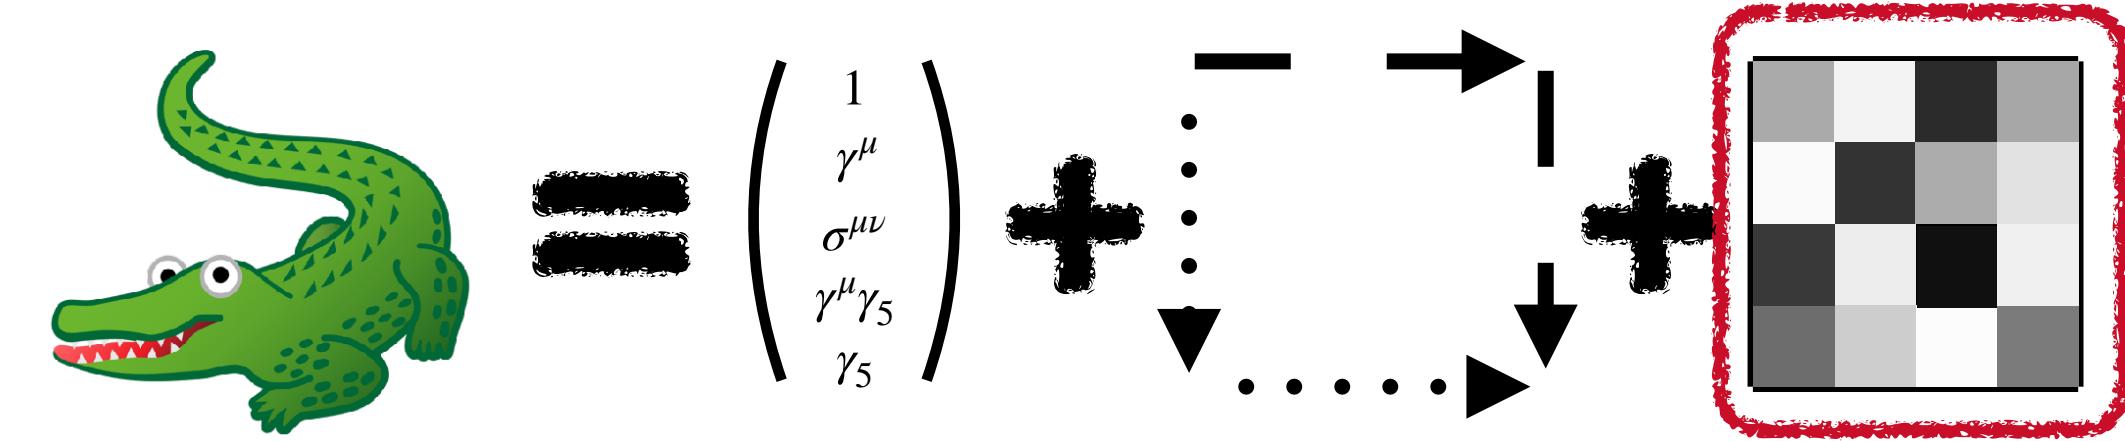


## L-GATr-ing all transformer layers

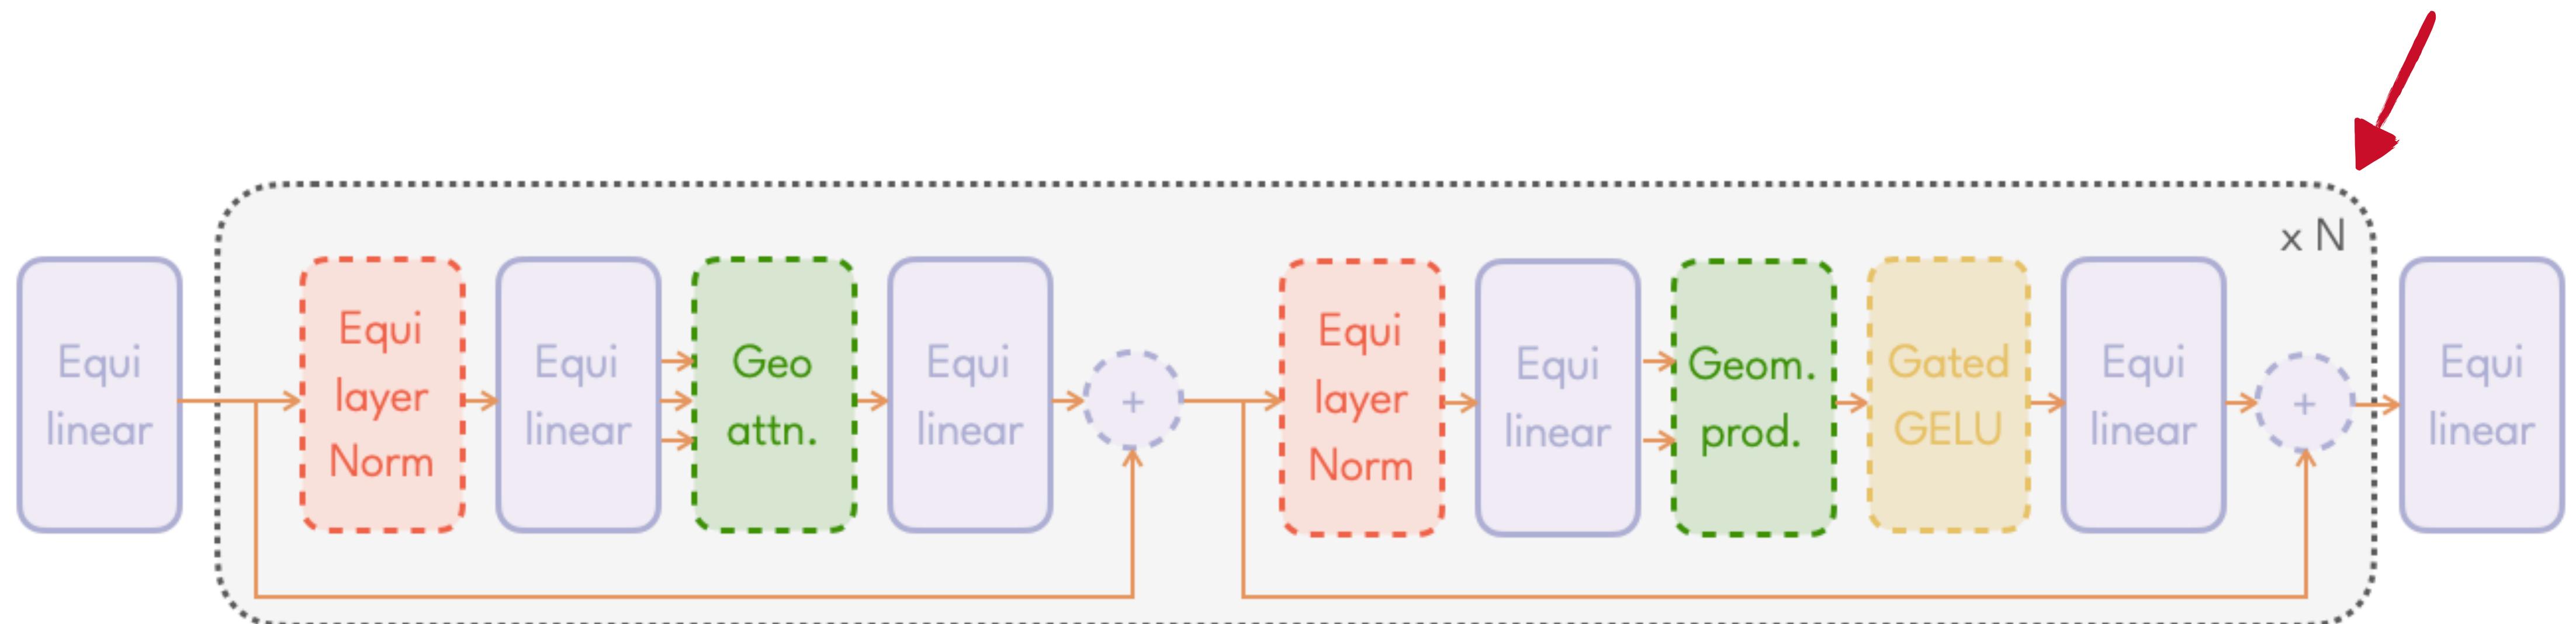
Layer type	Transformer	L-GATr
Linear( $x$ )	$vx + w$	$\sum_{k=0}^4 v_k \langle x \rangle_k \left( + \sum_{k=0}^4 w_k \gamma^5 \langle x \rangle_k \right)$
Attention( $q, k, v$ ) <sub>ic</sub>	$\sum_{j=1}^{n_t} \text{Softmax}_j \left( \sum_{c'=1}^{n_c} \frac{q_{ic'} k_{jc'}}{\sqrt{n_c}} \right) v_{jc}$	$\sum_{j=1}^{n_t} \text{Softmax}_j \left( \sum_{c'=1}^{n_c} \frac{\langle q_{ic'}, k_{jc'} \rangle}{\sqrt{16n_c}} \right) v_{jc}$
LayerNorm( $x$ )	$x \left[ \frac{1}{n_c} \sum_{c=1}^{n_c} x_c^2 + \epsilon \right]^{-1/2}$	$x \left[ \frac{1}{n_c} \sum_{c=1}^{n_c} \sum_{k=0}^4 \left  \langle \langle x_c \rangle_k, \langle x_c \rangle_k \rangle \right  + \epsilon \right]^{-1/2}$
Activation( $x$ )	GELU( $x$ )	GELU( $\langle x \rangle_0$ ) $x$
GP( $x, y$ )	—	$xy$

# L-GATr

## Full architecture



**L-GATr blocks**  
can be stacked  
to large depth



Equivariant Self-Attention

Equivariant MLP

## USAGE

[Quickstart](#)[Spacetime Geometric Algebra](#)[Attention Backends](#)[Lorentz Symmetry Breaking](#)

## REFERENCE

[API Reference](#)

# Quickstart

This page sets you up for building and running L-GATr models.

## Installation

Before using the package, install it via pip:

```
pip install lgatr
```

Alternatively, if you're developing locally:

```
git clone https://github.com/heidelberg-hepml/lgatr.git
cd lgatr
pip install -e .
```

## Building L-GATr

You can construct a simple `LGATr` model as follows:

```
from lgatr import LGATr
attention = dict(num_heads=2)
mlp = dict()
```

```
from lgatr import LGATr
attention = dict(num_heads=2)
mlp = dict()
lgatr = LGATr(
    in_mv_channels=1,
    out_mv_channels=1,
    hidden_mv_channels=8,
    in_s_channels=0,
    out_s_channels=0,
    hidden_s_channels=16,
    attention=attention,
    mlp=mlp,
    num_blocks=2,
)
```

```
from lgatr.interface import embed_vector, extract_scalar
multivector = embed_vector(p)
print(multivector.shape) # torch.Size([128, 20, 1, 16])
```

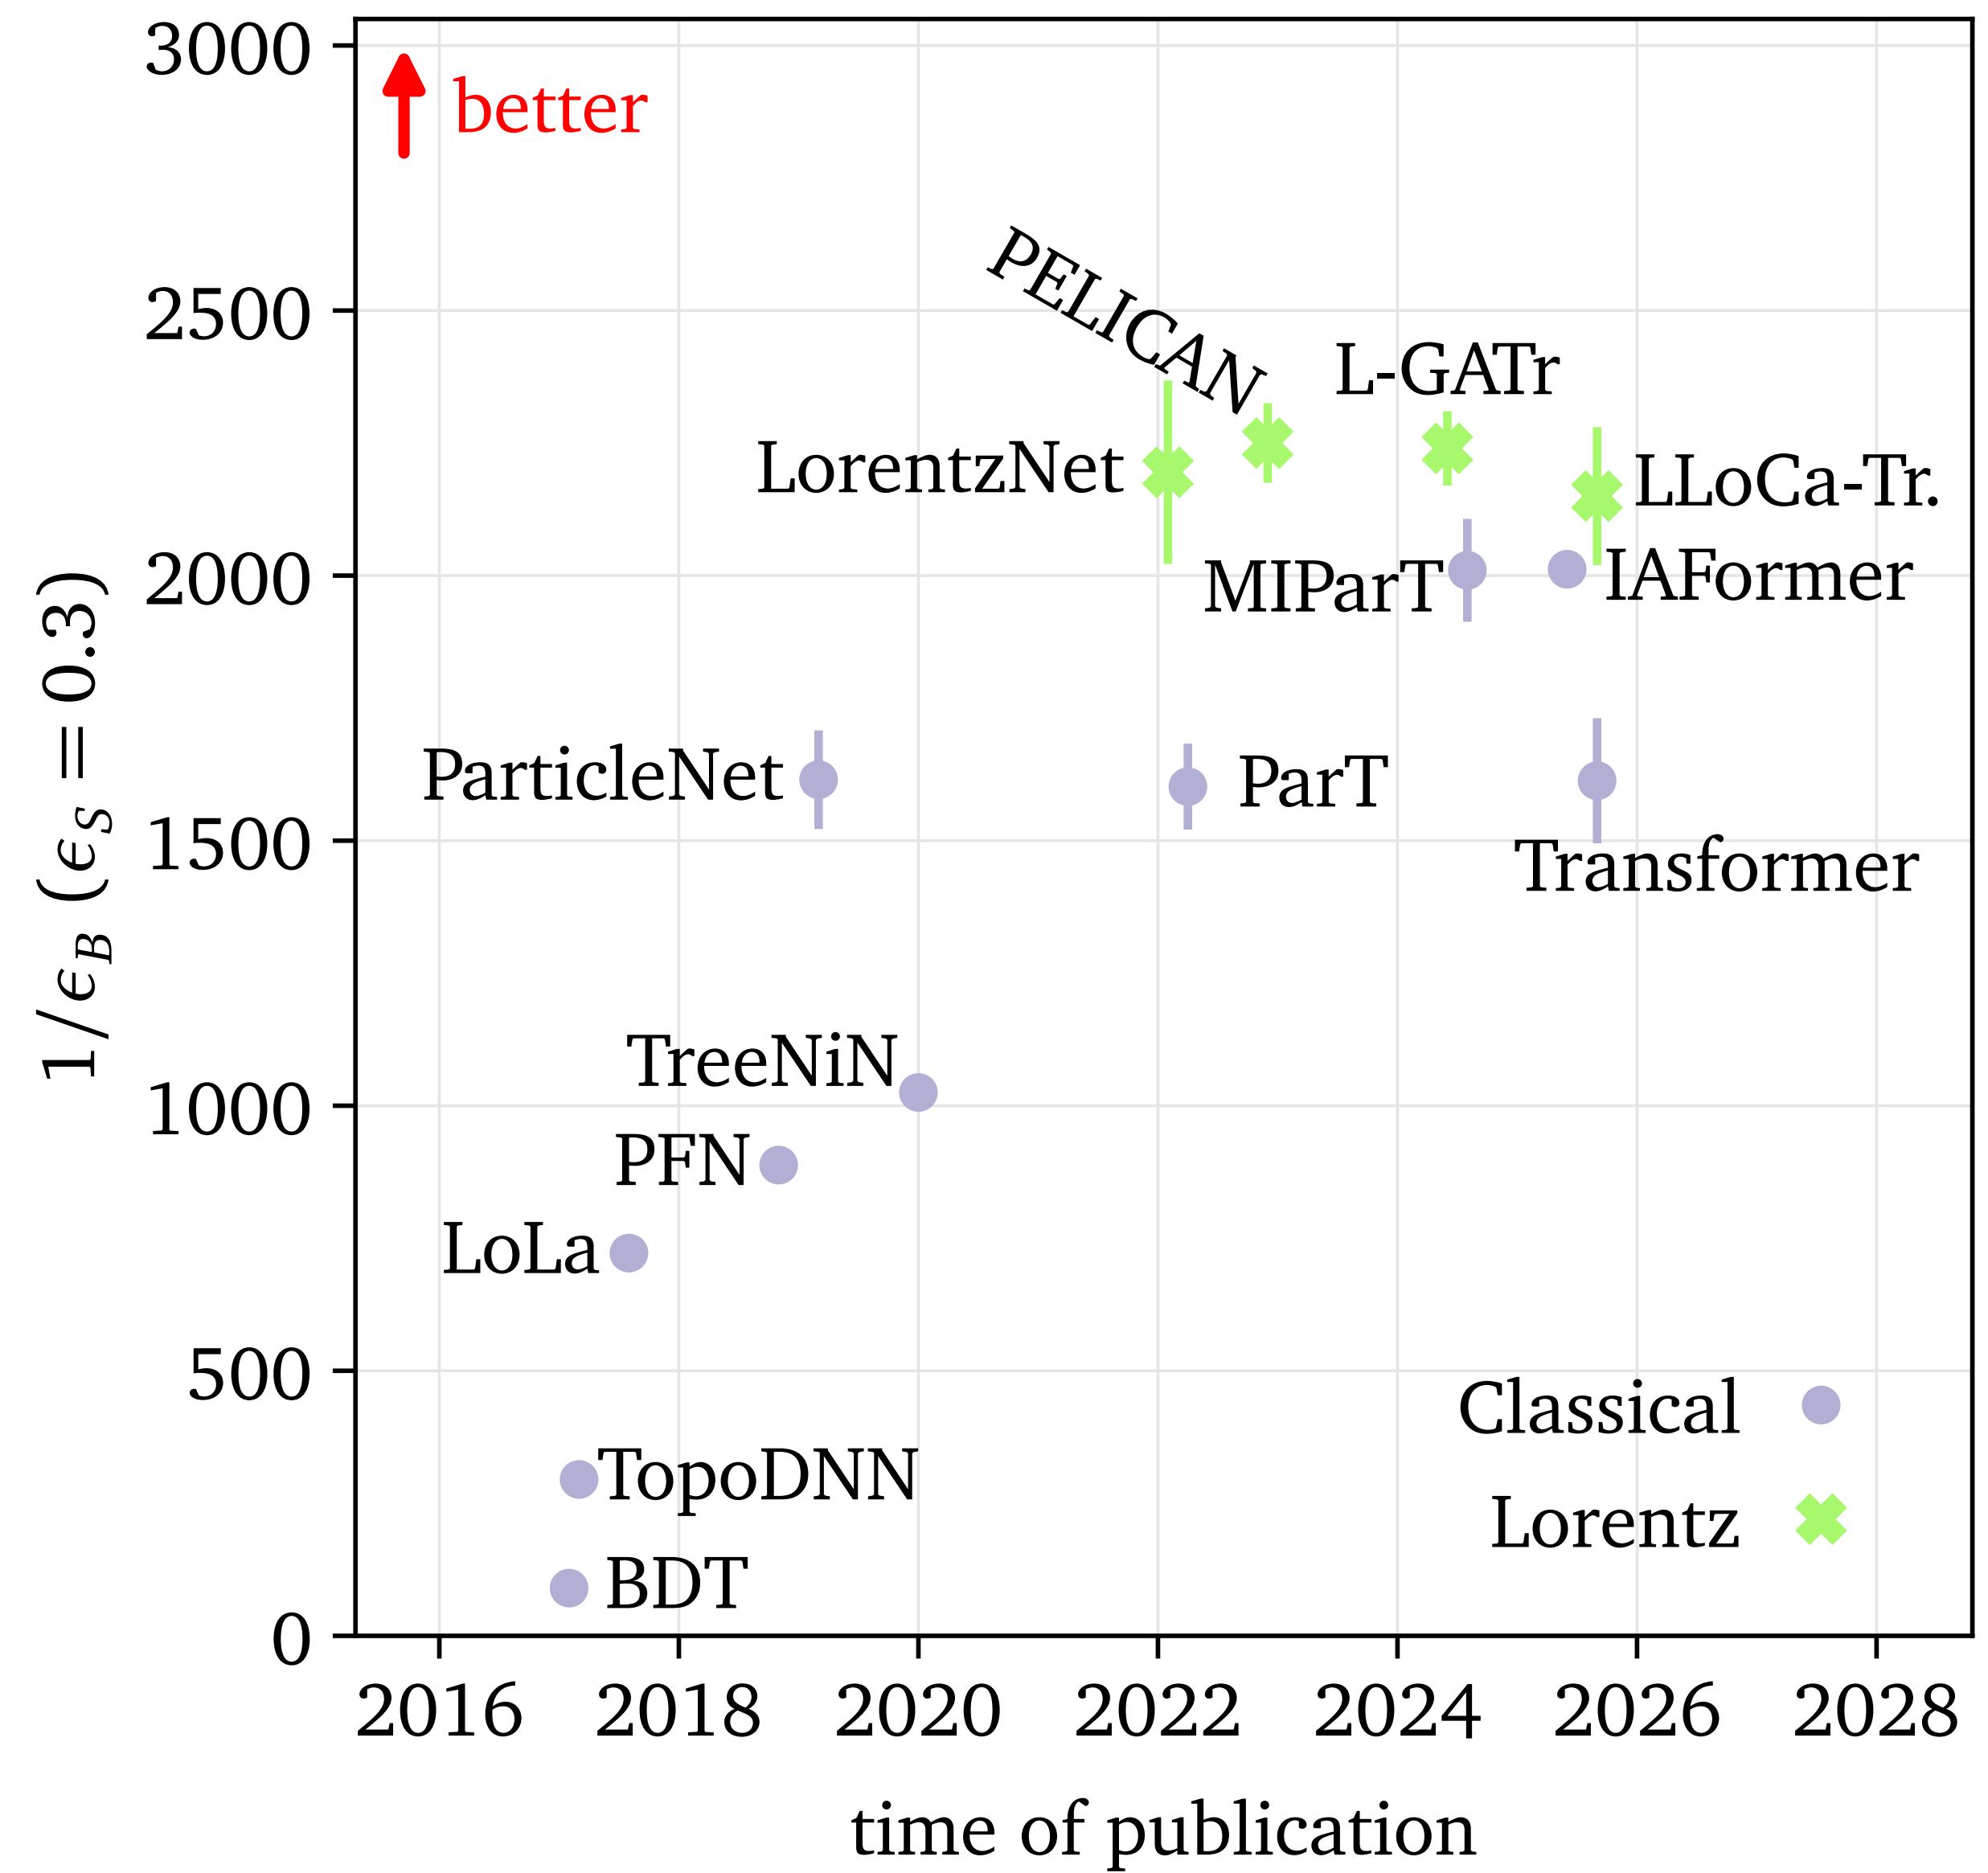
```
output_mv, output_s = lgatr(multivectors=multivector, scalars=None)
out = extract_scalar(output_mv)
print(out.shape) # torch.Size([128, 20, 1, 1])
```

# Jet Tagging

## The history of top tagging

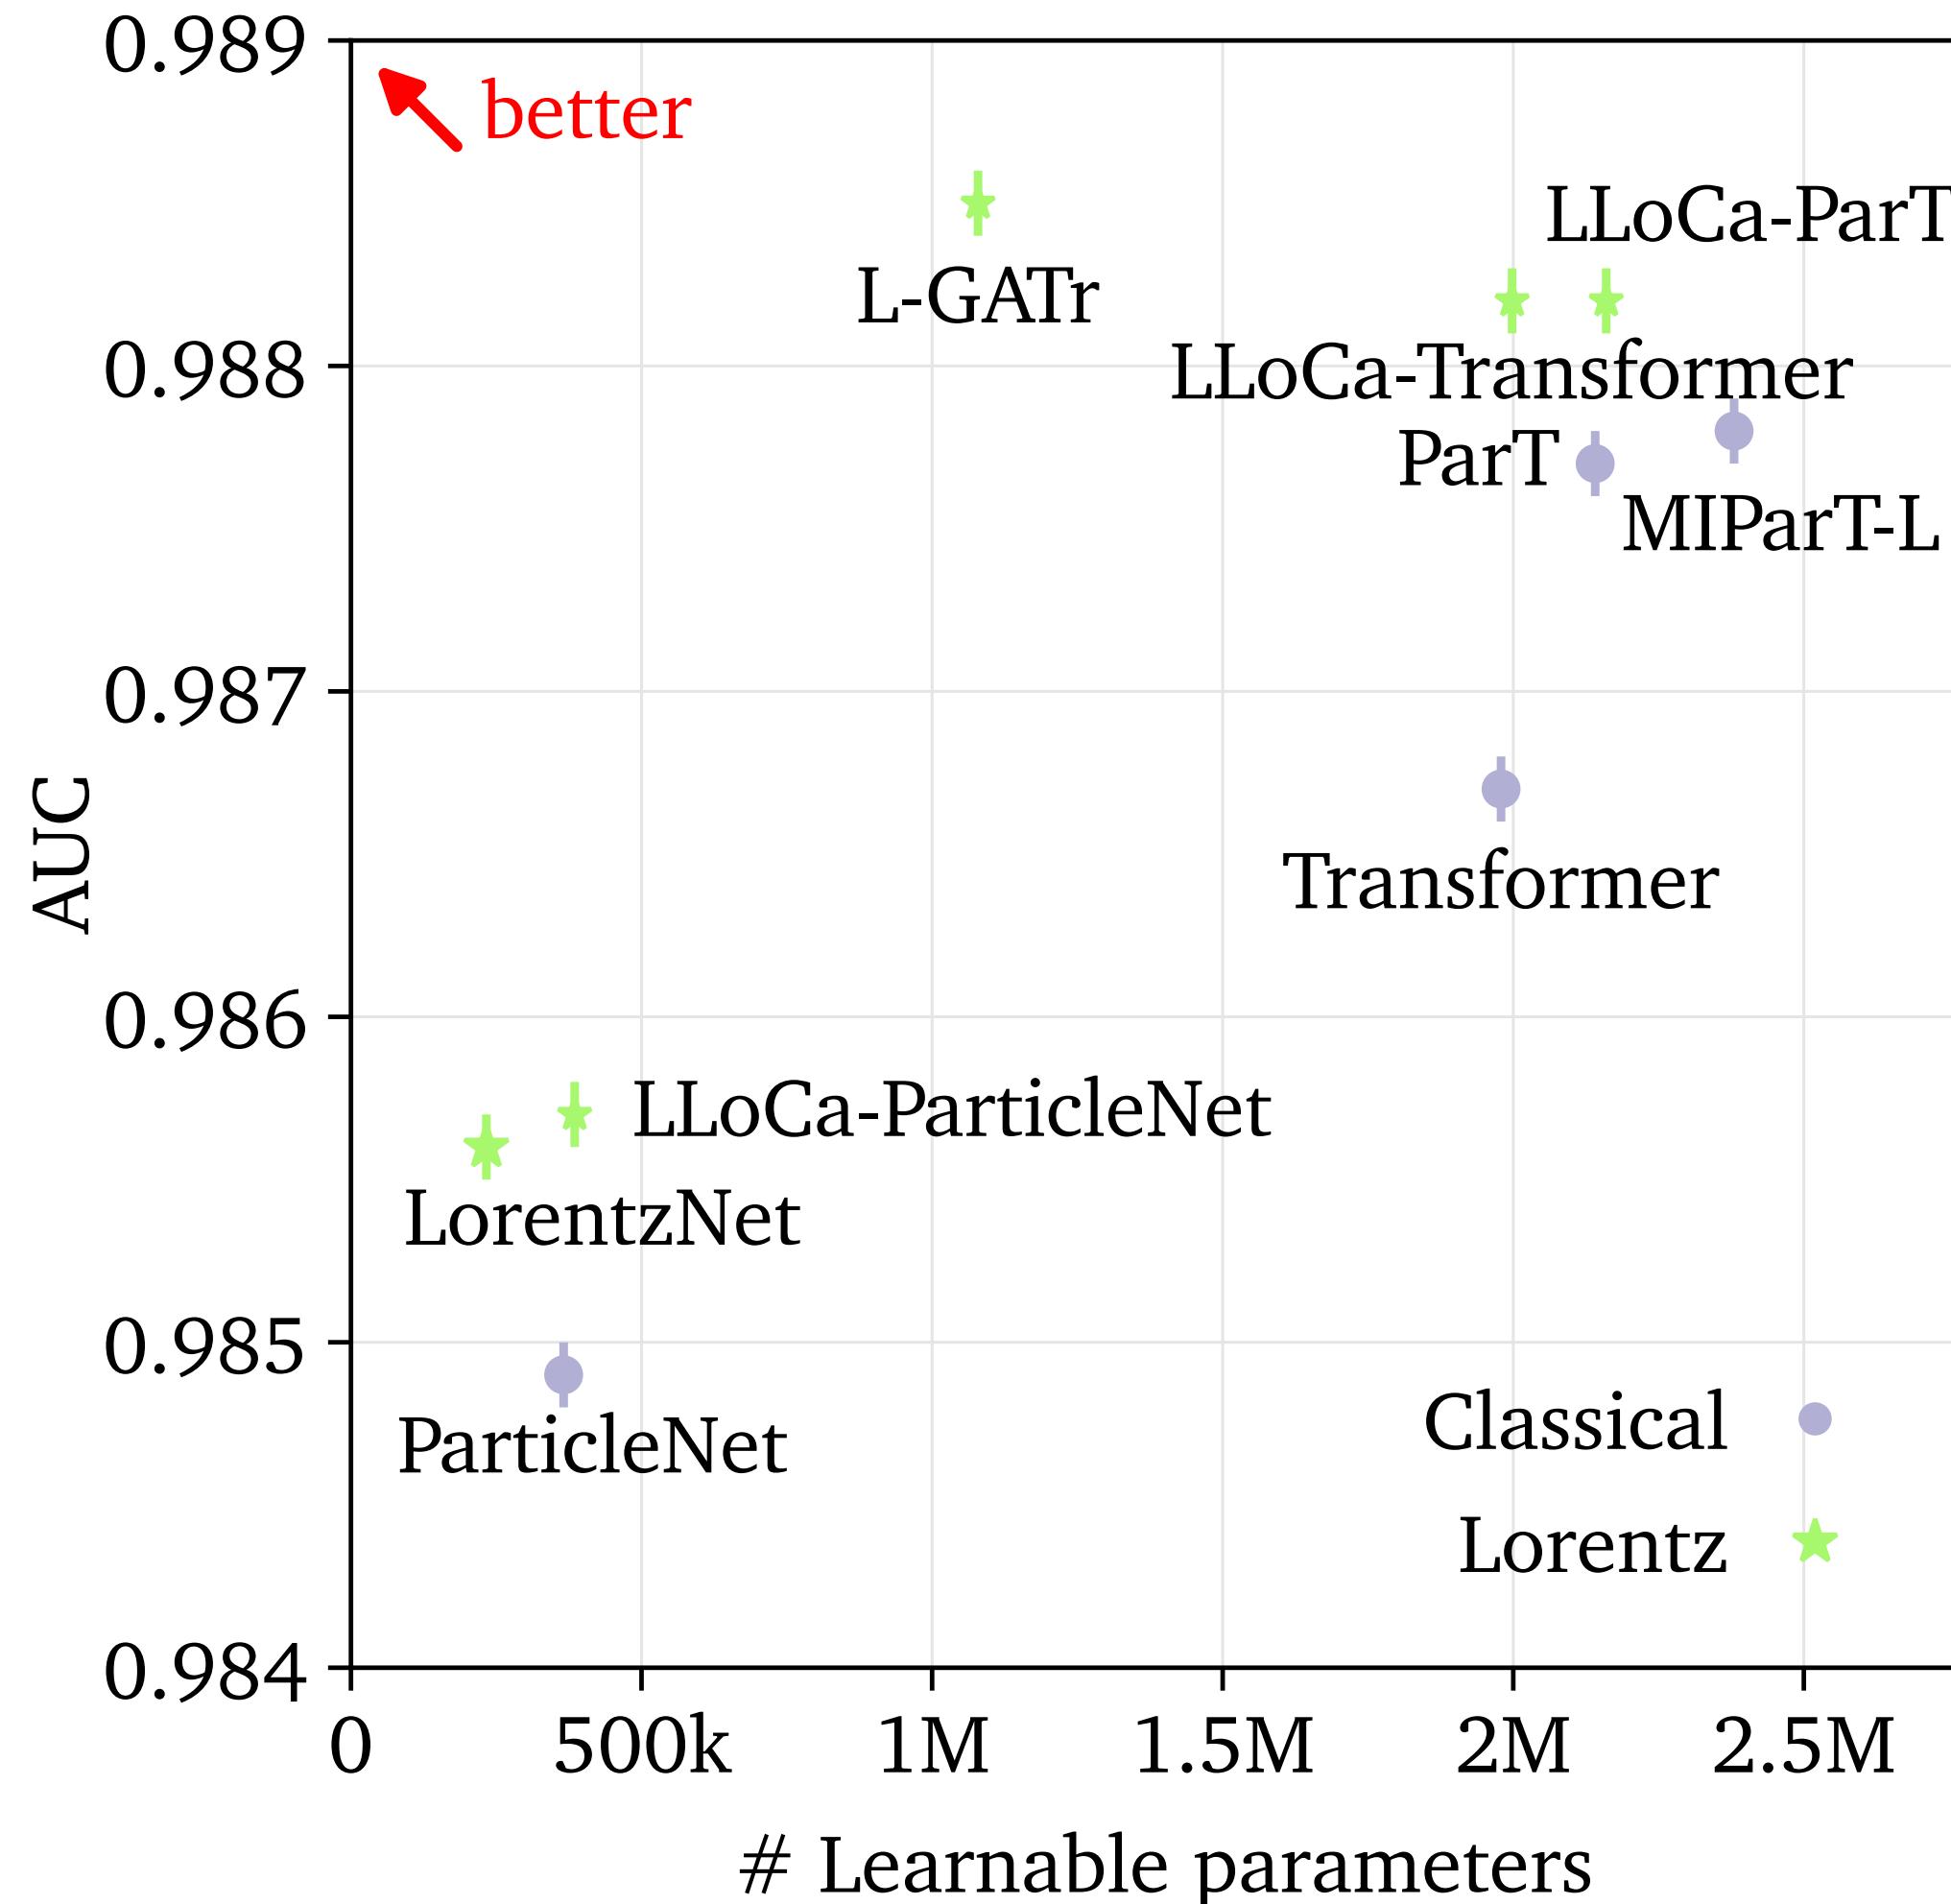
All top-performing taggers are Lorentz-equivariant

Graph networks and transformers perform similarly well



# Jet Tagging

## Training on 100M jets



Transformers outscale  
graph networks on  
large datasets

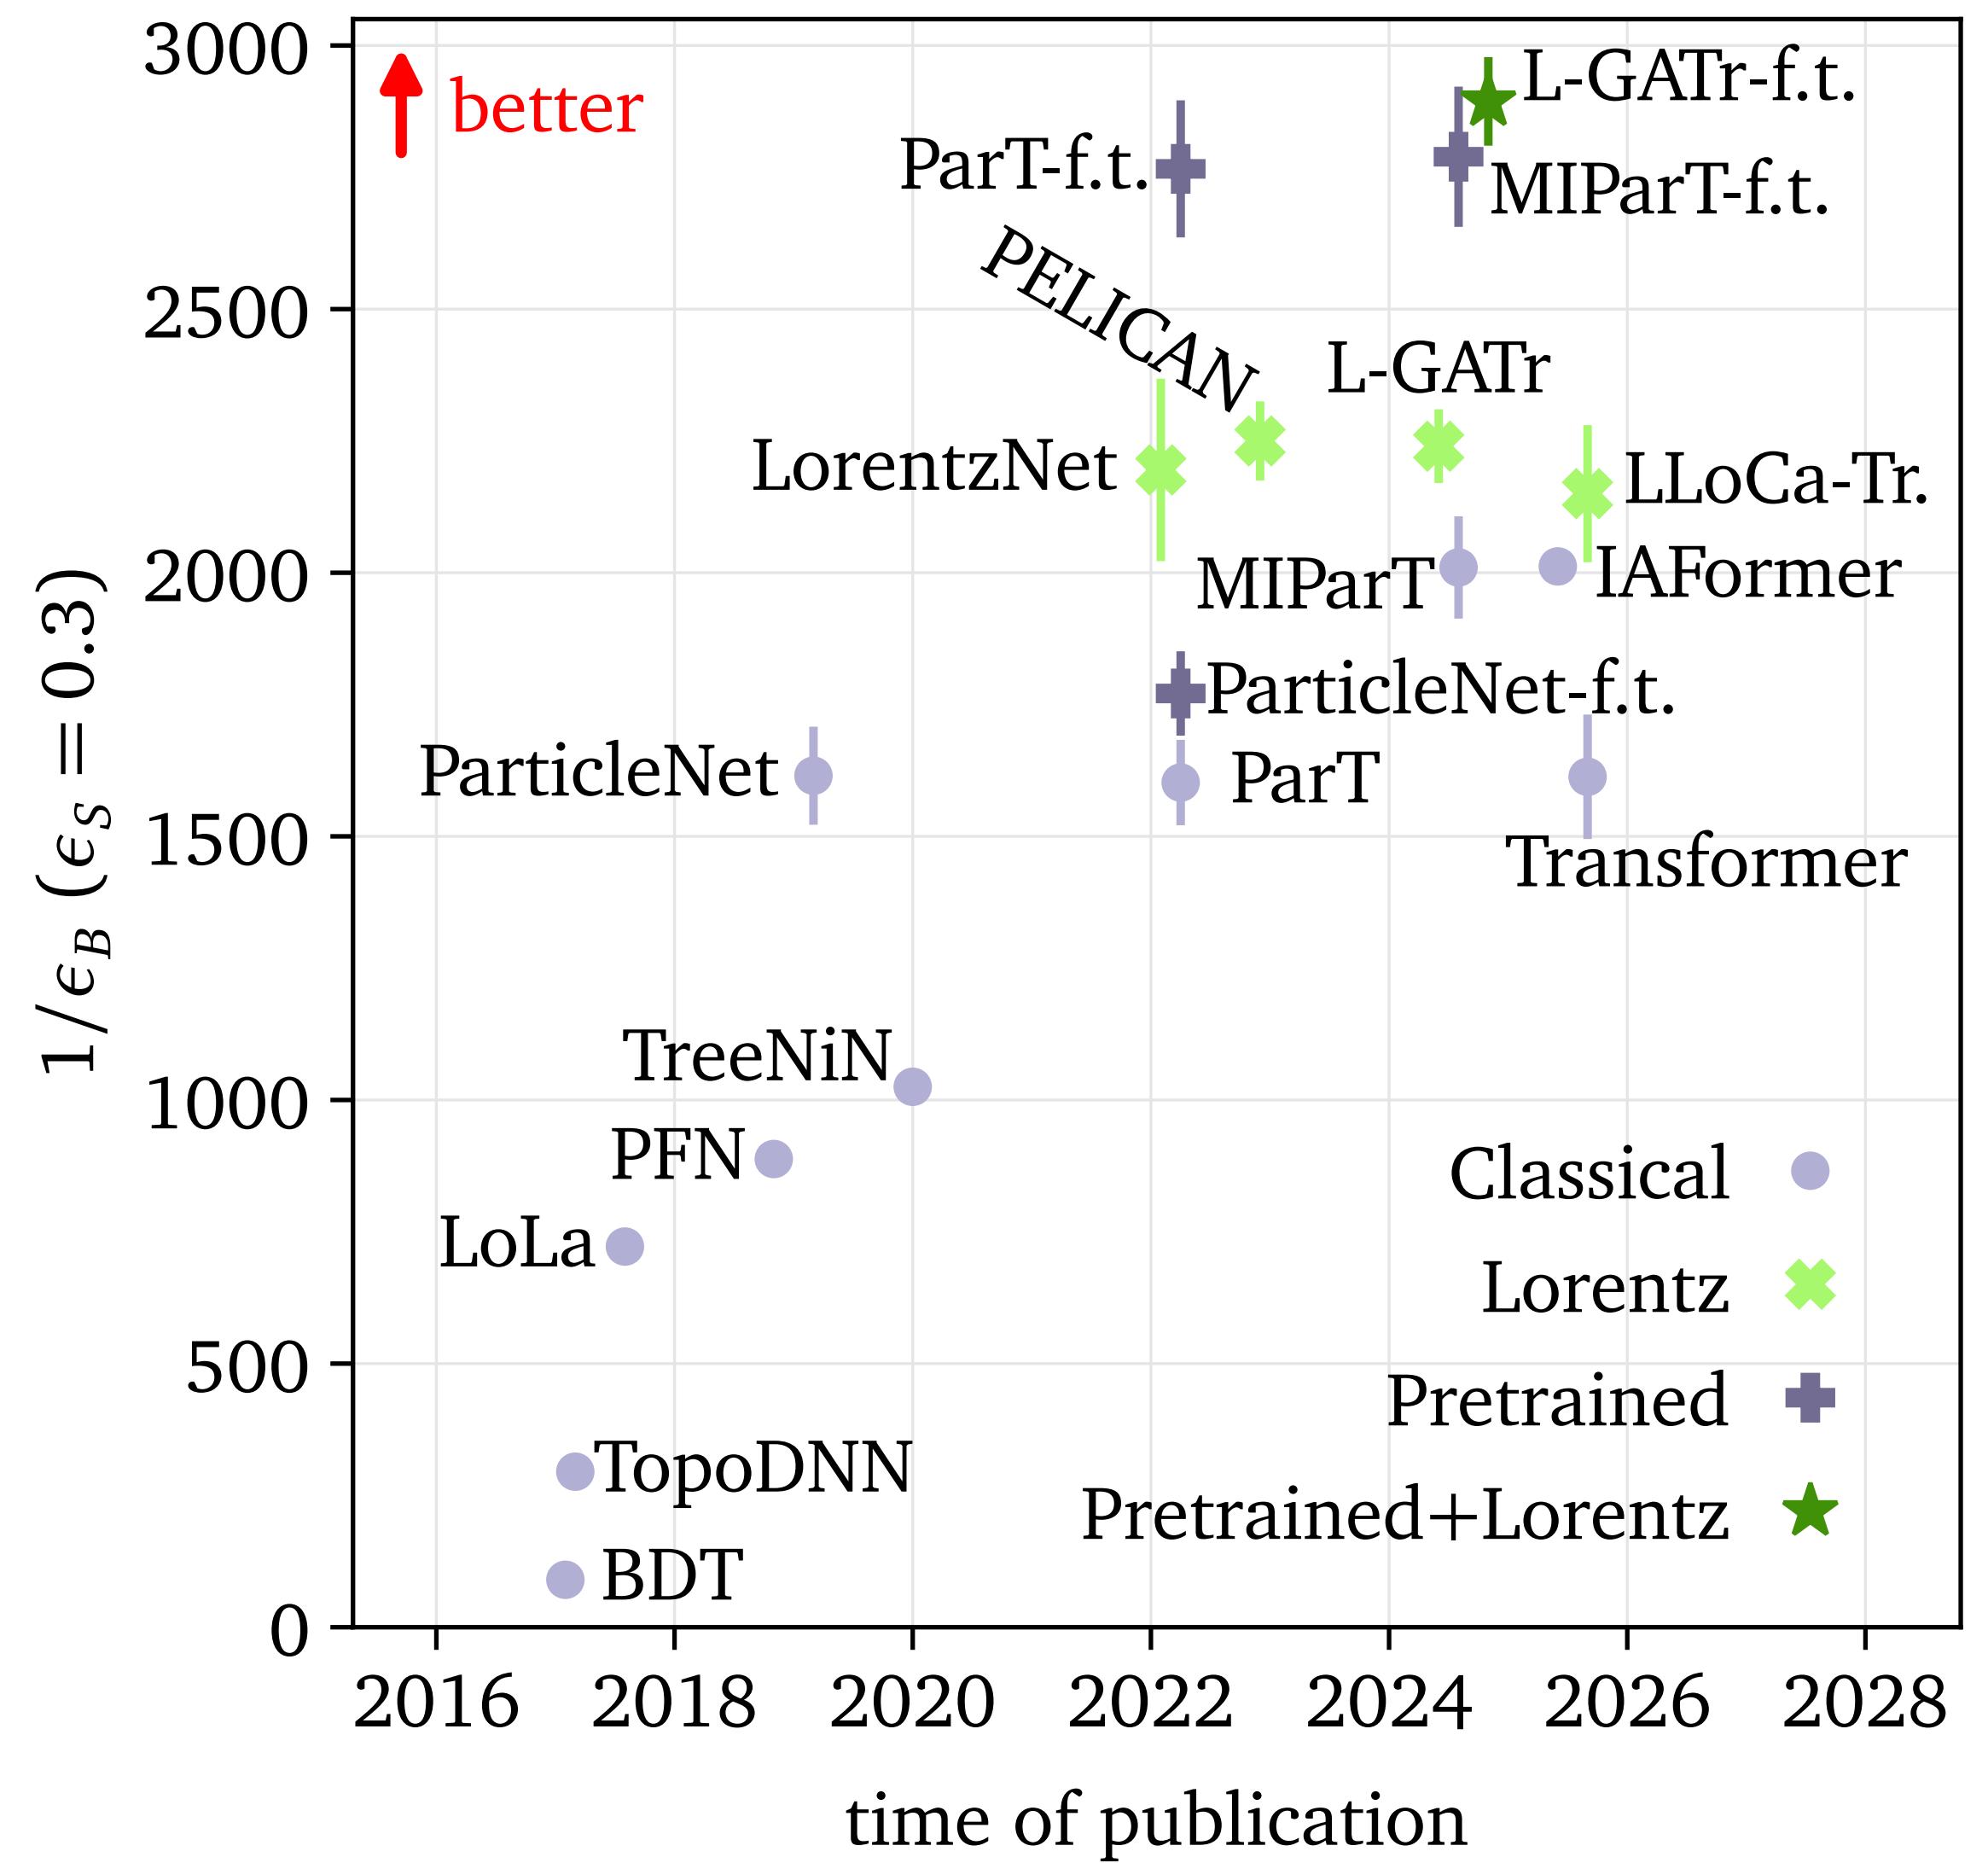
Lorentz-equivariance  
improves performance  
also at scale

# Jet Tagging

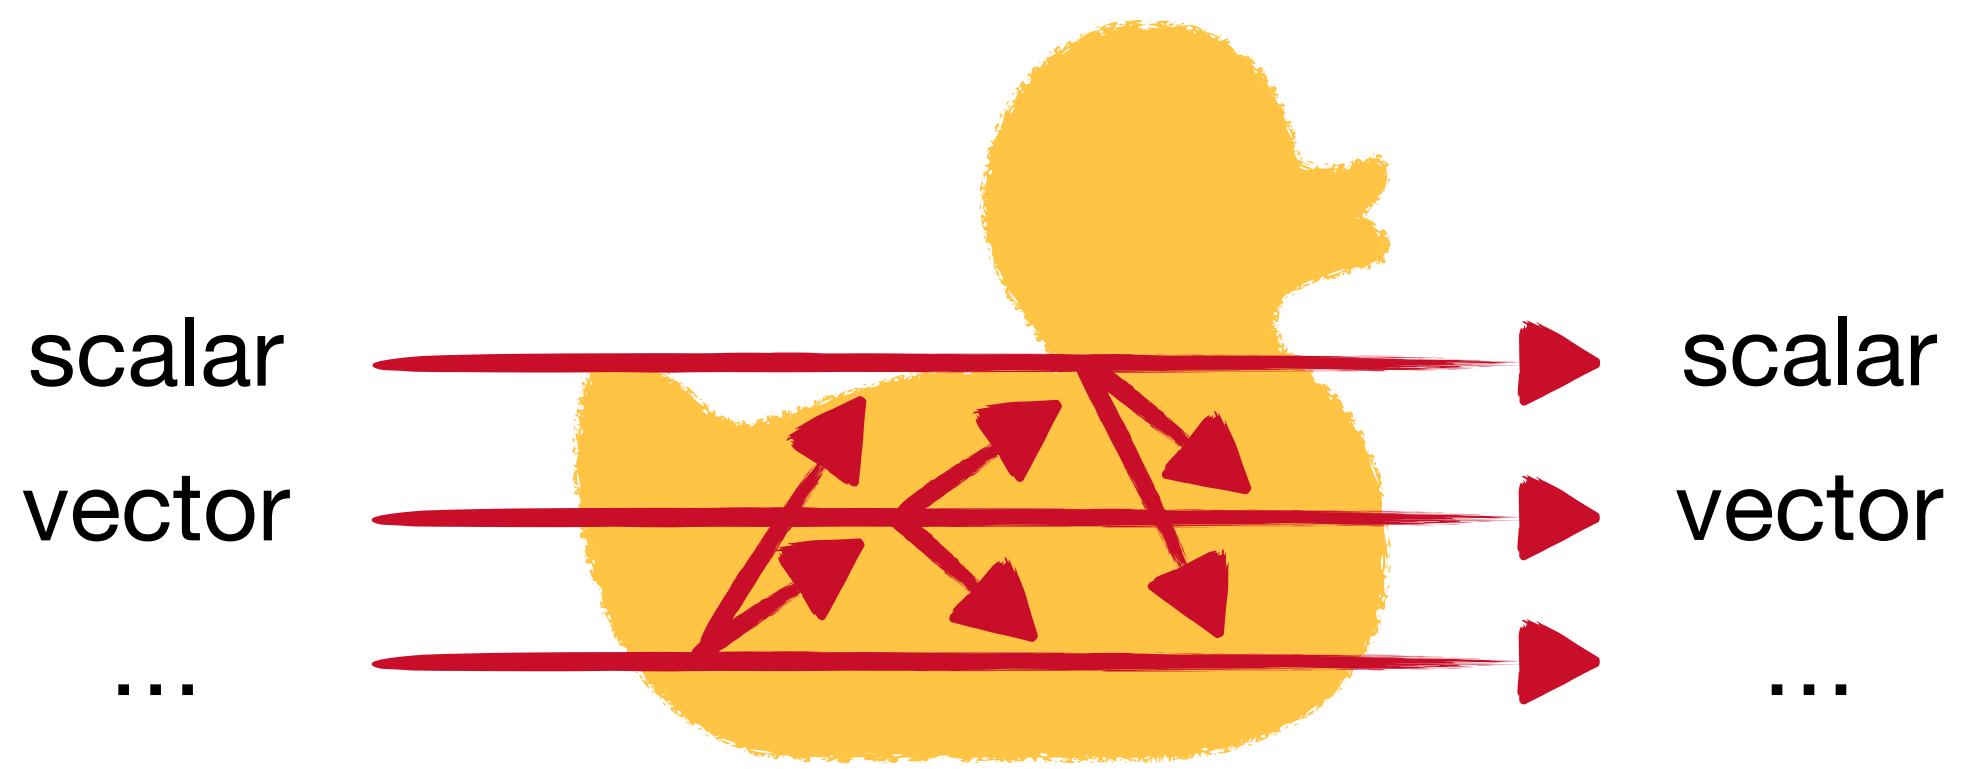
## Transfer learning

L-GATr-f.t. combines the benefits of pretraining and Lorentz equivariance

You can go for big + smart!



# Beyond specialised layers



**Efficient Lorentz-equivariance?**

How to make my **domain-specific** architecture Lorentz-equivariant?

## PELICAN

arXiv:2211.00454  
arXiv:2307.16506

## LorentzNet

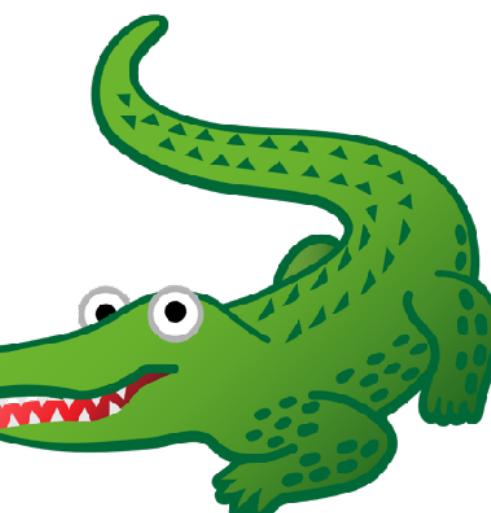
arXiv:2201.08187

## L-GATr

arXiv:2405.14806

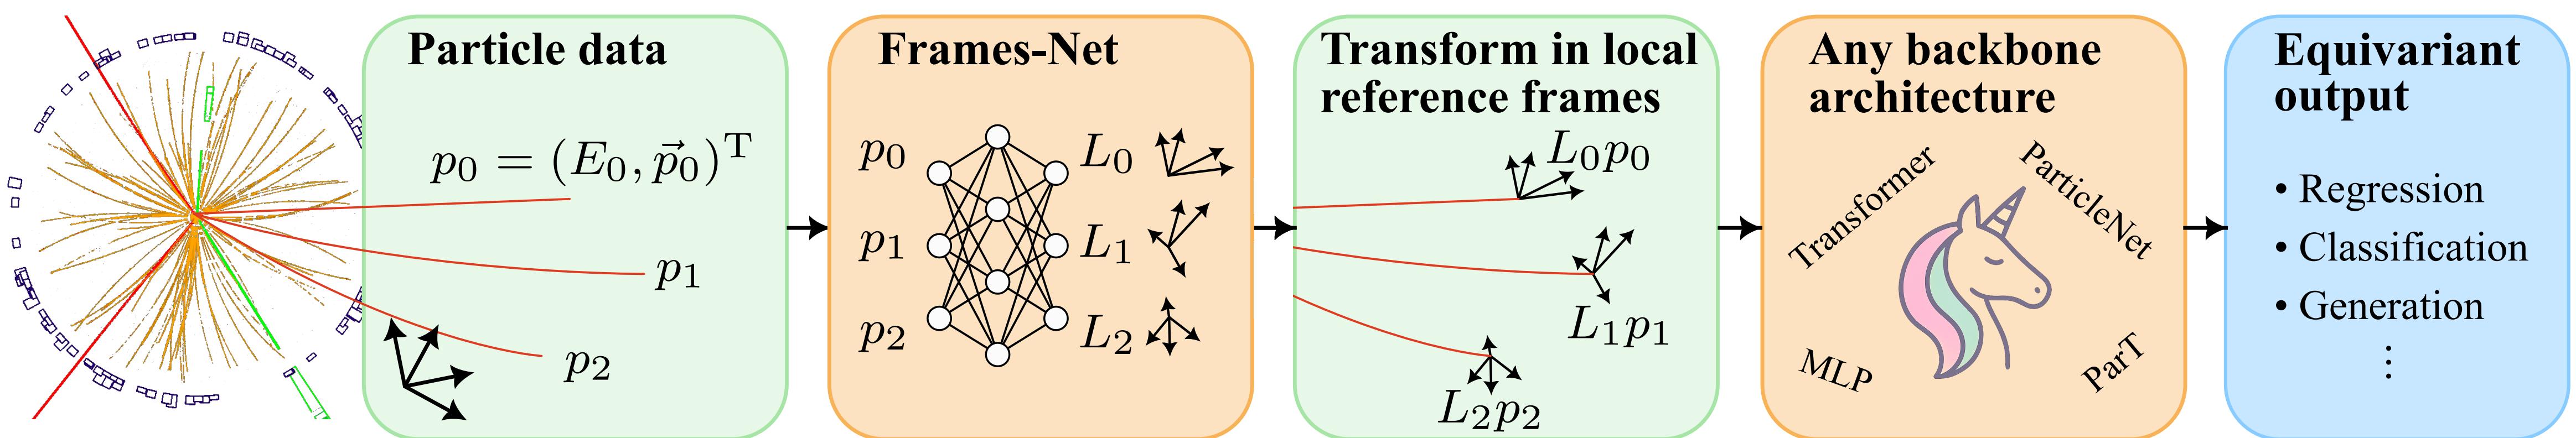
arXiv:2411.00446

pip install lgatr



What about **higher-order** internal representations?

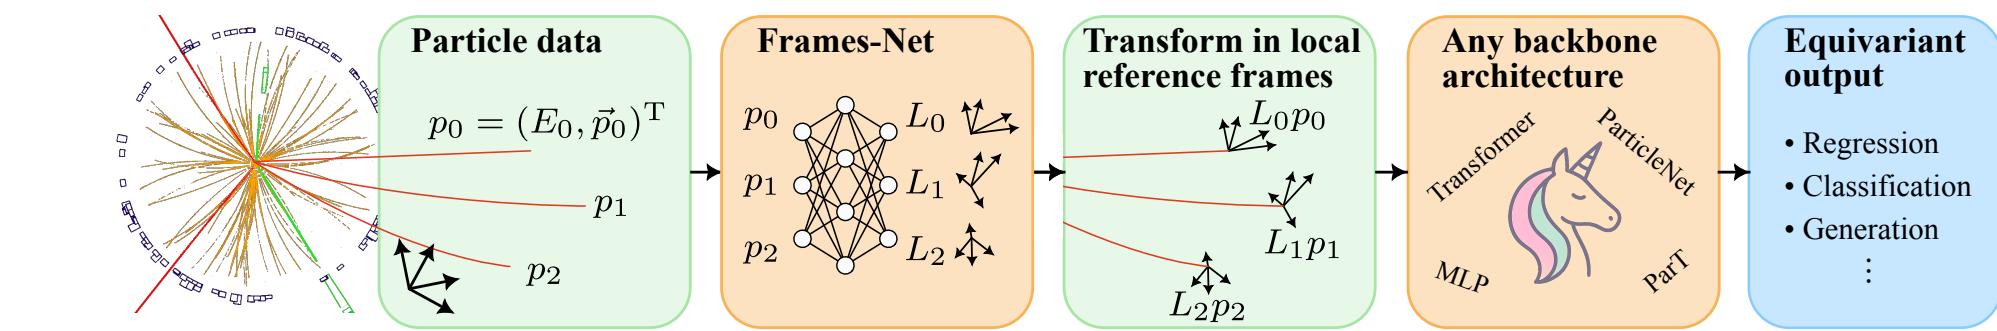
# How to make any network Lorentz-equivariant:



**Lorentz Local Canonicalization**

# LLoCa

## Local frames



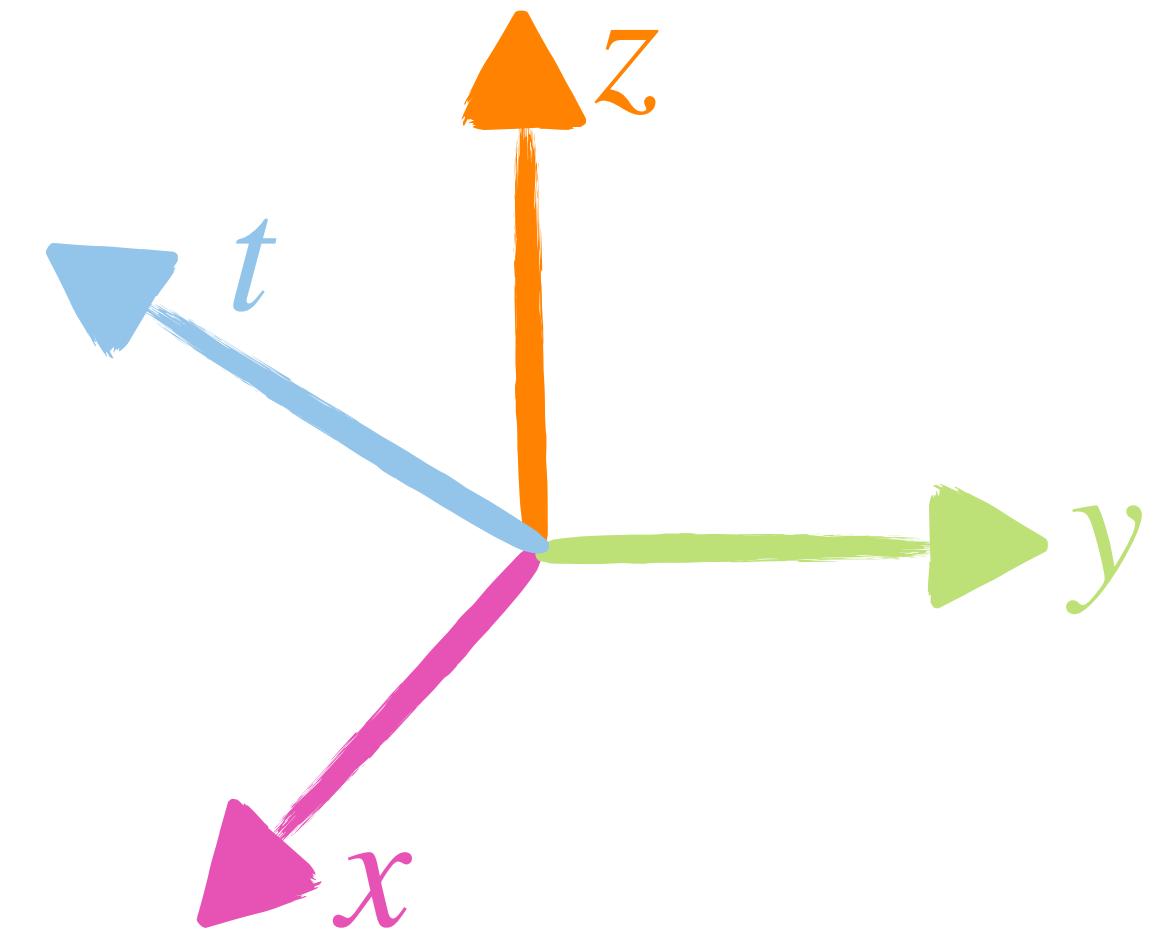
Features  $x_L = Lx$  in local frame are **invariant** by definition:

$$x_L = Lx \xrightarrow{\Lambda} L'x' = L'\Lambda x \stackrel{!}{=} x_L \quad \Rightarrow \quad L' = L\Lambda^{-1}$$

Achieve this by stacking **Lorentz vectors**  $u_a^\mu$ ,  $a = 0,1,2,3$ :

$$L = \begin{pmatrix} u_0^T g \\ u_1^T g \\ u_2^T g \\ u_3^T g \end{pmatrix} \xrightarrow{\Lambda} L' = \begin{pmatrix} u_0^T \Lambda^T g \\ u_1^T \Lambda^T g \\ u_2^T \Lambda^T g \\ u_3^T \Lambda^T g \end{pmatrix} = L\Lambda^{-1}$$

$$\sum_{\mu, \nu} u_a^\mu u_b^\nu g_{\mu\nu} = g_{ab}$$



# LLoCa

## How to construct local frames

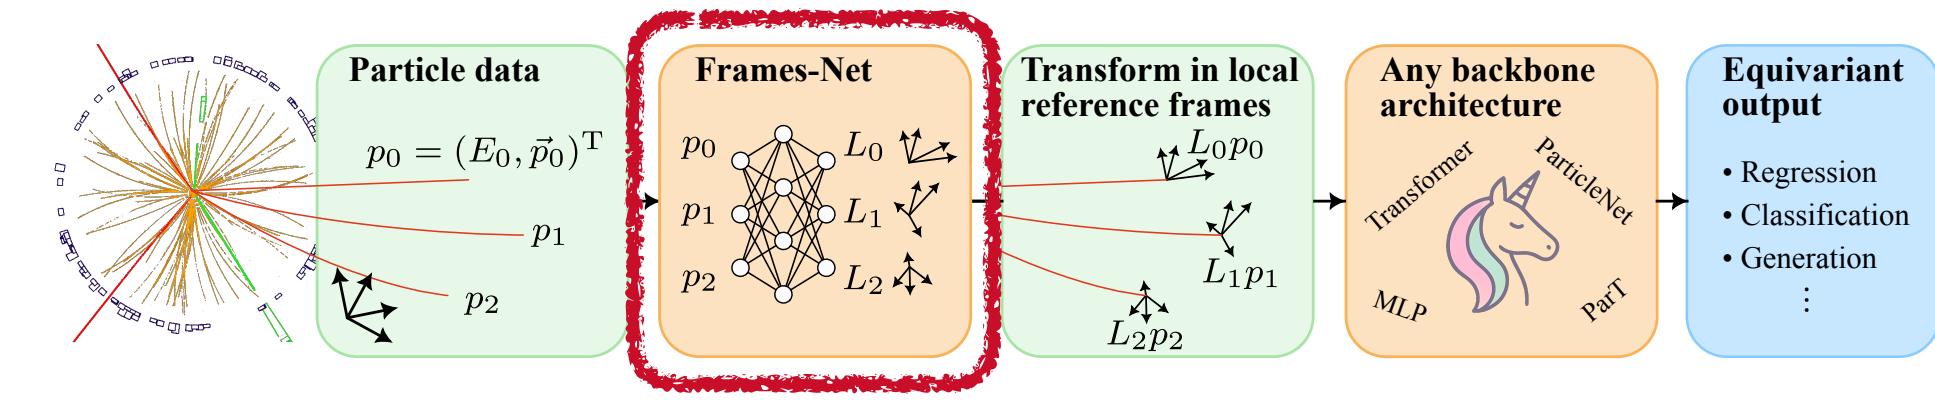
1) Predict 3 vectors  $v_{i,0}, v_{i,1}, v_{i,2}$  with simple **Lorentz-equivariant network**

$$v_{i,k} = \sum_{j=1}^N \text{softmax}\left(\varphi_k(s_i, s_j, \langle p_i, p_j \rangle)\right)(p_i + p_j) \quad \text{for } k = 0, 1, 2.$$



Frames-Net (MLP)

2) Orthonormalize  $v_{i,0}, v_{i,1}, v_{i,2}$  with **Gram-Schmidt algorithm** to obtain  $u_{i,a}$



# LLoCa

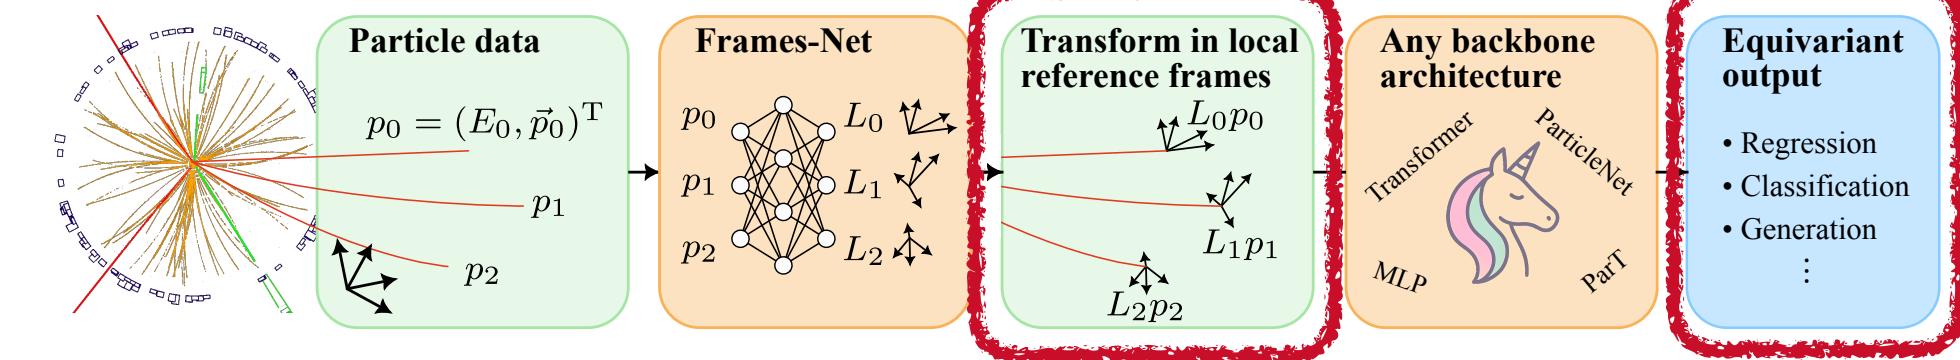
## Equivariance from canonicalisation

Latent features in the local frame  $f_L = \rho(L)f$  are **invariant**

$$f_L \xrightarrow{\Lambda} f'_L = \rho(L')f' = \rho(L\Lambda^{-1})\rho(\Lambda)f = \rho(L\Lambda^{-1}\Lambda)f = \rho(L)f = f_L$$

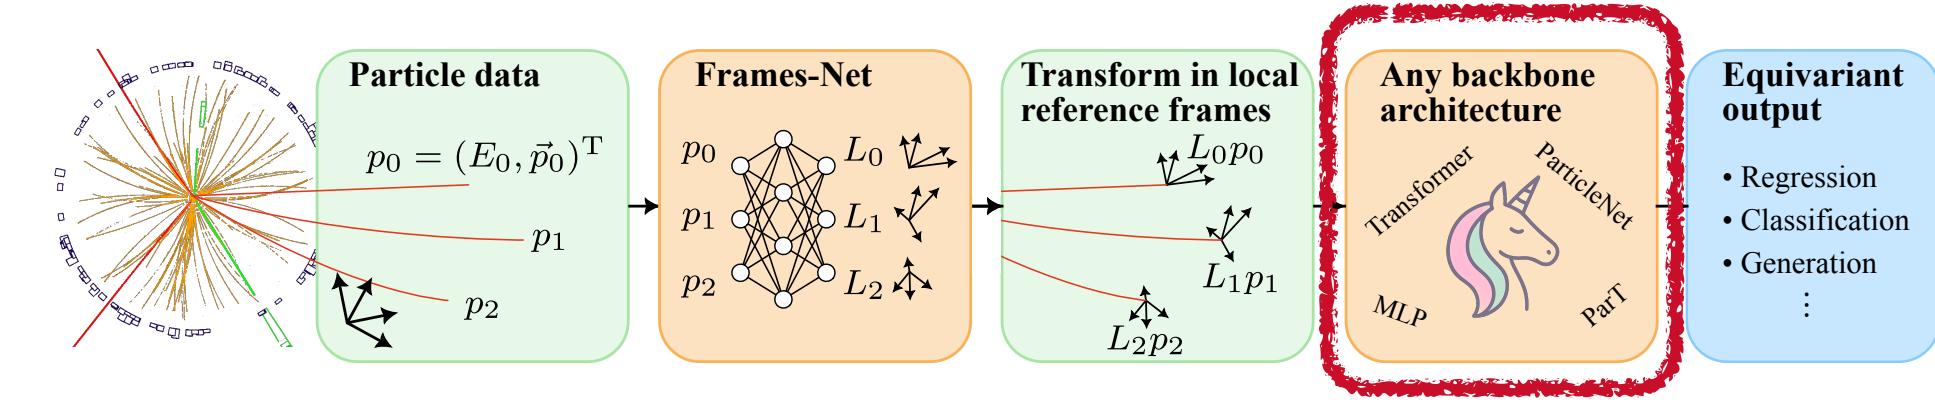
Output features in the global frame  $y = \rho(L^{-1})y_L$  are **equivariant**

$$y \xrightarrow{\Lambda} y' = \rho(L'^{-1})y'_L = \rho(\Lambda L^{-1})y_L = \rho(\Lambda)\rho(L^{-1})y_L = \rho(\Lambda)y$$

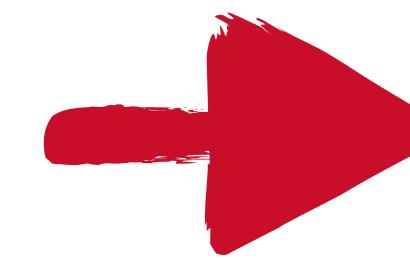


# LLoCa

## Message-passing between local frames



**Standard message-passing**

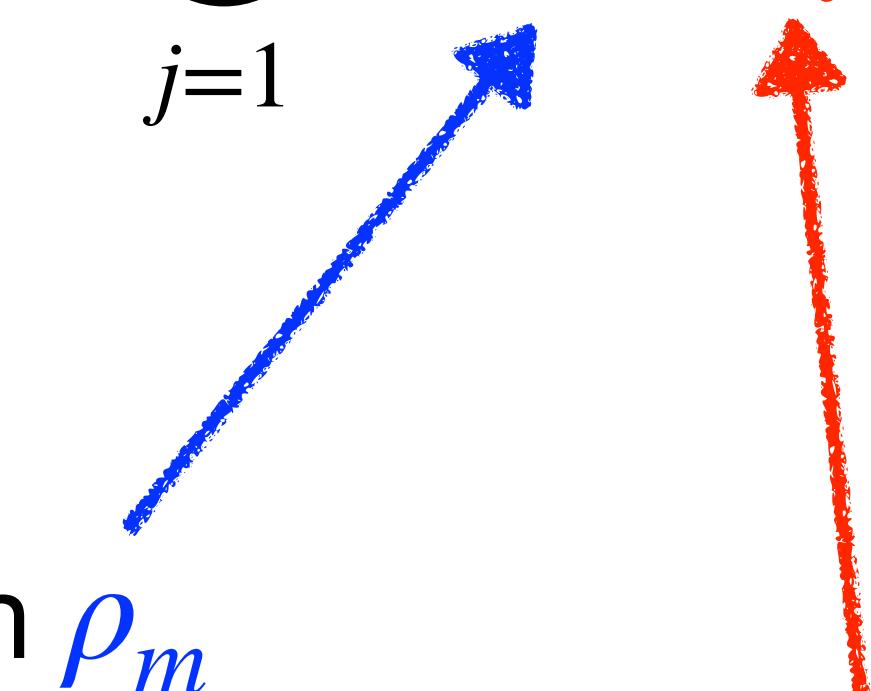


**Tensorial message-passing**  
between different local frames

$$f_i^{\text{updated}} = \psi\left(f_i, \bigoplus_{j=1}^N \phi(m_j)\right)$$

$$f_{i,L_i}^{\text{updated}} = \psi\left(f_{i,L_i}, \bigoplus_{j=1}^N \phi\left(\rho_m(L_i L_j^{-1}) m_{j,L_j}\right)\right)$$

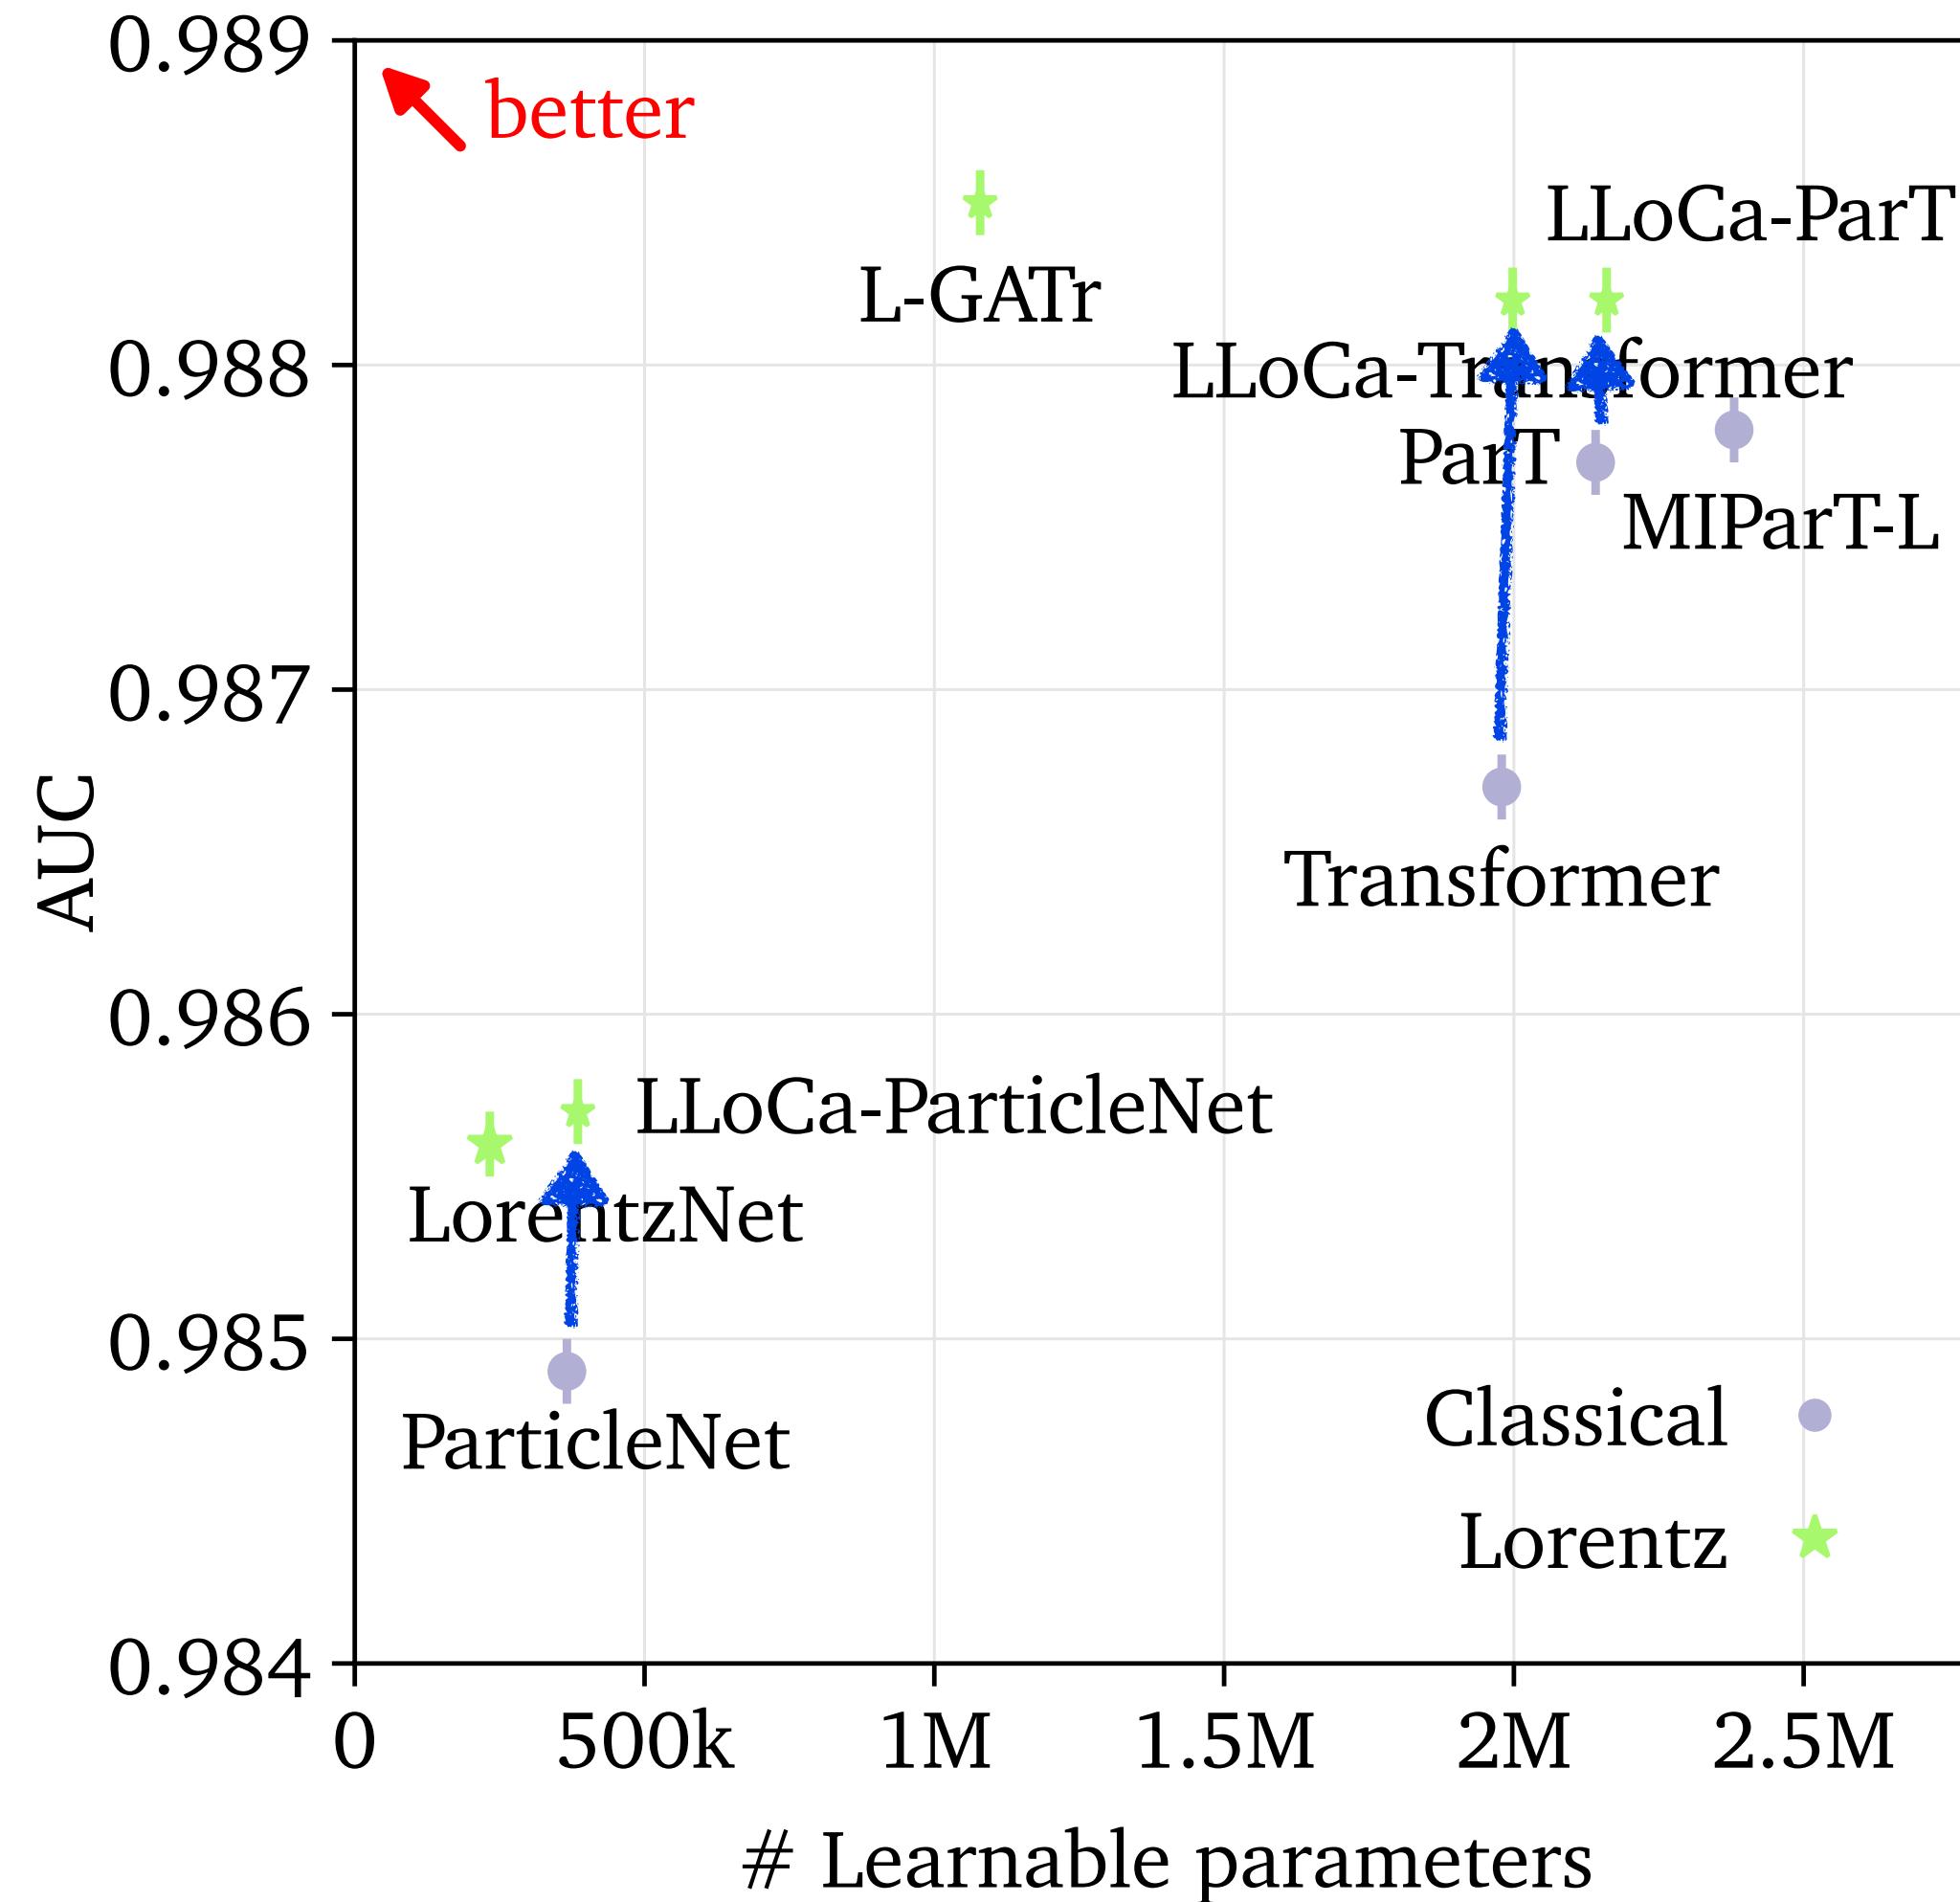
Message representation  $\rho_m$   
is a hyperparameter



Invariant frame  
transformation matrix  $L_i L_j^{-1}$

# Jet Tagging

Make any network Lorentz-equivariant

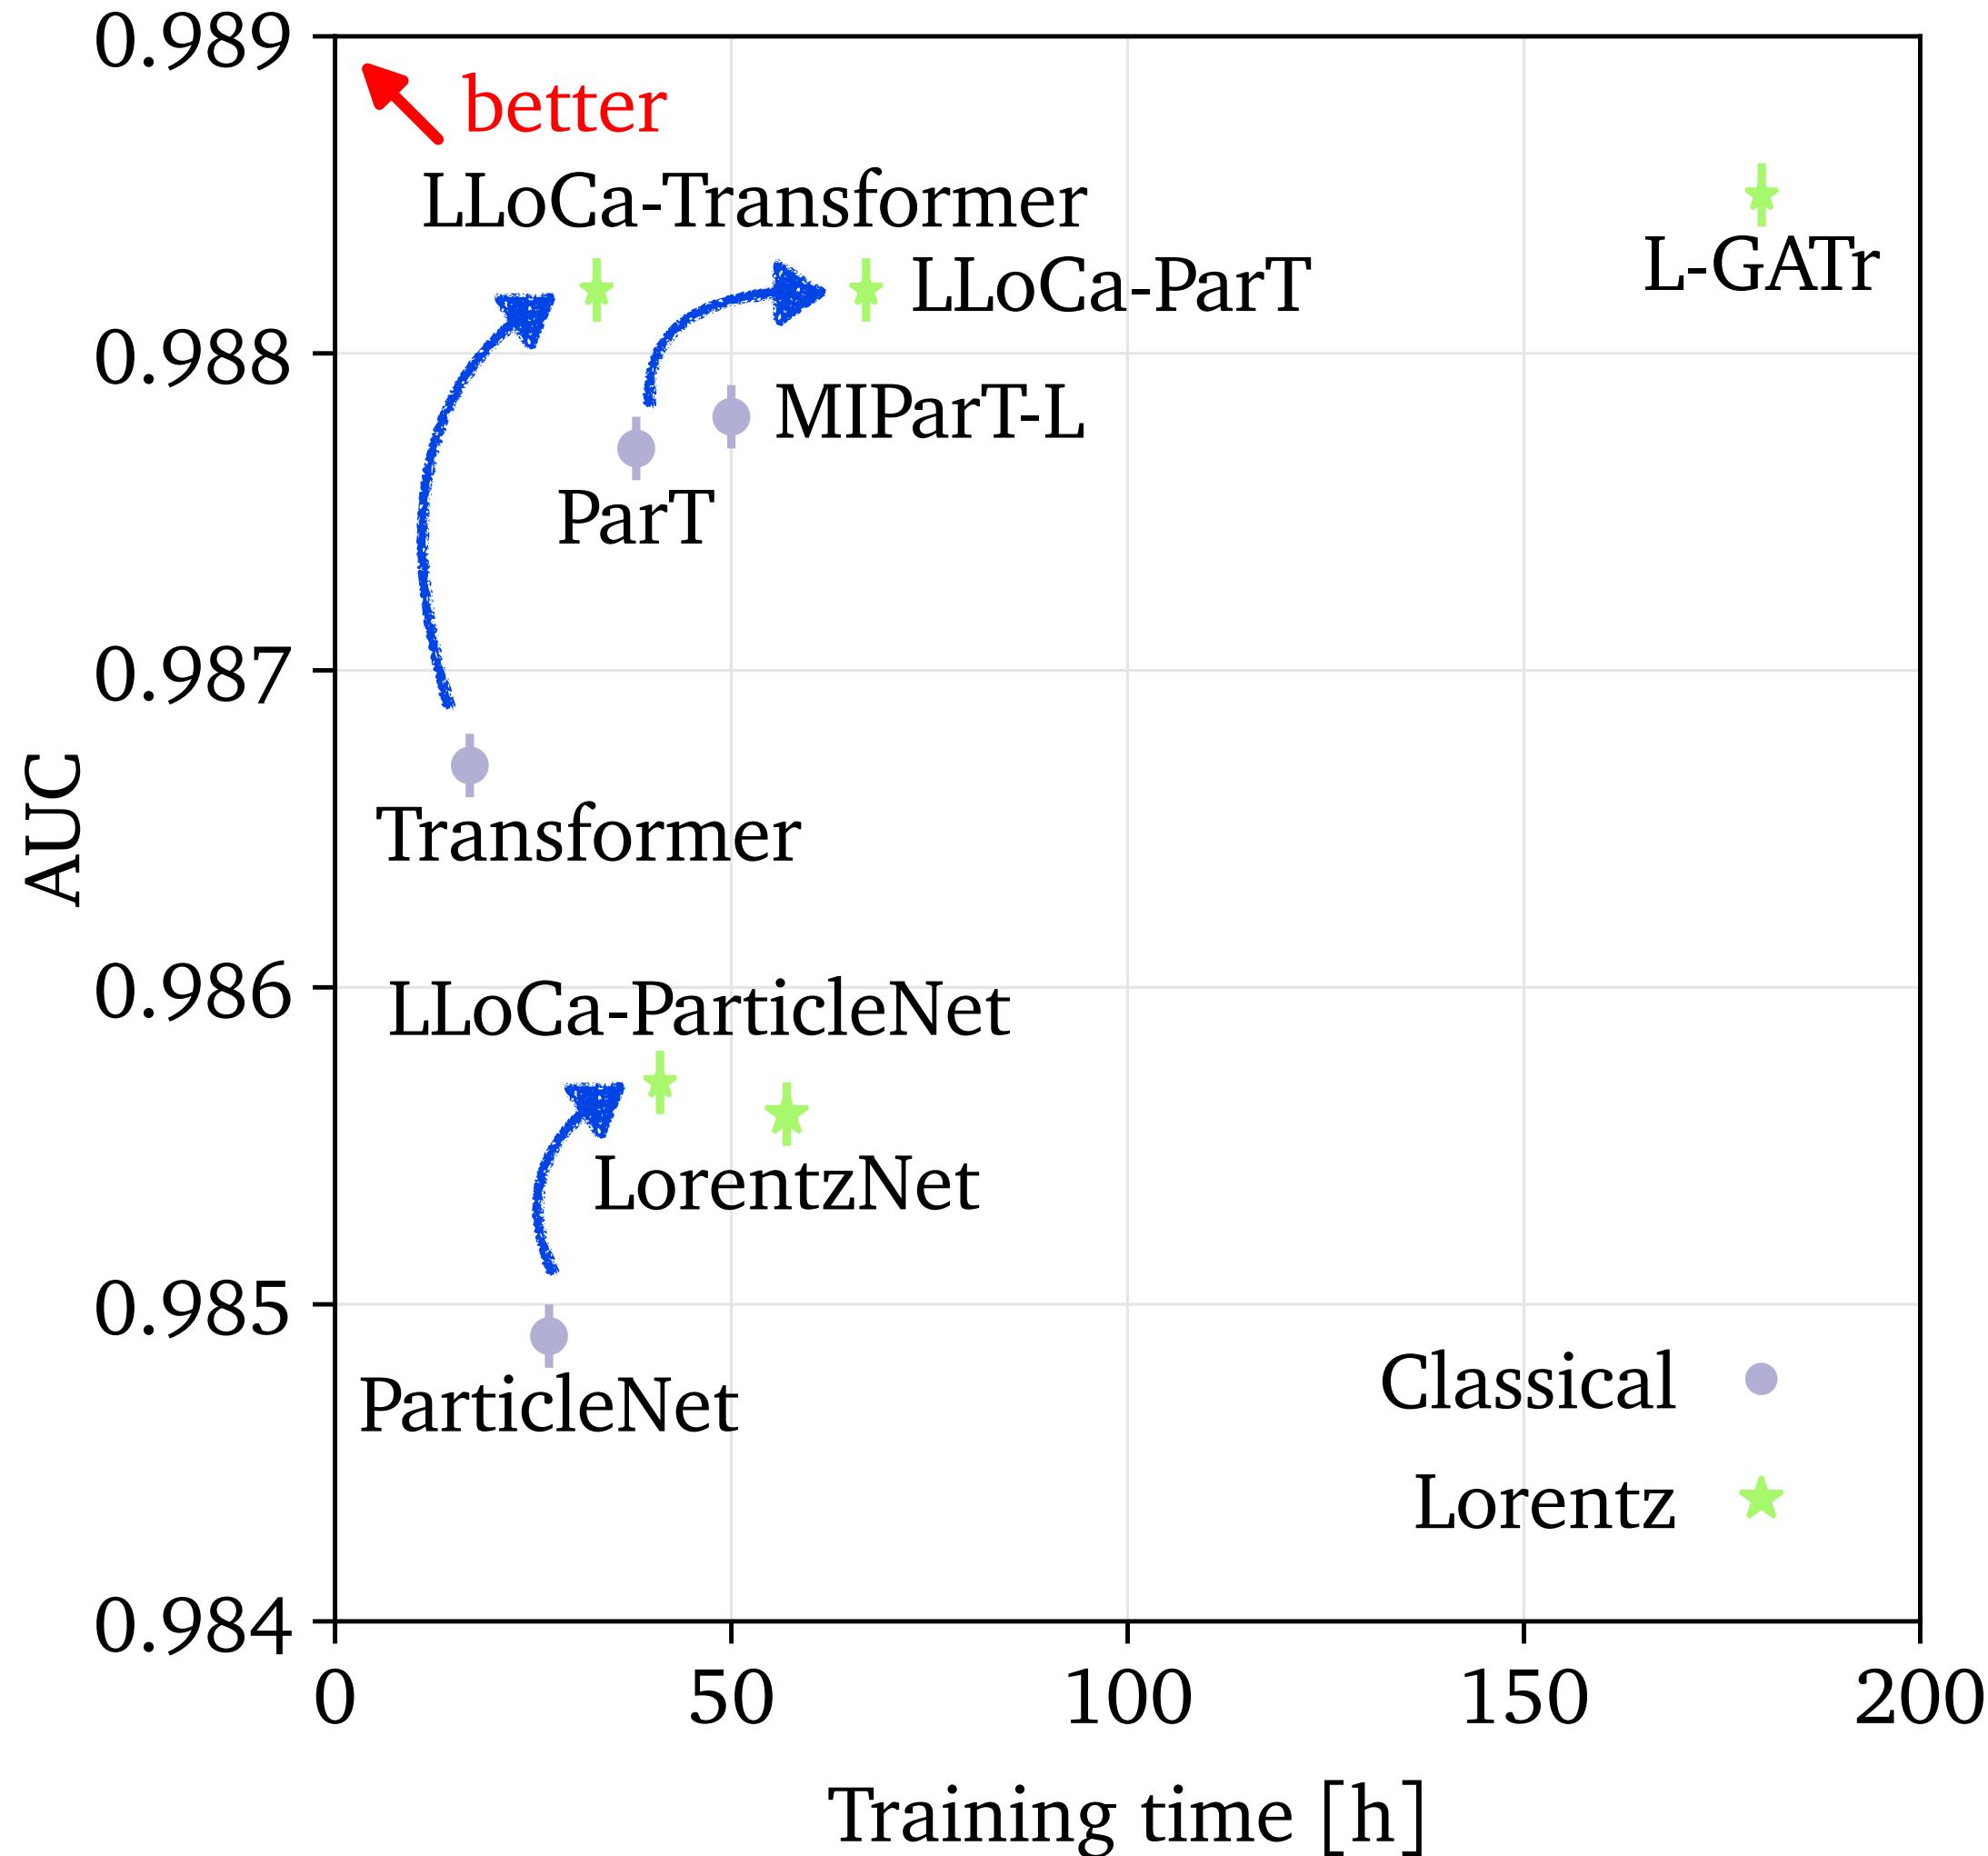


LLoCa upgrades established taggers to be Lorentz-equivariant

LLoCa-ParT is not better than LLoCa-Transformer

# Jet Tagging

## Efficient Lorentz-equivariance

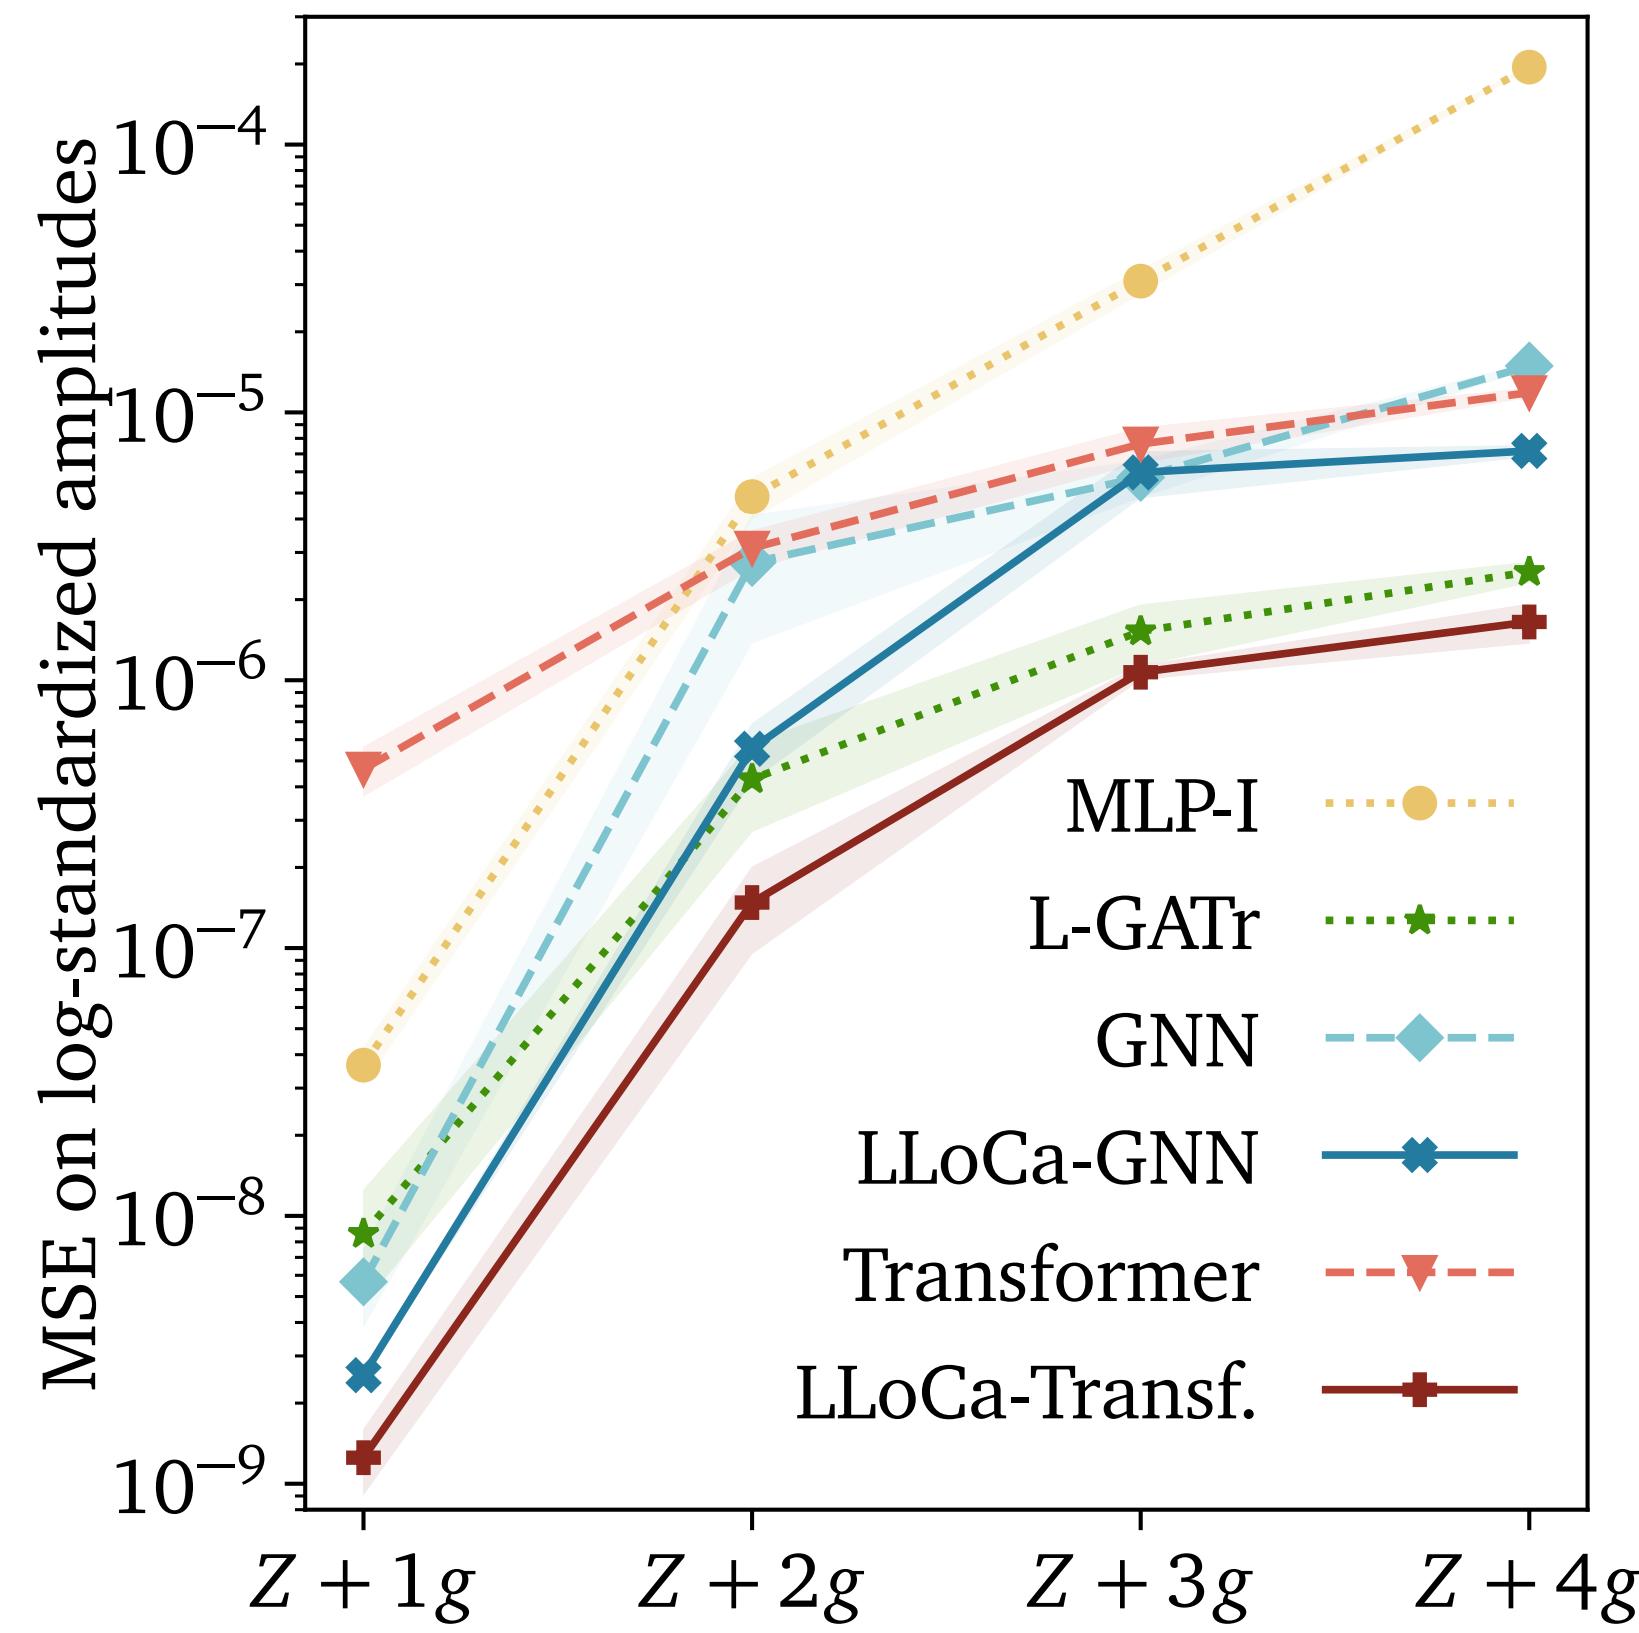


LLoCa costs only ~2x extra training time

LLoCa-Transformer is better and faster than Part

# Amplitude regression

## Efficient Lorentz-equivariance



L-GATr and LLoCa-Transformer  
achieve similar performance

Model	$MSE \times 10^{-6} (\downarrow)$	FLOPs	Time
MLP-I [38]	$195.0 \pm 2$	0.1M	0.6h
LGATr [38]	$2.5 \pm 0.2$	1160.0M	19.5h
GNN	$14.9 \pm 0.3$	9.9M	1.3h
LLoCa-GNN	$7.2 \pm 0.3$	10.6M	2.3h
Transformer	$11.9 \pm 0.5$	10.6M	3.5h
LLoCa-Transformer	$\underline{1.5 \pm 0.1}$	11.3M	4.5h

... but LLoCa-Transformer is 4x faster  
and requires 100x less FLOPs!

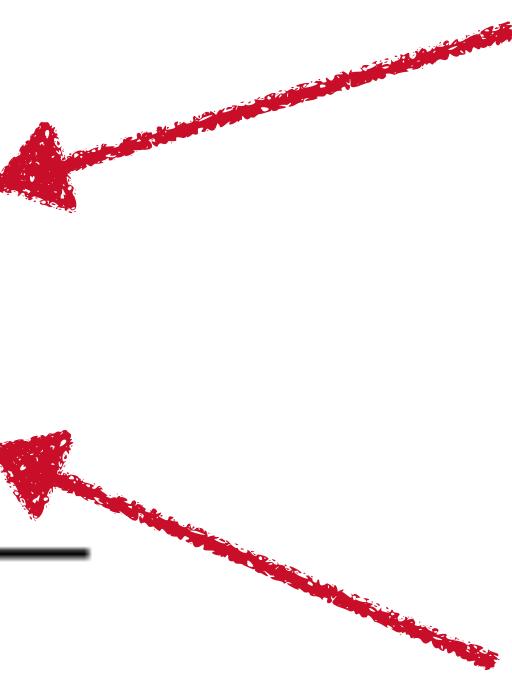
# Amplitude regression

Arbitrary internal representations  $\rho_m$

Method	MSE ( $\times 10^{-6}$ )
Non-equivariant	$11.9 \pm 0.8$
Global canonical.	$5.9 \pm 0.4$
LLoCa (16 scalars)	$54.0 \pm 3$
LLoCa (single 2-tensor)	$2.4 \pm 0.3$
LLoCa (4 vectors)	$2.0 \pm 0.2$
LLoCa (8 scalars, 2 vectors)	<b><math>1.5 \pm 0.1</math></b>

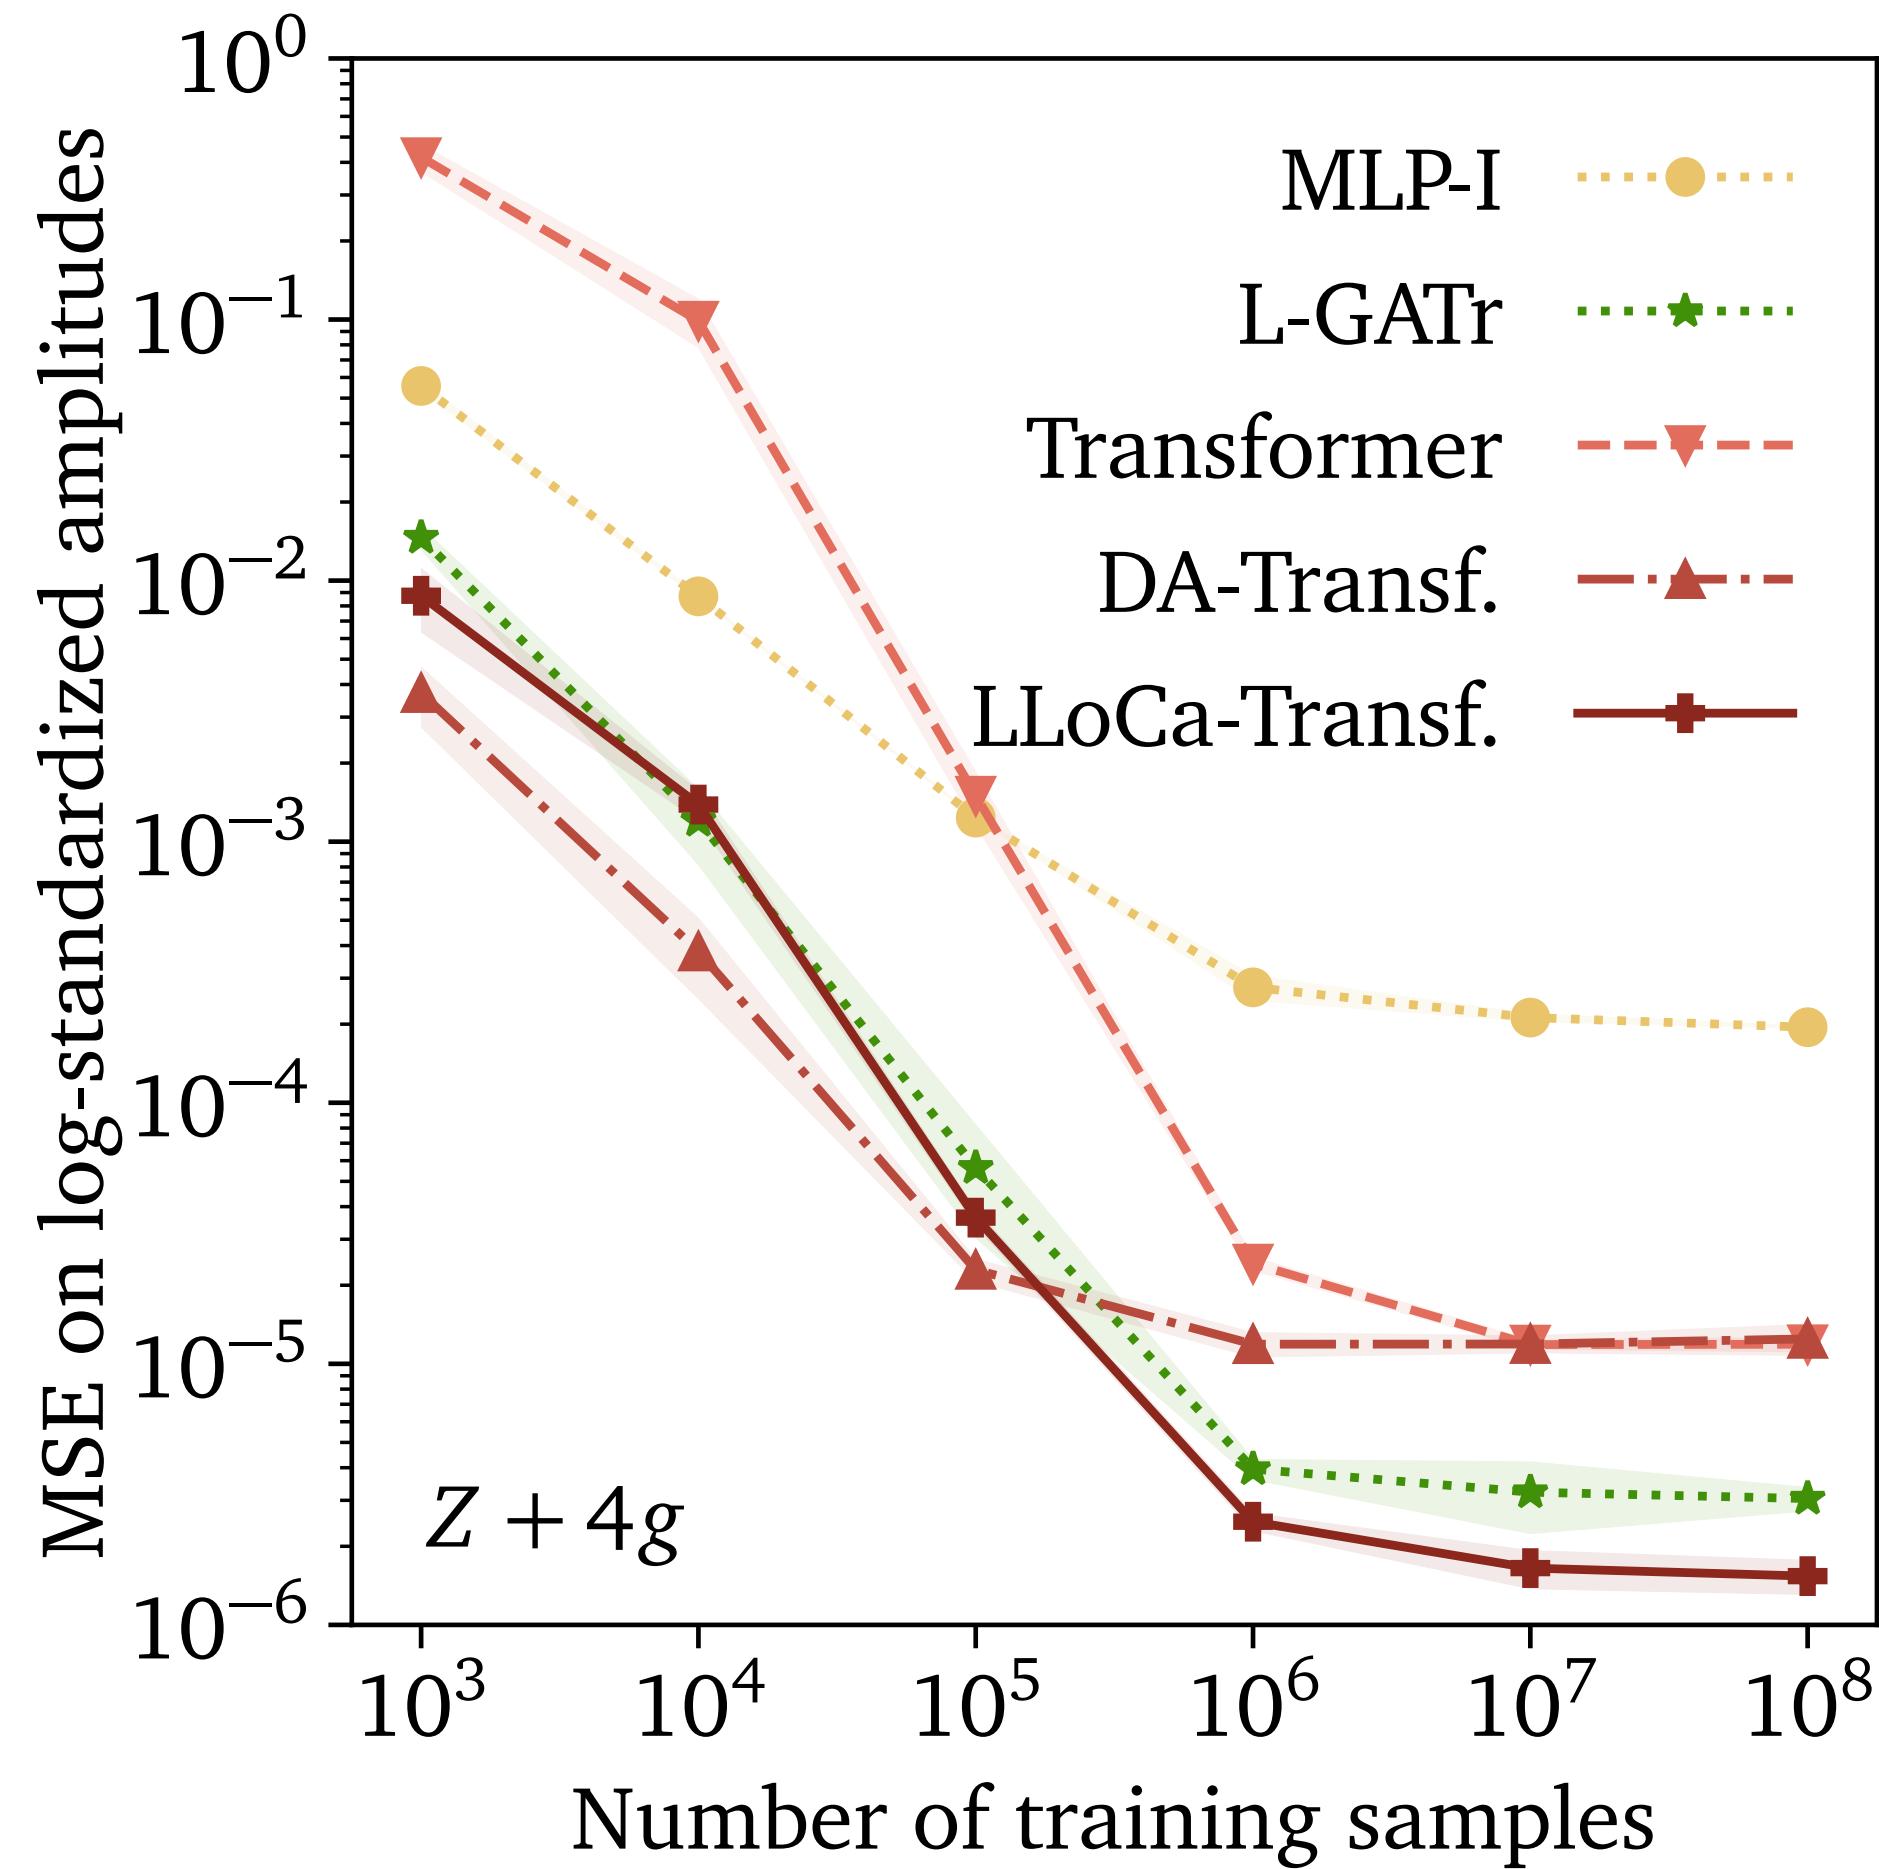
LLoCa handles arbitrary tensorial representations

Across our experiments,  
a mix of scalar and vector  
representations works best.



# Amplitude regression

## Equivariance vs data augmentation



Data augmentation (DA)  
emerges from LLoCa as  
random global frames

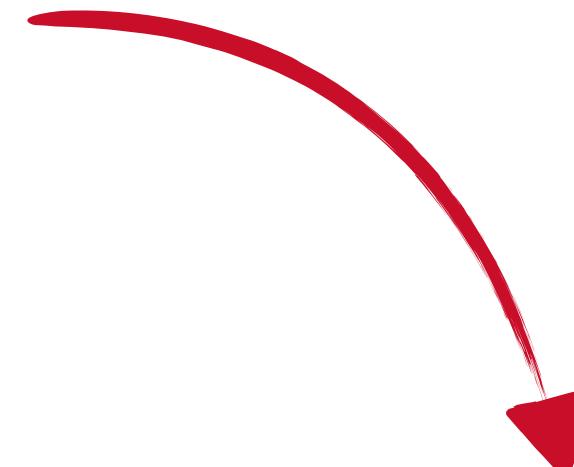
Data augmentation beats  
Lorentz-equivariance in  
the limit of little training data

# Symmetry breaking

## Architecture vs input level



The jet tagging score is  
**not Lorentz-equivariant**



Symmetry breaking at the  
**architecture level**

- Hard to do with specialised layers
- Possible in LLoCa by fixing  $v_{i,0}$

Symmetry breaking at the  
**input level with reference vectors**

- Beam reference vector:  
 $x^V = (0,0,0,1)$
- Time reference vector:  
 $x^V = (1,0,0,0)$

All Lorentz-equivariant taggers  
use symmetry breaking

# Jet Tagging

## Lorentz symmetry breaking

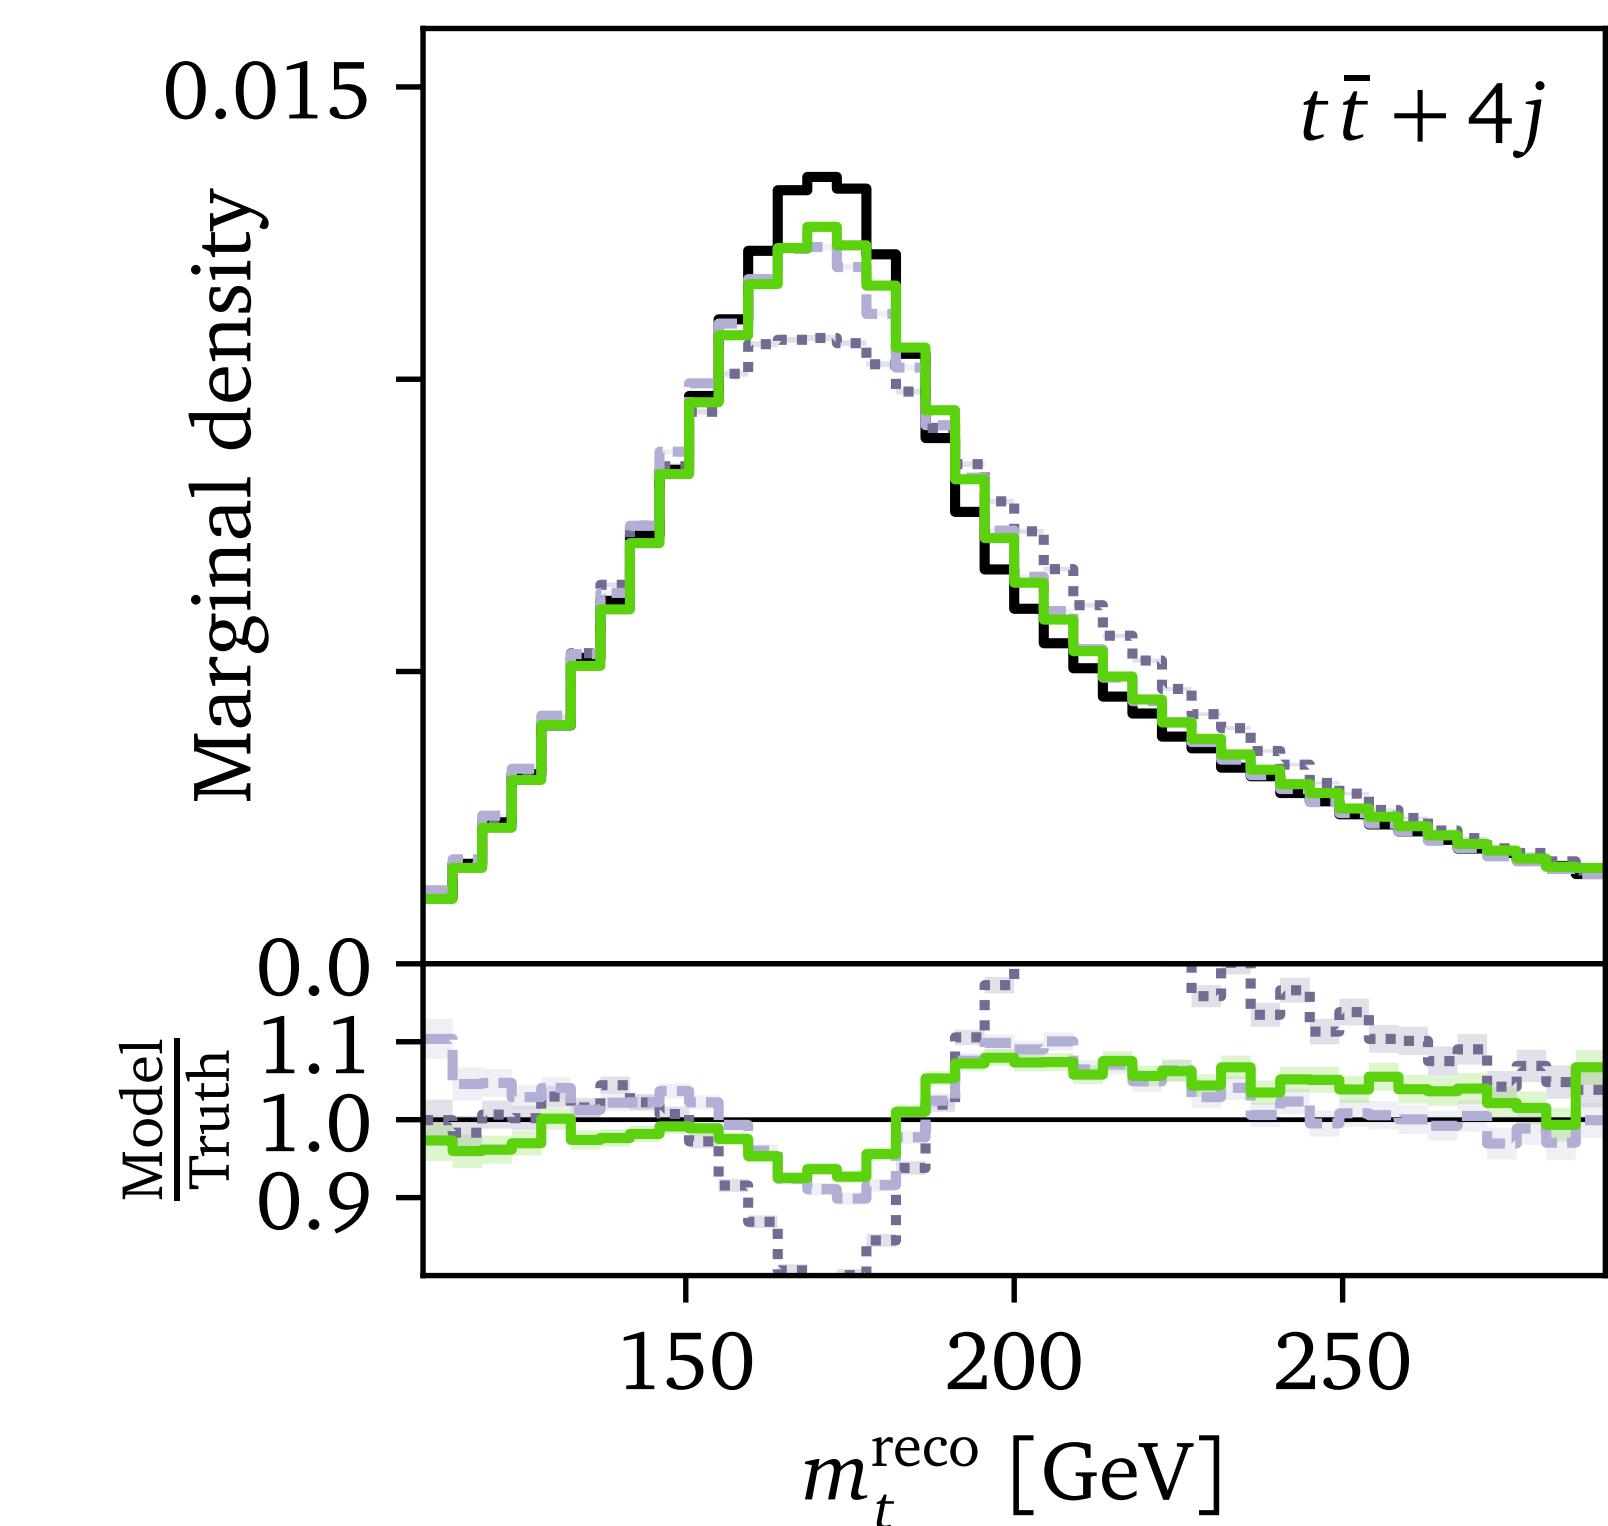
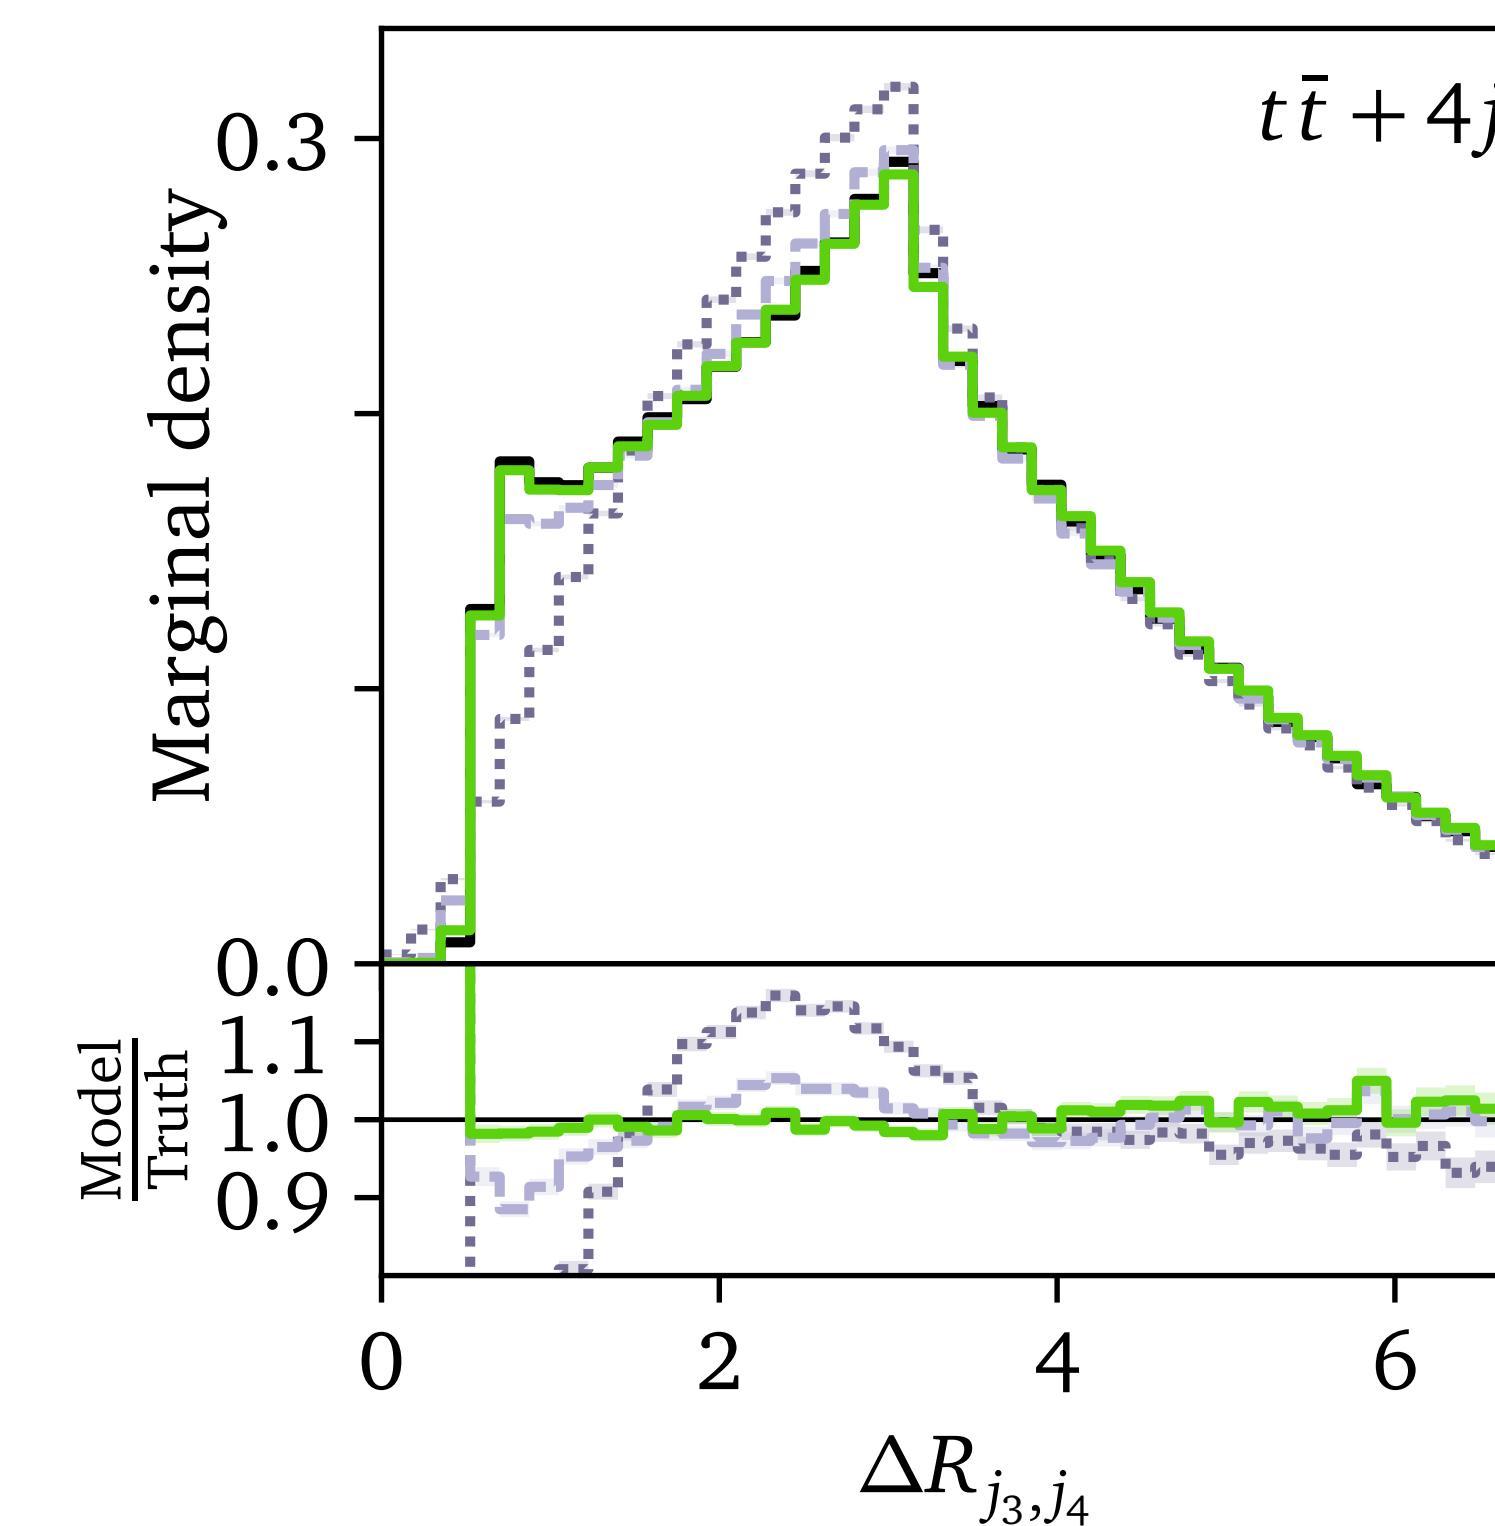
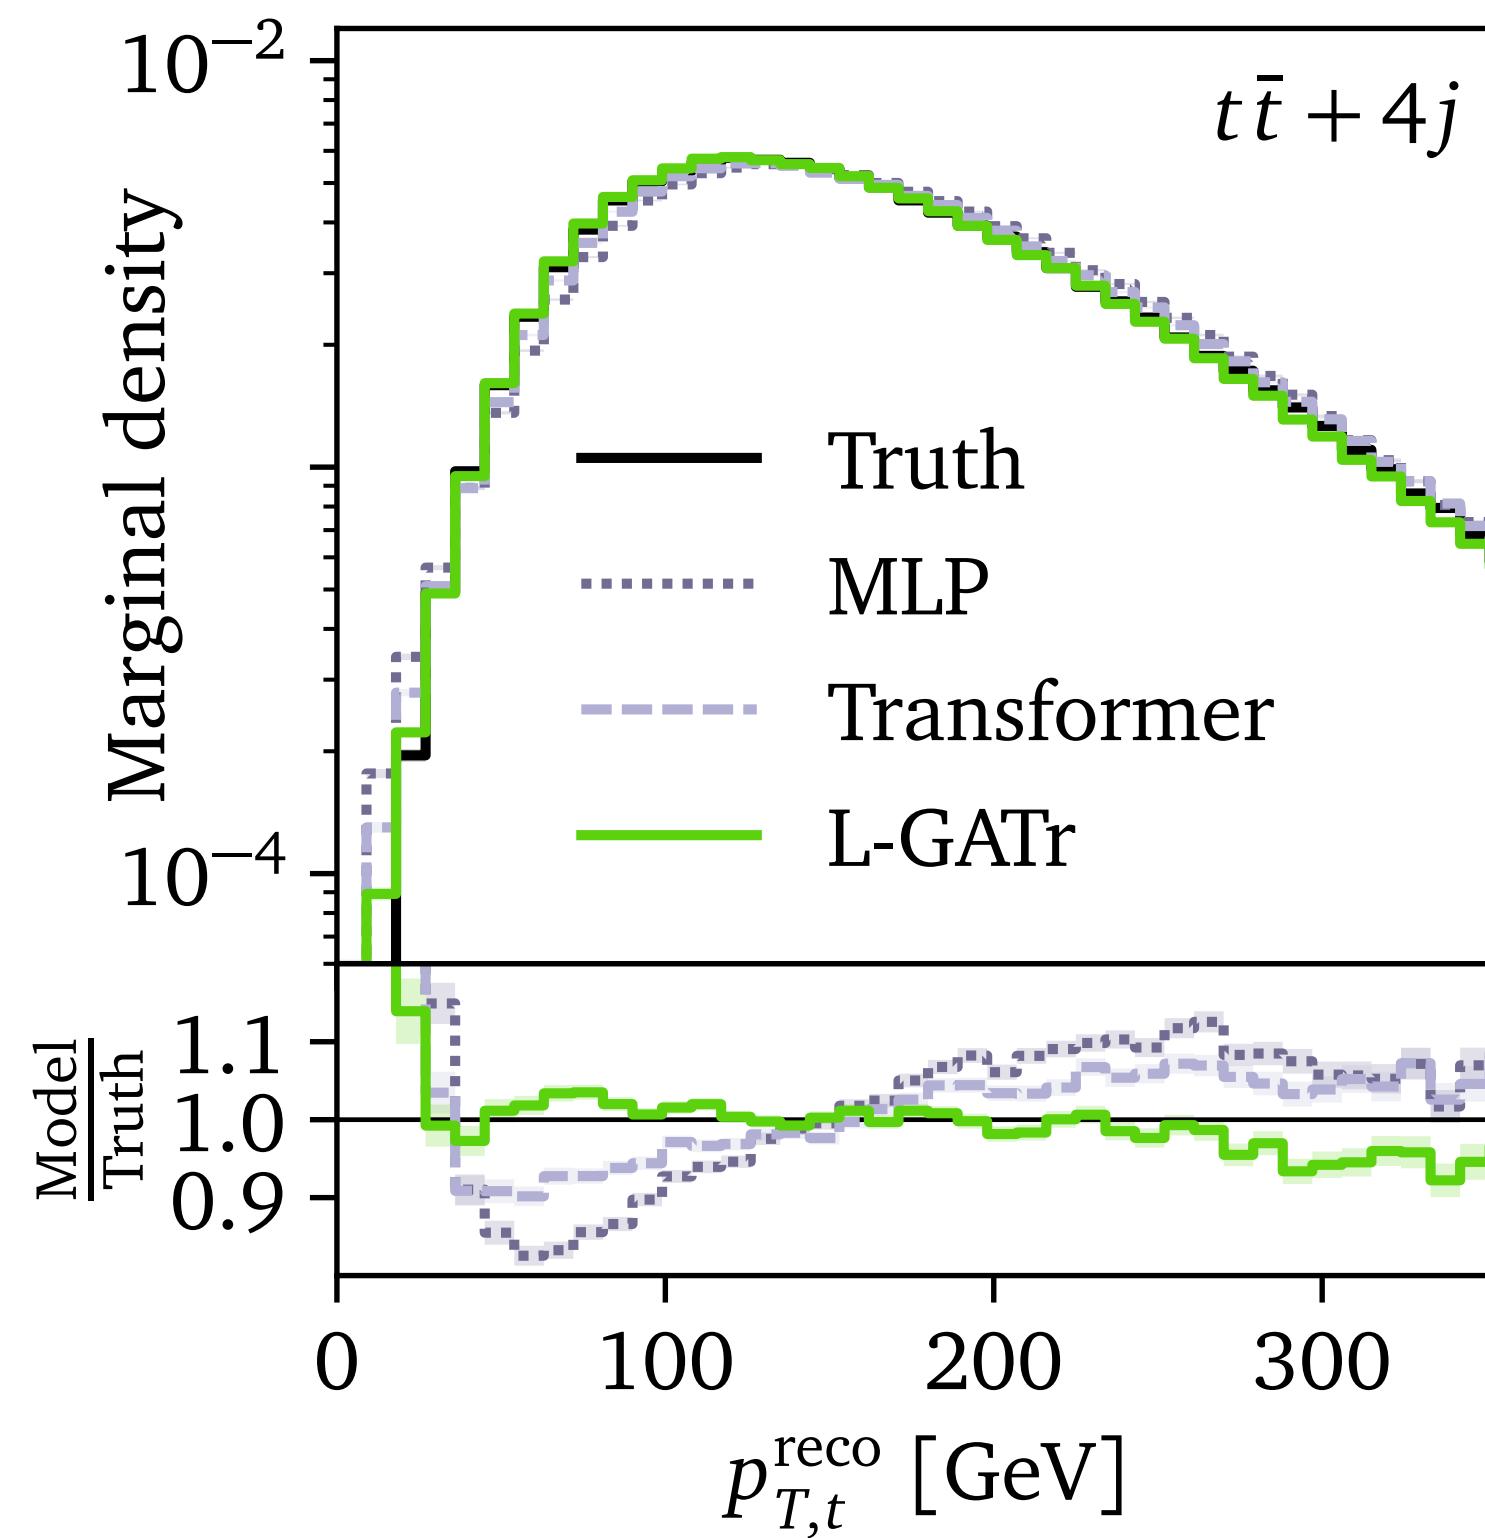
Symmetry breaking Network	Architecture Accuracy AUC	Input Accuracy AUC
Non-equivariant	0.855	0.9867 0.864 0.9882
SO(2)-equivariant	0.862	0.9878 0.864 0.9882
Lorentz-equivariant	0.856	0.9870 0.856 0.9870

= LLoCa-Transformer

Lorentz-equivariance + symmetry breaking  
outperforms SO(2)-equivariance!

# Event generation

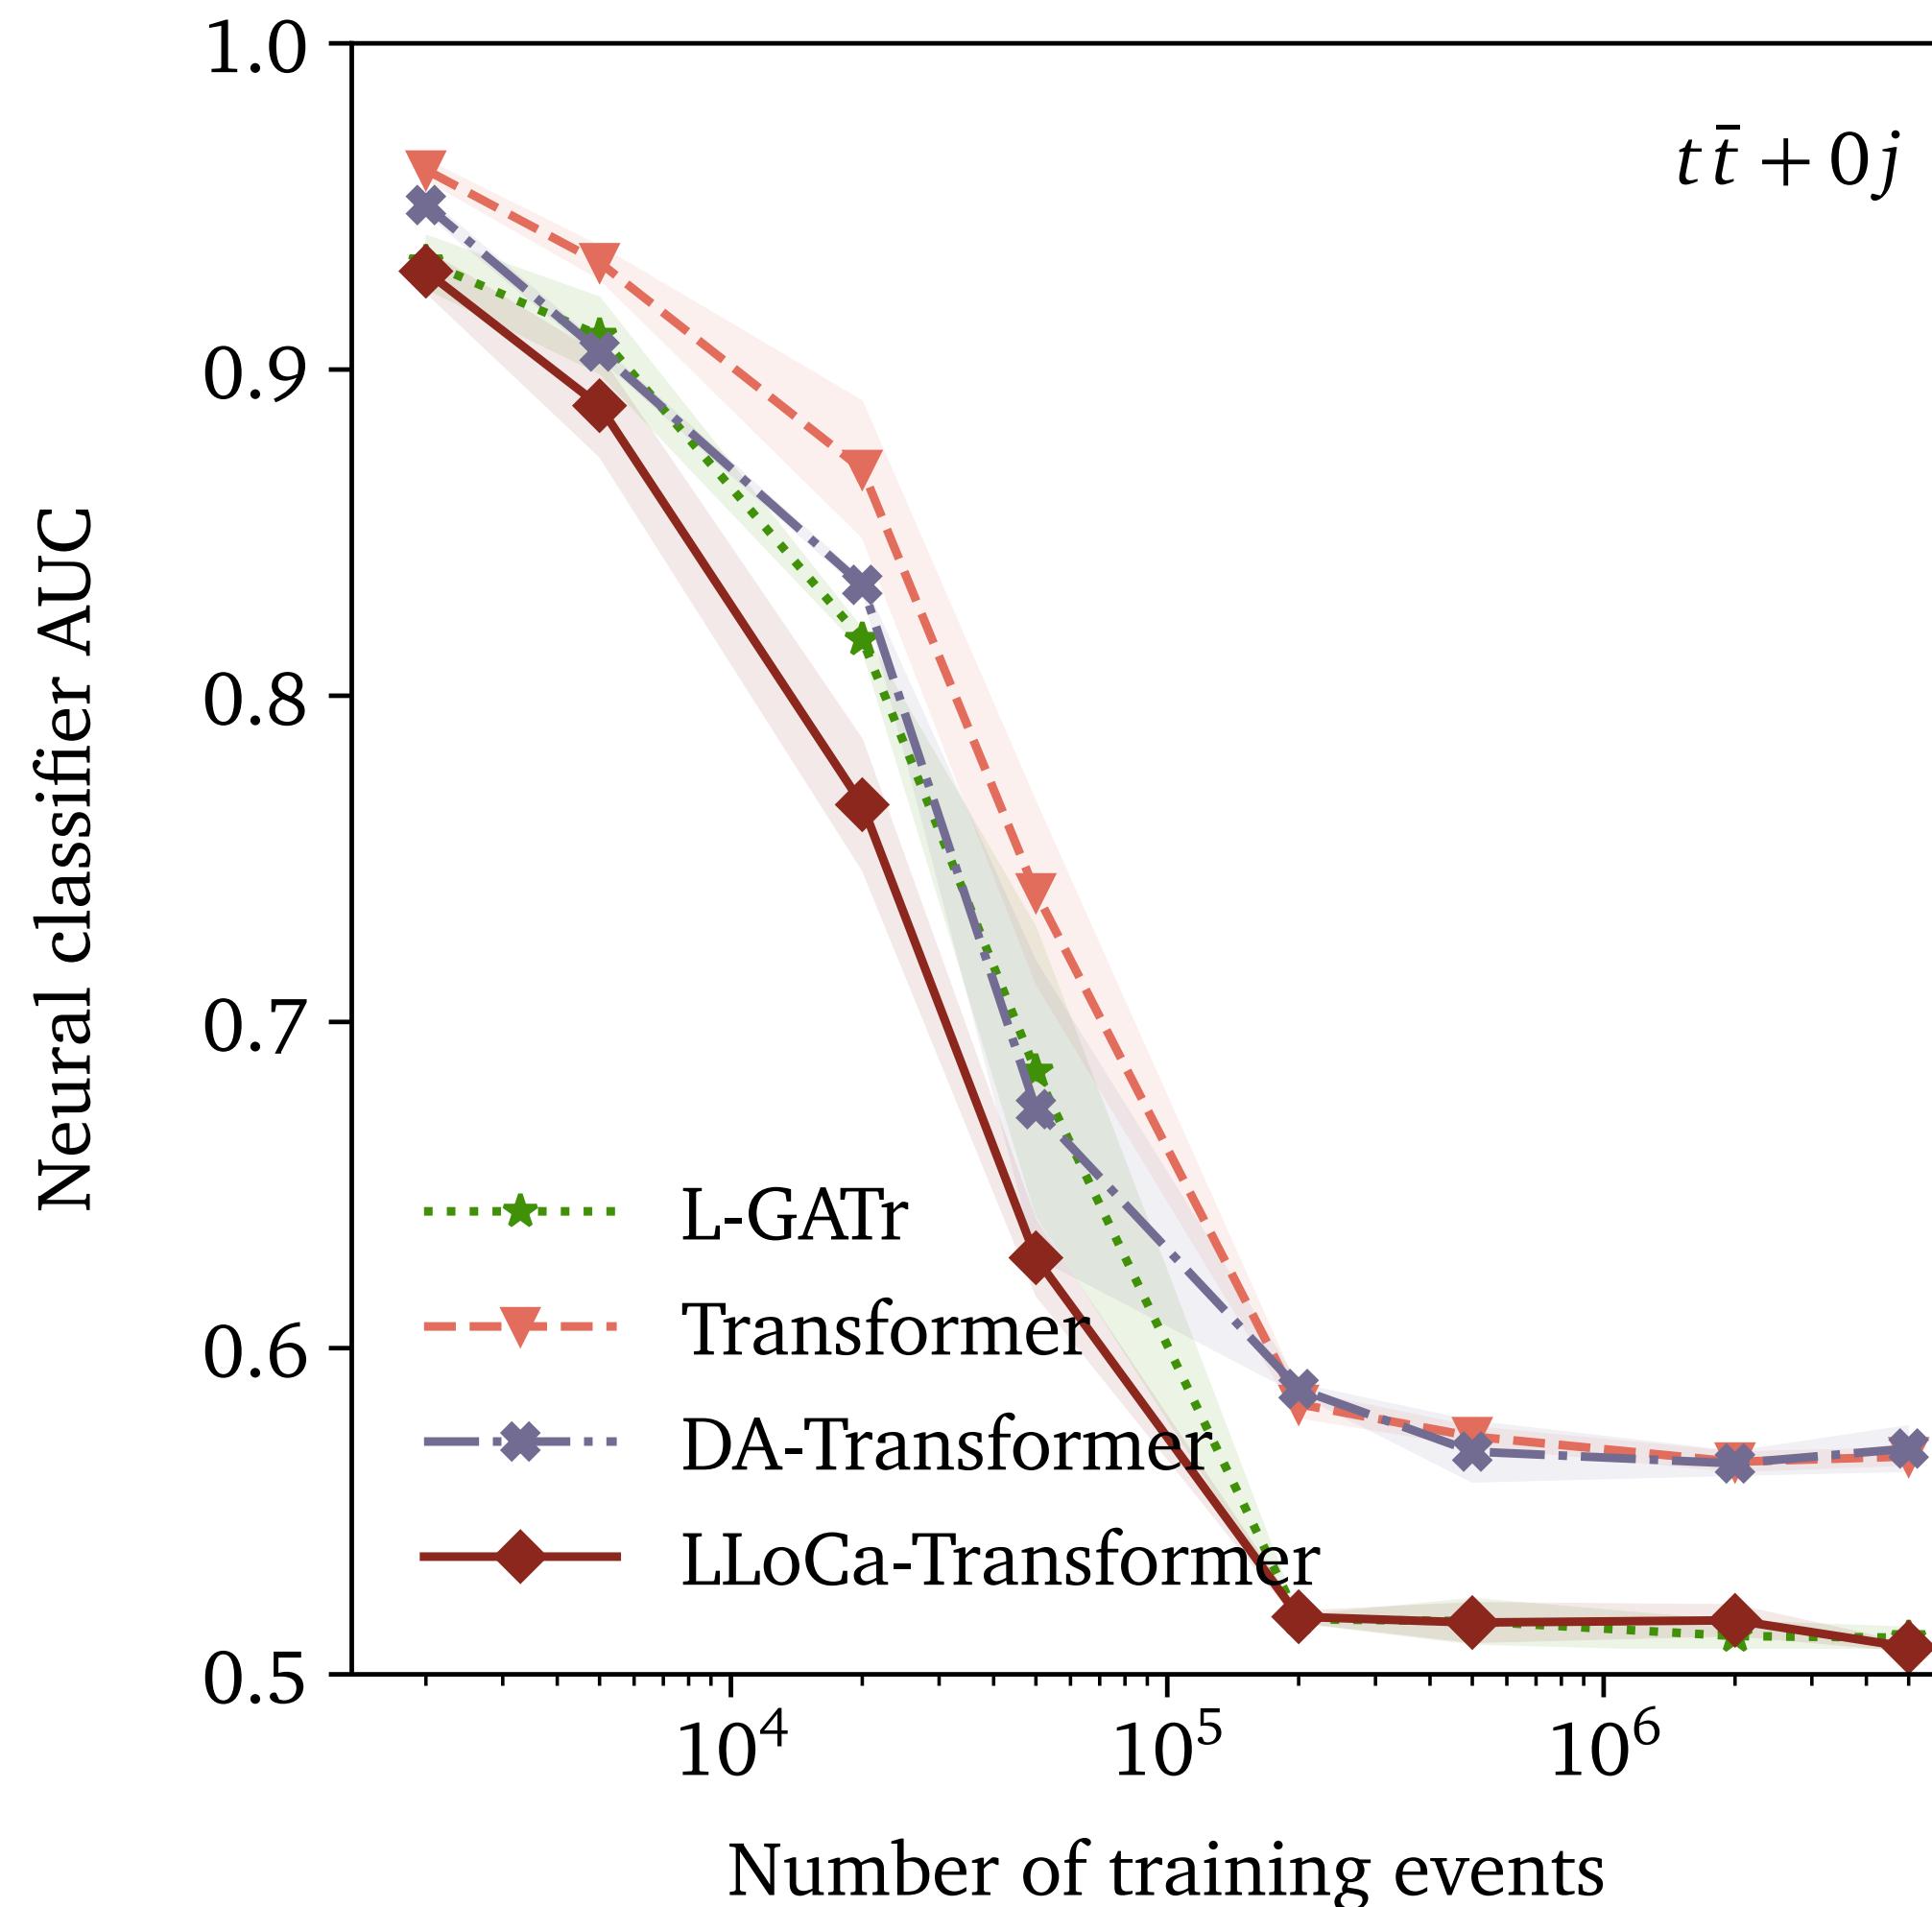
## Kinematic distributions



Equivariance helps,  
especially for angular correlations

# Event generation

## Neural classifier metric

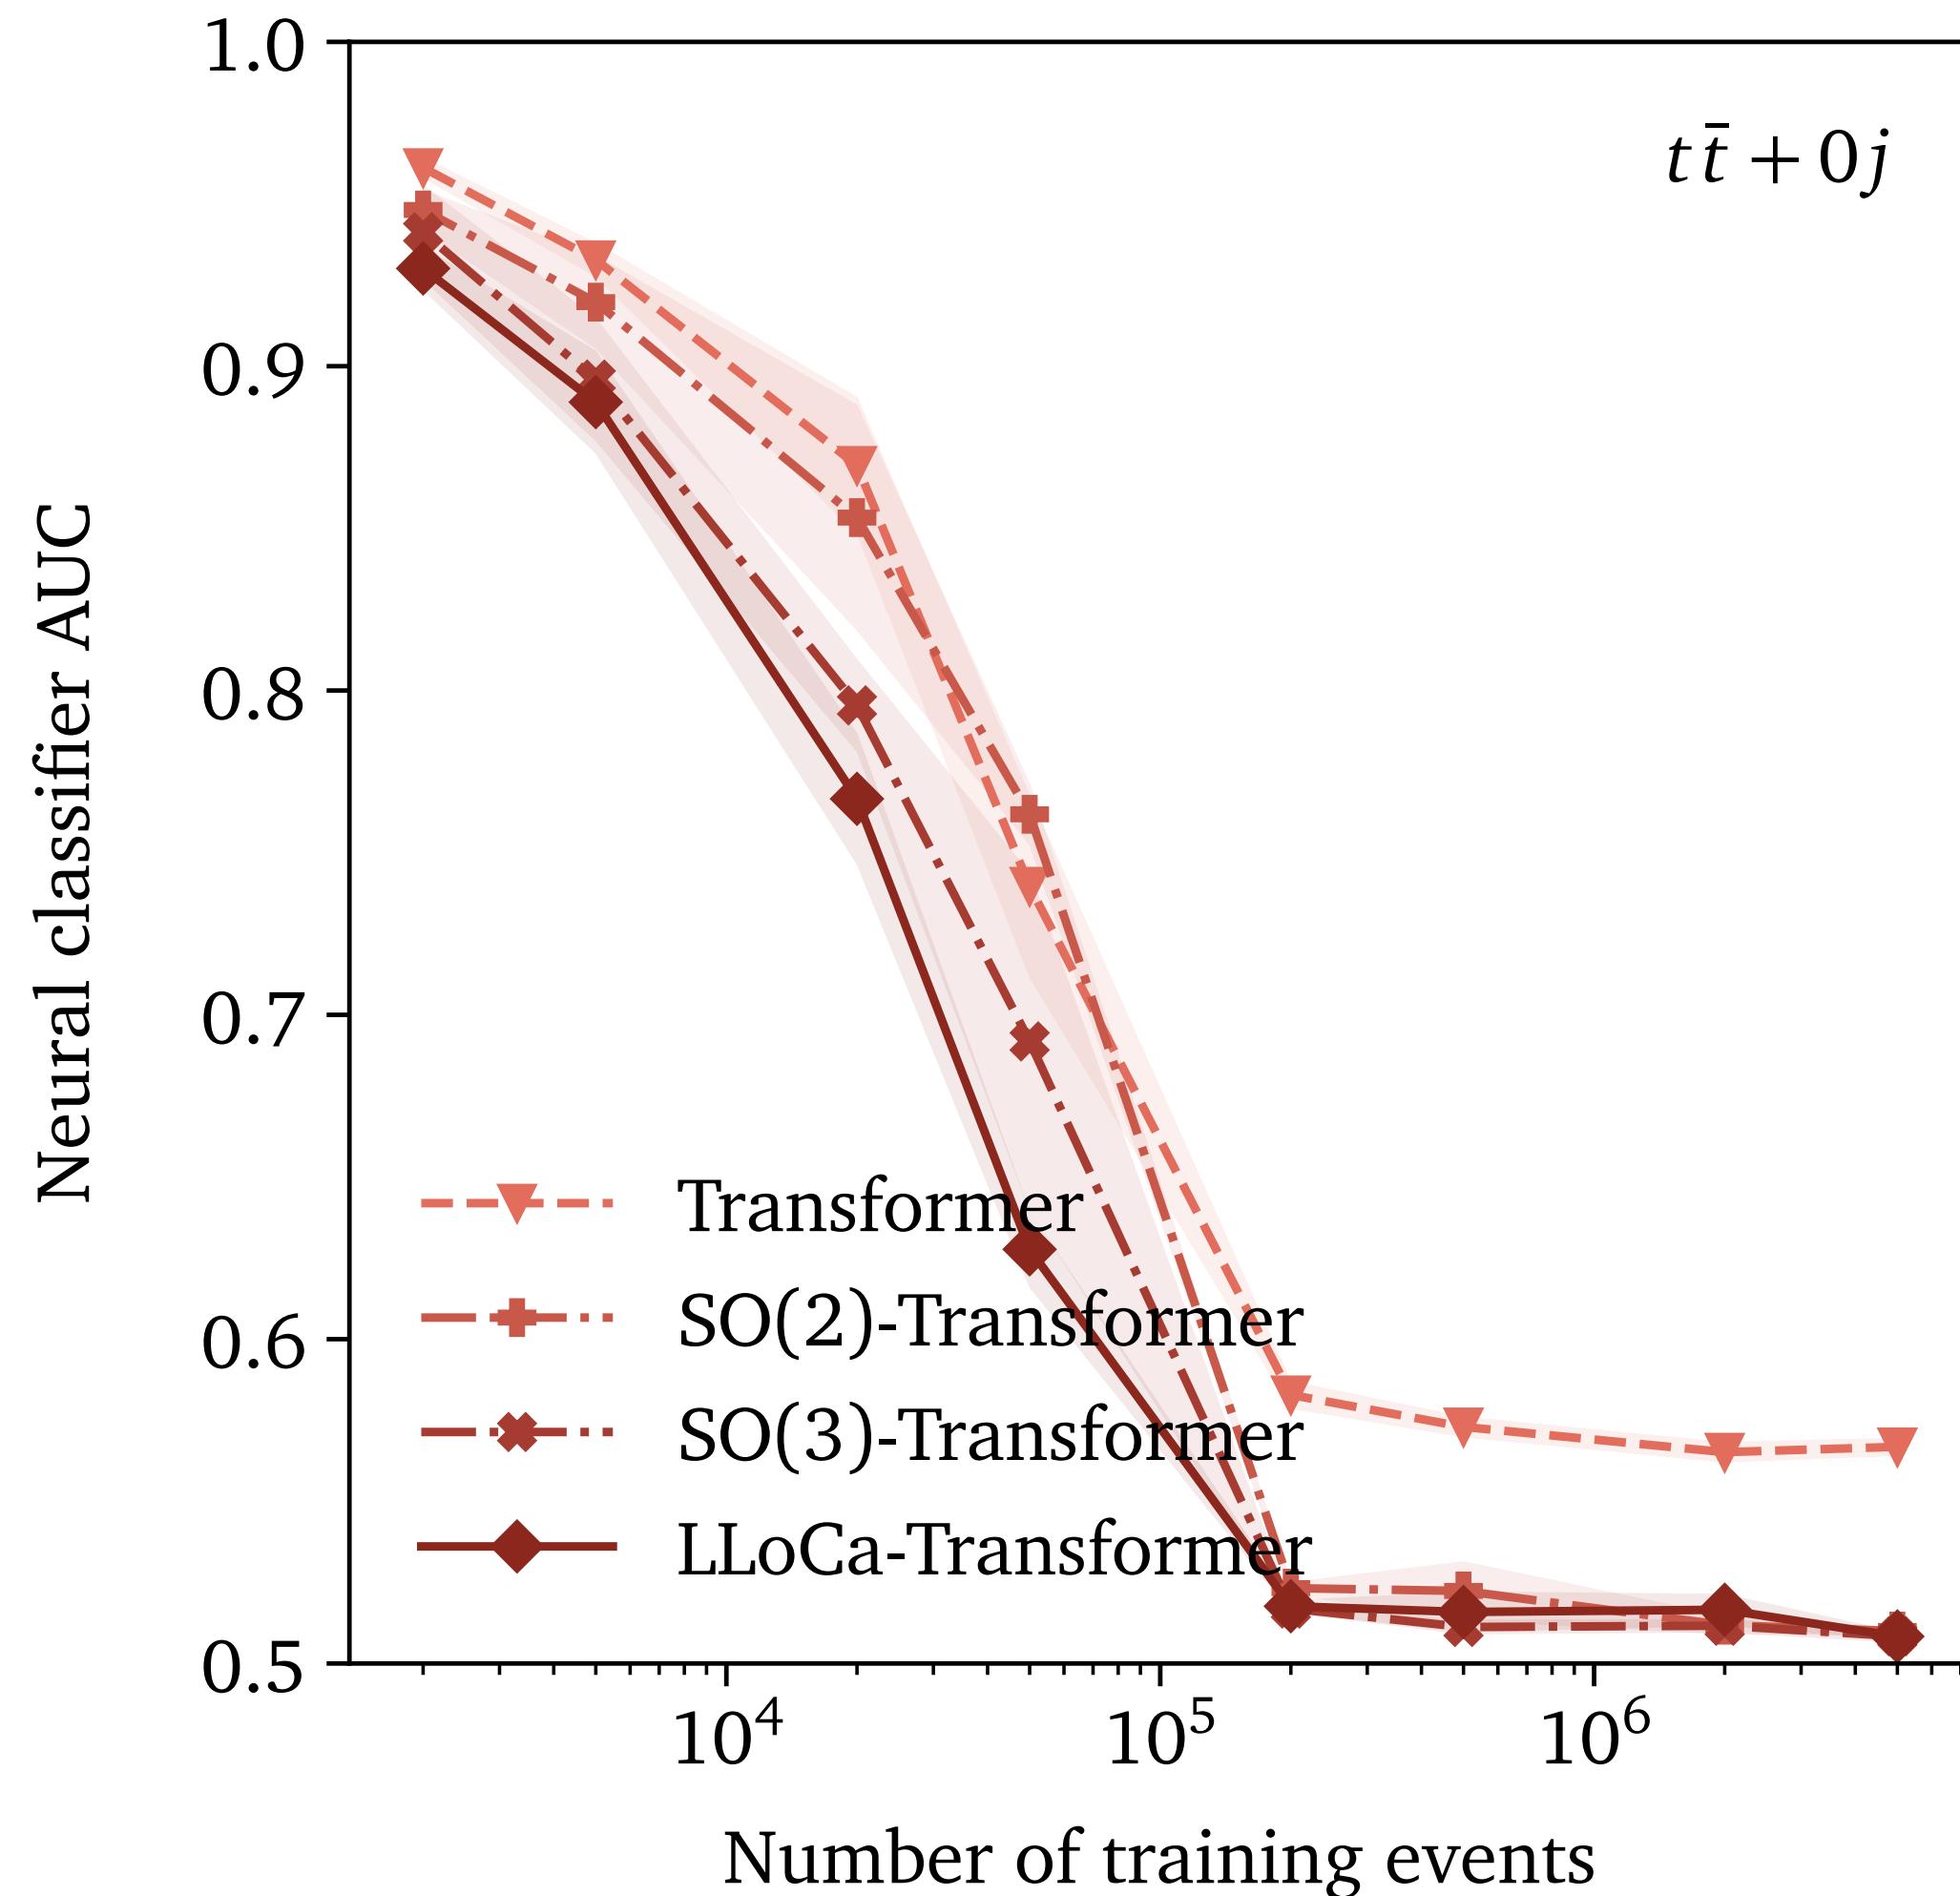


L-GATr and LLoCa-Transformer generate samples that challenge the classifier

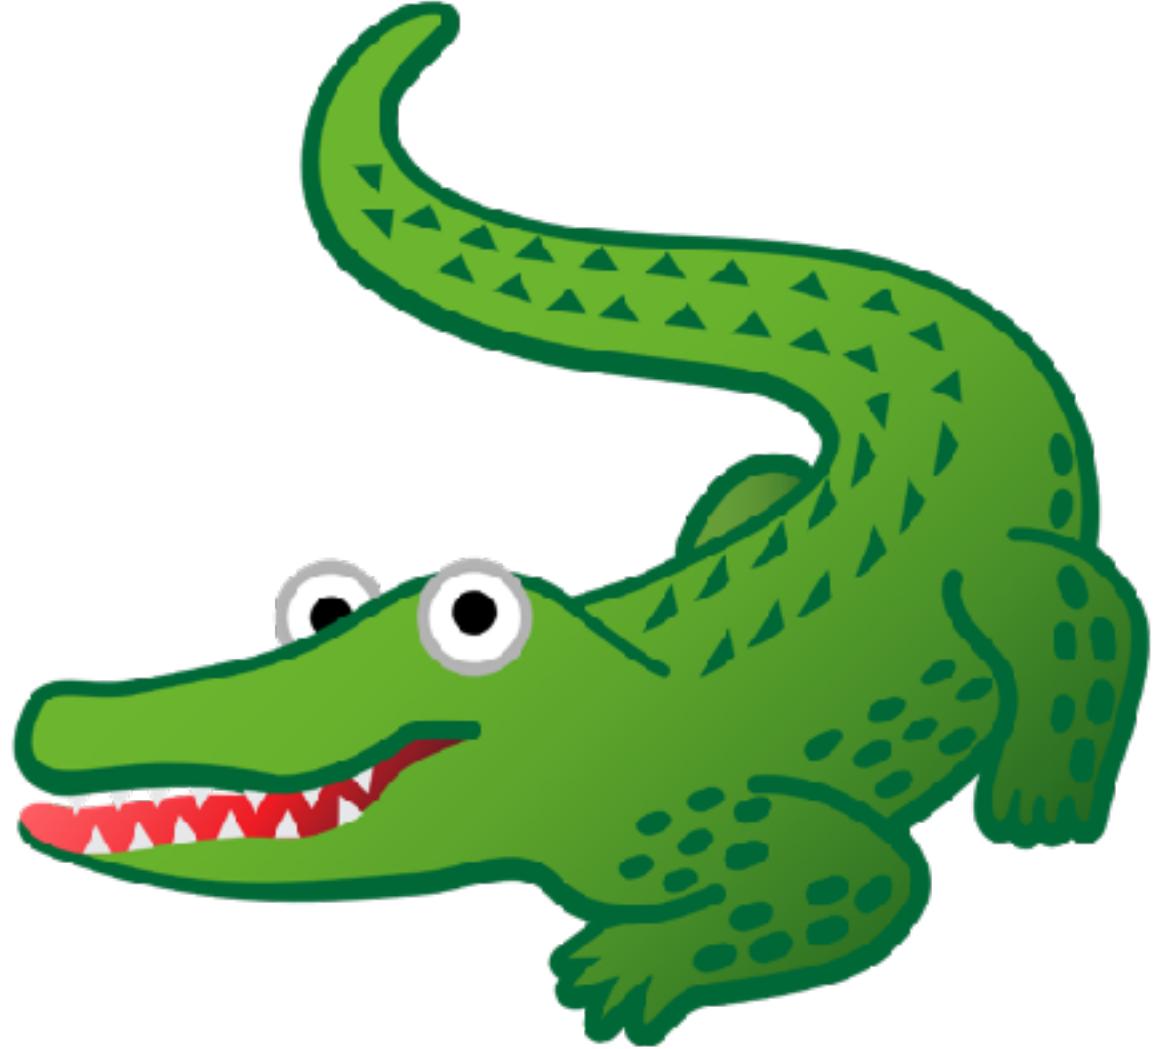
Data augmentation helps for small datasets, but not at scale

# Event generation

## Symmetry breaking

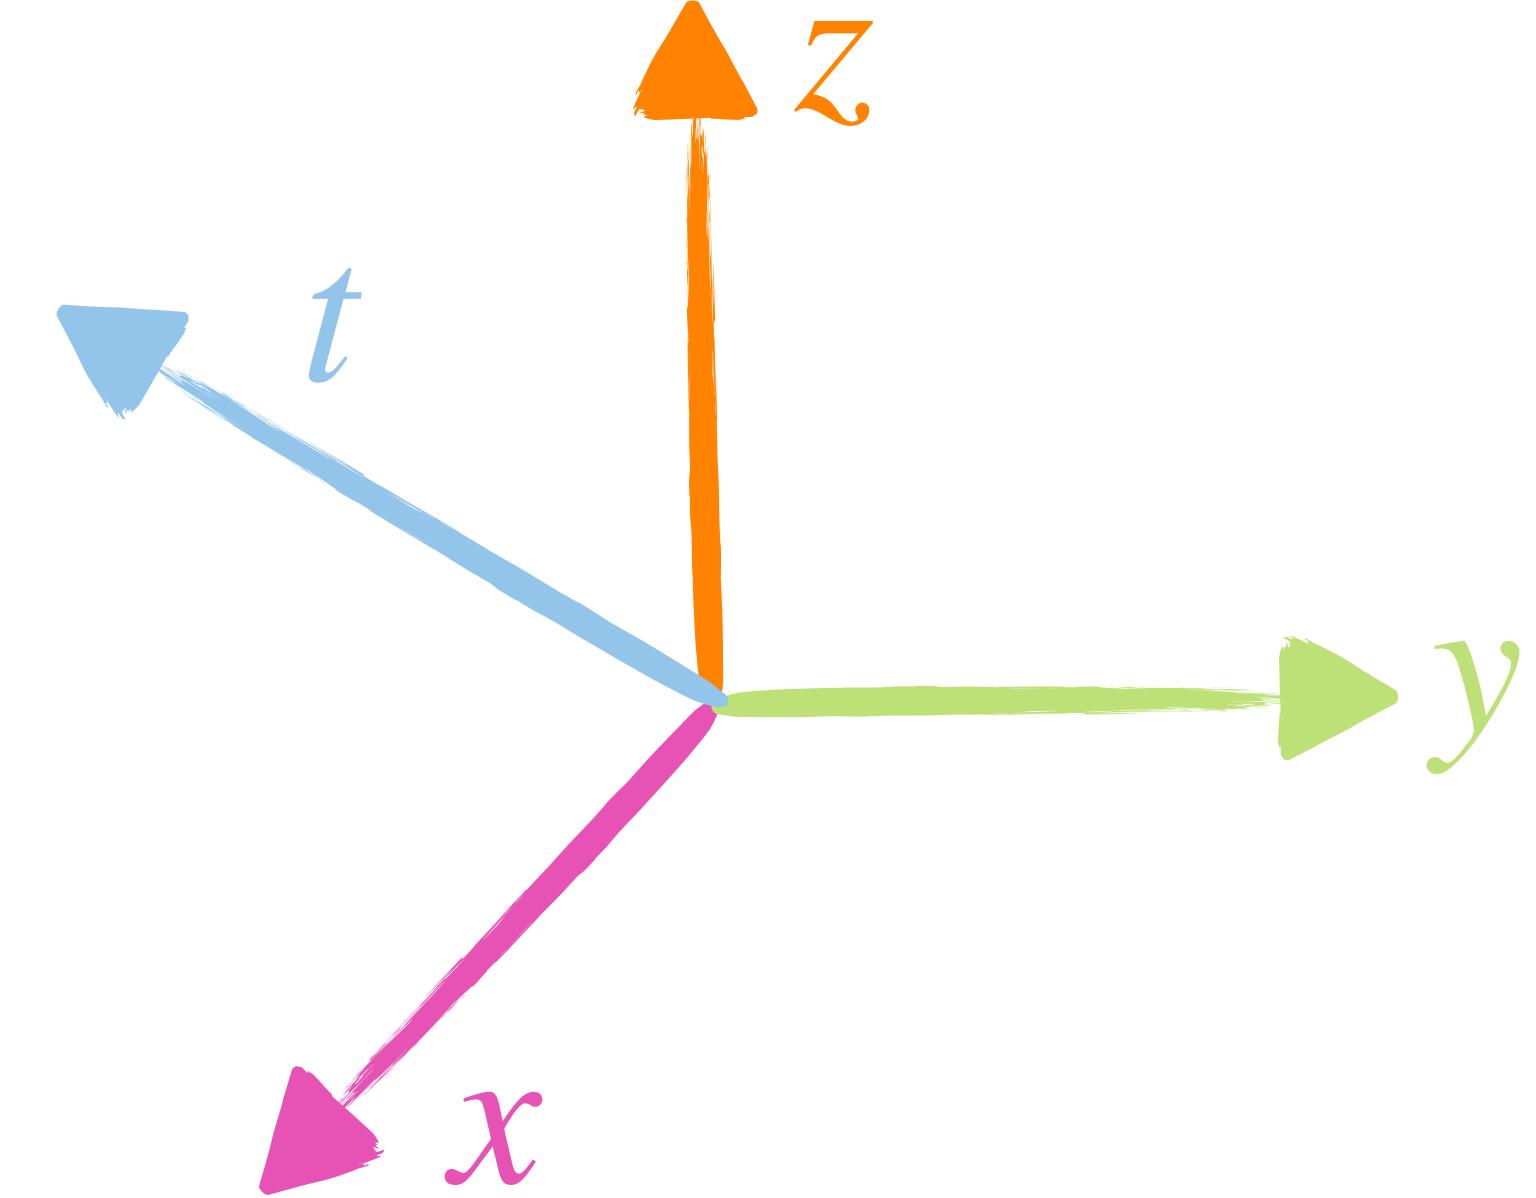


Symmetry breaking at the input and architecture level yield the same result



# Lorentz-Equivariant Geometric Algebra Transformer

- ✓ pip install lgatr
- ✓ Easy to use



**Lorentz Local  
Canonicalization**

- ✓ Resource efficient
- ✓ Flexible



Victor Bresó



Pim de Haan



Tilman Plehn



Huilin Qu



Jesse Thaler



Johann Brehmer

Lorentz-Equivariant Geometric Algebra  
Transformers for High-Energy Physics  
NeurIPS 2024, arXiv:2405.14806

A Lorentz-Equivariant Transformer  
for all of the LHC  
SciPost Physics 2025, arXiv:2411.00446



Luigi Favaro



Peter Lippmann



Sebastian Pitz



Gerrit Gerhartz



Tilman Plehn



Huilin Qu



Fred Hamprecht

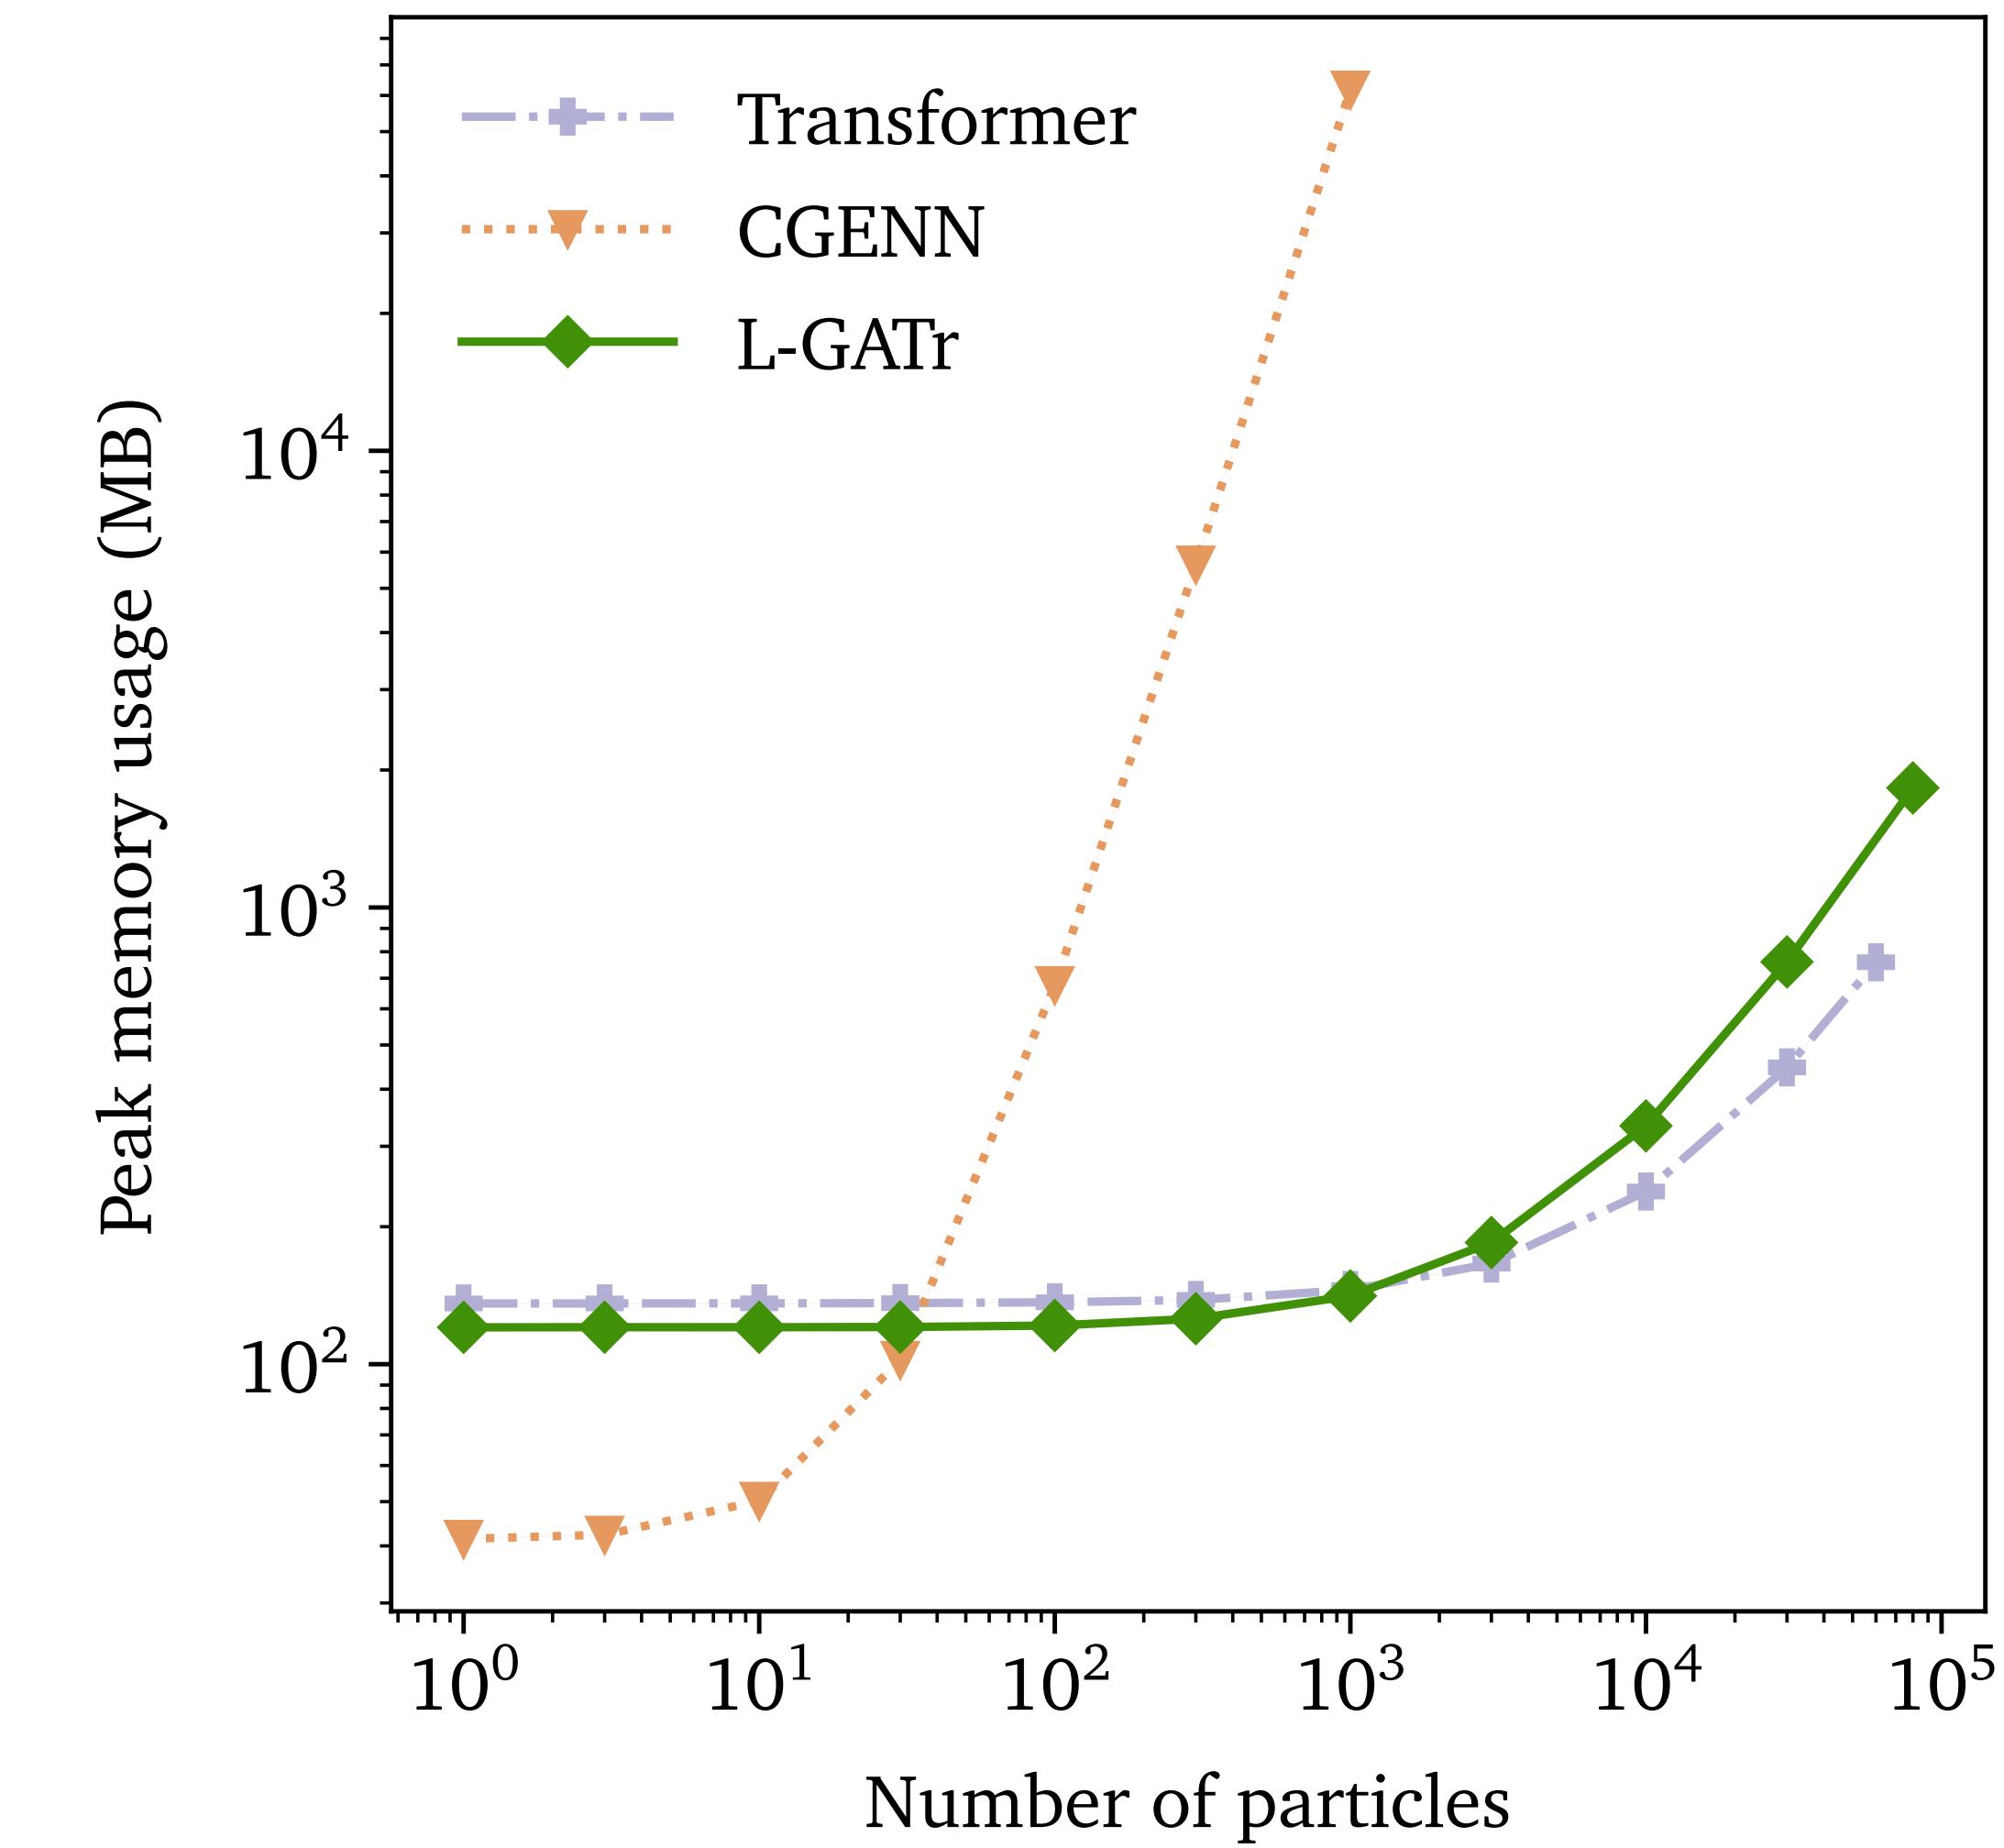
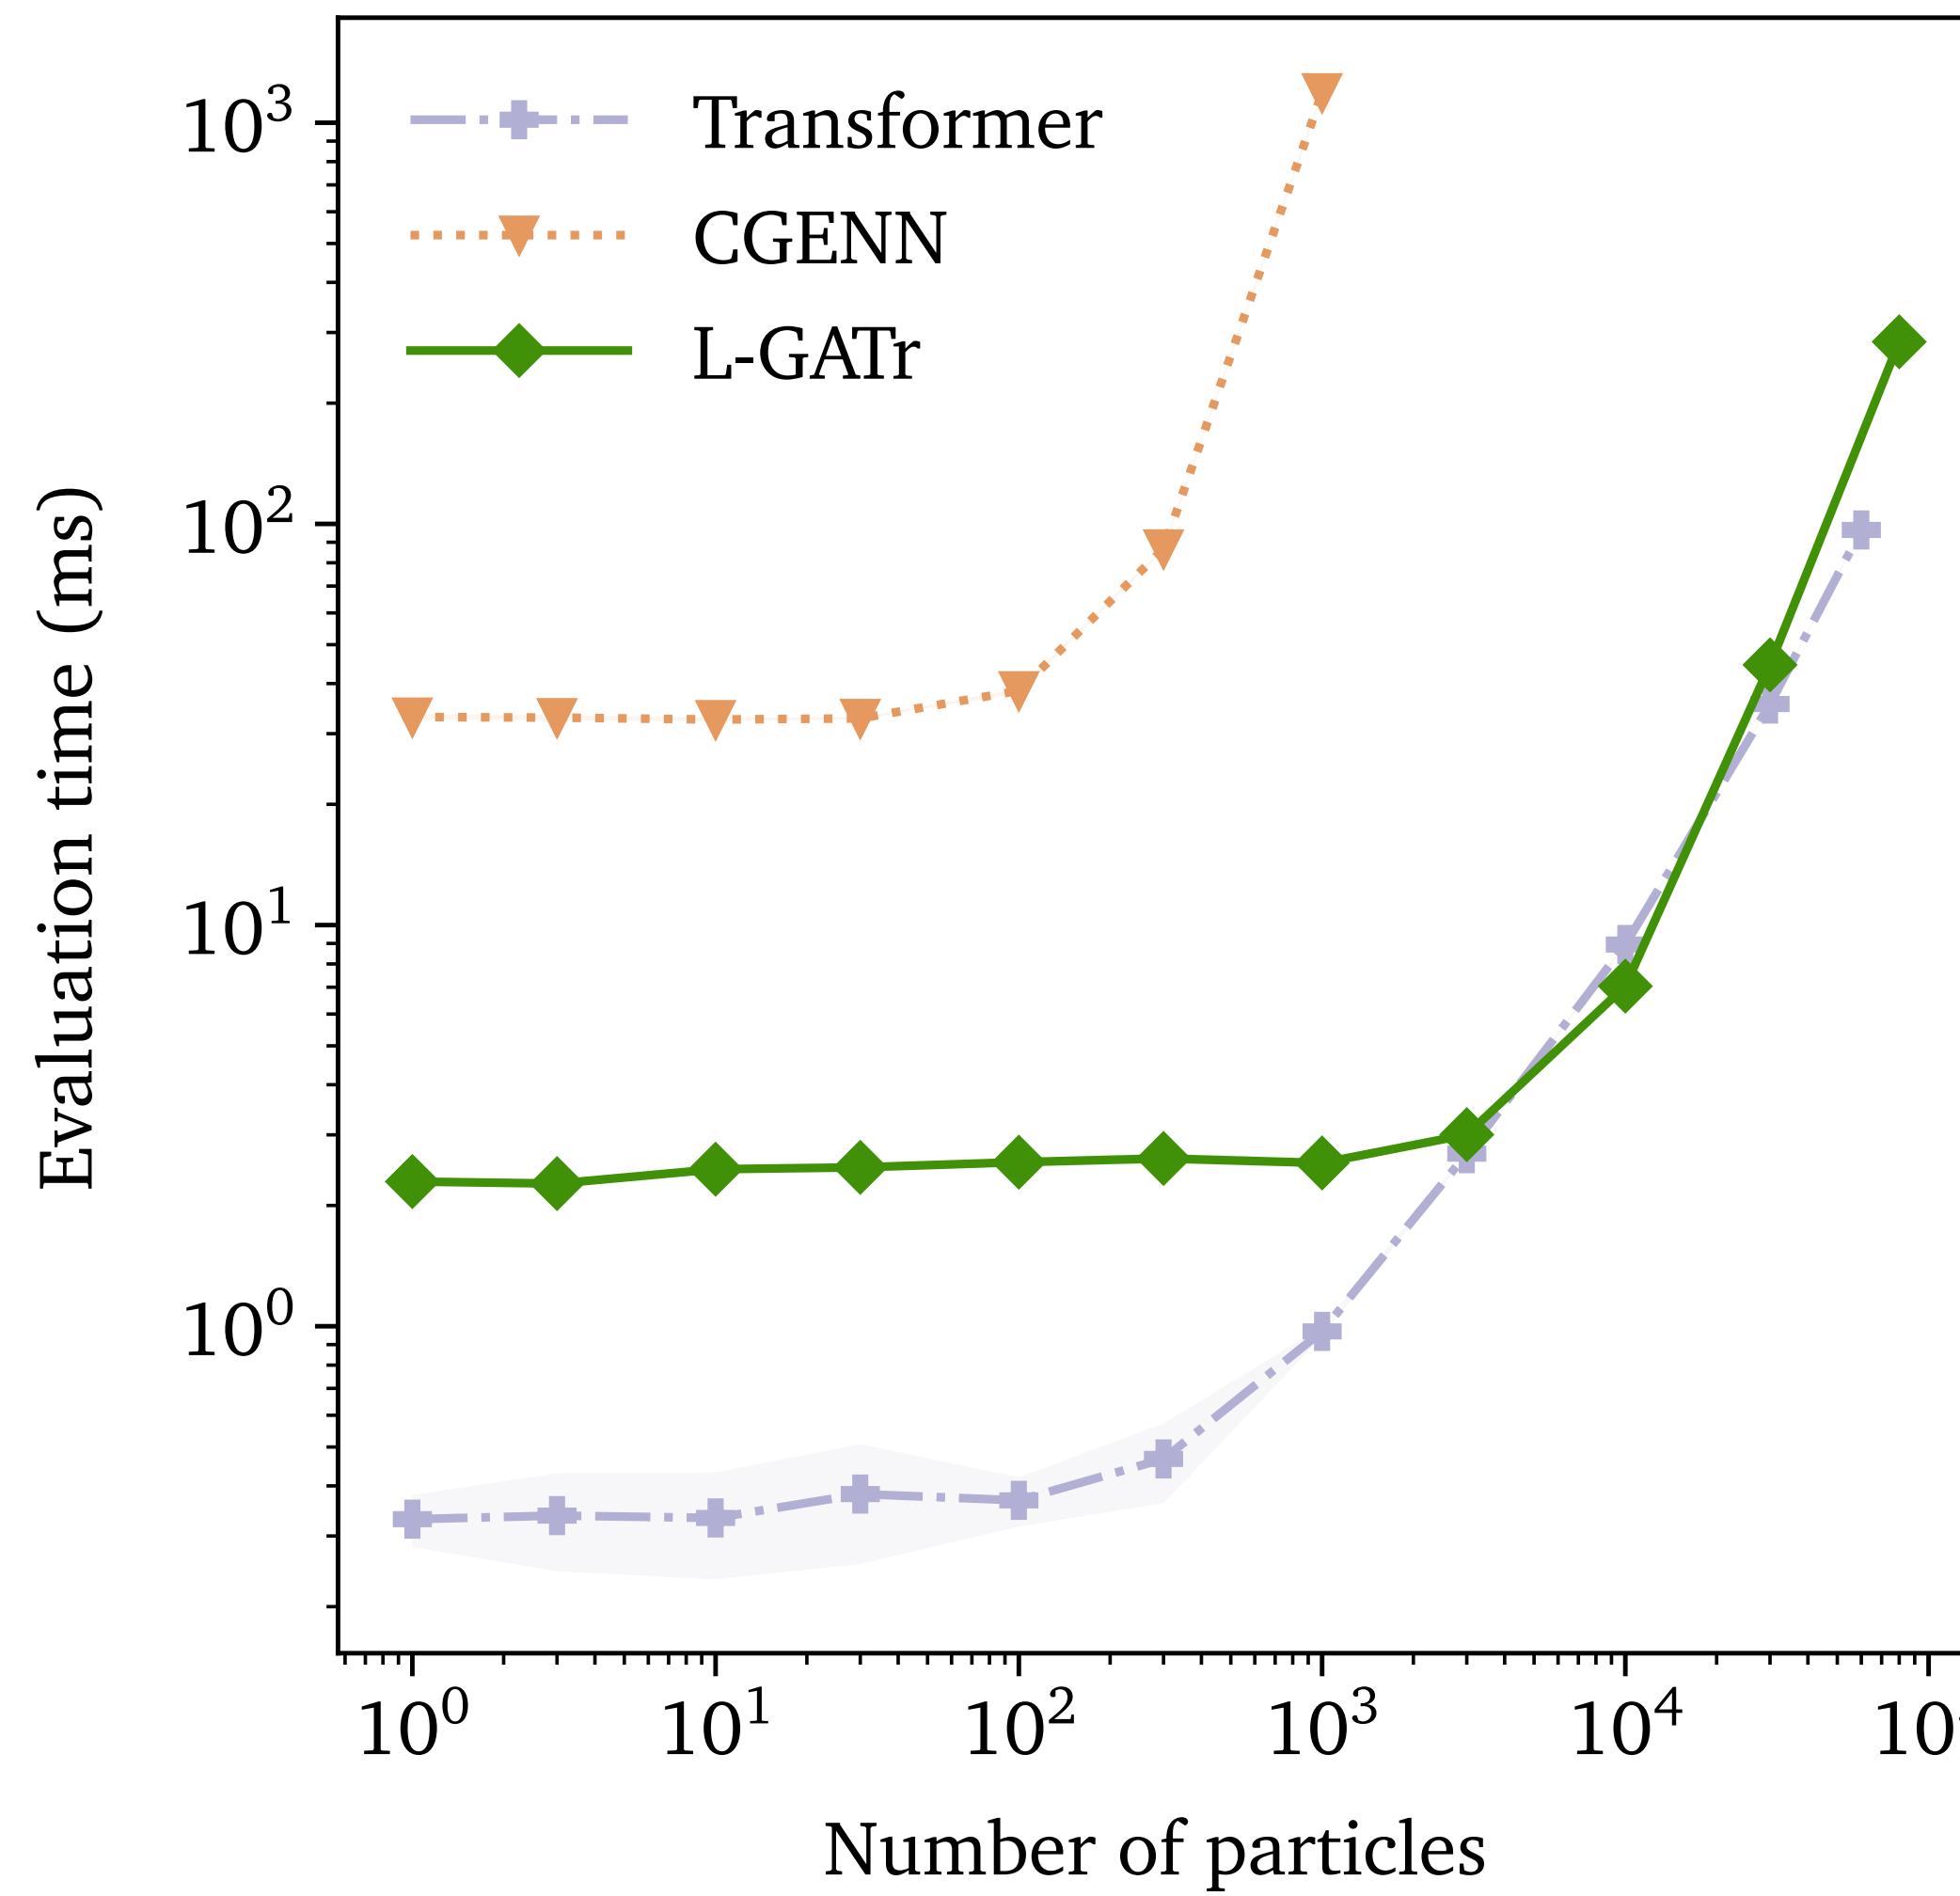
Lorentz Local Canonicalization:  
How to make any network Lorentz-equivariant  
Under review at NeurIPS 2025, arXiv:2505.20280

Lorentz-Equivariance without Limitations  
arXiv:2508.14898

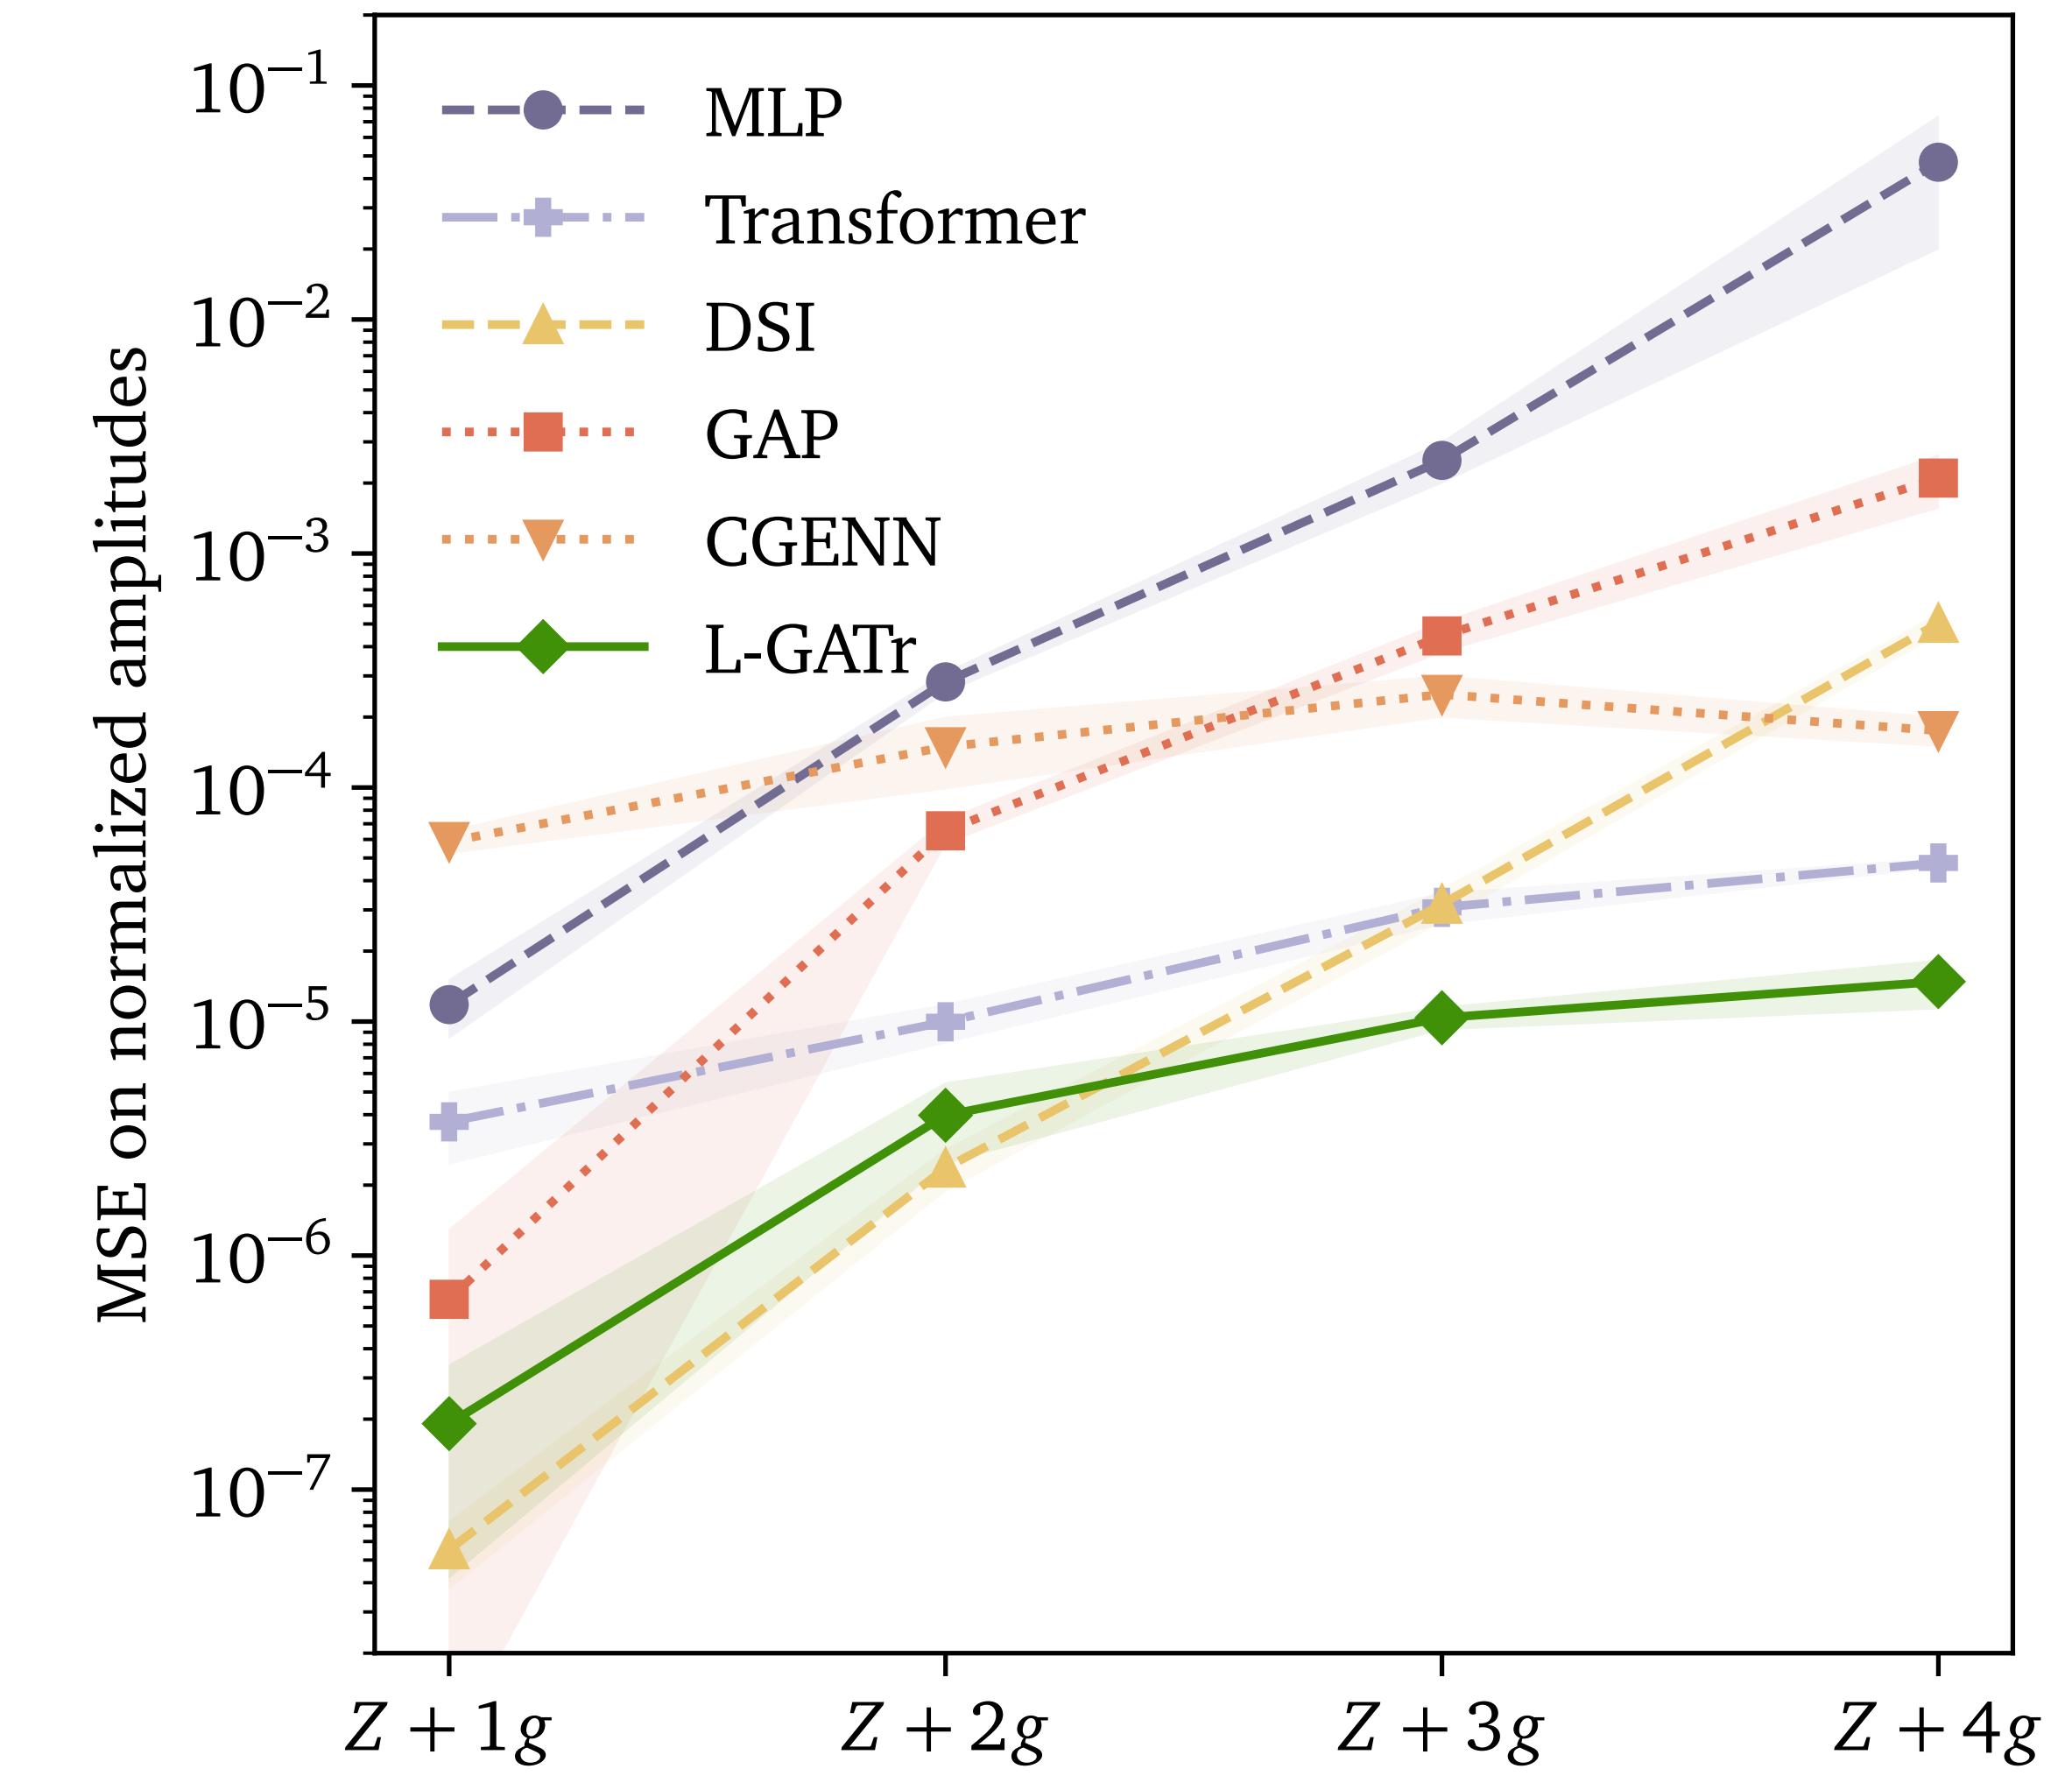
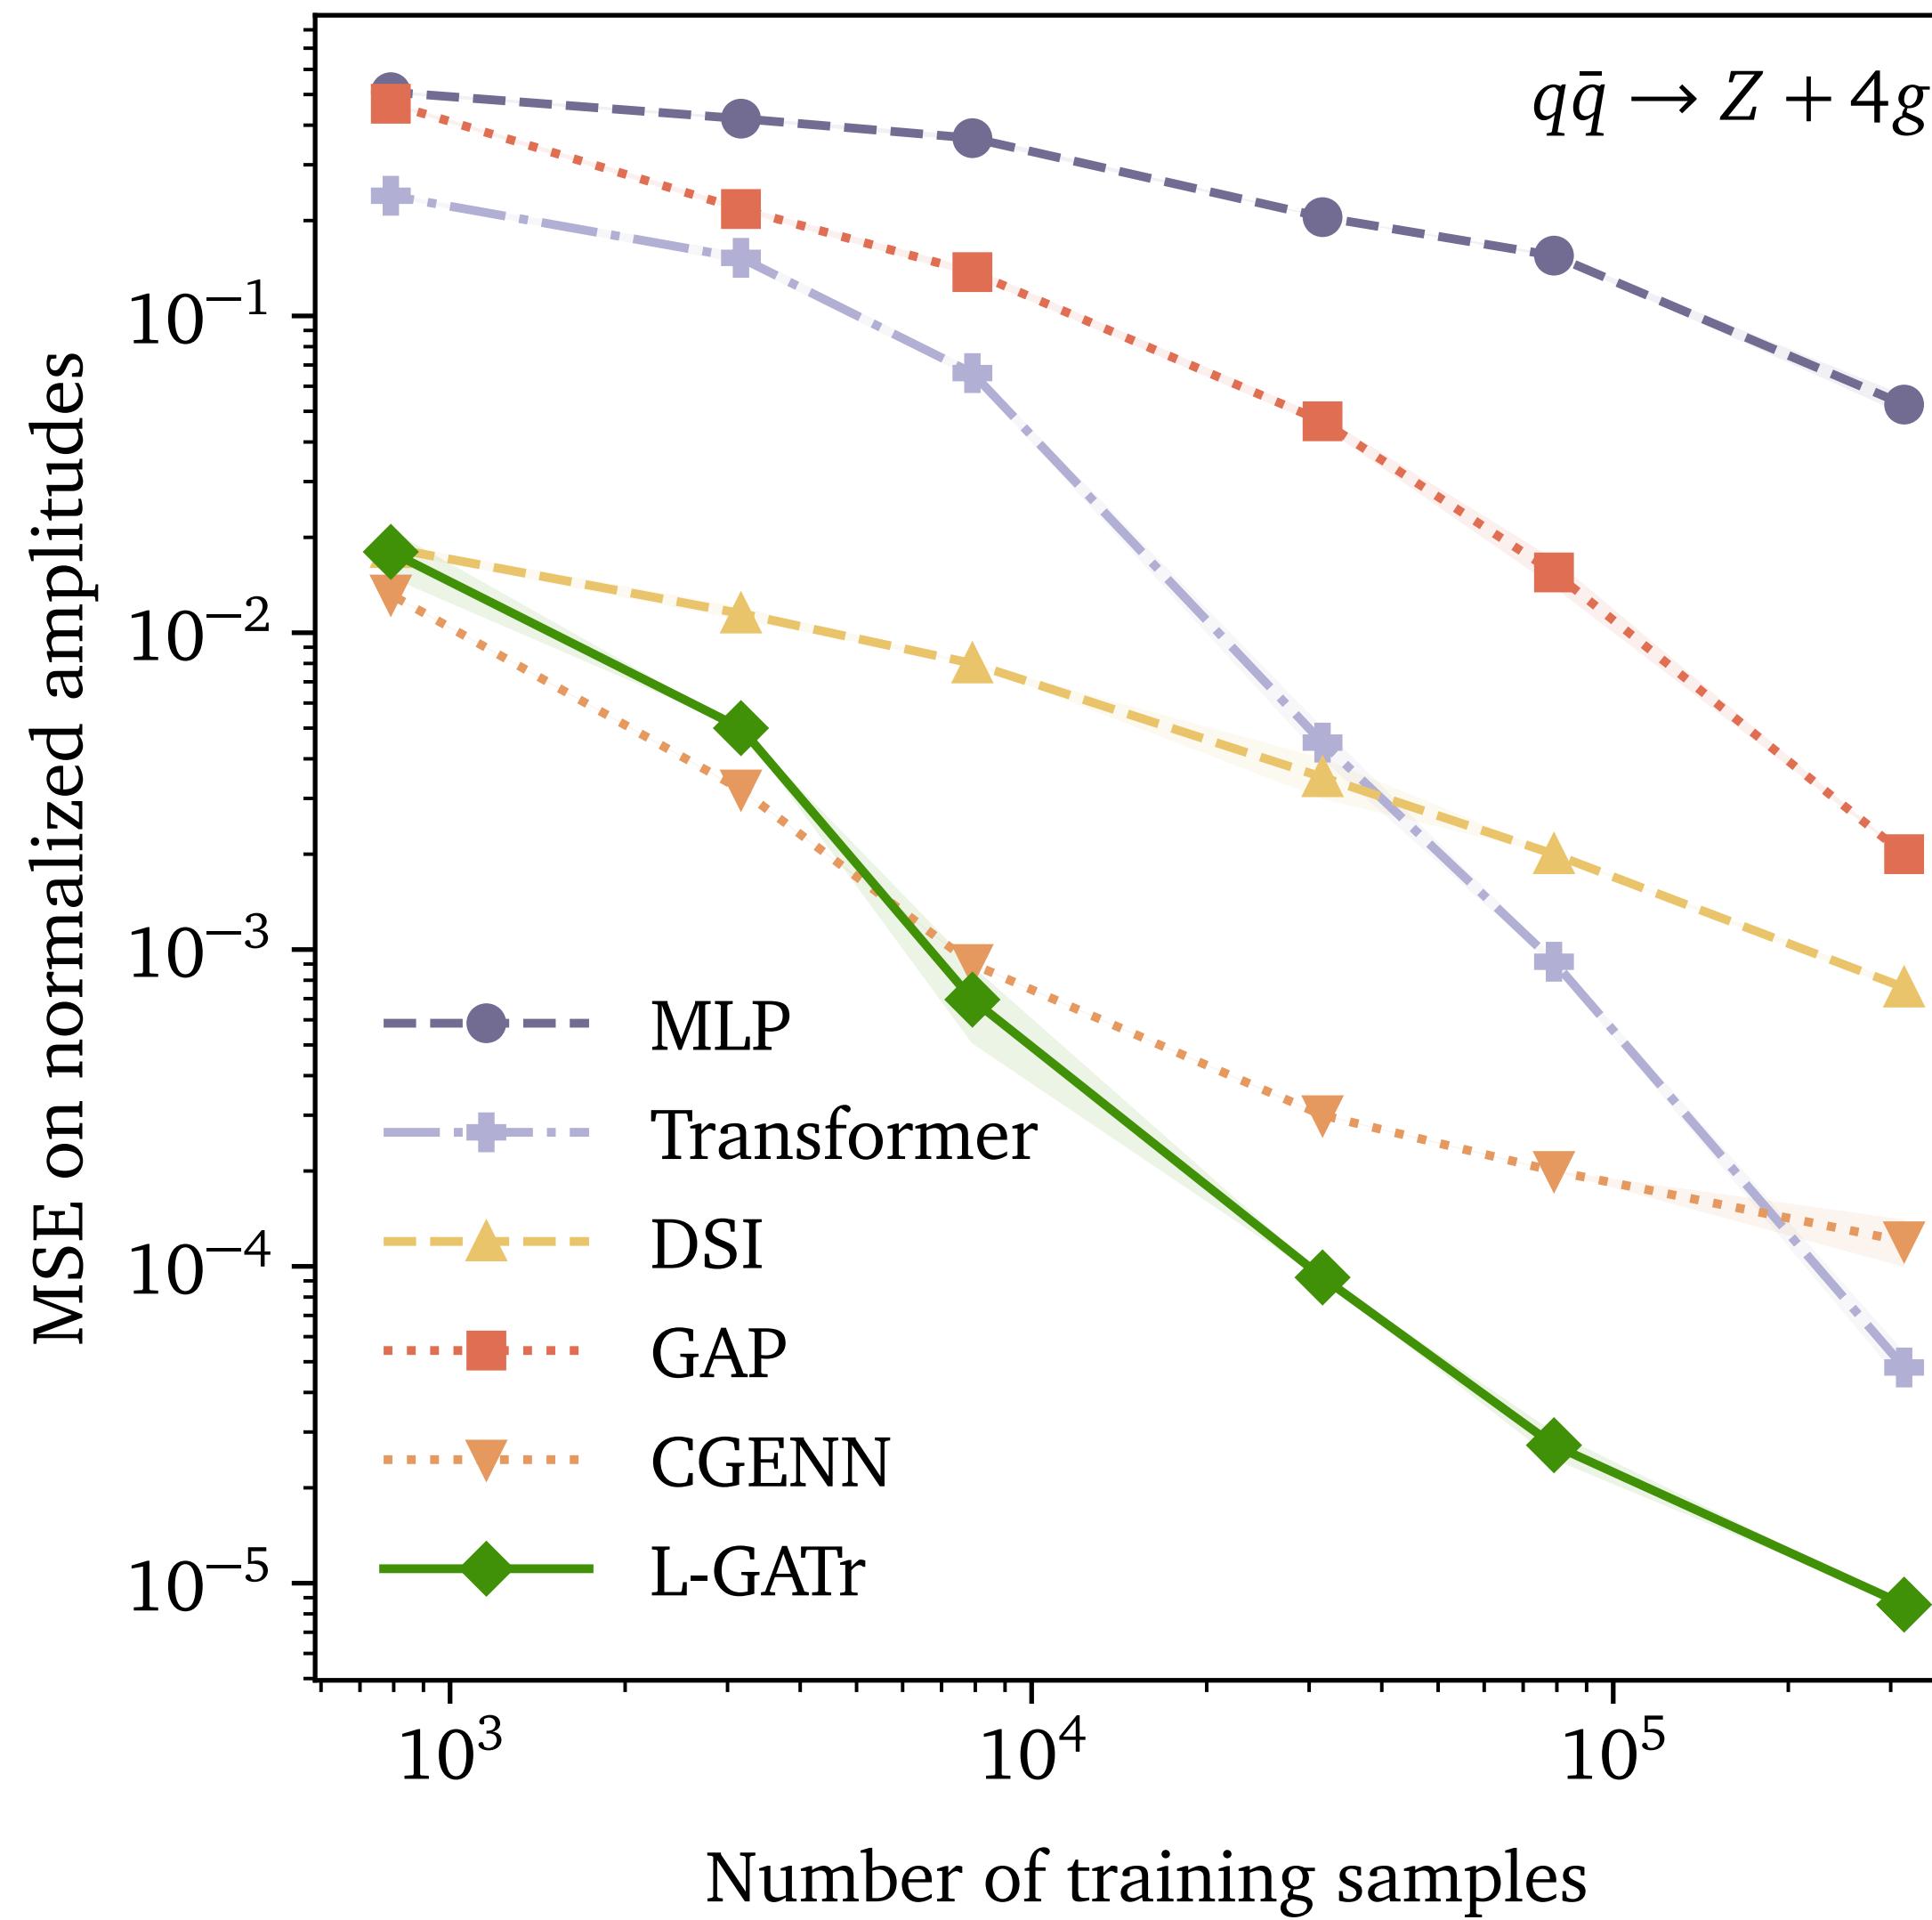
# Bonus material

# L-GATr

Transformers scale better than graphs



# Amplitude regression

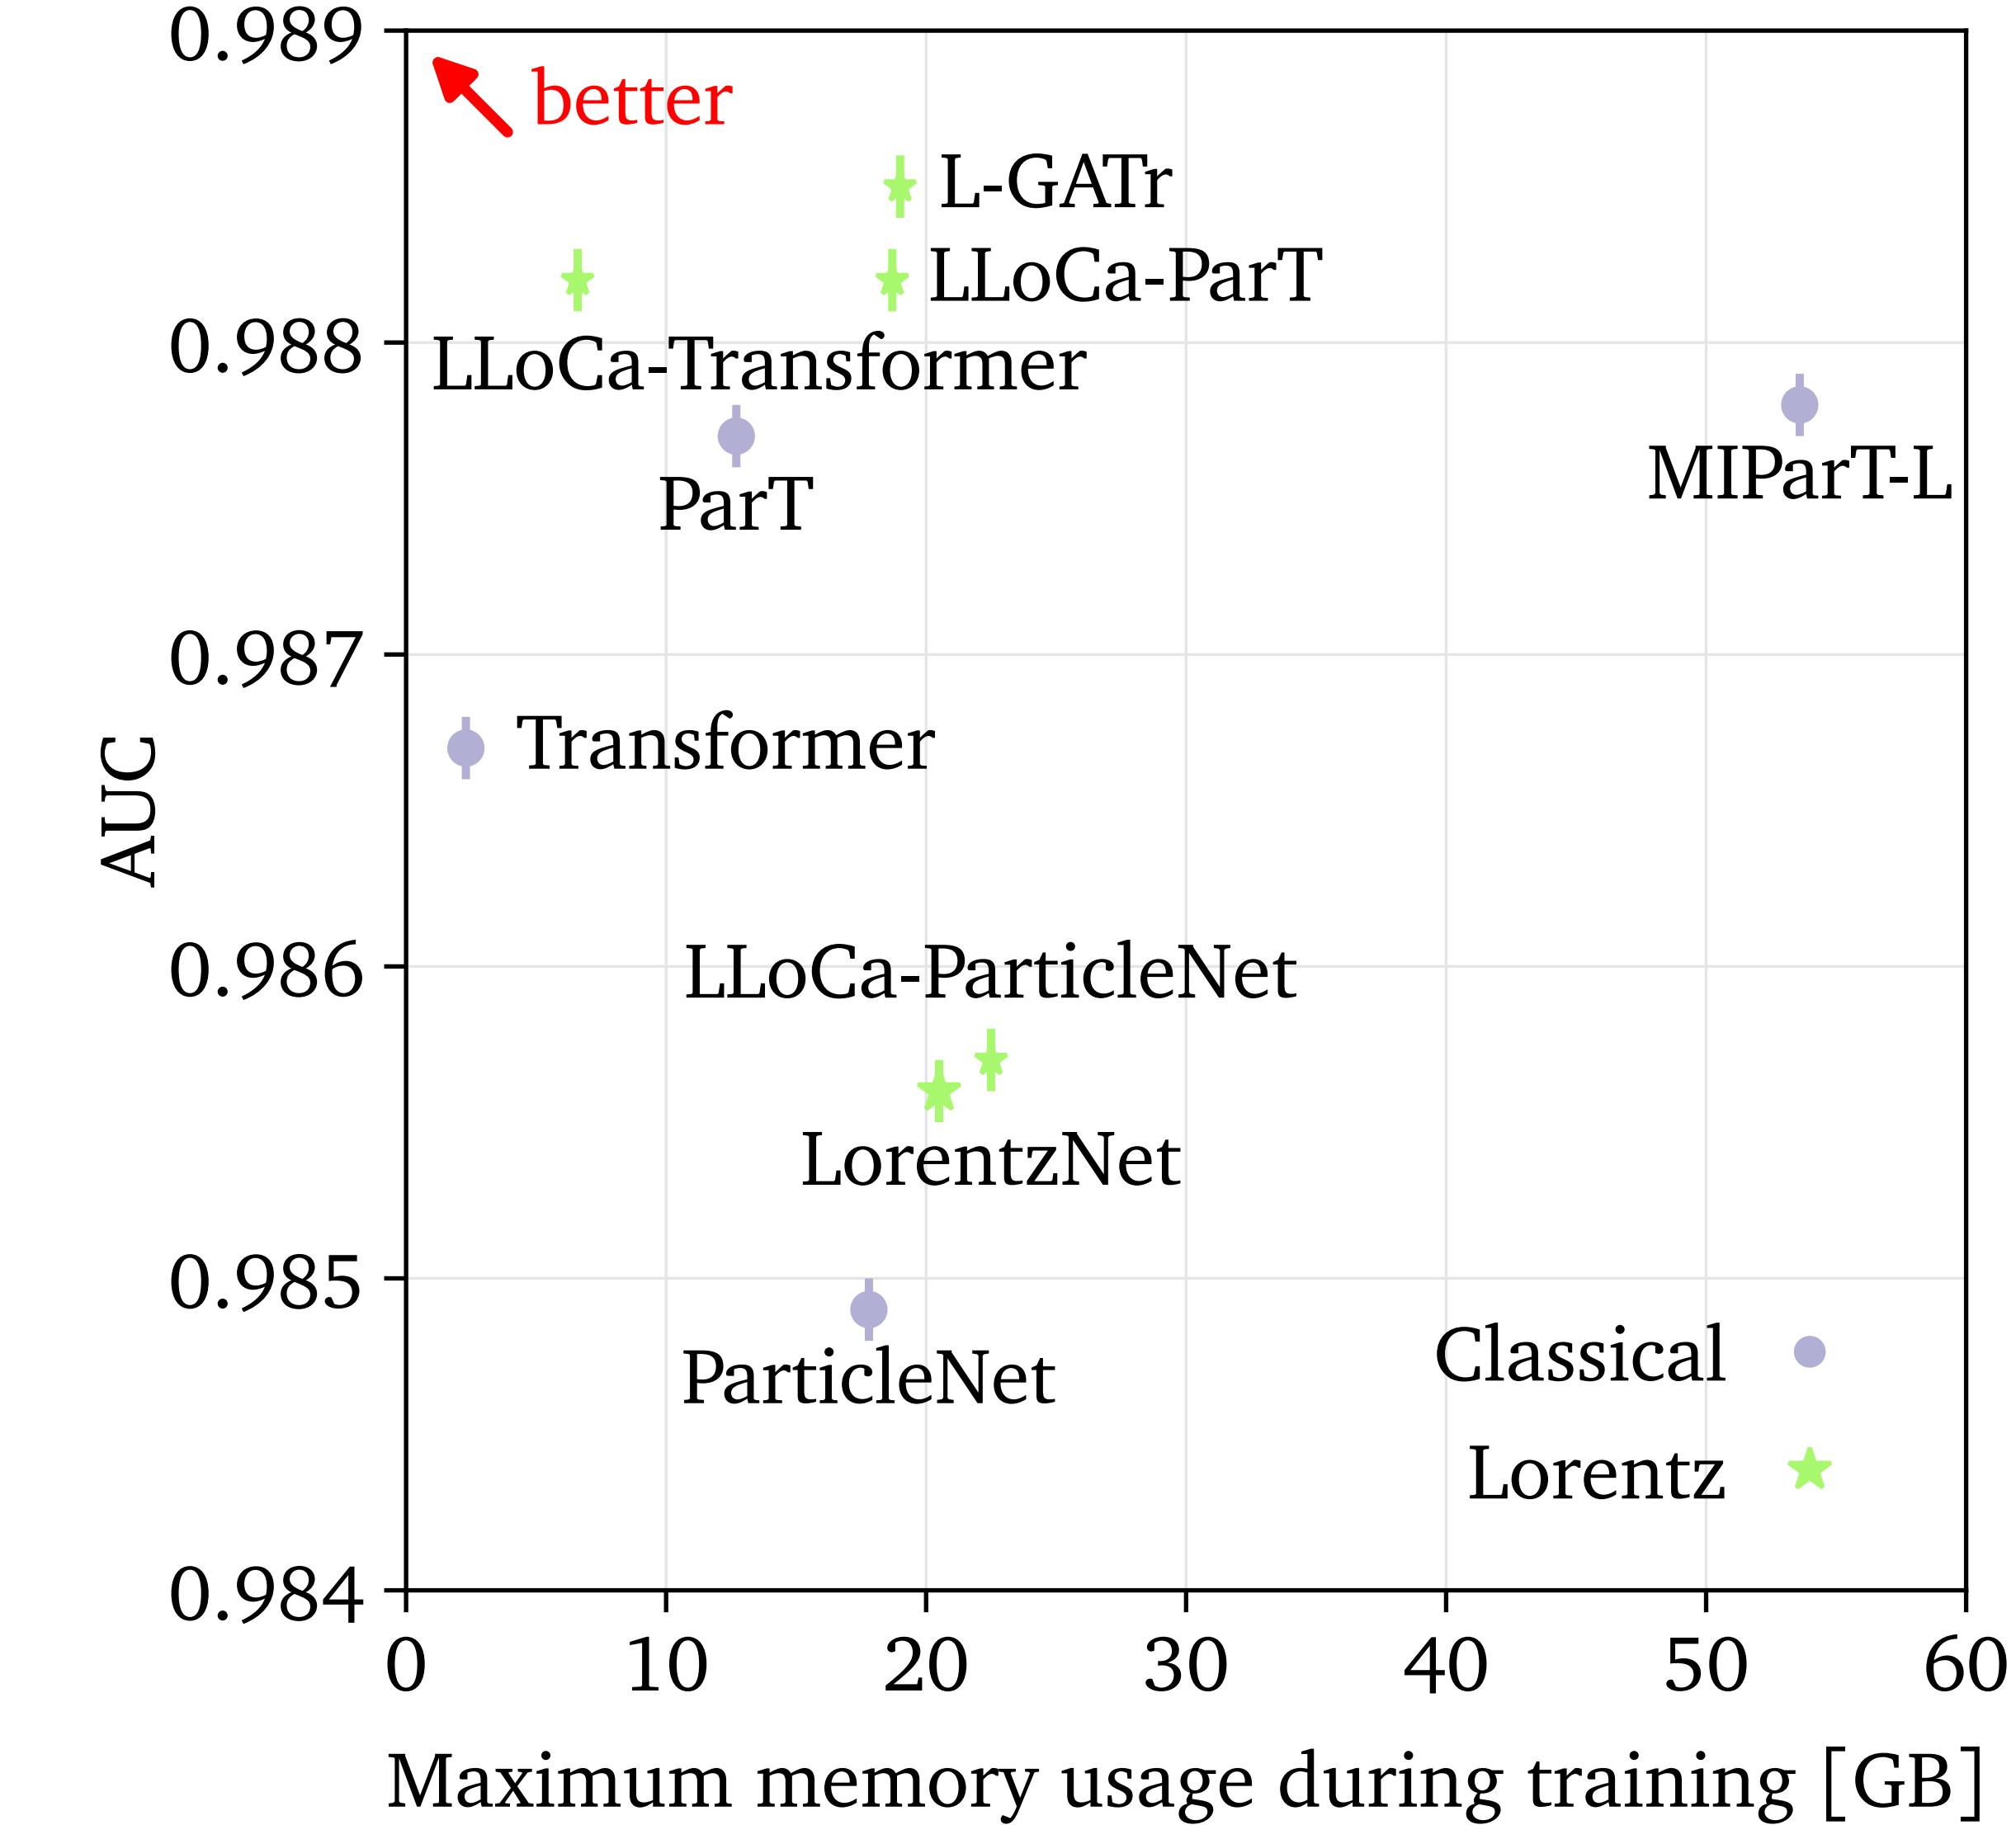
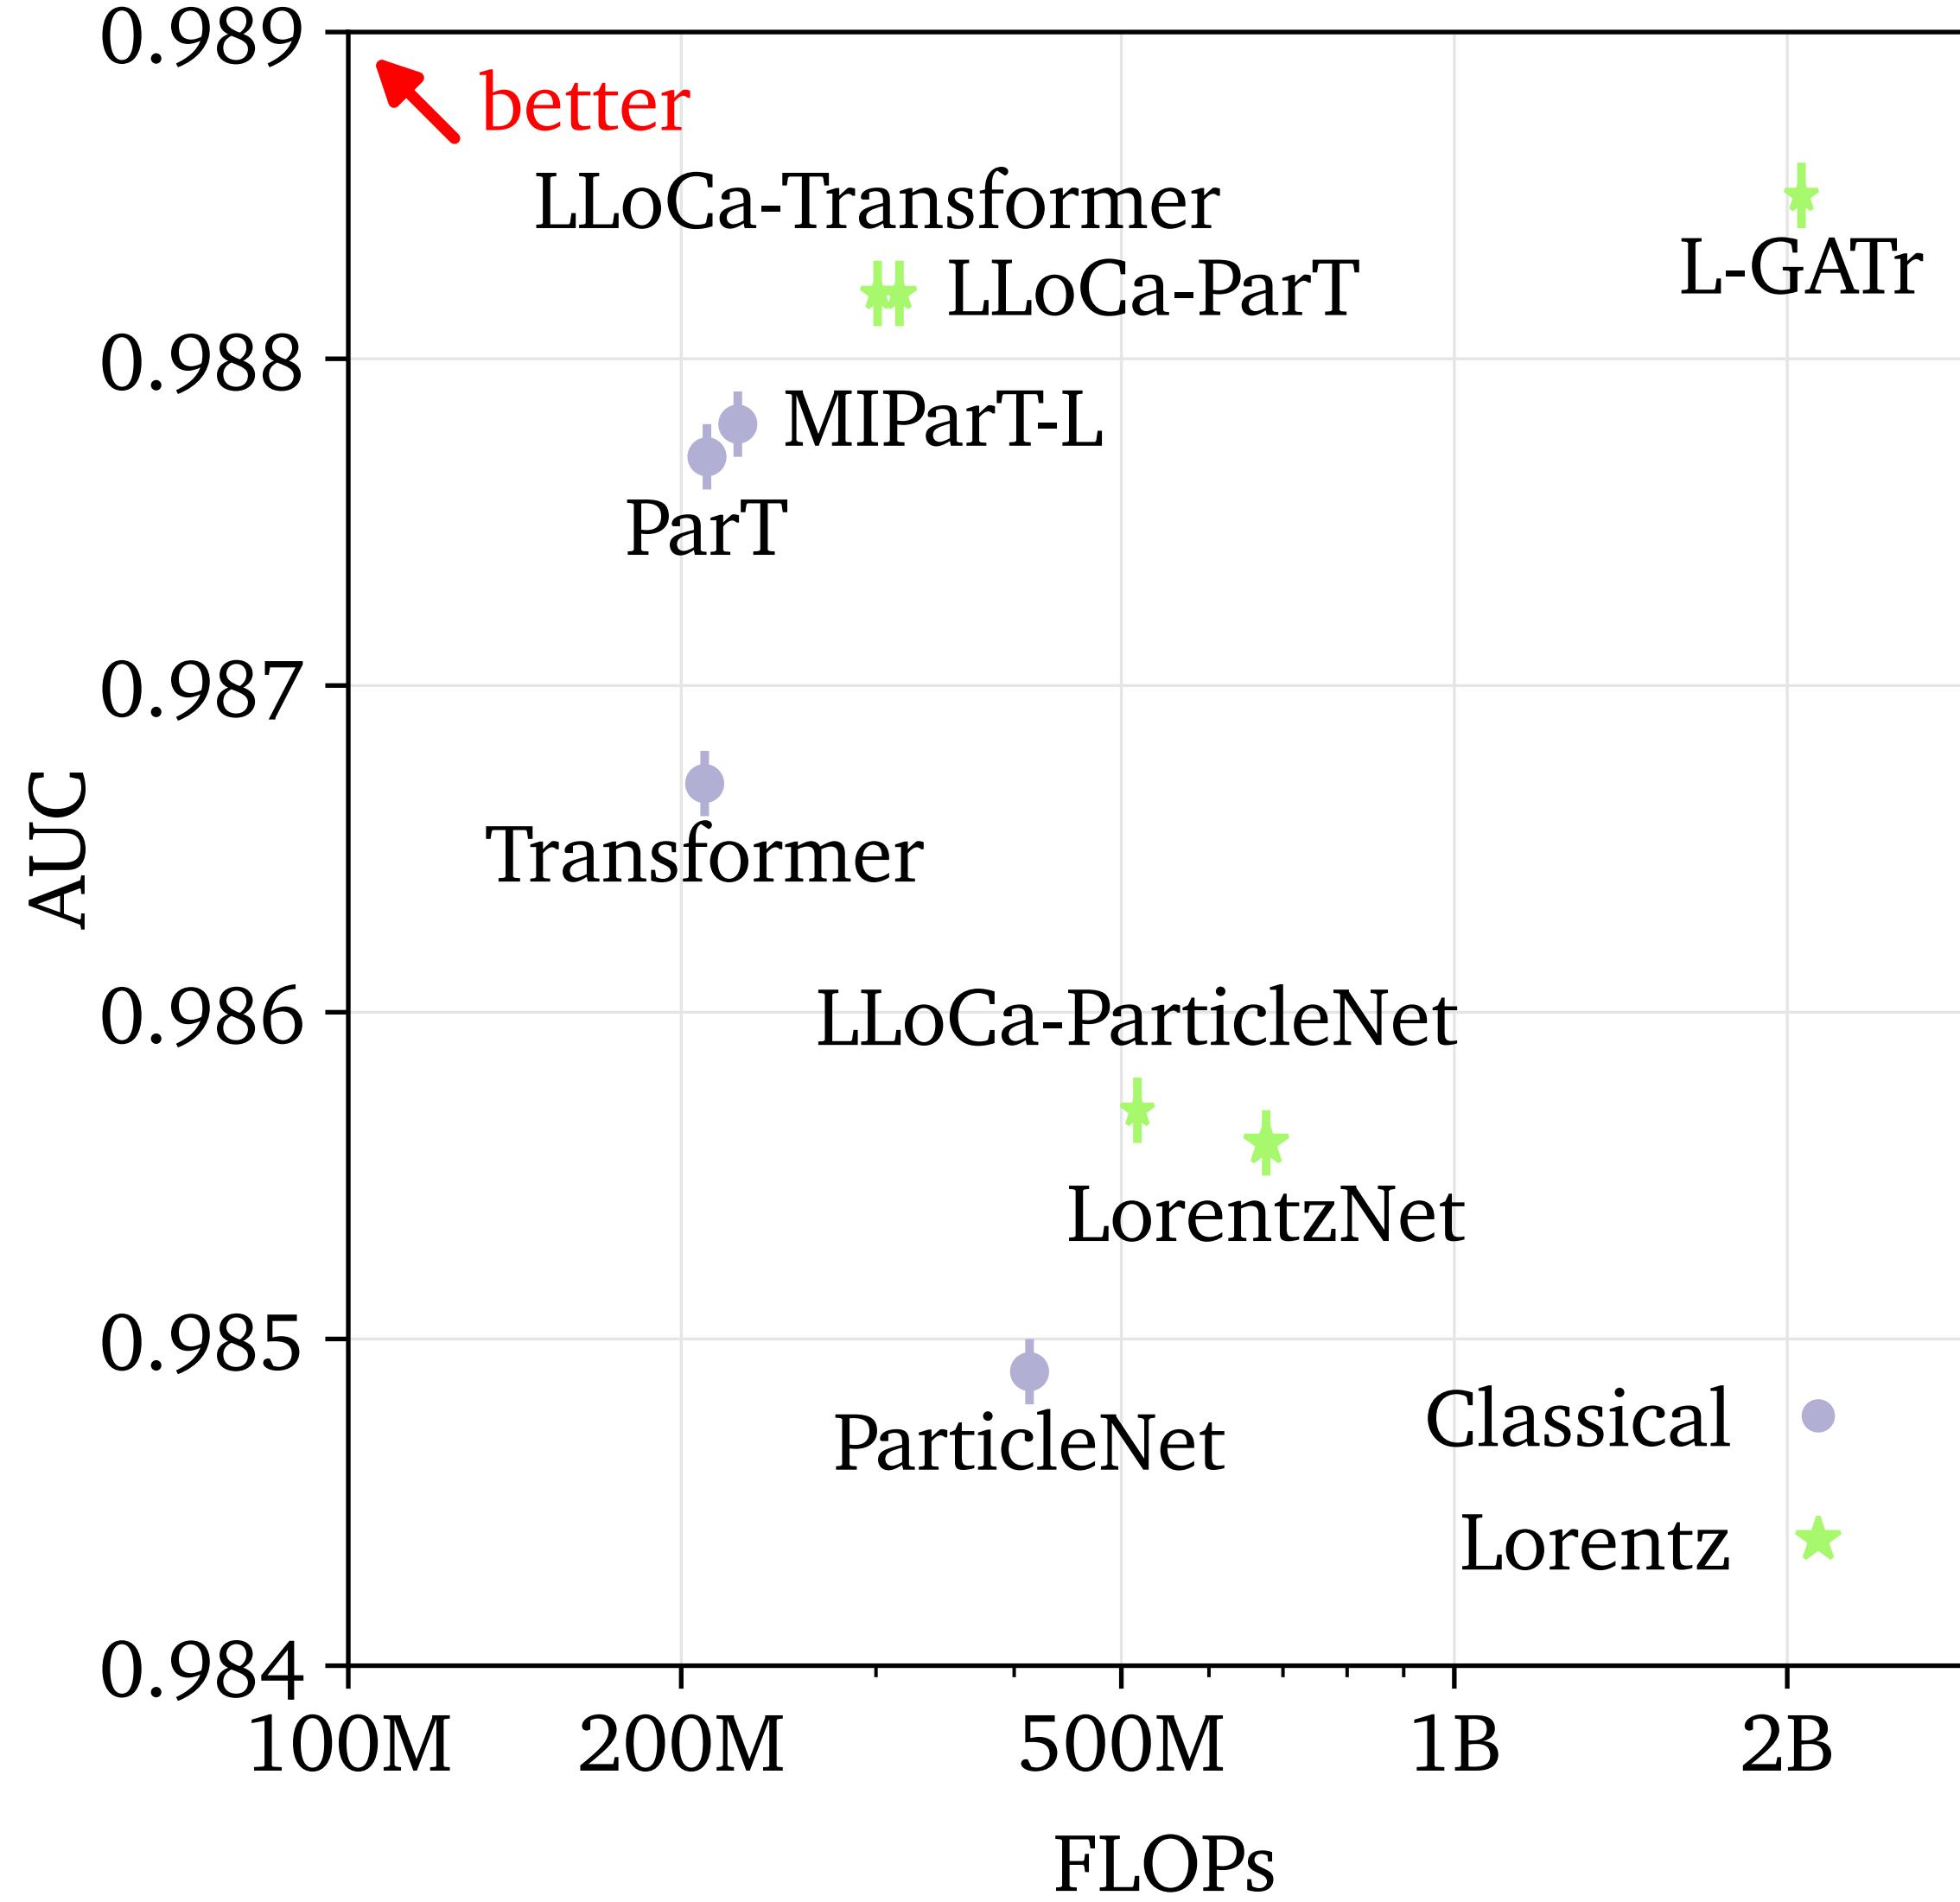


Network	Accuracy	AUC	$1/\epsilon_B$ ( $\epsilon_S = 0.5$ )	$1/\epsilon_B$ ( $\epsilon_S = 0.3$ )
TopoDNN [54]	0.916	0.972	–	$295 \pm 5$
LoLa [9]	0.929	0.980	–	$722 \pm 17$
<i>N</i> -subjettiness [55]	0.929	0.981	–	$867 \pm 15$
PFN [56]	0.932	0.9819	$247 \pm 3$	$888 \pm 17$
TreeNiN [57]	0.933	0.982	–	$1025 \pm 11$
ParticleNet [58]	0.940	0.9858	$397 \pm 7$	$1615 \pm 93$
ParT [59]	0.940	0.9858	$413 \pm 16$	$1602 \pm 81$
MIParT [60]	0.942	0.9868	$505 \pm 8$	$2010 \pm 97$
LorentzNet* [10]	0.942	0.9868	$498 \pm 18$	$2195 \pm 173$
CGENN* [14]	0.942	0.9869	500	2172
PELICAN* [42]	$0.9426 \pm 0.0002$	$0.9870 \pm 0.0001$	–	$2250 \pm 75$
L-GATr* [35]	$0.9423 \pm 0.0002$	$0.9870 \pm 0.0001$	$540 \pm 20$	$2240 \pm 70$
ParticleNet-f.t. [60]	0.942	0.9866	$487 \pm 9$	$1771 \pm 80$
ParT-f.t. [60]	0.944	0.9877	$691 \pm 15$	$2766 \pm 130$
MIParT-f.t. [60]	0.944	0.9878	$640 \pm 10$	$2789 \pm 133$
L-GATr-f.t.* (new)	$0.9442 \pm 0.0002$	$0.98792 \pm 0.00004$	$661 \pm 24$	$3005 \pm 186$

	All classes	$H \rightarrow b\bar{b}$	$H \rightarrow c\bar{c}$	$H \rightarrow gg$	$H \rightarrow 4q$	$H \rightarrow l\nu q\bar{q}'$	$t \rightarrow bq\bar{q}'$	$t \rightarrow bl\nu$	$W \rightarrow q\bar{q}'$	$Z \rightarrow q\bar{q}$	
	Accuracy	AUC	Rej <sub>50%</sub>	Rej <sub>50%</sub>	Rej <sub>50%</sub>	Rej <sub>99%</sub>	Rej <sub>50%</sub>	Rej <sub>99.5%</sub>	Rej <sub>50%</sub>	Rej <sub>50%</sub>	
ParticleNet [58]	0.844	0.9849	7634	2475	104	954	3339	10526	11173	347	283
ParT [59]	0.861	0.9877	10638	4149	123	1864	5479	32787	15873	543	402
MIParT [60]	0.861	0.9878	10753	4202	123	1927	5450	31250	16807	542	402
L-GATr	0.865	0.9884	12195	4819	128	2304	5764	37736	19231	580	427

Model	Accuracy ( $\uparrow$ )	AUC ( $\uparrow$ )	Time	FLOPs
PFN [25]	0.772	0.9714	3h	3M
P-CNN [38]	0.809	0.9789	3h	12M
MIParT [45]	0.861	0.9878	—	—
LorentzNet* [19]	0.847	0.9856	57h	676M
L-GATr* [9]	<b>0.866</b>	<b>0.9885</b>	180h	2060M
ParticleNet [34]	0.844	0.9849	27h	413M
LLoCa-ParticleNet*	<u>0.848</u>	<u>0.9857</u>	41h	517M
ParT [35]	0.861	0.9877	38h	211M
LLoCa-ParT*	<u>0.864</u>	<u>0.9882</u>	67h	315M
Transformer	0.855	0.9867	17h	210M
LLoCa-Transformer*	<u>0.864</u>	<u>0.9882</u>	33h	301M

Beam	Time	Embedding	AUC	$1/\epsilon_B$ ( $\epsilon_S = 0.3$ )
–	–	Token	$\times$	$0.9846 \pm 0.0002$ $1494 \pm 62$
$x_3^V = \pm 1$	–	Token	$\times$	$0.9854 \pm 0.0005$ $1684 \pm 221$
$x_{12}^B = x_{13}^B = x_{23}^B = 1$	$x_0^V = 1$	Token	$\times$	$0.9864 \pm 0.0002$ $2076 \pm 75$
–	$x_0^V = 1$	Token	$\times$	$0.9865 \pm 0.0001$ $2179 \pm 123$
$x_{12}^B = 1$	$x_0^V = 1$	Channel	$\times$	$0.9866 \pm 0.0001$ $2150 \pm 82$
$x_0^V = 1, x_3^V = \pm 1$	$x_0^V = 1$	Token	$\times$	$0.9869 \pm 0.0001$ $2282 \pm 174$
$x_3^V = \pm 1$	$x_0^V = 1$	Token	$\times$	$0.9869 \pm 0.0001$ $2293 \pm 141$
–	–	Token	$\checkmark$	$0.9869 \pm 0.0001$ $2179 \pm 156$
$x_{12}^B = 1$	$x_0^V = 1$	Token	$\checkmark$	$0.9870 \pm 0.0001$ $2173 \pm 83$
$x_{12}^B = 1$	$x_0^V = 1$	Token	$\times$	$0.9870 \pm 0.0001$ $2240 \pm 70$

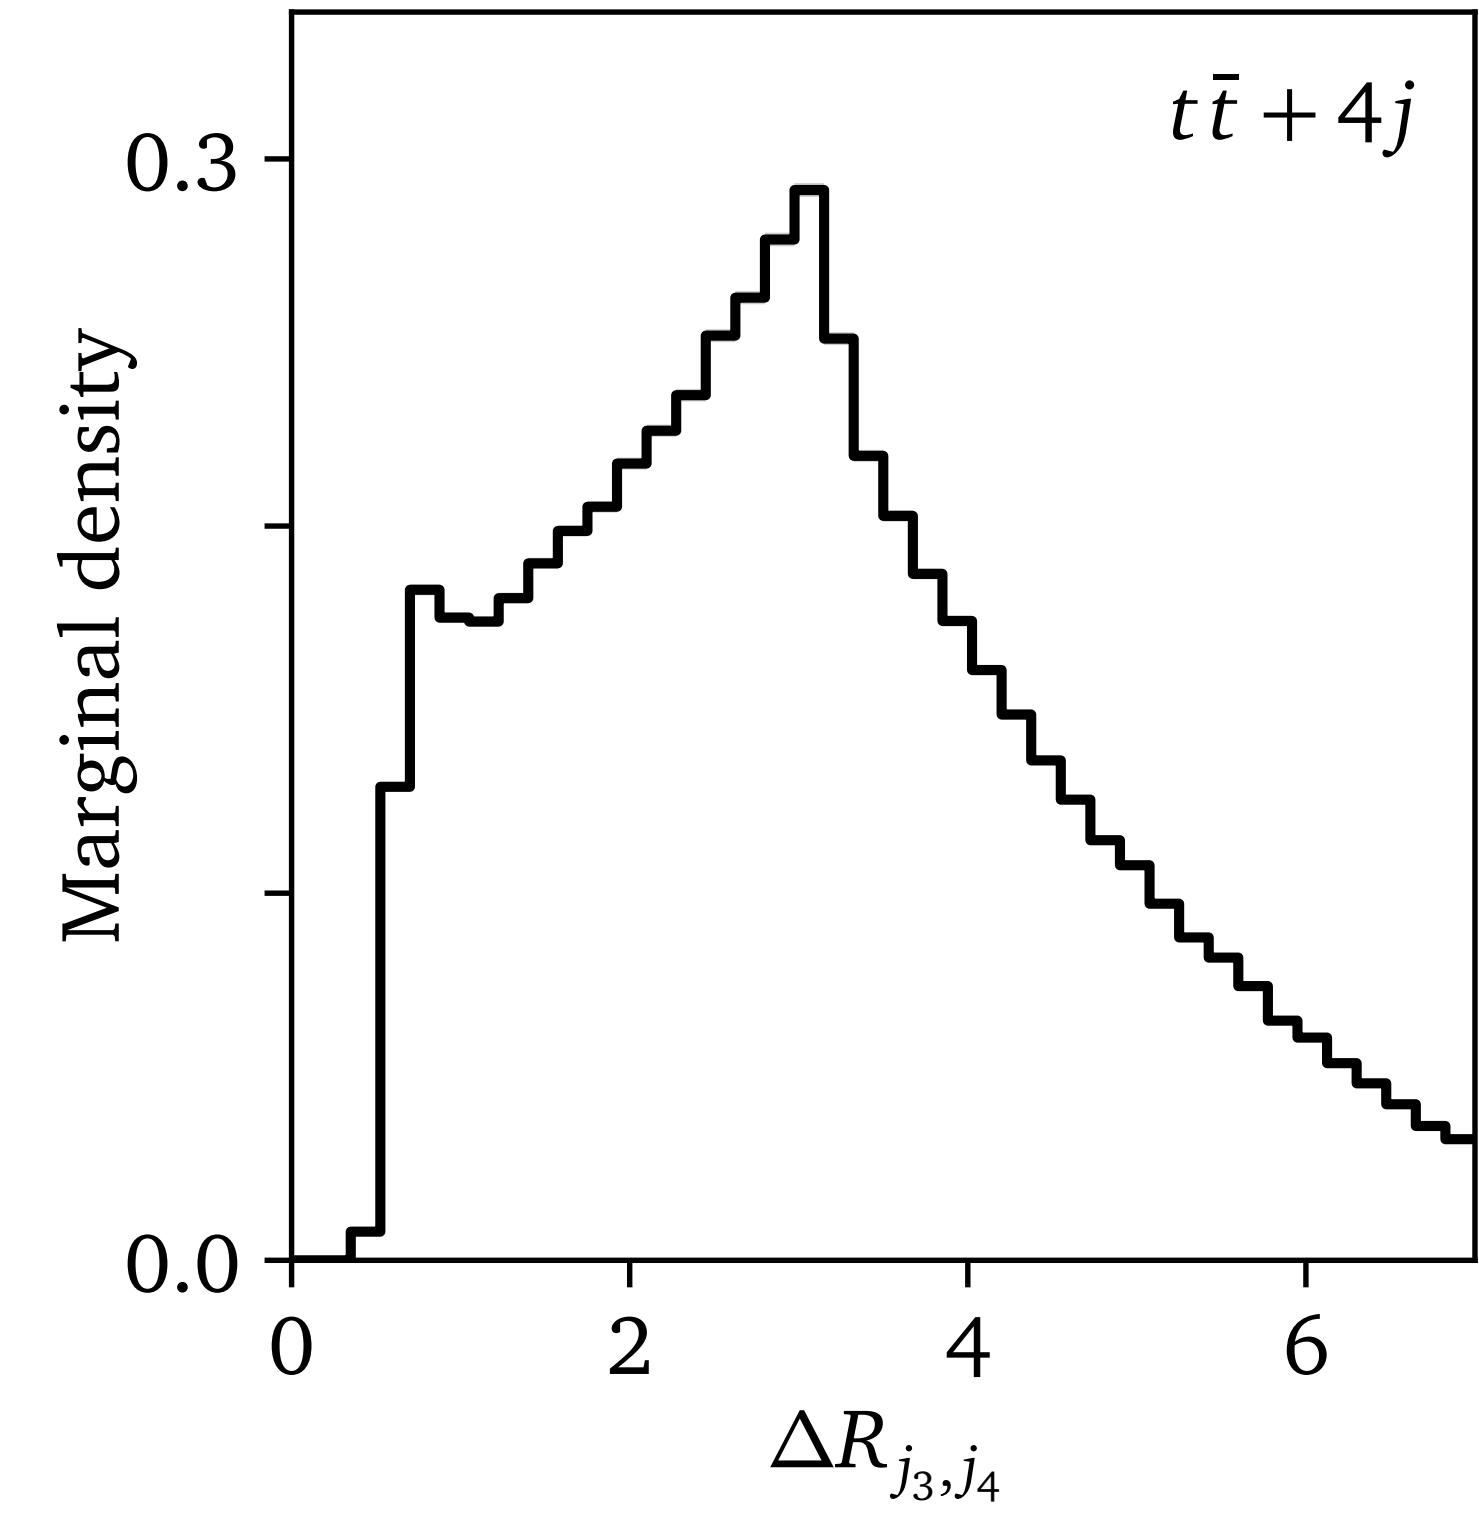
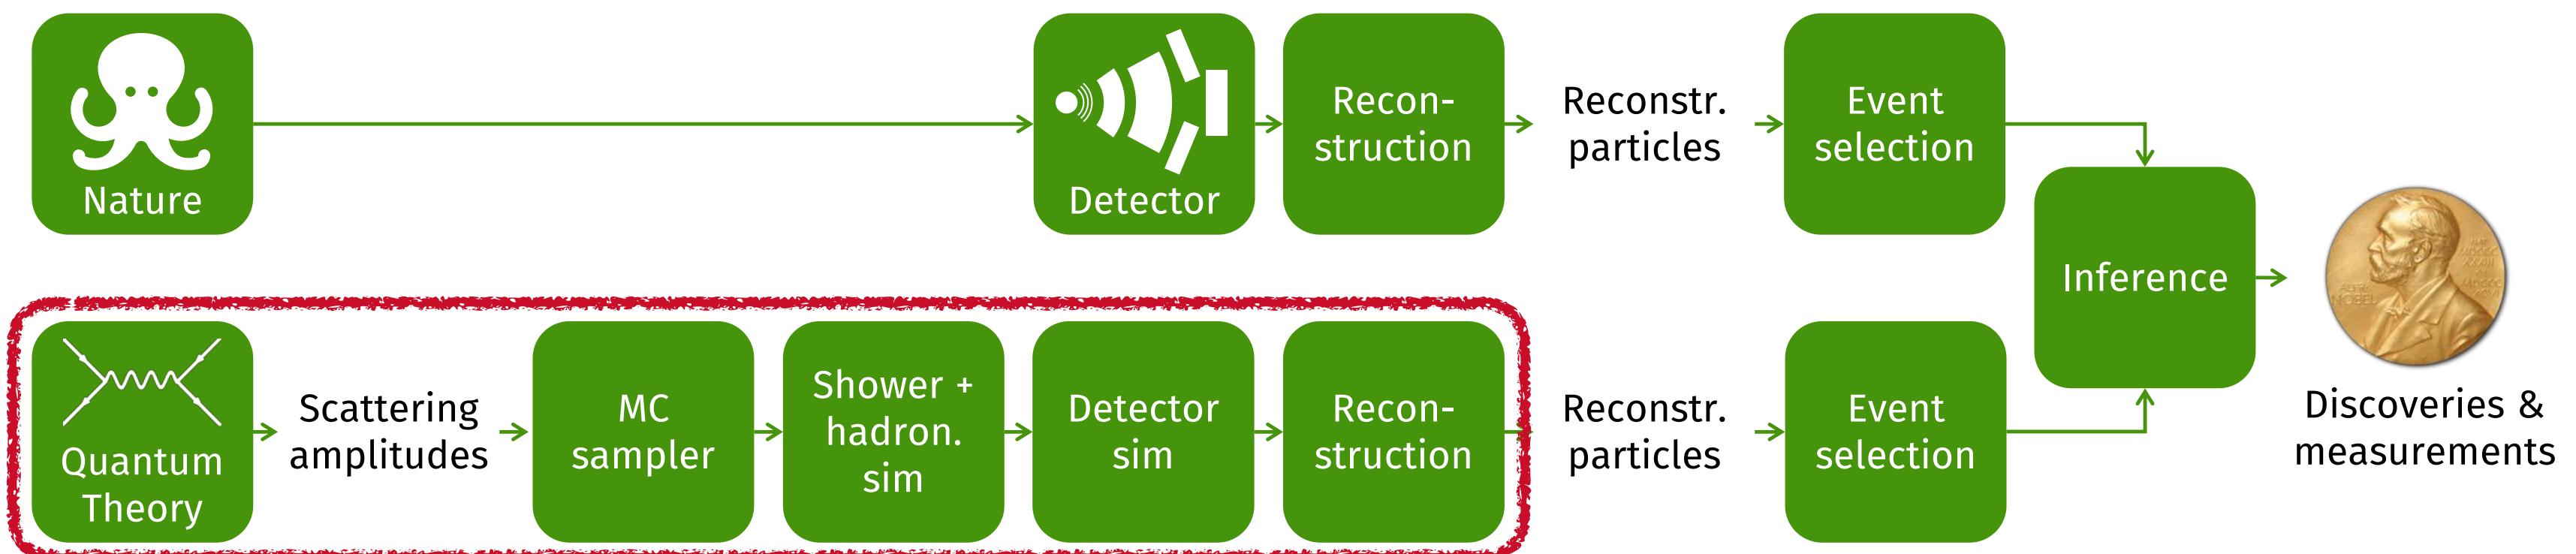


# Event generation

## Task

Dataset:  $pp \rightarrow t_h\bar{t}_h + nj, n = 0\dots4$

- MadGraph + Pythia + Delphes + Reconstruction
- Challenging features:  $m_t, m_W, \Delta R_{jj} > 0.5$
- Symmetry breaking with reference multivectors required at reconstruction level



# Event generation

## Conditional Flow Matching

Continuous normalising flows (CNFs)  
connect a simple base density  
to a complex target density  
through a neural differential equation

$$\frac{d}{dt}x = v_t(x)$$

arXiv:1806.07366

Conditional flow matching (CFM)  
is a simple way to train CNFs  
by comparing the learned velocity  $v_t(x)$   
to a conditional target velocity  $u_t(x | x_1)$

$$\mathcal{L} = \left\langle (v_t(x) - u_t(x | x_1))^2 \right\rangle$$

arXiv:2210.02747

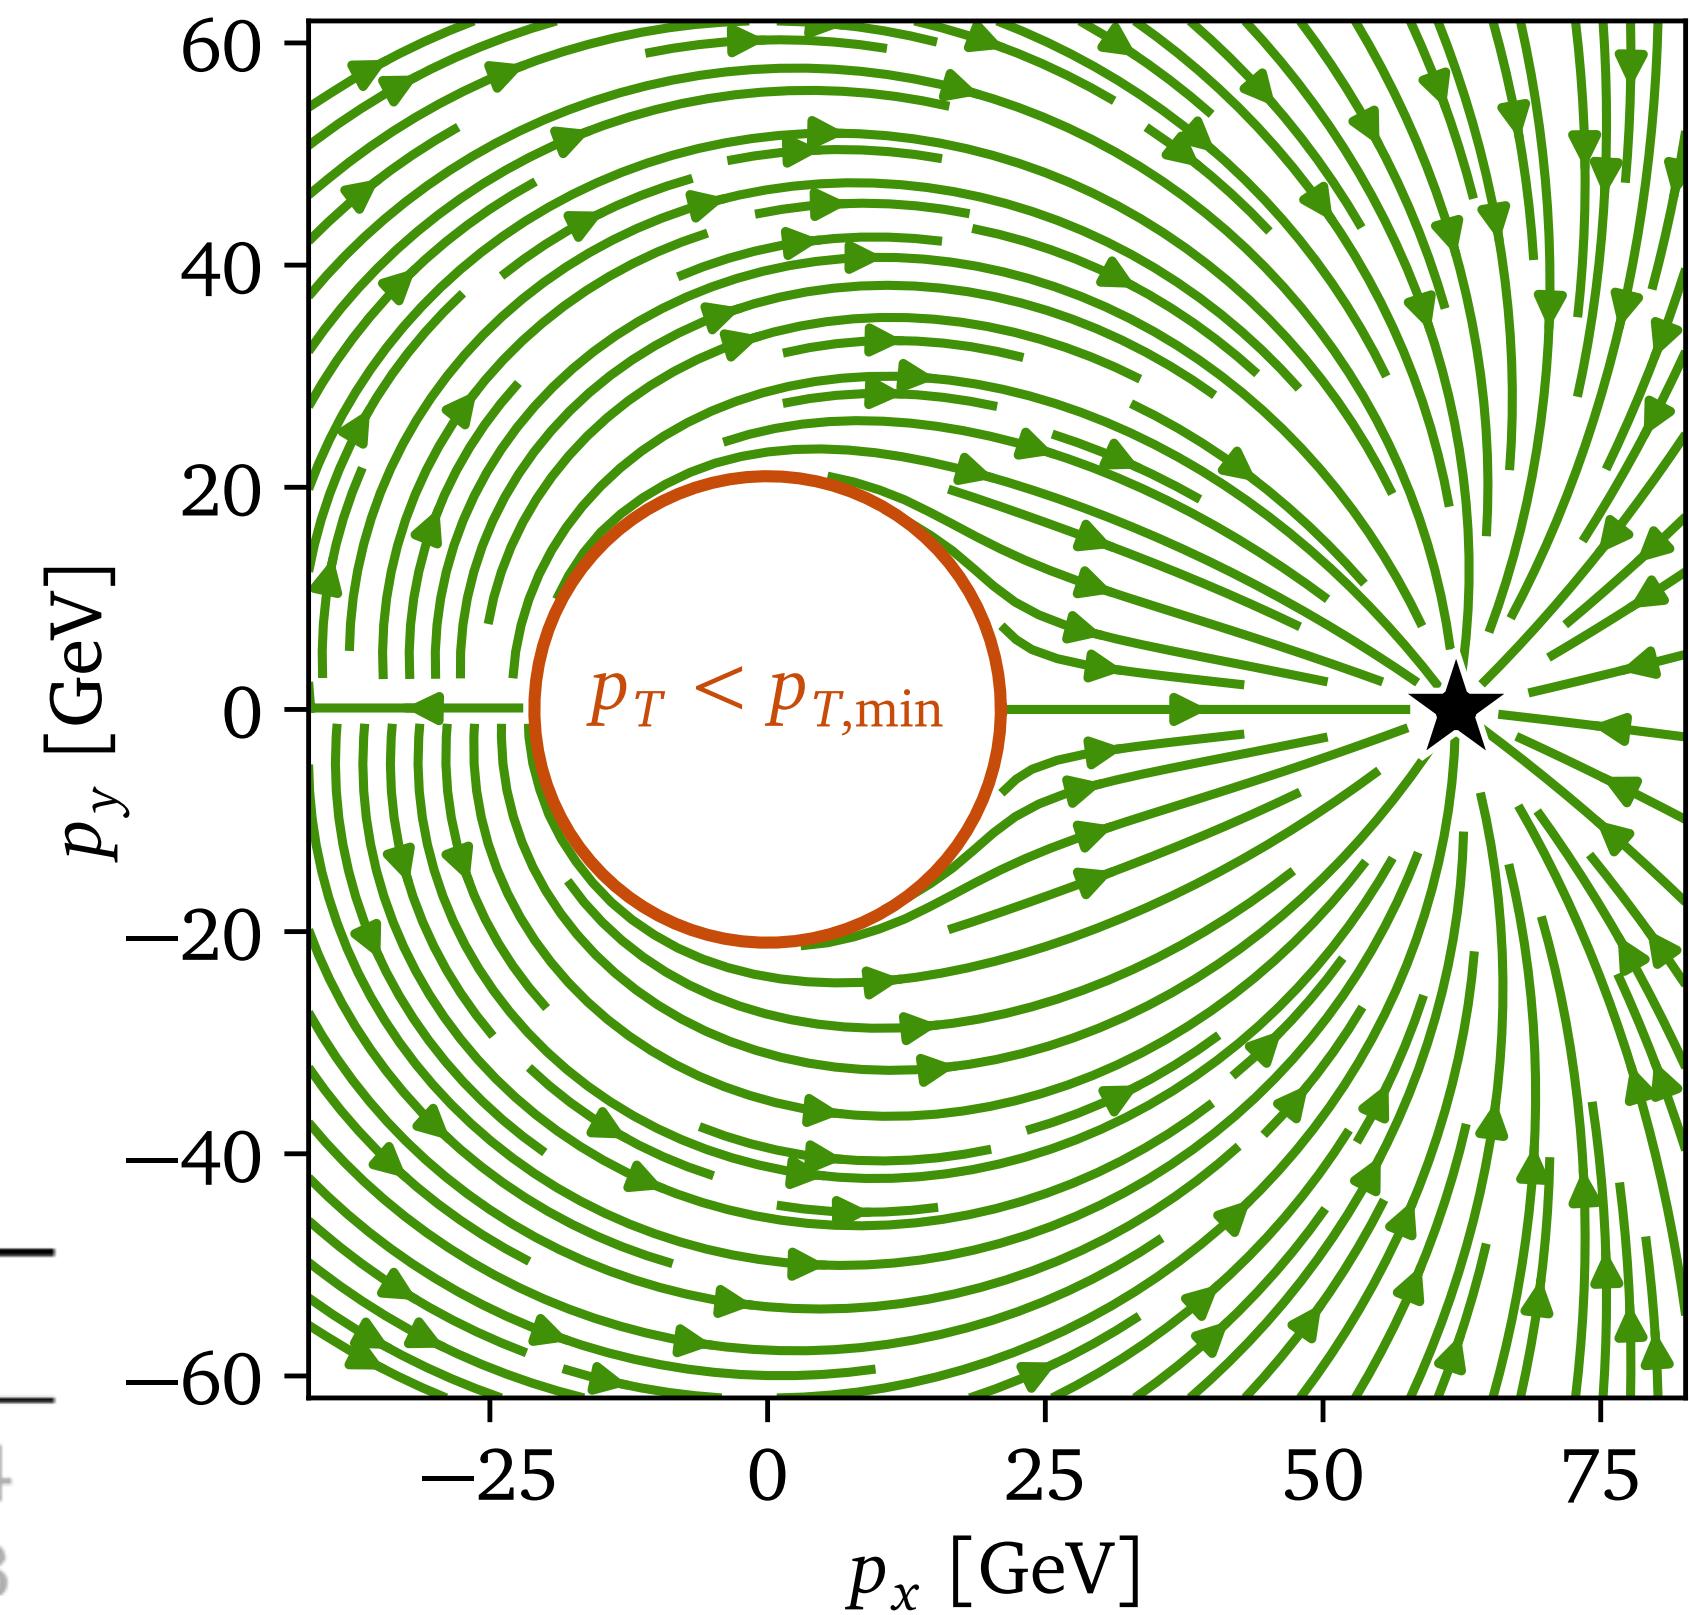
How to pick the target velocity  $u_t(x | x_1)$ ?

# Event generation

## Physics-inspired target trajectories

Straight trajectories in ‘modified jet momenta’  $x$ :

$$p = \begin{pmatrix} E \\ p_x \\ p_y \\ p_z \end{pmatrix} \rightarrow f^{-1}(p) = x = \begin{pmatrix} x_p \\ x_m \\ x_\eta \\ x_\phi \end{pmatrix} = \begin{pmatrix} \log(p_T - p_{T,\min}) \\ \log m^2 \\ \eta \\ \phi \end{pmatrix}$$



Data	Architecture	Base distribution	Periodic	Neg. log-likelihood	AUC
$p$	L-GATr	rejection sampling	✓	$-30.80 \pm 0.17$	$0.945 \pm 0.004$
$x$	MLP	rejection sampling	✓	$-32.13 \pm 0.05$	$0.780 \pm 0.003$
$x$	L-GATr	rejection sampling	✗	$-32.57 \pm 0.05$	$0.530 \pm 0.017$
$x$	L-GATr	no rejection sampling	✓	$-32.58 \pm 0.04$	$0.523 \pm 0.014$
$x$	L-GATr	rejection sampling	✓	$-32.65 \pm 0.04$	$0.515 \pm 0.009$

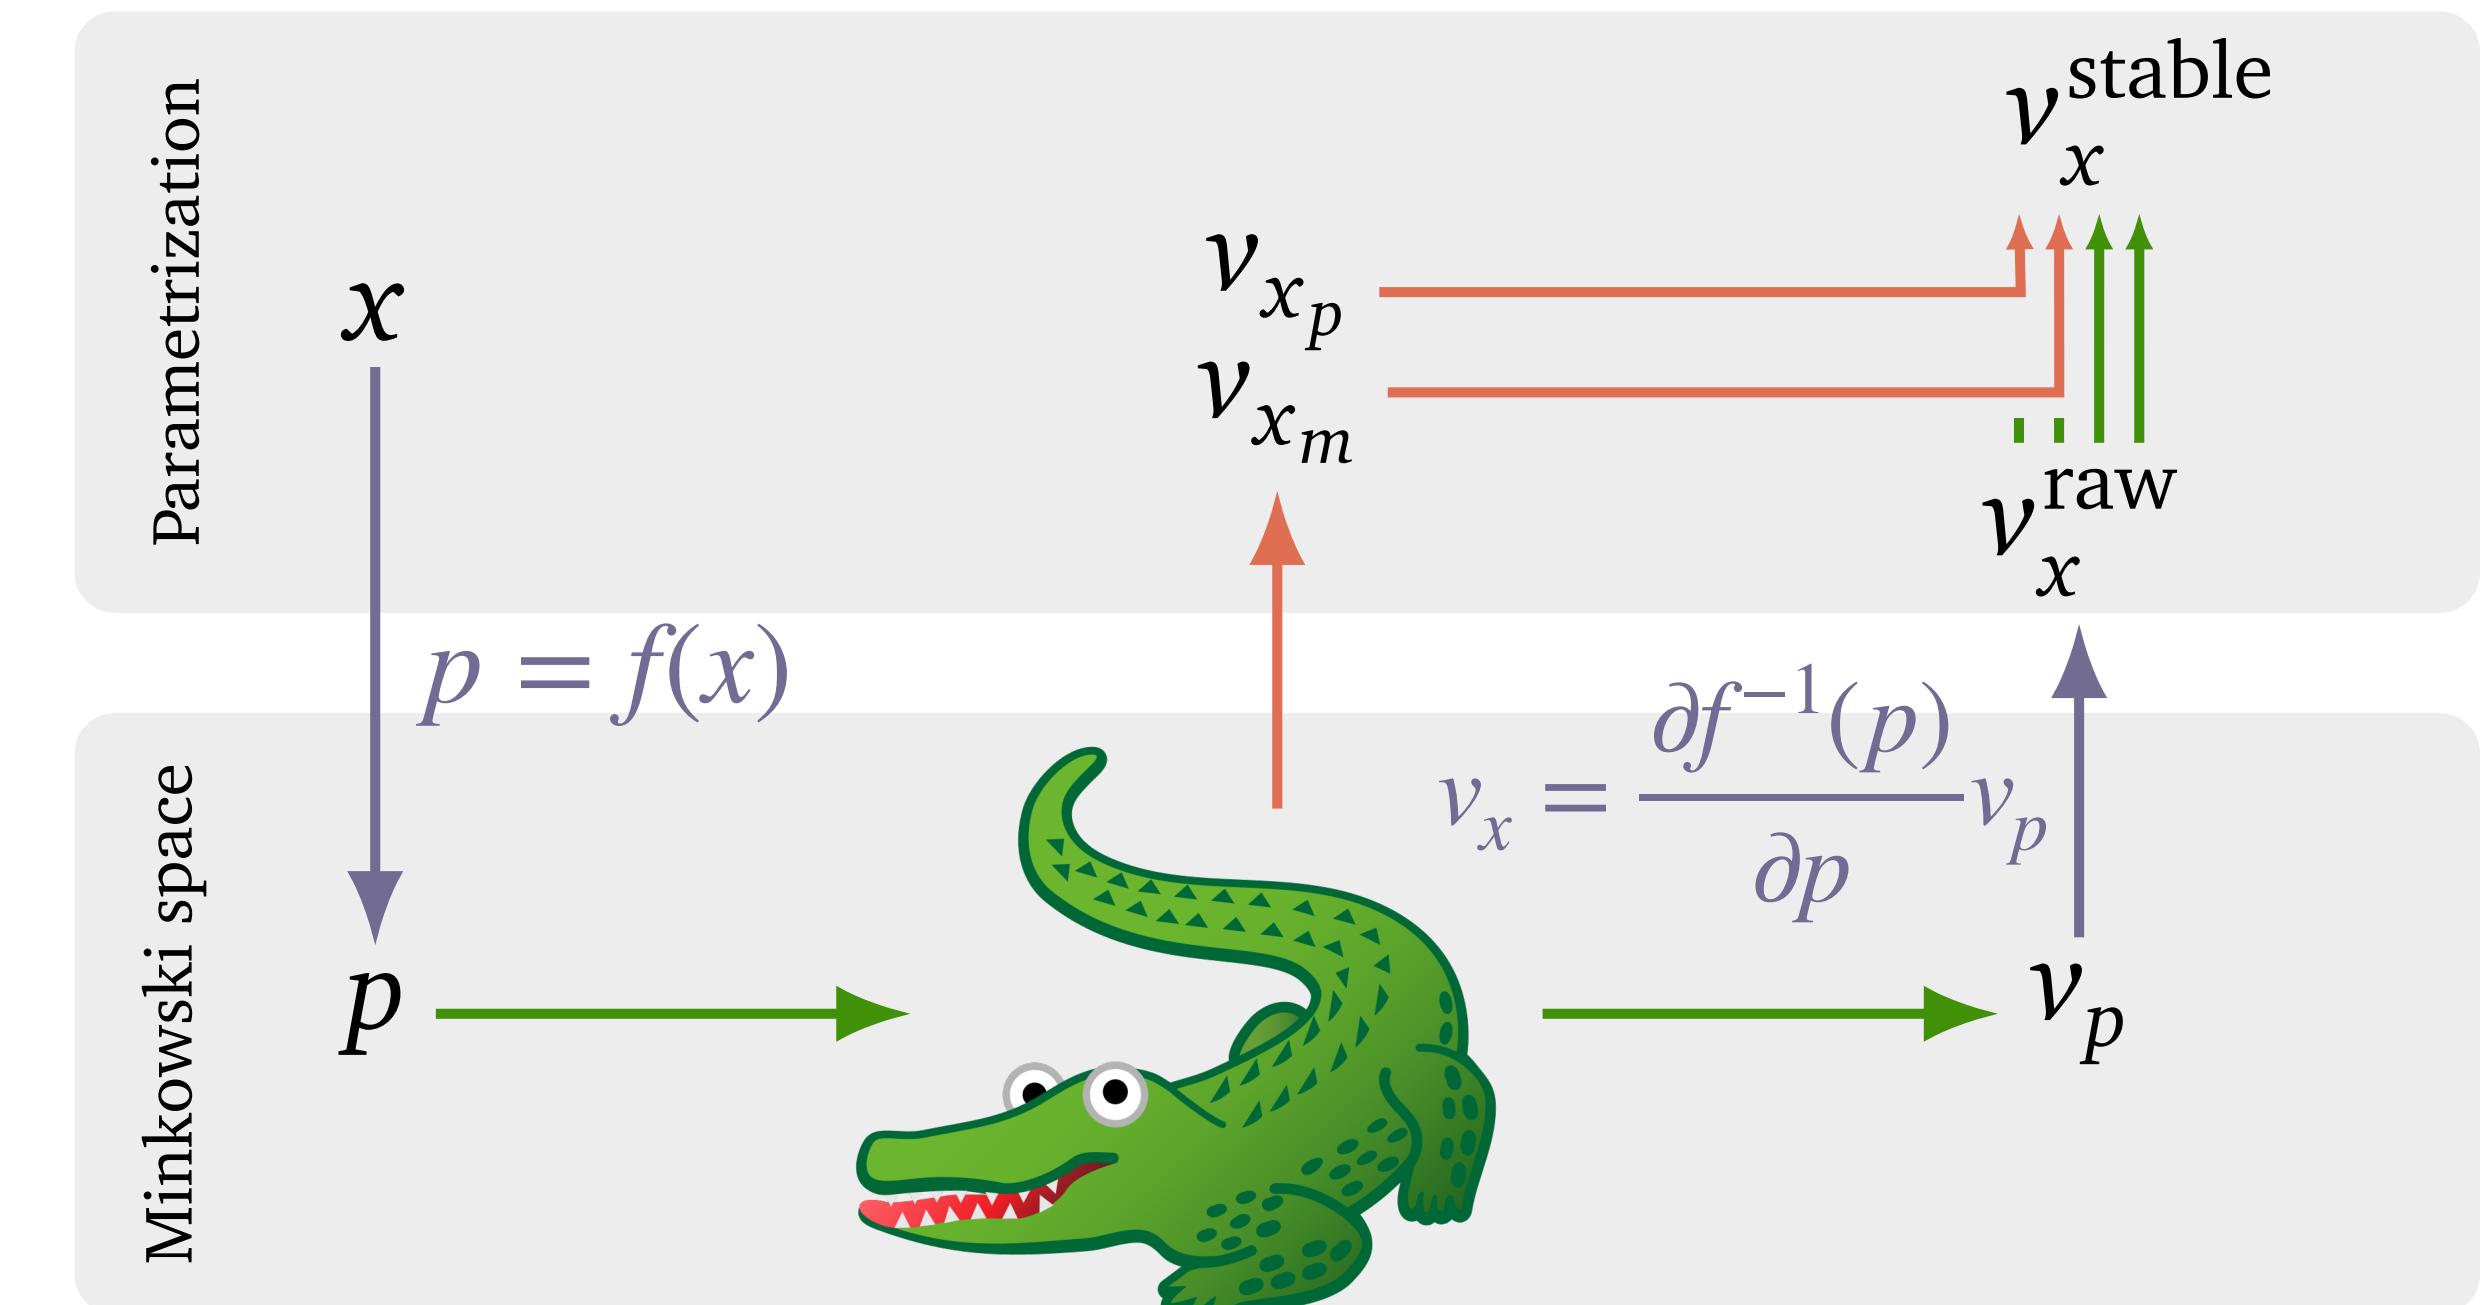
→ default

# Event generation

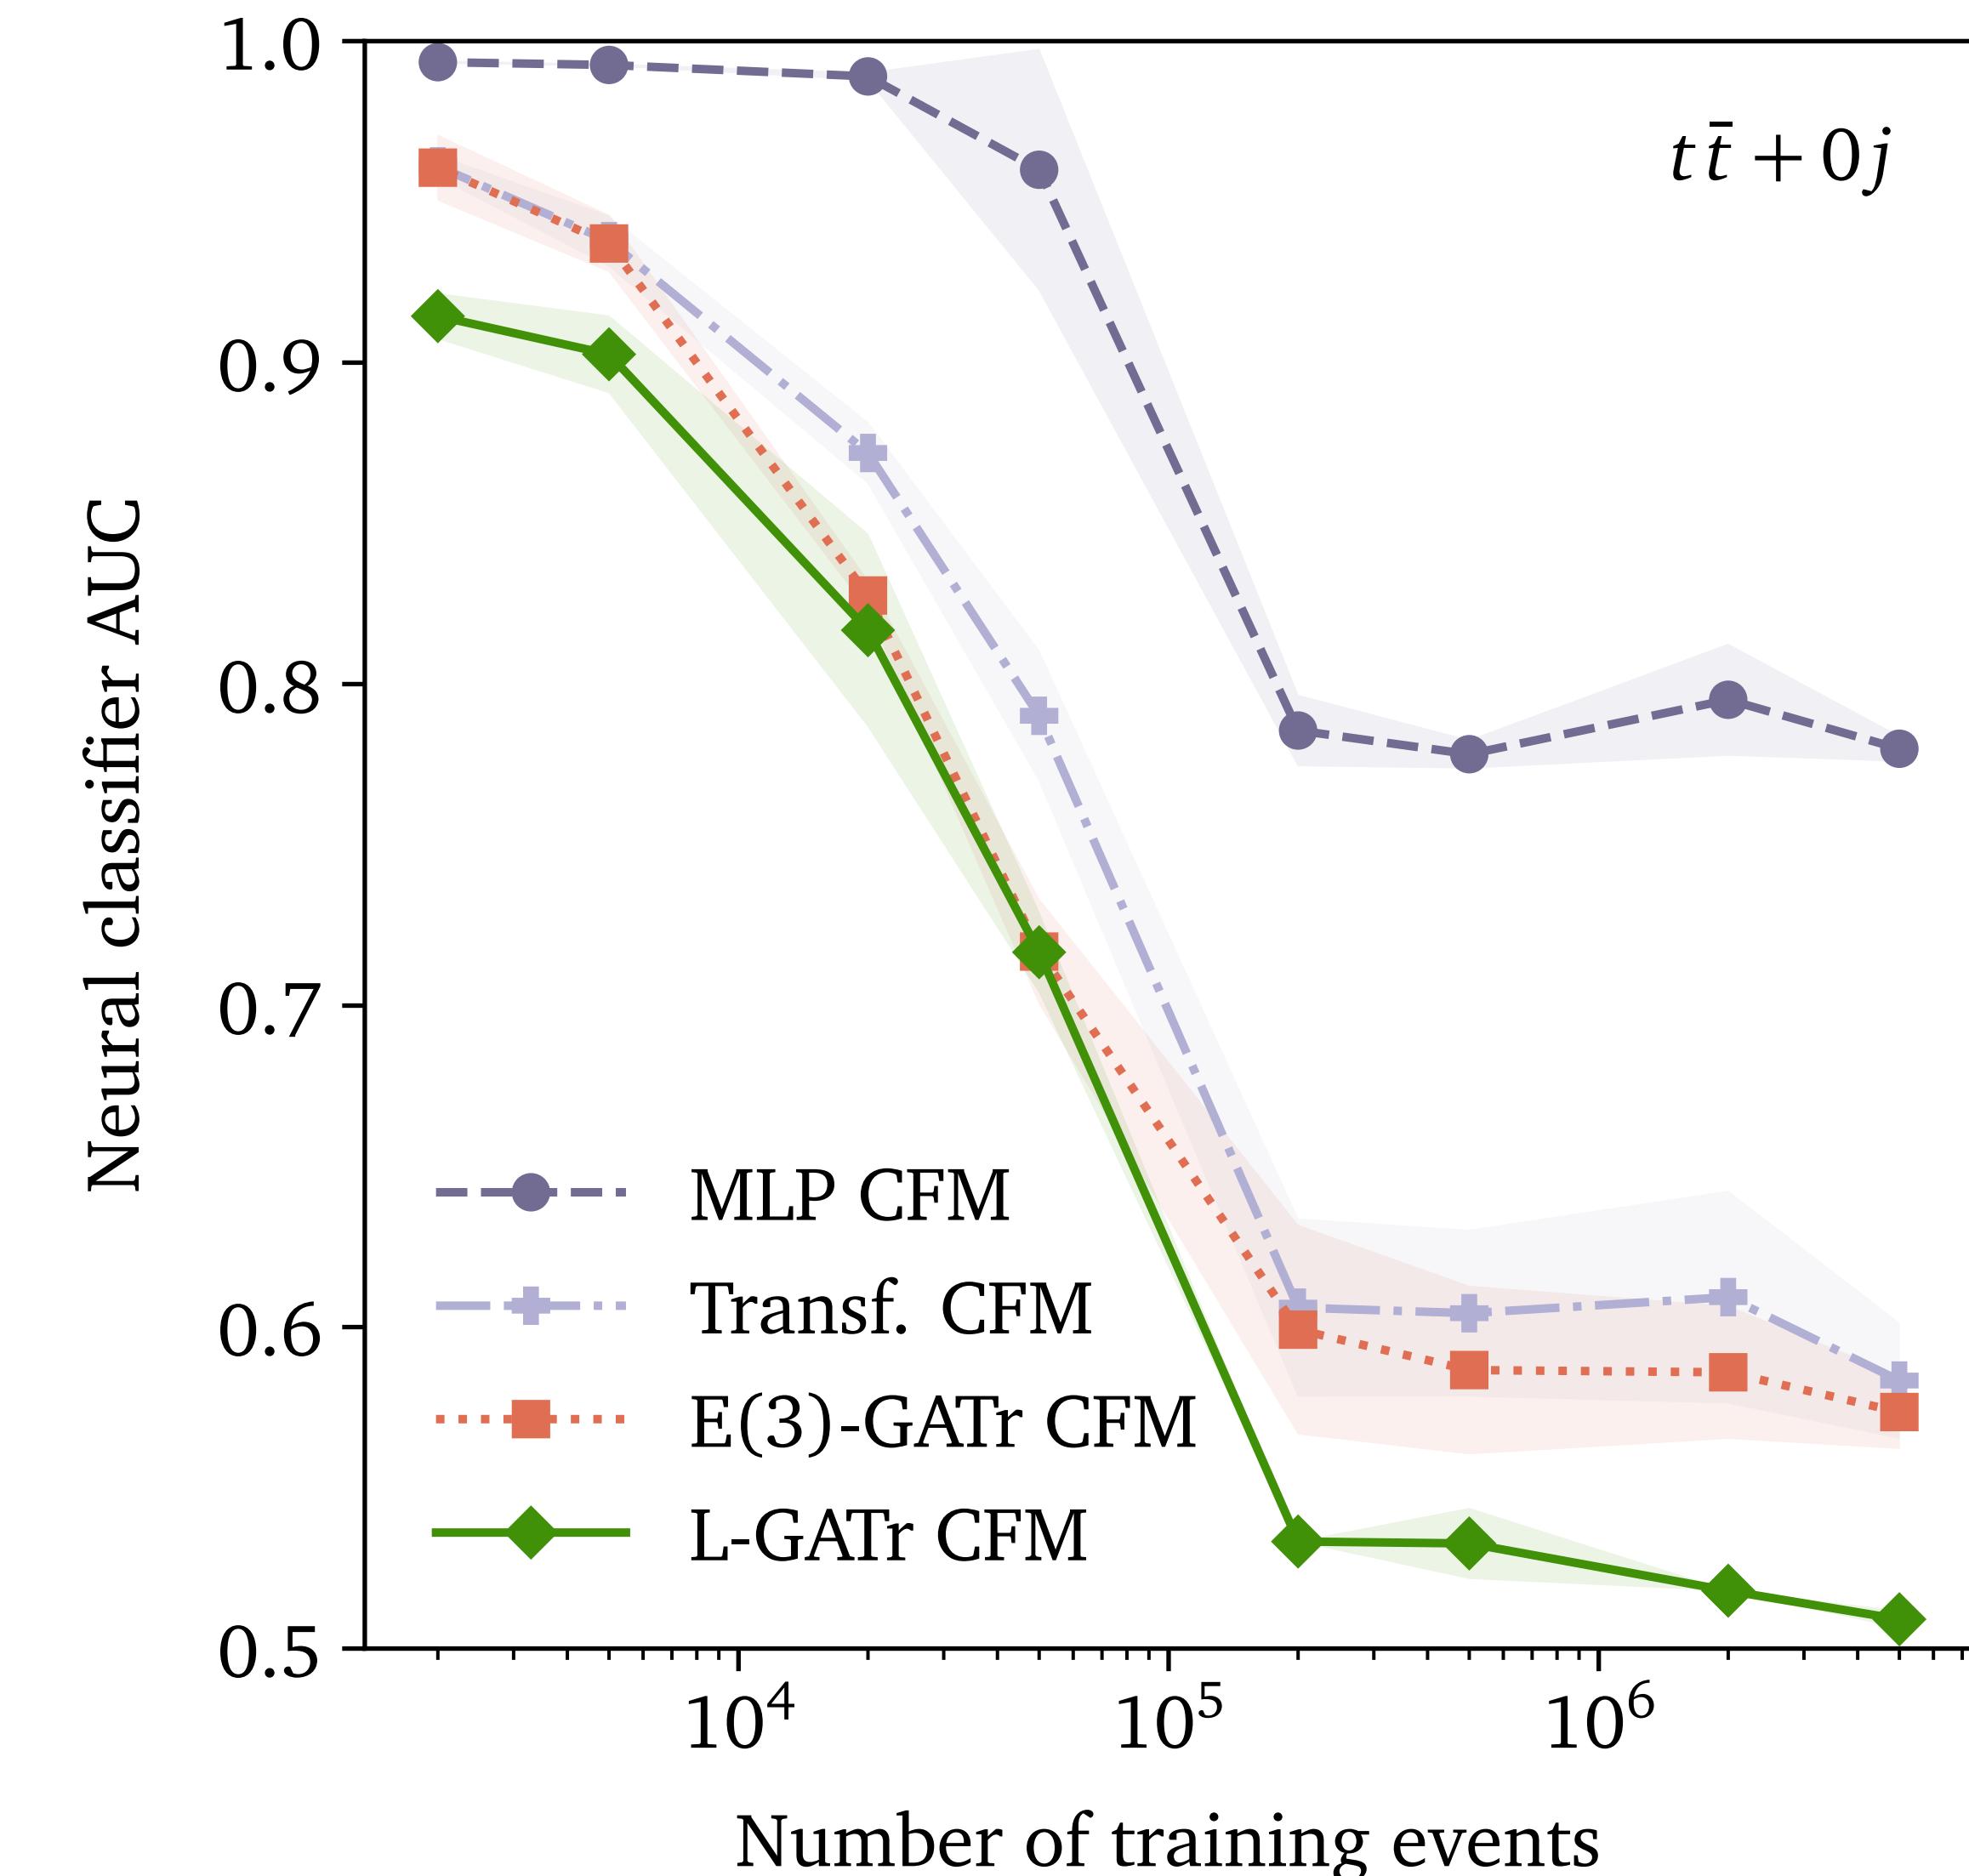
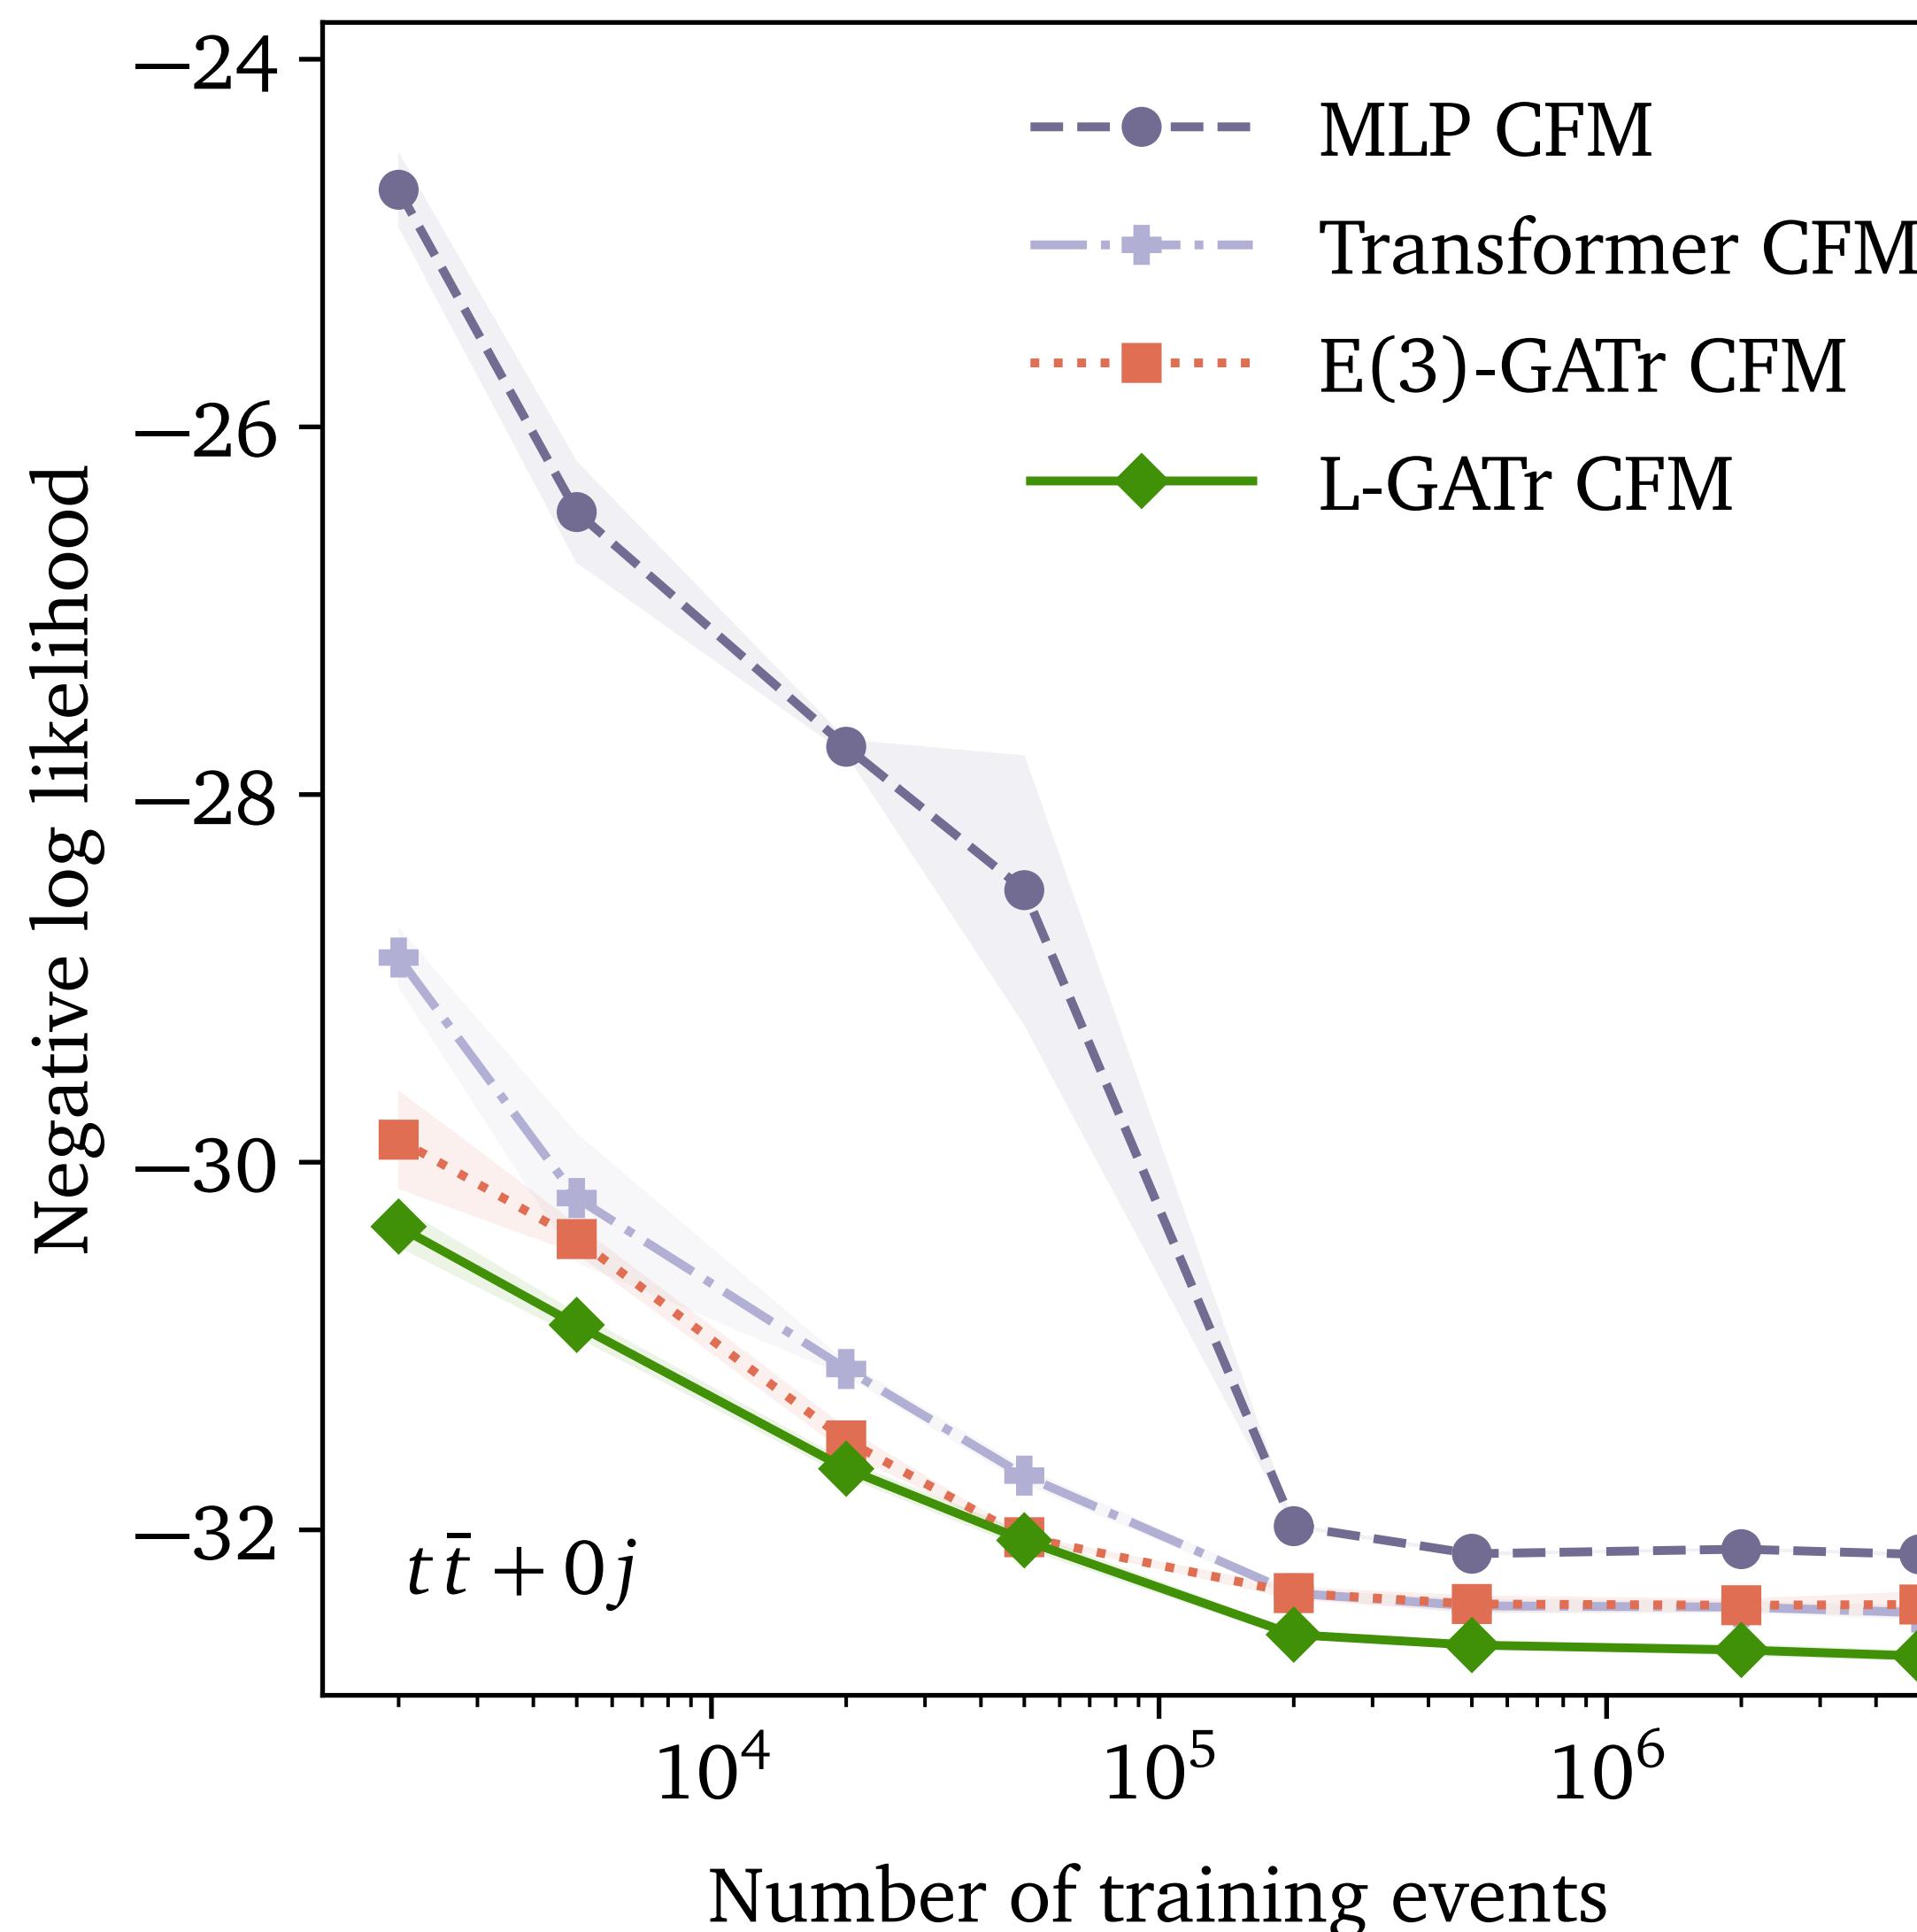
How to build an equivariant CFM field  $v_x(x)$ ?

Extend standard CFM workflow  
with L-GATr:

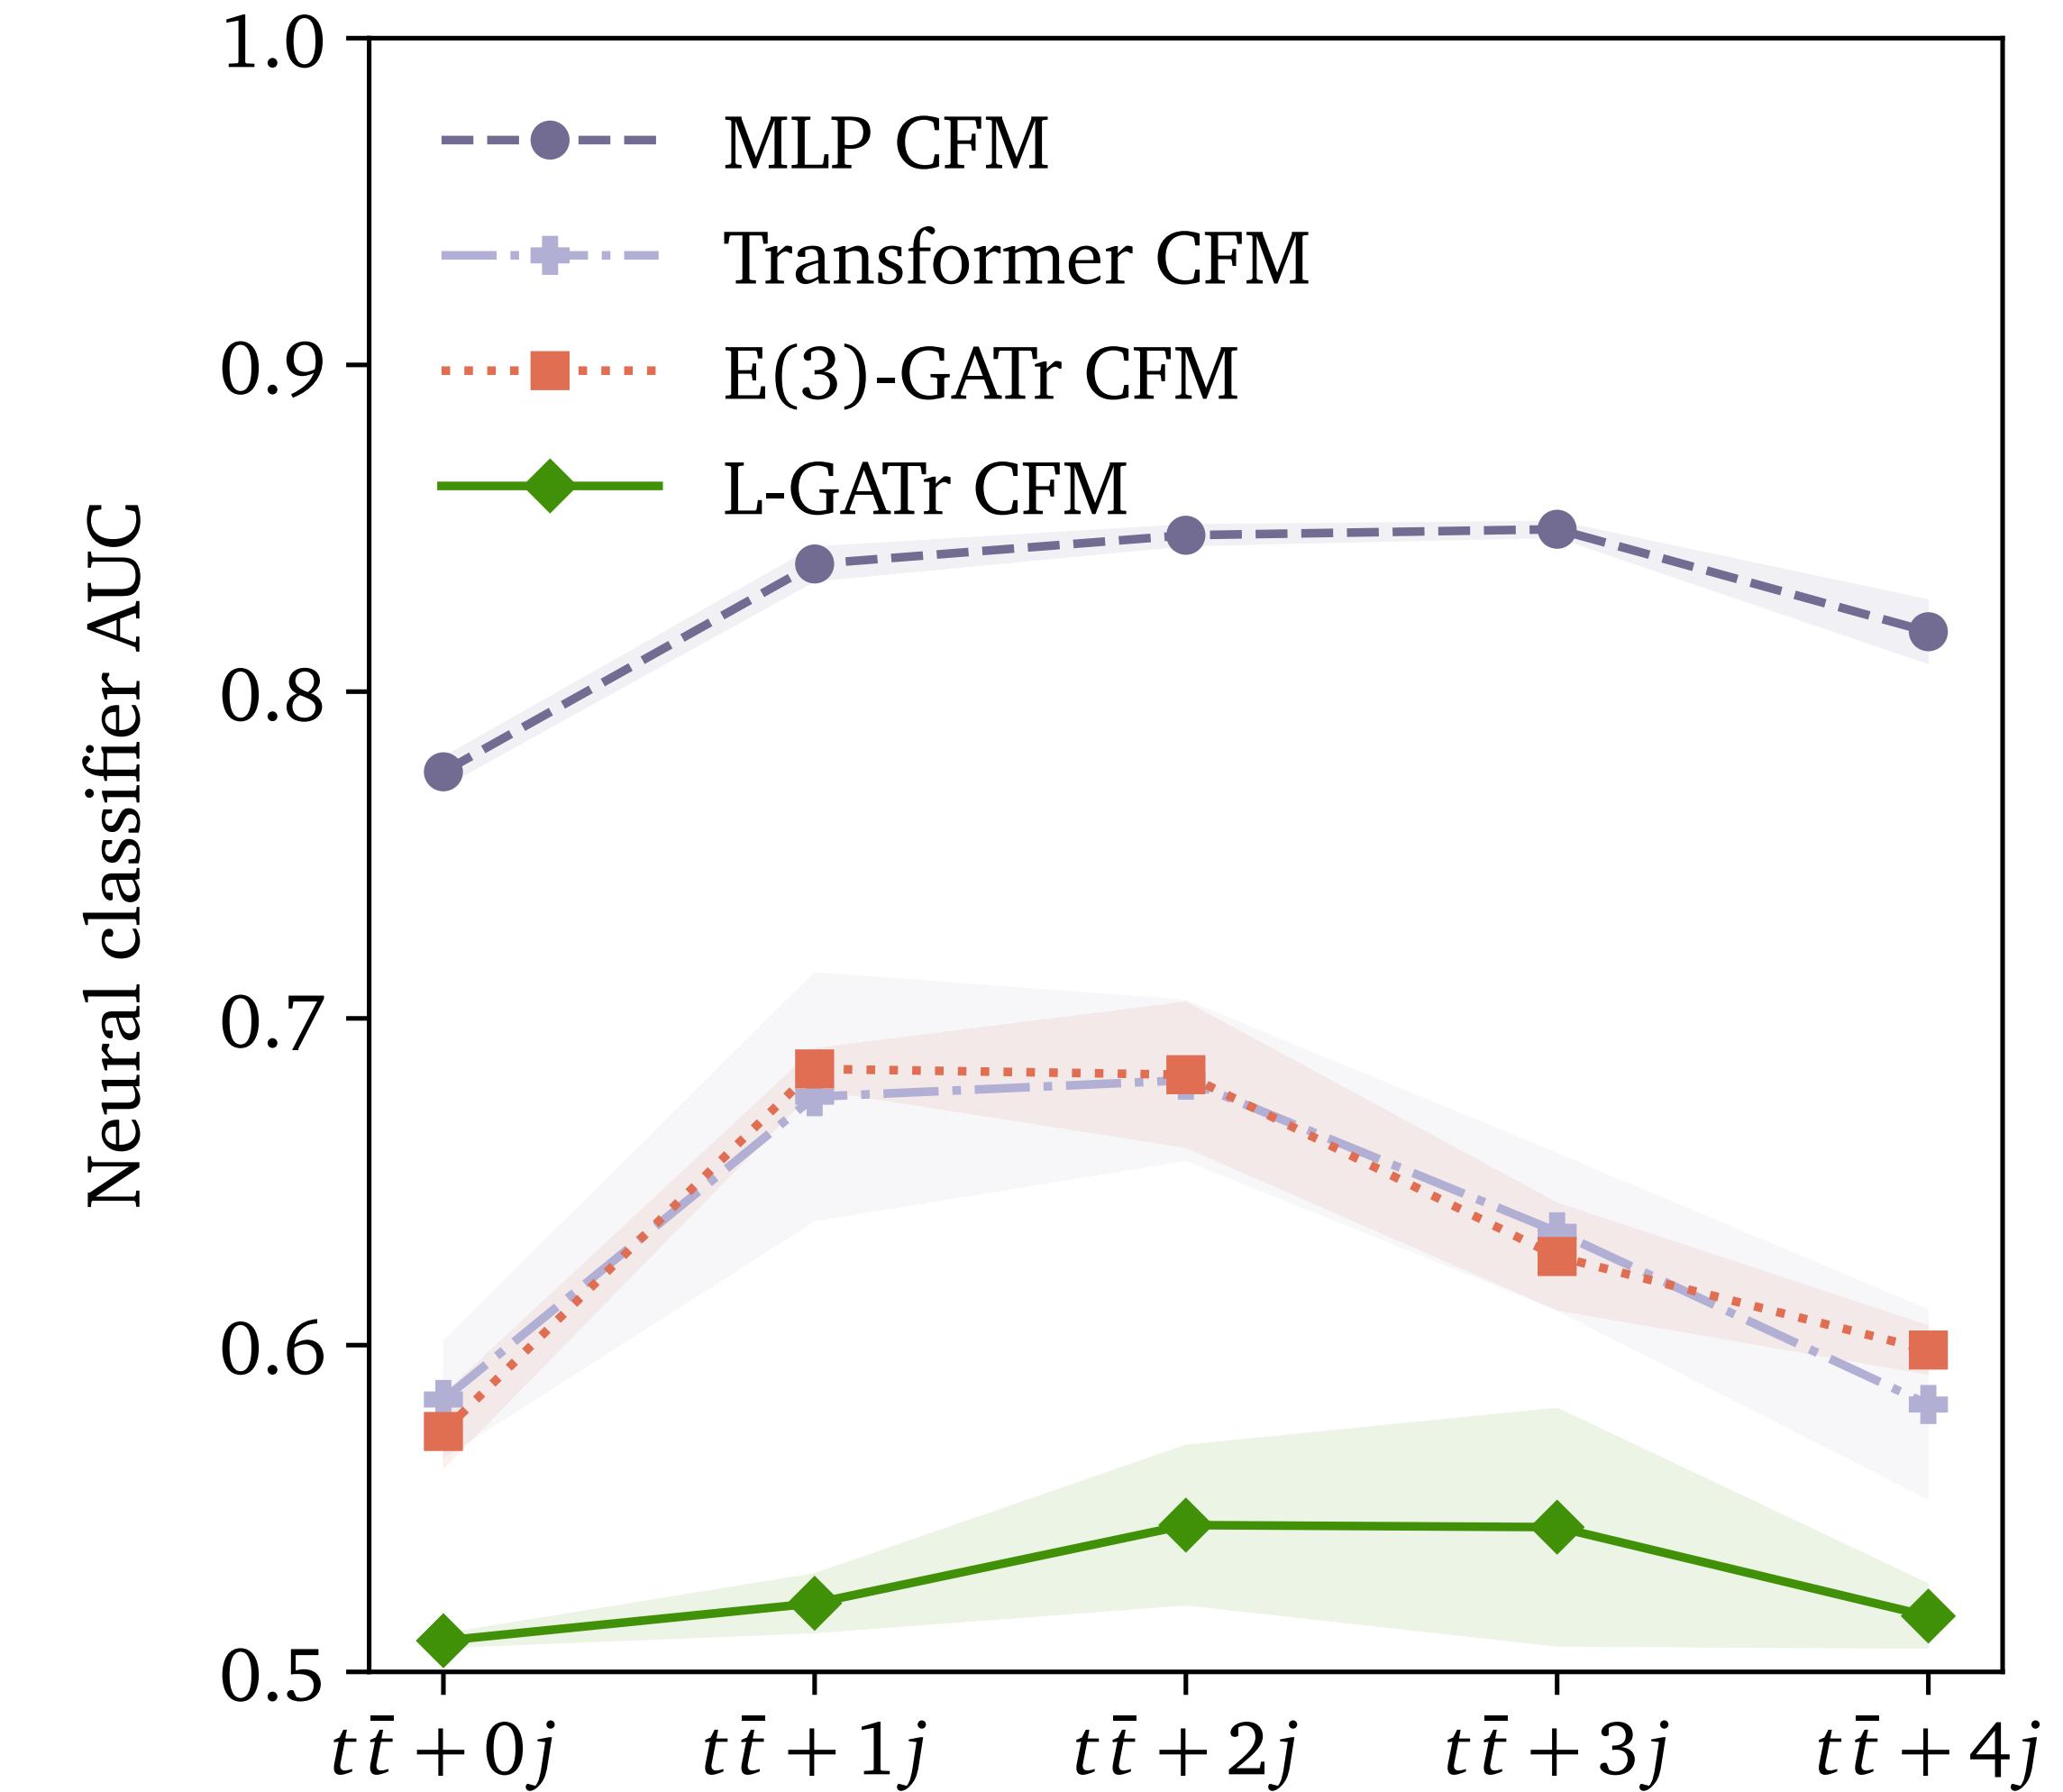
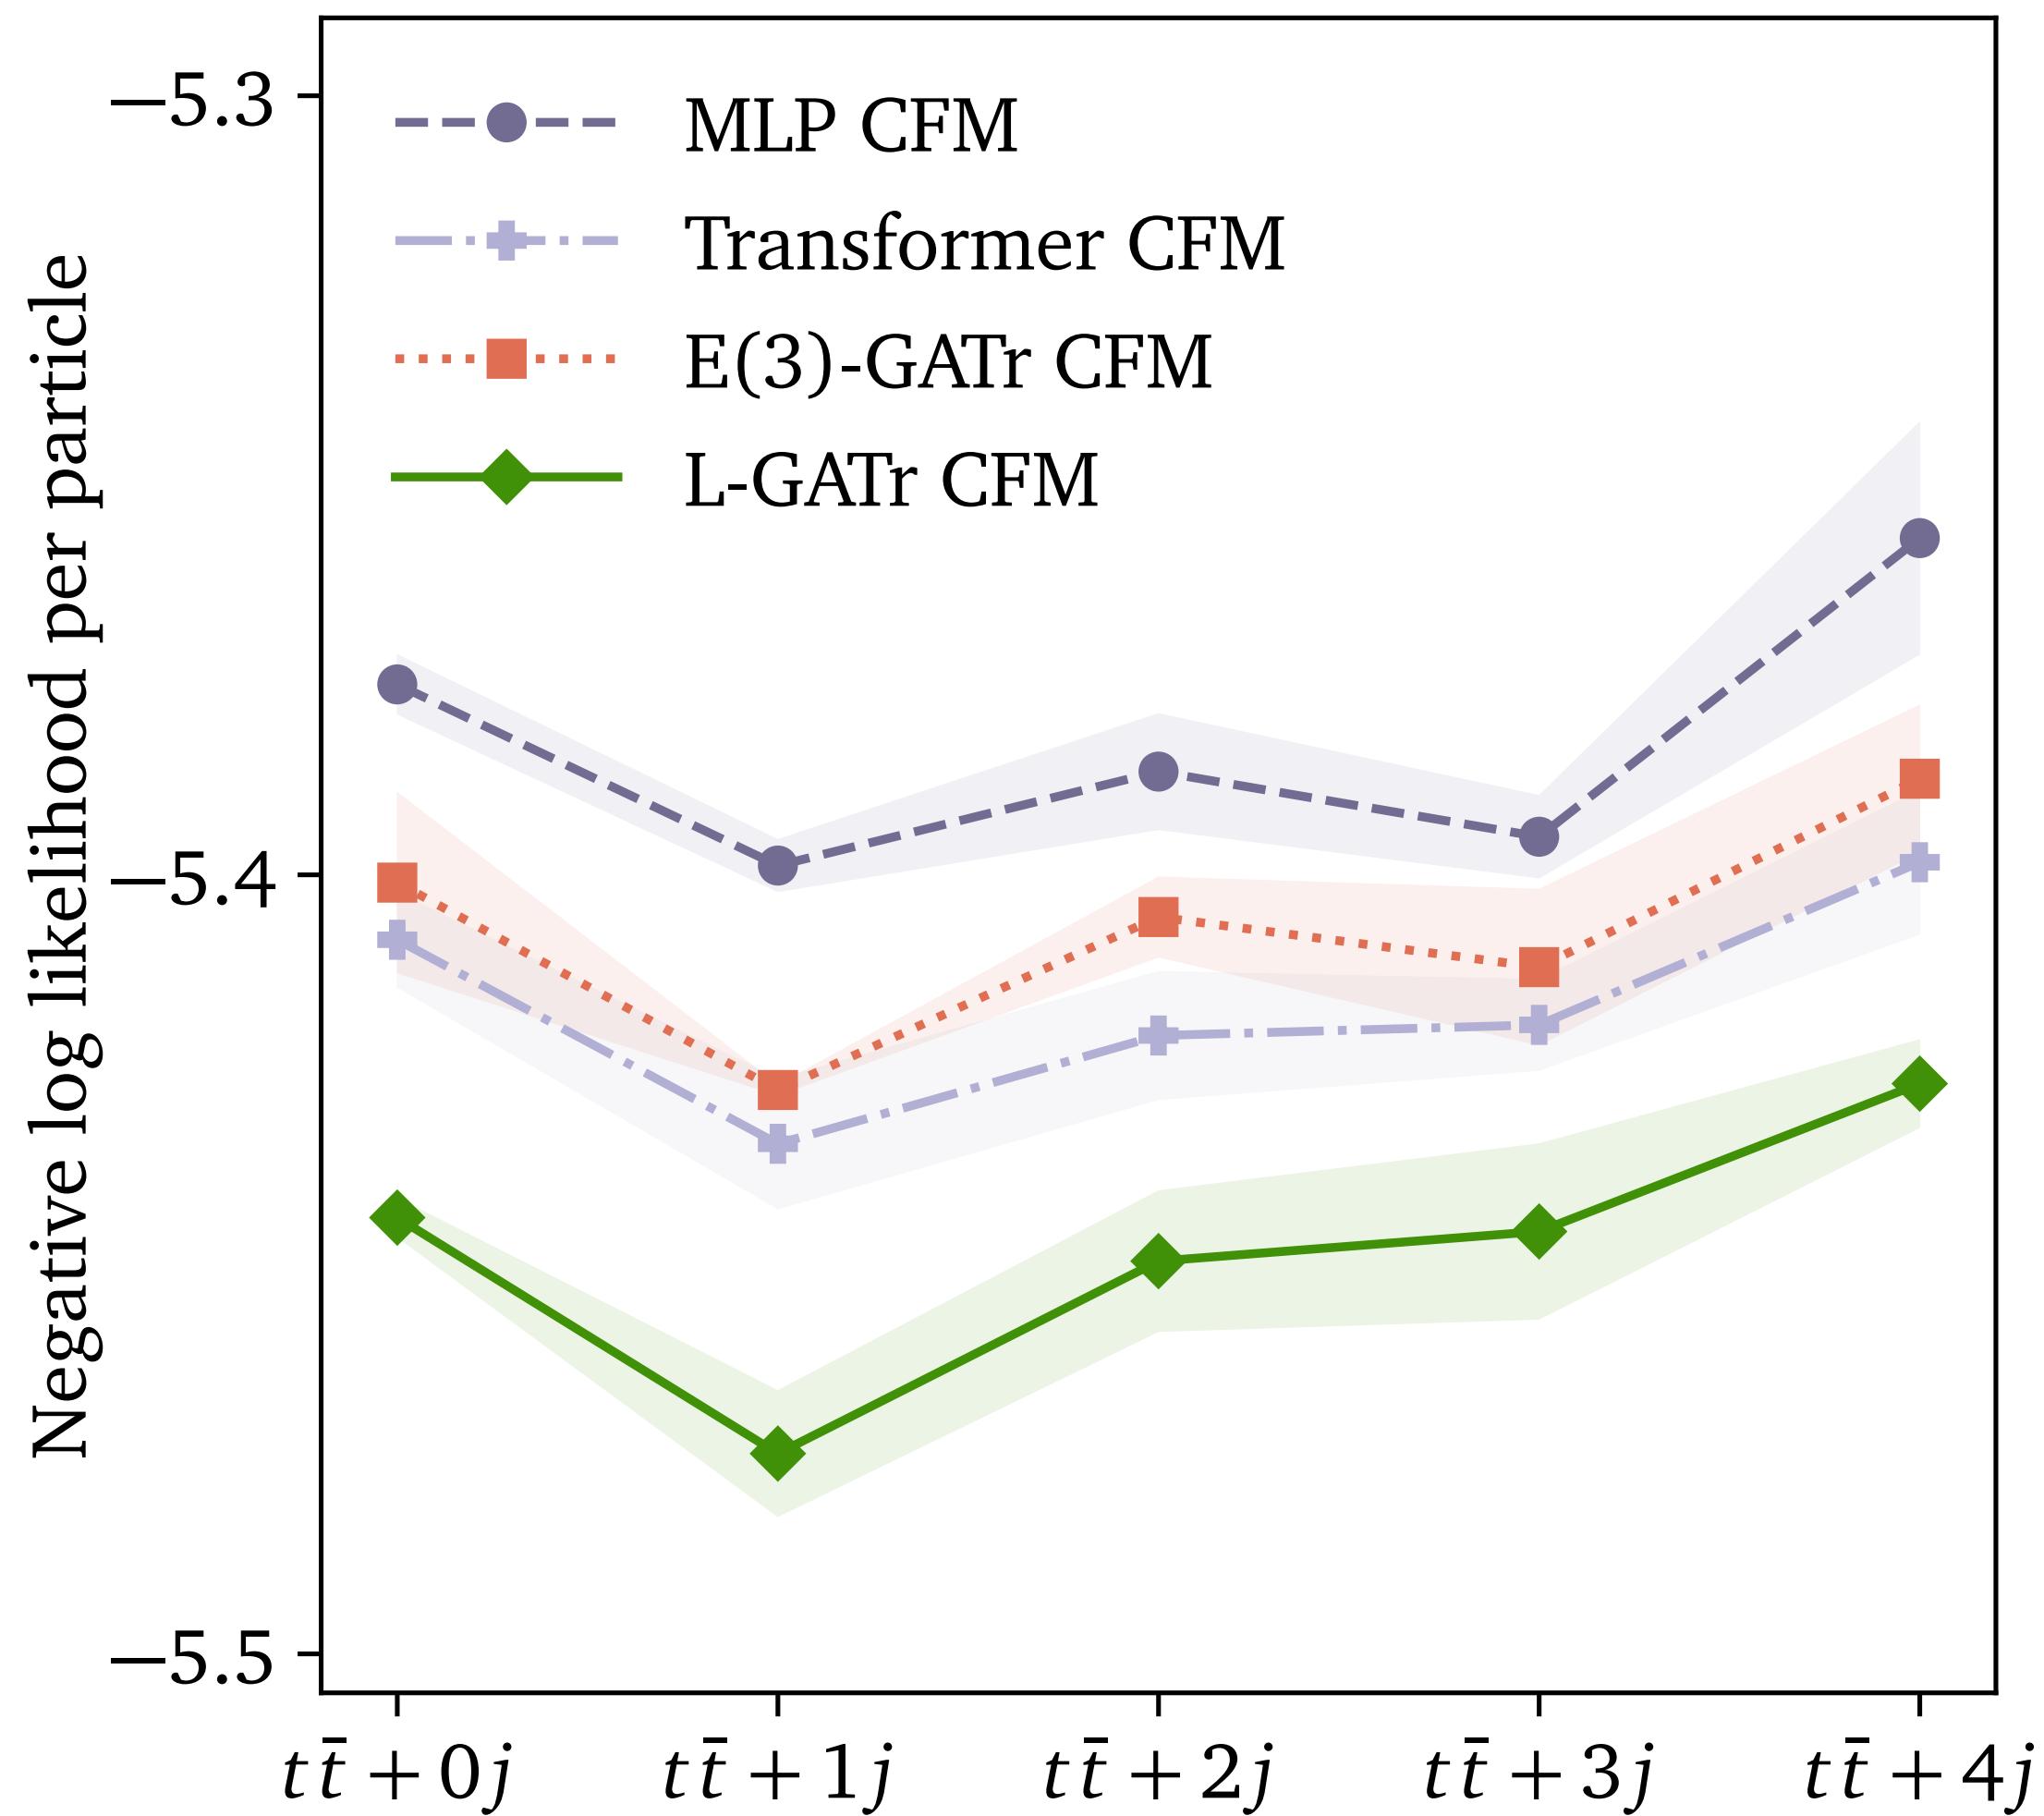
- Transformations  $f(x)$  between Minkowski space  $p$  and the parametrization  $x$
- Equivariant L-GATr operations using multivectors
- Symmetry-breaking operations using scalars (required for numerical stability)



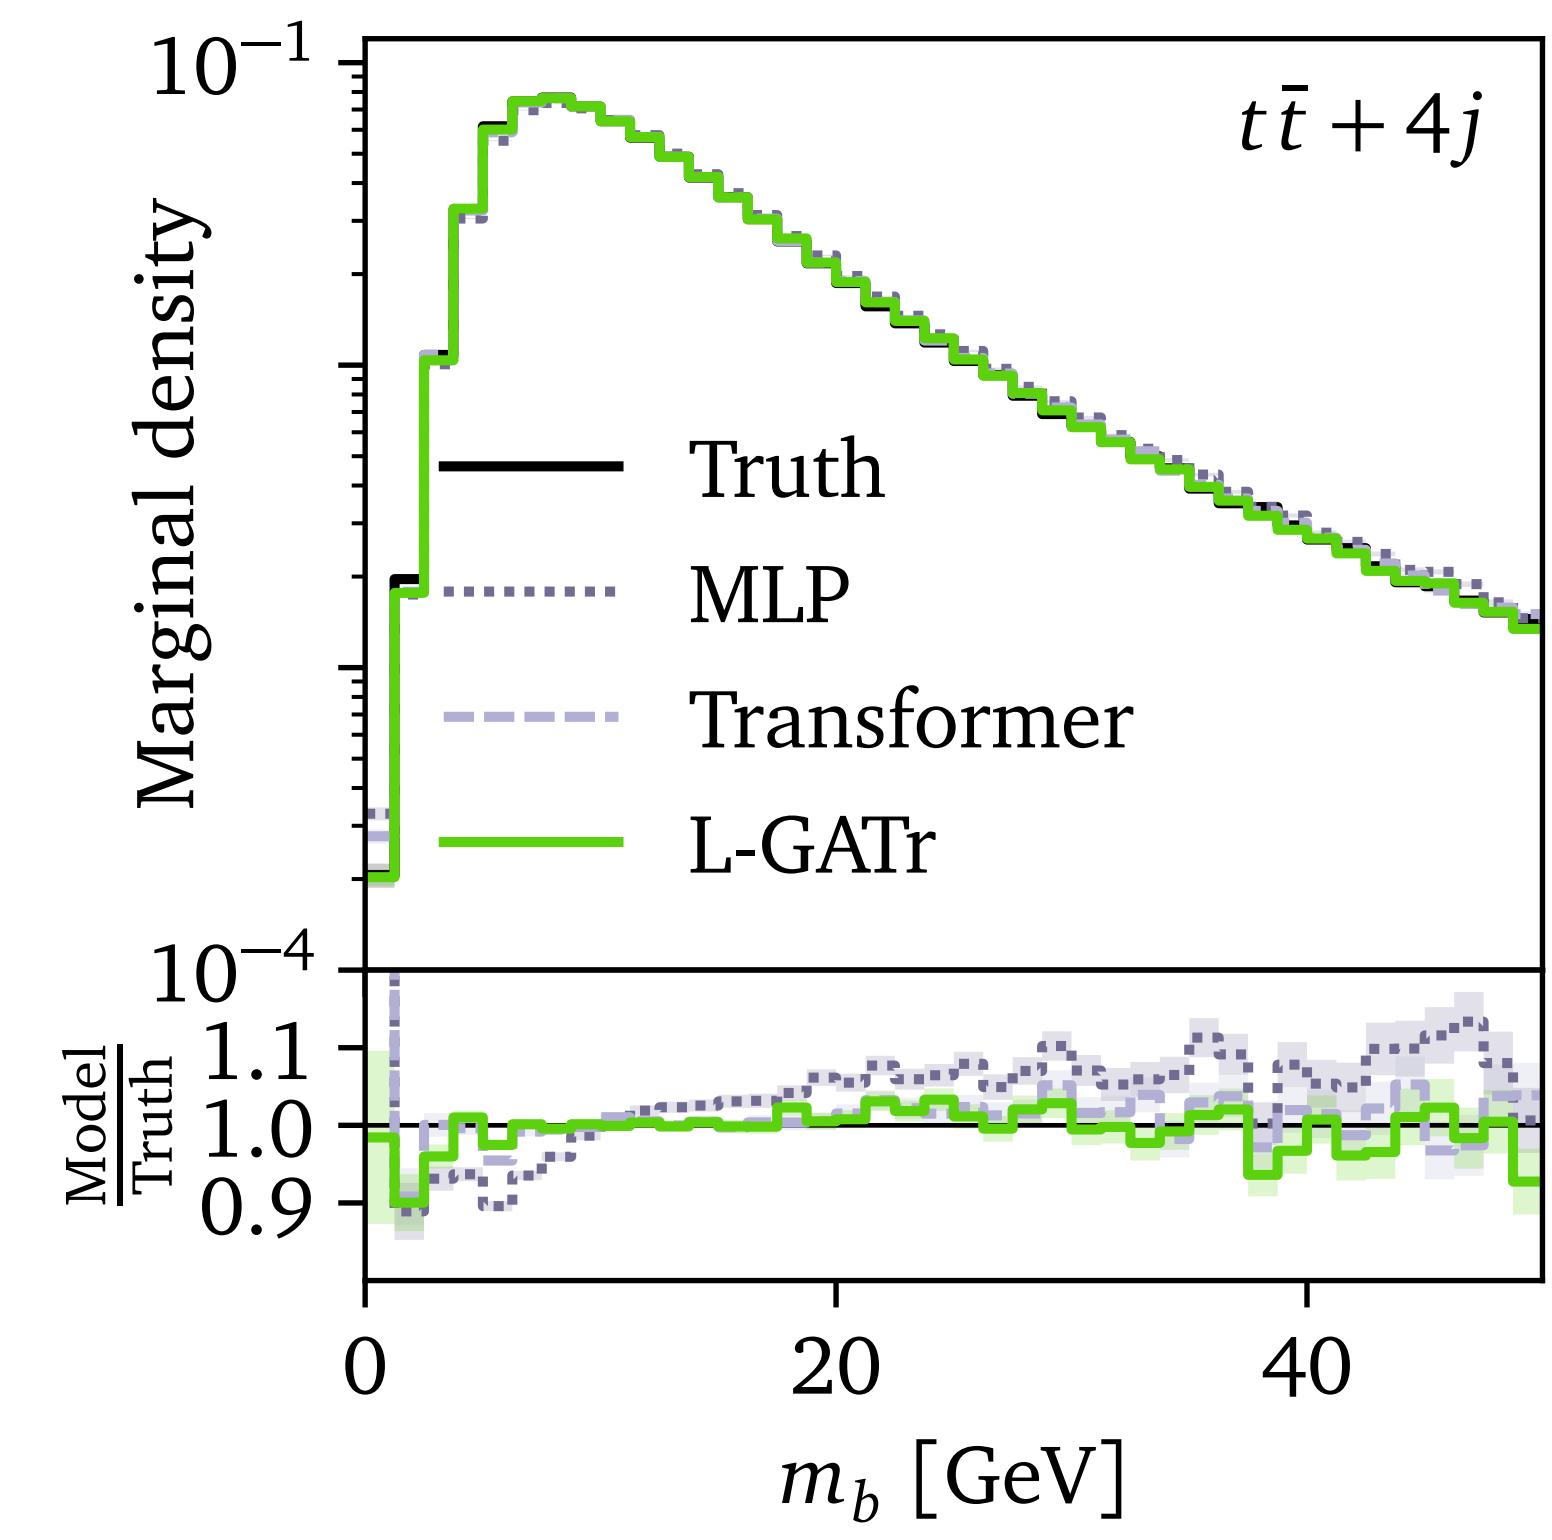
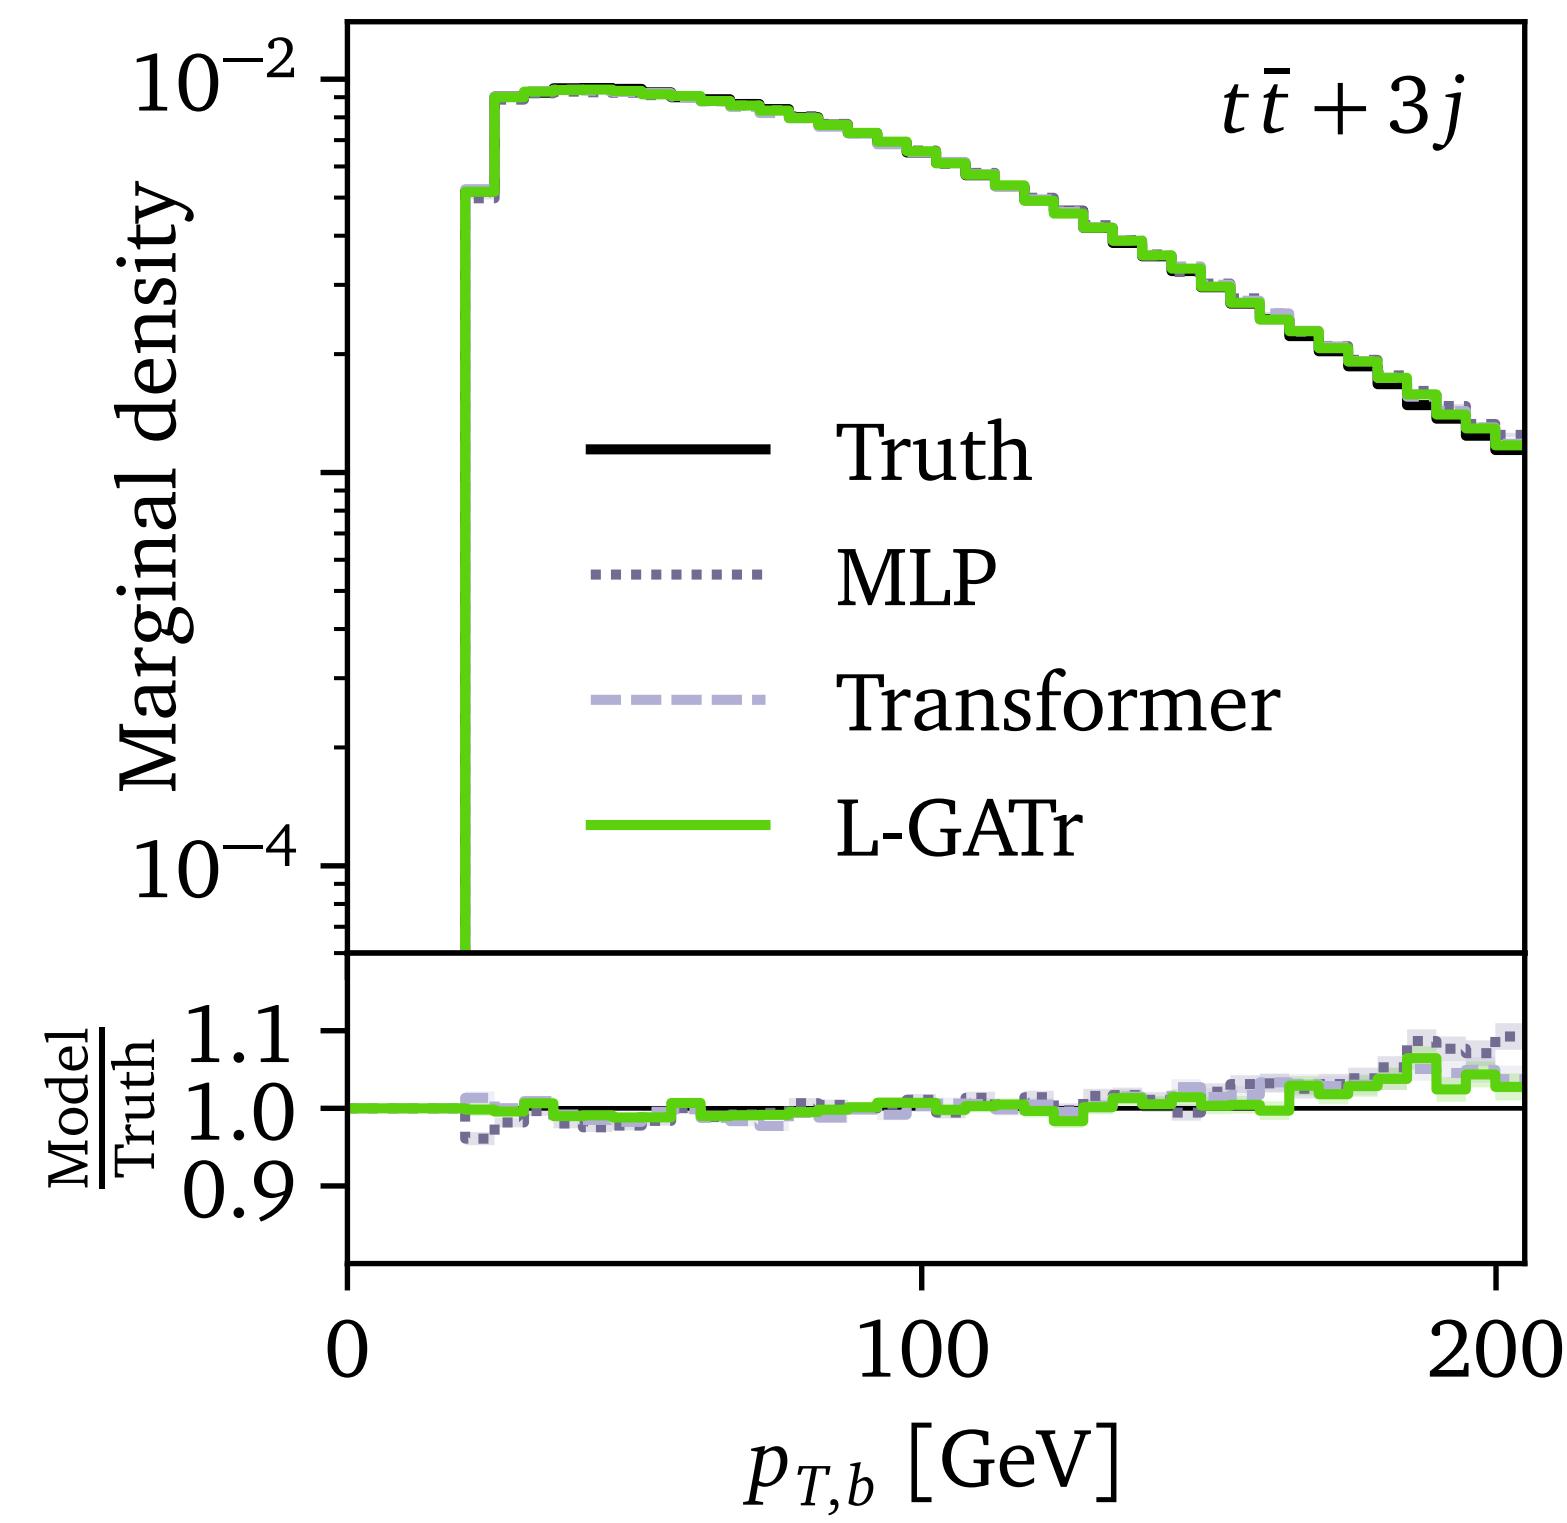
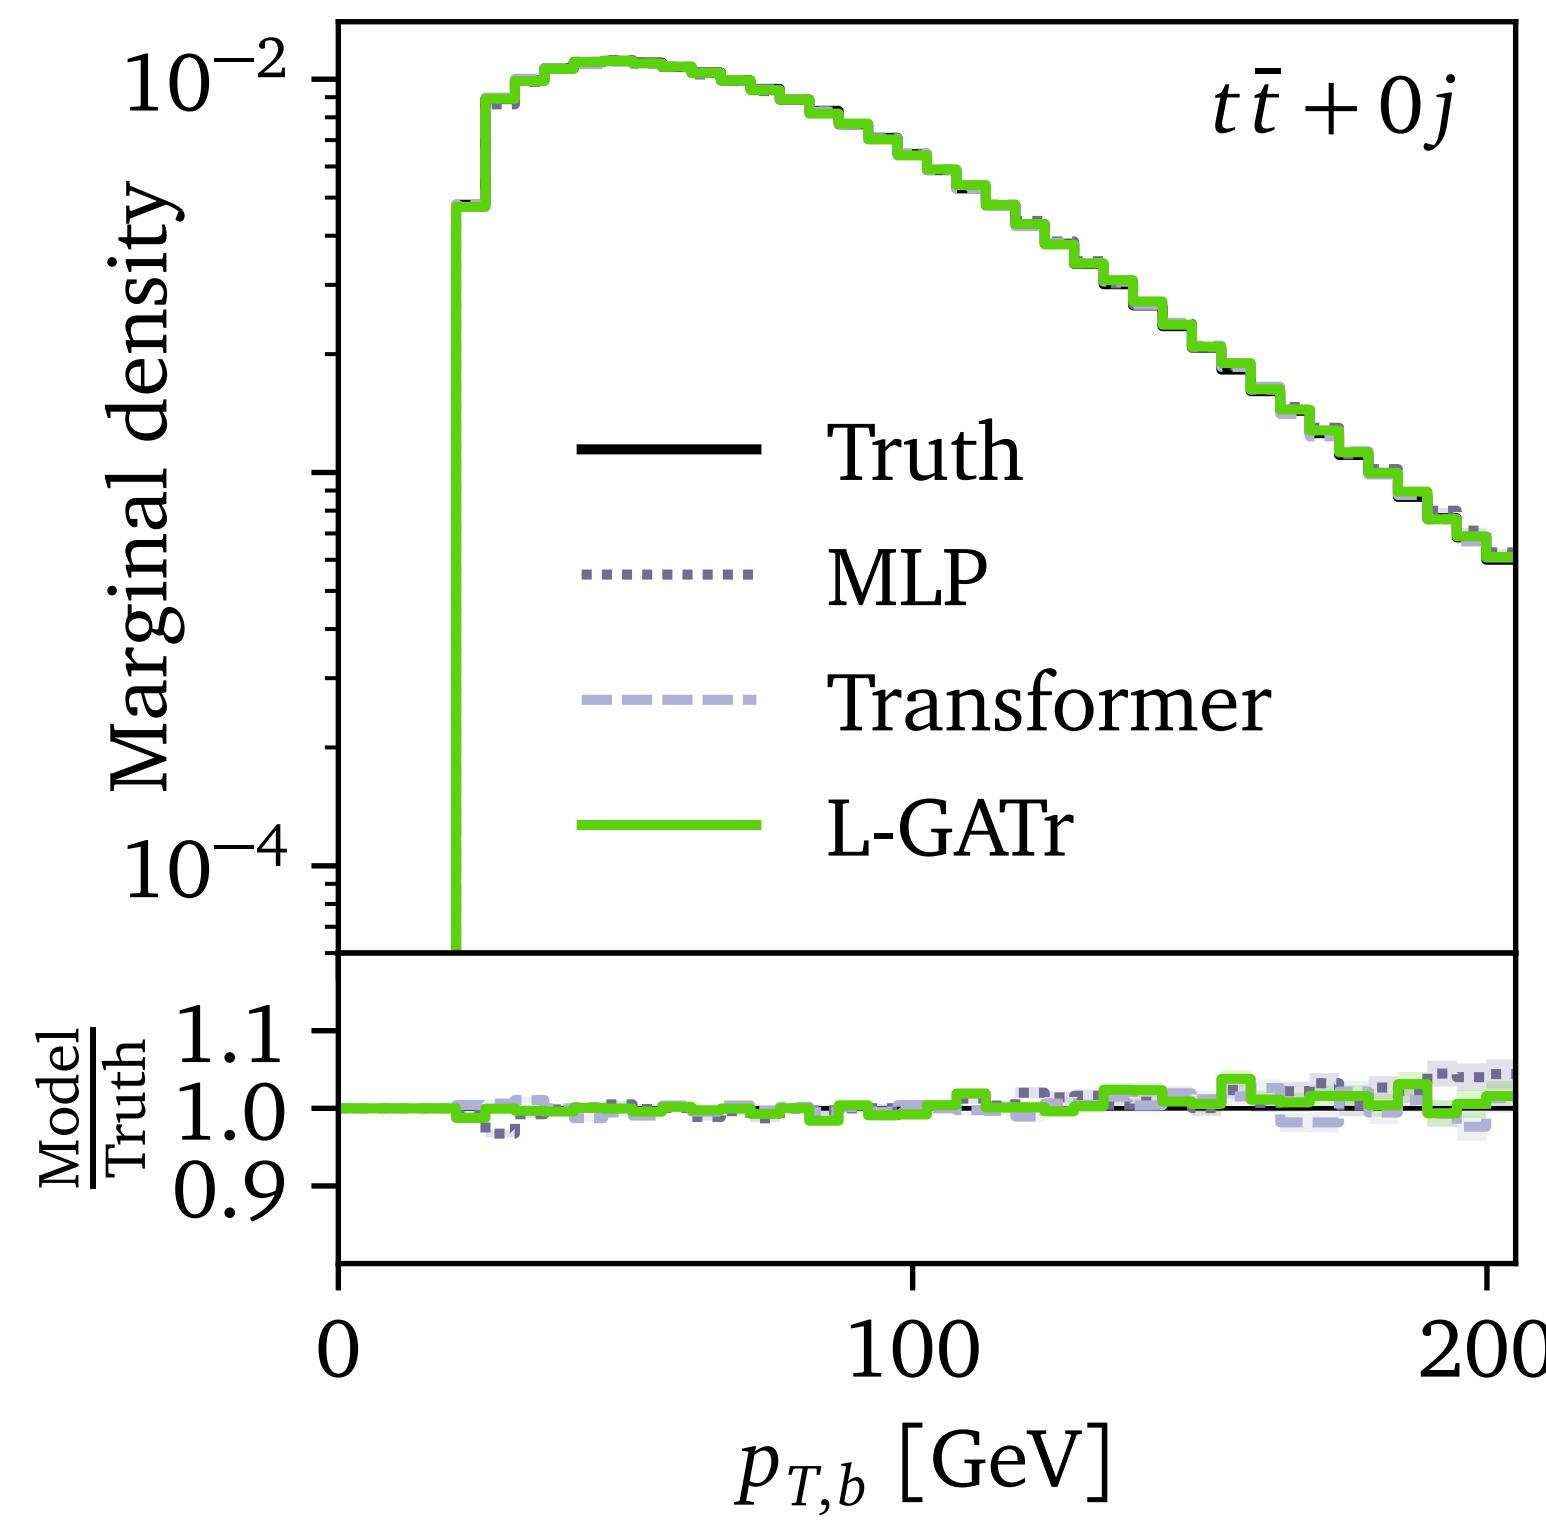
# Event generation



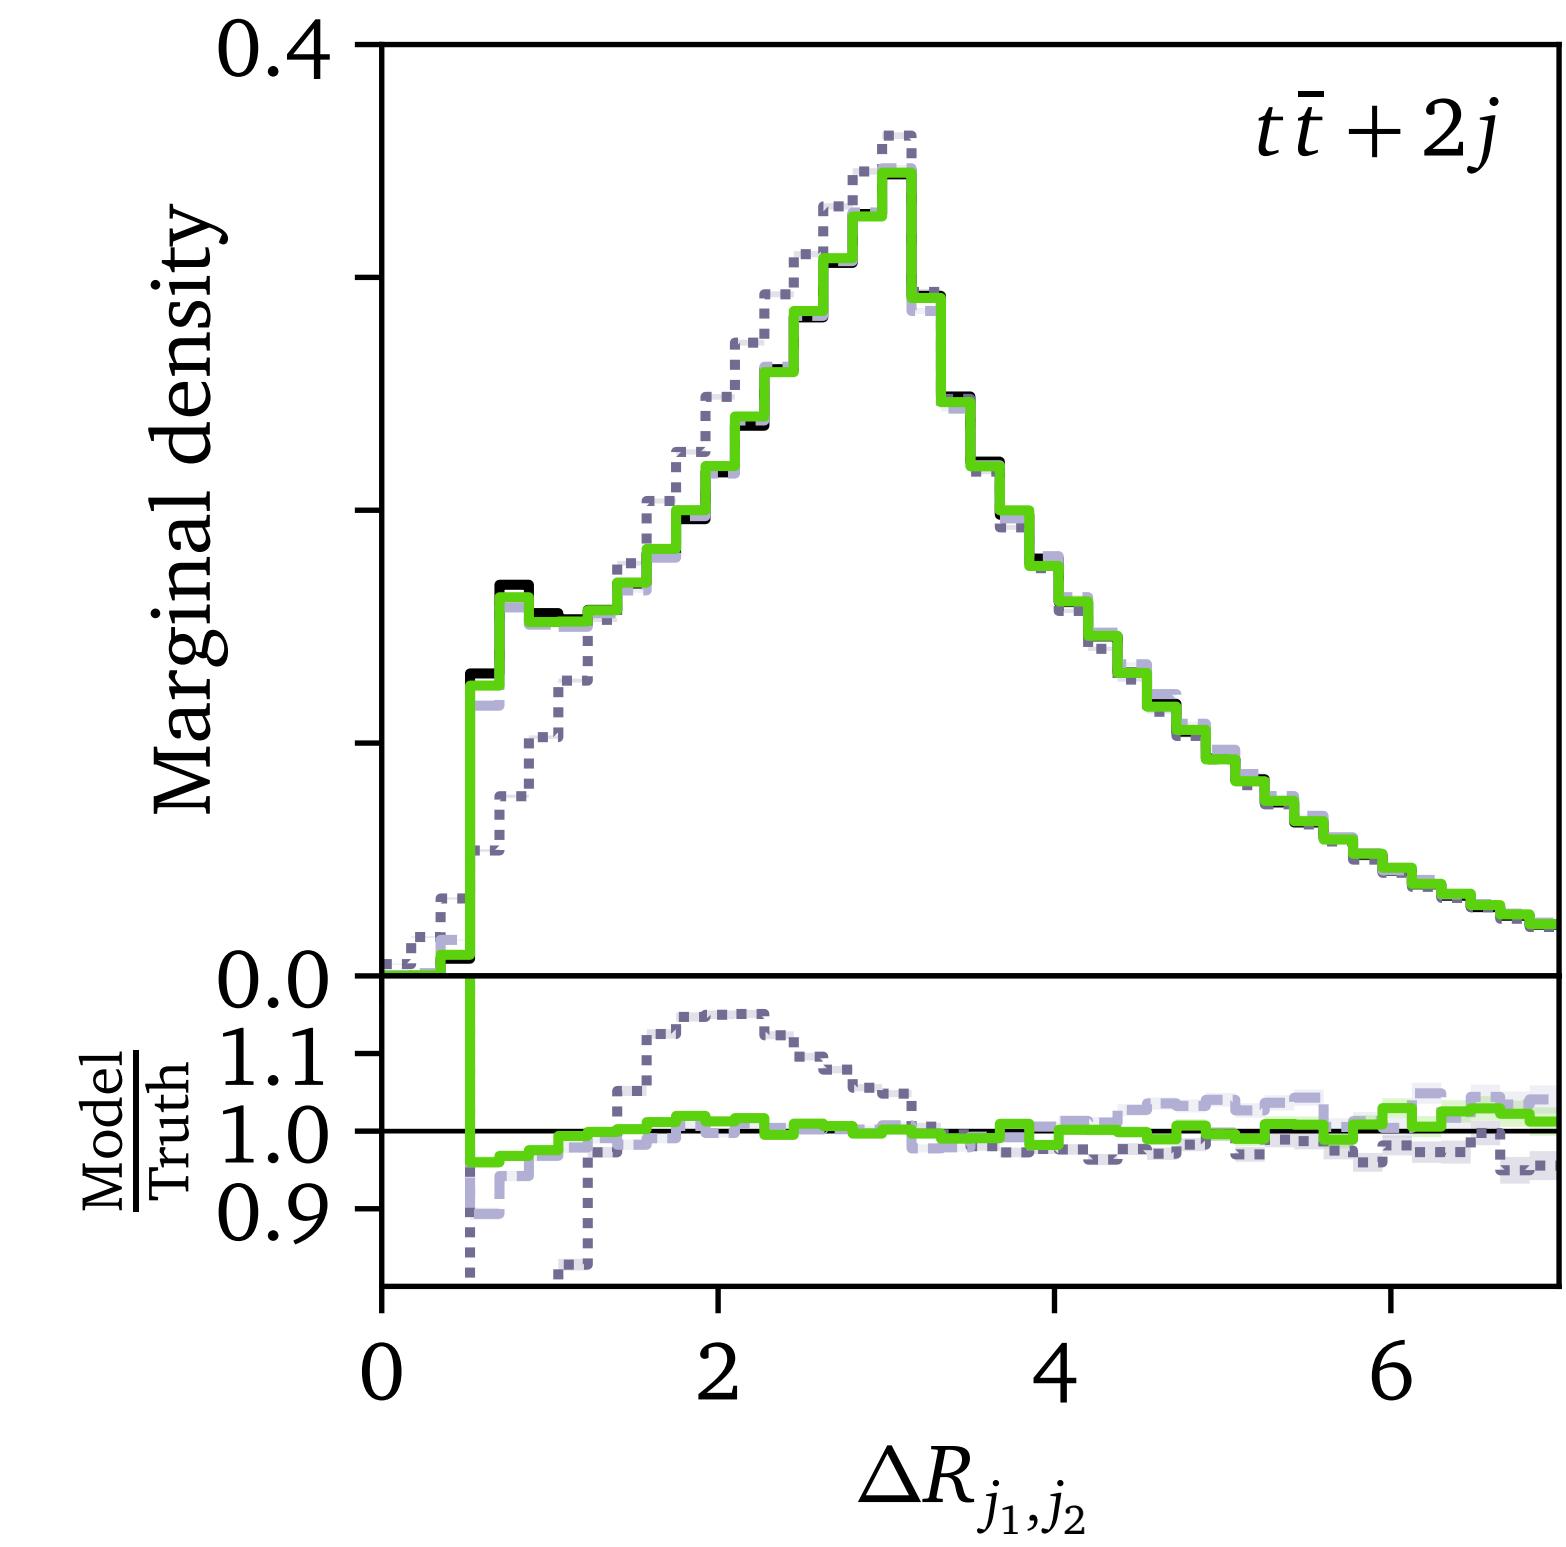
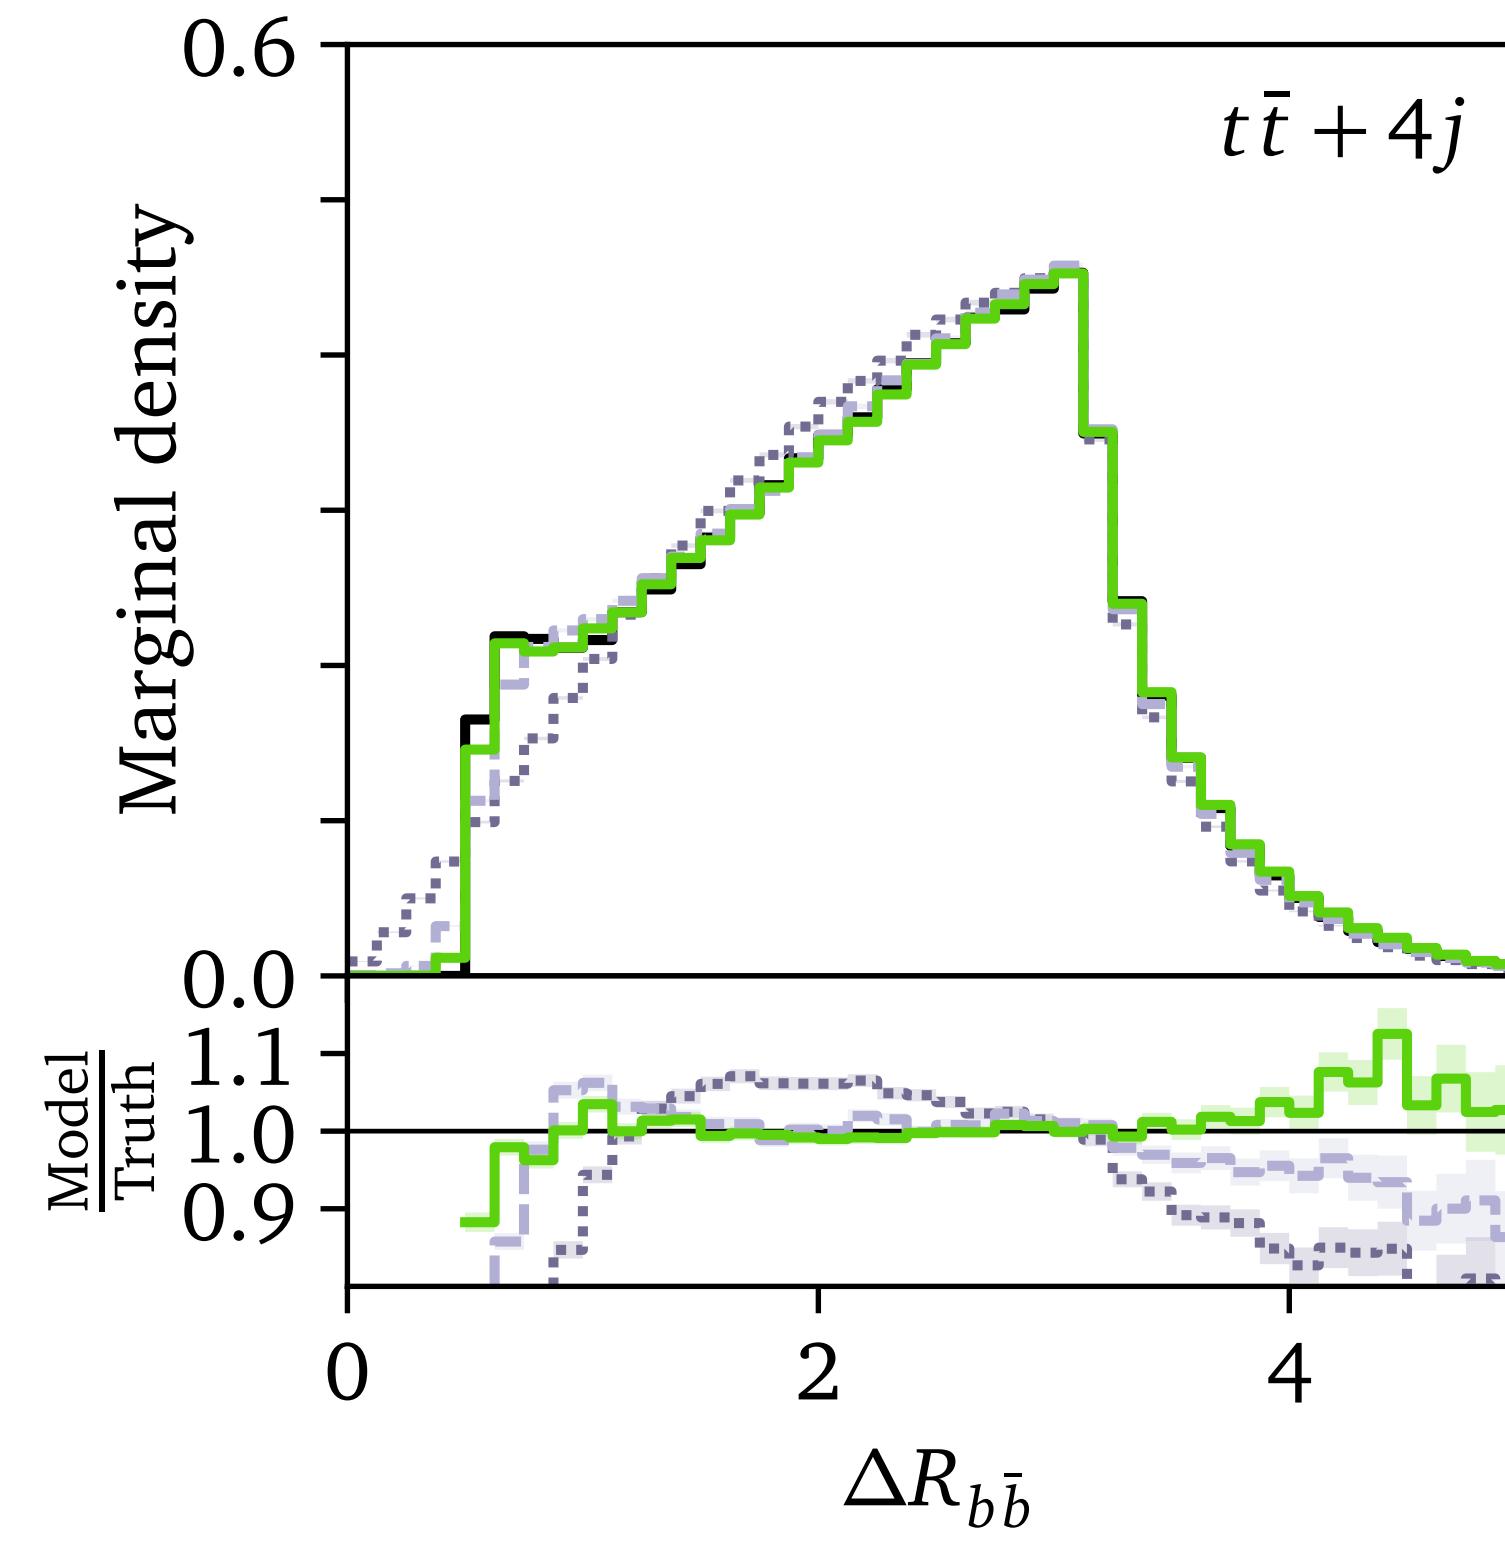
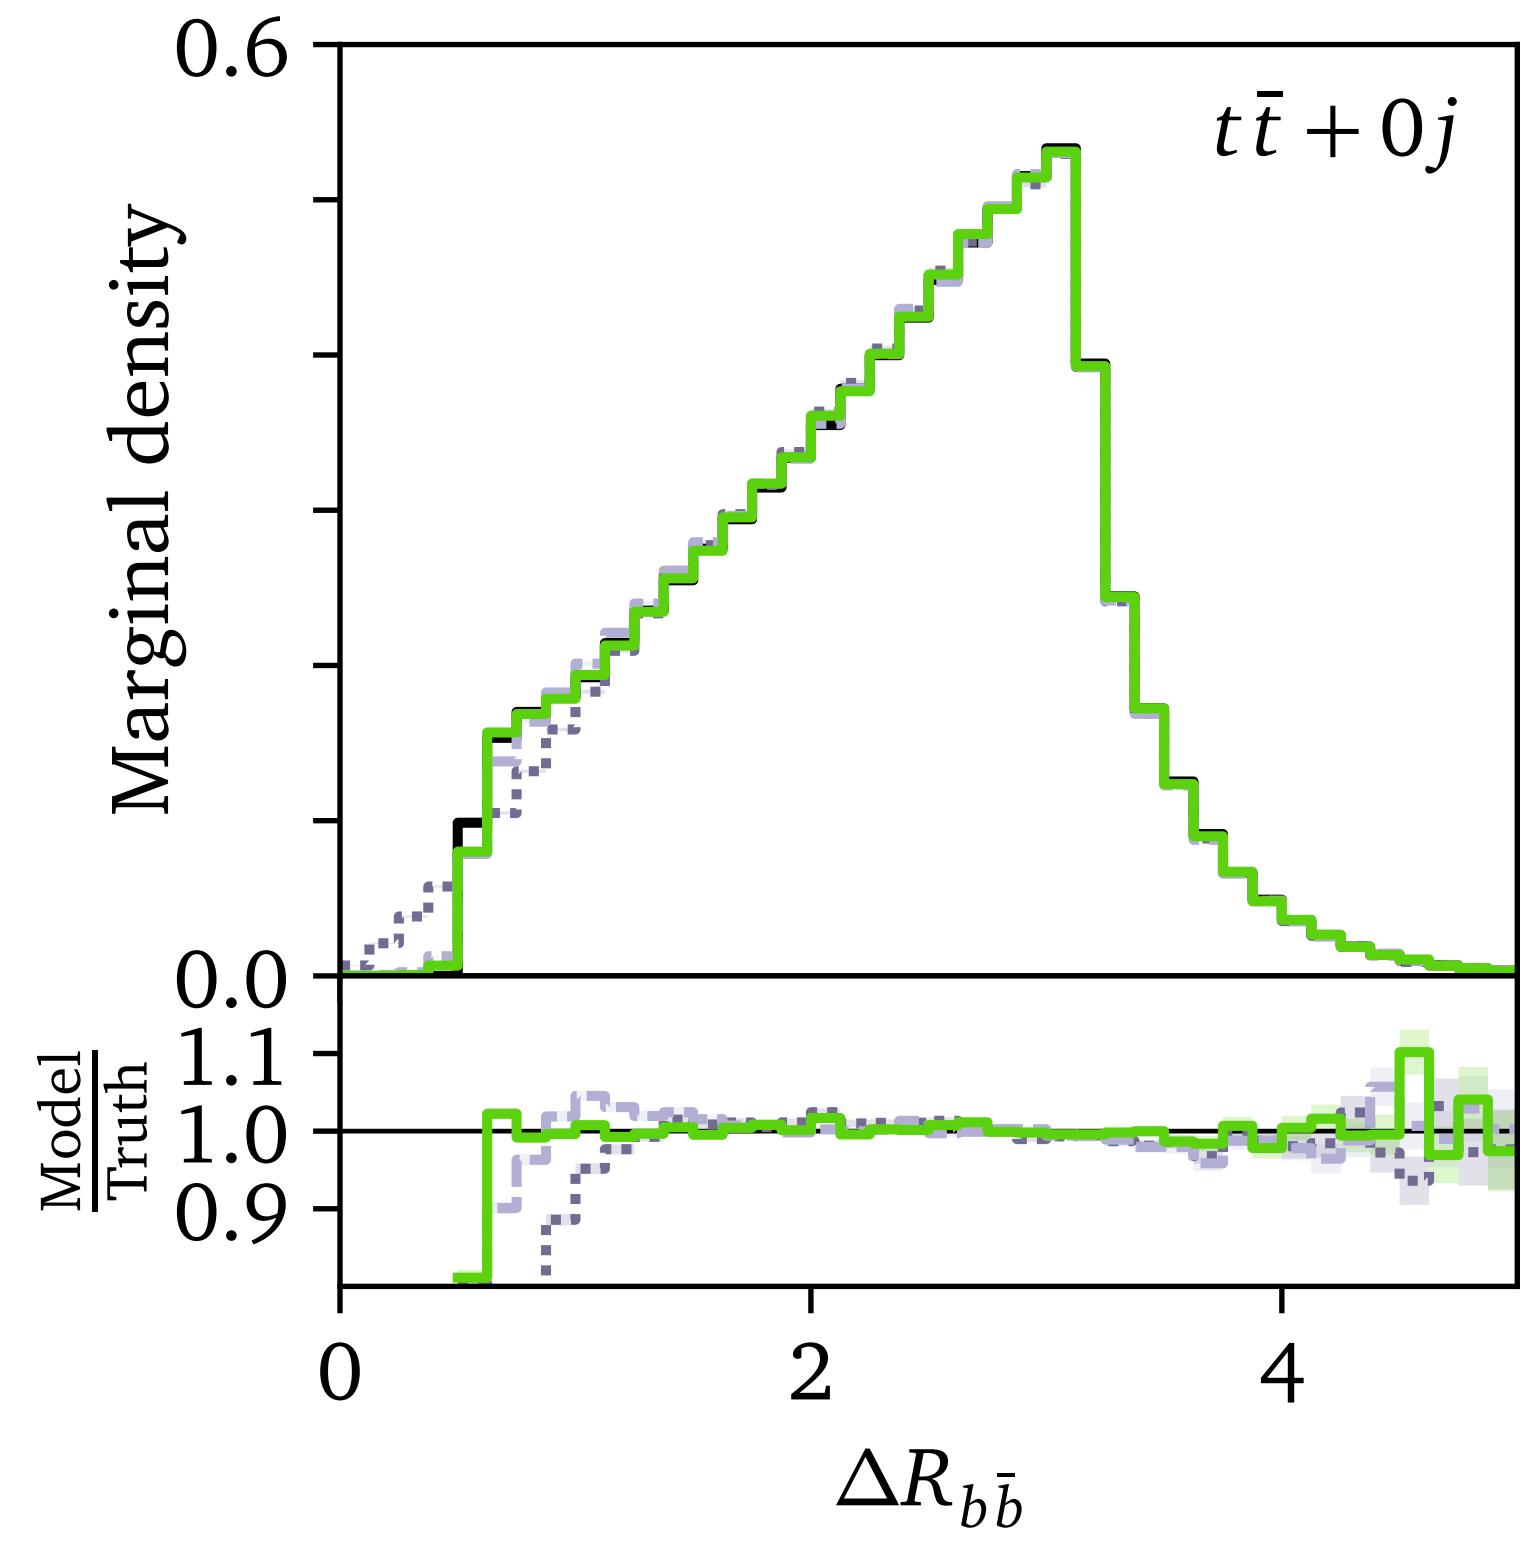
# Event generation



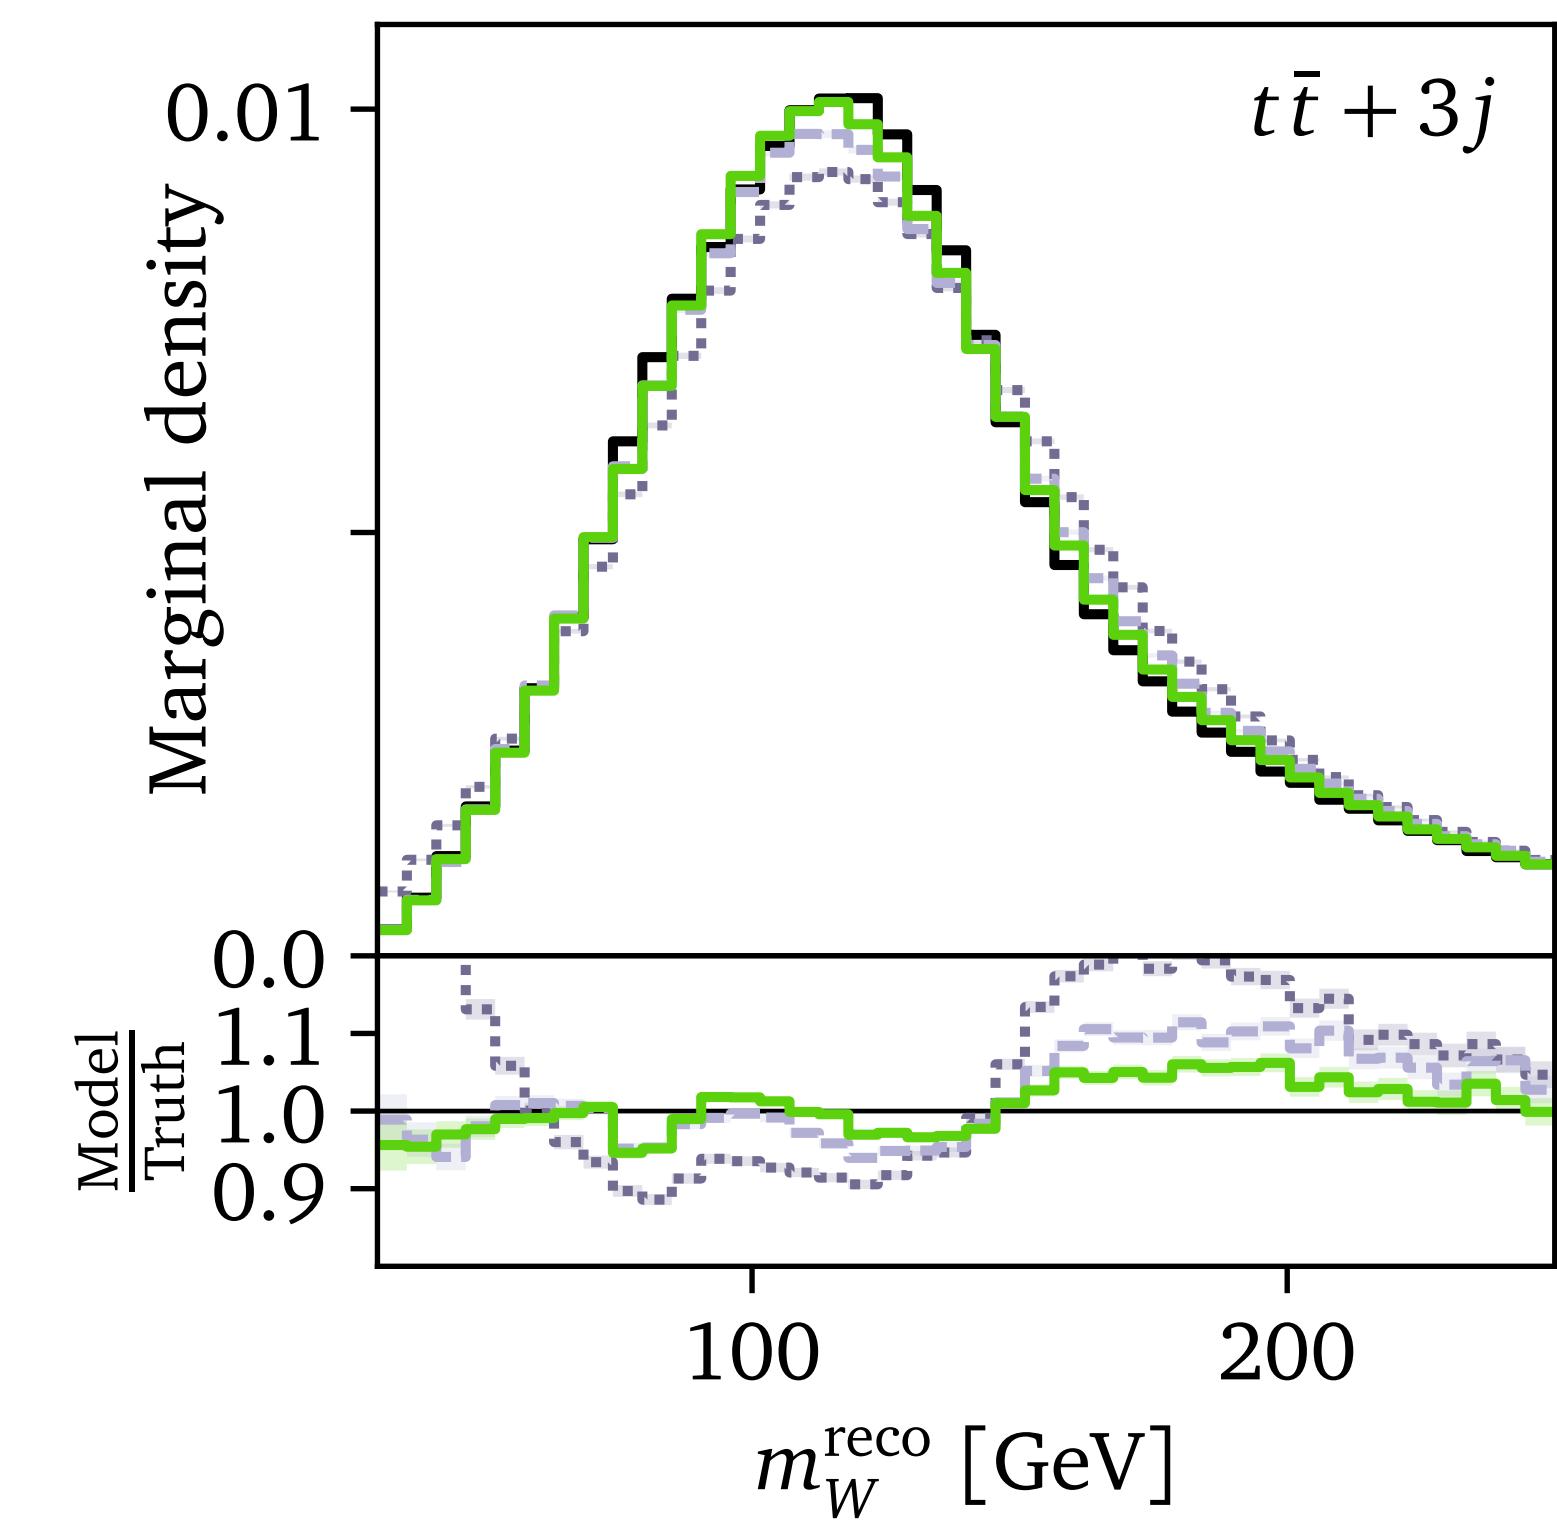
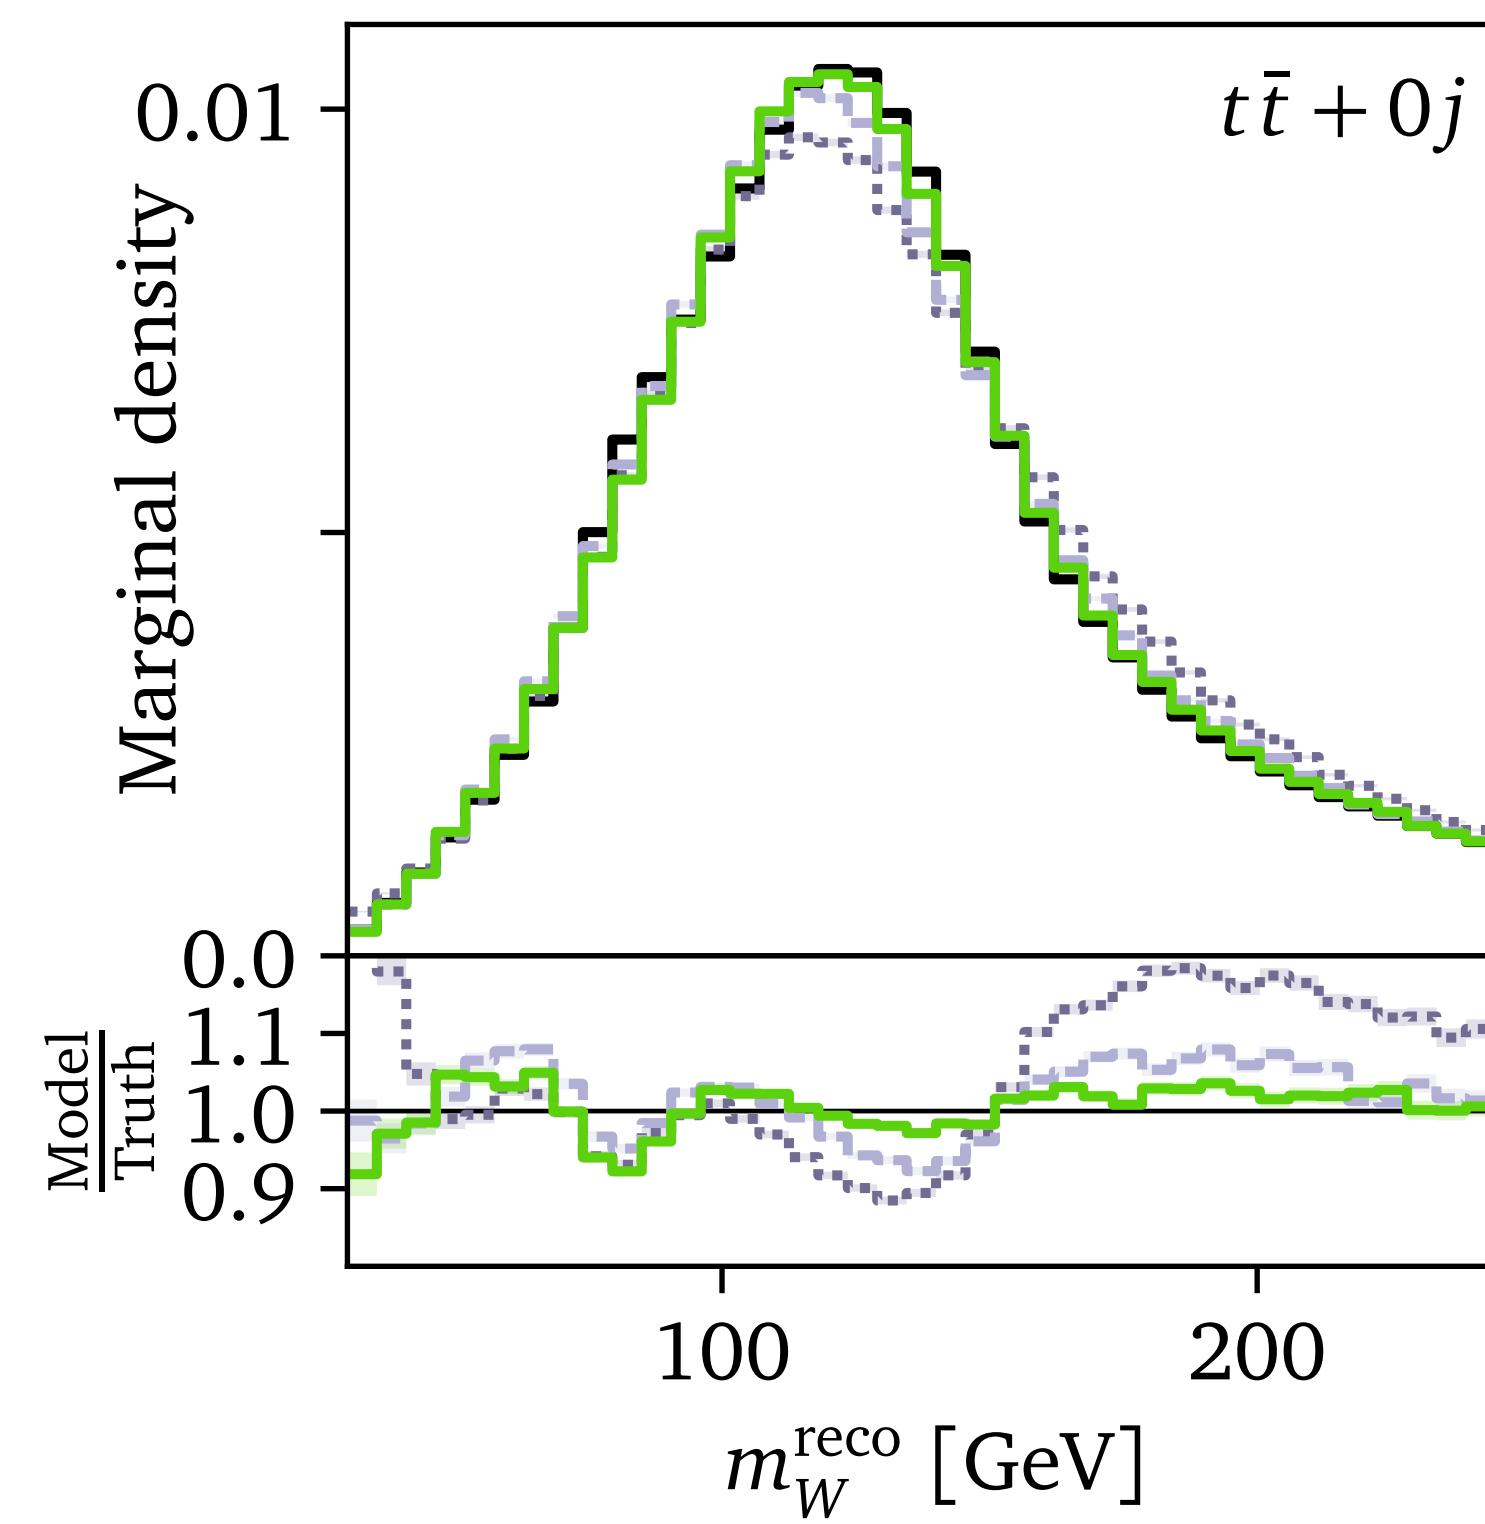
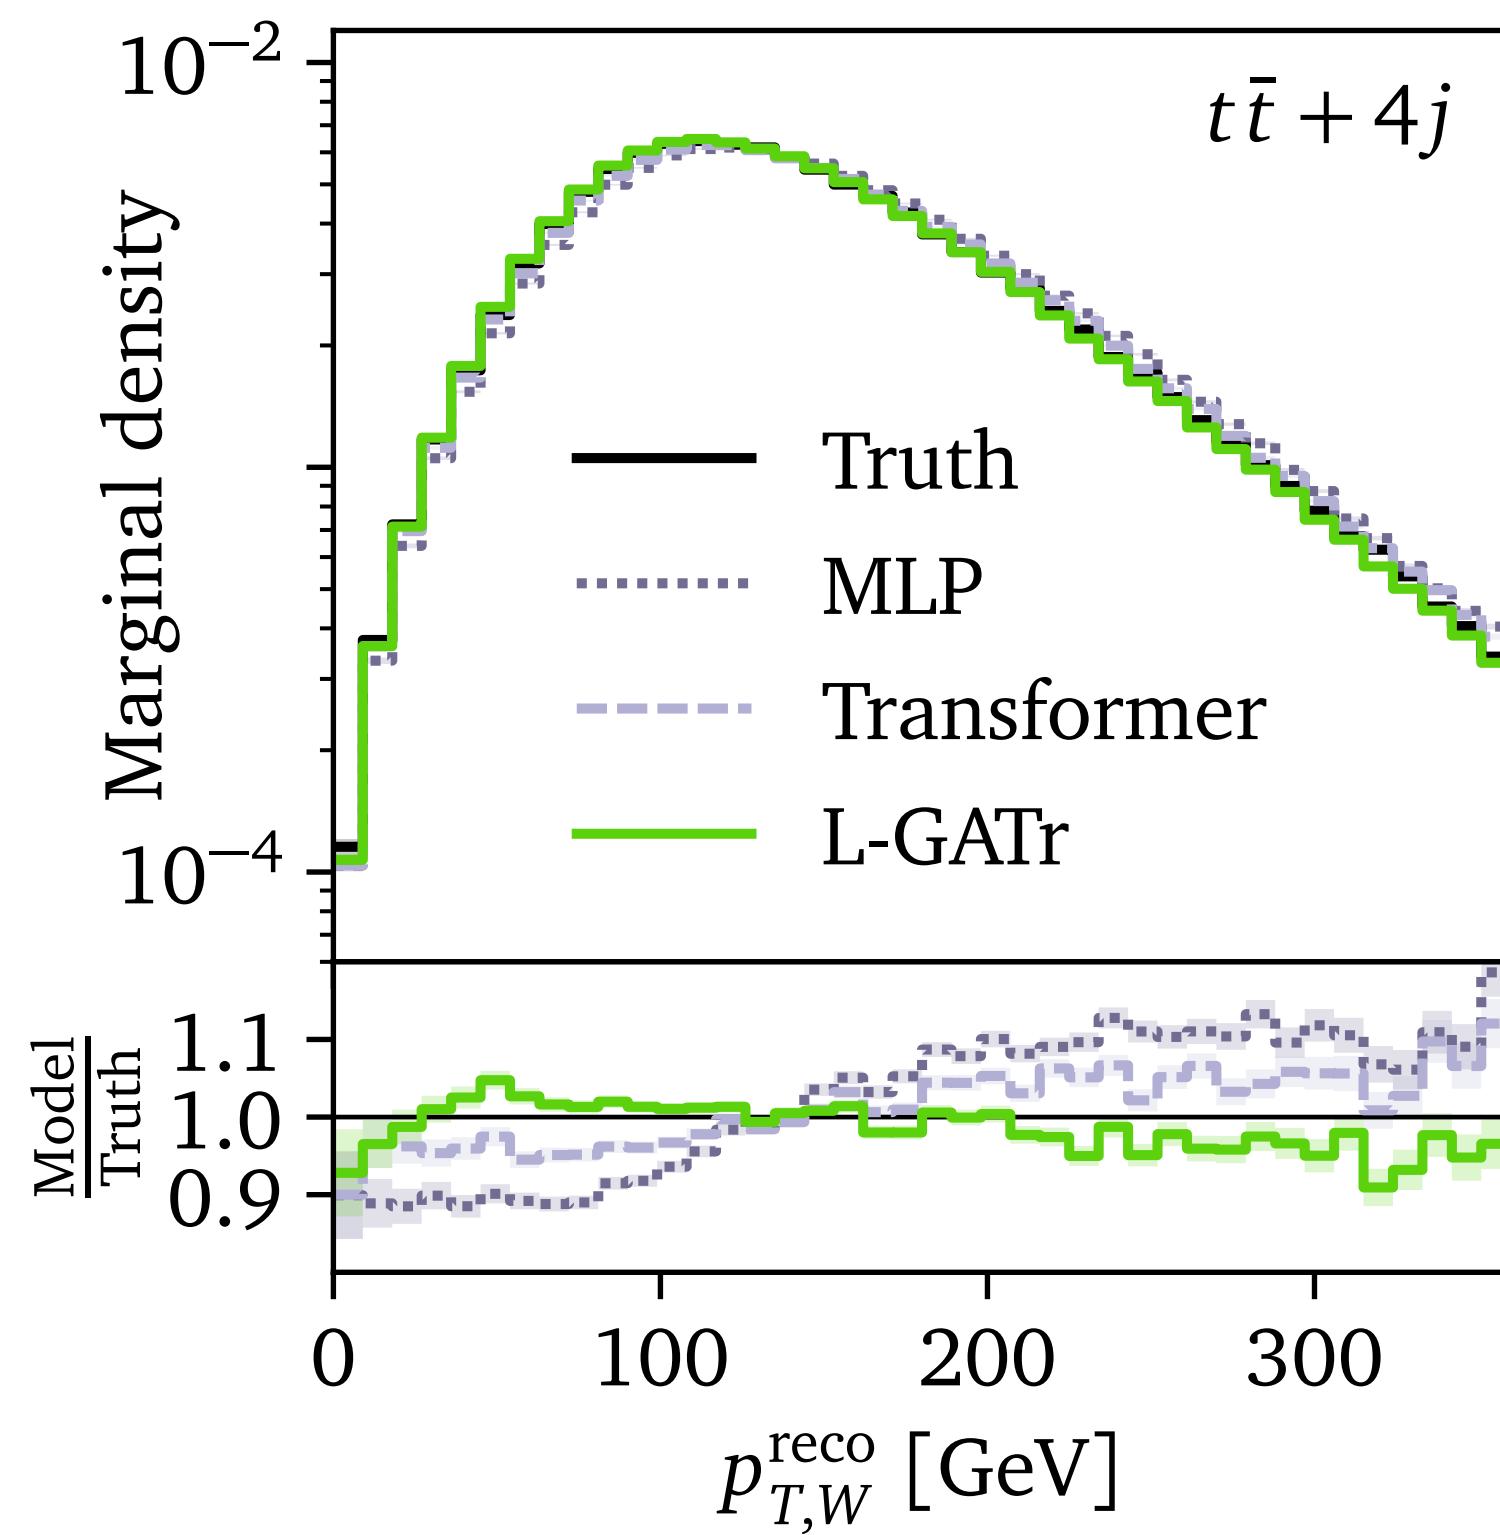
# Event generation



# Event generation



# Event generation



# Event generation

